

Freescale SemiconductorMPC603EUM/AD
Q2/02
REV 3

MPC603e RISC Microprocessor User's Manual



Home Page:

www.freescale.com

email:

support@freescale.com

USA/Europe or Locations Not Listed:

Freescale Semiconductor

Technical Information Center, CH370

1300 N. Alma School Road

Chandler, Arizona 85224

(800) 521-6274

480-768-2130

support@freescale.com

Europe, Middle East, and Africa:

Freescale Halbleiter Deutschland GmbH

Technical Information Center

Schatzbogen 7

81829 Muenchen, Germany

+44 1296 380 456 (English)

+46 8 52200080 (English)

+49 89 92103 559 (German)

+33 1 69 35 48 48 (French)

support@freescale.com

Japan:

Freescale Semiconductor Japan Ltd.

Headquarters

ARCO Tower 15F

1-8-1, Shimo-Meguro, Meguro-ku

Tokyo 153-0064, Japan

0120 191014

+81 2666 8080

support.japan@freescale.com

Asia/Pacific:

Freescale Semiconductor Hong Kong Ltd.

Technical Information Center

2 Dai King Street

Tai Po Industrial Estate,

Tai Po, N.T., Hong Kong

+800 2666 8080

support.asia@freescale.com

For Literature Requests Only:

Freescale Semiconductor

Literature Distribution Center

P.O. Box 5405

Denver, Colorado 80217

(800) 441-2447

303-675-2140

Fax: 303-675-2150

LDCForFreescaleSemiconductor

@hibbertgroup.com

Information in this document is provided solely to enable system and software implementers to use Freescale Semiconductor products. There are no express or implied copyright licenses granted hereunder to design or fabricate any integrated circuits or integrated circuits based on the information in this document.

Freescale Semiconductor reserves the right to make changes without further notice to any products herein. Freescale Semiconductor makes no warranty, representation or guarantee regarding the suitability of its products for any particular purpose, nor does Freescale Semiconductor assume any liability arising out of the application or use of any product or circuit, and specifically disclaims any and all liability, including without limitation consequential or incidental damages. "Typical" parameters which may be provided in Freescale Semiconductor data sheets and/or specifications can and do vary in different applications and actual performance may vary over time. All operating parameters, including "Typicals" must be validated for each customer application by customer's technical experts. Freescale Semiconductor does not convey any license under its patent rights nor the rights of others. Freescale Semiconductor products are not designed, intended, or authorized for use as components in systems intended for surgical implant into the body, or other applications intended to support or sustain life, or for any other application in which the failure of the Freescale Semiconductor product could create a situation where personal injury or death may occur. Should Buyer purchase or use Freescale Semiconductor products for any such unintended or unauthorized application, Buyer shall indemnify and hold Freescale Semiconductor and its officers, employees, subsidiaries, affiliates, and distributors harmless against all claims, costs, damages, and expenses, and reasonable attorney fees arising out of, directly or indirectly, any claim of personal injury or death associated with such unintended or unauthorized use, even if such claim alleges that Freescale Semiconductor was negligent regarding the design or manufacture of the part.



Overview	1
Programming Model	2
Instruction and Data Cache Operation	3
Exceptions	4
Memory Management	5
Instruction Timing	6
Signal Descriptions	7
System Interface Operation	8
Power Management	9
PowerPC Instruction Set Listings	A
Instructions Not Implemented	B
Revision History	C
Glossary of Terms and Abbreviations	GLO
Index	IND

- 1** Overview
- 2** Programming Model
- 3** Instruction and Data Cache Operation
- 4** Exceptions
- 5** Memory Management
- 6** Instruction Timing
- 7** Signal Descriptions
- 8** System Interface Operation
- 9** Power Management
- A** PowerPC Instruction Set Listings
- B** Instructions Not Implemented
- C** Revision History
- GLO** Glossary of Terms and Abbreviations
- IND** Index

Contents

Paragraph Number	Title	Page Number
	About This Book	
	Chapter 1 Overview	
1.1	Overview.....	1-1
1.1.1	Features	1-3
1.1.2	System Design and Programming Considerations.....	1-6
1.1.2.1	Hardware Features	1-6
1.1.2.1.1	Replacement of XATS Signal by CSE1 Signal	1-6
1.1.2.1.2	Addition of Half-Clock Bus Multipliers.....	1-7
1.1.2.2	Software Features	1-7
1.1.2.2.1	16-Kbyte Instruction and Data Caches	1-7
1.1.2.2.2	Clock Configuration Available in HID1 Register	1-7
1.1.2.2.3	Performance Enhancements	1-7
1.1.3	Instruction Unit	1-8
1.1.3.1	Instruction Queue and Dispatch Unit	1-8
1.1.3.2	Branch Processing Unit (BPU).....	1-9
1.1.4	Independent Execution Units.....	1-9
1.1.4.1	Integer Unit (IU)	1-9
1.1.4.2	Floating-Point Unit (FPU)	1-10
1.1.4.3	Load/Store Unit (LSU)	1-10
1.1.4.4	System Register Unit (SRU).....	1-10
1.1.4.5	Completion Unit	1-11
1.1.5	Memory Subsystem Support.....	1-11
1.1.5.1	Memory Management Units (MMUs).....	1-11
1.1.5.2	Cache Units.....	1-12
1.1.6	Processor Bus Interface	1-13
1.1.7	System Support Functions	1-14
1.1.7.1	Power Management	1-14
1.1.7.2	Time Base/Decrementer	1-15
1.1.7.3	IEEE 1149.1 (JTAG)/COP Test Interface	1-15
1.1.7.4	Clock Multiplier.....	1-15
1.2	PowerPC Architecture Implementation	1-15

Contents

Paragraph Number	Title	Page Number
1.3	Implementation-Specific Information.....	1-16
1.3.1	Programming Model	1-17
1.3.1.1	Processor Version Register (PVR).....	1-17
1.3.1.2	Hardware Implementation Register 0 (HID0)	1-17
1.3.1.3	General-Purpose Registers (GPRs).....	1-19
1.3.1.4	Floating-Point Registers (FPRs)	1-19
1.3.1.5	Condition Register (CR)	1-19
1.3.1.6	Floating-Point Status and Control Register (FPSCR)	1-19
1.3.1.7	Machine State Register (MSR)	1-20
1.3.1.8	Segment Registers (SRs)	1-20
1.3.1.9	Special-Purpose Registers (SPRs)	1-20
1.3.1.9.1	User-Level SPRs.....	1-20
1.3.1.9.2	Supervisor-Level SPRs.....	1-21
1.3.2	Instruction Set and Addressing Modes	1-22
1.3.2.1	PowerPC Instruction Set and Addressing Modes	1-22
1.3.2.2	Implementation-Specific Instruction Set	1-24
1.3.3	Cache Implementation	1-24
1.3.3.1	PowerPC Cache Characteristics	1-24
1.3.3.2	Implementation-Specific Cache Implementation	1-24
1.3.4	Exception Model.....	1-26
1.3.4.1	PowerPC Exception Model.....	1-26
1.3.4.2	Implementation-Specific Exception Model.....	1-28
1.3.5	Memory Management.....	1-30
1.3.5.1	PowerPC Memory Management.....	1-30
1.3.5.2	Implementation-Specific Memory Management	1-31
1.3.6	Instruction Timing	1-32
1.3.7	System Interface	1-33
1.3.7.1	Memory Accesses	1-34
1.3.7.2	Signals.....	1-35
1.3.7.3	Signal Configuration.....	1-36

Chapter 2 Programming Model

2.1	Register Set.....	2-1
2.1.1	PowerPC Register Set.....	2-2
2.1.2	Implementation-Specific Registers.....	2-6
2.1.2.1	Hardware Implementation Registers (HID0 and HID1).....	2-6
2.1.2.2	Data and Instruction TLB Miss Address Registers (DMISS and IMISS).....	2-10
2.1.2.3	Data and Instruction TLB Compare Registers (DCMP and ICMP)	2-10

Contents

Paragraph Number	Title	Page Number
2.1.2.4	Primary and Secondary Hash Address Registers (HASH1 and HASH2)	2-11
2.1.2.5	Required Physical Address Register (RPA).....	2-12
2.1.2.6	Instruction Address Breakpoint Register (IABR).....	2-12
2.2	Operand Conventions	2-13
2.2.1	Floating-Point Execution Models—UISA.....	2-13
2.2.2	Data Organization in Memory and Data Transfers.....	2-14
2.2.3	Alignment and Misaligned Accesses.....	2-14
2.2.4	Floating-Point Operands	2-15
2.2.5	Effect of Operand Placement on Performance	2-15
2.3	Instruction Set Summary	2-15
2.3.1	Classes of Instructions	2-17
2.3.1.1	Definition of Boundedly Undefined	2-17
2.3.1.2	Defined Instruction Class	2-17
2.3.1.3	Illegal Instruction Class	2-18
2.3.1.4	Reserved Instruction Class	2-19
2.3.2	Addressing Modes	2-19
2.3.2.1	Memory Addressing	2-19
2.3.2.2	Memory Operands	2-19
2.3.2.3	Effective Address Calculation	2-20
2.3.2.4	Synchronization	2-20
2.3.2.4.1	Context Synchronization	2-21
2.3.2.4.2	Execution Synchronization.....	2-21
2.3.2.4.3	Instruction-Related Exceptions	2-21
2.3.3	Instruction Set Overview	2-22
2.3.4	PowerPC UISA Instructions	2-22
2.3.4.1	Integer Instructions	2-22
2.3.4.1.1	Integer Arithmetic Instructions.....	2-23
2.3.4.1.2	Integer Compare Instructions	2-24
2.3.4.1.3	Integer Logical Instructions	2-24
2.3.4.1.4	Integer Rotate and Shift Instructions	2-25
2.3.4.2	Floating-Point Instructions	2-26
2.3.4.2.1	Floating-Point Arithmetic Instructions.....	2-26
2.3.4.2.2	Floating-Point Multiply-Add Instructions	2-27
2.3.4.2.3	Floating-Point Rounding and Conversion Instructions	2-28
2.3.4.2.4	Floating-Point Compare Instructions.....	2-28
2.3.4.2.5	Floating-Point Status and Control Register Instructions	2-28
2.3.4.2.6	Floating-Point Move Instructions	2-29
2.3.4.3	Load and Store Instructions	2-29
2.3.4.3.1	Self-Modifying Code	2-30
2.3.4.3.2	Integer Load and Store Address Generation.....	2-30

Contents

Paragraph Number	Title	Page Number
2.3.4.3.3	Register Indirect Integer Load Instructions	2-30
2.3.4.3.4	Integer Store Instructions.....	2-31
2.3.4.3.5	Integer Load and Store with Byte-Reverse Instructions.....	2-32
2.3.4.3.6	Integer Load and Store Multiple Instructions	2-32
2.3.4.3.7	Integer Load and Store String Instructions.....	2-33
2.3.4.3.8	Floating-Point Load and Store Address Generation.....	2-34
2.3.4.3.9	Floating-Point Load Instructions	2-34
2.3.4.3.10	Floating-Point Store Instructions.....	2-35
2.3.4.4	Branch and Flow Control Instructions.....	2-36
2.3.4.4.1	Branch Instruction Address Calculation.....	2-36
2.3.4.4.2	Branch Instructions.....	2-37
2.3.4.4.3	Condition Register Logical Instructions.....	2-37
2.3.4.5	Trap Instructions	2-38
2.3.4.6	Processor Control Instructions.....	2-38
2.3.4.6.1	Move To/From Condition Register Instructions.....	2-38
2.3.4.7	Memory Synchronization Instructions—UISA	2-38
2.3.5	PowerPC VEA Instructions	2-40
2.3.5.1	Processor Control Instructions.....	2-40
2.3.5.2	Memory Synchronization Instructions—VEA	2-41
2.3.5.3	Memory Control Instructions—VEA	2-41
2.3.5.4	External Control Instructions.....	2-43
2.3.6	PowerPC OEA Instructions	2-43
2.3.6.1	System Linkage Instructions.....	2-43
2.3.6.2	Processor Control Instructions—OEA	2-44
2.3.6.2.1	Move To/From Machine State Register Instructions.....	2-44
2.3.6.2.2	Move To/From Special-Purpose Register Instructions	2-44
2.3.6.3	Memory Control Instructions—OEA	2-45
2.3.6.3.1	Supervisor-Level Cache Management Instruction	2-45
2.3.6.3.2	Segment Register Manipulation Instructions	2-46
2.3.6.3.3	Translation Lookaside Buffer Management Instructions	2-46
2.3.7	Recommended Simplified Mnemonics.....	2-47
2.3.8	Implementation-Specific Instructions.....	2-47

Chapter 3

Instruction and Data Cache Operation

3.1	Instruction Cache Organization and Control	3-3
3.1.1	Instruction Cache Organization	3-3
3.1.2	Instruction Cache Fill Operations	3-4
3.1.3	Instruction Cache Control.....	3-4
3.1.3.1	Instruction Cache Invalidation.....	3-4
3.1.3.2	Instruction Cache Disabling	3-4

Contents

Paragraph Number	Title	Page Number
3.1.3.3	Instruction Cache Locking.....	3-5
3.2	Data Cache Organization and Control	3-5
3.2.1	Data Cache Organization	3-5
3.2.2	Data Cache Fill Operations.....	3-6
3.2.3	Data Cache Control.....	3-6
3.2.3.1	Data Cache Invalidation	3-6
3.2.3.2	Data Cache Disabling	3-7
3.2.3.3	Data Cache Locking	3-7
3.2.3.4	Data Cache Operations and Address Broadcasts.....	3-7
3.2.4	Data Cache Touch Load Support	3-8
3.3	Basic Data Cache Operations	3-8
3.3.1	Data Cache Fill	3-8
3.3.2	Data Cache Cast-Out Operation	3-8
3.3.3	Cache Block Push Operation.....	3-9
3.4	Data Cache Transactions on Bus	3-9
3.4.1	Single-Beat Transactions	3-9
3.4.2	Burst Transactions	3-9
3.4.3	Access to Direct-Store Segments.....	3-10
3.5	Memory Management/Cache Access Mode Bits—W, I, M, and G.....	3-10
3.5.1	Write-Through Attribute (W)	3-11
3.5.2	Caching-Inhibited Attribute (I).....	3-12
3.5.3	Memory Coherency Attribute (M).....	3-12
3.5.4	Guarded Attribute (G).....	3-13
3.5.5	W, I, and M Bit Combinations	3-13
3.5.5.1	Out-of-Order Execution and Guarded Memory	3-14
3.5.5.2	Effects of Out-of-Order Data Accesses	3-14
3.5.5.3	Effects of Out-of-Order Instruction Fetches.....	3-15
3.6	Cache Coherency—MEI Protocol	3-15
3.6.1	MEI State Definitions	3-16
3.6.2	MEI State Diagram	3-16
3.6.3	MEI Hardware Considerations	3-17
3.6.4	Coherency Precautions	3-19
3.6.4.1	Coherency in Single-Processor Systems	3-19
3.6.5	Load and Store Coherency Summary	3-19
3.6.6	Atomic Memory References	3-20
3.6.7	Cache Reaction to Specific Bus Operations	3-20
3.6.8	Operations Causing ARTRY Assertion	3-21
3.6.9	Enveloped High-Priority Cache Block Push Operation	3-22
3.7	Cache Control Instructions	3-22
3.7.1	Data Cache Block Invalidate (dcbi) Instruction.....	3-24
3.7.2	Data Cache Block Touch (dcbt) Instruction.....	3-24

Contents

Paragraph Number	Title	Page Number
3.7.3	Data Cache Block Touch for Store (dcbtst) Instruction.....	3-24
3.7.4	Data Cache Block Clear to Zero (dcbz) Instruction	3-24
3.7.5	Data Cache Block Store (dcbst) Instruction.....	3-25
3.7.6	Data Cache Block Flush (dcbf) Instruction.....	3-25
3.7.7	Enforce In-Order Execution of I/O Instruction (eieio)	3-25
3.7.8	Instruction Cache Block Invalidate (icbi) Instruction	3-26
3.7.9	Instruction Synchronize (isync) Instruction	3-26
3.8	Bus Operations Caused by Cache Control Instructions.....	3-26
3.9	Bus Interface	3-27
3.10	MEI State Transactions	3-28

Chapter 4 Exceptions

4.1	Exception Classes	4-2
4.1.1	Exception Priorities.....	4-6
4.1.2	Summary of Front-End Exception Handling	4-8
4.2	Exception Processing	4-9
4.2.1	Exception Processing Registers	4-9
4.2.1.1	SRR0 and SRR1 Bit Settings.....	4-9
4.2.1.2	MSR Bit Settings	4-11
4.2.2	Enabling and Disabling Exceptions	4-13
4.2.3	Steps for Exception Processing.....	4-13
4.2.4	Setting MSR[RI]	4-14
4.2.5	Returning from an Exception Handler	4-14
4.3	Process Switching	4-15
4.4	Exception Latencies	4-15
4.5	Exception Definitions	4-16
4.5.1	Reset Exceptions (0x00100)	4-17
4.5.1.1	Hard Reset and Power-On Reset	4-17
4.5.1.2	Soft Reset.....	4-18
4.5.2	Machine Check Exception (0x00200)	4-19
4.5.2.1	Machine Check Exception Enabled (MSR[ME] = 1).....	4-20
4.5.2.2	Checkstop State (MSR[ME] = 0)	4-20
4.5.3	DSI Exception (0x00300)	4-21
4.5.4	ISI Exception (0x00400).....	4-23
4.5.5	External Interrupt (0x00500)	4-23
4.5.6	Alignment Exception (0x00600)	4-24
4.5.6.1	Integer Alignment Exceptions	4-25
4.5.6.2	Load/Store Multiple Alignment Exceptions	4-26
4.5.7	Program Exception (0x00700).....	4-27
4.5.7.1	IEEE Floating-Point Exception Program Exceptions.....	4-27

Contents

Paragraph Number	Title	Page Number
4.5.7.2	Illegal, Reserved, and Unimplemented Instructions Program Exceptions.....	4-28
4.5.8	Floating-Point Unavailable Exception (0x00800)	4-28
4.5.9	Decrementer Exception (0x00900)	4-28
4.5.10	System Call Exception (0x00C00)	4-29
4.5.11	Trace Exception (0x00D00).....	4-29
4.5.11.1	Single-Step Instruction Trace Mode	4-30
4.5.11.2	Branch Trace Mode	4-30
4.5.12	Instruction TLB Miss Exception (0x01000).....	4-30
4.5.13	Data TLB Miss on Load Exception (0x01100).....	4-31
4.5.14	Data TLB Miss on Store Exception (0x01200)	4-32
4.5.15	Instruction Address Breakpoint Exception (0x01300)	4-32
4.5.16	System Management Interrupt (0x01400)	4-34

Chapter 5 Memory Management

5.1	MMU Features	5-2
5.1.1	Memory Addressing	5-3
5.1.2	MMU Organization.....	5-3
5.1.3	Address Translation Mechanisms	5-8
5.1.4	Memory Protection Facilities.....	5-10
5.1.5	Page History Information.....	5-11
5.1.6	General Flow of MMU Address Translation	5-11
5.1.6.1	Real Addressing Mode and Block Address Translation Selection	5-11
5.1.6.2	Page Address Translation Selection	5-12
5.1.7	MMU Exceptions Summary	5-14
5.1.8	MMU Instructions and Register Summary	5-17
5.2	Real Addressing Mode.....	5-19
5.3	Block Address Translation.....	5-20
5.4	Memory Segment Model	5-20
5.4.1	Page History Recording	5-21
5.4.1.1	Referenced Bit	5-22
5.4.1.2	Changed Bit	5-22
5.4.1.3	Scenarios for Referenced and Changed Bit Recording	5-23
5.4.2	Page Memory Protection	5-24
5.4.3	TLB Description	5-24
5.4.3.1	TLB Organization	5-25
5.4.3.2	TLB Entry Invalidiation.....	5-26
5.4.4	Page Address Translation Summary	5-27
5.5	Page Table Search Operation.....	5-27
5.5.1	Page Table Search Operation—Conceptual Flow	5-27

Contents

Paragraph Number	Title	Page Number
5.5.2	Implementation-Specific Table Search Operation.....	5-31
5.5.2.1	Resources for Table Search Operations	5-31
5.5.2.1.1	Data and Instruction TLB Miss Address Registers (DMISS and IMISS).....	5-33
5.5.2.1.2	Data and Instruction TLB Compare Registers (DCMP and ICMP).....	5-34
5.5.2.1.3	Primary and Secondary Hash Address Registers (HASH1 and HASH2)	5-34
5.5.2.1.4	Required Physical Address Register (RPA)	5-35
5.5.2.2	Software Table Search Operation	5-36
5.5.2.2.1	Flow for Example Exception Handlers	5-36
5.5.2.2.2	Code for Example Exception Handlers	5-40
5.5.3	Page Table Updates.....	5-47
5.5.4	Segment Register Updates	5-47

Chapter 6 Instruction Timing

6.1	Terminology and Conventions.....	6-1
6.2	Instruction Timing Overview.....	6-3
6.3	Timing Considerations	6-7
6.3.1	General Instruction Flow	6-8
6.3.2	Instruction Fetch Timing.....	6-10
6.3.2.1	Cache Arbitration.....	6-10
6.3.2.2	Cache Hit	6-10
6.3.2.3	Cache Miss.....	6-13
6.3.3	Instruction Dispatch and Completion Considerations	6-15
6.3.3.1	Rename Register Operation	6-15
6.3.3.2	Instruction Serialization.....	6-16
6.3.3.3	Execution Unit Considerations	6-17
6.4	Execution Unit Timings	6-17
6.4.1	Branch Processing Unit Execution Timing.....	6-17
6.4.1.1	Branch Folding	6-17
6.4.1.2	Static Branch Prediction	6-18
6.4.1.2.1	Predicted Branch Timing Examples	6-19
6.4.2	Integer Unit Execution Timing	6-21
6.4.3	Floating-Point Unit Execution Timing	6-21
6.4.4	Load/Store Unit Execution Timing.....	6-22
6.4.5	System Register Unit Execution Timing	6-22
6.5	Memory Performance Considerations	6-22
6.5.1	Copy-Back Mode	6-23
6.5.2	Write-Through Mode	6-23
6.5.3	Cache-Inhibited Accesses	6-23

Contents

Paragraph Number	Title	Page Number
6.6	Instruction Scheduling Guidelines.....	6-24
6.6.1	Branch, Dispatch, and Completion Unit Resource Requirements.....	6-25
6.6.1.1	Branch Resolution Resource Requirements	6-25
6.6.1.2	Dispatch Unit Resource Requirements	6-25
6.6.1.3	Completion Unit Resource Requirements	6-26
6.7	Instruction Latency Summary.....	6-26

Chapter 7 Signal Descriptions

7.1	Signal Configuration.....	7-2
7.2	Signal Descriptions	7-2
7.2.1	Address Bus Arbitration Signals.....	7-2
7.2.1.1	Bus Request (\overline{BR})—Output	7-3
7.2.1.2	Bus Grant (\overline{BG})—Input	7-4
7.2.1.3	Address Bus Busy (ABB).....	7-4
7.2.1.3.1	Address Bus Busy (ABB)—Output.....	7-5
7.2.1.3.2	Address Bus Busy (ABB)—Input	7-5
7.2.2	Address Transfer Start Signals.....	7-5
7.2.2.1	Transfer Start (\overline{TS})	7-6
7.2.2.1.1	Transfer Start (\overline{TS})—Output	7-6
7.2.2.1.2	Transfer Start (\overline{TS})—Input	7-6
7.2.3	Address Transfer Signals	7-6
7.2.3.1	Address Bus (A[0:31]).....	7-6
7.2.3.1.1	Address Bus (A[0:31])—Output	7-7
7.2.3.1.2	Address Bus (A[0:31])—Input	7-7
7.2.3.2	Address Bus Parity (AP[0:3])	7-7
7.2.3.2.1	Address Bus Parity (AP[0:3])—Output.....	7-7
7.2.3.2.2	Address Bus Parity (AP[0:3])—Input	7-8
7.2.3.3	Address Parity Error (\overline{APE})—Output	7-8
7.2.4	Address Transfer Attribute Signals.....	7-8
7.2.4.1	Transfer Type (TT[0:4]).....	7-9
7.2.4.1.1	Transfer Type (TT[0:4])—Output	7-9
7.2.4.1.2	Transfer Type (TT[0:4])—Input	7-9
7.2.4.2	Transfer Size (TSIZ[0:2])—Output	7-12
7.2.4.3	Transfer Burst (\overline{TBST})	7-12
7.2.4.3.1	Transfer Burst (\overline{TBST})—Output.....	7-12
7.2.4.3.2	Transfer Burst (\overline{TBST})—Input	7-13
7.2.4.4	Transfer Code (TC[0:1])—Output.....	7-13
7.2.4.5	Cache Inhibit (\overline{CI})—Output.....	7-13
7.2.4.6	Write-Through (\overline{WT})—Output	7-14
7.2.4.7	Global (\overline{GBL}).....	7-14

Contents

Paragraph Number	Title	Page Number
7.2.4.7.1	Global (<u>GBL</u>)—Output.....	7-14
7.2.4.7.2	Global (<u>GBL</u>)—Input	7-14
7.2.4.8	Cache Set Entry (CSE[0:1])—Output	7-14
7.2.5	Address Transfer Termination Signals.....	7-15
7.2.5.1	Address Acknowledge (<u>AACK</u>)—Input.....	7-15
7.2.5.2	Address Retry (<u>ARTRY</u>).....	7-15
7.2.5.2.1	Address Retry (<u>ARTRY</u>)—Output	7-15
7.2.5.2.2	Address Retry (<u>ARTRY</u>)—Input	7-16
7.2.6	Data Bus Arbitration Signals	7-16
7.2.6.1	Data Bus Grant (<u>DBG</u>)—Input	7-17
7.2.6.2	Data Bus Write Only (<u>DBWO</u>)—Input	7-17
7.2.6.3	Data Bus Busy (<u>DBB</u>)	7-18
7.2.6.3.1	Data Bus Busy (<u>DBB</u>)—Output	7-18
7.2.6.3.2	Data Bus Busy (<u>DBB</u>)—Input	7-18
7.2.7	Data Transfer Signals.....	7-18
7.2.7.1	Data Bus (DH[0:31], DL[0:31])	7-18
7.2.7.1.1	Data Bus (DH[0:31], DL[0:31])—Output	7-19
7.2.7.1.2	Data Bus (DH[0:31], DL[0:31])—Input.....	7-19
7.2.7.2	Data Bus Parity (DP[0:7])	7-19
7.2.7.2.1	Data Bus Parity (DP[0:7])—Output	7-19
7.2.7.2.2	Data Bus Parity (DP[0:7])—Input.....	7-20
7.2.7.3	Data Parity Error (<u>DPE</u>)—Output	7-20
7.2.7.4	Data Bus Disable (<u>DBDIS</u>)—Input	7-21
7.2.8	Data Transfer Termination Signals	7-21
7.2.8.1	Transfer Acknowledge (<u>TA</u>)—Input.....	7-21
7.2.8.2	Data Retry (<u>DRTRY</u>)—Input.....	7-22
7.2.8.3	Transfer Error Acknowledge (<u>TEA</u>)—Input	7-22
7.2.9	System Status Signals	7-23
7.2.9.1	Interrupt (<u>INT</u>)—Input.....	7-23
7.2.9.2	System Management Interrupt (<u>SMI</u>)—Input	7-23
7.2.9.3	Machine Check Interrupt (<u>MCP</u>)—Input.....	7-24
7.2.9.4	Checkstop Input (<u>CKSTP_IN</u>)—Input	7-24
7.2.9.5	Checkstop Output (<u>CKSTP_OUT</u>)—Output	7-24
7.2.9.6	Reset Signals.....	7-25
7.2.9.6.1	Hard Reset (<u>HRESET</u>)—Input	7-25
7.2.9.6.2	Soft Reset (<u>SRESET</u>)—Input	7-25
7.2.9.7	Processor Status Signals	7-26
7.2.9.7.1	Quiescent Request (<u>QREQ</u>).....	7-26
7.2.9.7.2	Quiescent Acknowledge (<u>QACK</u>)	7-26
7.2.9.7.3	Reservation (<u>RSRV</u>)—Output	7-27
7.2.9.7.4	Time Base Enable (TBEN)—Input	7-27

Contents

Paragraph Number	Title	Page Number
7.2.9.7.5	TLBI Sync (TLBISYNC)	7-27
7.2.10	COP/Scan Interface.....	7-27
7.2.11	Pipeline Tracking Support	7-28
7.2.12	Clock Signals	7-29
7.2.12.1	System Clock (SYSCLK)—Input.....	7-29
7.2.12.2	Test Clock (CLK_OUT)—Output	7-30
7.2.12.3	PLL Configuration (PLL_CFG[0:3])—Input	7-30
7.2.13	Power and Ground Signals	7-32

Chapter 8 System Interface Operation

8.1	Overview.....	8-1
8.1.1	Operation of the Instruction and Data Caches.....	8-2
8.1.2	Operation of the System Interface	8-4
8.1.2.1	Optional 32-Bit Data Bus Mode	8-5
8.1.3	Direct-Store Accesses	8-6
8.2	Memory Access Protocol	8-6
8.2.1	Arbitration Signals.....	8-7
8.2.2	Address Pipelining and Split-Bus Transactions.....	8-8
8.3	Address Bus Tenure	8-9
8.3.1	Address Bus Arbitration	8-9
8.3.2	Address Transfer	8-11
8.3.2.1	Address Bus Parity	8-12
8.3.2.2	Address Transfer Attribute Signals.....	8-12
8.3.2.2.1	Transfer Type (TT[0:4]) Signals.....	8-12
8.3.2.2.2	Transfer Size (TSIZ[0:2]) Signals	8-13
8.3.2.3	Burst Ordering During Data Transfers	8-13
8.3.2.4	Effect of Alignment in Data Transfers (64-Bit Bus)	8-14
8.3.2.5	Effect of Alignment in Data Transfers (32-Bit Bus)	8-16
8.3.2.5.1	Alignment of External Control Instructions	8-18
8.3.2.6	Transfer Code (TC[0:1]) Signals	8-19
8.3.3	Address Transfer Termination	8-19
8.4	Data Bus Tenure.....	8-21
8.4.1	Data Bus Arbitration	8-21
8.4.1.1	Using the DBB Signal	8-22
8.4.2	Data Bus Write Only	8-22
8.4.3	Data Transfer	8-23
8.4.4	Data Transfer Termination	8-24
8.4.4.1	Normal Single-Beat Termination.....	8-25
8.4.4.2	Normal Burst Termination.....	8-26
8.4.4.3	Data Transfer Termination Due to a Bus Error.....	8-27

Contents

Paragraph Number	Title	Page Number
8.4.5	Memory Coherency—MEI Protocol	8-28
8.5	Timing Examples	8-30
8.6	Optional Bus Configurations	8-37
8.6.1	32-Bit Data Bus Mode	8-37
8.6.2	No-DRTRY Mode.....	8-39
8.6.3	Reduced-Pinout Mode	8-40
8.7	Interrupt, Checkstop, and Reset Signals	8-40
8.7.1	External Interrupts	8-40
8.7.2	Checkstops	8-40
8.7.3	Reset Inputs.....	8-41
8.7.4	System Quiesce Control Signals.....	8-41
8.8	Processor State Signals	8-41
8.8.1	Support for the Iwarx/stwcx. Instruction Pair	8-41
8.8.2	TLBISYNC Input	8-42
8.9	IEEE 1149.1-Compliant Interface.....	8-42
8.9.1	IEEE 1149.1 Interface Description	8-42
8.10	Using DBWO.....	8-43

Chapter 9 Power Management

9.1	Overview.....	9-1
9.2	Dynamic Power Management.....	9-1
9.3	Programmable Power Modes.....	9-2
9.3.1	Power Management Modes	9-3
9.3.1.1	Full-Power Mode with DPM Disabled	9-3
9.3.1.2	Full-Power Mode with DPM Enabled	9-3
9.3.1.3	Doze Mode.....	9-3
9.3.1.4	Nap Mode	9-4
9.3.1.5	Sleep Mode	9-5
9.3.2	Power Management Software Considerations.....	9-6
9.4	Example Code Sequence for Entering Processor Sleep Mode	9-6

Appendix A PowerPC Instruction Set Listings

A.1	Instructions Sorted by Mnemonic	A-1
A.2	Instructions Sorted by Opcode.....	A-8
A.3	Instructions Grouped by Functional Categories	A-15
A.4	Instructions Sorted by Form	A-25
A.5	Instruction Set Legend	A-36

Contents

Paragraph Number	Title	Page Number
	Appendix B Instructions Not Implemented	
	Appendix C Revision History	
C.1	Revision Changes From Revision 2 to Revision 3	C-1
C.2	Revision Changes From Revision 1 to Revision 2	C-2
	Glossary of Terms and Abbreviations	
	Index	

Contents

Figures

Figure Number	Title	Page Number
1-1	MPC603e Microprocessor Block Diagram.....	1-2
1-2	Programming Model—Registers	1-18
1-3	Data Cache Organization	1-25
1-4	System Interface	1-34
1-5	Signal Groups	1-37
2-1	Programming Model—Registers	2-3
2-2	Hardware Implementation Register 0 (HID0)	2-6
2-3	Hardware Implementation Register 1 (HID1)	2-10
2-4	DMISS and IMISS Registers	2-10
2-5	DCMP and ICMP Registers.....	2-11
2-6	HASH1 and HASH2 Registers	2-11
2-7	Required Physical Address Register (RPA)	2-12
2-8	Instruction Address Breakpoint Register (IABR).....	2-12
3-1	Instruction Cache Organization	3-4
3-2	Data Cache Organization	3-6
3-3	Double-Word Address Ordering—Critical Double Word First	3-10
3-4	MEI Cache Coherency Protocol—State Diagram (WIM = 001).....	3-17
3-5	Bus Interface Address Buffers	3-28
4-1	Machine Status Save/Restore Register 0 (SSR0)	4-9
4-2	Machine Status Save/Restore Register 1 (SSR1)	4-9
4-3	Machine State Register (MSR)	4-11
5-1	MMU Conceptual Block Diagram—32-Bit Implementations.....	5-5
5-2	IMMU Block Diagram.....	5-6
5-3	DMMU Block Diagram	5-7
5-4	Address Translation Types	5-9
5-5	General Flow of Address Translation (Real Addressing Mode and Block)	5-12
5-6	General Flow of Page and Direct-Store Interface Address Translation	5-13
5-7	Segment Register and TLB Organization	5-25
5-8	Page Address Translation Flow for 32-Bit Implementations—TLB Hit.....	5-28
5-9	Primary Page Table Search—Conceptual Flow	5-30
5-10	Secondary Page Table Search Flow—Conceptual Flow	5-31
5-11	DMISS and IMISS Registers	5-34
5-12	DCMP and ICMP Registers.....	5-34
5-13	HASH1 and HASH2 Registers	5-35
5-14	Required Physical Address (RPA) Register	5-35

Figures

Figure Number	Title	Page Number
5-15	Flow for Example Software Table Search Operation	5-37
5-16	Check and Set R and C Bit Flow	5-38
5-17	Page Fault Setup Flow	5-39
5-18	Setup for Protection Violation Exceptions	5-40
6-1	Pipelined Execution Unit	6-4
6-2	Superscalar/Pipeline Diagram.....	6-5
6-3	MPC603e Microprocessor Pipeline Stages.....	6-7
6-4	Instruction Flow Diagram.....	6-9
6-5	Instruction Timing—Cache Hit	6-11
6-6	Instruction Timing—Cache Miss.....	6-14
6-7	Branch Instruction Timing.....	6-20
7-1	Signal Groups	7-3
7-2	IEEE 1149.1-Compliant Boundary Scan Interface	7-28
8-1	MPC603e Microprocessor Block Diagram.....	8-3
8-2	Overlapping Tenures on the Bus for a Single-Beat Transfer.....	8-6
8-3	Address Bus Arbitration	8-10
8-4	Address Bus Arbitration Showing Bus Parking	8-10
8-5	Address Bus Transfer.....	8-12
8-6	Snooped Address Cycle with <u>ARTRY</u>	8-20
8-7	Data Bus Arbitration	8-22
8-8	Normal Single-Beat Read Termination	8-25
8-9	Normal Single-Beat Write Termination	8-25
8-10	Normal Burst Transaction.....	8-26
8-11	Termination with <u>DRTRY</u>	8-27
8-12	Read Burst with <u>TA</u> Wait States and DRTRY	8-28
8-13	MEI Cache Coherency Protocol—State Diagram (WIM = 001).....	8-30
8-14	Fastest Single-Beat Reads.....	8-31
8-15	Fastest Single-Beat Writes.....	8-32
8-16	Single-Beat Reads Showing Data-Delay Controls	8-33
8-17	Single-Beat Writes Showing Data-Delay Controls	8-34
8-18	Burst Transfers with Data-Delay Controls	8-35
8-19	Use of Transfer Error Acknowledge (<u>TEA</u>)	8-36
8-20	32-Bit Data Bus Transfer (Eight-Beat Burst)	8-38
8-21	32-Bit Data Bus Transfer (Two-Beat Burst with <u>DRTRY</u>)	8-38
8-22	<u>DBWO</u> Transaction	8-43

Tables

Table Number	Title	Page Number
i	Acronyms and Abbreviated Terms.....	xxxii
ii	Terminology Conventions.....	xxxv
iii	Instruction Field Conventions.....	xxxvi
1-1	CSE[0:1] Signals	1-7
1-2	Generated SRR1[KEY] Bit	1-8
1-3	Additional/Changed HID0 Bits	1-19
1-4	Exception Classifications	1-28
1-5	Exceptions and Conditions	1-28
2-1	MSR[POW] and MSR[TGPR] Bits.....	2-5
2-2	HID0 Bit Functions	2-7
2-3	HID0[BCLK] and HID0[ECLK] CLK_OUT Configuration	2-9
2-4	HID1 Bit Settings	2-10
2-5	DCMP and ICMP Bit Settings	2-11
2-6	HASH1 and HASH2 Bit Settings.....	2-11
2-7	RPA Bit Settings	2-12
2-8	Instruction Address Breakpoint Register Bit Settings.....	2-13
2-9	Memory Operands	2-14
2-10	Integer Arithmetic Instructions.....	2-23
2-11	Integer Compare Instructions	2-24
2-12	Integer Logical Instructions.....	2-24
2-13	Integer Rotate Instructions	2-25
2-14	Integer Shift Instructions	2-26
2-15	Floating-Point Arithmetic Instructions	2-27
2-16	Floating-Point Multiply-Add Instructions	2-27
2-17	Floating-Point Rounding and Conversion Instructions	2-28
2-18	Floating-Point Compare Instructions	2-28
2-20	Floating-Point Move Instructions.....	2-29
2-21	Integer Load Instructions.....	2-31
2-22	Integer Store Instructions	2-31
2-23	Integer Load and Store with Byte-Reverse Instructions.....	2-32
2-24	Integer Load and Store Multiple Instructions.....	2-33
2-25	Integer Load and Store String Instructions.....	2-34
2-26	Floating-Point Load Instructions	2-35
2-27	Floating-Point Store Instructions.....	2-35
2-28	Branch Instructions.....	2-37

Tables

Table Number	Title	Page Number
2-29	Condition Register Logical Instructions.....	2-37
2-30	Trap Instructions.....	2-38
2-31	Move To/From Condition Register Instructions.....	2-38
2-32	Memory Synchronization Instructions—UISA.....	2-40
2-33	Move From Time Base Instruction.....	2-40
2-34	Memory Synchronization Instructions—VEA	2-41
2-35	User-Level Cache Instructions	2-42
2-36	External Control Instructions	2-43
2-37	System Linkage Instructions	2-43
2-38	Move To/From Machine State Register Instructions.....	2-44
2-39	Move To/From Special-Purpose Register Instructions.....	2-44
2-40	Implementation-Specific SPR Encodings (mfspr)	2-44
2-41	Segment Register Manipulation Instructions	2-46
2-42	Translation Lookaside Buffer Management Instructions	2-47
3-1	Combinations of W, I, and M Bits	3-13
3-2	MEI State Definitions	3-16
3-3	CSE[0:1] Signal Encoding	3-18
3-4	Memory Coherency Actions on Load Operations	3-19
3-5	Memory Coherency Actions on Store Operations.....	3-20
3-6	Response to Bus Transactions	3-20
3-7	Bus Operations Caused by Cache Control Instructions (WIM = 001)	3-26
3-8	MEI State Transitions	3-29
4-1	Exception Classifications	4-3
4-2	Exceptions and Conditions	4-4
4-3	Exception Priorities	4-6
4-4	SRR1 Bit Settings for Machine Check Exceptions	4-10
4-5	SRR1 Bit Settings for Software Table Search Operations	4-10
4-6	MSR Bit Settings	4-11
4-7	IEEE Floating-Point Exception Mode Bits	4-13
4-8	MSR Setting Due to Exception	4-16
4-9	Settings Caused by Hard Reset.....	4-17
4-10	Soft Reset Exception—Register Settings	4-19
4-11	Machine Check Exception—Register Settings	4-20
4-12	DSI Exception—Register Settings	4-21
4-13	External Interrupt—Register Settings	4-24
4-14	Alignment Interrupt—Register Settings	4-25
4-15	Access Types	4-26
4-16	Trace Exception—Register Settings	4-30
4-17	Instruction and Data TLB Miss Exceptions—Register Settings	4-31
4-18	Instruction Address Breakpoint Exception—Register Settings	4-32
4-19	Breakpoint Action for Multiple Modes Enabled for the Same Address.....	4-33

Tables

Table Number	Title	Page Number
4-20	System Management Interrupt—Register Settings	4-34
5-1	MMU Features Summary	5-2
5-2	Access Protection Options for Pages.....	5-10
5-3	Translation Exception Conditions	5-15
5-4	Other MMU Exception Conditions	5-16
5-5	Instruction Summary—MMU Control	5-17
5-6	MMU Registers	5-18
5-7	Table Search Operations to Update History Bits—TLB Hit Case	5-21
5-8	Model for Guaranteed R and C Bit Settings.....	5-24
5-9	Implementation-Specific Resources for Table Search Operations	5-32
5-10	Implementation-Specific SRR1 Bits	5-33
5-11	DCMP and ICMP Bit Settings	5-34
5-12	HASH1 and HASH2 Bit Settings.....	5-35
5-13	RPA Bit Settings	5-35
6-1	Branch Instructions.....	6-26
6-2	System Register Instructions	6-27
6-3	Condition Register Logical Instructions.....	6-27
6-4	Integer Instructions	6-28
6-5	Floating-Point Instructions	6-30
6-6	Load and Store Instructions.....	6-31
7-1	Transfer Encoding for the Bus Master.....	7-9
7-2	Snoop Hit Response	7-10
7-3	Implementation-Specific Transfer Encoding.....	7-11
7-4	Data Transfer Size	7-12
7-5	Encodings for TC[0:1] Signals	7-13
7-6	Data Bus Lane Assignments.....	7-19
7-7	DP[0:7] Signal Assignments	7-20
7-8	Pipeline Tracking Outputs	7-28
7-9	CLK_OUT Signal Configuration	7-30
7-10	PLL Configuration.....	7-31
8-1	Timing Diagram Legend.....	8-5
8-2	Transfer Size Signal Encodings.....	8-13
8-3	Burst Ordering—64-Bit Bus.....	8-14
8-4	Burst Ordering—32-Bit Bus.....	8-14
8-5	Aligned Data Transfers (64-Bit Bus).....	8-15
8-6	Misaligned Data Transfers (4-Byte Examples)	8-16
8-7	Aligned Data Transfers (32-Bit Bus Mode)	8-17
8-8	Misaligned 32-Bit Data Bus Transfer (4-Byte Examples)	8-18
8-9	Transfer Code Encoding	8-19
8-10	CSE[0:1] Signals	8-30
8-11	IEEE Interface Pin Descriptions.....	8-42

Tables

Table Number	Title	Page Number
9-1	MPC603e Programmable Power Modes	9-3
A-1	Complete Instruction List Sorted by Mnemonic.....	A-1
A-2	Complete Instruction List Sorted by Opcode.....	A-8
A-3	Integer Arithmetic Instructions	A-15
A-4	Integer Compare Instructions	A-16
A-5	Integer Logical Instructions	A-16
A-6	Integer Rotate Instructions	A-16
A-7	Integer Shift Instructions.....	A-17
A-8	Floating-Point Arithmetic Instructions	A-17
A-9	Floating-Point Multiply-Add Instructions	A-18
A-10	Floating-Point Rounding and Conversion Instructions.....	A-18
A-11	Floating-Point Compare Instructions	A-18
A-12	Floating-Point Status and Control Register Instructions.....	A-18
A-13	Integer Load Instructions	A-19
A-14	Integer Store Instructions	A-19
A-15	Integer Load and Store with Byte-Reverse Instructions	A-20
A-16	Integer Load and Store Multiple Instructions	A-20
A-17	Integer Load and Store String Instructions	A-20
A-18	Memory Synchronization Instructions	A-21
A-19	Floating-Point Load Instructions	A-21
A-20	Floating-Point Store Instructions	A-21
A-21	Floating-Point Move Instructions	A-22
A-22	Branch Instructions	A-22
A-23	Condition Register Logical Instructions	A-22
A-24	System Linkage Instructions	A-22
A-25	Trap Instructions	A-23
A-26	Processor Control Instructions	A-23
A-27	Cache Management Instructions	A-23
A-28	Segment Register Manipulation Instructions	A-24
A-29	Lookaside Buffer Management Instructions	A-24
A-30	External Control Instructions	A-24
A-31	I-Form	A-25
A-32	B-Form	A-25
A-33	SC-Form	A-25
A-34	D-Form.....	A-25
A-35	DS-Form.....	A-27
A-36	X-Form.....	A-27
A-37	XL-Form	A-31
A-38	XFX-Form.....	A-32
A-39	XFL-Form	A-32
A-40	XS-Form.....	A-32

Tables

Table Number	Title	Page Number
A-41	XO-Form	A-33
A-42	A-Form	A-33
A-43	M-Form	A-34
A-44	MD-Form	A-35
A-45	MDS-Form	A-35
A-46	PowerPC Instruction Set Legend	A-36
B-1	32-Bit Instructions Not Implemented by the MPC603e	B-1
B-2	64-Bit Instructions Not Implemented by the MPC603e	B-1
B-3	64-Bit SPR Encoding Not Implemented by the MPC603e	B-2



Tables

About This Book

The primary objective of this user's manual is to define the functionality of the MPC603e microprocessor for use by software and hardware developers.

It is important to note that this book is to be used with the *PowerPC Microprocessor Family: The Programming Environments*, referred to as the *Programming Environments Manual*. Contact your sales representative to obtain a copy. Because the PowerPC architecture is designed to be flexible to support a broad range of processors, the *Programming Environments Manual* provides a general description of features common to PowerPC processors and indicates those features that are optional or that may be implemented differently in the design of each processor.

This document describes in detail the MPC603e features not defined by the architecture. This document and the *Programming Environments Manual* distinguish between the three levels, or programming environments, of the PowerPC architecture, which are as follows:

- User instruction set architecture (UISA)—The UISA defines the architecture level to which user-level software should conform. The UISA defines the base user-level instruction set, user-level registers, data types, memory conventions, and the memory and programming models seen by application programmers.
- Virtual environment architecture (VEA)—The VEA, which is the smallest component of the PowerPC architecture, defines additional user-level functionality that falls outside typical user-level software requirements. The VEA describes the memory model for an environment in which multiple processors or other devices can access external memory, and defines aspects of the cache model and cache control instructions from a user-level perspective. The resources defined by the VEA are particularly useful for optimizing memory accesses and managing resources in an environment in which other processors and devices can access external memory.

Implementations that conform to the VEA also conform to the UISA but may not necessarily adhere to the OEA.

- Operating environment architecture (OEA)—The OEA defines supervisor-level resources typically required by an operating system. The OEA defines the memory management model, supervisor-level registers, and exception model.

Implementations that conform to the OEA also conform to the UISA and VEA.

Note that some resources are defined more generally at one level in the architecture and more specifically at another. For example, conditions that cause a floating-point exception are defined by the UIISA, while the exception mechanism is defined by the OEA.

Because it is important to distinguish between the levels of the architecture to ensure compatibility across multiple platforms, those distinctions are shown clearly throughout this book.

For ease in reference, topics in this book are arranged to build on one another, beginning with a description and complete summary of MPC603e-specific registers and progressing to more specialized topics such as MPC603e-specific details regarding the cache, exception, and memory management models. As such, chapters may include information from multiple levels of the architecture. For example, the discussion of the cache model uses information from both the VEA and the OEA.

The PowerPC Architecture: A Specification for a New Family of RISC Processors defines the architecture from the perspective of the three programming environments and remains the defining document for the PowerPC architecture.

The information in this book is subject to change without notice, as described in the disclaimers on the title page. As with any technical documentation, it is the readers' responsibility to be sure they are using the most recent version of the documentation. For updates to this document, refer to: <http://www.freescale.com>.

Audience

This manual is intended for system software and hardware developers and applications programmers who want to develop products using the MPC603e microprocessors. It is assumed that the reader understands operating systems, microprocessor system design, basic principles of RISC processing, and details of the PowerPC architecture.

Organization

The following summary briefly describes the major sections of this manual:

- Chapter 1, “Overview,” is useful for readers who want a general understanding of the features and functions of the PowerPC architecture and the MPC603e. This chapter describes the flexible nature of the PowerPC architecture definition and provides an overview of how the PowerPC architecture defines the register set, operand conventions, addressing modes, instruction set, cache model, exception model, and memory management model.
- Chapter 2, “Programming Model,” provides a brief synopsis of the registers implemented on the MPC603e, operand conventions, an overview of the PowerPC addressing modes, and a list of the instructions implemented by the MPC603e. Instructions are organized by function.

- Chapter 3, “Instruction and Data Cache Operation,” provides a discussion of the cache and memory model as implemented on the MPC603e.
- Chapter 4, “Exceptions,” describes the exception model defined in the OEA and the specific exception model implemented on the MPC603e.
- Chapter 5, “Memory Management,” describes MPC603e implementation of the memory management unit specified by the OEA.
- Chapter 6, “Instruction Timing,” provides information about latencies, interlocks, special situations, and various conditions to help make programming more efficient. This chapter is of special interest to software engineers and system designers.
- Chapter 7, “Signal Descriptions,” provides descriptions of individual signals of the MPC603e.
- Chapter 8, “System Interface Operation,” describes signal timings for various operations. It also provides information for interfacing to the MPC603e.
- Chapter 9, “Power Management,” provides information about power saving modes for the MPC603e.
- Appendix A, “PowerPC Instruction Set Listings,” lists all the PowerPC instructions while indicating those instructions that are not implemented by the MPC603e; it also includes the instructions that are specific to the MPC603e. Instructions are grouped according to mnemonic, opcode, function, and form. Also included is a quick reference table that contains general information, such as the architecture level, privilege level, and form, and indicates if the instruction is 64-bit and optional.
- Appendix B, “Instructions Not Implemented,” provides a list of PowerPC instructions not implemented by the MPC603e.
- Appendix C, “Revision History,” lists the major differences between Revision 1, Revision 2, and Revision 3 of the *MPC603e RISC Microprocessor User’s Manual*.
- This manual also includes a glossary and an index.

Suggested Reading

This section lists additional reading that provides background for the information in this manual as well as general information about the PowerPC architecture.

General Information

The following documentation, available through Morgan-Kaufmann Publishers, 340 Pine Street, Sixth Floor, San Francisco, CA, provides useful information about the PowerPC architecture and computer architecture in general:

- *The PowerPC Architecture: A Specification for a New Family of RISC Processors*, Second Edition, by International Business Machines, Inc.

For updates to the specification, see <http://www.austin.ibm.com/tech/ppc-chg.html>.

- *PowerPC Microprocessor Common Hardware Reference Platform: A System Architecture*, by Apple Computer, Inc., International Business Machines, Inc. and Freescale Semiconductor, Inc.
- *Computer Architecture: A Quantitative Approach*, Second Edition, by John L. Hennessy and David A. Patterson
- *Computer Organization and Design: The Hardware/Software Interface*, Second Edition, David A. Patterson and John L. Hennessy
- *Inside Macintosh: PowerPC System Software*, Addison-Wesley Publishing Company, One Jacob Way, Reading, MA, 01867; Tel. (800) 282-2732 (U.S.A.), (800) 637-0029 (Canada), (716) 871-6555 (International).

Related Documentation

Freescale documentation is available from the sources listed on the back cover of this manual; the document order numbers are included in parentheses for ease in ordering:

- *Programming Environments Manual for 32-Bit Implementations of the PowerPC Architecture* (MPCFPE32B/AD)—Describes resources defined by the PowerPC architecture.
- User's manuals—These books provide details about individual implementations and are intended for use with the *Programming Environments Manual*.
- Addenda/errata to user's manuals—Because some processors have follow-on parts an addendum is provided that describes the additional features and functionality changes. These addenda are intended for use with the corresponding user's manual.
- *Implementation Variances Relative to Rev. 1 of The Programming Environments Manual* is available at <http://www.freescale.com>
- Hardware specifications—Hardware specifications provide specific data regarding bus timing, signal behavior, and AC, DC, and thermal characteristics, as well as other design considerations.
- Technical Summaries—Each device has a technical summary that provides an overview of its features. This document is roughly the equivalent to the overview (Chapter 1) of an implementation's user's manual.
- *The Programmer's Reference Guide for the PowerPC Architecture* (MPCPRG/D)—This concise reference includes the register summary, memory control model, exception vectors, and the PowerPC instruction set.
- *The Programmer's Pocket Reference Guide for the PowerPC Architecture* (MPCPRGREF/D)—This foldout card provides an overview of PowerPC registers, instructions, and exceptions for 32-bit implementations.
- Application notes—These short documents contain useful information about specific design issues useful to programmers and engineers working with Freescale processors.

- Documentation for support chips—These include the following:
 - *MPC106 PCI Bridge/Memory Controller User's Manual* (MPC106UM/AD)
 - *MPC107 PCI Bridge/Memory Controller Technical Summary* (MPC107TS/D)
 - *MPC107 PCI Bridge/Memory Controller User's Manual* (MPC107UM/AD)

Additional literature is published as new processors become available. For a current list of documentation, refer to: <http://www.mot.com/semiconductors>.

Conventions

This document uses the following notational conventions:

cleared/set	When a bit takes the value zero, it is said to be cleared; when it takes a value of one, it is said to be set.
mnemonics	Instruction mnemonics are shown in lowercase bold.
<i>italics</i>	Italics indicate variable command parameters, for example, bcctrx . Book titles in text are set in italics.
0x0	Prefix to denote hexadecimal number
0b0	Prefix to denote binary number
rA, rB	Instruction syntax used to identify a source GPR
rA 0	Contents of a specified GPR or the value zero
rD	Instruction syntax used to identify a destination GPR
frA, frB, frC	Instruction syntax used to identify a source FPR
frD	Instruction syntax used to identify a destination FPR
REG[FIELD]	Abbreviations for registers are shown in uppercase text. Specific bits, fields, or ranges appear in brackets. For example, MSR[LE] refers to the little-endian mode enable bit in the machine state register.
x	In some contexts, such as signal encodings, an unitalicized x indicates a don't care.
<i>x</i>	An italicized x indicates an alphanumeric variable.
<i>n</i>	An italicized n indicates an numeric variable.
<u>ARTRY</u>	A bar over a signal name indicates that the signal is active low.
¬	NOT logical operator
&	AND logical operator
	OR logical operator

0 0 0 0

Indicates reserved bits or bit fields in a register. Although these bits can be written to as ones or zeros, they are always read as zeros.

Acronyms and Abbreviations

Table i contains acronyms and abbreviations that are used in this document.

Table i. Acronyms and Abbreviated Terms

Term	Meaning
ALU	Arithmetic logic unit
BAT	Block address translation
BHT	Branch history table
BIST	Built-in self test
BIU	Bus interface unit
BPU	Branch processing unit
BSDL	Boundary-scan description language
BTIC	Branch target instruction cache
BUC	Bus unit controller
BUID	Bus unit ID
CAR	Cache address register
CIA	Current instruction address
CMOS	Complementary metal-oxide semiconductor
COP	Common on-chip processor
CQ	Completion queue
CR	Condition register
CRTRY	Cache retry queue
CTR	Count register
DABR	Data address breakpoint register
DAR	Data address register
DBAT	Data BAT
DCMP	Data TLB compare
DEC	Decrementer register
DLL	Delay-locked loop
DMISS	Data TLB miss address
DMMU	Data MMU
DSISR	Register used for determining the source of a DS1 exception
DTLB	Data translation lookaside buffer

Table i. Acronyms and Abbreviated Terms (continued)

Term	Meaning
EA	Effective address
EAR	External access register
ECC	Error checking and correction
FIFO	First-in-first-out
FPR	Floating-point register
FPSCR	Floating-point status and control register
FPU	Floating-point unit
GPR	General-purpose register
HASH1	Primary hash address
HASH2	Secondary hash address
IABR	Instruction address breakpoint register
IBAT	Instruction BAT
ICMP	Instruction TLB compare
IEEE	Institute for Electrical and Electronics Engineers
IMISS	Instruction TLB miss address
IQ	Instruction queue
ITLB	Instruction translation lookaside buffer
IU	Integer unit
JTAG	Joint Test Action Group
L2	Secondary cache (level 2 cache)
LIFO	Last-in-first-out
LR	Link register
LRU	Least recently used
LSB	Least-significant byte
lsb	Least-significant bit
LSU	Load/store unit
MEI	Modified/exclusive/invalid
MESI	Modified/exclusive/shared/invalid—cache coherency protocol
MMU	Memory management unit
MQ	MQ register
MSB	Most-significant byte
msb	Most-significant bit
MSR	Machine state register

Table i. Acronyms and Abbreviated Terms (continued)

Term	Meaning
NaN	Not a number
No-op	No operation
OEA	Operating environment architecture
PEM	<i>Programming Environments Manual</i>
PID	Processor identification tag
PIM	<i>Programming Interface Manual</i>
PIR	Processor identification register
PLL	Phase-locked loop
POR	Power-on reset
POWER	Performance Optimized with Enhanced RISC architecture
PTE	Page table entry
PTEG	Page table entry group
PVR	Processor version register
RAW	Read-after-write
RISC	Reduced instruction set computing
RPA	Required physical address
RTL	Register transfer language
RWITM	Read with intent to modify
SDR1	Register that specifies the page table base address for virtual-to-physical address translation
SIMM	Signed immediate value
SLB	Segment lookaside buffer
SMI	System management interrupt
SPR	Special-purpose register
SR _n	Segment register
SRR0	Machine status save/restore register 0
SRR1	Machine status save/restore register 1
SRU	System register unit
TAP	Test access port
TB	Time base facility
TBL	Time base lower register
TBU	Time base upper register
TLB	Translation lookaside buffer
TTL	Transistor-to-transistor logic

Table i. Acronyms and Abbreviated Terms (continued)

Term	Meaning
UIMM	Unsigned immediate value
UISA	User instruction set architecture
UTLB	Unified translation lookaside buffer
UUT	Unit under test
VEA	Virtual environment architecture
WAR	Write-after-read
WAW	Write-after-write
WIMG	Write-through/caching-inhibited/memory-coherency enforced/guarded bits
XATC	Extended address transfer code
XER	Register used for indicating conditions such as carries and overflows for integer operations

Terminology Conventions

Table ii describes terminology conventions used in this manual and the equivalent terminology used in the PowerPC architecture specification.

Table ii. Terminology Conventions

The Architecture Specification	This Manual
Data storage interrupt (DSI)	DSI exception
Extended mnemonics	Simplified mnemonics
Fixed-point unit (FXU)	Integer unit (IU)
Instruction storage interrupt (ISI)	ISI exception
Interrupt	Exception
Privileged mode (or privileged state)	Supervisor-level privilege
Problem mode (or problem state)	User-level privilege
Real address	Physical address
Relocation	Translation
Storage (locations)	Memory
Storage (the act of)	Access
Store in	Write back
Store through	Write through

Table iii describes instruction field notation used in this manual.

Table iii. Instruction Field Conventions

The Architecture Specification	Equivalent to:
BA, BB, BT	crbA, crbB, crbD (respectively)
BF, BFA	crfD, crfS (respectively)
D	d
DS	ds
FLM	FM
FRA, FRB, FRC, FRT, FRS	frA, frB, frC, frD, frS (respectively)
FXM	CRM
RA, RB, RT, RS	rA, rB, rD, rS (respectively)
SI	SIMM
U	IMM
UI	UIMM
/, //, ///	0...0 (shaded)

Chapter 1 Overview

This chapter provides an overview of features for the MPC603e microprocessor and PowerPC architecture, and information about how the MPC603e implementation complies with the architectural definitions. Note that the MPC603e microprocessor is implemented in both a 2.5-volt version (PID 0007t MPC603e microprocessor, abbreviated as PID7t-603e) and a 3.3-volt version (PID 0006 MPC603e microprocessor, abbreviated as PID6-603e). Note that the PID6-603e is end-of-life and not recommended for new designs.

1.1 Overview

This section describes the details of the MPC603e, provides a block diagram showing the major functional units (see Figure 1-1), and briefly describes how these units interact. Any differences between the PID6-603e and PID7t-603e implementations are noted.

The MPC603e is a low-power implementation of this microprocessor family of reduced instruction set computing (RISC) microprocessors. The MPC603e implements the 32-bit portion of the PowerPC architecture, which defines 32-bit effective addresses, integer data types of 8, 16, and 32 bits, and floating-point data types of 32 and 64 bits.

The MPC603e is a superscalar processor that can issue and retire as many as three instructions per clock cycle. Instructions can execute out of program order for increased performance; however, the MPC603e makes completion appear sequential.

The MPC603e integrates five execution units—an integer unit (IU), a floating-point unit (FPU), a branch processing unit (BPU), a load/store unit (LSU), and a system register unit (SRU). The ability to execute five instructions in parallel and the use of simple instructions with rapid execution times yield high efficiency and throughput for MPC603e-based systems. Most integer instructions execute in one clock cycle. On the MPC603e, the FPU is pipelined so a single-precision multiply-add instruction can be issued and completed every clock cycle.

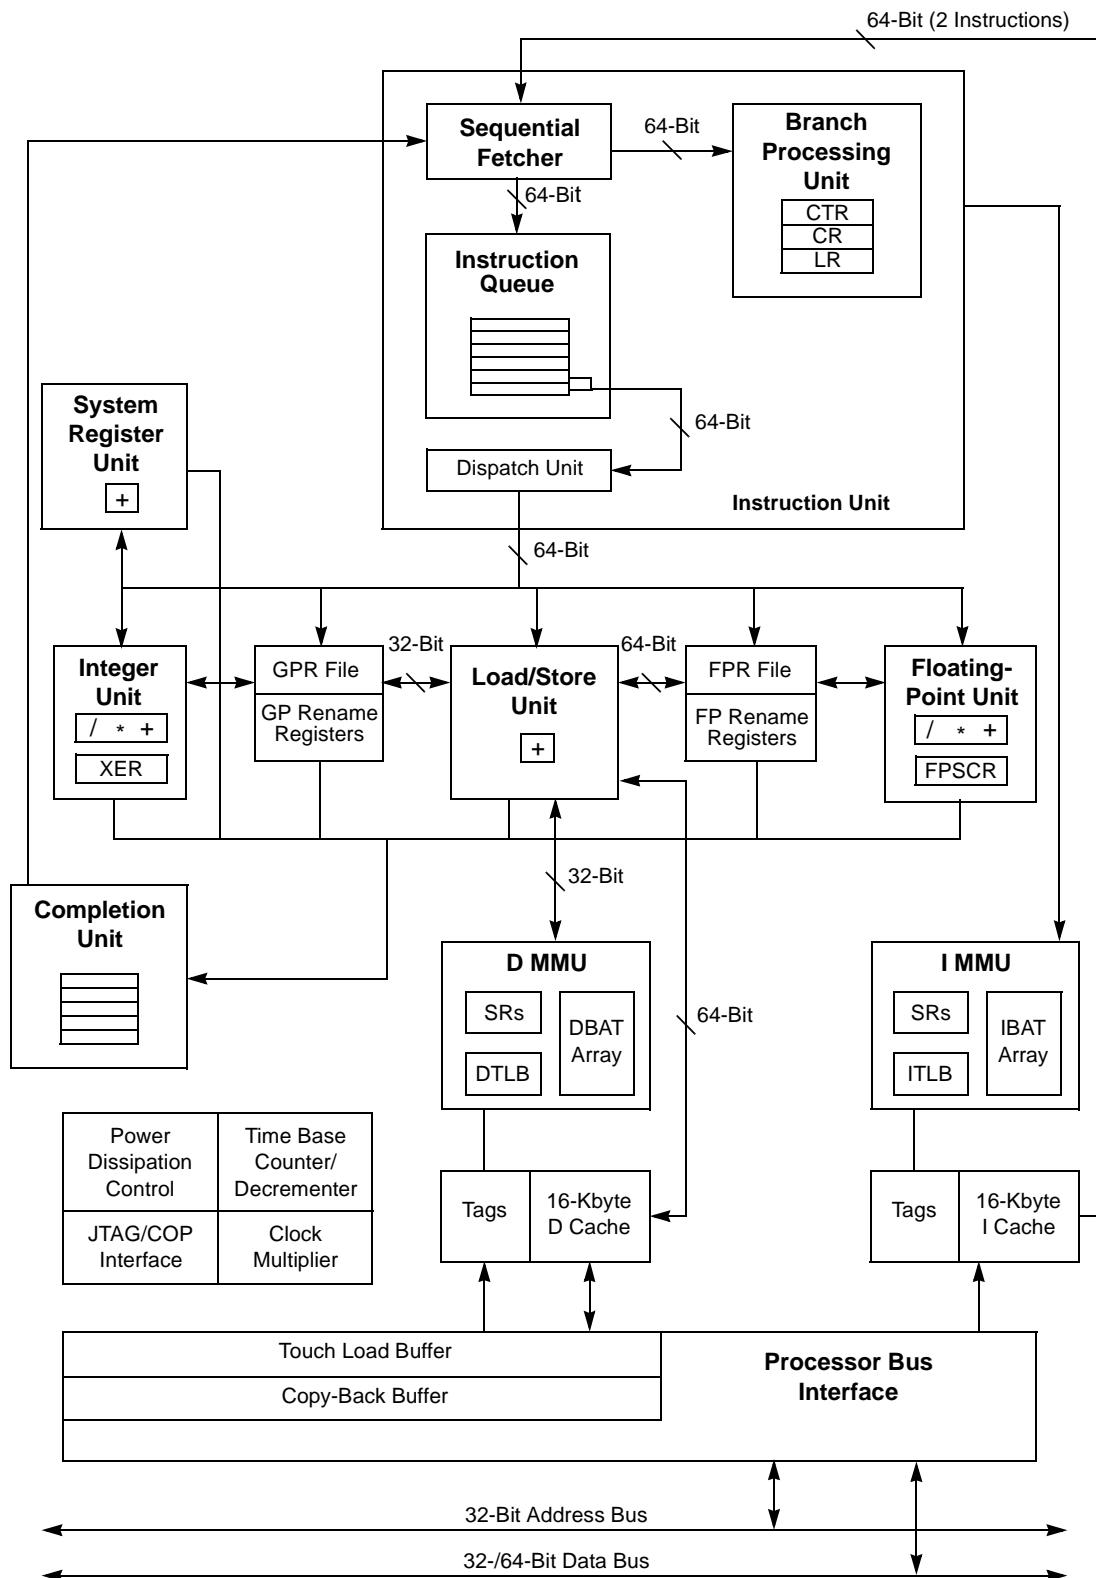


Figure 1-1. MPC603e Microprocessor Block Diagram

The MPC603e provides independent on-chip, 16-Kbyte, four-way set-associative, physically addressed caches for instructions and data, and on-chip instruction and data memory management units (MMUs). The MMUs contain 64-entry, two-way set-associative, data and instruction translation lookaside buffers (DTLB and ITLB) that provide support for demand-paged virtual-memory address translation and variable-sized block translation. The TLBs and caches use a least recently used (LRU) replacement algorithm. The MPC603e also supports block address translation through the use of two independent instruction and data block address translation (IBAT and DBAT) arrays of four entries each. Effective addresses are compared simultaneously with all four entries in the BAT array during block translation. In accordance with the PowerPC architecture, if an effective address hits in both the TLB and BAT array, the BAT translation takes priority.

The MPC603e has a selectable 32- or 64-bit data bus and a 32-bit address bus. The MPC603e interface protocol allows multiple masters to compete for system resources through a central external arbiter. The MPC603e provides a three-state (exclusive, modified, and invalid) coherency protocol which is a compatible subset of a four-state (modified/exclusive/shared/invalid) MESI protocol. This protocol operates coherently in systems that contain four-state caches. The MPC603e supports single-beat and burst data transfers for memory accesses and supports memory-mapped I/O operations.

The MPC603e is fabricated using an advanced CMOS process technology and is fully compatible with TTL devices.

1.1.1 Features

This section describes the major features of the MPC603e noting where the PID6-603e and PID7t-603e implementations differ:

- High-performance, superscalar microprocessor
 - As many as three instructions issued and retired per clock
 - As many as five instructions in execution per clock
 - Single-cycle execution for most instructions
 - Pipelined FPU for all single-precision and most double-precision operations
- Five independent execution units and two register files
 - BPU featuring static branch prediction
 - A 32-bit IU
 - Fully IEEE 754-compliant FPU for both single- and double-precision operations
 - LSU for data transfer between data cache and GPRs and FPRs
 - SRU that executes condition register (CR), special-purpose register (SPR), and integer add/compare instructions

- Thirty-two GPRs for integer operands
- Thirty-two FPRs for single- or double-precision operands
- High instruction and data throughput
 - Zero-cycle branch capability (branch folding)
 - Programmable static branch prediction on unresolved conditional branches
 - Instruction fetch unit capable of fetching two instructions per clock from the instruction cache
 - A six-entry instruction queue (IQ) that provides lookahead capability
 - Independent pipelines with feed-forwarding that reduces data dependencies in hardware
 - 16-Kbyte data cache and 16-Kbyte instruction cache—four-way set-associative, physically addressed, LRU replacement algorithm
 - Cache write-back or write-through operation programmable on a per page or per block basis
 - BPU that performs CR lookahead operations
 - Address translation facilities for 4-Kbyte page size, variable block size, and 256-Mbyte segment size
 - A 64-entry, two-way set-associative ITLB and DTLB
 - Four-entry data and instruction BAT arrays providing 128-Kbyte to 256-Mbyte blocks
 - Software table search operations and updates supported through fast trap mechanism
 - 52-bit virtual address; 32-bit physical address
- Facilities for enhanced system performance
 - A 32- or 64-bit split-transaction external data bus with burst transfers
 - Support for one-level address pipelining and out-of-order bus transactions
 - Hardware support for misaligned little-endian accesses (PID7t-603e)
- Integrated power management
 - Low-power 2.5-volt and 3.3-volt designs
 - Internal processor/bus clock multiplier ratios as follows:
 - 1/1, 1.5/1, 2/1, 2.5/1, 3/1, 3.5/1, and 4/1 (PID6-603e)
 - 2/1, 2.5/1, 3/1, 3.5/1, 4/1, 4.5/1, 5/1, 5.5/1, and 6/1 (PID7t-603e)
 - Three power-saving modes: doze, nap, and sleep
 - Automatic dynamic power reduction when internal functional units are idle
- In-system testability and debugging features through JTAG boundary-scan capability

Features specific to the PID7t-603e follow:

- Enhancements to the register set
 - The PID7t-603e adds two HID0 bits:
 - The address bus enable (ABE) bit, HID0[28], gives the PID7t-603e microprocessor the ability to broadcast **dcbf**, **dcbi**, and **dcbst** onto the 60x bus.
 - The instruction fetch enable M (IFEM) bit, HID0[24], allows the PID7t-603e to reflect the value of the M bit onto the 60x bus during instruction translation.
 - The Run_N counter register (Run_N) has been extended from 16 to 32 bits
- Enhancements to cache implementation
 - The instruction cache is blocked only until the critical load completes (hit under reloads allowed)
 - The critical double word is simultaneously written to the cache and forwarded to the requesting unit, thus minimizing stalls due to load delays.
 - Provides for an optional data cache operation broadcast feature (enabled by HID0[ABE]) that allows for correct system management using an external copy-back L2 cache.
 - All of the cache control instructions (**icbi**, **dcbi**, **dcbf**, and **dcbst**, excluding **dcbz**) require that HID0[ABE] be enabled in order to execute.
- Exceptions
 - The PID7t-603e now offers hardware support for misaligned little-endian accesses. Little-endian load/store accesses that are not on a word boundary, with the exception of strings and multiples, generate exceptions under the same circumstances as big-endian accesses.
 - The PID7t-603e removed misalignment support for **eciwx** and **ecowx** graphics instructions. These instructions cause an alignment exception if the access is not on a word boundary.
- Bus clock—New bus multipliers of 4.5x, 5x, 5.5x, and 6x that are selected by the unused encodings of PLL_CFG[0:3]. Bus multipliers of 1x and 1.5x are not supported by PID7t-603e.
- Power management—Internal voltage supply changed from 3.3 volts to 2.5 volts. The core logic of the chip now uses a 2.5-volt supply.
- Instruction timing
 - The integer divide instructions, **divwu[o][.]** and **divw[o][.]**, execute in 20 clock cycles; execution of these instructions in the PID6-603e takes 37 clock cycles.
 - Support for single-cycle store
 - An adder/comparator added to system register unit that allows dispatch and execution of multiple integer add and compare instructions on each cycle.

Figure 1-1 provides a block diagram of the MPC603e that shows how the execution units—IU, FPU, BPU, LSU, and SRU—operate independently and in parallel. Note that this is a conceptual diagram and does not attempt to show how these features are physically implemented on the chip.

The MPC603e provides address translation and protection facilities, including an ITLB, DTLB, and instruction and data BAT arrays. Instruction fetching and issuing is handled in the instruction unit. Translation of addresses for cache or external memory accesses are handled by the MMUs. Both units are discussed in more detail in Section 1.1.3, “Instruction Unit,” and Section 1.1.5.1, “Memory Management Units (MMUs).”

1.1.2 System Design and Programming Considerations

The MPC603e is built on the low-power dissipation, low cost, and the high performance attributes of the MPC603. It also provides the system designer additional capabilities through higher processor clock speeds (to 100 MHz), increases in cache size (16-Kbyte instruction and data caches) and set-associativity (four-way), and greater system clock flexibility. The following sections describe the differences between the MPC603 and MPC603e that affect the system designer and programmer already familiar with the operation of the MPC603.

The design enhancements to the MPC603e are described in the following sections as changes that can require a modification to the hardware or software configuration of a system designed for the MPC603.

1.1.2.1 Hardware Features

The following hardware features of the MPC603e may require modifications to MPC603 systems.

1.1.2.1.1 Replacement of XATS Signal by CSE1 Signal

The MPC603e employs four-way set-associativity for both the instruction and data caches, in place of the two-way set-associativity used in the MPC603. This change requires the use of an additional cache set entry (CSE1) signal to indicate which member of the cache set is being loaded during a cache line fill. CSE1 on the MPC603e is in the same pin location as XATS on the MPC603. Note that XATS is no longer needed by the MPC603e because support for access to direct-store segments has been removed.

Table 1-1 shows the CSE[0:1] signal encoding indicating the cache set element selected during a cache load operation.

Table 1-1. CSE[0:1] Signals

CSE[0:1]	Cache Set Element
00	Set 0
01	Set 1
10	Set 2
11	Set 3

1.1.2.1.2 Addition of Half-Clock Bus Multipliers

Some of the reserved clock configuration signal settings of the MPC603 are redefined to allow more flexible selection of higher internal and bus clock frequencies. The MPC603e provides programmable internal processor clock rates of 1x, 1.5x, 2x, 2.5x, 3x, 3.5x, and 4x multiples of the externally supplied clock frequency. For additional information, refer to the appropriate device-specific hardware specifications.

1.1.2.2 Software Features

The features of the MPC603e described in the following sections affect software originally written for the MPC603.

1.1.2.2.1 16-Kbyte Instruction and Data Caches

The MPC603e instruction and data caches are 16 Kbytes, twice the size of the MPC603 caches. The increase in cache size may require modification of cache flush routines. The increase is also reflected in four-way set-associativity of both caches in place of the two-way set-associativity in the MPC603.

1.1.2.2.2 Clock Configuration Available in HID1 Register

HID1[0–3] provides software read-only access to the configuration of the PLL_CFG signals. HID1 is not implemented in the MPC603.

1.1.2.2.3 Performance Enhancements

The following enhancements improve performance without requiring changes to software (other than compiler optimization) or hardware designed for the MPC603:

- Support for single-cycle store
- Addition of adder/comparator in the SRU allows dispatch and execution of multiple integer add and compare instructions on each cycle.
- Addition of SRR1[KEY] to provide information about memory protection violations prior to page table search operations. This bit is set when the combination of the settings in the appropriate SR[Kx] and in the MSR[PR] bit indicate that when

the PTE[PP] bits are either 00 or 01, a protection violation exists. If this is the case for a data write operation with a DTLB miss, the changed (C) bit in the page tables should not be updated (see Table 1-2). This reduces the time required to execute the page table search routine because the software no longer has to explicitly read both SR[Kx] and MSR[PR] to determine whether a protection violation exists before updating the C bit.

Table 1-2. Generated SRR1[KEY] Bit

Segment Register [Ks, Kp]	MSR[PR]	SRR1[KEY] Generated on DTLB Misses
0x	0	0
x0	1	0
1x	0	1
x1	1	1

Note: SRR1[KEY] indicates a protection violation if the PTE[pp] bits are 00 or 01.

1.1.3 Instruction Unit

As shown in Figure 1-1, the MPC603e instruction unit, containing a fetch unit, instruction queue, dispatch unit, and BPU, provides centralized control of instruction flow to the execution units. The instruction unit determines the address of the next instruction to be fetched based on information from the sequential fetcher and from the BPU.

The instruction unit fetches the instructions from the instruction cache into the instruction queue. The BPU receives branch instructions from the fetcher and uses static branch prediction to allow to fetching from a predicted instruction stream while a conditional branch is evaluated. The BPU folds out for unconditional branch instructions and conditional branch instructions unaffected by instructions in the execution pipeline.

Instructions issued beyond a predicted branch cannot complete execution until the branch is resolved, preserving the programming model of sequential execution. If any of these are branch instructions, they are decoded but not issued. Instructions to be executed by the FPU, IU, LSU, and SRU are issued and allowed to progress up to the register write-back stage. Write-back is allowed when a correctly predicted branch is resolved, and execution continues along the predicted path.

If branch prediction is incorrect, the instruction unit flushes all predicted path instructions, and instructions are issued from the correct path.

1.1.3.1 Instruction Queue and Dispatch Unit

The instruction queue (IQ), shown in Figure 1-1, holds as many as six instructions and loads up to two instructions from the instruction unit during a single cycle. The instruction fetch unit continuously loads as many instructions as space in the IQ allows. Instructions

are dispatched to their respective execution units from the dispatch unit at a maximum rate of two instructions per cycle. Dispatching is facilitated to the IU, FPU, LSU, and SRU by the provision of a reservation station at each unit. The dispatch unit performs source and destination register dependency checking, determines dispatch serializations, and inhibits subsequent instruction dispatching as required.

For a more detailed overview of instruction dispatch, see Section 1.3.6, “Instruction Timing.”

1.1.3.2 Branch Processing Unit (BPU)

The BPU receives branch instructions from the fetch unit and performs CR lookahead operations on conditional branches to resolve them early, achieving the effect of a zero-cycle branch in many cases.

The BPU uses a bit in the instruction encoding to predict the direction of the conditional branch. Therefore, when an unresolved conditional branch instruction is encountered, the MPC603e fetches instructions from the predicted target stream until the conditional branch is resolved.

The BPU contains an adder to compute branch target addresses and three user-control registers—the link register (LR), the count register (CTR), and the CR. The BPU calculates the return pointer for subroutine calls and saves it into the LR for certain types of branch instructions. The LR also contains the branch target address for the Branch Conditional to Link Register (**bclrx**) instruction. The CTR contains the branch target address for the Branch Conditional to Count Register (**bctrx**) instruction. The contents of the LR and CTR can be copied to or from any GPR. Because the BPU uses dedicated registers rather than GPRs or FPRs, execution of branch instructions is largely independent from execution of integer and floating-point instructions.

1.1.4 Independent Execution Units

The PowerPC architecture’s support for independent execution units allows implementation of processors with out-of-order instruction execution. For example, because branch instructions do not depend on GPRs or FPRs, branches can often be resolved early, eliminating stalls caused by taken branches.

The four other execution units and completion unit are described in the following sections.

1.1.4.1 Integer Unit (IU)

The IU executes all integer instructions. The IU executes one integer instruction at a time, performing computations with its arithmetic logic unit (ALU), multiplier, divider, and XER register. Most integer instructions are single-cycle instructions. The 32 GPRs hold integer operands. Stalls due to contention for GPRs are minimized by the automatic allocation of

rename registers. The MPC603e writes the contents of the rename registers to the appropriate GPR when integer instructions are retired by the completion unit.

1.1.4.2 Floating-Point Unit (FPU)

The FPU contains a single-precision multiply-add array and the floating-point status and control register (FPSCR). The multiply-add array allows the MPC603e to efficiently implement multiply and multiply-add operations. The FPU is pipelined so that single- and double-precision instructions can be issued back-to-back. The 32 FPRs are provided to support floating-point operations. Stalls due to contention for FPRs are minimized by the automatic allocation of rename registers. The MPC603e writes the contents of the rename registers to the appropriate FPR when floating-point instructions are retired by the completion unit.

The MPC603e supports all IEEE 754 floating-point data types (normalized, denormalized, NaN, zero, and infinity) in hardware, eliminating the latency incurred by software exception routines. (The term, ‘exception’ is also referred to as ‘interrupt’ in the architecture specification.)

1.1.4.3 Load/Store Unit (LSU)

The LSU executes all load and store instructions and provides the data transfer interface between the GPRs, FPRs, and the cache/memory subsystem. The LSU calculates effective addresses, performs data alignment, and provides sequencing for load/store string and multiple instructions.

Load and store instructions are issued and executed in program order; however, the memory accesses can occur out of order. Synchronizing instructions are provided to enforce strict ordering.

Cacheable loads, when free of data dependencies, can execute out of order with a maximum throughput of one per cycle and a two-cycle total latency. Data returned from the cache is held in a rename register until the completion logic commits the value to a GPR or FPR. Stores cannot be executed in a predicted manner and are held in the store queue until the completion logic signals that the store operation is to be completed to memory. The MPC603e executes store instructions with a maximum throughput of one per cycle and a three-cycle total latency. The time required to perform the actual load or store depends on whether the operation involves the cache, system memory, or an I/O device.

1.1.4.4 System Register Unit (SRU)

The SRU executes various system-level instructions, including condition register logical operations and move to/from special-purpose register instructions. It also executes integer add/compare instructions. In order to maintain system state, most instructions executed by the SRU are completion-serialized; that is, the instruction is held for execution in the SRU

until all prior instructions issued have completed. Results from completion-serialized instructions executed by the SRU are not available or forwarded for subsequent instructions until the instruction completes.

1.1.4.5 Completion Unit

The completion unit tracks instructions in program order from dispatch through execution and then completes. Completing an instruction commits the MPC603e to any architectural register changes caused by that instruction. In-order completion ensures the correct architectural state when the MPC603e must recover from a mispredicted branch or any exception.

Instruction state and other information required for completion is kept in a five-entry FIFO completion queue. A single completion queue entry is allocated for each instruction once it enters the execution unit from the dispatch unit. An available completion queue entry is a required resource for dispatch; if no completion entry is available, dispatch stalls. A maximum of two instructions per cycle are completed in order from the queue.

1.1.5 Memory Subsystem Support

The MPC603e provides both separate instruction and data caches and MMUs. The MPC603e also provides an efficient processor bus interface to facilitate access to main memory and other bus subsystems. The memory subsystem support functions are described in the following sections.

1.1.5.1 Memory Management Units (MMUs)

The MPC603e MMUs support up to 4 Petabytes (2^{52}) of virtual memory and 4 Gigabytes (2^{32}) of physical memory (referred to as real memory in the architecture specification) for instruction and data. The MMUs also control access privileges for these spaces on block and page granularities. Referenced and changed status is maintained by the processor for each page to assist implementation of a demand-paged virtual memory system. A key bit is implemented to provide information about memory protection violations prior to page table search operations.

The LSU calculates effective addresses (EAs) for data loads and stores, performs data alignment to and from cache memory, and provides the sequencing for load and store string and multiple word instructions. The instruction unit calculates effective addresses for instruction fetching.

After an EA is generated, its higher-order bits are translated by the appropriate MMU into physical address bits. The lower-order EA bits are the same on the physical address and form the index into the four-way set-associative tag array. After translating the address, the MMU passes the higher-order physical address bits to the cache and the cache lookup completes. For caching-inhibited accesses or accesses that miss in the cache, the

untranslated lower-order address bits are concatenated with the translated higher-order address bits; the resulting 32-bit physical address is then used by the memory unit and the system interface to access external memory.

The MMU also directs the address translation and enforces the protection hierarchy programmed by the operating system in relation to the supervisor/user privilege level of the access and in relation to whether the access is a load or store.

For instruction fetches, the IMMU looks for the address in the ITLB and in the IBAT array. If an address hits both, the IBAT array translation is used. Data accesses cause a lookup in the DTLB and DBAT array. In most cases, the translation is in a TLB and the physical address bits are readily available to the on-chip cache. The DBAT also is chosen if the translation is in both a DBAT and TLB.

When the EA misses in the TLBs, the MPC603e provides hardware assistance for software to perform a search of the translation tables in memory. The hardware assist consists of the following features:

- Automatic storage of the missed effective address in IMISS and DMISS
- Automatic generation of the primary and secondary hashed real address of the page table entry group (PTEG), which are readable from the HASH1 and HASH2 register locations.

The HASH data is generated from the contents of the IMISS or DMISS register. The register that is selected depends on the miss (instruction or data) that was last acknowledged.

- Automatic generation of the first word of the page table entry (PTE) of the tables being searched
- A real page address (RPA) register that matches the format of the lower word of the PTE
- TLB access instructions (**tbl1i** and **tbl1d**) that are used to load an address translation into the instruction or data TLBs
- Shadow registers for GPR0–GPR3 that allow miss code to execute without corrupting the state of any of the existing GPRs. Shadow registers are used only for servicing a TLB miss.

See Section 1.3.5.2, “Implementation-Specific Memory Management,” for more information about memory management for the MPC603e.

1.1.5.2 Cache Units

The MPC603e provides independent 16-Kbyte, four-way set-associative instruction and data caches. The cache block is 32 bytes long. The caches adhere to a write-back policy, but the MPC603e allows control of cacheability, write policy, and memory coherency at the page and block levels. The caches use an LRU replacement policy.

As shown in Figure 1-1, the caches provide a 64-bit interface to the instruction fetch unit and LSU. The surrounding logic selects, organizes, and forwards the requested information to the requesting unit. Write operations to the cache can be performed on a byte basis, and a complete read-modify-write operation to the cache can occur in each cycle.

The load/store and instruction fetch units provide the caches with the address of the data or instruction to be fetched. In the case of a cache hit, the cache returns two words to the requesting unit.

Because the data cache tags are single-ported, simultaneous load or store and snoop accesses cause resource contention. Snoop accesses have the highest priority and are given first access to the tags, unless the snoop access coincides with a tag write; in this case the snoop is retried and must re-arbitrate for cache access. Loads or stores deferred due to snoop accesses are performed on the clock cycle following the snoop.

1.1.6 Processor Bus Interface

Because the caches are on-chip, write-back caches, the most common transactions are burst-read memory operations, burst-write memory operations, and single-beat (noncacheable or write-through) memory read and write operations. There can also be address-only operations, variants of the burst and single-beat operations, (for example, global memory operations that are snooped and atomic memory operations), and address retry activity (for example, when a snooped read access hits a modified cache block).

Memory accesses can occur in single-beat (1–8 bytes) and four-beat burst (32 bytes) data transfers when the bus is configured as 64 bits, and in single-beat (1–4 bytes), two-beat (8 bytes), and eight-beat (32 bytes) data transfers when the bus is configured as 32 bits. The address and data buses operate independently to support pipelining and split transactions during memory accesses. The MPC603e can pipeline its own transactions to a depth of one level.

Access to the system interface is granted through an external arbitration mechanism that allows devices to compete for bus mastership. This arbitration is flexible, allowing the MPC603e to be integrated into systems that implement various fairness and bus parking procedures to avoid arbitration overhead.

Typically, memory accesses are weakly ordered—sequences of operations, including load/store string and multiple instructions, do not necessarily complete in the order they begin—maximizing the efficiency of the bus without sacrificing coherency of the data. The MPC603e allows read operations to precede store operations (except when a dependency exists, or in cases where a noncacheable access is performed), and provides support for a write operation to proceed a previously queued read data tenure (for example, allowing a snoop push to be enveloped by the address and data tenures of a read operation). Because the processor can dynamically optimize run-time ordering of load/store traffic, overall performance is improved.

1.1.7 System Support Functions

The MPC603e implements several support functions that include power management, time base/decrementer registers for system timing tasks, an IEEE 1149.1 (JTAG)/common on-chip processor (COP) test interface, and a phase-locked loop (PLL) clock multiplier. These system support functions are described in the following sections.

1.1.7.1 Power Management

The MPC603e provides four power modes, selectable by setting the appropriate control bits in the machine state register (MSR) and hardware implementation register 0 (HID0). The four power modes are as follows:

- Full-power—This is the default power state of the MPC603e. The MPC603e is fully powered and the internal functional units are operating at the full processor clock speed. If the dynamic power management mode is enabled, functional units that are idle will automatically enter a low-power state without affecting performance, software execution, or external hardware.
- Doze—All the functional units of the MPC603e are disabled except for the time base/decrementer registers and the bus snooping logic. When the processor is in doze mode, an external asynchronous interrupt, system management interrupt, decrementer exception, hard or soft reset, or machine check brings the MPC603e into the full-power state. The MPC603e in doze mode maintains the PLL in a fully powered state and locked to the system external clock input (SYSCLK) so a transition to the full-power state takes only a few processor clock cycles.
- Nap—The nap mode further reduces power consumption by disabling bus snooping, leaving only the time base register and the PLL in a powered state. The MPC603e returns to the full-power state upon receipt of an external asynchronous interrupt, system management interrupt, decrementer exception, hard or soft reset, or machine check input (\overline{MCP}) signal. A return to full-power state from a nap state takes only a few processor clock cycles.
- Sleep—Sleep mode reduces power consumption to a minimum by disabling all internal functional units; then external system logic may disable the PLL and SYSCLK. Returning the MPC603e to the full-power state requires the enabling of the PLL and SYSCLK, followed by the assertion of an external asynchronous interrupt, system management interrupt, hard or soft reset, or \overline{MCP} signal after the time required to relock the PLL.

The PID7t-603e implementation offers the following enhancements to the MPC603e family:

- Lower-power design
- 2.5-volt core and 3.3-volt I/O

1.1.7.2 Time Base/Decrementer

The time base is a 64-bit register (accessed as two 32-bit registers) that is incremented once every four bus clock cycles; external control of the time base is provided through the time base enable (TBEN) signal. The decrementer is a 32-bit register that generates a decrementer interrupt exception after a programmable delay. The contents of the decrementer register are decremented once every four bus clock cycles, and the decrementer exception is generated as the count passes through zero.

1.1.7.3 IEEE 1149.1 (JTAG)/COP Test Interface

The MPC603e provides IEEE 1149.1 and COP functions for facilitating board testing and chip debugging. The IEEE 1149.1 test interface provides a means for boundary-scan testing the MPC603e and the attached board. The COP function shares the IEEE 1149.1 test port, providing a means for executing test routines, and facilitating chip and software debugging.

1.1.7.4 Clock Multiplier

The internal clocking of the MPC603e is generated from and synchronized to the external clock signal, SYSCLK, by means of a voltage-controlled oscillator-based PLL. The PLL provides programmable internal processor clock rates of 1x, 1.5x, 2x, 2.5x, 3x, 3.5x, and 4x multiples of the externally supplied clock frequency. The bus clock is the same frequency and is synchronous with SYSCLK. The configuration of the PLL can be read by software from the hardware implementation register 1 (HID1).

1.2 PowerPC Architecture Implementation

The PowerPC architecture consists of the following layers, and adherence to the PowerPC architecture can be measured in terms of which of the following levels of the architecture is implemented:

- User instruction set architecture (UISA)—Defines the base user-level instruction set, user-level registers, data types, floating-point exception model, memory models for a uniprocessor environment, and programming model for a uniprocessor environment.
- Virtual environment architecture (VEA)—Describes the memory model for a multiprocessor environment, defines cache control instructions, and describes other aspects of virtual environments. Implementations that conform to the VEA also adhere to the UISA, but may not necessarily adhere to the OEA.
- Operating environment architecture (OEA)—Defines the memory management model, supervisor-level registers, synchronization requirements, and exception model. Implementations that conform to the OEA also adhere to the UISA and VEA.

The PowerPC architecture allows a wide range of designs for such features as cache and system interface implementations.

1.3 Implementation-Specific Information

The PowerPC architecture is derived from the IBM POWER architecture (Performance Optimized with Enhanced RISC architecture). The PowerPC architecture shares the benefits of the POWER architecture optimized for single-chip implementations. The PowerPC architecture design facilitates parallel instruction execution and is scaleable to take advantage of future technological gains.

This section describes the PowerPC architecture in general and specific details about the implementation of the MPC603e as a low-power, 32-bit member of this family. The main topics addressed are as follows:

- Section 1.3.1, “Programming Model,” describes the registers for the operating environment architecture common among processors that implement the PowerPC architecture and describes the programming model. It also describes the additional registers that are unique to the MPC603e.
- Section 1.3.2, “Instruction Set and Addressing Modes,” describes the PowerPC instruction set and addressing modes for the OEA, and defines and describes the PowerPC instructions implemented in the MPC603e.
- Section 1.3.3, “Cache Implementation,” describes the cache model that is defined generally for processors that implement the PowerPC architecture by the VEA. It also provides specific details about the MPC603e cache implementation.
- Section 1.3.4, “Exception Model,” describes the exception model of the OEA and the differences in the MPC603e exception model.
- Section 1.3.5, “Memory Management,” describes generally the conventions for memory management among these processors. This section also describes the MPC603e implementation of the 32-bit PowerPC memory management specification.
- Section 1.3.6, “Instruction Timing,” provides a general description of the instruction timing provided by the superscalar, parallel execution supported by the PowerPC architecture and the MPC603e.
- Section 1.3.7, “System Interface,” describes the signals implemented on the MPC603e.

The MPC603e is a high-performance, superscalar microprocessor. The PowerPC architecture allows optimizing compilers to schedule instructions to maximize performance through efficient use of the PowerPC instruction set and register model. The multiple, independent execution units allow compilers to optimize instruction throughput. Compilers that take advantage of the flexibility of the PowerPC architecture can additionally optimize system performance.

The following sections summarize the features of the MPC603e, including both those that are defined by the architecture and those that are unique to the various MPC603e implementations.

Specific features of the MPC603e are listed in Section 1.1.1, “Features.”

1.3.1 Programming Model

The PowerPC architecture defines register-to-register operations for most computational instructions. Source operands for these instructions are accessed from the registers or are provided as immediate values embedded in the instruction opcode. The three-register instruction format allows specification of a target register distinct from the two-source operands. Load and store instructions transfer data between registers and memory.

The MPC603e has two levels of privilege—supervisor mode of operation (typically used by the operating system) and user mode of operation (used by the application software). The programming models incorporate 32 GPRs, 32 FPRs, special-purpose registers (SPRs), and several miscellaneous registers. The MPC603e microprocessor also has its own unique set of hardware implementation (HID) registers.

Having access to privileged instructions, registers, and other resources allows the operating system to control the application environment (providing virtual memory and protecting operating system and critical machine resources). Instructions that control the state of the MPC603e, the address translation mechanism, and supervisor registers can be executed only when the processor is operating in supervisor mode.

Figure 1-2 shows all the MPC603e registers available at the user and supervisor level. The numbers to the right of the SPRs indicate the number that is used in the syntax of the instruction operands for the move to/from SPR instructions.

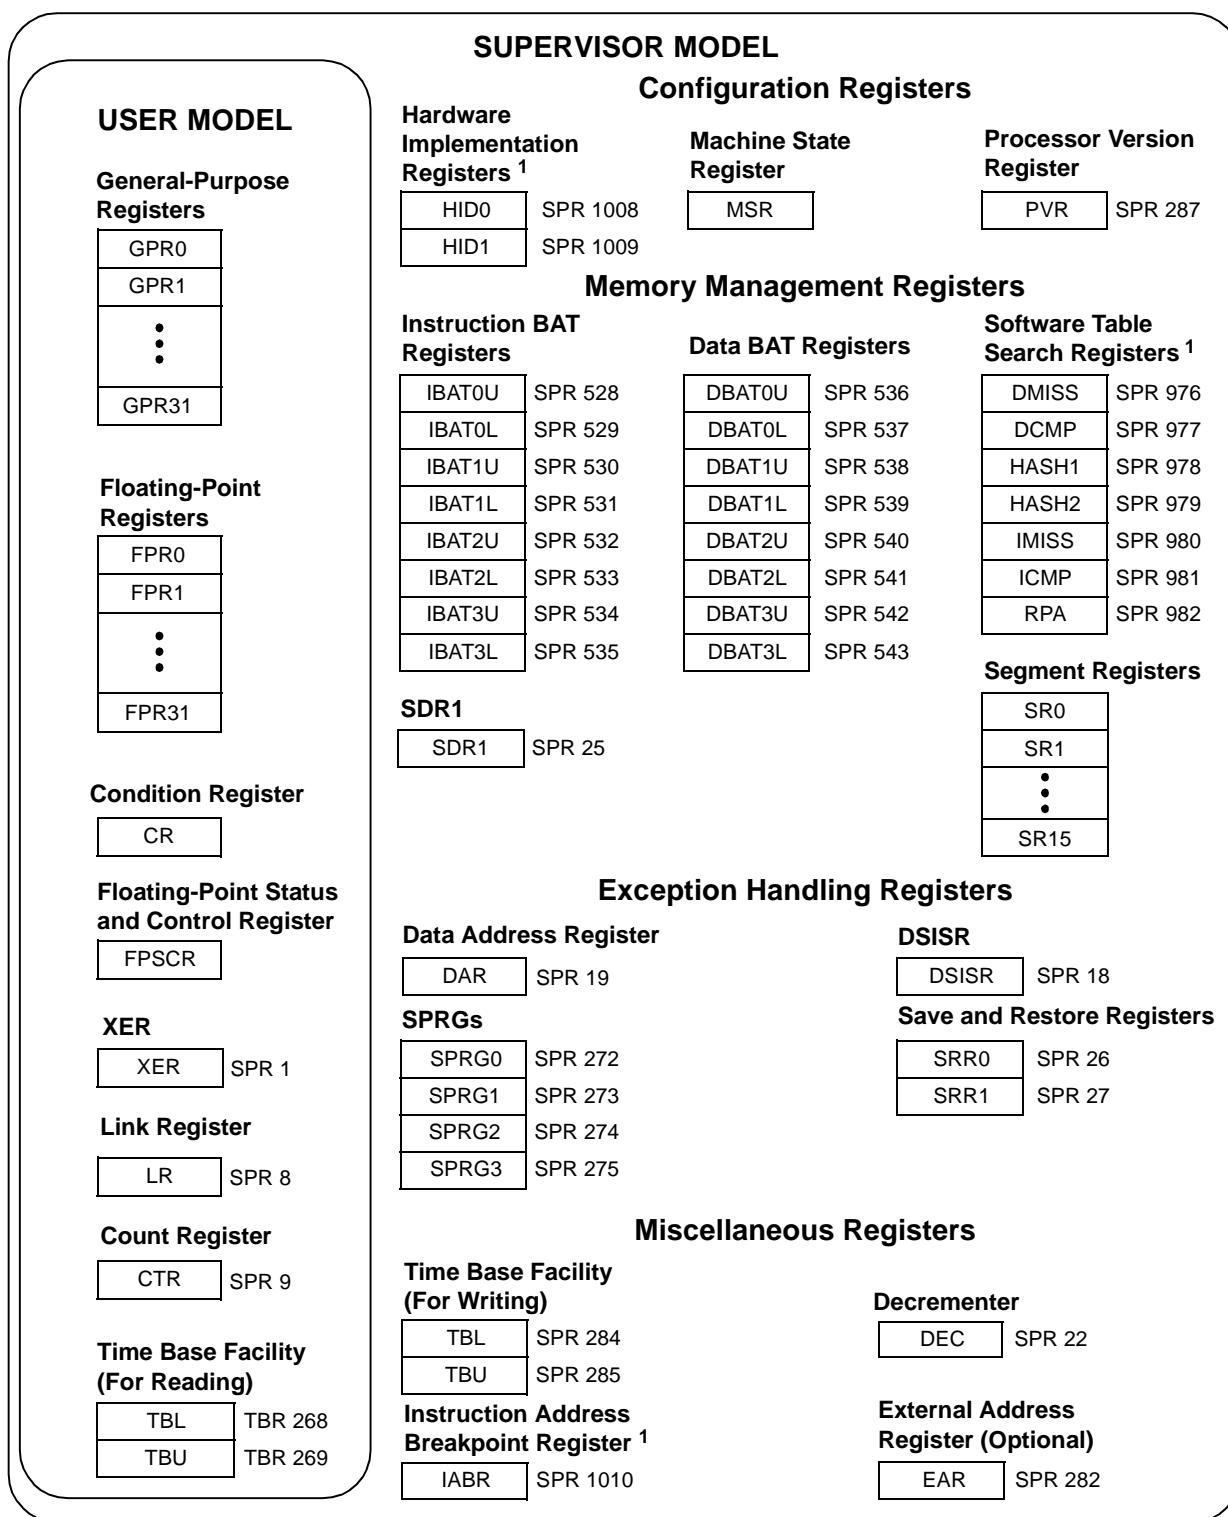
The following sections describe the PID7t-603e implementation-specific features as they apply to registers.

1.3.1.1 Processor Version Register (PVR)

The processor version number is 6 for the PID6-603e and 7 for the PID7t-603e. The processor revision level starts at 0x0100 and changes for each chip revision. The revision level is updated on all silicon revisions.

1.3.1.2 Hardware Implementation Register 0 (HID0)

PID7t-603e (designated by PVR level 0x0200) defines additional bits in the hardware implementation register 0 (HID0), a supervisor-level register that provides the means for enabling MPC603e checkstops and features, and allows software to read the configuration of the PLL configuration signals.



¹ These registers are MPC603e-specific (PID6-603e and PID7t-603e). They may not be supported by other processors of this family.

Figure 1-2. Programming Model—Registers

The HID0 bits with changed bit assignments are shown in Table 1-3. The HID0 bits that are not shown here are implemented as shown in Section 2.1.2.1, “Hardware Implementation Registers (HID0 and HID1).”

Table 1-3. Additional/Changed HID0 Bits

Bits	Description
24	Enable M bit on bus for instruction fetches (IFEM) (PID7t-603e only). 0 M bit disabled. Instruction fetches are treated as nonglobal on the bus. 1 Instruction fetches reflect the M bit from the WIM settings.
25–26	Reserved
28	Address broadcast enable. Controls whether certain address-only operations (such as cache operations) are broadcast on the 60x bus. 0 Address-only operations affect only local caches and are not broadcast. 1 Address-only operations are broadcast on the 60x bus. Affected instructions are dcbi , dcbf , and dcbst . Note that these cache control instruction broadcasts are not snooped by the PID7t-603e. Refer to Section 1.3.3, “Cache Implementation.”
29–30	Reserved

1.3.1.3 General-Purpose Registers (GPRs)

The PowerPC architecture defines 32 user-level GPRs. These registers are either 32 bits wide in 32-bit microprocessors or 64 bits wide in 64-bit microprocessors. The GPRs serve as the data source or destination for all integer instructions.

1.3.1.4 Floating-Point Registers (FPRs)

The PowerPC architecture also defines 32 user-level, 64-bit FPRs. The FPRs serve as the data source or destination for floating-point instructions. These registers can contain data objects of either single- or double-precision floating-point formats.

1.3.1.5 Condition Register (CR)

The CR is a 32-bit user-level register that provides a mechanism for testing and branching. It consists of eight 4-bit fields that reflect the results of certain operations, such as move, integer and floating-point comparisons, arithmetic, and logical operations.

1.3.1.6 Floating-Point Status and Control Register (FPSCR)

The user-level FPSCR contains all floating-point exception signal bits, exception summary bits, exception enable bits, and rounding control bits needed for compliance with the IEEE 754 standard.

1.3.1.7 Machine State Register (MSR)

The MSR is a supervisor-level register that defines the state of the processor. The contents of this register are saved when an exception is taken and restored when the exception handling completes. The MPC603e implements the MSR as a 32-bit register.

1.3.1.8 Segment Registers (SRs)

For memory management, 32-bit processors implement sixteen 32-bit SRs. To speed access, the MPC603e implements the SRs as two arrays; a main array (for data memory accesses) and a shadow array (for instruction memory accesses). Loading a segment entry with the Move to Segment Register (**mtsr**) instruction loads both arrays.

1.3.1.9 Special-Purpose Registers (SPRs)

The OEA defines numerous SPRs that serve a variety of functions, such as providing controls, indicating status, configuring the processor, and performing special operations. During normal execution, a program can access the registers, as shown in Figure 1-2, depending on the program's access privilege (supervisor or user, determined by the privilege-level bit, MSR[PR]). Note that GPRs and FPRs are accessed through operands that are part of the instructions. Access to registers can be explicit (that is, through the use of specific instructions for that purpose such as Move to Special-Purpose Register (**mtspr**) and Move from Special-Purpose Register (**mfspr**) instructions) or implicit, as the part of the execution of an instruction. Some registers are accessed both explicitly and implicitly.

In the MPC603e, all SPRs are 32 bits wide.

1.3.1.9.1 User-Level SPRs

The following MPC603e SPRs are accessible by user-level software:

- Link register (LR)—The LR can be used to provide the branch target address and to hold the return address after branch and link instructions. The LR is 32 bits wide in 32-bit implementations.
- Count register (CTR)—The CTR is decremented and tested automatically as a result of branch-and-count instructions. The CTR is 32 bits wide in 32-bit implementations.
- XER register—The 32-bit XER contains the summary overflow bit, integer carry bit, overflow bit, and a field specifying the number of bytes to be transferred by a Load String Word Indexed (**lswx**) or Store String Word Indexed (**stswx**) instruction.

1.3.1.9.2 Supervisor-Level SPRs

The MPC603e also contains SPRs that can be accessed only by supervisor-level software. These registers consist of the following:

- The DSISR defines the cause of data access and alignment exceptions.
- The data address register (DAR) holds the address of an access after an alignment or DSI exception.
- Decrementer register (DEC) is a 32-bit decrementing counter that provides a mechanism for causing a decrementer exception after a programmable delay.
- SDR1 specifies the page table format used in virtual-to-physical address translation for pages. (Note that physical address is referred to as real address in the architecture specification.)
- The machine status save/restore register 0 (SRR0) is used for saving the address of the instruction that caused the exception, and the address to return to when a Return from Interrupt (**rfi**) instruction is executed.
- The machine status save/restore register 1 (SRR1) is used to save machine status on exceptions and to restore machine status when an **rfi** instruction is executed.
- The SPRG0–SPRG3 registers are provided for operating system use.
- The external access register (EAR) controls access to the external control facility through the External Control In Word Indexed (**eciwx**) and External Control Out Word Indexed (**ecowx**) instructions.
- The time base register (TB) is a 64-bit register that maintains the time of day and operates interval timers. It consists of two 32-bit fields—time base upper (TBU) and time base lower (TBL).
- The processor version register (PVR) is a read-only register that identifies the version (model) and revision level of the processor.
- Block address translation (BAT) arrays—The PowerPC architecture defines 16 BAT registers—four pairs of data BATs (DBATs) and four pairs of instruction BATs (IBATs). See Figure 1-2 for a list of the SPR numbers for the BAT arrays.

The following supervisor-level SPRs are implementation-specific (not defined in the PowerPC architecture):

- DMISS and IMISS are read-only registers that are loaded automatically on an instruction or data TLB miss.
- HASH1 and HASH2 contain the physical addresses of the primary and secondary page table entry groups (PTEGs).
- ICMP and DCMP contain a duplicate of the first word in the page table entry (PTE) for which the table search is looking.
- The required physical address (RPA) register is loaded by the processor with the second word of the correct PTE during a page table search.

- The hardware implementation (HID0 and HID1) registers provide the means for enabling MPC603e checkstops and features, and allows software to read the configuration of the PLL configuration signals.
- The instruction address breakpoint register (IABR) is loaded with an instruction address that is compared to instruction addresses in the dispatch queue. When an address match occurs, an instruction address breakpoint exception is generated.

Figure 1-2 shows all the MPC603e registers available at the user and supervisor level. The numbers to the right of the SPRs indicate the number that is used in the syntax of the instruction operands to access the register.

1.3.2 Instruction Set and Addressing Modes

The following sections describe the PowerPC instruction set and addressing modes in general.

1.3.2.1 PowerPC Instruction Set and Addressing Modes

All PowerPC instructions are encoded as single-word (32-bit) opcodes. Instruction formats are consistent among all instruction types, permitting efficient decoding to occur in parallel with operand accesses. This fixed instruction length and consistent format greatly simplifies instruction pipelining.

The PowerPC instructions are divided into the following categories:

- Integer instructions—These include computational and logical instructions.
 - Integer arithmetic instructions
 - Integer compare instructions
 - Integer logical instructions
 - Integer rotate and shift instructions
- Floating-point instructions—These include floating-point computational instructions, as well as instructions that affect the FPSCR.
 - Floating-point arithmetic instructions
 - Floating-point multiply/add instructions
 - Floating-point rounding and conversion instructions
 - Floating-point compare instructions
 - Floating-point status and control instructions
- Load/store instructions—These include integer and floating-point load and store instructions.
 - Integer load and store instructions
 - Integer load and store multiple instructions

- Floating-point load and store
- Primitives used to construct atomic memory operations (**lwarx** and **stwcx.** instructions)
- Flow control instructions—These include branching instructions, condition register logical instructions, trap instructions, and other instructions that affect the instruction flow.
 - Branch and trap instructions
 - Condition register logical instructions
- Processor control instructions—These instructions are used for synchronizing memory accesses and management of caches, TLBs, and the segment registers.
 - Move to/from SPR instructions
 - Move to/from MSR
 - Synchronize
 - Instruction synchronize
- Memory control instructions—These instructions provide control of caches, TLBs, and segment registers.
 - Supervisor-level cache management instructions
 - User-level cache instructions
 - Segment register manipulation instructions
- Translation lookaside buffer management instructions

Note that this grouping of instructions does not indicate the execution unit that executes a particular instruction or group of instructions.

Integer instructions operate on byte, half-word, and word operands. Floating-point instructions operate on single-precision (one word) and double-precision (one double word) floating-point operands. The PowerPC architecture uses instructions that are 4 bytes long and word-aligned. It provides for byte, half-word, and word operand loads and stores between memory and a set of 32 GPRs. It also provides for word and double-word operand loads and stores between memory and a set of 32 FPRs.

Computational instructions do not modify memory. To use a memory operand in a computation and then modify the same or another memory location, the memory contents must be loaded into a register, modified, and then written back to the target location with distinct instructions.

The MPC603e follows the program flow when it is in the normal execution state. However, the flow of instructions can be interrupted directly by the execution of an instruction or by an asynchronous event. Either kind of exception may cause one of several components of the system software to be invoked.

1.3.2.2 Implementation-Specific Instruction Set

The MPC603e instruction set is defined as follows:

- The MPC603e provides hardware support for all 32-bit PowerPC instructions.
- The MPC603e provides two implementation-specific instructions used for software table search operations following TLB misses:
 - Load Data TLB Entry (**tlbld**)
 - Load Instruction TLB Entry (**tlbli**)
- The MPC603e implements the following instructions which are defined as optional by the PowerPC architecture:
 - External Control In Word Indexed (**eciwx**)
 - External Control Out Word Indexed (**ecowx**)
 - Floating Select (**fsel**)
 - Floating Reciprocal Estimate Single-Precision (**fres**)
 - Floating Reciprocal Square Root Estimate (**frsq rte**)
 - Store Floating-Point as Integer Word (**stfiwx**)

1.3.3 Cache Implementation

The following sections describe the general cache characteristics as implemented in the PowerPC architecture, and the MPC603e implementation, specifically. PID7t-603e specific information is noted where applicable.

1.3.3.1 PowerPC Cache Characteristics

The PowerPC architecture does not define hardware aspects of cache implementations. The MPC603e controls the following memory access modes on a page or block basis:

- Write-back/write-through mode
- Caching-inhibited mode
- Memory coherency

Note that in the MPC603e, a cache block is defined as eight words. The VEA defines cache management instructions that provide a means by which the application programmer can affect the cache contents.

1.3.3.2 Implementation-Specific Cache Implementation

The MPC603e has two 16-Kbyte, four-way set-associative (instruction and data) caches. The caches are physically addressed, and the data cache can operate in either write-back or write-through mode as specified by the PowerPC architecture.

The data cache is configured as 128 sets of four blocks each. Each block consists of 32 bytes, 2 state bits, and an address tag. The 2 state bits implement the three-state MEI (modified/exclusive/invalid) protocol. Each block contains eight 32-bit words. Note that the PowerPC architecture defines the term ‘block’ as the cacheable unit. For the MPC603e, the block size is equivalent to a cache line. A block diagram of the data cache organization is shown in Figure 1-3.

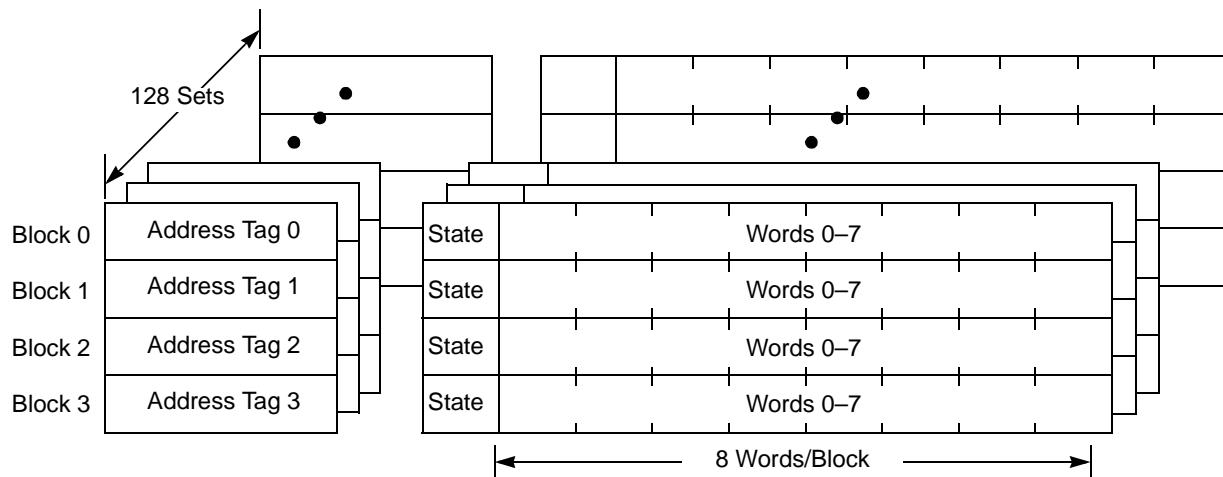


Figure 1-3. Data Cache Organization

The instruction cache also consists of 128 sets of four blocks, and each block consists of 32 bytes, an address tag, and a valid bit. The instruction cache may not be written to, except through a block fill operation. In the PID7t-603e, the instruction cache is blocked only until the critical load completes. The PID7t-603e supports instruction fetching from other instruction cache lines following the forwarding of the critical first double word of a cache line load operation. Successive instruction fetches from the cache line being loaded are forwarded, and accesses to other instruction cache lines can proceed during the cache line load operation. The instruction cache is not snooped, and cache coherency must be maintained by software. A fast hardware invalidation capability is provided to support cache maintenance. The organization of the instruction cache is very similar to the data cache shown in Figure 1-3.

Each cache block contains eight contiguous words from memory that are loaded from an 8-word boundary (that is, bits A[27–31] of the effective addresses are zero); thus, a cache block never crosses a page boundary. Misaligned accesses across a page boundary can incur a performance penalty.

The MPC603e cache blocks are loaded in four beats of 64 bits each when the MPC603e is configured with a 64-bit data bus. When the MPC603e is configured with a 32-bit bus, cache block loads are performed with eight beats of 32 bits each. The burst load is performed as critical-double-word-first. The data cache is blocked to internal accesses until the load completes; the instruction cache allows sequential fetching during a cache block

load. In the PID7t-603e, the critical-double-word is simultaneously written to the cache and forwarded to the requesting unit, thus minimizing stalls due to load delays.

To ensure coherency among caches in a multiprocessor (or multiple caching-device) implementation, the MPC603e implements the MEI protocol. The following three states indicate the state of the cache block:

- Modified—The cache block is modified with respect to system memory; that is, data for this address is valid only in the cache and not in system memory.
- Exclusive—This cache block holds valid data that is identical to the data at this address in system memory. No other cache has this data.
- Invalid—This cache block does not hold valid data.

Cache coherency is enforced by on-chip bus snooping logic. Because the MPC603e data cache tags are single-ported, a simultaneous load or store and snoop access represents a resource contention. The snoop access is given first access to the tags. The load or store then occurs on the clock following the snoop.

1.3.4 Exception Model

This section describes the PowerPC exception model and the MPC603e implementation, specifically. PID7t-603e-specific information is noted where applicable.

1.3.4.1 PowerPC Exception Model

The PowerPC exception mechanism allows the processor to change to supervisor state as a result of external signals, errors, or unusual conditions arising in the execution of instructions, and differs from the arithmetic exceptions defined by the IEEE for floating-point operations. When exceptions occur, information about the state of the processor is saved to certain registers and the processor begins execution at an address (exception vector) predetermined for each exception type. Processing of exceptions occurs in supervisor mode.

Although multiple exception conditions can map to a single exception vector, a more specific condition may be determined by examining a register associated with the exception—for example, the DSISR and the FPSCR. Additionally, some exception conditions can be explicitly enabled or disabled by software.

The PowerPC architecture requires that exceptions be handled in program order; therefore, although a particular implementation may recognize exception conditions out of order, they are presented strictly in order. When an instruction-caused exception is recognized, any unexecuted instructions that appear earlier in the instruction stream, including any that have not yet entered the execute stage, are required to complete before the exception is taken. Any exceptions caused by those instructions are handled first. Likewise, exceptions that are asynchronous and precise are recognized when they occur, but are not handled until the

instruction currently in the completion stage successfully completes execution or generates an exception, and the completed store queue is emptied.

Unless a catastrophic condition causes a system reset or machine check exception, only one exception is handled at a time. If, for example, a single instruction encounters multiple exception conditions, those conditions are handled sequentially. After the exception handler handles an exception, the instruction execution continues until the next exception condition is encountered. However, in many cases there is no attempt to re-execute the instruction. This method of recognizing and handling exception conditions sequentially guarantees that exceptions are recoverable.

Exception handlers should save the information stored in SRR0 and SRR1 early to prevent the program state from being lost due to a system reset or machine check exception or to an instruction-caused exception in the exception handler, and before enabling external interrupts.

The PowerPC architecture supports four types of exceptions:

- Synchronous, precise—These are caused by instructions. All instruction-caused exceptions are handled precisely; that is, the machine state at the time the exception occurs is known and can be completely restored. This means that (excluding the trap and system call exceptions) the address of the faulting instruction is provided to the exception handler and neither the faulting instruction nor subsequent instructions in the code stream will complete execution before the exception is taken. Once the exception is processed, execution resumes at the address of the faulting instruction (or at an alternate address provided by the exception handler). When an exception is taken due to a trap or system call instruction, execution resumes at an address provided by the handler.
- Synchronous, imprecise—The PowerPC architecture defines two imprecise floating-point exception modes, recoverable and nonrecoverable. Even though the MPC603e provides a means to enable the imprecise modes, it implements these modes identically to the precise mode (that is, all enabled floating-point enabled exceptions are always precise on the MPC603e).
- Asynchronous, maskable—The external, system management interrupt (SMI), and decrementer interrupts are maskable asynchronous exceptions. When these exceptions occur, their handling is postponed until the next instruction, and any exceptions associated with that instruction, completes execution. If there are no instructions in the execution units, the exception is taken immediately on determination of the correct restart address (for loading SRR0).
- Asynchronous, nonmaskable—There are two nonmaskable asynchronous exceptions: system reset and the machine check exception. These exceptions may not be recoverable, or may provide a limited degree of recoverability. All exceptions report recoverability through MSR[RI].

1.3.4.2 Implementation-Specific Exception Model

As specified by the PowerPC architecture, all MPC603e exceptions can be described as either precise or imprecise and either synchronous or asynchronous. Asynchronous exceptions (some of which are maskable) are caused by events external to the processor's execution; synchronous exceptions, which are all handled precisely by the MPC603e, are caused by instructions. The MPC603e exception classes are shown in Table 1-4.

Table 1-4. Exception Classifications

Synchronous/Asynchronous	Precise/Imprecise	Exception Type
Asynchronous, nonmaskable	Imprecise	Machine check System reset
Asynchronous, maskable	Precise	External interrupt Decrementer System management interrupt
Synchronous	Precise	Instruction-caused exceptions

Although exceptions have other characteristics as well, such as whether they are maskable or nonmaskable, the distinctions shown in Table 1-4 define categories of exceptions that the MPC603e handles uniquely. Note that Table 1-4 includes no synchronous imprecise instructions. While the PowerPC architecture supports imprecise handling of floating-point exceptions, the MPC603e implements floating-point exception modes as precise exceptions.

The MPC603e exceptions, and conditions that cause them, are listed in Table 1-5.

Table 1-5. Exceptions and Conditions

Exception Type	Vector Offset (hex)	Causing Conditions
Reserved	00000	—
System reset	00100	A system reset is caused by the assertion of either <u>SRESET</u> or <u>HRESET</u> .
Machine check	00200	A machine <u>check</u> is caused by the assertion of <u>TEA</u> during a data bus transaction, assertion of <u>MCP</u> , or an address or data parity error.
DSI	00300	The cause of a DSI exception can be determined by the bit settings in the DSISR, listed as follows: 1 Set if the translation of an attempted access is not found in the primary hash table entry group (HTEG), or in the rehashed secondary HTEG, or in the range of a DBAT register; otherwise cleared. 4 Set if a memory access is not permitted by the page or DBAT protection mechanism; otherwise cleared. 5 Set by an eciwx or ecowx instruction if the access is to an address that is marked as write-through, or execution of a load/store instruction that accesses a direct-store segment. 6 Set for a store operation and cleared for a load operation. 11 Set if eciwx or ecowx is used and EAR[E] is cleared.

Table 1-5. Exceptions and Conditions (continued)

Exception Type	Vector Offset (hex)	Causing Conditions
ISI	00400	An ISI exception is caused when an instruction fetch cannot be performed for any of the following reasons: <ul style="list-style-type: none"> The effective (logical) address cannot be translated. That is, there is a page fault for this portion of the translation, so an ISI exception must be taken to load the PTE (and possibly the page) into memory. The fetch access is to a direct-store segment (indicated by SRR1[3] set). The fetch access violates memory protection (indicated by SRR1[4] set). If the key bits (Ks and Kp) in the segment register and the PP bits in the PTE are set to prohibit read access, instructions cannot be fetched from this location.
External interrupt	00500	An external interrupt is caused when MSR[EE] = 1 and the INT signal is asserted.
Alignment	00600	An alignment exception is caused when the MPC603e cannot perform a memory access for any of the reasons described below: <ul style="list-style-type: none"> The operand of a floating-point load or store instruction is not word-aligned. The operand of lmw, stmw, lwarx, and stwcx. instructions are not aligned. The operand of a single-register load or store operation is not aligned, and the MPC603e is in little-endian mode (PID6-603e only). The execution of a floating-point load or store instruction to a direct-store segment. The operand of a load, store, load multiple, store multiple, load string, or store string instruction crosses a segment boundary into a direct-store segment, or crosses a protection boundary. Execution of a misaligned eciwx or ecowx instruction (PID7t-603e only). The instruction is lmw, stmw, lswi, lswx, stswi, stswx, and the MPC603e is in little-endian mode. The operand of dcbz is in memory that is write-through-required or caching-inhibited.
Program	00700	A program exception is caused by one of the following exception conditions, which correspond to bit settings in SRR1 and arise during execution of an instruction: <ul style="list-style-type: none"> Floating-point enabled exception—A floating-point enabled exception condition is generated when the following condition is met: $(MSR[FE0] \mid MSR[FE1]) \& FPSCR[FEX]$ is 1. FPSCR[FEX] is set by the execution of a floating-point instruction that causes an enabled exception or by the execution of one of the ‘move to FPSCR’ instructions that results in both an exception condition bit and its corresponding enable bit being set in the FPSCR. Illegal instruction—An illegal instruction program exception is generated when execution of an instruction is attempted with an illegal opcode or illegal combination of opcode and extended opcode fields (including PowerPC instructions not implemented in the MPC603e), or when execution of an optional instruction not provided in the MPC603e is attempted (these do not include those optional instructions that are treated as no-ops). Privileged instruction—A privileged instruction type program exception is generated when the execution of a privileged instruction is attempted and the MSR register user privilege bit, MSR[PR], is set. In the MPC603e, this exception is generated for mtspr or mfspr with an invalid SPR field if SPR[0] = 1 and MSR[PR] = 1. This may not be true for all processors. Trap—A trap type program exception is generated when any of the conditions specified in a trap instruction is met.

Chapter 1. Overview

Table 1-5. Exceptions and Conditions (continued)

Exception Type	Vector Offset (hex)	Causing Conditions
Floating-point unavailable	00800	A floating-point unavailable exception is caused by an attempt to execute a floating-point instruction (including floating-point load, store, and move instructions) when the floating-point available bit is disabled (MSR[FP] = 0).
Decrementer	00900	The decrementer exception occurs when DEC[31] changes from 0 to 1. Must also be enabled with MSR[EE].
Reserved	00A00–00BFF	—
System call	00C00	A system call exception occurs when a System Call (sc) instruction is executed.
Trace	00D00	A trace exception is taken when MSR[SE] =1 or when the currently completing instruction is a branch and MSR[BE] =1.
Reserved	00E00	The MPC603e does not generate an exception to this vector. Other processors may use this vector for floating-point assist exceptions.
Reserved	00E10–00FFF	—
Instruction translation miss	01000	An instruction translation miss exception is caused when an effective address for an instruction fetch cannot be translated by the ITLB.
Data load translation miss	01100	A data load translation miss exception is caused when an effective address for a data load operation cannot be translated by the DTLB.
Data store translation miss	01200	A data store translation miss exception is caused when an effective address for a data store operation cannot be translated by the DTLB, or where a DTLB hit occurs, and the change bit in the PTE must be set due to a data store operation.
Instruction address breakpoint	01300	An instruction address breakpoint exception occurs when the address (bits 0–29) in the IABR matches the next instruction to complete in the completion unit, and IABR[30] is set.
System management interrupt	01400	A system management interrupt is caused when MSR[EE] = 1 and the SMI input signal is asserted.
Reserved	01500–02FFF	—

1.3.5 Memory Management

The following sections describe the memory management features of the PowerPC architecture and the MPC603e implementation, respectively.

1.3.5.1 PowerPC Memory Management

The primary functions of the MMU are to translate logical (effective) addresses to physical addresses for memory accesses, and to provide access protection on blocks and pages of memory.

There are two types of accesses generated by the MPC603e that require address translation—instruction accesses and data accesses to memory generated by load and store instructions.

The PowerPC MMU and exception model support demand-paged virtual memory. Virtual memory management permits execution of programs larger than the size of physical memory; demand-paged implies that individual pages are loaded into physical memory from system memory only when they are first accessed by an executing program.

The hashed page table is a variable-sized data structure that defines the mapping between virtual page numbers and physical page numbers. The page table size is a power of two, and its starting address is a multiple of its size.

The page table contains a number of page table entry groups (PTEGs). A PTEG contains 8 page table entries (PTEs) of 8 bytes each; therefore, each PTEG is 64 bytes long. PTEG addresses are entry points for table search operations.

Address translations are enabled by setting bits in the MSR—MSR[IR] enables instruction address translations and MSR[DR] enables data address translations.

1.3.5.2 Implementation-Specific Memory Management

The instruction and data memory management units in the MPC603e provide 4 Gbytes of logical address space accessible to supervisor and user programs with a 4-Kbyte page size and 256-Mbyte segment size. Block sizes range from 128 Kbytes to 256 Mbytes and are software selectable. In addition, the MPC603e uses an interim 52-bit virtual address and hashed page tables for generating 32-bit physical addresses. The MMUs in the MPC603e rely on the exception processing mechanism for the implementation of the paged virtual memory environment and for enforcing protection of designated memory areas.

Instruction and data TLBs provide address translation in parallel with the on-chip cache access, incurring no additional time penalty in the event of a TLB hit. A TLB is a cache of the most recently used page table entries. Software is responsible for maintaining the consistency of the TLB with memory. The MPC603e TLBs are 64-entry, two-way set-associative caches that contain instruction and data address translations. The MPC603e provides hardware assist for software table search operations through the hashed page table on TLB misses. Supervisor software can invalidate TLB entries selectively.

The MPC603e also provides independent four-entry BAT arrays for instructions and data that maintain address translations for blocks of memory. These entries define blocks that can vary from 128 Kbytes to 256 Mbytes. The BAT arrays are maintained by system software.

As specified by the PowerPC architecture, the hashed page table is a variable-sized data structure that defines the mapping between virtual page numbers and physical page numbers. The page table size is a power of two, and its starting address is a multiple of its size.

Also as specified by the PowerPC architecture, the page table contains a number of PTEGs. A PTEG contains 8 PTEs of 8 bytes each; therefore, each PTEG is 64 bytes long. PTEG addresses are entry points for table search operations.

1.3.6 Instruction Timing

The MPC603e is a pipelined superscalar processor. The processing of an instruction is reduced into discrete stages by a pipelined processor. Because the processing of an instruction is broken into a series of stages, an instruction does not require the entire resources of an execution unit. For example, after an instruction completes the decode stage, it can pass on to the next stage, while the subsequent instruction can advance into the decode stage. This improves the throughput of the instruction flow. For example, it may take three cycles for a floating-point instruction to complete, but if there are no stalls in the floating-point pipeline, a series of floating-point instructions can have a throughput of one instruction per cycle.

The MPC603e instruction pipeline has four major pipeline stages, described as follows:

- The fetch pipeline stage primarily involves retrieving instructions from the memory system and determining the location of the next instruction fetch. Additionally, if possible, the BPU decodes branches during the fetch stage and folds out branch instructions before the dispatch stage.
- The dispatch pipeline stage is responsible for decoding the instructions supplied by the instruction fetch stage, and determining which of the instructions are eligible to be dispatched in the current cycle. In addition, the source operands of the instructions are read from the appropriate register file and dispatched with the instruction to the execute pipeline stage. At the end of the dispatch pipeline stage, the dispatched instructions and their operands are latched by the appropriate execution unit.
- In the execute pipeline stage, each execution unit with an executable instruction executes the selected instruction (perhaps over multiple cycles), writes the instruction's result into the appropriate rename register, and notifies the completion stage when the execution has finished. In the case of an internal exception, the execution unit reports the exception to the completion/writeback pipeline stage and discontinues instruction execution until the exception is handled. The exception is not signaled until that instruction is the next to be completed. Execution of most floating-point instructions is pipelined within the FPU allowing up to three instructions to be executing in the FPU concurrently. The FPU pipeline stages are multiply, add, and round-convert. The LSU has two pipeline stages. The first stage is for effective address calculation and MMU translation, and the second is for accessing data in the cache.
- The complete/writeback pipeline stage maintains the correct architectural machine state and transfers the contents of the rename registers to the GPRs and FPRs as

instructions are retired. If the completion logic detects an instruction causing an exception, all following instructions are canceled, their execution results in rename registers are discarded, and instructions are fetched from the correct instruction stream.

A superscalar processor issues multiple independent instructions into multiple pipelines allowing instructions to execute in parallel. The MPC603e has five independent execution units, one each for integer instructions, floating-point instructions, branch instructions, load/store instructions, and system register instructions. The IU and the FPU each have dedicated register files for maintaining operands (GPRs and FPRs, respectively), allowing integer and floating-point calculations to occur simultaneously without interference. Integer division performance of the PID7t-603e has been improved, with the **divwux** and **divwx** instructions executing in 20 clock cycles instead of the 37 cycles required in the PID6-603e.

The MPC603e provides support for single-cycle store and it provides an adder/comparator in the system register unit that allows the dispatch and execution of multiple integer add and compare instructions on each cycle. Refer to Chapter 6, “Instruction Timing,” for more information.

Because the PowerPC architecture can be applied to such a wide variety of implementations, instruction timing among various processors varies accordingly.

1.3.7 System Interface

The system interface is specific for each processor implementation.

The MPC603e provides a versatile system interface that allows for a wide range of implementations. The interface includes a 32-bit address bus, a 32- or 64-bit data bus, and 56 control and information signals (see Figure 1-4). The system interface allows for address-only transactions, as well as address and data transactions. The MPC603e control and information signals include the address arbitration, address start, address transfer, transfer attribute, address termination, data arbitration, data transfer, data termination, and processor state signals. Test and control signals provide diagnostics for selected internal circuits.

The system interface supports bus pipelining, allowing the address tenure of one transaction to overlap the data tenure of another. The extent of the pipelining depends on external arbitration and control circuitry. Similarly, the MPC603e supports split-bus transactions for systems with multiple potential bus masters—one device can have mastership of the address bus while another has mastership of the data bus. Allowing multiple bus transactions to occur simultaneously increases the available bus bandwidth for other activity, and as a result, improves performance.

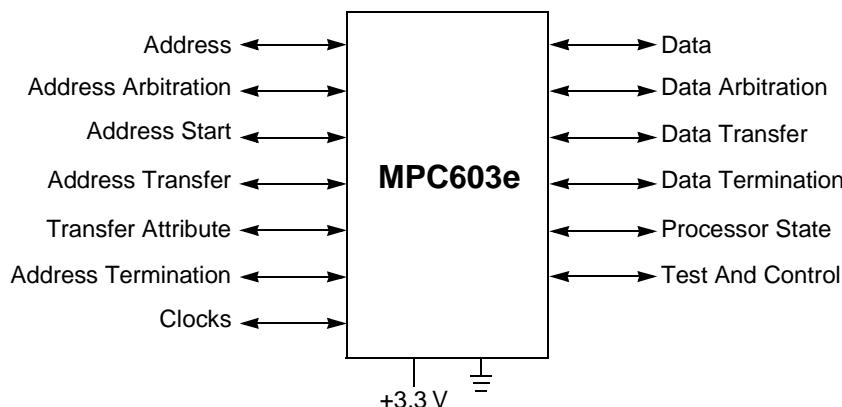


Figure 1-4. System Interface

The MPC603e supports multiple masters through a bus arbitration scheme that allows various devices to compete for the shared bus resource. Arbitration logic can implement priority protocols, such as fairness, and can park masters to avoid arbitration overhead. The MEI protocol ensures coherency among multiple devices and system memory. Also, the MPC603e on-chip caches, TLBs, and optional second-level caches can be controlled externally.

The MPC603e clocking structure allows the bus to operate at integer multiples of the processor cycle time.

The following sections describe the MPC603e bus support for memory operations. Note that some signals perform different functions depending on the addressing protocol used.

1.3.7.1 Memory Accesses

The MPC603e bus is configured at power-up to either a 32- or 64-bit width.

- When the processor is configured with a 32-bit data bus, memory accesses allow transfer sizes of 8, 16, 24, or 32 bits in one bus clock cycle. Data transfers occur in either single-beat transactions, two-beat or eight-beat burst transactions, with a single-beat transaction transferring as many as 32 bits. Single- or double-beat transactions are caused by noncached accesses that access memory directly (that is, reads and writes when caching is disabled, caching-inhibited accesses, and stores in write-through mode). Eight-beat burst transactions, which always transfer an entire cache block (32 bytes), are initiated when a line is read from or written to memory.
- When the MPC603e is configured with a 64-bit data bus, memory accesses allow transfer sizes of 8, 16, 24, 32, 40, 48, 56, or 64 bits in one bus clock cycle. Data transfers occur in either single-beat transactions or four-beat burst transactions. Single-beat transactions are caused by noncached accesses that access memory directly (that is, reads and writes when caching is disabled, caching-inhibited accesses, and stores in write-through mode). Four-beat burst transactions, which

always transfer an entire cache block (32 bytes), are initiated when a line is read from or written to memory.

1.3.7.2 Signals

The MPC603e signals are grouped as follows:

- Address arbitration signals—The MPC603e uses these signals to arbitrate for address bus mastership.
- Address transfer start signals—These signals indicate that a bus master has begun a transaction on the address bus.
- Address transfer signals—These signals, consisting of address bus, address parity, and address parity error signals, are used to transfer the address and to ensure the integrity of the transfer.
- Transfer attribute signals—These signals provide information about the type of transfer, such as the transfer size and whether the transaction is bursted, write-through, or caching-inhibited.
- Address transfer termination signals—These signals are used to acknowledge the end of the address phase of the transaction. They also indicate whether a condition exists that requires the address phase to be repeated.
- Data arbitration signals—The MPC603e uses these signals to arbitrate for data bus mastership.
- Data transfer signals—These signals, consisting of data bus, data parity, and data parity error signals, are used to transfer the data and to ensure the integrity of the transfer.
- Data transfer termination signals—Data termination signals are required after each data beat in a data transfer. In a single-beat transaction, the data termination signals also indicate the end of the tenure. In burst accesses, the data termination signals apply to individual beats and indicate the end of the tenure only after the final data beat. They also indicate whether a condition exists that requires the data phase to be repeated.
- System status signals—These signals include the interrupt signal, checkstop signals, and soft- and hard-reset signals. They are used to interrupt and, under various conditions, to reset the processor.
- Processor state signals—These signals indicate the state of the reservation coherency bit, enable the time base, provide machine quiescence control, and can be used to cause a machine halt on execution of a **tlbsync** instruction.
- JTAG/COP interface signals—The JTAG (IEEE 1149.1) interface and common on-chip processor (COP) unit provides a serial interface to the system for performing monitoring and boundary tests.
- Test interface signals—These signals are used for production testing.

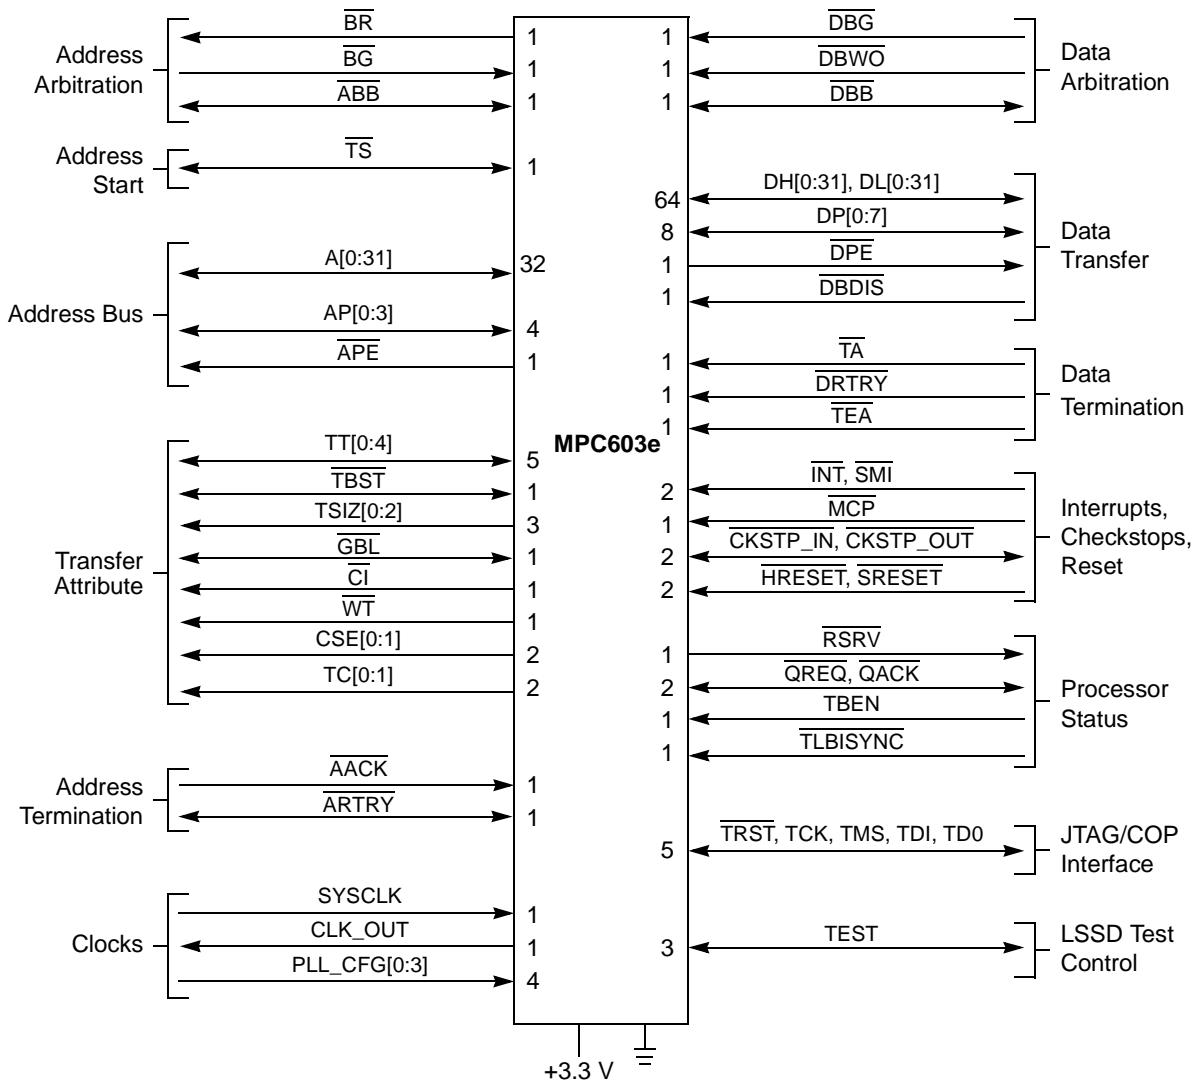
- Clock signals—These signals determine the system clock frequency and can be used to synchronize multiprocessor systems.

NOTE

A bar over a signal name indicates that the signal is active low—for example, $\overline{\text{ARTRY}}$ (address retry) and $\overline{\text{TS}}$ (transfer start). Active-low signals are referred to as asserted (active) when they are low and negated when they are high. Signals that are not active low, such as AP[0:3] (address bus parity signals) and TT[0:4] (transfer type signals) are referred to as asserted when they are high and negated when they are low.

1.3.7.3 Signal Configuration

Figure 1-5 illustrates the MPC603e logical pin configuration, showing how the signals are grouped.


Figure 1-5. Signal Groups



Implementation-Specific Information

Freescale Semiconductor, Inc.

Freescale Semiconductor, Inc.

MPC603e RISC Microprocessor User's Manual

**For More Information On This Product,
Go to: www.freescale.com**

Chapter 2

Programming Model

This chapter describes the PowerPC programming model with respect to the MPC603e microprocessor. It consists of three major sections that describe the following:

- Registers implemented in the MPC603e
- Operand conventions
- MPC603e instruction set

2.1 Register Set

This section describes the register organization in the MPC603e as defined by the three levels of the PowerPC architecture— user instruction set architecture (UISA), virtual environment architecture (VEA), and operating environment architecture (OEA), as well as the MPC603e implementation-specific registers. Full descriptions of the basic register set defined by the PowerPC architecture are provided in Chapter 2, “Register Set,” in the *Programming Environments Manual*.

The PowerPC architecture defines register-to-register operations for all computational instructions. Source data for these instructions is accessed from the on-chip registers or is provided as an immediate value embedded in the opcode. The three-register instruction format allows specification of a target register distinct from the two source registers, thus preserving the original data for use by other instructions and reducing the number of instructions required for certain operations. Data is transferred between memory and registers with explicit load and store instructions only.

Note that there may be registers common to other processors of this family that are not implemented in the MPC603e. When the MPC603e detects special-purpose register (SPR) encodings other than those defined in this document, it either takes an exception or it treats the instruction as a no-op. (Note that exceptions are referred to as interrupts in the architecture specification.) Conversely, some SPRs in the MPC603e may not be implemented in other processors or may not be implemented in the same way.

2.1.1 PowerPC Register Set

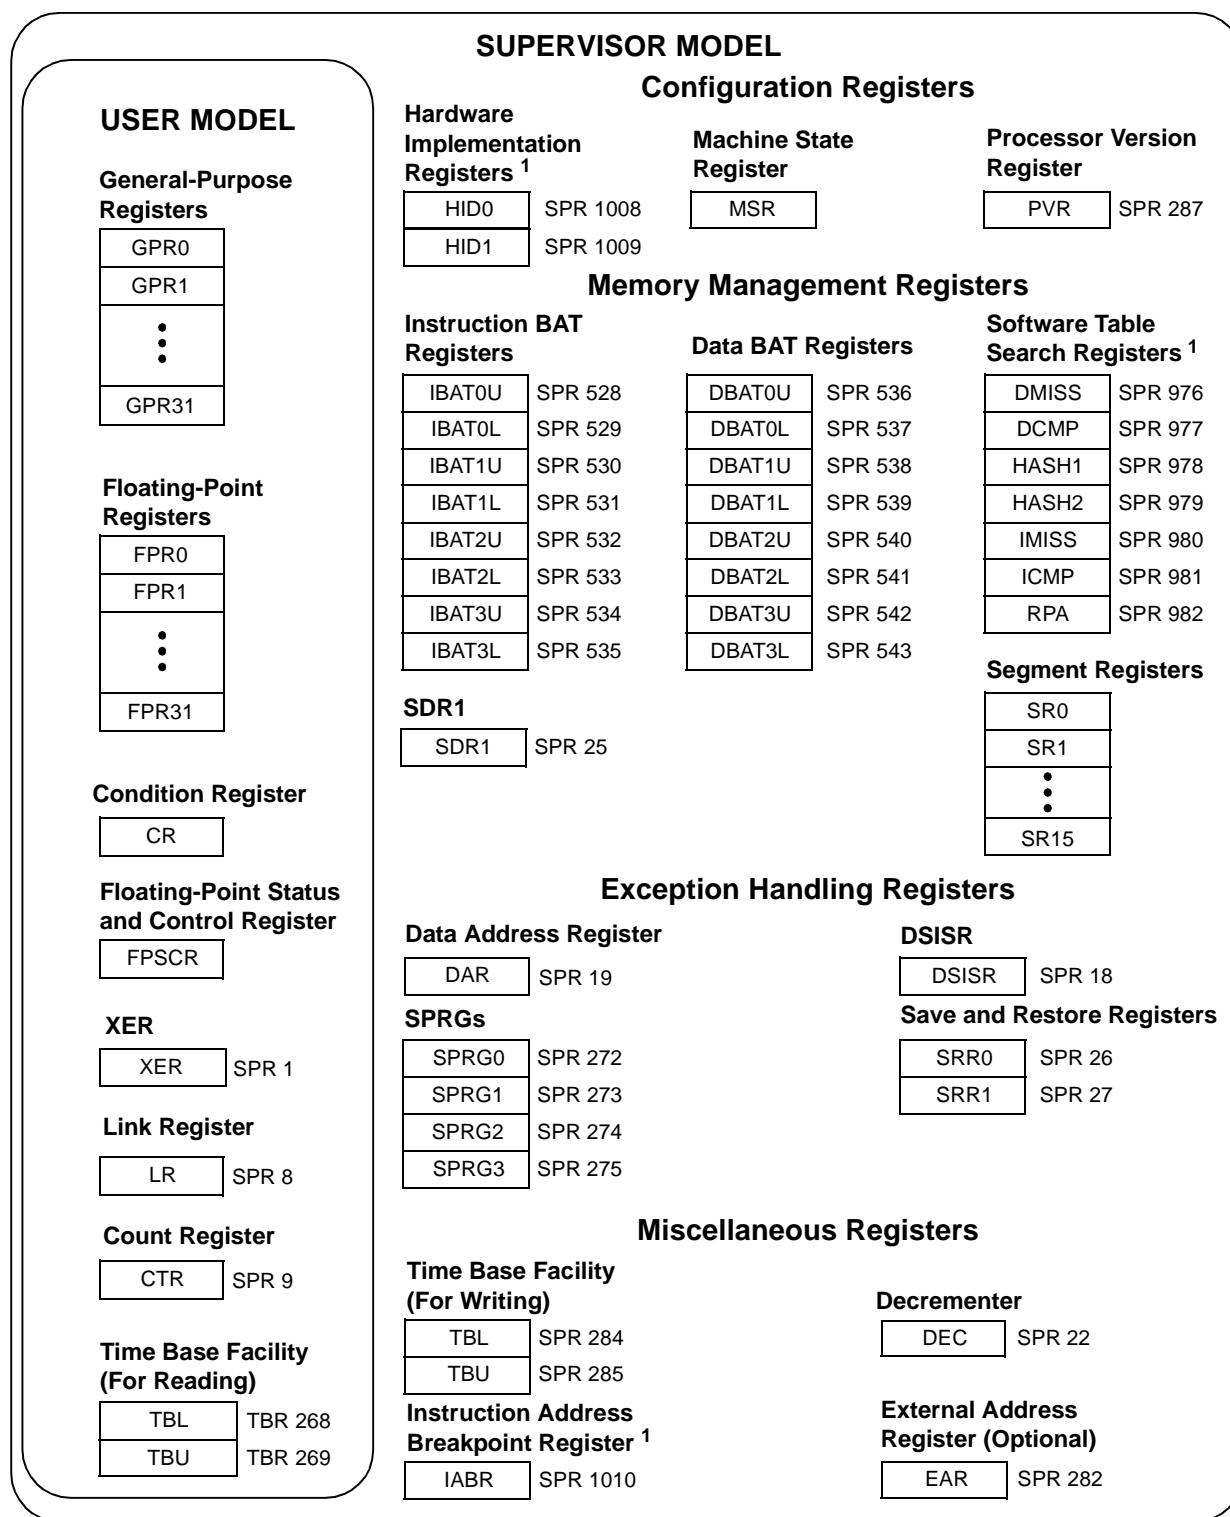
The PowerPC UISA registers, shown in Figure 2-1, can be accessed by either user- or supervisor-level instructions (the architecture specification refers to user- and supervisor-level as problem state and privileged state, respectively). The general-purpose registers (GPRs) and floating-point registers (FPRs) are accessed through instruction operands. Access to registers can be explicit (that is, through the use of specific instructions for that purpose such as the **mtspr** and **mfsp** instructions) or implicit as part of the execution (or side effect) of an instruction. Some registers are accessed both explicitly and implicitly.

The number to the right of the register name indicates the number that is used in the syntax of the instruction operands to access the register (for example, the number used to access the XER is SPR1).

For more information on the PowerPC register set, refer to Chapter 2, “Register Set,” in the *Programming Environments Manual*.

The MPC603e user-level registers are described as follows:

- User-level registers (UISA)—The user-level registers can be accessed by all software with either user or supervisor privileges. The user-level register set includes the following:
 - General-purpose registers (GPRs). The GPR file consists of thirty-two 32-bit GPRs designated as GPR0–GPR31. This register file serves as the data source or destination for all integer instructions and provides data for generating addresses.
 - Floating-point registers (FPRs). The FPR file consists of thirty-two 64-bit FPRs designated as FPR0–FPR31, which serves as the data source or destination for all floating-point instructions. These registers can contain data objects of either single- or double-precision floating-point format.
- Condition register (CR). The CR consists of eight 4-bit fields, CR0–CR7, that reflect the results of certain arithmetic operations and provides a mechanism for testing and branching.
- Floating-point status and control register (FPSCR). The FPSCR contains all floating-point exception signal bits, exception summary bits, exception enable bits, and rounding control bits needed for compliance with the IEEE 754 standard.



¹ These registers are MPC603e-specific (PID6-603e and PID7t-603e). They may not be supported by other processors of this family.

Figure 2-1. Programming Model—Registers

The remaining user-level registers are SPRs. Note that the PowerPC architecture provides a separate mechanism for accessing SPRs (the **mtspr** and **mfsp** instructions). These instructions are commonly used to explicitly access certain registers, while other SPRs may be accessed as the side effect of executing other instructions.

- XER register (XER). The 32-bit XER indicates overflow and carries for integer operations. It is set implicitly by many instructions.
- Link register (LR). The 32-bit LR provides the branch target address for the Branch Conditional to Link Register (**bclrx**) instruction and can optionally be used to hold the logical address (referred to as the effective address in the architecture specification) of the instruction that follows a branch and link instruction, typically used for linking to subroutines.
- Count register (CTR). The 32-bit CTR can be used to hold a loop count that can be decremented during execution of appropriately coded branch instructions. It can also provide the branch target address for the Branch Conditional to Count Register (**bcctrx**) instruction.
- User-level registers (VEA)—The VEA introduces the time base facility (TB) for reading. The TB is a 64-bit register pair whose contents are incremented once every four bus clock cycles. The TB consists of two 32-bit registers—time base upper (TBU) and time base lower (TBL). Note that the time base registers are read-only in user state.

The MPC603e supervisor-level registers are described as follows:

- Supervisor-level registers (OEA)—The OEA defines the registers an operating system uses for memory management, configuration, and exception handling. The PowerPC architecture defines the following supervisor-level registers:
 - Configuration registers
 - Machine state register (MSR). The MSR defines the state of the processor. The MSR can be modified by the Move to Machine State Register (**mtmsr**), System Call (**sc**), and Return from Exception (**rfi**) instructions. It can be read by the Move from Machine State Register (**mfmsr**) instruction.

Implementation Note—The MPC603e defines MSR[13] as the power management enable (POW) bit and MSR[14] as the temporary GPR remapping (TGPR) bit. These bits are described in Table 2-1.

- Processor version register (PVR). This read-only register identifies the version (model) and revision level of this processor.

Implementation Note—The processor version number is 6 for the PID6-603e and 7 for the PID7t-603e. The processor revision level starts at 0x0100 and changes for each chip revision. The revision level is updated on all silicon revisions.

Table 2-1. MSR[POW] and MSR[TGPR] Bits

Bits	Name	Description
13	POW	Power management enable (MPC603e-specific). Controls the programmable power modes only; it has no effect on dynamic power management (DPM). MSR[POW] may be altered with an mtmsr instruction only. Also, when altering the POW bit, software may alter only this bit in the MSR and no others. The mtmsr instruction must be followed by a context-synchronizing instruction. See Chapter 9, “Power Management,” for more information on power management. 0 Disables programmable power modes (normal operation mode). 1 Enables programmable power modes (nap, doze, or sleep mode).
14	TGPR	Temporary GPR remapping (MPC603e-specific). The contents of GPR0–GPR3 remain unchanged while MSR[TGPR] = 1. Attempts to use GPR4–GPR31 with MSR[TGPR] = 1 yield undefined results. When this bit is set, all instruction accesses to GPR0–GPR3 are mapped to TGPR0–TGPR3, respectively. TGPR is set when an instruction TLB miss, data TLB miss on load, or data TLB miss on store exception is taken. TGPR is cleared by an rfi instruction. 0 Normal operation 1 TGPR mode. GPR0–GPR3 are remapped to TGPR0–TGPR3 for use by TLB miss routines.

- Memory management registers
 - Block-address translation (BAT) registers. The MPC603e includes eight block-address translation registers (BATs): four pairs of instruction BATs (IBAT0U–IBAT3U and IBAT0L–IBAT3L) and four pairs of data BATs (DBAT0U–DBAT3U and DBAT0L–DBAT3L). Figure 2-1 lists SPR numbers for the BAT registers.
 - SDR1. The SDR1 register specifies the page table base address used in virtual-to-physical address translation. (Note that physical address is referred to as real address in the architecture specification.)
 - Segment registers (SRs). The PowerPC OEA defines sixteen 32-bit segment registers (SR0–SR15). The fields in the segment register are interpreted differently depending on the value of bit 0.
- Exception handling registers
 - Data address register (DAR). After a data access or an alignment exception, the DAR is set to the effective address generated by the faulting instruction.
 - SPRG0–SPRG3. The SPRG0–SPRG3 registers are provided for operating system use.
 - DSISR. The DSISR defines the cause of data access and alignment exceptions.
 - Machine status save/restore register 0 (SRR0). The SRR0 is used to save machine status on exceptions and to restore machine status when an **rfi** instruction is executed.
 - Machine status save/restore register 1 (SRR1). The SRR1 is used to save machine status on exceptions and to restore machine status when an **rfi** instruction is executed.

Implementation Note—The MPC603e implements the KEY bit (bit 12) in the SRR1 register to simplify the table search software. For more information refer to Chapter 5, “Memory Management.”

- Miscellaneous registers

- The time base facility (TB) for writing. The TB is a 64-bit register pair that can be used to provide time-of-day or interval timing. It consists of two 32-bit registers—time base upper (TBU) and time base lower (TBL). The TB is incremented once every four clock cycles on the MPC603e.
- Decrementer (DEC). The DEC register is a 32-bit decrementing counter that provides a mechanism for causing a decrementer exception after a programmable delay. The DEC is decremented once every four bus clock cycles.
- External access register (EAR). The EAR is a 32-bit register used in conjunction with the **eciwx** and **ecowx** instructions. Although the PowerPC architecture specifies that EAR[26–31] are used to select a device, the MPC603e implements only bits 28–31. Note that EAR and the **eciwx** and **ecowx** instructions are optional in the PowerPC architecture and may not be supported in all processors that implement the OEA.

2.1.2 Implementation-Specific Registers

The MPC603e defines the DMISS, IMISS, DCMP, ICMP, HASH1, HASH2, RPA, HID0, HID1, and IABR SPRs for software table search operations. These registers should be accessed only when address translation is disabled (MSR[IR] and MSR[DR] are both zero). For a complete discussion, refer to Section 5.5.2, “Implementation-Specific Table Search Operation.” Also, the HID0, HID1, and IABR SPRs are defined and described in this section. These registers can be accessed by supervisor-level instructions only using the SPR numbers shown in Figure 2-1.

2.1.2.1 Hardware Implementation Registers (HID0 and HID1)

The HID0 and HID1 registers, shown in Figure 2-2 and Figure 2-3, respectively, define enable bits for various MPC603e-specific features.

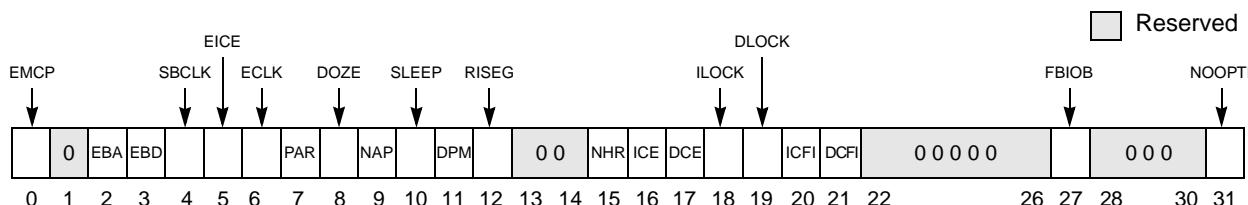


Figure 2-2. Hardware Implementation Register 0 (HID0)

Table 2-2 shows the bit definitions for HID0.

Table 2-2. HID0 Bit Functions

Bits	Name	Function
0	EMCP	Enable <u>MCP</u> . The primary purpose of this bit is to mask out further machine check exceptions caused by assertion of MCP, similar to how MSR[EE] can mask external interrupts. 0 Masks MCP. Asserting MCP does not generate a machine check exception or a checkstop. 1 Asserting MCP causes checkstop if MSR[ME] = 0 or a machine check exception if ME = 1.
1	—	Reserved, do not clear.
2	EBA	Enable 60x bus address parity checking. EBA and EBD allow the processor to operate with memory subsystems that do not generate parity. 0 Disables address parity checking. 1 Allows a address parity error to cause a checkstop if MSR[ME] = 0 or a machine check exception if MSR[ME] = 1.
3	EBD	Enable 60x bus data parity checking. EBA and EBD allow the processor to operate with memory subsystems that do not generate parity. 0 Disables data parity checking. 1 Allows a data parity error to cause a checkstop if MSR[ME] = 0 or a machine check exception if MSR[ME] = 1.
4	BCLK	<u>CLK_OUT</u> output enable and clock type selection. Used in conjunction with HID0[ECLK] and HRESET to configure CLK_OUT. See Table 2-3.
5	EICE	Enables in-circuit emulator outputs for pipeline tracking. See Section 7.2.11, “Pipeline Tracking Support.”
6	ECLK	<u>CLK_OUT</u> output enable and clock type selection. Used in conjunction with HID0[BCLK] and the HRESET signal to configure CLK_OUT. See Table 2-3.
7	PAR	Disable precharge of <u>ARTRY</u> . 0 Precharge of ARTRY enabled. 1 Alters bus protocol slightly by preventing the processor from driving <u>ARTRY</u> to high (negated) state. If this is done, the system must restore the signals to the high state.
8	DOZE ¹	Doze mode enable. Operates in conjunction with MSR[POW]. 0 Doze mode disabled. 1 Doze mode enabled. Doze mode is invoked by setting MSR[POW] while this bit is set. In doze mode, the PLL, time base, and snooping remain active.
9	NAP ¹	Nap mode enable. Operates in conjunction with MSR[POW]. 0 Nap mode disabled. 1 Nap mode enabled. Doze mode is invoked by setting MSR[POW] while this bit is set. In nap mode, the PLL and the time base remain active.
10	SLEEP ¹	Sleep mode enable. Operates in conjunction with MSR[POW]. 0 Sleep mode disabled. 1 Sleep mode enabled. Sleep mode is invoked by setting MSR[POW] while this bit is set. <u>QREQ</u> is asserted to indicate that the processor is ready to enter sleep mode. If the system logic determines that the processor may enter sleep mode, the quiesce acknowledge signal, <u>QACK</u> , is asserted back to the processor. Once QACK assertion is detected, the processor enters sleep mode after several processor clocks. At this point, the system logic may turn off the PLL by first configuring PLL_CFG[0:3] to PLL bypass mode, then disabling SYSCLK.
11	DPM ¹	Dynamic power management enable. 0 Dynamic power management is disabled. 1 Functional units enter a low-power mode automatically if the unit is idle. This does not affect operational performance and is transparent to software or any external hardware.

Chapter 2. Programming Model

Table 2-2. HID0 Bit Functions (continued)

Bits	Name	Function
12–15	—	Reserved, should be cleared.
16	ICE ²	<p>Instruction cache enable.</p> <p>0 The instruction cache is neither accessed nor updated. All pages are accessed as if they were marked cache-inhibited (WIM = x1x). Potential cache accesses from the bus (snoop and cache operations) are ignored. In the disabled state for the L1 caches, the cache tag state bits are ignored and all accesses are propagated to the 60x bus as single-beat transactions. For those transactions, however, CI reflects the original state determined by address translation regardless of cache disabled status. ICE is zero at power-up.</p> <p>1 The instruction cache is enabled.</p>
17	DCE	<p>Data cache enable.</p> <p>0 The data cache is neither accessed nor updated. All pages are accessed as if they were marked cache-inhibited (WIM = x1x). Potential cache accesses from the bus (snoop and cache operations) are ignored. In the disabled state for the L1 caches, the cache tag state bits are ignored and all accesses are propagated to the 60x bus as single-beat transactions. For those transactions, however, CI reflects the original state determined by address translation regardless of cache disabled status. DCE is zero at power-up.</p> <p>1 The data cache is enabled.</p>
18	ILOCK	<p>Instruction cache lock.</p> <p>0 Normal operation.</p> <p>1 Instruction cache is locked. A locked cache supplies data normally on a hit, but the access is treated as a cache-inhibited transaction on a miss. On a miss, the transaction to the 60x bus is single-beat; however, CI still reflects the original state as determined by address translation independent of cache locked or disabled status.</p> <p>To prevent locking during a cache access, an isync instruction must precede the setting of ILOCK.</p>
19	DLOCK	<p>Data cache lock.</p> <p>0 Normal operation.</p> <p>1 Data cache is locked. A locked cache supplies data normally on a hit but is treated as a cache-inhibited transaction on a miss. On a miss, the transaction to the 60x bus is single-beat; however, CI still reflects the original state as determined by address translation independent of cache locked or disabled status. A snoop hit to a locked L1 data cache performs as if the cache were not locked. A cache block invalidated by a snoop remains invalid until the cache is unlocked.</p> <p>To prevent locking during a cache access, a sync instruction must precede the setting of DLOCK.</p>
20	ICFI	<p>Instruction cache flash invalidate.</p> <p>0 The instruction cache is not invalidated. The bit is cleared when the invalidation operation begins (usually the next cycle after the write operation to the register). The instruction cache must be enabled for the invalidation to occur.</p> <p>1 An invalidate operation is issued that marks the state of each instruction cache block as invalid without writing back modified cache blocks to memory. Cache access is blocked during this time. Bus accesses to the cache are signaled as a miss during invalidate-all operations. Setting ICFI clears all the valid bits of the blocks and the PLRU bits to point to way L0 of each set.</p> <p>For MPC603e processors, the proper use of the ICFI and DCFI bits is to set them and clear them with two consecutive mtspr operations.</p>

Table 2-2. HID0 Bit Functions (continued)

Bits	Name	Function
21	DCFI	Data cache flash invalidate. 0 The data cache is not invalidated. The bit is cleared when the invalidation operation begins (usually the next cycle after the write operation to the register). The data cache must be enabled for the invalidation to occur. 1 An invalidate operation is issued that marks the state of each data cache block as invalid without writing back modified cache blocks to memory. Cache access is blocked during this time. Bus accesses to the cache are signaled as a miss during invalidate-all operations. Setting DCFI clears all the valid bits of the blocks and the PLRU bits to point to way L0 of each set. For MPC603e processors, the proper use of the ICFI and DCFI bits is to set them and clear them with two consecutive mtspr operations.
22–23	—	Reserved, should be cleared.
24	IFEM	Enable M bit on 60x bus for instruction fetches (PID7t-603e only). 0 M bit not reflected on bus for instruction fetches. Instruction fetches are treated as nonglobal on the bus 1 Instruction fetches reflect the M bit from the WIM settings.
25–26	—	Reserved, should be cleared.
27	FBIOD	Force branch indirect on bus. 0 Register indirect branch targets are fetched normally. 1 Forces register indirect branch targets to be fetched externally.
28	ABE	Address broadcast enable. Controls whether certain address-only operations (such as cache operations) are broadcast on the 60x bus. 0 Address-only operations affect only local caches and are not broadcast. 1 Address-only operations are broadcast on the 60x bus. Affected instructions are dcbi , dcbf , and dcbst . Note that these cache control instruction broadcasts are not snooped by the PID7t-603e. Refer to Section 3.2.3, “Data Cache Control,” for more information.
29–30	—	Reserved
31	NOOPTI	No-op the data cache touch instructions. 0 The dcbt and dcbtst instructions are enabled. 1 The dcbt and dcbtst instructions are no-operated globally.

¹ See Chapter 9, “Power Management.”² See Chapter 3, “Instruction and Data Cache Operation.”

Table 2-3 shows how HID0[BCLK], HID0[ECLK], and **HRESET** are used to configure CLK_OUT. See Section 7.2.12.2, “Test Clock (CLK_OUT)—Output,” for more information.

Table 2-3. HID0[BCLK] and HID0[ECLK] CLK_OUT Configuration

HRESET	HID0[ECLK]	HID0[BCLK]	CLK_OUT
Asserted	x	x	Bus
Negated	0	0	High impedance
Negated	0	1	Core clock frequency
Negated	1	0	Bus
Negated	1	1	Core clock frequency

Chapter 2. Programming Model

HID0 can be accessed with **mtspr** and **mfspr** using SPR1008.

PC0	PC1	PC2	PC3		0 0	31	
0	1	2	3	4			

Figure 2-3. Hardware Implementation Register 1 (HID1)

Table 2-4 shows the bit definitions for HID1.

Table 2-4. HID1 Bit Settings

Bits	Name	Description
0	PC0	PLL configuration bit 0 (read-only)
1	PC1	PLL configuration bit 1 (read-only)
2	PC2	PLL configuration bit 2 (read-only)
3	PC3	PLL configuration bit 3 (read-only)
4–31	—	Reserved, should be cleared.

Note: The clock configuration bits reflect the state of the PLL_CFG[0:3] signals.

HID1 can be accessed with **mfspr** using SPR1009.

2.1.2.2 Data and Instruction TLB Miss Address Registers (DMISS and IMISS)

DMISS and IMISS, shown in Figure 2-4, are loaded automatically upon a data or instruction TLB miss. DMISS and IMISS contain the effective page address of the access that caused the TLB miss exception. The contents are used by the MPC603e when calculating the values of HASH1 and HASH2 and by the **tlbld** and **tblli** instructions when loading a new TLB entry. Note that the MPC603e always loads DMISS with a big-endian address, even when MSR[LE] is set. These registers are read and write to the software.

	Effective Page Address	0	31
--	------------------------	---	----

Figure 2-4. DMISS and IMISS Registers

2.1.2.3 Data and Instruction TLB Compare Registers (DCMP and ICMP)

DCMP and ICMP, shown in Figure 2-5, contain the first word in the required PTE. The contents are constructed automatically from the contents of the segment registers and the effective address (DMISS or IMISS) when a TLB miss exception occurs. Each PTE read from the tables during the table search process should be compared with this value to determine if the PTE is a match. Upon execution of a **tlbld** or **tblli** instruction, the upper

25 bits of the DCMP or ICMP register and 11 bits of the effective address are loaded into the first word of the selected TLB entry. These registers are read and write to the software.

V	VSID	0	API
0 1		24 25 26	31

Figure 2-5. DCMP and ICMP Registers

Table 2-5 describes the bit settings for the DCMP and ICMP registers.

Table 2-5. DCMP and ICMP Bit Settings

Bits	Name	Description
0	V	Valid bit. Set by the processor on a TLB miss exception.
1–24	VSID	Virtual segment ID. Copied from VSID field of corresponding segment register.
25	—	Reserved, should be cleared.
26–31	API	Abbreviated page index. Copied from API of effective address.

2.1.2.4 Primary and Secondary Hash Address Registers (HASH1 and HASH2)

HASH1 and HASH2, shown in Figure 2-6, contain the physical addresses of the primary and secondary PTEGs for the access that caused the TLB miss exception. For convenience, the MPC603e automatically constructs the full physical address by routing SDR1[0–6] into HASH1 and HASH2 and clearing the lower 6 bits. These read-only registers are constructed from the DMISS or IMISS contents (the register choice is determined by which miss most recently occurred).

HTABORG[0-6]	Hashed Page Address	0 0 0 0 0 0
0	6 7	25 26
31		

Figure 2-6. HASH1 and HASH2 Registers

Table 2-6 describes the bit settings of the HASH1 and HASH2 registers.

Table 2-6. HASH1 and HASH2 Bit Settings

Bits	Name	Description
0–6	HTABORG[0–6]	Copy of the upper 7 bits of the HTABORG field from SDR1
7–25	Hashed page address	Address bits 7–25 of the PTEG to be searched
26–31	—	Reserved

2.1.2.5 Required Physical Address Register (RPA)

During a page table search operation, the software must load the RPA, shown in Figure 2-7, with the second word of the correct PTE. When the **tlbld** or **tblli** instruction is executed, the RPA and DMISS or IMISS register are merged and loaded into the selected TLB entry. The referenced (R) bit is ignored when the write occurs (no location exists in the TLB entry for this bit). The RPA register is read and write to the software.

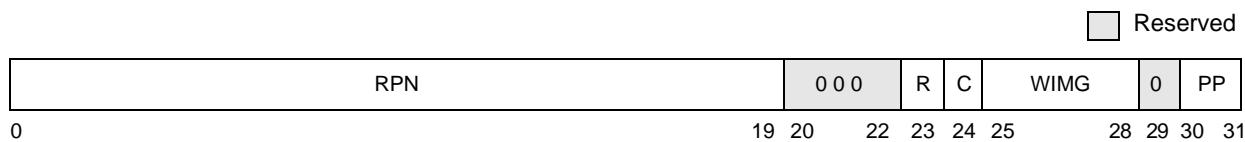


Figure 2-7. Required Physical Address Register (RPA)

Table 2-7 describes the bit settings of the RPA register.

Table 2-7. RPA Bit Settings

Bits	Name	Description
0–19	RPN	Physical page number from PTE
20–22	—	Reserved
23	R	Referenced bit from PTE
24	C	Changed bit from PTE
25–28	WIMG	Memory/cache access attribute bits
29	—	Reserved
30–31	PP	Page protection bits from PTE

2.1.2.6 Instruction Address Breakpoint Register (IABR)

The IABR, shown in Figure 2-8, controls the instruction address breakpoint exception. IABR[CEA] holds an effective address to which each instruction is compared. The exception is enabled by setting IABR[BE]. The exception is taken when there is an instruction address breakpoint match on the next instruction to complete. The instruction tagged with the match does not complete before the breakpoint exception is taken.

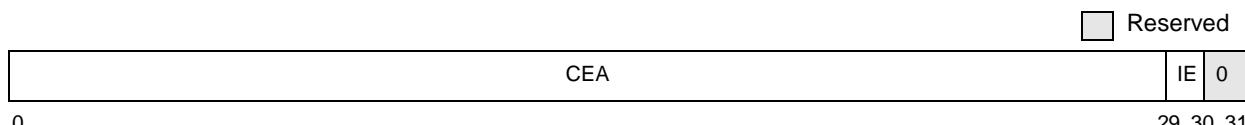


Figure 2-8. Instruction Address Breakpoint Register (IABR)

The bits in the IABR are defined in Table 2-8.

Table 2-8. Instruction Address Breakpoint Register Bit Settings

Bits	Description
0–29	Word address to be compared
30	IABR enabled. Setting this bit indicates that the IABR exception is enabled.
31	Reserved

2.2 Operand Conventions

This section describes the operand conventions as they are represented in two levels of the PowerPC architecture. It also provides detailed descriptions of conventions used for storing values in registers and memory, accessing the MPC603e registers, and representation of data in these registers.

2.2.1 Floating-Point Execution Models—UISA

The IEEE 754 standard includes 64- and 32-bit arithmetic. The standard requires that single-precision arithmetic be provided for single-precision operands. The standard permits double-precision arithmetic instructions to have either (or both) single-precision or double-precision operands, but states that single-precision arithmetic instructions should not accept double-precision operands.

The PowerPC UISA follows these guidelines:

- Double-precision arithmetic instructions may have single-precision operands but always produce double-precision results.
- Single-precision arithmetic instructions require all operands to be single-precision and always produce single-precision results.

For arithmetic instructions, conversions from double- to single-precision must be done explicitly by software, while conversions from single- to double-precision are done implicitly.

All implementations of the Power PC architecture provide the equivalent of the following execution models to ensure that identical results are obtained. The definition of the arithmetic instructions for infinities, denormalized numbers, and NaNs follow conventions described in the following sections.

Although the double-precision format specifies an 11-bit exponent, exponent arithmetic uses two additional bit positions to avoid potential transient overflow conditions. An extra bit is required when denormalized double-precision numbers are prenormalized. A second bit is required to permit computation of the adjusted exponent value in the following examples when the corresponding exception enable bit is one:

- Underflow during multiplication using a denormalized factor

- Overflow during division using a denormalized divisor

2.2.2 Data Organization in Memory and Data Transfers

Bytes in memory are numbered consecutively starting with 0. Each number is the address of the corresponding byte.

Memory operands may be bytes, half words, words, or double words, or, for the load/store multiple and move assist instructions, a sequence of bytes or words. The address of a memory operand is the address of its first byte (that is, of its lowest-numbered byte). Operand length is implicit for each instruction.

2.2.3 Alignment and Misaligned Accesses

The operand of a single-register memory access instruction has a natural alignment boundary equal to the operand length. In other words, the natural address of an operand is an integral multiple of the operand length. A memory operand is said to be aligned if it is aligned at its natural boundary; otherwise it is misaligned.

Operands for single-register memory access instructions have the characteristics shown in Table 2-9. (Although not permitted as memory operands, quad words are shown because quad-word alignment is desirable for certain memory operands.)

Table 2-9. Memory Operands

Operand	Length	Addr[28–31] If Aligned
Byte	8 bits	xxxx
Half word	2 bytes	xxx0
Word	4 bytes	xx00
Double word	8 bytes	x000
Quad word	16 bytes	0000

Note: An x in an address bit position indicates that the bit can be 0 or 1 independent of the state of other address bits.

The concept of alignment is also applied more generally to data in memory. For example, a 12-byte data item is said to be word-aligned if its address is a multiple of 4.

Implementation Notes—The following describes how the MPC603e handles alignment and misaligned accesses:

- The MPC603e provides hardware support for some misaligned memory accesses. However, misaligned accesses suffer a performance degradation compared to aligned accesses of the same type.

- The MPC603e does not provide hardware support for floating-point load/store operations that are not word-aligned. In such a case, the MPC603e invokes an alignment exception and the exception handler must break up the misaligned access. For this reason, floating-point single- and double-word accesses should always be word-aligned. Note that a floating-point double-word access on a word-aligned boundary requires an extra cycle to complete.

Any half-word, word, double-word, and string reference access that crosses an alignment boundary must be broken into multiple discrete accesses. For string accesses, the hardware makes no attempt to get aligned to reduce the number of accesses. (Multiple word accesses are architecturally required to be aligned.) The resulting performance degradation depends upon how well each individual access behaves with respect to the memory hierarchy. At a minimum, additional cache access cycles are required. More dramatically, each discrete access to a noncacheable page involves an individual bus operation that reduces the effective bus bandwidth.

The frequent use of misaligned accesses is discouraged because they can compromise the overall performance.

2.2.4 Floating-Point Operands

The MPC603e provides hardware support for all single- and double-precision floating-point operations for most value representations and all rounding modes. The PowerPC architecture provides for hardware to implement a floating-point system as defined in ANSI/IEEE Standard 754-1985, *IEEE Standard for Binary Floating Point Arithmetic*. For detailed information about the floating-point execution model, refer to Chapter 3, “Operand Conventions,” in the *Programming Environments Manual*.

2.2.5 Effect of Operand Placement on Performance

The VEA states that the placement (location and alignment) of operands in memory affect the relative performance of memory accesses. The best performance is guaranteed if memory operands are aligned on natural boundaries. To obtain the best performance from the MPC603e, the programmer should assume the performance model described in Chapter 3, “Operand Conventions,” in the *Programming Environments Manual*.

2.3 Instruction Set Summary

This section describes instructions and addressing modes defined for the MPC603e. These instructions are divided into the following functional categories:

- Integer instructions—These include arithmetic and logical instructions. For more information, see Section 2.3.4.1, “Integer Instructions.”

- Floating-point instructions—These include floating-point arithmetic instructions, as well as instructions that affect the floating-point status and control register (FPSCR). For more information, see Section 2.3.4.2, “Floating-Point Instructions.”
- Load and store instructions—These include integer and floating-point load and store instructions. For more information, see Section 2.3.4.3, “Load and Store Instructions.”
- Flow control instructions—These include branching instructions, condition register logical instructions, and other instructions that affect the instruction flow. For more information, see Section 2.3.4.4, “Branch and Flow Control Instructions.”
- Trap instructions—These are used to test for a specified set of conditions; see Section 2.3.4.5, “Trap Instructions.”
- Processor control instructions—These are used for synchronizing memory accesses and managing caches, TLBs, and segment registers. For more information, see Section 2.3.4.6, “Processor Control Instructions,” Section 2.3.5.1, “Processor Control Instructions,” and Section 2.3.6.2, “Processor Control Instructions—OEA.”
- Memory synchronization instructions—These are used for synchronizing memory accesses. See Section 2.3.4.7, “Memory Synchronization Instructions—UISA” and Section 2.3.5.2, “Memory Synchronization Instructions—VEA.”
- Memory control instructions—These provide control of caches, TLBs, and segment registers. For more information, see Section 2.3.5.3, “Memory Control Instructions—VEA,” and Section 2.3.6.3, “Memory Control Instructions—OEA.”
- System linkage instructions—These include the System Call (**sc**) and Return from Interrupt (**rff**) instructions. See Section 2.3.6.1, “System Linkage Instructions.”
- External control instructions—These include instructions for use with special input/output devices. See Section 2.3.5.4, “External Control Instructions.”

Note that this grouping of instructions does not necessarily indicate the execution unit that processes a particular instruction or group of instructions. This information, which is useful in taking full advantage of the MPC603e superscalar parallel instruction execution, is provided in Chapter 8, “Instruction Set,” of the *Programming Environments Manual*.

Integer instructions operate on word operands. Floating-point instructions operate on single- and double-precision floating-point operands. The PowerPC instructions are 4-byte words. The UISA provides for byte, half-word, and word operand loads and stores between memory and a set of 32 GPRs. It also provides for word and double-word operand loads and stores between memory and a set of 32 FPRs.

Arithmetic and logical instructions do not read or modify memory. To use the contents of a memory location in a computation and then modify the same or another memory location, the memory contents must be loaded into a register, modified, and then written to the target location using load and store instructions.

The description of each instruction includes the mnemonic and a formatted list of operands. To simplify assembly language programming, a set of simplified mnemonics (extended mnemonics in the architecture specification) and symbols is provided for some of the frequently-used instructions; see Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* for a complete list of simplified mnemonic examples.

2.3.1 Classes of Instructions

The MPC603e instructions belong to one of the following three classes:

- Defined
- Illegal
- Reserved

Note that although the definitions of these terms are consistent among the processors of this family, the assignment of these classifications is not. For example, an instruction that is specific to 64-bit implementations is considered defined for 64-bit implementations but illegal for 32-bit implementations such as the MPC603e.

The class is determined by examining the primary opcode and the extended opcode, if any. If either is not that of a defined instruction or of a reserved instruction, the instruction is illegal.

In future versions of the PowerPC architecture, instruction codings that are now illegal may become assigned to instructions in the architecture or may be reserved by being assigned to processor-specific instructions.

2.3.1.1 Definition of Boundedly Undefined

If instructions are encoded with incorrectly set bits in reserved fields, the results on execution can be said to be boundedly undefined. If a user-level program executes the incorrectly coded instruction, the resulting undefined results are bounded in that a spurious change from user to supervisor state is not allowed, and the level of privilege exercised by the program in relation to memory access and other system resources cannot be exceeded. Boundedly undefined results for a given instruction may vary between implementations, and between execution attempts in the same implementation.

2.3.1.2 Defined Instruction Class

Defined instructions are guaranteed to be supported in all implementations of the Power PC architecture, except as stated in the instruction descriptions in Chapter 8, “Instruction Set,” in the *Programming Environments Manual*. The MPC603e provides hardware support for all instructions defined for 32-bit implementations.

A processor of this family invokes the illegal instruction error handler (part of the program exception) when it encounters PowerPC instructions that have not been implemented. The instructions can be emulated in software, as required.

A defined instruction can have invalid forms, as described in the following section.

2.3.1.3 Illegal Instruction Class

Illegal instructions are grouped into the following categories:

- Instructions not defined in the PowerPC architecture. These opcodes are available for future extensions of the PowerPC architecture; that is, future versions of the architecture may define any of these instructions to perform new functions.

The following primary opcodes are defined as illegal but may be used in future extensions to the architecture:

1, 4, 5, 6, 9, 22, 56, 57, 60, 61

- Instructions defined in the PowerPC architecture but not implemented in a specific implementation. For example, instructions that can be executed on 64-bit processors are considered illegal by 32-bit processors.

The following primary opcodes are defined for 64-bit implementations only and are illegal on the MPC603e:

2, 30, 58, 62

- All unused extended opcodes are illegal. The unused extended opcodes can be determined from information in Section A.2, “Instructions Sorted by Opcode,” and Section 2.3.1.4, “Reserved Instruction Class.” Notice that extended opcodes for instructions that are defined only for 64-bit implementations are illegal in 32-bit implementations, and vice versa.

The following primary opcodes have unused extended opcodes:

17, 19, 31, 59, 63

(primary opcodes 30 and 62 are illegal for all 32-bit implementations, but as 64-bit opcodes they have some unused extended opcodes)

- An instruction consisting entirely of zeros is guaranteed to be an illegal instruction. This increases the probability that an attempt to execute data or uninitialized memory invokes the system illegal instruction error handler (a program exception). Note that if only the primary opcode consists of all zeros, the instruction is considered a reserved instruction. This is further described in Section 2.3.1.4, “Reserved Instruction Class.”

An attempt to execute an illegal instruction invokes the illegal instruction error handler (a program exception) but has no other effect. Section 4.5.7, “Program Exception (0x00700),” describes illegal and invalid instruction exceptions.

Except for an instruction consisting entirely of binary zeros, illegal instructions are available for further additions to the PowerPC architecture.

2.3.1.4 Reserved Instruction Class

Reserved instructions are allocated to specific implementation-dependent purposes not defined by the PowerPC architecture. An attempt to execute an unimplemented reserved instruction invokes the illegal instruction error handler (a program exception). See Section 4.5.7, “Program Exception (0x00700),” for additional information about illegal and invalid instruction exceptions.

The following types of instructions are included in this class:

- Implementation-specific instructions (for example, Load Data TLB Entry (**tblld**) and Load Instruction TLB Entry (**tblli**) instructions)
- Optional instructions defined by the PowerPC architecture but not implemented by the MPC603e (for example, Floating Square Root (**fsqrt**) and Floating Square Root Single (**fsqrts**) instructions)

2.3.2 Addressing Modes

This section provides an overview of conventions for addressing memory and calculating effective addresses as defined by the PowerPC architecture for 32-bit implementations. For more detailed information, see “Conventions” in Chapter 4, “Addressing Modes and Instruction Set Summary,” of the *Programming Environments Manual*.

2.3.2.1 Memory Addressing

A program references memory using the effective (logical) address computed by the processor when it executes a memory access or branch instruction or when it fetches the next sequential instruction.

Bytes in memory are numbered consecutively starting with zero. Each number is the address of the corresponding byte.

2.3.2.2 Memory Operands

Memory operands may be bytes, half words, words, or double words, or, for the load/store multiple and load/store string instructions, a sequence of bytes or words. The address of a memory operand is the address of its first byte (that is, of its lowest-numbered byte). Operand length is implicit for each instruction. The PowerPC architecture supports both big-endian and little-endian byte ordering. The default byte and bit ordering is big-endian. See “Byte Ordering” in Chapter 3, “Operand Conventions,” in the *Programming Environments Manual* for more information about big-endian and little-endian byte ordering.

The operand of a single-register memory access instruction has a natural alignment boundary equal to the operand length. In other words, the ‘natural’ address of an operand is an integral multiple of the operand length. A memory operand is said to be aligned if it is aligned at its natural boundary; otherwise it is misaligned. For a detailed discussion about memory operands, see Chapter 3, “Operand Conventions,” in the *Programming Environments Manual*.

2.3.2.3 Effective Address Calculation

An effective address (EA) is the 32-bit sum computed by the processor when executing a memory access or branch instruction or when fetching the next sequential instruction. For a memory access instruction, if the sum of the effective address and the operand length exceeds the maximum effective address, the memory operand is considered to wrap around from the maximum effective address through effective address 0, as described in the following paragraphs.

Effective address computations for both data and instruction accesses use 32-bit unsigned binary arithmetic. A carry from bit 0 is ignored.

Load and store operations have three categories of effective address generation:

- Register indirect with immediate index mode
- Register indirect with index mode
- Register indirect mode

Section 2.3.4.3.2, “Integer Load and Store Address Generation,” describes effective address generation for load and store operations.

Branch instructions have three categories of effective address generation:

- Immediate
- Link register indirect
- Count register indirect

Section 2.3.4.4.1, “Branch Instruction Address Calculation,” describes branch instruction effective address generation.

2.3.2.4 Synchronization

The synchronization described in this section refers to the state of the processor performing the synchronization.

2.3.2.4.1 Context Synchronization

The System Call (**sc**) and Return from Interrupt (**rfi**) instructions perform context synchronization by allowing previously issued instructions to complete before performing a change in context. Execution of one of these instructions ensures the following:

- No higher priority exception exists (**sc**).
- All previous instructions have completed to a point where they can no longer cause an exception. If a prior memory access instruction causes direct-store error exceptions, the results are guaranteed to be determined before this instruction is executed.
- Previous instructions complete execution in the context (privilege, protection, and address translation) under which they were issued.
- The instructions following the **sc** or **rfi** instruction execute in the context established by these instructions.

2.3.2.4.2 Execution Synchronization

An instruction is execution synchronizing if all previously initiated instructions appear to have completed before the instruction is initiated or, in the case of the Synchronize (**sync**) and Instruction Synchronize (**isync**) instructions, before the instruction completes. For example, the Move to Machine State Register (**mtmsr**) instruction is execution synchronizing. It ensures that all preceding instructions have completed execution and will not cause an exception before the instruction executes but does not ensure subsequent instructions execute in the newly established environment. For example, if the **mtmsr** sets MSR[PR], unless an **isync** immediately follows the **mtmsr** instruction, a privileged instruction could be executed or privileged access could be performed without causing an exception even though MSR[PR] indicates user mode.

2.3.2.4.3 Instruction-Related Exceptions

There are two kinds of exceptions in the MPC603e—those caused directly by the execution of an instruction and those caused by an asynchronous event. Either may cause components of the system software to be invoked.

Exceptions can be caused directly by the execution of an instruction as follows:

- An attempt to execute an illegal instruction causes the illegal instruction (program exception) handler to be invoked. An attempt by a user-level program to execute the supervisor-level instructions listed below causes the privileged instruction (program exception) handler to be invoked. The MPC603e provides the following supervisor-level instructions: **dcbi**, **mfmsr**, **mfspr**, **mfsr**, **mfsrin**, **mtmsr**, **mtspr**, **mtsri**, **mtsri**, **rfi**, **tlbie**, **tlbsync**, **tblld**, and **tblli**. Note that the privilege level of the **mfspr** and **mtspr** instructions depends on the SPR encoding.

- An attempt to access memory that is not available (page fault) causes the ISI exception handler to be invoked.
- An attempt to access memory with an effective address alignment that is invalid for the instruction causes the alignment exception handler to be invoked.
- The execution of an **sc** instruction invokes the system call exception handler that permits a program to request the system to perform a service.
- The execution of a trap instruction invokes the program exception trap handler.
- The execution of a floating-point instruction when floating-point instructions are disabled or unavailable invokes the floating-point unavailable exception handler.
- The execution of an instruction that causes a floating-point exception while exceptions are enabled in the MSR invokes the program exception handler.

Exceptions caused by asynchronous events are described in Chapter 4, “Exceptions.”

2.3.3 Instruction Set Overview

This section provides a brief overview of the PowerPC instructions implemented in the MPC603e and highlights any special information with respect to how the MPC603e implements a particular instruction. Note that the categories used in this section correspond to those used in Chapter 4, “Addressing Modes and Instruction Set Summary,” in the *Programming Environments Manual*. These categorizations are somewhat arbitrary and are provided for the convenience of the programmer and do not necessarily reflect the PowerPC architecture specification.

Note that some of the instructions have the following optional features:

- CR Update—The dot (.) suffix on the mnemonic enables the update of the CR.
- Overflow option—The **o** suffix indicates that the overflow bit in the XER is enabled.

2.3.4 PowerPC UIA Instructions

The UIA includes the base user-level instruction set (excluding a few user-level cache control, synchronization, and time base instructions), user-level registers, programming model, data types, and addressing modes. This section discusses the instructions defined in the UIA.

2.3.4.1 Integer Instructions

This section describes the integer instructions. These consist of the following:

- Integer arithmetic instructions
- Integer compare instructions

- Integer logical instructions
- Integer rotate and shift instructions

Integer instructions use the content of the GPRs as source operands and place results into GPRs, into the XER, and into condition register (CR) fields.

2.3.4.1.1 Integer Arithmetic Instructions

Table 2-10 lists the integer arithmetic instructions for the MPC603e.

Table 2-10. Integer Arithmetic Instructions

Name	Mnemonic	Operand Syntax
Add	add (add. addo addo.)	rD,rA,rB
Add Carrying	addc (addc. addco addco.)	rD,rA,rB
Add Extended	adde (adde. addeo addeo.)	rD,rA,rB
Add Immediate	addi	rD,rA,SIMM
Add Immediate Carrying	addic	rD,rA,SIMM
Add Immediate Carrying and Record	addic.	rD,rA,SIMM
Add Immediate Shifted	addis	rD,rA,SIMM
Add to Minus One Extended	addme (addme. addmeo addmeo.)	rD,rA
Add to Zero Extended	addze (addze. addzeo addzeo.)	rD,rA
Divide Word	divw (divw. divwo divwo.)	rD,rA,rB
Divide Word Unsigned	divwu (divwu. divwuo divwuo.)	rD,rA,rB
Multiply High Word	mulhw (mulhw.)	rD,rA,rB
Multiply High Word Unsigned	mulhwu (mulhwu.)	rD,rA,rB
Multiply Low	mullw (mullw. mullwo mullwo.)	rD,rA,rB
Multiply Low Immediate	mulli	rD,rA,SIMM
Negate	neg (neg. nego nego.)	rD,rA
Subtract From	subf (subf. subfo subfo.)	rD,rA,rB
Subtract From Carrying	subfc (subfc. subfco subfco.)	rD,rA,rB
Subtract From Extended	subfe (subfe. subfeo subfeo.)	rD,rA,rB
Subtract From Immediate Carrying	subfic	rD,rA,SIMM
Subtract From Minus One Extended	subfme (subfme. subfmeo subfmeo.)	rD,rA
Subtract From Zero Extended	subfze (subfze. subfzeo subfzeo.)	rD,rA

Although there is no Subtract Immediate instruction, its effect can be achieved by using an **addi** instruction with the immediate operand negated. Simplified mnemonics are provided that include this negation. The **subf** instructions subtract the second operand (**rA**) from the third operand (**rB**). Simplified mnemonics are provided in which the third operand is

subtracted from the second operand. See Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* for examples.

2.3.4.1.2 Integer Compare Instructions

The integer compare instructions algebraically or logically compare the contents of **rA** with either the UIMM operand, the SIMM operand, or the contents of **rB**. The comparison is signed for the **cmpi** and **cmp** instructions, and unsigned for the **cmpli** and **cmpl** instructions. Table 2-11 lists the integer compare instructions.

Table 2-11. Integer Compare Instructions

Name	Mnemonic	Operand Syntax
Compare	cmp	crfD,L,rA,rB
Compare Immediate	cmpi	crfD,L,rA,SIMM
Compare Logical	cmpl	crfD,L,rA,rB
Compare Logical Immediate	cmpli	crfD,L,rA,UIMM

The **crfD** operand can be omitted if the result of the comparison is to be placed in CR0. Otherwise, the target CR field must be specified in the instruction **crfD** field.

For more information refer to Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual*.

2.3.4.1.3 Integer Logical Instructions

The logical instructions shown in Table 2-12 perform bit-parallel operations. Logical instructions with the CR update enabled and instructions **andi.** and **andis.** set CR field CR0 to characterize the result of the logical operation. These fields are set as if the sign-extended low-order 32 bits of the result were algebraically compared to zero. Logical instructions without CR update and the remaining logical instructions do not modify the CR. Logical instructions do not affect the XER[SO], XER[OV], and XER[CA] bits.

For simplified mnemonics examples for the integer logical operations see Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual*.

Table 2-12. Integer Logical Instructions

Name	Mnemonic	Operand Syntax
AND	and (and.)	rA,rS,rB
AND Immediate	andi.	rA,rS,UIMM
AND Immediate Shifted	andis.	rA,rS,UIMM
AND with Complement	andc (andc.)	rA,rS,rB
Count Leading Zeros Word	cntlzw (cntlzw.)	rA,rS
Equivalent	eqv (eqv.)	rA,rS,rB

Table 2-12. Integer Logical Instructions (continued)

Name	Mnemonic	Operand Syntax
Extend Sign Byte	extsb (extsb.)	rA,rS
Extend Sign Half Word	extsh (extsh.)	rA,rS
NAND	nand (nand.)	rA,rS,rB
NOR	nor (nor.)	rA,rS,rB
OR	or (or.)	rA,rS,rB
OR Immediate	ori	rA,rS,UIMM
OR Immediate Shifted	oris	rA,rS,UIMM
OR with Complement	orc (orc.)	rA,rS,rB
XOR	xor (xor.)	rA,rS,rB
XOR Immediate	xori	rA,rS,UIMM
XOR Immediate Shifted	xoris	rA,rS,UIMM

2.3.4.1.4 Integer Rotate and Shift Instructions

Rotation operations are performed on data from a GPR, and the result, or a portion of the result, is returned to a GPR. See Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* for a complete list of simplified mnemonics that allows simpler coding of often-used functions such as clearing the leftmost or rightmost bits of a register, left justifying or right justifying an arbitrary field, and simple rotates and shifts.

Integer rotate instructions rotate the contents of a register. The result of the rotation is either inserted into the target register under control of a mask (if a mask bit is 1, the associated bit of the rotated data is placed into the target register; and if the mask bit is 0, the associated bit in the target register is unchanged), or ANDed with a mask before being placed into the target register.

The integer rotate instructions are listed in Table 2-13.

Table 2-13. Integer Rotate Instructions

Name	Mnemonic	Operand Syntax
Rotate Left Word Immediate then AND with Mask	rlwinm (rlwinm.)	rA,rS,SH,MB,ME
Rotate Left Word Immediate then Mask Insert	rlwimi (rlwimi.)	rA,rS,SH,MB,ME
Rotate Left Word then AND with Mask	rlwnm (rlwnm.)	rA,rS,rB,MB,ME

The integer shift instructions perform left and right shifts. Immediate-form logical (unsigned) shift operations are obtained by specifying masks and shift values for certain rotate instructions. Simplified mnemonics are provided, making coding of such shifts simpler and easier to understand.

Multiple-precision shifts can be programmed as shown in Appendix C, “Multiple-Precision Shifts,” in the *Programming Environments Manual*.

The integer shift instructions are listed in Table 2-14.

Table 2-14. Integer Shift Instructions

Name	Mnemonic	Operand Syntax
Shift Left Word	slw (slw.)	rA,rS,rB
Shift Right Algebraic Word	sraw (sraw.)	rA,rS,rB
Shift Right Algebraic Word Immediate	srawi (srawi.)	rA,rS,SH
Shift Right Word	srw (srw.)	rA,rS,rB

2.3.4.2 Floating-Point Instructions

This section describes the floating-point instructions, which include the following:

- Floating-point arithmetic instructions
- Floating-point multiply-add instructions
- Floating-point rounding and conversion instructions
- Floating-point compare instructions
- Floating-point status and control register instructions
- Floating-point move instructions

See Section 2.3.4.3, “Load and Store Instructions,” for information about floating-point loads and stores.

The PowerPC architecture supports a floating-point system as defined in the IEEE 754 standard, but requires software support to conform with that standard. All floating-point operations conform to the IEEE 754 standard, except if software sets the non-IEEE mode bit (NI) in the FPSCR. The MPC603e is in the nonnormalized mode when the NI bit is set in the FPSCR. If a denormalized result is produced, a default result of zero is generated. The generated zero has the same sign as the denormalized number. The MPC603e performs single- and double-precision floating-point operations compliant with the IEEE 754 floating-point standard.

Implementation Note—Single-precision denormalized results require two additional processor clock cycles to round. When loading or storing a single-precision denormalized number, the load/store unit may take up to 24 processor clock cycles to convert between the internal double-precision format and the external single-precision format.

2.3.4.2.1 Floating-Point Arithmetic Instructions

The floating-point arithmetic instructions are listed in Table 2-15.

Table 2-15. Floating-Point Arithmetic Instructions

Name	Mnemonic	Operand Syntax
Floating Add (Double-Precision)	fadd (fadd.)	frD,frA,frB
Floating Add Single	fadds (fadds.)	frD,frA,frB
Floating Divide (Double-Precision)	fdiv (fdiv.)	frD,frA,frB
Floating Divide Single	fdivs (fdivs.)	frD,frA,frB
Floating Multiply (Double-Precision)	fmul (fmul.)	frD,frA,frC
Floating Multiply Single	fmuls (fmuls.)	frD,frA,frC
Floating Reciprocal Estimate Single	fres (fres.)	frD,frB
Floating Reciprocal Square Root Estimate	frsqrt (frsqrt.)	frD,frB
Floating Select	fsel (fsel.)	frD,frA,frC,frB
Floating Subtract (Double-Precision)	fsub (fsub.)	frD,frA,frB
Floating Subtract Single	fsubs (fsubs.)	frD,frA,frB

2.3.4.2.2 Floating-Point Multiply-Add Instructions

These instructions combine multiply and add operations without an intermediate rounding operation. The fractional part of the intermediate product is 106 bits wide, and all 106 bits take part in the add/subtract portion of the instruction.

The floating-point multiply-add instructions are listed in Table 2-16.

Table 2-16. Floating-Point Multiply-Add Instructions

Name	Mnemonic	Operand Syntax
Floating Multiply-Add (Double-Precision)	fmadd (fmadd.)	frD,frA,frC,frB
Floating Multiply-Add Single	fma (fma.)	frD,frA,frC,frB
Floating Multiply-Subtract (Double-Precision)	fmsub (fmsub.)	frD,frA,frC,frB
Floating Multiply-Subtract Single	fms (fms.)	frD,frA,frC,frB
Floating Negative Multiply-Add (Double-Precision)	fnmadd (fnmadd.)	frD,frA,frC,frB
Floating Negative Multiply-Add Single	fnma (fnma.)	frD,frA,frC,frB
Floating Negative Multiply-Subtract (Double-Precision)	fnmsub (fnmsub.)	frD,frA,frC,frB
Floating Negative Multiply-Subtract Single	fnms (fnms.)	frD,frA,frC,frB

Implementation Note—Single-precision multiply-type instructions operate faster than their double-precision equivalents. See Chapter 6, “Instruction Timing,” for more information.

2.3.4.2.3 Floating-Point Rounding and Conversion Instructions

The Floating Round to Single-Precision (**frsp**) instruction is used to truncate a 64-bit double-precision number to a 32-bit single-precision floating-point number. The floating-point conversion instructions convert a 64-bit double-precision floating-point number to a 32-bit signed integer number.

The PowerPC architecture defines bits 0–31 of floating-point register **frD** as undefined when executing the Floating Convert to Integer Word (**fctiw**) and Floating Convert to Integer Word with Round toward Zero (**fctiwz**) instructions.

Examples of uses of these instructions to perform various conversions can be found in Appendix D, “Floating-Point Models,” in the *Programming Environments Manual*. The floating-point rounding instructions are shown in Table 2-17.

Table 2-17. Floating-Point Rounding and Conversion Instructions

Name	Mnemonic	Operand Syntax
Floating Convert to Integer Word	fctiw (fctiw.)	frD,frB
Floating Convert to Integer Word with Round Toward Zero	fctiwz (fctiwz.)	frD,frB
Floating Round to Single-Precision	frsp (frsp.)	frD,frB

2.3.4.2.4 Floating-Point Compare Instructions

Floating-point compare instructions compare the contents of two floating-point registers. The comparison ignores the sign of zero (that is $+0 = -0$). The floating-point compare instructions are listed in Table 2-18.

Table 2-18. Floating-Point Compare Instructions

Name	Mnemonic	Operand Syntax
Floating Compare Ordered	fcmo	crfD,frA,frB
Floating Compare Unordered	fcmpu	crfD,frA,frB

2.3.4.2.5 Floating-Point Status and Control Register Instructions

Every FPSCR instruction appears to synchronize the effects of all floating-point instructions executed by a given processor. Executing an FPSCR instruction ensures that all floating-point instructions previously initiated by the given processor appear to have completed before the FPSCR instruction is initiated and that no subsequent floating-point instructions appear to be initiated by the given processor until the FPSCR instruction has completed. The FPSCR instructions are listed in Table 2-19.

Table 2-19. Floating-Point Status and Control Register Instructions

Name	Mnemonic	Operand Syntax
Move from FPSCR	mffs (mffs.)	frD
Move to Condition Register from FPSCR	mcrfs	crfD,crfS
Move to FPSCR Bit 0	mtfsb0 (mtfsb0.)	crbD
Move to FPSCR Bit 1	mtfsb1 (mtfsb1.)	crbD
Move to FPSCR Field Immediate	mtfsfi (mtfsfi.)	crfD,IMM
Move to FPSCR Fields	mtfsf (mtfsf.)	FM,frB

Implementation Note—The architecture notes that, in some implementations, the Move to FPSCR Fields (**mtfsfx**) instruction may perform more slowly when only a portion of the fields are updated as opposed to all of the fields. This is not the case in the MPC603e.

2.3.4.2.6 Floating-Point Move Instructions

Floating-point move instructions copy data from one floating-point register to another. The floating-point move instructions do not modify the FPSCR. The CR update option in these instructions controls the placing of result status into CR1. Floating-point move instructions are listed in Table 2-20.

Table 2-20. Floating-Point Move Instructions

Name	Mnemonic	Operand Syntax
Floating Absolute Value	fabs (fabs.)	frD,frB
Floating Move Register	fmr (fmr.)	frD,frB
Floating Negate	fneg (fneg.)	frD,frB
Floating Negative Absolute Value	fnabs (fnabs.)	frD,frB

2.3.4.3 Load and Store Instructions

Load and store instructions are issued and translated in program order; however, the accesses can occur out of order. Synchronizing instructions are provided to enforce strict ordering. This section describes the load and store instructions of the MPC603e, which consist of the following:

- Integer load instructions
- Integer store instructions
- Integer load and store with byte-reverse instructions
- Integer load and store multiple instructions
- Integer load and store string instructions
- Floating-point load instructions
- Floating-point store instructions

2.3.4.3.1 Self-Modifying Code

When a processor modifies a memory location that may be contained in the instruction cache, software must ensure that memory updates are visible to the instruction fetching mechanism. This can be achieved by the following instruction sequence:

dcbst	update memory
sync	wait for update
icbi	remove (invalidate) copy in instruction cache
isync	remove copy in own instruction buffer

These operations are required because the data cache is a write-back cache. Since instruction fetching bypasses the data cache, changes to items in the data cache may not be reflected in memory until the fetch operations complete.

Special care must be taken to avoid coherency paradoxes in systems that implement unified secondary caches, and designers should carefully follow the guidelines for maintaining cache coherency that are provided in the VEA, and discussed in Chapter 5, “Cache Model and Memory Coherency,” in the *Programming Environments Manual*. Because the MPC603e does not broadcast the M bit for instruction fetches, external caches are subject to coherency paradoxes.

2.3.4.3.2 Integer Load and Store Address Generation

Integer load and store operations generate effective addresses using register indirect with immediate index mode, register indirect with index mode, or register indirect mode. See Section 2.3.2.3, “Effective Address Calculation.” Note that the MPC603e is optimized for load and store operations that are aligned on natural boundaries, and operations that are not naturally aligned may suffer performance degradation. Refer to Section 4.5.6.1, “Integer Alignment Exceptions.”

2.3.4.3.3 Register Indirect Integer Load Instructions

For integer load instructions, the byte, half word, word, or double word addressed by the EA is loaded into **rD**. Many integer load instructions have an update form, in which **rA** is updated with the generated effective address. For these forms, the EA is placed into **rA** and the memory element (byte, half word, word, or double word) addressed by EA is loaded into **rD**.

Implementation Note—In some implementations, the load half word algebraic instructions (**lha** and **lhax**) and the load with update (**lbzu**, **lbzux**, **lhzu**, **lhzux**, **lhau**, **lhaux**, **lwu**, and **lwux**) instructions may execute with greater latency than other types of load instructions. In the MPC603e, these instructions operate with the same latency as other load instructions.

Table 2-21 lists the integer load instructions.

Table 2-21. Integer Load Instructions

Name	Mnemonic	Operand Syntax
Load Byte and Zero	lbz	rD,d(rA)
Load Byte and Zero Indexed	lbzx	rD,rA,rB
Load Byte and Zero with Update	lbzu	rD,d(rA)
Load Byte and Zero with Update Indexed	lbzux	rD,rA,rB
Load Half Word Algebraic	lha	rD,d(rA)
Load Half Word Algebraic Indexed	lhax	rD,rA,rB
Load Half Word Algebraic with Update	lhau	rD,d(rA)
Load Half Word Algebraic with Update Indexed	lhaux	rD,rA,rB
Load Half Word and Zero	lhz	rD,d(rA)
Load Half Word and Zero Indexed	lhzx	rD,rA,rB
Load Half Word and Zero with Update	lhzu	rD,d(rA)
Load Half Word and Zero with Update Indexed	lhzux	rD,rA,rB
Load Word and Zero	lwz	rD,d(rA)
Load Word and Zero Indexed	lwzx	rD,rA,rB
Load Word and Zero with Update	lwzu	rD,d(rA)
Load Word and Zero with Update Indexed	lwzux	rD,rA,rB

2.3.4.3.4 Integer Store Instructions

For integer store instructions, the contents of **rS** are stored into the byte, half word, word, or double word in memory addressed by the effective address. Many store instructions have an update form, in which **rA** is updated with the EA. For these forms, the following rules apply:

- If **rA** ≠ 0, the EA is placed into **rA**.
- If **rS** = **rA**, the contents of **rS** are copied to the target memory element, then the generated EA is placed into **rA** (**rS**).

The MPC603e defines store with update instructions with **rA** = 0 and integer store instructions with the CR update option enabled (Rc field, bit 31, in the instruction encoding = 1) to be invalid forms. Table 2-22 provides a list of the integer store instructions for the MPC603e.

Table 2-22. Integer Store Instructions

Name	Mnemonic	Operand Syntax
Store Byte	stb	rS,d(rA)
Store Byte Indexed	stbx	rS,rA,rB
Store Byte with Update	stbu	rS,d(rA)
Store Byte with Update Indexed	stbux	rS,rA,rB

Table 2-22. Integer Store Instructions (continued)

Name	Mnemonic	Operand Syntax
Store Half Word	sth	rS,d(rA)
Store Half Word Indexed	sthx	rS,rA,rB
Store Half Word with Update	sthu	rS,d(rA)
Store Half Word with Update Indexed	sthux	rS,rA,rB
Store Word	stw	rS,d(rA)
Store Word Indexed	stwx	rS,rA,rB
Store Word with Update	stwu	rS,d(rA)
Store Word with Update Indexed	stwux	rS,rA,rB

2.3.4.3.5 Integer Load and Store with Byte-Reverse Instructions

Table 2-23 describes integer load and store with byte-reverse instructions. When used in a system operating with the default big-endian byte order, these instructions have the effect of loading and storing data in little-endian order. Likewise, when used in a system operating with little-endian byte order, these instructions have the effect of loading and storing data in big-endian order. For more information about big-endian and little-endian byte ordering, see “Byte Ordering” in Chapter 3, “Operand Conventions,” in the *Programming Environments Manual*.

Table 2-23. Integer Load and Store with Byte-Reverse Instructions

Name	Mnemonic	Operand Syntax
Load Half Word Byte-Reverse Indexed	lhbrx	rD,rA,rB
Load Word Byte-Reverse Indexed	lwbrx	rD,rA,rB
Store Half Word Byte-Reverse Indexed	sthbrx	rS,rA,rB
Store Word Byte-Reverse Indexed	stwbrx	rS,rA,rB

Implementation Note—In some implementations, load byte-reverse instructions (**lhbrx** and **lwbrx**) may have greater latency than other load instructions; however, these instructions operate with the same latency as other load instructions in the MPC603e.

2.3.4.3.6 Integer Load and Store Multiple Instructions

The integer load/store multiple instructions are used to move blocks of data to and from the GPRs. In some implementations, these instructions are likely to have greater latency and take longer to execute, perhaps much longer, than a sequence of individual load or store instructions that produce the same results.

Implementation Notes—The following describes the MPC603e implementation of the load/store multiple instruction:

- The load multiple and store multiple instructions may have operands that require memory accesses crossing a 4-Kbyte page boundary. As a result, these instructions may be interrupted by a DS1 exception associated with the address translation of the second page. In this case, the MPC603e performs some or all of the memory references from the first page, and none of the memory references from the second page before taking the exception. On return from the DS1 exception, the load or store multiple instruction will re-execute from the beginning. For additional information, refer to “DS1 Exception (0x00300)” in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.
- The PowerPC architecture defines the load multiple word (**lmw**) instruction with **rA** in the range of registers to be loaded as an invalid form. It defines the load multiple and store multiple instructions with misaligned operands (that is, the EA is not a multiple of four) to cause an alignment exception. The MPC603e defines the load multiple word (**lmw**) instruction with **rA** in the range of registers to be loaded as an invalid form.
- The PowerPC architecture describes some preferred instruction forms for the integer load and store multiple instructions that may perform better than other forms in some implementations. None of these preferred forms affect instruction performance in the MPC603e.
- When the MPC603e is operating with little-endian byte order, execution of a load or store multiple instruction causes the system alignment error handler to be invoked; see “Byte Ordering” in Chapter 3, “Operand Conventions,” in the *Programming Environments Manual* for more information. Table 2-24 lists the integer load and store multiple instructions for the MPC603e.

Table 2-24. Integer Load and Store Multiple Instructions

Name	Mnemonic	Operand Syntax
Load Multiple Word	lmw	rD,d(rA)
Store Multiple Word	stmw	rS,d(rA)

2.3.4.3.7 Integer Load and Store String Instructions

The integer load and store string instructions allow movement of data from memory to registers or from registers to memory without concern for alignment. These instructions can be used for a short move between arbitrary memory locations or to initiate a long move between misaligned memory fields.

When the MPC603e is operating with little-endian byte order, execution of a load or store string instruction causes the system alignment error handler to be invoked; see “Byte

Ordering” in Chapter 3, “Operand Conventions,” in the *Programming Environments Manual* for more information.

Table 2-25 lists the integer load and store string instructions.

Table 2-25. Integer Load and Store String Instructions

Name	Mnemonic	Operand Syntax
Load String Word Immediate	lswi	rD,rA,NB
Load String Word Indexed	lswx	rD,rA,rB
Store String Word Immediate	stswi	rS,rA,NB
Store String Word Indexed	stswx	rS,rA,rB

Load string and store string instructions may involve operands that are not word-aligned. As described in “Alignment Exception (0x00600)” in Chapter 6, “Exceptions,” in the *Programming Environments Manual*, a misaligned string operation suffers a performance penalty compared to a word-aligned operation of the same type.

When a string operation crosses a 4-Kbyte boundary, the instruction may be interrupted by a DSI exception associated with the address translation of the second page. In this case, the MPC603e performs some or all memory references from the first page and none from the second before taking the exception. On return from the DSI exception, the load or store string instruction will re-execute from the beginning. For more information, refer to “DSI Exception (0x00300)” in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.

Implementation Note—If **rA** is in the range of registers to be loaded for a Load String Word Immediate (**lswi**) instruction or if either **rA** or **rB** is in the range of registers to be loaded for a Load String Word Indexed (**lswx**) instruction, the PowerPC architecture defines the instruction to be of an invalid form. In addition, the **lswx** and **stswx** instructions that specify a string length of zero are defined to be invalid by the PowerPC architecture. However, none of these cases hold true for the MPC603e—the MPC603e treats these cases as valid forms.

2.3.4.3.8 Floating-Point Load and Store Address Generation

Floating-point load and store operations generate effective addresses using the register indirect with immediate index addressing mode and register indirect with index addressing mode (details are described below). Floating-point loads and stores are not supported for direct-store accesses. The use of the floating-point load and store operations for direct-store accesses results in a DSI exception.

2.3.4.3.9 Floating-Point Load Instructions

Separate floating-point load instructions are used for single-precision and double-precision operands. Because FPRs support only double-precision format, the FPU converts

single-precision data to double-precision format before loading the operands into the target FPR. This conversion is described fully in “Floating-Point Load Instructions” in Appendix D, “Floating-Point Models,” in the *Programming Environments Manual*.

Implementation Note—The PowerPC architecture defines load with update instructions with **rA = 0** as an invalid form; however, the MPC603e treats this case as a valid form.

Table 2-26 provides a list of the floating-point load instructions.

Table 2-26. Floating-Point Load Instructions

Name	Mnemonic	Operand Syntax
Load Floating-Point Double	lfd	frD,d(rA)
Load Floating-Point Double Indexed	lfdx	frD,rA,rB
Load Floating-Point Double with Update	lfdw	frD,d(rA)
Load Floating-Point Double with Update Indexed	lfdwx	frD,rA,rB
Load Floating-Point Single	lfs	frD,d(rA)
Load Floating-Point Single Indexed	lfsx	frD,rA,rB
Load Floating-Point Single with Update	lfsu	frD,d(rA)
Load Floating-Point Single with Update Indexed	lfsux	frD,rA,rB

2.3.4.3.10 Floating-Point Store Instructions

There are three basic forms of the store instruction—single-precision, double-precision, and integer. The integer form is supported by the optional **stfiwx** instruction. Because the FPRs support only double-precision format for floating-point data, the FPU converts double-precision data to single-precision format before storing the operands. The conversion steps are described in “Floating-Point Store Instructions” in Appendix D, “Floating-Point Models,” in the *Programming Environments Manual*.

Implementation Note—The PowerPC architecture defines store with update instructions with **rA = 0** as an invalid form; however, the MPC603e treats this case as valid.

Table 2-27 lists the floating-point store instructions.

Table 2-27. Floating-Point Store Instructions

Name	Mnemonic	Operand Syntax
Store Floating-Point as Integer Word Indexed	stfiwx	frS,rA,rB
Store Floating-Point Double	stfd	frS,d(rA)
Store Floating-Point Double Indexed	stfdx	frS,rA,rB
Store Floating-Point Double with Update	stfdw	frS,d(rA)
Store Floating-Point Double with Update Indexed	stfdwx	frS,rA,rB
Store Floating-Point Single	stfs	frS,d(rA)

Table 2-27. Floating-Point Store Instructions (continued)

Name	Mnemonic	Operand Syntax
Store Floating-Point Single Indexed	stfsx	frS,rA,rB
Store Floating-Point Single with Update	stfsu	frS,d(rA)
Store Floating-Point Single with Update Indexed	stfsux	frS,rA,rB

2.3.4.4 Branch and Flow Control Instructions

Branch instructions are executed by the branch processing unit (BPU). The BPU receives branch instructions from the fetch unit and performs CR look-ahead operations on conditional branches to resolve them early, achieving the effect of a zero-cycle branch in many cases.

Some branch instructions can redirect instruction execution conditionally based on the value of bits in the CR. When the branch processor encounters one of these instructions, it scans the execution pipelines to determine whether an instruction in progress may affect the particular CR bit. If no interlock is found, the branch can be resolved immediately by checking the bit in the CR and taking the action defined for the branch instruction.

If an interlock is detected, the branch is considered unresolved and the direction of the branch is predicted using static branch prediction as described in “Conditional Branch Control” in Chapter 4, “Addressing Modes and Instruction Set Summary,” in the *Programming Environments Manual*. The interlock is monitored while instructions are fetched for the predicted branch. When the interlock is cleared, the branch processor determines whether the prediction was correct, based on the value of the CR bit. If the prediction is correct, the branch is considered completed and instruction fetching continues. If the prediction is incorrect, the fetched instructions are purged, and instruction fetching continues along the alternate path. See Chapter 8, “Instruction Timing,” in the *Programming Environments Manual* for more information about how branches are executed.

2.3.4.4.1 Branch Instruction Address Calculation

Branch instructions can change the instruction sequence. Instruction addresses are always assumed to be word aligned; the processor ignores the two low-order bits of the generated branch target address.

Branch instructions compute the effective address (EA) of the next instruction address using the following addressing modes:

- Branch relative
- Branch conditional to relative address
- Branch to absolute address

- Branch conditional to absolute address
- Branch conditional to link register
- Branch conditional to count register

2.3.4.4.2 Branch Instructions

Table 2-28 lists the branch instructions provided by the processors that implement the PowerPC architecture. To simplify assembly language programming, a set of simplified mnemonics and symbols is provided for the most frequently used forms of branch conditional, compare, trap, rotate and shift, and certain other instructions. See Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* for a list of simplified mnemonic examples.

Table 2-28. Branch Instructions

Name	Mnemonic	Operand Syntax
Branch	b (ba bl bla)	target_addr
Branch Conditional	bc (bca bcl bcla)	BO,BI,target_addr
Branch Conditional to Count Register	bcctr (bcctrl)	BO,BI
Branch Conditional to Link Register	bclr (bclrl)	BO,BI

2.3.4.4.3 Condition Register Logical Instructions

Condition register logical instructions, shown in Table 2-29, and the Move Condition Register Field (**mcrf**) instruction are also defined as flow control instructions, although they are executed by the system register unit (SRU). Most instructions executed by the SRU are completion-serialized to maintain system state; that is, the instruction is held for execution in the SRU until all prior instructions issued have completed.

Table 2-29. Condition Register Logical Instructions

Name	Mnemonic	Operand Syntax
Condition Register AND	crand	crbD,crbA,crbB
Condition Register AND with Complement	crandc	crbD,crbA,crbB
Condition Register Equivalent	creqv	crbD,crbA,crbB
Condition Register NAND	crnand	crbD,crbA,crbB
Condition Register NOR	crnor	crbD,crbA,crbB
Condition Register OR	cror	crbD,crbA,crbB
Condition Register OR with Complement	crorc	crbD,crbA,crbB
Condition Register XOR	crxor	crbD,crbA,crbB
Move Condition Register Field	mcrf	crfD,crfS

Note that if the LR update option is enabled for any of these instructions, these forms of the instructions are invalid in the MPC603e.

2.3.4.5 Trap Instructions

The trap instructions shown in Table 2-30 are provided to test for a specified set of conditions. If any of the conditions tested by a trap instruction are met, the system trap handler is invoked. If the tested conditions are not met, instruction execution continues normally.

Table 2-30. Trap Instructions

Name	Mnemonic	Operand Syntax
Trap Word	tw	TO,rA,rB
Trap Word Immediate	twi	TO,rA,SIMM

See Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* for a complete set of simplified mnemonics.

2.3.4.6 Processor Control Instructions

UISA-level processor control instructions are used to read from and write to the condition register (CR).

2.3.4.6.1 Move To/From Condition Register Instructions

Table 2-31 lists the instructions provided by the MPC603e for reading from or writing to the CR.

Table 2-31. Move To/From Condition Register Instructions

Name	Mnemonic	Operand Syntax
Move from Condition Register	mfcr	rD
Move to Condition Register Fields	mtcrf	CRM,rS
Move to Condition Register from XER	mcrxr	crfD

2.3.4.7 Memory Synchronization Instructions—UISA

Memory synchronization instructions control the order in which memory operations are completed with respect to asynchronous events and the order in which memory operations are seen by other processors or memory access mechanisms. See Chapter 3, “Instruction and Data Cache Operation,” for additional information about these instructions and about related aspects of memory synchronization.

The **sync** instruction delays execution of subsequent instructions until previous instructions have completed to the point that they can no longer cause an exception and until all

previous memory accesses are performed globally; the **sync** operation is not broadcast onto the MPC603e bus interface. Additionally, all load and store cache/bus activities initiated by prior instructions are completed. Touch load operations (**dcbt** and **dcbtst**) are required to complete at least through address translation but are not required to complete on the bus.

The functions performed by the **sync** instruction normally take a significant amount of time to complete; as a result, frequent use of this instruction may adversely affect performance. In addition, the number of cycles required to complete a **sync** instruction depends on system parameters and on the processor's state when the instruction is issued.

The proper paired use of the **Iwarx** and **stwcx.** instructions allows programmers to emulate common semaphore operations such as test and set, compare and swap, exchange memory, and fetch and add. Examples of these operations can be found in Appendix E, "Synchronization Programming Examples," in the *Programming Environments Manual*. Typically, the **Iwarx** instruction should be paired with an **stwcx.** instruction with the same effective address used for both instructions of the pair. Note that the reservation granularity is 32 bytes.

The concept behind the use of the **Iwarx** and **stwcx.** instructions is that a processor may load a semaphore from memory, compute a result based on the value of the semaphore, and conditionally store it back to the same location (only if that location has not been modified since it was first read), and determine if the store was successful. The conditional store is performed, based on the existence of a reservation established by the preceding **Iwarx** instruction. If the reservation exists when the store is executed, the store is performed which sets a bit in the CR. If the reservation does not exist when the store is executed, the target memory location is not modified and a bit is cleared in the CR.

If the store was successful, the sequence of instructions from the read of the semaphore to the store that updated the semaphore appear to have been executed atomically (that is, no other processor or mechanism modified the semaphore location between the read and the update), thus providing the equivalent of a real atomic operation. However, in reality, other processors may have read from the location during this operation. In the MPC603e, the reservations are made on behalf of aligned 32-byte sections of the memory address space.

The **Iwarx** and **stwcx.** instructions require the EA to be aligned. Exception handling software should not attempt to emulate a misaligned **Iwarx** or **stwcx.** instruction, because there is no correct way to define the address associated with the reservation.

In general, the **Iwarx** and **stwcx.** instructions should be used only in system programs, which can be invoked by application programs as needed.

At most, one reservation exists simultaneously on any processor. The address associated with the reservation can be changed by a subsequent **Iwarx** instruction. The conditional store is performed, based on the existence of a reservation established by the preceding **Iwarx** regardless of whether the address generated by the **Iwarx** matches that generated by

the **stwcx.** instruction. A reservation held by the processor is cleared by one of the following:

- Executing an **stwcx.** instruction to any address
- Attempt by some other device to modify a location in the reservation granularity (32 bytes)

The **Iwarx** and **stwcx.** instructions to write-through memory do not cause a DSI exception.

Table 2-32 lists the UISA memory synchronization instructions for the MPC603e.

Table 2-32. Memory Synchronization Instructions—UISA

Name	Mnemonic	Operand Syntax
Load Word and Reserve Indexed	Iwarx	rD,rA,rB
Store Word Conditional Indexed	stwcx.	rS,rA,rB
Synchronize	sync	—

2.3.5 PowerPC VEA Instructions

The VEA describes the semantics of the memory model that can be assumed by software processes, and includes descriptions of the cache model, cache-control instructions, address aliasing, and other related issues.

2.3.5.1 Processor Control Instructions

The VEA defines the Move from Time Base (**mftb**) instruction for reading the contents of the time base register. The **mftb** is a user-level instruction, as shown in Table 2-33.

Table 2-33. Move From Time Base Instruction

Name	Mnemonic	Operand Syntax
Move from Time Base	mftb	rD, TBR

Simplified mnemonics are provided for the **mftb** instruction so it can be coded with the TBR name as part of the mnemonic rather than requiring it to be coded as an operand. The **mftb** instruction serves as both a basic and simplified mnemonic. Assemblers recognize an **mftb** mnemonic with two operands as the basic form, and an **mftb** mnemonic with one operand as the simplified form. Simplified mnemonics are also provided for Move from Time Base Upper (**mftbu**), a variant of the **mftb** instruction rather than of **mfsp**. The MPC603e ignores the extended opcode differences between **mftb** and **mfsp** by ignoring bit 25 of both instructions and treating them identically. Refer to Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual*.

2.3.5.2 Memory Synchronization Instructions—VEA

Memory synchronization instructions control the order in which memory operations are performed with respect to asynchronous events, and the order in which memory operations are seen by other processors or memory access mechanisms. See Chapter 3, “Instruction and Data Cache Operation,” for additional information about these instructions and about related aspects of memory synchronization.

Implementation Notes—The following describes how the MPC603e handles memory synchronization in the VEA:

- The Instruction Synchronize (**isync**) instruction causes the MPC603e to discard all prefetched instructions, wait for any preceding instructions to complete, and then branch to the next sequential instruction (having the effect of clearing the pipeline behind the **isync** instruction).
- The Enforce In-Order Execution of I/O (**eieio**) instruction is used to ensure memory reordering of noncacheable memory access. Because the MPC603e does not reorder noncacheable memory accesses, the **eieio** instruction is treated as a no-op.

Table 2-34 lists the VEA memory synchronization instructions for the MPC603e.

Table 2-34. Memory Synchronization Instructions—VEA

Name	Mnemonic	Operand Syntax
Enforce In-Order Execution of I/O	eieio	—
Instruction Synchronize	isync	—

2.3.5.3 Memory Control Instructions—VEA

Memory control instructions include the following types:

- Cache management instructions
- Segment register manipulation instructions
- Translation lookaside buffer management instructions

This section describes the user-level cache management instructions defined by the VEA. See Section 2.3.6.3, “Memory Control Instructions—OEA,” for information about supervisor-level cache, segment register manipulation, and translation lookaside buffer management instructions.

The instructions listed in Table 2-35 provide user-level programs the ability to manage on-chip caches when they exist.

Table 2-35. User-Level Cache Instructions

Name	Mnemonic	Operand Syntax
Data Cache Block Flush	dcbf	rA,rB
Data Cache Block Set to Zero	dcbz	rA,rB
Data Cache Block Store	dcbst	rA,rB
Data Cache Block Touch	dcbt	rA,rB
Data Cache Block Touch for Store	dcbtst	rA,rB
Instruction Cache Block Invalidate	icbi	rA,rB

As with other memory-related instructions, the effect of the cache management instructions on memory are weakly ordered. If the programmer needs to ensure that cache or other instructions have been performed with respect to all other processors and system mechanisms, a **sync** instruction must be placed in the program following those instructions.

Note that when data address translation is disabled (MSR[DR] = 0), the Data Cache Block Set to Zero (**dcbz**) instruction allocates a cache block in the cache and may not verify that the physical address is valid. If a cache block is created for an invalid physical address, a machine check condition may result when an attempt is made to write that cache block back to memory. The cache block could be written back as a result of the execution of an instruction that causes a cache miss and the invalid addressed cache block is the target for replacement or a Data Cache Block Store (**dcbst**) instruction.

Note that any cache control instruction that generates an effective address that corresponds to a direct-store segment (SR[T] = 1) is treated as a no-op.

Table 2-35 lists the cache instructions that are accessible to user-level programs.

Note that incoherency may occur if a write-through store is followed by a **dcbz** instruction that is, in turn, followed by a snoop, all to the same cache block. This occurs when the logical address for the **dcbz** and the write-through store are different but aliased to the same physical page.

To avoid potential adverse effects, **dcbz** should not address write-through memory that can be accessed through multiple logical addresses. Explicit store instructions that write all zeroes should be used instead.

Note that broadcasting a sequence of **dcbz** instructions may cause snoop accesses to be retried indefinitely, which may cause the snoop originator to time out or the snooped transaction to not complete. This can be avoided by disabling the broadcasting of **dcbz** by marking the memory space being addressed by the **dcbz** instruction as not global in the BAT or PTE.

Note that incoherency may occur if the following sequence of accesses hits the same cache block: a write-through, a **dcbz** instruction, a snoop. This occurs when the logical address for the **dcbz** and the write-through store are different but aliased to the same physical page.

2.3.5.4 External Control Instructions

The **eciwx** instruction provides an alternative way to map special devices. The MMU translation of the EA is not used to select the special device, as it is used in loads and stores. Rather, it is used as an address operand that is passed to the device over the address bus. Four other signals (the burst and size signals on the 60x bus) are used to select the device; these four signals output the 4-bit resource ID (RID) field in the EAR register. The **eciwx** instruction also loads a word from the data bus that is output by the special device. Executing these instructions when **MSR[DR] = 0** causes a programming error, and the physical address on the bus is undefined. Executing these instructions to a direct-store segment causes a DSI exception. The external control instructions are listed in Table 2-36.

Table 2-36. External Control Instructions

Name	Mnemonic	Operand Syntax
External Control In Word Indexed	eciwx	rD,rA,rB
External Control Out Word Indexed	ecowx	rS,rA,rB

2.3.6 PowerPC OEA Instructions

The OEA includes the structure of the memory management model, supervisor-level registers, and the exception model.

2.3.6.1 System Linkage Instructions

This section describes the system linkage instructions (see Table 2-37). The **sc** instruction is a user-level instruction that permits a user program to call on the system to perform a service and causes the processor to take an exception. The Return from Interrupt (**rfi**) instruction is a supervisor-level instruction that is useful for returning from an exception handler.

Table 2-37. System Linkage Instructions

Name	Mnemonic	Operand Syntax
Return from Interrupt	rfi	—
System Call	sc	—

2.3.6.2 Processor Control Instructions—OEA

Processor control instructions are used to read from and write to the condition register (CR), machine state register (MSR), and special-purpose registers (SPRs), and to read from the time base register (TBU or TBL).

2.3.6.2.1 Move To/From Machine State Register Instructions

Table 2-38 lists the instructions provided by the MPC603e for reading from or writing to the MSR.

Table 2-38. Move To/From Machine State Register Instructions

Name	Mnemonic	Operand Syntax
Move from Machine State Register	mfmsr	rD
Move to Machine State Register	mtmsr	rS

2.3.6.2.2 Move To/From Special-Purpose Register Instructions

Simplified mnemonics are provided for the **mtspr** and **mfspr** instructions so they can be coded with the SPR name as part of the mnemonic rather than as a numeric operand. See Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* for simplified mnemonic examples. The **mtspr** and **mfspr** instructions are shown in Table 2-39.

Table 2-39. Move To/From Special-Purpose Register Instructions

Name	Mnemonic	Operand Syntax
Move from Special-Purpose Register	mfspr	rD,SPR
Move to Special-Purpose Register	mtspr	SPR,rS

For **mtspr** and **mfspr** instructions, the SPR number coded in assembly language does not appear directly as a 10-bit binary number in the instruction. The number coded is split into two 5-bit halves that are reversed in the instruction encoding, with the high-order 5 bits appearing in bits 16–20 of the instruction encoding and the low-order 5 bits in bits 11–15.

If the SPR field contains any value other than one of the values shown in Table 2-40, either the program exception handler is invoked or the results are boundedly undefined.

Table 2-40. Implementation-Specific SPR Encodings (mfspr)

SPR ¹			Register Name
Decimal	spr[5–9]	spr[0–4]	
976	11110	10000	DMISS
977	11110	10001	DCMP
978	11110	10010	HASH1

Table 2-40. Implementation-Specific SPR Encodings (mfspr) (continued)

SPR ¹			Register Name
Decimal	spr[5–9]	spr[0–4]	
979	11110	10011	HASH2
980	11110	10100	IMISS
981	11110	10101	ICMP
982	11110	10110	RPA
1008	11111	10000	HID0
1009	11111	10001	HID1
1010	11111	10010	IABR

¹ Note that the order of the two 5-bit halves of the SPR number is reversed compared with actual instruction coding.

For **mtspr** and **mfspr** instructions, the SPR number coded in assembly language does not appear directly as a 10-bit binary number in the instruction. The number coded is split into two 5-bit halves that are reversed in the instruction, with the high-order 5 bits appearing in bits 16–20 of the instruction and the low-order 5 bits in bits 11–15.

Implementation Note—The MPC603e ignores the extended opcode differences between **mftb** and **mfspr** by ignoring TB[25] and treating both instructions identically.

2.3.6.3 Memory Control Instructions—OEA

This section describes memory control instructions, which include the following types:

- Cache management instructions
- Segment register manipulation instructions
- Translation lookaside buffer management instructions

2.3.6.3.1 Supervisor-Level Cache Management Instruction

The supervisor-level cache management instruction in the PowerPC architecture, **dcbi**, should not be used on the MPC603e. The user-level **dcbf** instruction, described in Section 2.3.5.3, “Memory Control Instructions—VEA,” and in Section 3.7, “Cache Control Instructions,” should be used when the program needs to invalidate cache blocks. Note that the **dcbf** instruction causes modified blocks to be flushed to system memory if they are the target of a **dcbf** instruction, whereas, by definition in the PowerPC architecture, the **dcbi** instruction only invalidates modified blocks.

2.3.6.3.2 Segment Register Manipulation Instructions

The instructions listed in Table 2-41 provide access to the segment registers for the MPC603e. These instructions operate completely independent of the MSR[IR] and MSR[DR] bit settings. Refer to “Synchronization Requirements for Special Registers and TLBs” in Chapter 2, “Register Set,” in the *Programming Environments Manual* for serialization requirements and other recommended precautions to observe when manipulating the segment registers.

Table 2-41. Segment Register Manipulation Instructions

Name	Mnemonic	Operand Syntax
Move from Segment Register	mfsr	rD,SR
Move from Segment Register Indirect	mfsrin	rD,rB
Move to Segment Register	mtsrr	SR,rS
Move to Segment Register Indirect	mtsrin	rS,rB

2.3.6.3.3 Translation Lookaside Buffer Management Instructions

The address translation mechanism is defined in terms of segment descriptors and page table entries (PTEs) used by the processors to locate the effective-to-physical address mapping for a particular access. The PTEs reside in page tables in memory. As defined for 32-bit implementations by the PowerPC architecture, segment descriptors reside in 16 on-chip segment registers.

Implementation Note—The MPC603e provides the ability to invalidate a TLB entry. The TLB Invalidate Entry (**tlbie**) instruction invalidates the TLB entry indexed by the EA, and operates on both the instruction and data TLBs simultaneously invalidating four TLB entries (both sets in each TLB). The index corresponds to bits 15–19 of the EA. To invalidate all entries within both TLBs, 32 **tlbie** instructions should be issued, incrementing this field by one each time.

The MPC603e provides two implementation-specific instructions (**tlbld** and **tlbli**) that are used by software table search operations following TLB misses to load TLB entries on-chip.

For more information on **tlbld** and **tlbli** refer to Section 2.3.8, “Implementation-Specific Instructions.”

Note that the **tlbia** instruction is not implemented on the MPC603e.

Refer to Chapter 5, “Memory Management” for more information about the TLB operations for the MPC603e. Table 2-42 lists the TLB instructions.

Table 2-42. Translation Lookaside Buffer Management Instructions

Name	Mnemonic	Operand Syntax
Load Data TLB Entry	tlbld	rB
Load Instruction TLB Entry	tblli	rB
TLB Invalidate Entry	tlbie	rB
TLB Synchronize	tlbsync	—

Because the presence and exact semantics of the translation lookaside buffer management instructions is implementation-dependent, system software should incorporate uses of the instructions into subroutines to maximize compatibility with programs written for other processors.

For more information on the PowerPC instruction set, refer to Chapter 4, “Addressing Modes and Instruction Set Summary,” and Chapter 8, “Instruction Set,” in the *Programming Environments Manual*.

2.3.7 Recommended Simplified Mnemonics

To simplify assembly language programs, a set of simplified mnemonics is provided for some of the most frequently used operations (such as no-op, load immediate, load address, move register, and complement register). PowerPC compliant assemblers provide the simplified mnemonics listed in “Recommended Simplified Mnemonics” in Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual* and listed with some of the instruction descriptions in this chapter. Programs written to be portable across the various assemblers for the PowerPC architecture should not assume the existence of mnemonics not described in this document.

For a complete list of simplified mnemonics, see Appendix F, “Simplified Mnemonics,” in the *Programming Environments Manual*.

2.3.8 Implementation-Specific Instructions

This section provides a detailed look at the two MPC603e implementation-specific instructions—**tlbld** and **tblli**.

tlbld

Load Data TLB Entry

tlbld

Integer Unit

tlbld**rB** Reserved

31	0 0 0 0	0 0 0 0	B	978	0
0	5 6	10 11	15 16	20 21	30 31

EA ← (rB)

TLB entry created from DCMP and RPA

DTLB entry selected by EA[15-19] and SRR1[WAY] ← created TLB entry

The EA is the contents of rB. The **tlbld** instruction loads the contents of the data PTE compare (DCMP) and required physical address (RPA) registers into the first word of the selected data TLB entry. The specific DTLB entry to be loaded is selected by the EA and SRR1[WAY] bit.

The **tlbld** instruction should only be executed when address translation is disabled (MSR[IR] = 0 and MSR[DR] = 0).

Note that it is possible to execute the **tlbld** instruction when address translation is enabled; however, extreme caution should be used in doing so. If data address translation is set (MSR[DR] = 1) **tlbld** must be preceded by a **sync** instruction and succeeded by a context synchronizing instruction.

Also, note that care should be taken to avoid modification of the instruction TLB entries that translate current instruction prefetch addresses.

This is a supervisor-level instruction; it is also a MPC603e-specific instruction, and not part of the PowerPC instruction set.

Other registers altered:

- None

tlbli

Load Instruction TLB Entry

tlbld**rB****tlbli**

Integer Unit

Reserved

31	0 0 0 0	0 0 0 0	B	1010	0
0	5 6	10 11	15 16	20 21	30 31

EA ← (**rB**)

TLB entry created from ICMP and RPA

ITLB entry selected by EA[15-19] and SRR1[WAY] ← created TLB entry

The EA is the contents of **rB**. The **tlbli** instruction loads the contents of the instruction PTE compare (ICMP) and required physical address (RPA) registers into the first word of the selected instruction TLB entry. The specific ITLB entry to be loaded is selected by the EA and SRR1[WAY] bit.

The **tlbli** instruction should only be executed when address translation is disabled (MSR[IR] = 0 and MSR[DR] = 0).

Note that it is possible to execute the **tlbld** instruction when address translation is enabled; however, extreme caution should be used in doing so. If instruction address translation is set (MSR[IR] = 1), **tlbli** must be followed by a context synchronizing instruction such as **isync** or **rfi**.

Also, note that care should be taken to avoid modification of the instruction TLB entries that translate current instruction prefetch addresses.

This is a supervisor-level instruction; it is also a MPC603e-specific instruction, and not part of the PowerPC instruction set.

Other registers altered:

- None



Chapter 3

Instruction and Data Cache Operation

The MPC603e microprocessor provides two 16-Kbyte, four-way set-associative caches to allow the registers and execution units rapid access to instructions and data. Both the instruction and data caches are tightly coupled to the MPC603e bus interface unit (BIU) to allow efficient access to the system memory controller and other bus masters. The MPC603e load/store unit (LSU) is also directly coupled to the data cache to allow the efficient movement of data to and from the general-purpose and floating-point registers.

This chapter describes the organization of the cache, cache coherency protocols, cache control instructions, and various cache operations. It describes the interaction between the caches, the load/store unit, the instruction unit, and the memory subsystem.

Note that in this chapter, the term multiprocessor is used in the context of maintaining cache coherency. These multiprocessor devices could be actual processors or other devices that can access system memory, maintain their own caches, and function as bus masters requiring cache coherency.

Both the instruction and data caches have 32-byte blocks, and data cache blocks can be snooped or cast out when the cache block is reloaded. The data cache is designed to adhere to a write-back policy, but the MPC603e allows control of cacheability, write-back policy, and memory coherency at the page and block level. Both caches use a least recently used (LRU) replacement policy. Burst fill operations to the caches result from cache misses, or in the case of the data cache, cache block write-back operations to memory. Note that in the PowerPC architecture, the term ‘cache block’, or simply ‘block’ when used in the context of cache implementations, refers to the unit of memory at which coherency is maintained. For the MPC603e, the block size is equivalent to the eight-word cache line. This value may be different for other implementations that support the PowerPC architecture.

The data cache is configured as 128 sets of 4 blocks. Each block consists of 32 bytes, 2 state bits, and an address tag. The 2 state bits implement the three-state MEI (modified/exclusive/invalid) protocol, a coherent subset of the standard four-state MESI protocol. Cache coherency is enforced by on-chip bus snooping logic. Since the MPC603e data cache tags are single-ported, a simultaneous load or store and snoop access represent a resource contention. The snoop access is given first access to the tags. Load or store operations can be performed to the cache on the clock cycle immediately following a snoop access if the

snoop misses; snoop hits may block the data cache for two or more cycles, depending on whether a copy back to main memory is required.

The instruction cache also consists of 128 sets of 4 blocks, and each block consists of 32 bytes, an address tag, and a valid bit. The instruction cache is only written as a result of a block fill operation on a cache miss. In the PID7t-603e, the instruction cache is blocked only until the critical load completes. The PID7t-603e supports instruction fetching from other instruction cache lines following the forwarding of the critical first double word of a cache line load operation. Successive instruction fetches from the cache line being loaded are forwarded, and accesses to other instruction cache lines can proceed during the cache line load operation. The instruction cache is not snooped, and cache coherency must be maintained by software. A fast hardware invalidation capability is provided to support cache maintenance.

The load/store unit provides the data transfer interface between the data cache and the GPRs and FPRs. The LSU provides all logic required to calculate effective addresses, handle data alignment to and from the data cache, and provides sequencing for load and store string and multiple operations. As shown in Figure 1-1, the caches provide a 64-bit interface to the instruction fetcher and LSU. Write operations to the data cache can be performed on a byte, half-word, word, or double-word basis.

The MPC603e bus interface unit (BIU) receives requests for bus operations from the instruction and data caches, and executes the operations according to the MPC603e bus protocol. The BIU provides address queues, prioritization, and bus control logic. The BIU also captures snoop addresses for data cache, address queue, and memory reservation (**lwarx** and **stwcx.** instruction) operations. The BIU also contains a touch load address buffer used for address compares during load or store operations. All the data for the corresponding address queues (load and store data queues) is located in the data cache. The data queues are considered temporary storage for the cache and not part of the BIU.

On a cache miss, the MPC603e cache blocks are loaded in 4 beats of 64 bits each when the MPC603e is configured with a 64-bit data bus; when the MPC603e is configured with a 32-bit bus, cache block loads are performed with 8 beats of 32 bits each. The burst load is performed as critical double word first. The data cache is blocked to internal accesses until the load completes; the instruction cache allows sequential fetching during a cache block load. In the PID7t-603e, the critical double word is simultaneously written to the cache and forwarded to the requesting unit, thus minimizing stalls due to load delays. Note that the cache being filled cannot be accessed internally until the fill completes.

When address translation is enabled, the memory access is performed under the control of the page table entry used to translate the effective address. Each page table entry and BAT contains four mode control bits, W, I, M, and G, that specify the storage mode for all accesses translated using that particular page table entry. The W (write-through) and I (caching-inhibited) bits control how the processor executing the access uses its own cache. The M (memory coherence) bit specifies whether the processor executing the access must

use the MEI (modified, exclusive, or invalid) cache coherence protocol to ensure all copies of the addressed memory location are kept consistent. The G (guarded memory) bit controls whether out-of-order data and instruction fetching is permitted.³

The MPC603e maintains data cache coherency in hardware by coordinating activity between the data cache, memory system, and bus interface logic. As bus operations are performed on the bus by other bus masters, the MPC603e bus snooping logic monitors the addresses that are referenced. These addresses are compared with the addresses resident in the data cache. If there is a snoop hit, the MPC603e bus snooping logic responds to the bus interface with the appropriate snoop status (for example, an ARTRY). Additional snoop action may be forwarded to the cache as a result of a snoop hit in some cases (a cache push of modified data or cache block invalidation).

The MPC603e supports a fully-coherent 4-Gbyte physical memory address space. Bus snooping is used to drive the MEI three-state cache-coherency protocol that ensures the coherency of global memory with respect to the processor's cache. See Section 3.6.1, "MEI State Definitions."

This chapter describes the organization of the MPC603e on-chip instruction and data caches, the MEI cache coherency protocol, cache control instructions, various cache operations, and the interaction between the cache, LSU, and BIU. PID7t-603e specific information is noted where applicable.

3.1 Instruction Cache Organization and Control

The instruction fetcher accesses the instruction cache frequently in order to sustain the high throughput provided by the six-entry instruction queue.

3.1.1 Instruction Cache Organization

The instruction cache organization is shown in Figure 3-1. Each cache block contains eight contiguous words from memory that are loaded from an eight-word boundary (that is, bits A27–A31 of the effective addresses are zero); thus, a cache block never crosses a page boundary. Misaligned accesses across a page boundary can incur a performance penalty.

Note that address bits A20–A26 provide an index to select a set. Bits A27–A31 select a byte within a block. The tags consists of bits PA0–PA19. Address translation occurs in parallel, such that higher-order bits (the tag bits in the cache) are physical. Note that the replacement algorithm is strictly an LRU algorithm; that is, the least recently used block is filled with new instructions on a cache miss.

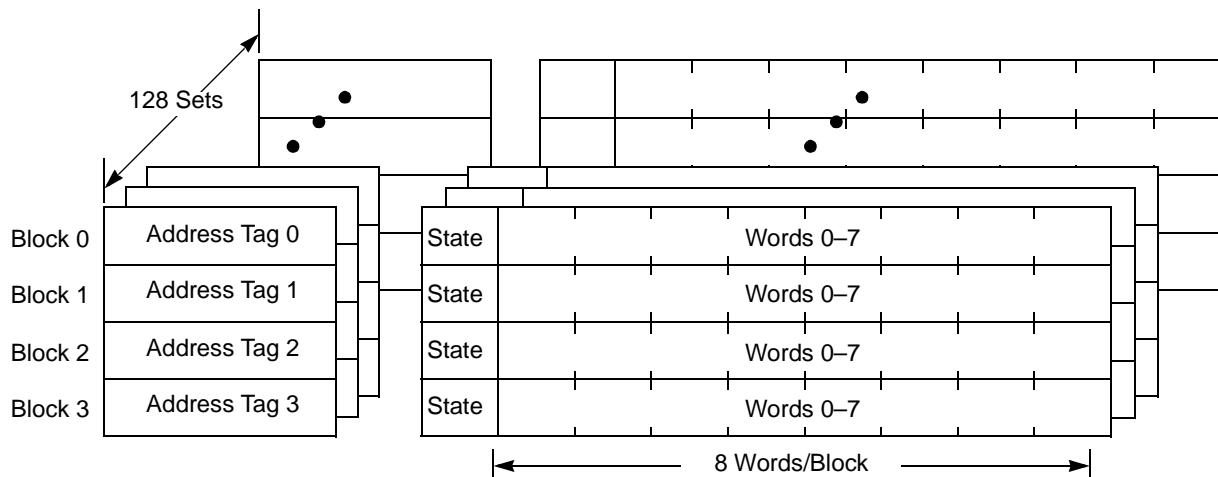


Figure 3-1. Instruction Cache Organization

3.1.2 Instruction Cache Fill Operations

The MPC603e instruction cache blocks are loaded in four 64-bit beats, with the critical double word loaded first. The instruction cache allows sequential fetching during a cache block load. On a cache miss, the critical and following double words read from memory are simultaneously written to the instruction cache and forwarded to the dispatch queue, thus minimizing stalls due to cache fill latency. There is no snooping of the instruction cache. In the PID7t-603e, the critical double word is simultaneously written to the cache and forwarded to the requesting unit, thus minimizing stalls due to load delays.

3.1.3 Instruction Cache Control

In addition to instruction cache control instructions, the MPC603e provides several HID0 bits to control invalidating, disabling, and locking the instruction cache. The WIMG bits in the page tables and the IBATs also affect the cacheability of pages and whether the pages are considered guarded.

3.1.3.1 Instruction Cache Invalidation

Although the MPC603e instruction cache is automatically invalidated during a power-on or hard reset, asserting \overline{SRESET} does not invalidate the instruction cache. Software can invalidate the contents of the instruction cache using the instruction cache flash invalidate control bit, HID0[ICFI]. Flash invalidation of the instruction cache is accomplished by setting and clearing ICFI with two consecutive **mtspr[HID0]** instructions.

3.1.3.2 Instruction Cache Disabling

The instruction cache may be disabled through the use of the instruction cache enable control bit, HID0[ICE]. When the instruction cache is in the disabled state, the cache tag

state bits are ignored and all accesses are propagated to the bus as single-beat transactions. The ICE bit is cleared during a power-on reset, causing the instruction cache to be disabled. To prevent the cache from being enabled or disabled in the middle of a data access, an **isync** instruction should be issued before changing the value of ICE.

3.1.3.3 Instruction Cache Locking

The contents of instruction cache may be locked through the use of HID0[ILOCK]. A locked instruction cache supplies instructions normally on a cache hit, but cache misses are treated as cache-inhibited accesses. The cache inhibited (\overline{CI}) signal is asserted if a cache access misses into a locked cache. The setting of the ILOCK bit must be preceded by an **isync** instruction to prevent the instruction cache from being locked during an instruction access.

3.2 Data Cache Organization and Control

The LSU transfers data between the data cache and the GPRs and FPRs and provides buffers for load and store bus operations. The data cache also provides storage for the cache tags required for memory coherency and performs the cache block replacement LRU function.

3.2.1 Data Cache Organization

The organization of the data cache is shown in Figure 3-2. Each cache block contains eight contiguous words from memory that are loaded from an eight-word boundary (that is, bits A27–A31 of the effective addresses are zero); thus, a cache block never crosses a page boundary. Misaligned accesses across a page boundary can incur a performance penalty.

Note that bits A20–A26 provide an index to select a set. Bits A27–A31 select a byte within a block. The tags consists of bits PA0–PA19. Address translation occurs in parallel, such that higher-order bits (the tag bits in the cache) are physical. Note that the replacement algorithm is strictly an LRU algorithm; that is, the least recently used block is filled with new data on a cache miss.

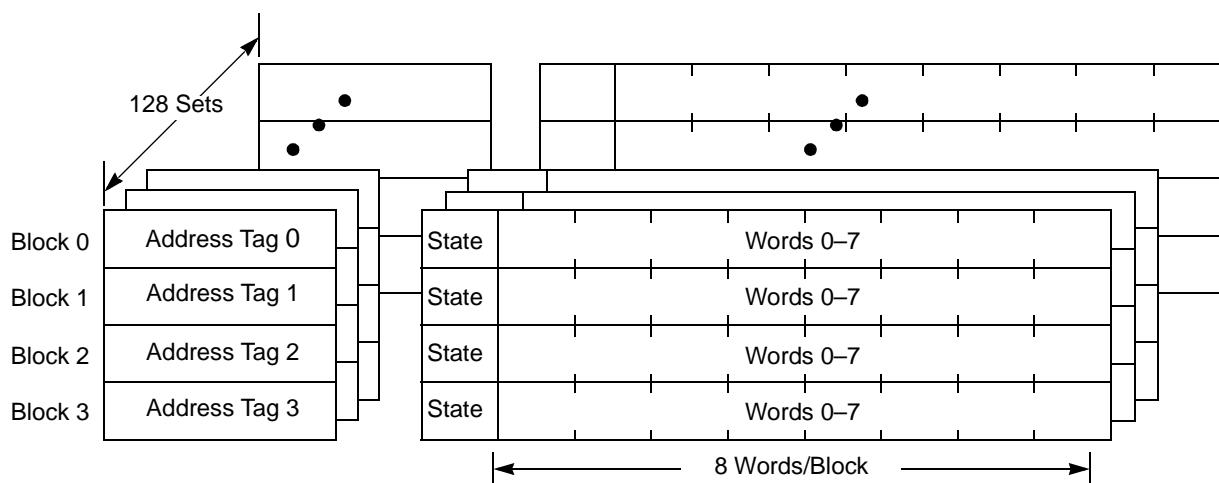


Figure 3-2. Data Cache Organization

3.2.2 Data Cache Fill Operations

When the MPC603e is configured with a 64-bit data bus, cache blocks are loaded in 4 beats of 64 bits each. When the MPC603e is configured with a 32-bit bus, cache block loads are performed with 8 beats of 32 bits each. The burst load is performed as critical double word first. The data cache is blocked to internal accesses until the load completes. In the PID7t-603e, the critical double word is simultaneously written to the cache and forwarded to the requesting unit, thus minimizing stalls due to load delays.

3.2.3 Data Cache Control

The MPC603e provides several means of data cache control through the use of the WIMG bits in the page tables, control bits in the HID0 register, and user- and supervisor-level cache control instructions. While memory page level cache control is provided by the WIMG bits, the on-chip data cache can be invalidated, disabled, locked, or broadcast by the control bits in the HID0 register described in this section. (Note that user- and supervisor-level are referred to as problem and privileged state, respectively, in the architecture specification.)

3.2.3.1 Data Cache Invalidation

While the data cache is automatically invalidated when the MPC603e is powered up and during a hard reset, assertion of the soft reset signal does not cause data cache invalidation. Software may invalidate the contents of the data cache using the data cache flash invalidate (DCFI) control bit in the HID0 register. Flash invalidation of the data cache is accomplished by setting and clearing the DCFI bit in two consecutive store operations.

3.2.3.2 Data Cache Disabling

The data cache may be disabled through the use of the data cache enable (DCE) control bit in the HID0 register. When the data cache is in the disabled state, the cache tag state bits are ignored, and all accesses are propagated to the bus as single-beat transactions. The DCE bit is cleared on power-up, causing the data cache to be disabled. To prevent the cache from being enabled or disabled in the middle of a data access, a **sync** instruction should be issued before changing the value of DCE.

Note that while snooping is not performed when the data cache is disabled, cache operations (caused by the **dcbz**, **dcbf**, **dcbst**, and **dcbi** instructions) are not affected by disabling the cache, causing potential coherency errors. An example of this would be a **dcbf** instruction that hits a modified cache block in the disabled cache, causing a copy back to memory of potentially stale data.

Regardless of the state of HID0[DCE], load and store operations are assumed to be weakly ordered. Thus, the LSU can perform load operations that occur later in the program ahead of store operations, even when the data cache is disabled. However, strongly ordered load and store operations can be enforced through the setting of the I bit (of the page WIMG bits) when address translation is enabled. Note that when address translation is disabled, the default WIMG bits cause the I bit to be cleared (accesses are assumed to be cacheable), and thus the accesses are weakly ordered. Refer to Section 3.5.2, “Caching-Inhibited Attribute (I),” for a description of the operation of the I bit and Section 5.2, “Real Addressing Mode,” for a description of the WIMG bits when address translation is disabled.

3.2.3.3 Data Cache Locking

The contents of the data cache may be locked through the HID0[DLOCK]. A locked data cache supplies data normally on a cache hit, but cache misses are treated as cache-inhibited accesses. The cache inhibited (\overline{CI}) signal is asserted if a cache access misses into a locked cache. The setting of DLOCK must be preceded by a **sync** instruction to prevent the cache from being locked during an access.

3.2.3.4 Data Cache Operations and Address Broadcasts

Executing a **dcbz** instruction generates an address-only broadcast on the bus. Additionally, if HID0[ABE] is set on a PID7t-603e processor, the execution of the **dcbf**, **dcbi**, and **dcbst** instructions also causes an address-only broadcast. The ability of the PID7t-603e to optionally perform address-only broadcasts when executing the **dcbi**, **dcbf**, and **dcbst** instructions allows the coherency management of an external copy-back L2 cache. Note that these cache control instruction broadcasts are not snooped by the PID7t-603e.

3.2.4 Data Cache Touch Load Support

Touch load operations allow an instruction stream to prefetch data from memory prior to a cache miss. The MPC603e supports touch load operations through a temporary cache block buffer located between the BIU and the data cache. The cache block buffer is essentially a floating cache block that is loaded by the BIU on a touch load operation, and is then read by a load instruction that requests that data. After a touch load completes on the bus, the BIU continues to compare the touch load address with subsequent load requests from the data cache. If the load address matches the touch load address in the BIU, the data is forwarded to the data cache from the touch load buffer, the read from memory is canceled, and the touch load address buffer is invalidated.

To avoid the storage of stale data in the touch load buffer, touch load requests that are mapped as write-through or caching-inhibited by the MMU are treated as no-ops by the BIU. Also, subsequent load instructions after a touch load that are mapped as write-through or caching-inhibited do not hit in the touch load buffer, and cause the touch load buffer to be invalidated on a matching address.

While the MPC603e provides only a single cache block buffer, other microprocessor implementations may provide buffering for more than one cache block. Programs written for other implementations may issue several **dcbt** or **dcbtst** instructions sequentially, reducing the performance if executed on the MPC603e. To improve performance in these situations, HID0[NOOPTI] (bit 31) can be set. This causes the **dcbt** and **dcbtst** instructions to be treated as no-ops, cause no bus activity, and incur only one processor clock cycle of execution latency. NOOPTI is cleared at a power-on reset, enabling the use of the **dcbt** and **dcbtst** instructions.

3.3 Basic Data Cache Operations

This section describes the three types of operations that can occur to the data cache, and how these operations are implemented in the MPC603e.

3.3.1 Data Cache Fill

A cache block is filled after a read miss or write miss (read-with-intent-to-modify) occurs in the cache. The cache block that corresponds to the missed address is updated by a burst transfer of the data from system memory. Note that if a read miss occurs in a system with multiple bus masters, and the data is modified in another cache, the modified data is first written to external memory before the cache fill occurs.

3.3.2 Data Cache Cast-Out Operation

The MPC603e uses an LRU replacement algorithm to determine which of the four possible cache locations should be used for a cache update on a cache miss. Adding a new block to

the cache causes any modified data associated with the least recently used element to be written back, or cast out, to system memory to maintain memory coherence.

3.3.3 Cache Block Push Operation

When a cache block in the MPC603e is snooped and hit by another bus master and the data is modified, the cache block must be written to memory and made available to the snooping device. The cache block that is hit, is pushed out onto the bus. The MPC603e supports two kinds of push operations—normal push operations and enveloped high-priority push operations, described in Section 3.6.9, “Enveloped High-Priority Cache Block Push Operation.”

3.4 Data Cache Transactions on Bus

The MPC603e transfers data to and from the data cache in single-beat transactions of two words, or in four-beat transactions of eight words which fill a cache block.

3.4.1 Single-Beat Transactions

Single-beat bus transactions can transfer from 1 to 8 bytes to or from the MPC603e. Single-beat transactions can be caused by cache write-through accesses, caching-inhibited accesses (I bit of the WIMG bits for the page is set), or accesses when the cache is disabled (HID0[DCE] bit is cleared), and can be misaligned.

3.4.2 Burst Transactions

Burst transactions on the MPC603e always transfer eight words of data at a time, and are aligned to a double-word boundary. The MPC603e transfer burst (TBST) output signal indicates to the system whether the current transaction is a single-beat transaction or four-beat burst transfer. Burst transactions have an assumed address order. For cacheable read operations or cacheable, non-write-through write operations that miss the cache, the MPC603e presents the double-word aligned address associated with the load or store instruction that initiated the transaction.

As shown in Figure 3-3, this quad word contains the address of the load or store that missed the cache. This minimizes latency by allowing the critical code or data to be forwarded to the processor before the rest of the block is filled. For all other burst operations, however, the entire block is transferred in order (oct-word aligned). Critical-double-word-first fetching on a cache miss applies to both the data and instruction cache.

MPC603e Cache Address

Bits 27:28

0 0	0 1	1 0	1 1
A	B	C	D

If the address requested is in double-word A, the address placed on the bus is that of double-word A, and the four data beats are ordered in the following manner:

Beat	0	1	2	3
	A	B	C	D

If the address requested is in double-word C, the address placed on the bus will be that of double-word C, and the four data beats are ordered in the following manner:

Beat	0	1	2	3
	C	D	A	B

Figure 3-3. Double-Word Address Ordering—Critical Double Word First

3.4.3 Access to Direct-Store Segments

The MPC603e does not provide support for access to direct-store segments. Operations attempting to access a direct-store segment will invoke a DSI exception. See Section 4.5.3, “DSI Exception (0x00300).”

3.5 Memory Management/Cache Access Mode Bits—W, I, M, and G

Some memory characteristics can be set on either a block or page basis by using the WIMG bits in the BAT registers or page table entry (PTE), respectively. The WIMG attributes control the following functionality:

- Write-through (W bit)
- Caching-inhibited (I bit)
- Memory coherency (M bit)
- Guarded memory (G bit)

These bits allow both uniprocessor and multiprocessor system designs to exploit numerous system-level performance optimizations.

Careless specification and use of these bits may create situations where coherency paradoxes are observed by the processor. In particular, this can happen when the state of these bits is changed without appropriate precautions being taken (for example, when

flushing the pages that correspond to the changed bits from the caches of all processors in the system is required, or when the address translations of aliased physical addresses (referred to as real addresses in the architecture specification) specify different values for any of the WIM bits). The MPC603e considers any of these cases to be a programming error that may compromise the coherency of memory. These paradoxes can occur within a single processor or across several devices, as described in Section 3.6.4.1, “Coherency in Single-Processor Systems.”

The WIMG attributes are programmed by the operating system for each page and block. The W and I attributes control how the processor performing an access uses its own cache. The M attribute ensures that coherency is maintained for all copies of the addressed memory location. The G attribute prevents out-of-order loading and prefetching from the addressed memory location.

When an access requires coherency, the processor performing the access must inform the coherency mechanisms throughout the system that the access requires memory coherency. The M attribute determines the kind of access performed on the bus (global or local).

The WIMG attributes occupy 4 bits in the BAT registers for block address translation and in the PTEs for page address translation. The WIMG bits are programmed as follows:

- The operating system uses the **mtspr** instruction to program the WIMG bits in the BAT registers for block address translation. The IBAT register pairs do not have a G bit and all accesses that use the IBAT register pairs are considered not guarded.
- The operating system writes the WIMG bits for each page into the PTEs in system memory as it sets up the page tables.

Note that for accesses performed with direct address translation (MSR[IR] = 0 or MSR[DR] = 0 for instruction or data access, respectively), the WIMG bits are automatically generated as 0b0011 (the data is write-back, caching is enabled, memory coherency is enforced, and memory is guarded).

3.5.1 Write-Through Attribute (W)

When an access is designated as write-through ($W = 1$), if the data is in the cache, a store operation updates the cached copy of the data. In addition, the update is written to the external memory location (as described below).

While the PowerPC architecture permits multiple store instructions to be combined for write-through accesses except when the store instructions are separated by a **sync** or **eieio** instruction, the MPC603e does not implement this ‘combined store’ capability. Note that a store operation that uses the write-through attribute may cause any part of valid data in the cache to be written back to main memory.

The definition of the external memory location to be written to, in addition to the on-chip cache, depends on the implementation of the memory system and can be illustrated by the following examples:

- RAM—The store is sent to the RAM controller to be written into the target RAM.
- I/O device—The store is sent to the memory-mapped I/O control hardware to be written to the target register or memory location.

In systems with multilevel caching, the store must be written to at least a depth in the memory hierarchy that is seen by all processors and devices.

Accesses that correspond to $W = 0$ are considered write-back. For this case, although the store operation is performed to the cache, it is only made to external memory when a copy-back operation is required. Use of the write-back mode ($W = 0$) can improve overall performance for areas of the memory space that are seldom referenced by other masters in the system.

3.5.2 Caching-Inhibited Attribute (I)

If $I = 1$, the memory access is completed by referencing the location in main memory, bypassing the on-chip cache. During the access, the addressed location is not loaded into the cache nor is the location allocated in the cache. It is considered a programming error if a copy of the target location of an access to caching-inhibited memory is resident in the cache. Software must ensure that the location has not been previously loaded into the cache, or, if it has, that it has been flushed from the cache.

The PowerPC architecture permits data accesses from more than one instruction to be combined for cache-inhibited operations, except when the accesses are separated by a **sync** instruction, or by an **eieio** instruction when the page or block is also designated as guarded. This ‘combined access’ capability is not implemented on the MPC603e. Note that the **eieio** is treated as a no-op by the MPC603e.

The caching-inhibited (I) bit in the MPC603e controls whether load and store operations are strongly or weakly ordered. If an I/O device requires load and store accesses to occur in program order, then the I bit for the page must be set.

3.5.3 Memory Coherency Attribute (M)

This attribute is provided to allow improved performance in systems where hardware-enforced coherency is relatively slow, and software is able to enforce the required coherency. When $M = 0$, the processor does not enforce data coherency. When $M = 1$, the processor enforces data coherency and the corresponding access is considered to be a global access.

When the M attribute is set, and the access is performed, the global signal is asserted to indicate that the access is global. Snooping devices affected by the access must then

respond to this global access if their data is modified by asserting \overline{ARTRY} , and updating the memory location.

Because instruction memory does not have to be consistent with data memory, the MPC603e ignores the M attribute for instruction accesses.

3.5.4 Guarded Attribute (G)

When the guarded bit is set, the memory area (block or page) is designated as guarded, meaning that the processor will perform out-of-order accesses to this area of memory, only as follows:

- Out-of-order load operations from guarded memory areas are performed only if the corresponding data is resident in the cache.
- The processor prefetches from guarded areas, but only when required, and only within the memory boundary dictated by the cache block. That is, if an instruction is certain to be required for execution by the program, it is fetched and the remaining instructions in the block may be prefetched, even if the area is guarded.

This setting can be used to protect certain memory areas from read accesses made by the processor that are not dictated directly by the program. If there are areas of memory that are not fully populated (in other words, there are holes in the memory map within this area), this setting can protect the system from undesired accesses caused by out-of-order load operations or instruction prefetches that could lead to the generation of the machine check exception. Also, the guarded bit can be used to prevent out-of-order load operations or prefetches from occurring to certain peripheral devices that produce undesired results when accessed in this way.

3.5.5 W, I, and M Bit Combinations

Table 3-1 summarizes the six combinations of the WIM bits. Note that either a zero or one setting for the G bit is allowed for each of these WIM bit combinations.

Table 3-1. Combinations of W, I, and M Bits

WIM Setting	Meaning
000	Data may be cached. Loads or stores whose target hits in the cache use that entry in the cache. Memory coherency is not enforced by hardware.
001	Data may be cached. Loads or stores whose target hits in the cache use that entry in the cache. Memory coherency is enforced by hardware.
010	Caching is inhibited. The access is performed to external memory, completely bypassing the cache. Memory coherency is not enforced by hardware.

Table 3-1. Combinations of W, I, and M Bits (continued)

WIM Setting	Meaning
011	Caching is inhibited. The access is performed to external memory, completely bypassing the cache. Memory coherency must be enforced by external hardware (processor provides hardware indication that access is global).
100	Data may be cached. Load operations whose target hits in the cache use that entry in the cache. Stores are written to external memory. The target location of the store may be cached and is updated on a hit. Memory coherency is not enforced by hardware.
101	Data may be cached. Load operations whose target hits in the cache use that entry in the cache. Stores are written to external memory. The target location of the store may be cached and is updated on a hit. Memory coherency is enforced by hardware.

3.5.5.1 Out-of-Order Execution and Guarded Memory

Out-of-order execution occurs when the MPC603e performs operations in advance in case the result is needed. Typically, these operations are performed by otherwise idle resources; thus if a result is not required, it is ignored and the out-of-order operation incurs no time penalty (typically).

Supervisor-level programs designate memory as guarded on a block or page level. Memory is designated as guarded if it may not be well-behaved with respect to out-of-order operations.

For example, the memory area that contains a memory-mapped I/O device may be designated as guarded if an out-of-order load or instruction fetch performed to such a device might cause the device to perform unexpected or incorrect operations. Another example of memory that should be designated as guarded is the area that corresponds to the device that resides at the highest implemented physical address (as it has no successor and out-of-order sequential operations such as instruction prefetching may result in a machine check exception). In addition, areas that contain holes in the physical memory space may be designated as guarded.

3.5.5.2 Effects of Out-of-Order Data Accesses

Most data operations may be performed out-of-order, as long as the machine appears to follow a simple sequential model. However, the following out-of-order operations do not occur:

- Out-of-order loading from guarded memory ($G = 1$) does not occur. However, when a load or store operation is required by the program, the entire cache block(s) containing the referenced data may be loaded into the cache.
- Out-of-order store operations that alter the state of the target location do not occur.

- No errors except machine check exceptions are reported due to the out-of-order execution of an instruction until it is known that execution of the instruction is required.

Machine check exceptions resulting solely from out-of-order execution (from nonguarded memory) may be reported. When an out-of-order instruction result is abandoned, only one side effect (other than a possible machine check) may occur—the referenced bit (R) in the corresponding page table entry (and TLB entry) can be set due to an out-of-order load operation. See Chapter 4, “Exceptions,” for more information on the machine check exception.

Thus, an out-of-order load or store instruction will not access guarded memory unless one of the following conditions exist:

- The target memory item is resident in an on-chip cache. In this case, the location may be accessed from the cache or main memory.
- The target memory item is cacheable ($I = 0$) and it is guaranteed that the load or store is in the execution path (assuming there are no intervening exceptions). In this case, the entire cache block containing the target may be loaded into the cache.

3.5.5.3 Effects of Out-of-Order Instruction Fetches

To avoid instruction fetch delay, the processor typically fetches instructions ahead of those currently being executed. Such instruction prefetching is said to be out-of-order in that prefetched instructions may not be executed due to intervening branches or exceptions.

During instruction prefetching, no errors except machine check exceptions are reported due to the out-of-order fetching of an instruction until it is known that execution of the instruction is required.

Machine check exceptions resulting solely from out-of-order execution (from nonguarded memory) may be reported. When an out-of-order instruction result is abandoned, only one side effect (other than a possible machine check) may occur—the referenced bit (R) in the corresponding page table entry (and TLB entry) can be set due to an out-of-order load operation. See Chapter 4, “Exceptions,” for more information on the machine check exception.

Instruction fetching from guarded memory is not permitted.

3.6 Cache Coherency—MEI Protocol

The primary objective of a coherent memory system is to provide the same image of memory to all devices using the system. Coherency allows synchronization and cooperative use of shared resources. Otherwise, multiple copies of a memory location, some containing stale values, could exist in a system resulting in errors when the stale values are used. Each potential bus master must follow rules for managing the state of its cache.

The MPC603e cache coherency protocol is a coherent subset of the standard MESI four-state cache protocol that omits the shared state. Since data cannot be shared, the MPC603e signals all cache block fills as if they were write misses (read-with-intent-to-modify), flushing the corresponding copies of the data in all caches external to the MPC603e prior to the MPC603e cache block fill operation. Following the cache block load, the MPC603e is the exclusive owner of the data and may write to it without a bus broadcast transaction.

To maintain this coherency, all global reads observed on the bus by the MPC603e are snooped as if they were writes, causing the MPC603e to write a modified cache block back to memory and invalidate the cache block, or simply invalidate the cache block if it is unmodified. The exception to this rule occurs when a snooped transaction is a caching-inhibited read (either burst or single-beat, where TT[0:4] = 0x1010; see Table 7-1 for clarification), in which case the MPC603e does not invalidate the snooped cache block. If the cache block is modified, the block is written back to memory, and the cache block is marked exclusive unmodified. If the cache block is marked exclusive unmodified when snooped, no bus action is taken, and the cache block remains in the exclusive unmodified state. This treatment of caching-inhibited reads decreases the possibility of data thrashing by allowing noncaching devices to read data without invalidating the entry from the MPC603e data cache.

3.6.1 MEI State Definitions

The MPC603e data cache characterizes each 32-byte block it contains as being in one of three MEI states. Addresses presented to the cache are indexed into the cache directory with bits A20–A26, and the upper-order 20 bits from the physical address translation (PA0–PA19) are compared against the indexed cache directory tags. If neither of the indexed tags matches, the result is a cache miss. If a tag matches, a cache hit occurred and the directory indicates the state of the cache block through two state bits kept with the tag. The three possible states for a cache block in the cache are the modified state (M), the exclusive state (E), and the invalid state (I). The three MEI states are defined in Table 3-2.

Table 3-2. MEI State Definitions

MEI State	Definition
Modified (M)	The addressed cache block is valid only in the cache. The cache block is modified with respect to system memory—that is, the modified data in the cache block has not been written back to memory.
Exclusive (E)	The addressed block is in this cache only. The data in this cache block is consistent with system memory.
Invalid (I)	This state indicates that the addressed cache block is not resident in the cache.

3.6.2 MEI State Diagram

The MPC603e provides dedicated hardware to provide memory coherency by snooping bus transactions. The address retry capability of the MPC603e enforces the MEI protocol, as

shown in Figure 3-4. Figure 3-4 assumes that the WIM bits for the page or block are set to 001; that is, write-back, caching-not-inhibited, and memory coherency enforced.

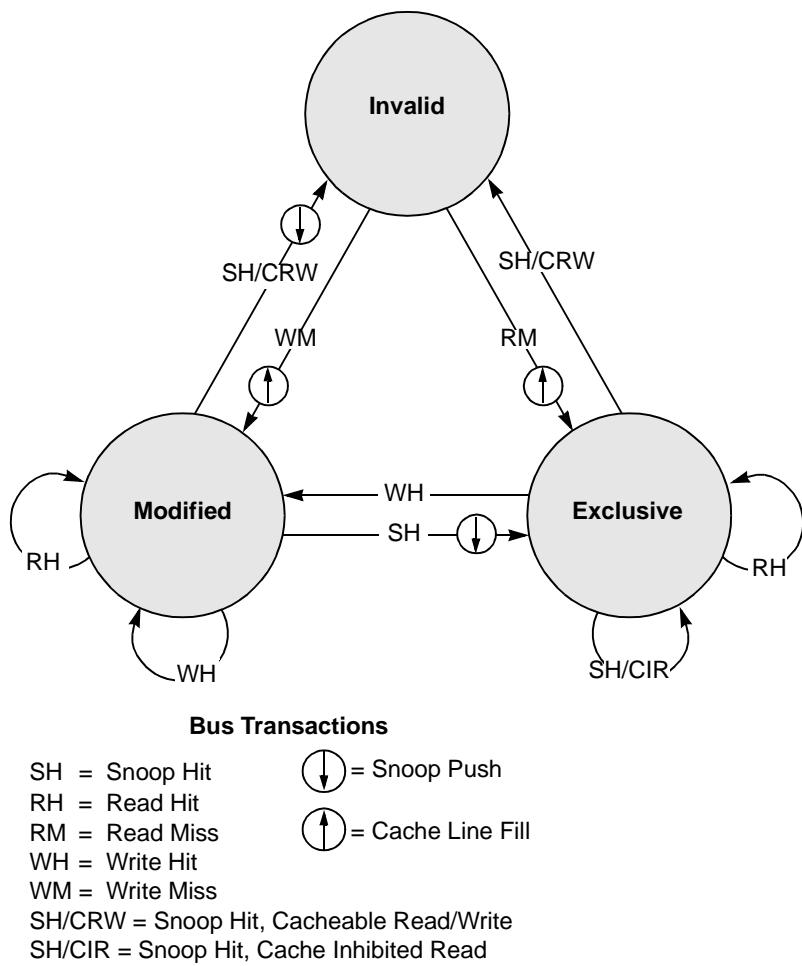


Figure 3-4. MEI Cache Coherency Protocol—State Diagram (WIM = 001)

Section 3.10, “MEI State Transactions,” provides a detailed list of MEI transitions for various operations and WIM bit settings.

3.6.3 MEI Hardware Considerations

While the MPC603e provides the hardware required to monitor bus traffic for coherency, the MPC603e data cache tags are single-ported, and a simultaneous load or store and snoop access represent a resource conflict. In general, the snoop access has highest priority and is given first access to the tags. The load or store access will then occur on the clock following the snoop. The snoop is not given priority into the tags when the snoop coincides with a tag write (for example, validation after a cache block load). In these situations, the snoop is retried and must re-arbitrate before the lookup is possible.

Occasionally, cache snoops cannot be serviced and must be retried. These retries occur if the cache is busy with a burst read or write when the snoop operation takes place.

Note that it is possible for a snoop to hit a modified cache block that is already in the process of being written to the copy-back buffer for replacement purposes. If this happens, the MPC603e retries the snoop, and raises the priority of the cast-out operation to allow it to go to the bus before the cache block fill.

The global ($\overline{\text{GBL}}$) signal, asserted as part of the address attribute field during a bus transaction, enables the snooping hardware of the MPC603e. Address bus masters assert $\overline{\text{GBL}}$ to indicate that the current transaction is a global access (that is, an access to memory shared by more than one device). If $\overline{\text{GBL}}$ is not asserted for the transaction, that transaction is not snooped by the MPC603e. Note that the $\overline{\text{GBL}}$ signal is not asserted for instruction fetches, and that $\overline{\text{GBL}}$ is asserted for all data read or write operations when using direct address translation. (Note that direct address translation is referred to as the real addressing mode, not the direct-store segment, in the architecture specification.)

Normally, $\overline{\text{GBL}}$ reflects the M-bit value specified for the memory reference in the corresponding translation descriptor(s). Care must be taken to minimize the number of pages marked as global, because the retry protocol enforces coherency and can use considerable bus bandwidth if a lot of data is shared. Therefore, available bus bandwidth can decrease as more traffic is marked global.

The MPC603e snoops a transaction if the transfer start ($\overline{\text{TS}}$) and $\overline{\text{GBL}}$ signals are asserted together in the same bus clock (this is a qualified snooping condition). No snoop update to the MPC603e cache occurs if the snooped transaction is not marked global. Also, because cache block cast-outs and snoop pushes do not require snooping, the $\overline{\text{GBL}}$ signal is not asserted for these operations.

When the MPC603e detects a qualified snoop condition, the address associated with the $\overline{\text{TS}}$ signal is compared with the cache tags. Snooping finishes if no hit is detected. If, however, the address hits in the cache, the MPC603e reacts according to the MEI protocol shown in Figure 3-4.

To facilitate external monitoring of the internal cache tags, the cache set entry signals (CSE[0:1]) represent in binary the cache set being replaced on read operations (including read-with-intent-to-modify operations). The CSE[0:1] signals do not apply for write operations to memory, or during noncacheable or touch load operations. Note that these signals are valid only for MPC603e burst operations. Table 3-3 shows the CSE[0:1] (cache set entry) encodings.

Table 3-3. CSE[0:1] Signal Encoding

CSE[0:1]	Cache Set Element
00	Set 0
01	Set 1
10	Set 2
11	Set 3

3.6.4 Coherency Precautions

The MPC603e supports a three-state coherency protocol that supports the modified, exclusive, and invalid (MEI) cache states. This protocol is a compatible subset of the MESI four-state protocol and operates coherently in systems that contain four-state caches. In addition, the MPC603e does not broadcast cache operations caused by cache instructions. They are intended for the management of the local cache but not for other caches in the system.

3.6.4.1 Coherency in Single-Processor Systems

The following situations concerning coherency can be encountered within a single-processor system:

- Load or store to a caching-inhibited page (WIM = 0bx1x) and a cache hit occurs.
Caching is inhibited for this page ($I = 1$)—Load or store operations to a caching-inhibited page that hit in the cache cause boundedly undefined results.
- Store to a page marked write-through (WIM = 0b10x) and a cache read hit to a modified cache block.
This page is marked as write-through ($W = 1$)—The MPC603e pushes the modified cache block to memory and the block remains marked modified (M).

Note that when WIM bits are changed, it is critical that the cache contents reflect the new WIM bit settings. For example, if a block or page that had allowed caching becomes caching-inhibited, software should ensure that the appropriate cache blocks are flushed to memory and invalidated.

3.6.5 Load and Store Coherency Summary

Table 3-4 provides a summary of memory coherency actions performed by the MPC603e on load operations. Noncacheable cases are not part of this table.

Table 3-4. Memory Coherency Actions on Load Operations

Cache State	Bus Operation	ARTRY	Action
M	None	Don't care	Read from cache
E	None	Don't care	Read from cache
I	Read	Negated	Load data and mark E
I	Read	Asserted	Retry read operation

Table 3-5 provides an overview of memory coherency actions on store operations. This table does not include noncacheable or write-through cases. The read-with-intent-to-modify (RWITM) examples involve selecting a replacement class and casting-out modified data that may have resided in that replacement class.

Table 3-5. Memory Coherency Actions on Store Operations

Cache State	Bus Operation	ARTRY	Action
M	None	Don't care	Modify cache
E	None	Don't care	Modify cache, mark M
I	RWITM	Negated	Load data, modify it, mark M
I	RWITM	Asserted	Retry the RWITM

3.6.6 Atomic Memory References

The Load Word and Reserve Indexed (**lwarx**) and Store Word Conditional Indexed (**stwcx.**) instructions provide an atomic update function for a single, aligned word of memory. While an **lwarx** instruction will normally be paired with an **stwcx.** instruction with the same effective address, an **stwcx.** instruction to any address will cancel the reservation. For detailed information on these instructions, refer to Chapter 2, “Programming Model,” in this book and Chapter 8, “Instruction Set,” in the *Programming Environments Manual*.

3.6.7 Cache Reaction to Specific Bus Operations

There are several bus transaction types defined for the MPC603e bus. The MPC603e must snoop these transactions and perform the appropriate action to maintain memory coherency as shown in Table 3-6. A processor may assert **ARTRY** for any bus transaction due to internal conflicts that prevent the appropriate snooping. The transactions in Table 3-6 correspond to the transfer type signals TT[0:4], described in Section 7.2.4.1, “Transfer Type (TT[0:4]).”

Table 3-6. Response to Bus Transactions

Snooped Transaction	MPC603e Response
Clean block	No action is taken
Flush block	No action is taken
Write-with-flush Write-with-flush-atomic	<p>Write-with-flush and write-with-flush-atomic operations occur after the processor issues a store or stwcx. instruction, respectively.</p> <ul style="list-style-type: none"> • If the addressed block is in the exclusive state, the address snoop forces the state of the addressed block to invalid. • If the addressed block is in the modified state, the address snoop causes ARTRY to be asserted and initiates a push of the modified block out of the cache and changes the state of the block to invalid. • The execution of an stwcx. instruction cancels the reservation associated with any address.
Kill block	The kill block operation is an address-only bus transaction initiated when a dcbz instruction is executed; when snooped by the MPC603e, the addressed cache block is invalidated if in the E state, or flushed to memory and invalidated if in the M state, and any associated reservation is canceled.

Table 3-6. Response to Bus Transactions (continued)

Snooped Transaction	MPC603e Response
Write-with-kill	In a write-with-kill operation, the processor snoops the cache for a copy of the addressed block. If one is found, an additional snoop action is initiated internally and the cache block is forced to the I state, killing modified data that may have been in the block. Any reservation associated with the block is also canceled.
Read Read-atomic	The read operation is used by most single-beat and burst read operations on the bus. All burst reads observed on the bus are snooped as if they were writes, causing the addressed cache block to be flushed. A read on the bus with the GBL signal asserted causes the following responses: <ul style="list-style-type: none"> If the addressed block in the cache is invalid, the MPC603e takes no action. If the addressed block in the cache is in the exclusive state, the block is invalidated. If the addressed block in the cache is in the modified state, the block is flushed to memory and the block is invalidated. If the snooped transaction is a caching-inhibited read, and the block in the cache is in the exclusive state, the snoop causes no bus activity and the block remains in the exclusive state. If the block is in the cache in the modified state, the MPC603e initiates a push of the modified block out to memory and marks the cache block as exclusive. Read-atomic operations appear on the bus in response to Iwarx instructions and generate the same snooping responses as read operations.
Read-with-intent-to-modify (RWITM) RWITM-atomic	A RWITM operation is issued to acquire exclusive use of a memory location for the purpose of modifying it. <ul style="list-style-type: none"> If the addressed block is invalid, the MPC603e takes no action. If the addressed block in the cache is in the exclusive state, the MPC603e initiates an additional snoop action to change the state of the cache block to invalid. If the addressed block in the cache is in the modified state, the block is flushed to memory and the block is invalidated. The RWITM-atomic operations appear on the bus in response to stwcx. instructions and are snooped like RWITM instructions.
sync	No action is taken
TLB invalidate	No action is taken

3.6.8 Operations Causing ARTRY Assertion

The following scenarios cause the MPC603e to assert the **ARTRY** signal:

- Snoop hits to a block in the M state (flush or clean)

This case is a normal snoop hit and will result in **ARTRY** being asserted if the snooped transaction was a ‘flush’ or ‘clean’ request. If the snooped transaction was a ‘kill’ request, **ARTRY** will not be asserted.

- Snoop attempt during the last **TA** of a cache line fill

In No-**DRTRY** mode, during the cycle that the last **TA** is asserted to the MPC603e on a cache line fill, the tag is being written to its new state by the MPC603e and is not accessible. This will result in a collision being signaled by asserting **ARTRY**. With **DRTRY** enabled, the cache tags are inaccessible to a snoop operation one cycle after the last **TA**.

- Snoop hit after the first \overline{TA} of a burst load operation
After the first \overline{TA} of a burst load operation, the data tags are committed to being written; snoop operations cannot be serviced until the load completes, thereby causing the assertion of \overline{ARTRY} .
- Snoop hits to line in the cast-out buffer
The MPC603e cast-out buffer is kept coherent with main memory, and snoop operations that hit in the cast-out buffer will cause the assertion of \overline{ARTRY} .
- Snoop attempt during cycles that **dcbz** instruction or load or store operation is updating the tag
During the execution of a **dcbz** instruction or during a load or store operation that requires a cache line cast out, the cache tags will be inaccessible during the first and last cycle of the operation.
- Snoop attempt during the cycle that a **dcbf** or **dcbst** instruction is updating the tag
If the EA of a **dcbf** or **dcbst** instruction hits in the cache, the tag will be changed to its new state. During that clock, the tag is not accessible and snoop transactions during that cycle will cause the assertion of \overline{ARTRY} .

3.6.9 Enveloped High-Priority Cache Block Push Operation

In cases where the MPC603e has completed the address tenure of a read operation, and then detects a snoop hit to a modified cache block by another bus master, the MPC603e provides a high-priority push operation. If the address snooped is the same as the address of the data to be returned by the read operation, \overline{ARTRY} is asserted one or more times until the data tenure of the read operation is completed. The cache block push transaction can be enveloped within the address and data tenures of a read operation. This feature prevents deadlocks in system organizations that support multiple memory-mapped buses.

More specifically, the MPC603e internally detects the scenario where a load request is outstanding and the processor has pipelined a write operation on top of the load. Normally, when the data bus is granted to the MPC603e, the resulting data bus tenure is used for the load operation. The enveloped high-priority cache block push feature defines a bus signal, the data bus write only qualifier (\overline{DBWO}), which, when asserted with a qualified data bus grant, indicates that the resulting data tenure should be used for the store operation instead. This signal is described in Section 8.10, “Using \overline{DBWO} .” Note that the enveloped copy-back operation is an internally pipelined bus operation.

3.7 Cache Control Instructions

Software must use the appropriate cache management instructions to ensure that caches are kept consistent when data is modified by the processor. When a processor alters a memory location that may be contained in an instruction cache, software must ensure that updates

to memory are visible to the instruction fetching mechanism. Although the instructions to enforce coherency vary among implementations and, hence, operating systems should provide a system service for this function, the following sequence is typical:

1. **dcbst** (update memory)
2. **sync** (wait for update)
3. **icbi** (invalidate copy in cache)
4. **isync** (invalidate copy in own instruction buffer)

These operations are necessary because the processor does not maintain instruction memory coherent with data memory. Software is responsible for enforcing coherency of instruction caches and data memory. Since instruction fetching may bypass the data cache, changes made to items in the data cache may not be reflected in memory until after the instruction fetch completes.

The PowerPC architecture defines instructions for controlling both the instruction and data caches when they exist. The MPC603e interprets the cache control instructions (**icbi**, **dcbi**, **dcbt**, **dcbz**, **dcbst**) as if they pertain only to the MPC603e caches. They are not intended for use in managing other caches in the system.

The **dcbz** instruction causes an address-only broadcast on the bus if the contents of the block are from a page marked global through the WIMG bits. This broadcast is performed for coherency reasons; the **dcbz** instruction is the only cache control instruction that can allocate and take new ownership of a line. Note that if the HID0[ABE] bit is set on a PID7t-603e processor, the execution of the **dcbf**, **dcbi**, and **dcbst** instructions will also cause an address-only broadcast. The **dcbz** instruction is also the only cache operation that is snooped by the MPC603e. The cache instructions are intended primarily for the management of the on-chip cache, and do not perform address-only broadcasts for the maintenance of other caches in the system. The ability of the PID7t-603e to optionally perform address-only broadcasts when executing the **dcbi**, **dcbf**, and the **dcbst** instructions allows the coherency management of an external copy-back L2 cache.

The other instructions do not broadcast either for the purpose of invalidating or flushing other caches in the system or for managing system resources. Any bus activity caused by these instructions is the direct result of performing the operation in the MPC603e cache. Note that a data access exception is generated if the effective address of a **dcbi**, **dcbst**, **dcbf**, or **dcbz** instruction cannot be translated due to the lack of a TLB entry. (Note that exceptions are referred to as interrupts in the architecture specification.)

Note that in the PowerPC architecture, the term ‘cache block’ or ‘block,’ when used in the context of cache implementations, refers to the unit of memory at which coherency is maintained. For the MPC603e, this is the eight-word cache line. This value may be different for other implementations that support the PowerPC architecture. In-depth descriptions of coding these instructions is provided in Chapter 3, “Addressing Modes and Instruction Set Summary,” and Chapter 10, “Instruction Set,” in the *Programming Environments Manual*.

3.7.1 Data Cache Block Invalidate (dcbi) Instruction

If the block containing the byte addressed by the EA is in the data cache, the cache block is invalidated regardless of whether the block is in the exclusive or modified state. If HID0[ABE] is set on a PID7t-603e when a **dcbi** instruction is executed, the PID7t-603e will perform an address-only bus transaction. The **dcbi** instruction can only be executed when the MPC603e is in the supervisor state.

3.7.2 Data Cache Block Touch (dcbt) Instruction

This instruction provides a method for improving performance through the use of software-initiated prefetch hints. The MPC603e performs the fetch when the address hits in the TLB or BAT registers, and when it is a permitted load access from the addressed page. The operation is treated similarly to a byte load operation with respect to coherency.

If the address translation does not hit in the TLB or BAT mechanism, or if it does not have load access permission, the instruction is treated as a no-op.

If the cache is locked or disabled, or if the access is to a page that is marked as guarded, the **dcbt** instruction is treated as a no-op.

If the access is directed to a write-through or caching-inhibited page, the instruction is treated as a no-op.

The **dcbt** instruction never affects the referenced or changed bits in the hashed page table.

A successful **dcbt** instruction affects the state of the TLB and cache LRU bits as defined by the LRU algorithm.

The touch load buffer will be marked invalid if the contents of the touch buffer have been moved to the cache, if any data cache management instruction has been executed, if a **dcbz** instruction is executed that matches the address of the cache block in the touch buffer, or if another **dcbt** instruction is executed.

3.7.3 Data Cache Block Touch for Store (dcbtst) Instruction

The **dcbtst** instruction, like the data cache block touch instruction (**dcbt**), allows software to prefetch a cache block in anticipation of a store operation (read-with-intent-to-modify).

3.7.4 Data Cache Block Clear to Zero (dcbz) Instruction

If the block containing the byte addressed by the EA is in the data cache, all bytes are cleared.

If the block containing the byte addressed by the EA is not in the data cache and the corresponding page is caching-allowed, the block is established in the data cache without fetching the block from main memory, and all bytes of the block are cleared. If the contents

of the cache block are from a page marked global through the WIM bits, an address-only bus transaction is run.

If the page containing the byte addressed by the EA is caching-inhibited or write-through, then the system alignment exception handler is invoked.

The **dcbz** instruction is treated as a store to the addressed byte with respect to address translation and protection.

3.7.5 Data Cache Block Store (dcbst) Instruction

If the block containing the byte addressed by the EA is in coherency-required mode, and a block containing the byte addressed by the EA is in the data cache of any processor and has been modified, the writing of it to main memory is initiated. On a PID7t-603e, if the cache block is unmodified, HID0[ABE] is set, and if the contents of the cache block are from a page marked global through the WIM bits, an address-only bus transaction is run.

The function of this instruction is independent of the write-through and caching-inhibited/caching-allowed modes of the block containing the byte addressed by the EA.

This instruction is treated as a load to the addressed byte with respect to address translation and protection.

3.7.6 Data Cache Block Flush (dcbf) Instruction

The action taken depends on the memory mode associated with the target, and on the state of the cache block. The following list describes the action taken for the various cases. These actions are executed regardless of whether the page containing the addressed byte is in caching-inhibited or caching-allowed mode. The following actions occur in both coherency-required (WIM = 0bxx1) and coherency-not-required mode (WIM = 0bxx0).

The **dcbf** instruction causes the following cache activity:

- Unmodified block—invalidates the block in the processor’s cache
- Modified block—copies the block to memory and invalidates data cache block
- Absent block—does nothing

The MPC603e treats this instruction as a load to the addressed byte with respect to address translation and protection.

3.7.7 Enforce In-Order Execution of I/O Instruction (eieio)

As defined by the PowerPC architecture, the **eieio** instruction provides an ordering function for the effects of load and store instructions executed by a given processor. Executing **eieio** ensures that all memory accesses previously initiated by the given processor are completed with respect to main memory before any memory accesses subsequently initiated by the

processor access main memory. The **eieio** instruction orders loads and stores to caching-inhibited memory only.

The **eieio** instruction is intended for use only in performing memory-mapped I/O operations. It enforces ‘strong’ ordering of cache-inhibited memory accesses during I/O operations between the processor and I/O devices.

When executed by the MPC603e, the **eieio** instruction is treated as a no-op; caching-inhibited load and store operations (inhibited by the WIMG bits for the page) are performed in strict program order.

3.7.8 Instruction Cache Block Invalidate (**icbi**) Instruction

The execution of an **icbi** instruction causes all four cache sets indexed by the EA to be marked invalid. No cache hit is required, and no MMU translation is performed.

3.7.9 Instruction Synchronize (**isync**) Instruction

The **isync** instruction waits for all previous instructions to complete and then discards any previously fetched instructions, causing subsequent instructions to be fetched (or refetched) from memory and to execute in the context established by the previous instructions. This instruction has no effect on other processors or on their caches.

3.8 Bus Operations Caused by Cache Control Instructions

Table 3-7 provides an overview of the bus operations initiated by cache control instructions. The cache control, TLB management, and synchronization instructions supported by the MPC603e may affect or be affected by the operation of the bus. None of the instructions will actively broadcast through address-only transactions on the bus (except for **dcbz**), and no broadcasts by other masters are snooped by the MPC603e (except for kills). The operation of the instructions, however, may indirectly cause bus transactions to be performed, or their completion may be linked to the bus. Table 3-7 summarizes how these instructions may operate with respect to the bus.

**Table 3-7. Bus Operations Caused by Cache Control Instructions
(WIM = 001)**

Operation	Cache State	Next Cache State	Bus Operations	Comment
sync	Don't care	No change	None	Waits for memory queues to complete bus activity
icbi	Don't care	I	None	—
dcbi	Don't care	I	None	—
dcbf	I, E	I	None	—

**Table 3-7. Bus Operations Caused by Cache Control Instructions
(WIM = 001) (continued)**

Operation	Cache State	Next Cache State	Bus Operations	Comment
dcbf	M	I	Write-with-kill	Block is pushed
dcbst	I, E	No change	None	—
dcbst	M	E	Write	Block is pushed
dcbz	I	M	Write-with-kill	—
dcbz	E, M	M	Kill block	Writes over modified data
dcbt	I	No change	Read	Fetched cache block is stored in touch load queue
dcbt	E, M	No change	None	—
dcbtst	I	No change	Read-with-intent-to-modify	Fetched cache block is stored in touch load queue
dcbtst	E, M	No change	None	—

Note that Table 3-7 assumes that the WIM bits are set to 001; that is, since the cache is operating in write-back mode, caching is permitted and coherency is enforced.

Table 3-7 does not include noncacheable or write-through cases, nor does it completely describe the mechanisms for the operations described. For more information, see Section 3.10, “MEI State Transactions.”

For detailed information on the cache control instructions, refer to Chapter 2, “Programming Model,” in this book and Chapter 8, “Instruction Set,” in the *Programming Environments Manual*. The MPC603e contains snooping logic to monitor the bus for these commands and the control logic required to keep the cache and the memory queues coherent. For additional details about the specific bus operations performed by the MPC603e, see Chapter 8, “System Interface Operation.”

3.9 Bus Interface

The bus interface buffers bus requests from the instruction and data caches, and executes the requests per the MPC603e bus protocol. It includes address register queues, prioritization logic, and bus control logic. The bus interface also captures snoop addresses for snooping in the cache and in the address register queues, snoops for reservations, and holds the touch load address for the cache. All data storage for the address register buffers (load and store data buffers) are located in the cache section. The data buffers are considered temporary storage for the cache and not part of the bus interface.

The general functions and features of the bus interface are as follows:

- Address register buffers that include:
 - Instruction cache load address buffer

- Data cache load address buffer
- Data cache touch load address buffer (associated data block buffer located in cache)
- Data cache castout/store address buffer (associated data line buffer located in cache)
- Data cache snoop copy-back address buffer (associated data line buffer located in cache)
- Reservation address buffer for snoop monitoring
- Pipeline collision detection for data cache buffers
- Reservation address snooping for **lwarx/stwex**. instructions
- One-level address pipelining
- Load ahead of store capability

Figure 3-5 is a conceptual block diagram of the bus interface. The address register queues hold transaction requests that the bus interface may issue on the bus independently of the other requests. The bus interface may have up to two transactions operating on the bus at any given time through the use of address pipelining.

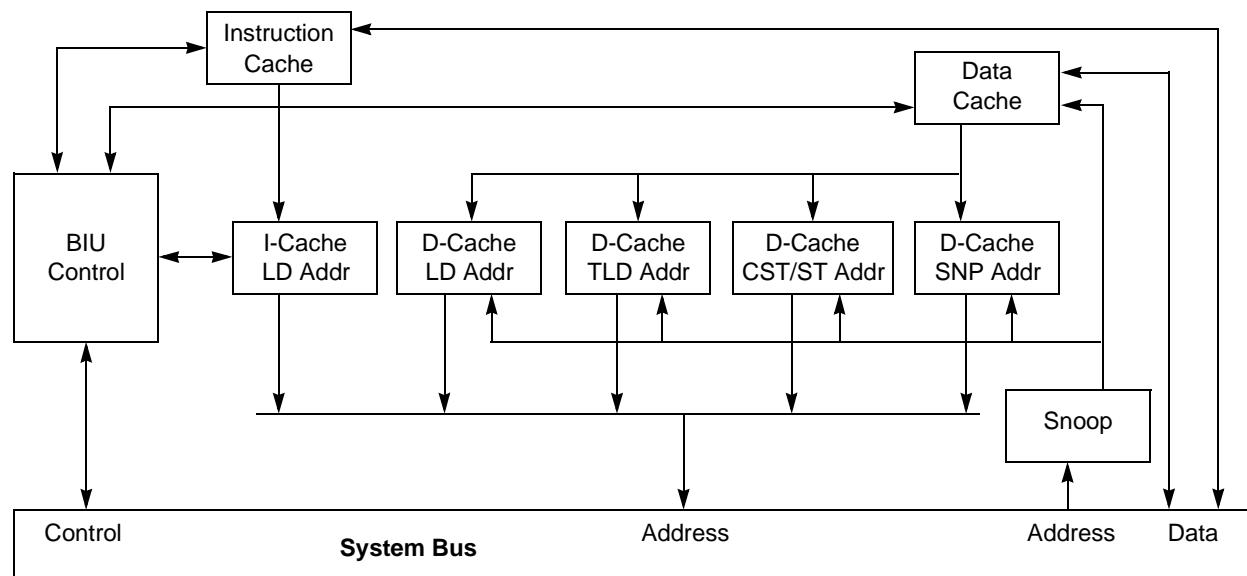


Figure 3-5. Bus Interface Address Buffers

For additional information about the MPC603e bus interface and the bus protocols, refer to Chapter 8, “System Interface Operation.”

3.10 MEI State Transactions

Table 3-8 shows MEI state transitions for various operations. Bus operations are described in Table 3-6.

Table 3-8. MEI State Transitions

Operation	Cache Operation	Bus Sync	WIM	Current State	Next State	Cache Actions	Bus Operation
Load (T = 0)	Read	No	x0x	I	Same	1 Cast out of modified block (as required)	Write-with-kill
						2 Pass four-beat read to memory queue	Read
Load (T = 0)	Read	No	x0x	E,M	Same	Read data from cache	—
Load (T = 0)	Read	No	x1x	I	Same	Pass single-beat read to memory queue	Read
Load (T = 0)	Read	No	x1x	E	I	CRTRY read	—
Load (T = 0)	Read	No	x1x	M	I	CRTRY read (push sector to write queue)	Write-with-kill
lwarx	Read	Acts like other reads but bus operation uses special encoding					
Store (T = 0)	Write	No	00x	I	Same	1 Cast out of modified block (if necessary)	Write-with-kill
						2 Pass RWITM to memory queue	RWITM
Store (T = 0)	Write	No	00x	E,M	M	Write data to cache	—
Store ≠ stwcx. (T = 0)	Write	No	10x	I	Same	Pass single-beat write to memory queue	Write-with-flush
Store ≠ stwcx. (T = 0)	Write	No	10x	E	Same	1 Write data to cache	—
						2 Pass single-beat write to memory queue	Write-with-flush
Store ≠ stwcx. (T = 0)	Write	No	10x	M	Same	1 CRTRY write	—
						2 Push block to write queue	Write-with-kill
Store (T = 0) or stwcx. (WIM = 10x)	Write	No	x1x	I	Same	Pass single-beat write to memory queue	Write-with-flush
Store (T = 0) or stwcx. (WIM = 10x)	Write	No	x1x	E	I	CRTRY write	—
Store (T = 0) or stwcx. (WIM = 10x)	Write	No	x1x	M	I	1 CRTRY write	—
						2 Push block to write queue	Write-with-kill
stwcx.	Conditional write	If the reserved bit is set, this operation is like other writes except the bus operation uses a special encoding.					

Table 3-8. MEI State Transitions (continued)

Operation	Cache Operation	Bus Sync	WIM	Current State	Next State	Cache Actions	Bus Operation
dcbf	Data cache block flush	No	xxx	I,E	Same	1 CRTRY dcbf	—
						2 Pass flush	Flush
				Same	I	3 State change only	—
dcbf	Data cache block flush	No	xxx	M	I	Push block to write queue	Write-with-kill
dcbst	Data cache block store	No	xxx	I,E	Same	1 CRTRY dcbst	—
						2 Pass clean	Clean
				Same	Same	3 No action	—
dcbst	Data cache block store	No	xxx	M	E	Push block to write queue	Write-with-kill
dcbz	Data cache block set to zero	No	x1x	x	x	Alignment trap	—
dcbz	Data cache block set to zero	No	10x	x	x	Alignment trap	—
dcbz	Data cache block set to zero	Yes	00x	I	Same	1 CRTRY dcbz	—
						2 Cast out of modified block	Write-with-kill
						3 Pass kill	Kill
				Same	M	4 Clear block	—
dcbz	Data cache block set to zero	No	00x	E,M	M	Clear block	—
dcbt	Data cache block touch	No	x1x	I	Same	Pass single-beat read to memory queue	Read
dcbt	Data cache block touch	No	x1x	E	I	CRTRY read	—
dcbt	Data cache block touch	No	x1x	M	I	1 CRTRY read	—
						2 Push block to write queue	Write-with-kill
dcbt	Data cache block touch	No	x0x	I	Same	1 Cast out of modified block (as required)	Write-with-kill
						2 Pass four-beat read to memory queue	Read
dcbt	Data cache block touch	No	x0x	E,M	Same	No action	—
Single-beat read	Reload dump 1	No	xxx	I	Same	Forward data_in	—

Table 3-8. MEI State Transitions (continued)

Operation	Cache Operation	Bus Sync	WIM	Current State	Next State	Cache Actions	Bus Operation
Four-beat read (double-word-aligned)	Reload dump	No	xxx	I	E	Write data_in to cache	—
Four-beat write (double-word-aligned)	Reload dump	No	xxx	I	M	Write data_in to cache	—
E→I	Snoop write or kill	No	xxx	E	I	State change only (committed)	—
M→I	Snoop kill	No	xxx	M	I	State change only (committed)	—
Push M→I	Snoop flush	No	xxx	M	I	Conditionally push	Write-with-kill
Push M→E	Snoop clean	No	xxx	M	E	Conditionally push	Write-with-kill
tlbie	TLB invalidate	No	xxx	x	x	1 CRTRY TLBI	—
						2 Pass TLBI	—
						3 No action	—
sync	Synchronization	No	xxx	x	x	1 CRTRY sync	—
						2 Pass sync	—
						3 No action	—

Note: Single-beat writes are not snooped in the write queue.



MEI State Transactions

Freescale Semiconductor, Inc.

Freescale Semiconductor, Inc.

MPC603e RISC Microprocessor User's Manual

**For More Information On This Product,
Go to: www.freescale.com**

Chapter 4

Exceptions

The PowerPC exception mechanism allows the processor to change to supervisor state as a result of external signals, errors, or unusual conditions arising in the execution of instructions, and differ from the arithmetic exceptions defined by the IEEE for floating-point operations. When exceptions (referred to as interrupts in the architecture specification) occur, information about the state of the processor is saved to certain registers and the processor begins execution at an address (exception vector) predetermined for each exception. Processing of exceptions occurs in supervisor mode.

Although multiple exception conditions can map to a single exception vector, a more specific condition may be determined by examining a register associated with the exception—for example, the DSISR or FPSCR. Additionally, certain exception conditions can be explicitly enabled or disabled by software.

The PowerPC architecture requires that exceptions be handled in program order; therefore, although a particular implementation may recognize exception conditions out of order, they are handled strictly in order with respect to the instruction stream. When an instruction-caused exception is recognized, any unexecuted instructions that appear earlier in the instruction stream, including any that have not yet entered the execute state, are required to complete before the exception is taken. Any exceptions caused by those instructions are handled first. Likewise, exceptions that are asynchronous and precise are recognized when they occur, but are not handled until the instruction currently in the completion stage successfully completes execution or generates an exception, and the completed store queue is emptied. An instruction is said to have completed when the results of that instruction's execution have been committed to the registers defined by the architecture (for example, the GPRs or FPRs, rather than rename buffers). If a single instruction encounters multiple exception conditions, those exceptions are taken and handled sequentially. Likewise, exceptions that are asynchronous are recognized when they occur, but are not handled until the next instruction to complete in program order successfully completes. Throughout this chapter, the phrase next instruction implies the next instruction to complete in program order.

Note that exceptions can occur while an exception handler routine is executing, and multiple exceptions can become nested. It is up to the exception handler to save the states to allow control to ultimately return to the original excepting program.

Unless a catastrophic condition causes a system reset or machine check exception, only one exception is handled at a time. If, for example, a single instruction encounters multiple exception conditions, those conditions are handled sequentially. After the exception handler handles an exception, the instruction execution continues until the next exception condition is encountered. However, in many cases there is no attempt to re-execute the instruction. This method of recognizing and handling exception conditions sequentially guarantees that exceptions are recoverable.

To prevent loss of state information, exception handlers should save the information stored in SRR0 and SRR1 soon after the exception is taken. This prevents loss of information due to a system reset or machine check exception or to an instruction-caused exception in the exception handler before enabling external interrupts.

In this chapter, the following terminology is used to describe the various stages of exception processing:

Recognition	Exception recognition occurs when the condition that can cause an exception is identified by the processor.
Taken	An exception is said to be taken when control of instruction execution is passed to the exception handler; that is, the context is saved and the instruction at the appropriate vector offset is fetched and the exception handler routing is executed in supervisor mode.
Handling	Exception handling is performed by the software linked to the appropriate vector offset. Exception handling is performed at supervisor-level.

4.1 Exception Classes

The PowerPC architecture supports four types of exceptions:

- Synchronous, precise—These are caused by instructions. All instruction-caused exceptions are handled precisely; that is, the machine state at the time the exception occurs is known and can be completely restored. This means that (excluding the trap and system call exceptions) the address of the faulting instruction is provided to the exception handler and that neither the faulting instruction nor subsequent instructions in the code stream will complete execution before the exception is taken. Once the exception is processed, execution resumes at the address of the faulting instruction (or at an alternate address provided by the exception handler). When an exception is taken due to a trap or system call instruction, execution resumes at an address provided by the handler.
- Synchronous, imprecise—The PowerPC architecture defines two imprecise floating-point exception modes, recoverable and nonrecoverable. Even though the MPC603e provides a means to enable the imprecise modes, it implements these

modes identically to the precise mode (that is, all enabled floating-point exceptions are always precise on the MPC603e).

- Asynchronous, maskable—The external interrupt ($\overline{\text{INT}}$) , system management interrupt (SMI), and decrementer exceptions are maskable asynchronous exceptions. When these exceptions occur, their handling is postponed until the next instruction, and any exceptions associated with that instruction completes execution. If there are no instructions in the execution units, the exception is taken immediately upon determination of the correct restart address (for loading SRR0).
- Asynchronous, nonmaskable—There are two nonmaskable asynchronous exceptions: system reset and the machine check exception. These exceptions may not be recoverable, or may provide a limited degree of recoverability. All exceptions report recoverability through the MSR[RI] bit.

The MPC603e exception classes are shown in Table 4-1.

Table 4-1. Exception Classifications

Synchronous/Asynchronous	Precise/Imprecise	Exception Type
Asynchronous, nonmaskable	Imprecise	Machine check System reset
Asynchronous, maskable	Precise	External interrupt Decrementer System management interrupt
Synchronous	Precise	Instruction-caused exceptions

Table 4-1 define categories of exceptions that the MPC603e handles uniquely. Note that Table 4-1 includes no synchronous imprecise exceptions. While the PowerPC architecture supports imprecise handling of floating-point exceptions, the MPC603e implements floating-point exception modes as precise exceptions. Although the PowerPC architecture specifies that the recognition of the machine check exception is nonmaskable, on the MPC603e the stimuli that cause this exception are maskable. For example, the machine check exception is caused by the assertion of $\overline{\text{TEA}}$, $\overline{\text{APE}}$, $\overline{\text{DPE}}$, or $\overline{\text{MCP}}$. However, $\overline{\text{MCP}}$, $\overline{\text{APE}}$, and $\overline{\text{DPE}}$ can be disabled by bits 0, 2, and 3, respectively, in HID0. Therefore, the machine check caused by asserting $\overline{\text{TEA}}$ is the only truly nonmaskable machine check exception.

The MPC603e exceptions, and conditions that cause them, are listed in Table 4-2. Exceptions that are specific to either the PID6-603e or PID7t-603e are also noted in Table 4-2.

Table 4-2. Exceptions and Conditions

Exception Type	Vector Offset (hex)	Causing Conditions
Reserved	00000	—
System reset	00100	A system reset is caused by the assertion of either <u>SRESET</u> or <u>HRESET</u> .
Machine check	00200	A machine check is caused by the assertion of the <u>TEA</u> signal during a data bus transaction, assertion of MCP, or an address or data parity error.
DSI	00300	The cause of a DSI exception can be determined by the bit settings in the DSISR, listed as follows: 1 Set if the translation of an attempted access is not found in the primary hash table entry group (HTEG), or in the rehashed secondary HTEG, or in the range of a DBAT register; otherwise cleared. 4 Set if a memory access is not permitted by the page or DBAT protection mechanism; otherwise cleared. 5 Set by an eciwx or ecowx instruction if the access is to an address that is marked as write-through, or execution of a load/store instruction that accesses a direct-store segment. 6 Set for a store operation and cleared for a load operation. 11 Set if eciwx or ecowx is used and EAR[E] is cleared.
ISI	00400	An ISI exception is caused when an instruction fetch cannot be performed for any of the following reasons: <ul style="list-style-type: none"> The effective (logical) address cannot be translated. That is, there is a page fault for this portion of the translation, so an ISI exception must be taken to load the PTE (and possibly the page) into memory. The fetch access is to a direct-store segment (indicated by SRR1[3] set). The fetch access violates memory protection (indicated by SRR1[4] set). If the key bits (Ks and Kp) in the segment register and the PP bits in the PTE are set to prohibit read access, instructions cannot be fetched from this location.
External interrupt	00500	An external interrupt is caused when MSR[EE] = 1 and the <u>INT</u> signal is asserted.
Alignment	00600	An alignment exception is caused when the MPC603e cannot perform a memory access for any of the reasons described below: <ul style="list-style-type: none"> The operand of a floating-point load or store instruction is not word-aligned. The operand of lmw, stmw, lwarx, and stwcx. instructions are not aligned. The operand of a single-register load or store operation is not aligned, and the MPC603e is in little-endian mode (PID6-603e only). The execution of a floating-point load or store instruction to a direct-store segment. The operand of a load, store, load multiple, store multiple, load string, or store string instruction crosses a segment boundary into a direct-store segment, or crosses a protection boundary. Execution of a misaligned eciwx or ecowx instruction (PID7t-603e only). The instruction is lmw, stmw, lswi, lswx, stswi, stswx, and the MPC603e is in little-endian mode. The operand of dcbz is in memory that is write-through-required or caching-inhibited.

Table 4-2. Exceptions and Conditions (continued)

Exception Type	Vector Offset (hex)	Causing Conditions
Program	00700	<p>A program exception is caused by one of the following exception conditions, which correspond to bit settings in SRR1 and arise during execution of an instruction:</p> <p>Floating-point enabled exception—A floating-point enabled exception condition is generated when the following condition is met: $(MSR[FEO] \mid MSR[FE1]) \& FPSCR[FEX]$ is 1.</p> <ul style="list-style-type: none"> FPSCR[FEX] is set by the execution of a floating-point instruction that causes an enabled exception or by the execution of one of the ‘move to FPSCR’ instructions that results in both an exception condition bit and its corresponding enable bit being set in the FPSCR. Illegal instruction—An illegal instruction program exception is generated when execution of an instruction is attempted with an illegal opcode or illegal combination of opcode and extended opcode fields (including PowerPC instructions not implemented in the MPC603e), or when execution of an optional instruction not provided in the MPC603e is attempted (these do not include those optional instructions that are treated as no-ops). Privileged instruction—A privileged instruction type program exception is generated when the execution of a privileged instruction is attempted and the MSR register user privilege bit, MSR[PR], is set. In the MPC603e, this exception is generated for mtspr or mfspr with an invalid SPR field if SPR[0] = 1 and MSR[PR] = 1. This may not be true for all processors that implement the PowerPC architecture. Trap—A trap type program exception is generated when any of the conditions specified in a trap instruction is met.
Floating-point unavailable	00800	A floating-point unavailable exception is caused by an attempt to execute a floating-point instruction (including floating-point load, store, and move instructions) when the floating-point available bit is cleared (MSR[FP] = 0).
Decrementer	00900	The decrementer exception occurs when DEC[31] changes from 0 to 1. This exception is enabled with MSR[EE].
Reserved	00A00–00BFF	—
System call	00C00	A system call exception occurs when a System Call (sc) instruction is executed.
Trace	00D00	A trace exception is taken when MSR[SE] =1 or when the currently completing instruction is a branch and MSR[BE] =1.
Reserved	00E00	The MPC603e does not generate an exception to this vector. Other processors that implement the PowerPC architecture may use this vector for floating-point assist exceptions.
Reserved	00E10–00FFF	—
Instruction translation miss	01000	An instruction translation miss exception is caused when the effective address for an instruction fetch cannot be translated by the ITLB.
Data load translation miss	01100	A data load translation miss exception is caused when the effective address for a data load operation cannot be translated by the DTLB.
Data store translation miss	01200	A data store translation miss exception is caused when the effective address for a data store operation cannot be translated by the DTLB, or where a DTLB hit occurs, and the change bit in the PTE must be set due to a data store operation.
Instruction address breakpoint	01300	An instruction address breakpoint exception occurs when the address (bits 0–29) in the IABR matches the next instruction to complete in the completion unit, and the IABR enable bit (bit 30) is set.

Chapter 4. Exceptions

Table 4-2. Exceptions and Conditions (continued)

Exception Type	Vector Offset (hex)	Causing Conditions
System management interrupt	01400	A system management interrupt is caused when MSR[EE] = 1 and the <u>SMI</u> input signal is asserted.
Reserved	01500–02FFF	—

Exceptions are roughly prioritized by exception class, as follows:

1. Nonmaskable, asynchronous exceptions have priority over all other exceptions—system reset and machine check exceptions (although the machine check exception condition can be disabled so the condition causes the processor to go directly into the checkstop state). These exceptions cannot be delayed, and do not wait for the completion of any precise exception handling.
2. Synchronous, precise exceptions are caused by instructions and are taken in strict program order.
3. Maskable asynchronous exceptions (for example, external interrupt and decrementer exceptions) are delayed until higher priority exceptions are taken.

System reset and machine check exceptions may occur at any time and are not delayed even if an exception is being handled. As a result, state information for the interrupted exception may be lost; therefore, these exceptions are typically nonrecoverable.

All other exceptions have lower priority than system reset and machine check exceptions, and the exception may not be taken immediately when it is recognized.

4.1.1 Exception Priorities

The exceptions are listed in Table 4-3 in order of highest to lowest priority.

Table 4-3. Exception Priorities

Exception Category	Priority	Exception	Cause
Asynchronous	0	System reset	HRESET or power-on reset
	1	Machine check	<u>TEA</u> , <u>MCP</u> , <u>APE</u> , or <u>DPE</u>
	2	System reset	SRESET
	3	System management interrupt	<u>SMI</u>
	4	External interrupt	<u>INT</u>
	5	Decrementer exception	Decrementer passed through 0x00000000
Instruction fetch	0	ITLB miss	Instruction TLB miss
	1	Instruction access	Instruction access exception

Table 4-3. Exception Priorities (continued)

Exception Category	Priority	Exception	Cause
Instruction dispatch/ execution	0	IABR	Instruction address breakpoint exception
	1	Program	Program exception due to the following: <ul style="list-style-type: none">• Illegal instruction• Privileged instruction• Trap
	2	System call	System call exception
	3	Floating-point unavailable	Floating-point unavailable exception
	4	Program	Program exception due to a floating-point enabled exception
	5	Alignment	Alignment exception due to the following: <ul style="list-style-type: none">• Floating-point not word-aligned• lmw, stmw, lwarx, or stwcx. not word-aligned• Little-endian access is misaligned• eciw or ecow with misaligned operand (PID7T-603e only)• Multiple or string access with little-endian bit set
	6	Data access	Data access exception due to a BAT page protection violation
	7	Data access	Data access exception due to the following: <ul style="list-style-type: none">• eciw, ecow, lwarx, or stwcx. to direct-store segment (bit 5 of DSISR)• Crossing from memory segment to direct-store segment (bit 0 of DSISR)• Crossing from direct-store segment to memory segment• Any access to direct-store, SR[T] = 1• eciw or ecow with EAR[E] = 0 (bit 11 of DSISR)
	8	DTLB miss	Data TLB miss exception due to: <ul style="list-style-type: none">• Store miss• Load miss
	9	Alignment	Alignment exception due to a dcbz to a write-through or caching-inhibited page
	10	Data access	Data access exception due to TLB page protection violation
	11	DTLB miss	Data TLB miss exception due to a change bit not set on a store operation
Post-instruction execution	0	Trace	Trace exception due to the following: <ul style="list-style-type: none">• MSR[SE] = 1• MSR[BE] = 1 for branches

Exception priorities are described in detail in “Exception Priorities,” in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.

4.1.2 Summary of Front-End Exception Handling

The following list of interrupt categories describes how the MPC603e handles exceptions up to the point of signaling the appropriate exception to occur. Note that a recoverable state is reached if the completed store queue is empty (drained, not canceled) and any instruction that is next in program order and has been signaled to complete has completed. If MSR[RI] is clear, the MPC603e is in a nonrecoverable state by default. Also, completion of an instruction is defined as performing all architectural register writes associated with that instruction, and then removing that instruction from the completion buffer queue.

- Asynchronous nonmaskable nonrecoverable—(system reset caused by the assertion of either HRESET or internally during power-on reset (POR)). These exceptions have highest priority and are taken immediately regardless of other pending exceptions or recoverability. A nonpredicted address is guaranteed.
- Asynchronous maskable nonrecoverable—(machine check). A machine check exception takes priority over any other pending exception except a nonrecoverable system reset caused by the assertion of either HRESET or internally during POR. A machine check exception is taken immediately regardless of recoverability. A machine check exception can occur only if the machine check enable bit, MSR[ME], is set. If MSR[ME] is cleared, the processor goes directly into checkstop state when a machine check exception condition occurs. A nonpredicted address is guaranteed.
- Asynchronous nonmaskable recoverable—(system reset caused by the assertion of SRESET). This interrupt takes priority over any other pending exceptions except nonrecoverable exceptions listed above. This exception is taken immediately when a recoverable state is reached.
- Asynchronous maskable recoverable—(system management interrupt, external interrupt, decrementer exception). Before handling this type of exception, the next instruction in program order must complete or except. If this action causes another type of exception, that exception is taken and the asynchronous maskable recoverable exception remains pending. Once an instruction can complete without causing an exception, further instruction completion is halted while the exception not taken remains pending. The exception is taken when a recoverable state is reached.
- Instruction fetch—(ITLB, ISI). When this type of exception is detected, dispatch is halted and the current instruction stream is allowed to drain. If completing any instructions in this stream causes an exception, that exception is taken and the instruction fetch exception is forgotten. Otherwise, as soon as the machine is empty and a recoverable state is reached, the instruction fetch exception is taken.
- Instruction dispatch/execution—(program, DSI, alignment, emulation trap, system call, DTLB miss on load or store, IABR). This type of exception is determined at dispatch or execution of an instruction. The exception remains pending until all instructions in program order before the exception-causing instruction are

completed. The exception is then taken without completing the exception-causing instruction. If any other exception condition is created in completing these previous instructions in the machine, that exception takes priority over the pending instruction dispatch/execution exception, which will then be forgotten.

- Post-instruction execution—(trace). This type of exception is generated following execution and completion of an instruction while a trace mode is enabled. If executing the instruction produces conditions for another type of interrupt, that exception is taken and the post-instruction execution exception is forgotten for that instruction.

4.2 Exception Processing

When an exception is taken, the processor uses the save/restore registers, SRR0 and SRR1, to save the contents of the machine state register for user-level mode and to identify where instruction execution should resume after the exception is handled.

4.2.1 Exception Processing Registers

The MPC603e implements the SRR0 and SRR1 registers that are used for saving processor state on an exception. Also, many exception parameters are determined by MSR bits.

4.2.1.1 SRR0 and SRR1 Bit Settings

When an exception occurs, SRR0 is set to point to the instruction at which instruction processing should resume when the exception handler returns control to the interrupted process. All instructions in the program flow preceding this one will have completed and no subsequent instruction will have completed. This may be the address of the instruction that caused the exception or the next one (as in the case of a system call exception). The instruction addressed can be determined from the exception type and status bits. This address is used to resume instruction processing in the interrupted process, typically when an **rfi** instruction is executed. The SRR0 register is shown in Figure 4-1.

SRR0 (Holds EA for Resuming Program Execution)	
0	31

Figure 4-1. Machine Status Save/Restore Register 0 (SSR0)

The save/restore register 1 (SRR1) is used to save machine status (the contents of the MSR) on exceptions and to restore those values when **rfi** is executed. SRR1 is shown in Figure 4-2.

Exception-Specific Information and MSR Bit Values	
0	31

Figure 4-2. Machine Status Save/Restore Register 1 (SSR1)

Typically, when an exception occurs, bits 0–15 of SRR1 are loaded with exception-specific information and bits 16–31 of MSR are placed into the corresponding bit positions of SRR1. The MPC603e loads SRR1 with specific bits for handling machine check exceptions, as shown in Table 4-4.

Table 4-4. SRR1 Bit Settings for Machine Check Exceptions

Bits	Name	Description
0	MSR[0]	Copy of MSR bit 0
1–4	—	Reserved
5–9	MSR[5–9]	Copy of MSR bits 5–9
10–11	—	Reserved
12	MCP	Machine check
13	TEA	TEA error
14	DPE	Data parity error
15	APE	Address parity error
16–31	MSR[16–31]	Copy of MSR bits 16–31

The MPC603e loads SRR1 with specific bits for handling the three TLB miss exceptions, as shown in Table 4-5.

Table 4-5. SRR1 Bit Settings for Software Table Search Operations

Bits	Name	Description
0–3	CRF0	Copy of condition register field 0 (CR0)
4	—	Reserved
5–9	MSR[5–9]	Copy of MSR bits 5–9
10–11	—	Reserved
12	KEY	TLB miss protection key
13	I/D	Instruction/data TLB miss 0 DTLB miss 1 ITLB miss
14	WAY	Bit 14 indicates which TLB associativity set should be replaced 0 Set 0 1 Set 1
15	S/L	Store/load protection instruction 0 Load miss 1 Store miss
16–31	MSR[16–31]	Copy of MSR bits 16–31

Note that in some implementations, every instruction fetch when $\text{MSR}[\text{IR}] = 1$ and every instruction execution requiring address translation when $\text{MSR}[\text{DR}] = 1$, may modify SRR1.

4.2.1.2 MSR Bit Settings

The MSR is shown in Figure 4-3. When an exception occurs, MSR bits, as described in Table 4-6, are altered as determined by the exception.

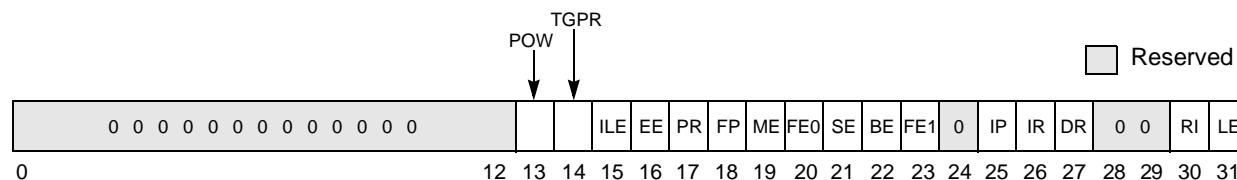


Figure 4-3. Machine State Register (MSR)

Table 4-6 shows the bit definitions for the MSR. Full function reserved bits are saved in SRR1 when an exception occurs; partial function reserved bits are not saved.

Table 4-6. MSR Bit Settings

Bits	Name	Description
0	—	Reserved. Full function.
1–4	—	Reserved. Partial function.
5–9	—	Reserved. Full function.
10–12	—	Reserved. Partial function.
13	POW	Power management enable (MPC603e-specific). 0 Disables programmable power modes (normal operation mode). 1 Enables programmable power modes (nap, doze, or sleep mode). This bit controls the programmable power modes only; it has no effect on dynamic power management (DPM). MSR[POW] may be altered with an mtmsr instruction only. Also, when altering the POW bit, software may alter only this bit in the MSR and no others. The mtmsr instruction must be followed by a context-synchronizing instruction. See Chapter 9, “Power Management,” for more information.
14	TGPR	Temporary GPR remapping (MPC603e-specific) 0 Normal operation. 1 TGPR mode. GPR0–GPR3 are remapped to TGPR0–TGPR3 for use by TLB miss routines. The contents of GPR0–GPR3 remain unchanged while MSR[TGPR] = 1. Attempts to use GPR4–GPR31 with MSR[TGPR] = 1 yield undefined results. Temporarily replaces TGPR0–TGPR3 with GPR0–GPR3 for use by TLB miss routines. The TGPR bit is set when either an instruction TLB miss, data read miss, or data write miss exception is taken. The TGPR bit is cleared by an rfi instruction.
15	ILE	Exception little-endian mode. When an exception occurs, this bit is copied into MSR[LE] to select the endian mode for the context established by the exception.
16	EE	External interrupt enable 0 The processor ignores external interrupts, system management interrupts, and decrementer interrupts. 1 The processor is enabled to take an external interrupt, system management interrupt, or decrementer interrupt.
17	PR	Privilege level 0 The processor can execute both user- and supervisor-level instructions. 1 The processor can only execute user-level instructions.

Table 4-6. MSR Bit Settings (continued)

Bits	Name	Description
18	FP	Floating-point available 0 The processor prevents dispatch of floating-point instructions, including floating-point loads, stores, and moves. 1 The processor can execute floating-point instructions, and can take floating-point enabled exception type program exceptions.
19	ME	Machine check enable 0 Machine check exceptions are disabled. 1 Machine check exceptions are enabled.
20	FE0	Floating-point exception mode 0 (see Table 4-7)
21	SE	Single-step trace enable 0 The processor executes instructions normally. 1 The processor generates a trace exception upon the successful completion of the next instruction.
22	BE	Branch trace enable 0 The processor executes branch instructions normally. 1 The processor generates a trace exception upon the successful completion of a branch instruction.
23	FE1	Floating-point exception mode 1 (see Table 4-7)
24	—	Reserved
25	IP	Exception prefix. The setting of this bit specifies whether an exception vector offset is prepended with Fs or 0s. In the following description, <i>n</i> is the offset of the exception. See Table 4-2. 0 Exceptions are vectored to the physical address 0x000n_ <i>nnnn</i> . 1 Exceptions are vectored to the physical address 0xFFFFn_ <i>nnnn</i> .
26	IR	Instruction address translation 0 Instruction address translation is disabled. 1 Instruction address translation is enabled. See Chapter 5, “Memory Management.”
27	DR	Data address translation 0 Data address translation is disabled. 1 Data address translation is enabled. See Chapter 5, “Memory Management.”
28–29	—	Reserved. Full function.
30	RI	Recoverable exception (for system reset and machine check exceptions) 0 Exception is not recoverable. 1 Exception is recoverable.
31	LE	Little-endian mode enable 0 The processor runs in big-endian mode. 1 The processor runs in little-endian mode.

The IEEE floating-point exception mode bits (FE0 and FE1) together define whether floating-point exceptions are handled precisely, imprecisely, or if they are taken at all. The possible settings and default conditions for the MPC603e are shown in Table 4-7. For further details, see Chapter 6, “Exceptions,” of the *Programming Environments Manual*.

MSR bits are guaranteed to be written to SRR1 when the first instruction of the exception handler is encountered.

Table 4-7. IEEE Floating-Point Exception Mode Bits

FE0	FE1	Mode
0	0	Floating-point exceptions disabled
0	1	Floating-point imprecise nonrecoverable ¹
1	0	Floating-point imprecise recoverable ¹
1	1	Floating-point precise mode

¹ Not implemented in the MPC603e.

4.2.2 Enabling and Disabling Exceptions

When a condition exists that may cause an exception to be generated, it must be determined whether the exception is enabled for that condition as follows:

- IEEE floating-point enabled exceptions (a type of program exception) are ignored when both MSR[FE0] and MSR[FE1] are cleared. If either of these bits are set, all IEEE enabled floating-point exceptions are taken and cause a program exception.
- Asynchronous, maskable exceptions (that is, the external, system management, and decrementer interrupts) are enabled by setting the MSR[EE] bit. When MSR[EE] = 0, recognition of these exception conditions is delayed. MSR[EE] is cleared automatically when an exception is taken, to delay recognition of conditions causing those exceptions.
- A machine check exception can occur only if the machine check enable bit, MSR[ME], is set. If MSR[ME] is cleared, the processor goes directly into checkstop state when a machine check exception condition occurs. Individual machine check exceptions can be enabled and disabled through bits in the HID0 register, as described in Table 2-2.
- System reset exceptions cannot be masked.

4.2.3 Steps for Exception Processing

After it is determined that the exception can be taken (by confirming that any instruction-caused exceptions occurring earlier in the instruction stream have been handled, and by confirming that the exception is enabled for the exception condition), the processor does the following:

1. The machine status save/restore register 0 (SRR0) is loaded with an instruction address that depends on the type of exception. See the individual exception description for details about how this register is used for specific exceptions.
2. SRR1[1–4, 10–15] are loaded with information specific to the exception type.

3. SRR1[5–9, 16–31] are loaded with a copy of the corresponding bits of the MSR.
4. The MSR is set as described in Table 4-6. The new values take effect beginning with the fetching of the first instruction of the exception-handler routine located at the exception vector address.

Note that MSR[IR] and MSR[DR] are cleared for all exception types; therefore, address translation is disabled for both instruction fetches and data accesses beginning with the first instruction of the exception-handler routine.

5. Instruction fetch and execution resumes, using the new MSR value, at a location specific to the exception type. The location is determined by adding the exception's vector (see Table 4-2) to the base address determined by MSR[IP]. If IP is cleared, exceptions are vectored to the physical address $0x000n_nnnn$. If IP is set, exceptions are vectored to the physical address $0xFFFn_nnnn$. For a machine check exception that occurs when $\text{MSR}[\text{ME}] = 0$ (machine check exceptions are disabled), the processor enters the checkstop state (the machine stops executing instructions). See Section 4.5.2, “Machine Check Exception (0x00200).”

4.2.4 Setting MSR[RI]

The operating system should handle MSR[RI] as follows:

- In the machine check and system reset exceptions—if SRR1[RI] is cleared, the exception is not recoverable. If it is set, the exception is recoverable with respect to the processor.
- In each exception handler—When enough state information has been saved that a machine check or system reset exception can reconstruct the previous state, set MSR[RI].
- In each exception handler—Clear MSR[RI], set the SRR0 and SRR1 registers appropriately, and then execute **rfi**.
- Note that the RI bit being set indicates that, with respect to the processor, enough processor state data is valid for the processor to continue, but it does not guarantee that the interrupted process can resume.

4.2.5 Returning from an Exception Handler

The Return from Interrupt (**rfi**) instruction performs context synchronization by allowing previously issued instructions to complete before returning to the interrupted process. In general, execution of the **rfi** instruction ensures the following:

- All previous instructions have completed to a point where they can no longer cause an exception. If a previous instruction causes a direct-store interface error exception, the results must be determined before this instruction is executed.

- Previous instructions complete execution in the context (privilege, protection, and address translation) under which they were issued.
- The **rfi** instruction copies SRR1 bits back into the MSR.
- The instructions following this instruction execute in the context established by this instruction.

For a complete description of context synchronization, refer to Chapter 6, “Exceptions,” of the *Programming Environments Manual*.

4.3 Process Switching

The operating system should execute one of the following when processes are switched:

- The **sync** instruction, which orders the effects of instruction execution. All instructions previously initiated appear to have completed before the **sync** instruction completes, and no subsequent instructions appear to be initiated until the **sync** instruction completes. For an example showing use of the **sync** instruction, see Chapter 2, “Register Set,” of the *Programming Environments Manual*.
- The **isync** instruction, which waits for all previous instructions to complete and then discards any fetched instructions, causing subsequent instructions to be fetched (or refetched) from memory and to execute in the context (privilege, translation, protection, etc.) established by the previous instructions.
- The **stwex.** instruction, to clear any outstanding reservations, which ensures that an **lwarx** instruction in the old process is not paired with an **stwex.** instruction in the new process.

The operating system should set the MSR[RI] bit as described in Section 4.2.4, “Setting MSR[RI].”

4.4 Exception Latencies

Latencies for taking various exceptions depend on the state of the machine when the exception conditions occur. This latency may be as short as one cycle, in which case an exception is signaled in the cycle following the appearance of the exception condition. The latencies are as follows:

- Hard reset and machine check—In most cases, a hard reset or machine check exception will have a single-cycle latency. A two- to three-cycle delay may occur only when a predicted instruction is next to complete, and the branch guess that forced this instruction to be predicted was resolved to be incorrect.
- Soft reset—The latency of a soft reset exception is affected by recoverability. The time to reach a recoverable state may depend on the time needed to complete or except an instruction at the point of completion, the time needed to drain the completed store queue, and the time waiting for a correct empty state so that a valid

MSR[IP] may be saved. For lower-priority externally-generated interrupts, a delay may be incurred waiting for another interrupt, generated while reaching a recoverable state, to be serviced.

Further delays are possible for other types of exceptions depending on the number and type of instructions that must be completed before those exceptions may be serviced. See Section 4.1.2, “Summary of Front-End Exception Handling,” to determine possible maximum latencies for different exceptions.

4.5 Exception Definitions

Table 4-8 shows all the types of exceptions that can occur with the MPC603e and the MSR bit settings when the processor transitions to supervisor mode. The state of these bits prior to the exception is typically stored in SRR1.

Table 4-8. MSR Setting Due to Exception

Exception Type	MSR Bit															
	POW	TGPR	ILE	EE	PR	FP	ME	FEO	SE	BE	FE1	IP	IR	DR	RI	LE
System reset	0	0	—	0	0	0	—	0	0	0	0	1	0	0	0	ILE
Machine check	0	0	—	0	0	0	0	0	0	0	0	—	0	0	0	ILE
DSI	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
ISI	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
External	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
Alignment	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
Program	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
Floating-point unavailable	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
Decrementer	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
System call	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
Trace exception	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
ITLB miss	0	1	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
DTLB miss on load	0	1	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
DTLB miss on store	0	1	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE
Instruction address breakpoint	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE

Table 4-8. MSR Setting Due to Exception (continued)

Exception Type	MSR Bit															
	POW	TGPR	ILE	EE	PR	FP	ME	FEO	SE	BE	FE1	IP	IR	DR	RI	LE
System management interrupt	0	0	—	0	0	0	—	0	0	0	0	—	0	0	0	ILE

Note:

0 Bit is cleared.

1 Bit is set.

ILE Bit is copied from the ILE bit in the MSR.

— Bit is not altered.

Reserved bits are read as if written as 0.

4.5.1 Reset Exceptions (0x00100)

The system reset exception is a nonmaskable, asynchronous exception signaled to the MPC603e either through the assertion of the reset signals (SRESET or HRESET) or internally during the power-on reset (POR) process. The assertion of the soft reset signal, SRESET, as described in Section 7.2.9.6.2, “Soft Reset (SRESET)—Input,” causes the soft reset exception to be taken and the physical base address of the handler is determined by the MSR[IP] bit. The assertion of the hard reset signal, HRESET, as described in Section 7.2.9.6.1, “Hard Reset (HRESET)—Input,” causes the hard reset exception to be taken and the physical address of the handler is always 0xFFFF0_0100.

4.5.1.1 Hard Reset and Power-On Reset

As described in Section 4.1.2, “Summary of Front-End Exception Handling,” the hard reset exception is a nonrecoverable, nonmaskable asynchronous exception (maskable interrupt). When HRESET is asserted or at power-on reset (POR), the MPC603e immediately branches to 0xFFFF0_0100 without attempting to reach a recoverable state. A hard reset has the highest priority of any exception, and is always nonrecoverable. Table 4-9 shows the state of the machine just before it fetches the first instruction of the system reset handler after a hard reset.

Table 4-9. Settings Caused by Hard Reset

Register	Setting	Register	Setting
GPRs	Unknown	PVR	0003000n
FPRs	Unknown	HID0	00000000
FPSCR	00000000	HID1	00000000
CR	All 0s	DMISS and IMISS	All 0s
SRs	Unknown	DCMP and ICMP	All 0s
MSR	00000040	RPA	All 0s

Table 4-9. Settings Caused by Hard Reset (continued)

Register	Setting	Register	Setting
XER	00000000	IABR	All 0s
TBU	00000000	DSISR	00000000
TBL	00000000	DAR	00000000
LR	00000000	DEC	FFFFFF
CTR	00000000	HASH1	00000000
SDR1	00000000	HASH2	00000000
SRR0	00000000	TLBs	Unknown
SRR1	00000000	Cache	All cache blocks invalidated
SPRGs	00000000	BATs	Unknown
Tag directory	All 0s. (However, LRU bits are initialized so each side of the cache has a unique LRU value.)		

The **HRESET** signal can be asserted for the following reasons:

- System power-on reset
- System reset from a panel switch

For information on the **HRESET** signal, see Section 7.2.9.6.1, “Hard Reset (**HRESET**)—Input.”

The following is also true after a hard reset operation:

- External checkstops are enabled.
- The on-chip test interface has given control of the I/Os to the rest of the chip for functional use.
- Since the reset exception has data and instruction translation disabled (MSR[DR] and MSR[IR] both cleared), the chip operates in real addressing mode as described in Section 5.2, “Real Addressing Mode.”

4.5.1.2 Soft Reset

As described in Section 4.1.2, “Summary of Front-End Exception Handling,” the soft reset exception is a type of system reset exception that is recoverable, nonmaskable, and asynchronous. When **SRESET** is asserted, the processor attempts to reach a recoverable state by allowing the next instruction to either complete or cause an exception, blocking the completion of subsequent instructions, and allowing the completed store queue to drain.

Unlike a hard reset, no registers or latches are initialized; however, the instruction cache is disabled (HID0[ICE] = 0). After **SRESET** is recognized as asserted, the processor begins fetching instructions from the system reset routine at offset 0x0100. When a soft reset

occurs, registers are set as shown in Table 4-10. A soft reset is recoverable provided that attaining the recoverable state does not cause a machine check exception. This interrupt case is third in priority, following hard reset and machine check.

When a soft reset occurs, registers are set as shown in Table 4-10 in addition to the clearing of HID0[ICE].

Table 4-10. Soft Reset Exception—Register Settings

Register	Setting Description						
SRR0	Set to the effective address of the instruction that the processor would have attempted to complete next if no exception conditions were present.						
SRR1	0–15 Cleared 16–31 Loaded from MSR[16–31]. Note that if the processor state is corrupted to the extent that execution cannot be reliably restarted, SRR1[30] is cleared.						
MSR	POW 0	EE 0	FEO 0	IR 0			
	TGPR 0	PR 0	SE 0	DR 0			
	ILE —	FP 0	BE 0	RI 0			
	IP —	ME —	FE1 0	LE	Set to value of ILE		

4.5.2 Machine Check Exception (0x00200)

The MPC603e conditionally initiates a machine check exception after detecting the assertion of the TEA or MCP signals on the MPC603e bus (assuming the machine check is enabled with MSR[ME] = 1). The assertion of one of these signals indicates that a bus error occurred and the system terminates the current transaction. One clock cycle after the signal is asserted, the data bus signals go to the high-impedance state; however, data entering the GPR or the cache is not invalidated. Note that if HID0[EMCP] is cleared, the processor ignores the assertion of the MCP signal.

A machine check exception also occurs when an address or data parity error is detected on the bus and the address or data parity error is enabled in HID0. See Section 2.1.2.1, “Hardware Implementation Registers (HID0 and HID1),” for more information.

Note that the MPC603e makes no attempt to force recoverability on a machine check; however, it does guarantee that the machine check exception is always taken immediately upon request, with a nonpredicted address saved in SRR0, regardless of the current machine state. Because pending stores in the completed store queue are not canceled when a machine check exception occurs, two consecutive stores that result in the assertion of TEA can cause the processor to checkstop. To prevent a checkstop in this case, a **sync** instruction must be placed between two stores that can result in assertion of TEA.

Software can use the machine check exception in a recoverable mode to probe memory. For this case, a **sync**, load, **sync** instruction sequence is used. If the load access results in a system error (for example, the assertion of TEA), the processor can handle this in a recoverable state. If the **sync** instruction is not used, a second access to the same address as the first load could cause the processor to enter the checkstop state.

If the MSR[ME] bit is set, the exception is recognized and handled; otherwise, the MPC603e attempts to enter an internal checkstop. Note that the resulting machine check exception has priority over any exceptions caused by the instruction that generated the bus operation.

Machine check exceptions are only enabled when MSR[ME] = 1; this is described in Section 4.5.2.1, “Machine Check Exception Enabled (MSR[ME] = 1).” If MSR[ME] = 0 and a machine check occurs, the processor enters the checkstop state. Checkstop state is described in 4.5.2.2, “Checkstop State (MSR[ME] = 0).”

4.5.2.1 Machine Check Exception Enabled (MSR[ME] = 1)

When a machine check exception is taken, registers are updated as shown in Table 4-11.

When a machine check exception is taken, instruction execution for the handler begins at offset 0x00200 from the physical base address indicated by MSR[IP].

In order to return to the main program, the exception handler should do the following:

1. SRR0 and SRR1 should be given the values to be used by the **rfi** instruction
2. Execute **rfi**

Table 4-11. Machine Check Exception—Register Settings

Register	Setting Description							
SRR0	Set to the address of the next instruction that would have been completed in the interrupted instruction stream. Neither this instruction nor any others beyond it will have been completed. All preceding instructions will have been completed.							
SRR1	0–11 Cleared 12 MCP—Machine check signal caused exception 13 TEA—Transfer error acknowledge signal caused exception 14 DPE—Data parity error condition (and signal assertion) caused exception 15 APE—Address parity error condition (and signal assertion) caused exception 16–31 Loaded from MSR[16–31]							
MSR	POW 0 EE 0 FE0 0 IR 0 TGPR 0 PR 0 SE 0 DR 0 ILE — FP 0 BE 0 RI 0 IP — ME 0 FE1 0 LE Set to value of ILE							

Note: When a machine check exception is taken, the exception handler should set MSR[ME] as soon as it is practical to handle another TEA assertion. Otherwise, subsequent TEA assertions cause the processor to automatically enter the checkstop state.

4.5.2.2 Checkstop State (MSR[ME] = 0)

When the MPC603e enters the checkstop state, it asserts the checkstop output signal, CKSTP_OUT. The following events will cause the MPC603e to enter the checkstop state:

- Machine check exception occurs with MSR[ME] cleared.
- External checkstop input, CKSTP_IN, is asserted.

When a processor is in the checkstop state, instruction processing is suspended and generally cannot be restarted without resetting the processor. The contents of all latches are frozen within two cycles upon entering the checkstop state so that the state of the processor can be analyzed as an aid in problem determination.

Note that not all processors that implement the PowerPC architecture provide the same level of error checking. The reasons a processor can enter checkstop state are implementation-dependent.

4.5.3 DSI Exception (0x00300)

A DSI exception occurs when no higher priority exception exists and a data memory access cannot be performed. The condition that caused the DSI exception can be determined by reading the DSISR register, a supervisor-level SPR (SPR18) that can be read by using the **mfspr** instruction. Bit settings are provided in Table 4-12. Table 4-12 also indicates the memory element that is saved to the DAR.

Table 4-12. DSI Exception—Register Settings

Register	Setting Description												
SRR0	Set to the effective address of the instruction that caused the exception.												
SRR1	0–15 Cleared 16–31 Loaded with MSR[16–31]												
MSR	POW 0	EE 0	FE0 0	IR 0									
	TGPR 0	PR 0	SE 0	DR 0									
	ILE —	FP 0	BE 0	RI 0									
	IP —	ME —	FE1 0	LE Set to value of ILE									
DSISR	0	Set if a load or store instruction results in a direct-store error exception.											
	1	Set by the data TLB miss exception handler if the translation of an attempted access is not found in the primary hash table entry group (HTEG), or in the rehashed secondary HTEG, or in the range of a DBAT register; otherwise cleared.											
	2–3	Cleared											
	4	Set if a memory access is not permitted by the page or BAT protection mechanism; otherwise cleared.											
	5	Set if the lwarx or stwcx . instruction is attempted to direct-store space.											
	6	Set for a store operation and cleared for a load operation.											
	7–31	Cleared											
DAR	Set to the effective address of a memory element as described in the following list:												
	<ul style="list-style-type: none"> • A byte in the first word accessed in the page that caused the DSI exception, for a byte, half word, or word memory access. • A byte in the first word accessed in the BAT area that caused the DSI exception for a byte, half word, or word access to a BAT area. • A byte in the block that caused the exception for icbi, dcbz, dcbst, dcbf, or dcbi instructions. • Any EA in the memory range addressed (for direct-store exceptions). 												

DSI exceptions can occur for any of the following reasons.

- The instruction is not supported for the type of memory addressed.
- Any access to a direct-store segment (SR[T] = 1).

- The attempted access violates the memory protection defined by SR[Ks,Kp], PTE[PP], or DBATn[PP].

Note that the OEA specifies an additional case that may cause a DS1 exception—when an effective address for a load, store, or cache operation cannot be translated by the TLBs. On the MPC603e, this condition causes a TLB miss exception instead. These scenarios are common among all processors that implement the PowerPC architecture. The following additional scenarios can cause a DS1 exception in the MPC603e:

- A bus error indicates crossing from a direct-store segment to a memory segment.
- The execution of any load/store instruction to a direct-store segment, SR[T] = 1.
- A data access crosses from a memory segment (SR[T] = 0) into a direct-store segment (SR[T] = 1).

DS1 exceptions can be generated by load/store instructions, and the cache control instructions (**dcbi**, **dcbz**, **dcbst**, and **dcbf**).

The MPC603e supports the crossing of page boundaries. However, if the second page has a translation error or protection violation associated with it, the MPC603e takes the DS1 exception in the middle of the instruction. In this case, the data address register (DAR) always points to a byte address in the first word of the offending page.

If an **stwex.** instruction has an effective address for which a normal store operation would cause a DS1 exception, the MPC603e takes the DS1 exception without checking for the reservation.

If the XER indicates that the byte count for an **lswi** or **stswi** instruction is zero, a DS1 exception does not occur, regardless of the effective address.

The condition that caused the exception is defined in the DSISR. These conditions also use the data address register (DAR) as shown in Table 4-12.

When a DS1 exception is taken, instruction execution for the handler begins at offset 0x00300 from the physical base address indicated by MSR[IP].

The architecture permits certain instructions to be partially executed when they cause a DS1 exception. These are as follows:

- Load multiple or load string instructions—some registers in the range of registers to be loaded may have been loaded.
- Store multiple or store string instructions—some bytes of memory in the range addressed may have been updated.

In these cases, the number of registers and amount of memory altered are instruction- and boundary-dependent. However, memory protection is not violated. Furthermore, if some of the data accessed is in direct-store space (SR[T] = 1) and the instruction is not supported for direct-store accesses, the locations in direct-store space are not accessed.

For update forms, the update register (**rA**) is not altered.

4.5.4 ISI Exception (0x00400)

The ISI exception is implemented as it is defined by the PowerPC architecture. An ISI exception occurs when no higher priority exception exists and an attempt to fetch the next instruction fails for any of the following reasons:

- If an instruction TLB miss fails to find the desired PTE, then a page fault is synthesized. The ITLB miss handler branches to the ISI exception handler to retrieve the translation from a storage device.
- An attempt is made to fetch an instruction from a direct-store segment while instruction translation is enabled (MSR[IR] = 1).
- An attempt is made to fetch an instruction from no-execute memory.
- An attempt is made to fetch an instruction from guarded memory when MSR[IR] = 1.
- The fetch access violates memory protection.

Register settings for this exception are described in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.

When an ISI exception is taken, instruction execution for the handler begins at offset 0x00400 from the physical base address indicated by MSR[IP].

4.5.5 External Interrupt (0x00500)

An external interrupt is signaled to the MPC603e by the assertion of the $\overline{\text{INT}}$ signal as described in Section 7.2.9.1, “Interrupt ($\overline{\text{INT}}$)—Input.” The interrupt may not be recognized if a higher priority exception occurs simultaneously or if the MSR[EE] bit is cleared when $\overline{\text{INT}}$ is asserted.

After the $\overline{\text{INT}}$ is recognized, the MPC603e generates a recoverable halt to instruction completion. The MPC603e allows the next instruction in program order to complete, including handling any exceptions that instruction may generate. However, the MPC603e blocks subsequent instructions from completing and allows any outstanding stores to occur to system memory. If any other exceptions are encountered in this process, they are taken first and the external interrupt is delayed until a recoverable halt is achieved. At this time, the MPC603e saves the state information and takes the external interrupt as defined by the PowerPC architecture.

The register settings for the external interrupt are shown in Table 4-13.

Table 4-13. External Interrupt—Register Settings

Register	Setting					
SRR0	Set to the effective address of the instruction that the processor would have attempted to execute next if no interrupt conditions were present.					
SRR1	0–15 Cleared 16–31 Loaded from MSR[16–31]					
MSR	POW 0	EE 0	FE0 0	IR 0	DR 0	RI 0
	TGPR 0	PR 0	SE 0	LE	Set to value of ILE	
	ILE —	FP 0	BE 0			
	IP —	ME —	FE1 0			

When an external interrupt is taken, instruction execution for the handler begins at offset 0x00500 from the physical base address indicated by MSR[IP].

The MPC603e only recognizes the interrupt condition ($\overline{\text{INT}}$ asserted) if the MSR[EE] bit is set; it ignores the interrupt condition if the MSR[EE] bit is cleared. To guarantee that the external interrupt is taken, the $\overline{\text{INT}}$ signal must be held asserted until the MPC603e takes the interrupt. If the $\overline{\text{INT}}$ signal is negated before the interrupt is taken, the MPC603e is not guaranteed to take an external interrupt. The interrupt handler must send a command to the device that asserted $\overline{\text{INT}}$, acknowledging the interrupt and instructing the device to negate $\overline{\text{INT}}$ before the handler re-enables recognition of external interrupts.

4.5.6 Alignment Exception (0x00600)

This section describes conditions that can cause alignment exceptions in the MPC603e. The MPC603e implements the alignment exception as it is defined in the PowerPC architecture. For information on bit settings and how exception conditions are detected, refer to the *Programming Environments Manual*. Note that the PowerPC architecture allows individual processors to determine whether an exception is required to handle various alignment conditions.

Similar to DSI exceptions, alignment exceptions use the SRR0 and SRR1 to save the machine state and the DSISR to determine the source of the exception. The MPC603e initiates an alignment exception when it detects any of the following conditions:

- The operand of a floating-point load or store operation is not word-aligned.
- The operand of an **lmw**, **stmw**, **lwarx**, or **stwcx** instruction is not word-aligned.
- A little-endian access (MSR[LE] = 1) is misaligned.
- A multiple or string access is attempted with the MSR[LE] bit set.
- The operand of a floating-point load or store operation is to a direct-store segment.
- The operand of an elementary, multiple or string load or store crosses a segment boundary with a change to the direct-store attribute (T bit different).
- The operand of a **eciwx** or **ecowx** instruction is not aligned (PID7t-603e only).

- The operand of a **dcbz** instruction is in a page that is write-through or caching-inhibited.

The register settings for alignment exceptions are shown in Table 4-13.

The architecture does not support the use of a misaligned EA by **lwarx** or **stwcx** instructions. If one of these instructions specifies a misaligned EA, the exception handler should not emulate the instruction, but should treat the occurrence as a programming error.

Table 4-14. Alignment Interrupt—Register Settings

Register	Setting						
SRR0	Set to the effective address of the instruction that caused the exception.						
SRR1	0–15 Cleared 16–31 Loaded from MSR[16–31]						
MSR	POW 0 EE 0 FE0 0 IR 0 TGPR 0 PR 0 SE 0 DR 0 ILE — FP 0 BE 0 RI 0 IP — ME — FE1 0 LE Set to value of ILE						
DSISR	0–11 Cleared 12–13 Cleared. (Note that these bits can be set by several 64-bit PowerPC instructions that are not supported in the MPC603e.) 14 Cleared 15–16 For instructions that use register indirect with index addressing—set to bits 29–30 of the instruction. For instructions that use register indirect with immediate index addressing—cleared. 17 For instructions that use register indirect with index addressing—set to bit 25 of the instruction. For instructions that use register indirect with immediate index addressing—set to bit 5 of the instruction. 18–21 For instructions that use register indirect with index addressing—set to bits 21–24 of the instruction. For instructions that use register indirect with immediate index addressing—set to bits 1–4 of the instruction. 22–26 Set to bits 6–10 (identifying either the source or destination) of the instruction. Undefined for dcbz . 27–31 Set to bits 11–15 of the instruction (rA) Set to either bits 11–15 of the instruction or to any register number not in the range of registers loaded by a valid form instruction, for lmw , lswi , and lswx instructions. Otherwise undefined.						
DAR	Set to the EA of the data access as computed by the instruction causing the alignment exception. When the operand of an lmw , stmw , lwarx , or stwcx instruction is not word-aligned, that address value +4 is stored into the DAR.						

4.5.6.1 Integer Alignment Exceptions

The MPC603e is optimized for load and store operations that are aligned on natural boundaries. Operations that are not naturally aligned may suffer performance degradation, depending on the type of operation, the boundaries crossed, and the mode that the processor is in during execution. More specifically, these operations may either cause an alignment exception or they may cause the processor to break the memory access into multiple, smaller accesses with respect to the cache and the memory subsystem.

The MPC603e can initiate an alignment exception for the access shown in Table 4-15. In this case, the appropriate range check is performed before the instruction begins execution. As a result, if an alignment exception is taken, it is guaranteed that no portion of the instruction has been executed.

Table 4-15. Access Types

MSR[DR]	SR[T]	Access Type
1	0	Page-address translation access

A page-address translation access occurs when MSR[DR] is set, SR[T] is cleared, and there is not a match in the BAT. Note the following points:

- The following is true for all loads and stores except strings/multiples:
 - Byte operands never cause an alignment exception.
 - Half-word operands can cause an alignment exception if the EA ends in 0xFFFF.
 - Word operands can cause an alignment exception if the EA ends in 0xFFD–FFF.
 - Double-word operands cause an alignment exception if the EA ends in 0xFF9–FFF.
- The **dcbz** instruction causes an alignment exception if the access is to a page or block with the W (write-through) or I (cache-inhibit) bit set in the TLB or BAT, respectively.

A misaligned memory access that does not cause an alignment exception will not perform as well as an aligned access of the same type. The resulting performance degradation due to misaligned accesses depends on how well each individual access behaves with respect to the memory hierarchy. At a minimum, additional cache access cycles are required that can delay other processor resources from using the cache. More dramatically, for an access to a noncacheable page, each discrete access involves individual processor bus operations that reduce the effective bandwidth of that bus.

Finally, note that when the MPC603e is in page address translation mode, there is no special handling for accesses that fall into BAT regions.

4.5.6.2 Load/Store Multiple Alignment Exceptions

Most alignment exceptions store the address as computed by the instruction in the DAR. However, when the operand of an **lmw**, **stmw**, **lwarx**, or **stwex**. instruction is not word-aligned that address value +4 is stored into the DAR.

4.5.7 Program Exception (0x00700)

The MPC603e implements the program exception as it is defined by the PowerPC architecture (OEA). A program exception occurs when no higher priority exception exists and one or more of the exception conditions defined in the OEA occur.

When a program exception is taken, instruction execution for the handler begins at offset 0x00700 from the physical base address indicated by MSR[IP]. The exception conditions are as follows:

- Floating-point enabled exception—These exceptions correspond to IEEE-defined exception conditions, such as overflows, and divide by zeros that may occur during the execution of a floating-point arithmetic instruction. As a group, these exceptions are enabled by the FE0 and FE1 bits in the MSR. Individual conditions are enabled by specific bits in the FPSCR. For general information about this exception, see the *Programming Environments Manual*. For more information about how these exceptions are implemented in the MPC603e, see Section 4.5.7.1, “IEEE Floating-Point Exception Program Exceptions.”
- Illegal instruction—An illegal instruction program exception is generated when execution of an instruction is attempted with an illegal opcode or illegal combination of opcode and extended opcode fields (including PowerPC instructions not implemented in the MPC603e). These do not include those optional instructions treated as no-ops.
- Privileged instruction—A privileged instruction type program exception is generated when the execution of a privileged instruction is attempted and the MSR register user privilege bit, MSR[PR], is set. In the MPC603e, this exception is generated for **mtspr** or **mfspr** with an invalid SPR field if SPR[0] = 1 and MSR[PR] = 1. This may not be true for all processors that implement the PowerPC architecture.
- Trap—A trap type program exception is generated when any of the conditions specified in a trap instruction is met.

4.5.7.1 IEEE Floating-Point Exception Program Exceptions

Floating-point exceptions are signaled by condition bits set in the floating-point status and control register (FPSCR). They can cause the system floating-point enabled exception handler to be invoked. The MPC603e handles all floating-point exceptions precisely. The MPC603e implements the FPSCR as it is defined by the PowerPC architecture; for more information about the FPSCR, see the *Programming Environments Manual*.

Floating-point operations that change exception sticky bits in the FPSCR may suffer a performance penalty. When an exception is disabled in the FPSCR and MSR[FE] = 0, updates to the FPSCR exception sticky bits are serialized at the completion stage. This serialization may result in a one- or two-cycle execution delay. The penalty is incurred only

when the exception bit is changed and not on subsequent operations with the same exception. See Chapter 6, “Instruction Timing,” for a full description of completion serialization.

When an exception is enabled in the FPSCR, the instruction traps to the emulation trap exception vector without updating the FPSCR or the target FPR. The emulation trap exception handler is required to complete the instruction. The emulation trap exception handler is invoked regardless of the FE setting in the MSR.

The two IEEE floating-point imprecise modes, defined by the PowerPC architecture when $\text{MSR}[\text{FE0}] \neq \text{MSR}[\text{FE1}]$, are treated as precise exceptions (that is, $\text{MSR}[\text{FE0}] = \text{MSR}[\text{FE1}] = 1$). This is regardless of the setting of $\text{MSR}[\text{NI}]$.

For the highest and most predictable floating-point performance, all exceptions should be disabled in the FPSCR and MSR. For more information about the program exception, see the *Programming Environments Manual*.

4.5.7.2 Illegal, Reserved, and Unimplemented Instructions Program Exceptions

In accordance with the PowerPC architecture, the MPC603e considers all instructions defined for 64-bit implementations and unimplemented optional instructions, such as **fsqrt**, **eciwx**, and **ecowx** as illegal and takes a program exception when one of these instructions is encountered. Likewise, if a supervisor-level instruction is encountered when the processor is in user-level mode, a privileged instruction-type program exception is taken.

The MPC603e implements some instructions, such as double-precision floating-point and load/store string instructions in software. These instructions take the MPC603e-specific emulation trap exception (0x01600) rather than a program exception.

4.5.8 Floating-Point Unavailable Exception (0x00800)

A floating-point unavailable exception occurs when no higher priority exception exists, an attempt is made to execute a floating-point instruction (including floating-point load, store, and move instructions), and the floating-point available bit in the MSR is disabled ($\text{MSR}[\text{FP}] = 0$). Register settings for this exception are described in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.

When a floating-point unavailable exception is taken, instruction execution for the handler begins at offset 0x00800 from the physical base address indicated by $\text{MSR}[\text{IP}]$.

4.5.9 Decrementer Exception (0x00900)

The MPC603e implements the decrementer interrupt exception as it is defined in the PowerPC architecture. A decrementer exception request is made when the decrementer counts down through zero. The request is held until there are no higher priority exceptions

and $\text{MSR[EE]} = 1$. At this point the decrementer exception is taken. If multiple decrementer exception requests are received before the first can be reported, only one exception is reported. The occurrence of a decrementer exception cancels the request. Register settings for this exception are described in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.

When a decrementer exception is taken, instruction execution for the handler begins at offset 0x00900 from the physical base address indicated by MSR[IP] .

4.5.10 System Call Exception (0x00C00)

The MPC603e implements the system call exception as it is defined by the PowerPC architecture. A system call exception request is made when a system call (**sc**) instruction is completed. If no higher priority exception exists, the system call exception is taken, with **SRR0** being set to the EA of the instruction following the **sc** instruction. Register settings for this exception are described in Chapter 6, “Exceptions,” in the *Programming Environments Manual*.

When a system call exception is taken, instruction execution for the handler begins at offset 0x00C00 from the physical base address indicated by MSR[IP] .

4.5.11 Trace Exception (0x00D00)

The trace exception is taken under one of the following conditions:

- When MSR[SE] is set, a single-step instruction trace exception is taken when no higher priority exception exists and any instruction (other than **rfi** or **isync**) is successfully completed. Note that other processors will take the trace exception on **isync** instructions (when MSR[SE] is set); the MPC603e does not take the trace exception on **isync** instructions. Single-step instruction trace mode is described in Section 4.5.11.1, “Single-Step Instruction Trace Mode.”
- When MSR[BE] is set, the branch trace exception is taken after each branch instruction is completed.
- The MPC603e deviates from the architecture by not taking trace exceptions on **isync** instructions. Single-step instruction trace mode is described in Section 4.5.11.2, “Branch Trace Mode.”

Successful completion implies that the instruction caused no other exceptions. A trace exception is never taken for an **sc** or trap instruction that takes a trap exception.

MSR[SE] and MSR[BE] are cleared when the trace exception is taken. In the normal use of this function, MSR[SE] and MSR[BE] are restored when the exception handler returns to the interrupted program using an **rfi** instruction.

Register settings for the trace mode are described in Table 4-16.

Table 4-16. Trace Exception—Register Settings

Register	Setting Description					
SRR0	Set to the address of the instruction following the one for which the trace exception was generated.					
SRR1	0–15 Cleared 16–31 Loaded from MSR[16–31]					
MSR	POW 0	EE 0	FE0 0	IR 0	DR 0	
	TGPR 0	PR 0	SE 0	RI 0	LE Set to value of ILE	
	ILE —	FP 0	BE 0			
	IP —	ME —	FE1 0			

Note that a trace or instruction address breakpoint exception condition generates a soft stop instead of an exception if soft stop has been enabled by the JTAG/COP logic. If trace and breakpoint conditions occur simultaneously, the breakpoint conditions receive higher priority.

When a trace exception is taken, instruction execution for the handler begins as offset 0x00D00 from the base address indicated by MSR[IP].

4.5.11.1 Single-Step Instruction Trace Mode

The single-step instruction trace mode is enabled by setting MSR[SE]. Encountering the single-step breakpoint causes one of the following actions:

- Trap to address vector 0x00D00
- Soft stop (wait for quiescence)

The single-step trace action traps after an instruction execution and completion.

4.5.11.2 Branch Trace Mode

The branch trace mode is enabled by setting MSR[BE]. Encountering the branch trace breakpoint causes one of the following actions:

- Trap to interrupt vector 0x00D00
- Soft stop
- Hard stop

The branch trace action is to trap after the completion of any branch instruction whenever MSR[BE] is set.

4.5.12 Instruction TLB Miss Exception (0x01000)

When the effective address for an instruction load, store, or cache operation cannot be translated by the ITLB, an instruction TLB miss exception is generated. Register settings for the instruction and data TLB miss exceptions are described in Table 4-17.

If the instruction TLB miss exception handler fails to find the desired PTE, then a page fault must be synthesized. The handler must restore the machine state and clear MSR[TGPR] before invoking the ISI exception (0x00400).

Software table search operations are discussed in Chapter 5, “Memory Management.”

When an instruction TLB miss exception is taken, instruction execution for the handler begins at offset 0x01000 from the physical base address indicated by MSR[IP].

4.5.13 Data TLB Miss on Load Exception (0x01100)

When the effective address for a data load or cache operation cannot be translated by the DTLB, a data TLB miss on load exception is generated. Register settings for the instruction and data TLB miss exceptions are described in Table 4-17.

If a data TLB miss exception handler fails to find the desired PTE, then a page fault must be synthesized. The handler must restore the machine state and clear MSR[TGPR] before invoking the DS1 exception (0x00300).

Table 4-17. Instruction and Data TLB Miss Exceptions—Register Settings

Register	Setting Description					
SRR0	Set to the address of the next instruction to be executed in the program for which the TLB miss exception was generated.					
SRR1	0–3 Loaded from condition register CR0 field 4–11 Cleared 12 KEY. Key for TLB miss (SR[Ks] or SR[Kp], depending on whether the access is a user or supervisor access). 13 D/I. Data or instruction access 0 = Data TLB miss 1 = Instruction TLB miss 14 WAY. Next TLB set to be replaced (set per LRU) 0 = Replace TLB associativity set 0 1 = Replace TLB associativity set 1 15 S/L. Store or load data access 0 = Data TLB miss on load 1 = Data TLB miss on store (or C = 0) 16–31 Loaded from MSR[16–31]					
MSR	POW 0 EE 0 FE0 0 IR 0 TGPR 1 PR 0 SE 0 DR 0 ILE — FP 0 BE 0 RI 0 IP — ME — FE1 0 LE Set to value of ILE					

Software table search operations are discussed in Chapter 5, “Memory Management.”

When a data TLB miss on load exception is taken, instruction execution for the handler begins at offset 0x01100 from the physical base address indicated by MSR[IP].

4.5.14 Data TLB Miss on Store Exception (0x01200)

When the effective address for a data store or cache operation cannot be translated by the DTLB, a data TLB miss on store exception is generated. The data TLB miss on store exception is also taken when the changed bit (C = 0) for a DTLB entry needs to be updated for a store operation. Register settings for the instruction and data TLB miss exceptions are described in Table 4-17.

If a data TLB miss exception handler fails to find the desired PTE, then a page fault must be synthesized. The handler must restore the machine state and clear MSR[TGPR] before invoking the DS1 exception (0x00300).

Software table search operations are discussed in Chapter 5, “Memory Management.”

When a data TLB miss on store exception is taken, instruction execution for the handler begins at offset 0x01200 from the physical base address indicated by MSR[IP].

4.5.15 Instruction Address Breakpoint Exception (0x01300)

The instruction address breakpoint is controlled by the IABR special purpose register. IABR[0–29] holds an effective address to which each instruction’s address is compared. The exception is enabled by setting IABR[30]. Note that the MPC603e ignores the translation enable bit (IABR[31]). The exception is taken when an instruction breakpoint address matches on the next instruction to complete. The instruction tagged with the match is not completed before the instruction address breakpoint exception is taken.

The breakpoint action can be one of the following:

- Trap to interrupt vector 0x01300 (default)
- Soft stop

The bit settings for when an instruction address breakpoint exception is taken are shown in Table 4-18.

Table 4-18. Instruction Address Breakpoint Exception—Register Settings

Register	Setting Description							
SRR0	Set to the address of the next instruction to be executed in the program for which the TLB miss exception was generated.							
SRR1	0–15 Cleared 16–31 Loaded from MSR[16–31]							
MSR	POW 0	EE	0	FE0	0	IR	0	
	TGPR 0	PR	0	SE	0	DR	0	
	ILE —	FP	0	BE	0	RI	0	
	IP —	ME	—	FE1	0	LE	Set to value of ILE	

The default breakpoint action is to trap before the execution of the matching instruction.

The soft stop feature can be enabled only through the COP interface. With soft stop enabled, the MPC603e stops in a restartable state, while with hard stop enabled, the MPC603e stops immediately without attempting to reach a restartable state. Upon restarting from a soft stop, the matching instructions are executed and completed unless it generates an exception. For soft stops, the next ten instructions that could have passed the IABR check can be monitored only by single-stepping the processor. When soft stops are used, the address compare must be separated by at least ten instructions.

If soft stop is enabled, only one soft stop is generated before completion of an instruction with an IABR match, regardless of whether a soft stop is generated before that instruction for any other reason, such as trace mode on for the preceding instruction or a COP soft stop request.

Table 4-19 shows the priority of actions taken when more than one mode is enabled for the same instruction.

Note that a trace or instruction address breakpoint exception condition generates a soft stop instead of an exception if soft stop has been enabled by the JTAG/COP logic. If trace and breakpoint conditions occur simultaneously, the breakpoint conditions receive higher priority.

The MPC603e requires that an **mtspr** instruction that updates the IABR be followed by a context-synchronizing instruction. If the **mtspr** instruction enables the instruction address breakpoint exception, the context-synchronizing instruction cannot generate a breakpoint response. The MPC603e also cannot block a breakpoint response on the context-synchronizing instruction if the breakpoint was disabled by the **mtspr** instruction. See “Synchronization Requirements for Special Registers and TLBs” in Chapter 2, “Register Set,” in the *Programming Environments Manual* for more information on this requirement.

Table 4-19. Breakpoint Action for Multiple Modes Enabled for the Same Address

IABR[IE]	MSR[BE]	MSR[SE]	First Action	Next Action	Comments
1	1	0	Instruction address	Trace (branch)	Enabling both modes is useful only if both trace and address breakpoint interrupts are needed.
1	0	1	Instruction address breakpoint	Trace (single-step)	Enabling both modes is useful only if different breakpoint actions are required.
0	1	1	Trace (branch)	None	The action for branch trace and single-step trace is the same. Enabling both trace modes is redundant except for hard stop on branches.
1	1	1	Instruction address breakpoint	Trace	Enabling all modes is redundant. This entry is for clarification only.

Section 2.1.2.6, “Instruction Address Breakpoint Register (IABR),” provides more information about the instruction breakpoint facility.

4.5.16 System Management Interrupt (0x01400)

The system management interrupt behaves like an external interrupt except for the signal asserted and the vector taken. A system management interrupt is signaled to the MPC603e by the assertion of the $\overline{\text{SMI}}$ signal. The interrupt may not be recognized if a higher priority exception occurs simultaneously or if MSR[EE] is cleared when $\overline{\text{SMI}}$ is asserted. Note that $\overline{\text{SMI}}$ takes priority over $\overline{\text{INT}}$ if they are recognized simultaneously.

After the $\overline{\text{SMI}}$ is detected (and provided that MSR[EE] is set), the MPC603e generates a recoverable halt to instruction completion. The MPC603e requires the next instruction in program order to complete or except, block completion of any following instructions, and allow the completed store queue to drain. If any higher priority exceptions are encountered in this process, they are taken first and the system management interrupt is delayed until a recoverable halt is achieved. At this time the MPC603e saves state information and takes the system management interrupt.

The register settings for the external interrupt exception are shown in Table 4-20.

Table 4-20. System Management Interrupt—Register Settings

Register	Setting Description					
SRR0	Set to the effective address of the instruction that the processor would have attempted to complete next if no interrupt conditions were present.					
SRR1	0–15 Cleared 16–31 Loaded from MSR[16–31]					
MSR	POW 0	EE 0	FE0 0	IR 0		
	TGPR 0	PR 0	SE 0	DR 0		
	ILE —	FP 0	BE 0	RI 0		
	IP —	ME —	FE1 0	LE Set to value of ILE		

When a system management interrupt is taken, instruction execution for the handler begins at offset 0x01400 from the physical base address indicated by MSR[IP].

The MPC603e recognizes the interrupt condition ($\overline{\text{SMI}}$ asserted) only if the MSR[EE] bit is set; otherwise, the interrupt condition is ignored. To guarantee that the external interrupt is taken, the $\overline{\text{SMI}}$ signal must be held active until the MPC603e takes the interrupt. If the $\overline{\text{SMI}}$ signal is negated before the interrupt is taken, the MPC603e is not guaranteed to take a system management interrupt. The interrupt handler must send a command to the device that asserted $\overline{\text{SMI}}$, acknowledging the interrupt and instructing the device to negate $\overline{\text{SMI}}$.

Chapter 5

Memory Management

This chapter describes the MPC603e microprocessor's implementation of the memory management unit (MMU) specifications provided by the PowerPC operating environment architecture (OEA) for processors that implement the PowerPC architecture. The MPC603e MMU implementation is very similar to that of the MPC603 microprocessor except that the MPC603e implements an extra key bit in the SRR1 register that simplifies the table search software. In addition, because the MPC603e does not support direct-store bus accesses, it causes a DS1 exception when a direct-store segment is encountered.

The primary function of the MMU processor is the translation of logical (effective) addresses to physical addresses (referred to as real addresses in the architecture specification) for memory accesses, and I/O accesses (I/O accesses are assumed to be memory-mapped). In addition, the MMU provides access protection on a segment, block, or page basis. This chapter describes the specific hardware used to implement the MMU model of the OEA in the MPC603e. Refer to Chapter 7, "Memory Management," in the *Programming Environments Manual* for a complete description of the conceptual model.

Two general types of accesses generated by processors that implement the PowerPC architecture require address translation—instruction accesses and data accesses to memory generated by load and store instructions. Generally, the address translation mechanism is defined in terms of segment descriptors and page tables used by processors to locate the effective-to-physical address mapping for instruction and data accesses. The segment information translates the effective address to an interim virtual address, and the page table information translates the virtual address to a physical address.

The segment descriptors, used to generate the interim virtual addresses, are stored as on-chip segment registers on 32-bit implementations (such as the MPC603e). In addition, two translation lookaside buffers (TLBs) are implemented on the MPC603e to keep recently-used page address translations on-chip. Although the PowerPC OEA describes one MMU (conceptually), the MPC603e hardware maintains separate TLBs and table search resources for instruction and data accesses that can be accessed independently (and simultaneously). Therefore, the MPC603e is described as having two MMUs, one for instruction accesses (IMMU) and one for data accesses (DMMU).

The block address translation (BAT) mechanism is a software-controlled array that stores the available block address translations on-chip. BAT array entries are implemented as

pairs of BAT registers that are accessible as supervisor-level special-purpose registers (SPRs). There are separate instruction and data BAT mechanisms, and in the MPC603e, they reside in the instruction and data MMUs, respectively.

The MMUs, together with the exception processing mechanism, provide the necessary support for the operating system to implement a paged virtual memory environment and for enforcing protection of designated memory areas. Exception processing is described in Chapter 4, “Exceptions.” Section 4.2, “Exception Processing,” describes the MSR which controls some of the critical functionality of the MMUs.

5.1 MMU Features

The MPC603e implements the memory management specification of the PowerPC OEA for 32-bit implementations. Thus, it provides 4 Gbytes of effective address space accessible to supervisor and user programs with a 4-Kbyte page size and 256-Mbyte segment size. In addition, the MMUs of 32-bit processors use an interim virtual address (52 bits) and hashed page tables in the generation of 32-bit physical addresses. These processors also have a BAT mechanism for mapping large blocks of memory. Block sizes range from 128 Kbytes to 256 Mbytes and are software-programmable.

The MPC603e completely implements all features required by the MMU specifications of the PowerPC architecture (OEA) for 32-bit implementations. Table 5-1 summarizes all MPC603e MMU features including the architectural features of PowerPC MMUs (defined by the OEA) for 32-bit processors and the implementation-specific features provided by the MPC603e.

Table 5-1. MMU Features Summary

Feature Category	Architecturally Defined/ MPC603e-Specific	Feature
Address ranges	Architecturally defined	2^{32} bytes of effective address
		2^{52} bytes of virtual address
		2^{32} bytes of physical address
Page size	Architecturally defined	4 Kbytes
Segment size	Architecturally defined	256 Mbytes
Block address translation	Architecturally defined	Range of 128 Kbytes–256 Mbytes sizes
		Implemented with IBAT and DBAT registers in BAT array
Memory protection	Architecturally defined	Segments selectable as no-execute
		Pages selectable as user/supervisor and read-only
		Blocks selectable as user/supervisor and read-only
Page history	Architecturally defined	Referenced and changed bits defined and maintained
Page address translation	Architecturally defined	Translations stored as PTEs in hashed page tables in memory
		Page table size determined by mask in SDR1 register

Table 5-1. MMU Features Summary (continued)

Feature Category	Architecturally Defined/ MPC603e-Specific	Feature
TLBs	Architecturally defined	Instructions for maintaining optional TLBs (tlbie instruction in MPC603e)
	MPC603e-specific	64-entry, two-way set-associative ITLB 64-entry, two-way set-associative DTLB
Segment descriptors	Architecturally defined	Stored as segment registers on-chip
Page table search support	MPC603e-specific	Three MMU exceptions defined: ITLB miss exception, DTLB miss on load exception, and DTLB miss on store (or C = 0) exception; MMU-related bits set in SRR1 for these exceptions.
		IMISS and DMISS registers (missed effective address) HASH1 and HASH2 registers (PTEG addr) ICMP and DCMP registers (for comparing PTEs) RPA register (for loading TLBs)
		tlbli rB instruction for loading ITLB entries tlbld rB instruction for loading DTLB entries
		Shadow registers for GPR0–GPR3 (can use r0–r3 in table search handler without corruption of r0–r3 in context that was previously executing)

5.1.1 Memory Addressing

A program references memory using the effective (logical) address computed by the processor when it executes a load, store, or cache instruction, and when it fetches the next instruction. The effective address is translated to a physical address according to the procedures described in Chapter 7, “Memory Management,” in the *Programming Environments Manual*, augmented with information in this chapter. The memory subsystem uses the physical address for the access.

For a complete discussion of effective address calculation, see Section 2.3.2.3, “Effective Address Calculation.”

5.1.2 MMU Organization

Figure 5-1 shows the conceptual organization of a PowerPC MMU in a 32-bit implementation; note that it does not describe the specific hardware used to implement the memory management function for a particular processor. Processors may optionally implement on-chip TLBs and may optionally support the automatic search of the page tables for PTEs. In addition, other hardware features (invisible to the system software) not depicted in the figure may be implemented.

Figure 5-2 and Figure 5-3 show the conceptual organization of the MPC603e instruction and data MMUs, respectively. The instruction addresses shown in Figure 5-2 are generated by the processor for sequential instruction fetches and addresses that correspond to a

change of program flow. Data addresses shown in Figure 5-3 are generated by load and store instructions and by cache instructions.

As shown in the figures, after an address is generated, the higher-order bits of the effective address, EA0–EA19 (or a smaller set of address bits, EA0–EAn, in the cases of blocks), are translated into physical address bits PA0–PA19. The lower-order address bits, A20–A31, are untranslated and, therefore, identical for both effective and physical addresses. After translating the address, the MMUs pass the resulting 32-bit physical address to the memory subsystem.

In addition to the higher-order address bits, the MMUs automatically keep an indicator of whether each access was generated as an instruction or data access and a supervisor/user indicator that reflects the state of the PR bit of the MSR when the effective address was generated. In addition, for data accesses, there is an indicator of whether the access is for a load or a store operation. This information is then used by the MMUs to appropriately direct the address translation and to enforce the protection hierarchy programmed by the operating system. Section 4.2, “Exception Processing,” describes the MSR, which controls some of the critical functionality of the MMUs.

The figures show how the A20–A26 address bits index into the on-chip instruction and data caches to select a cache set. The remaining physical address bits are then compared with the tag fields (comprised of bits PA0–PA19) of the four selected cache blocks to determine if a cache hit has occurred. In the case of a cache miss, the instruction or data access is then forwarded to the bus interface unit which then initiates an external memory access.

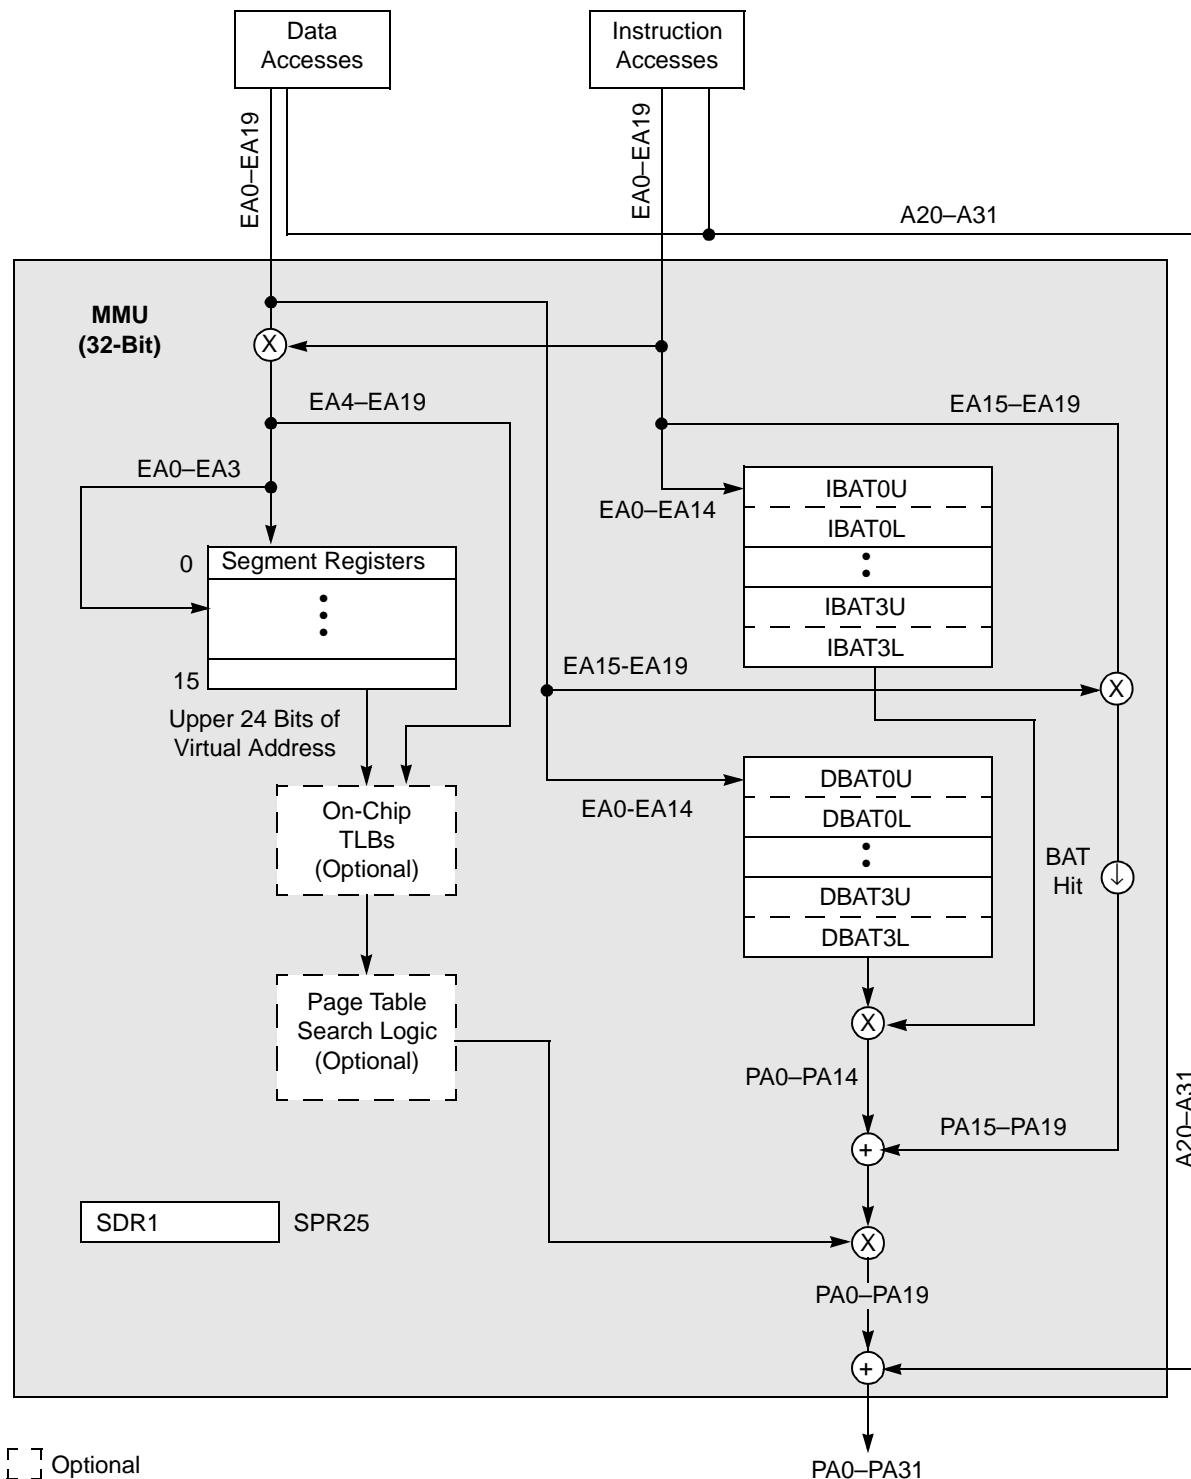


Figure 5-1. MMU Conceptual Block Diagram—32-Bit Implementations

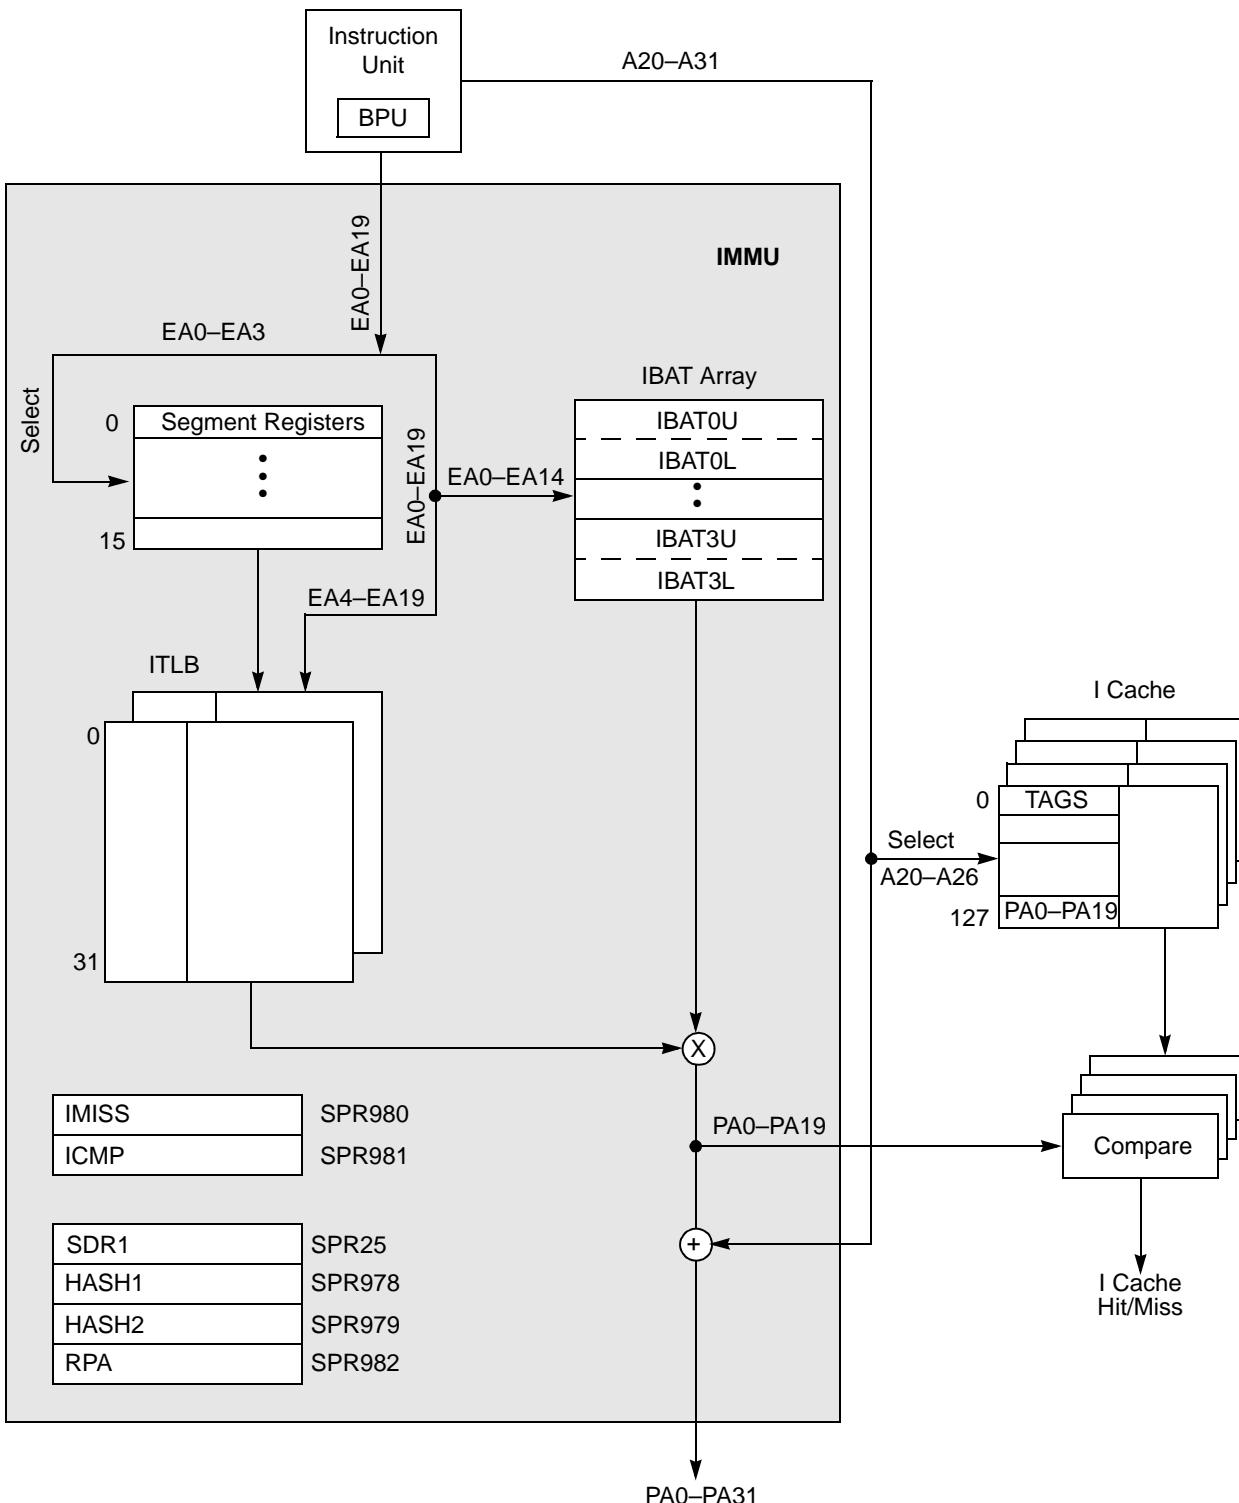


Figure 5-2. IMMU Block Diagram

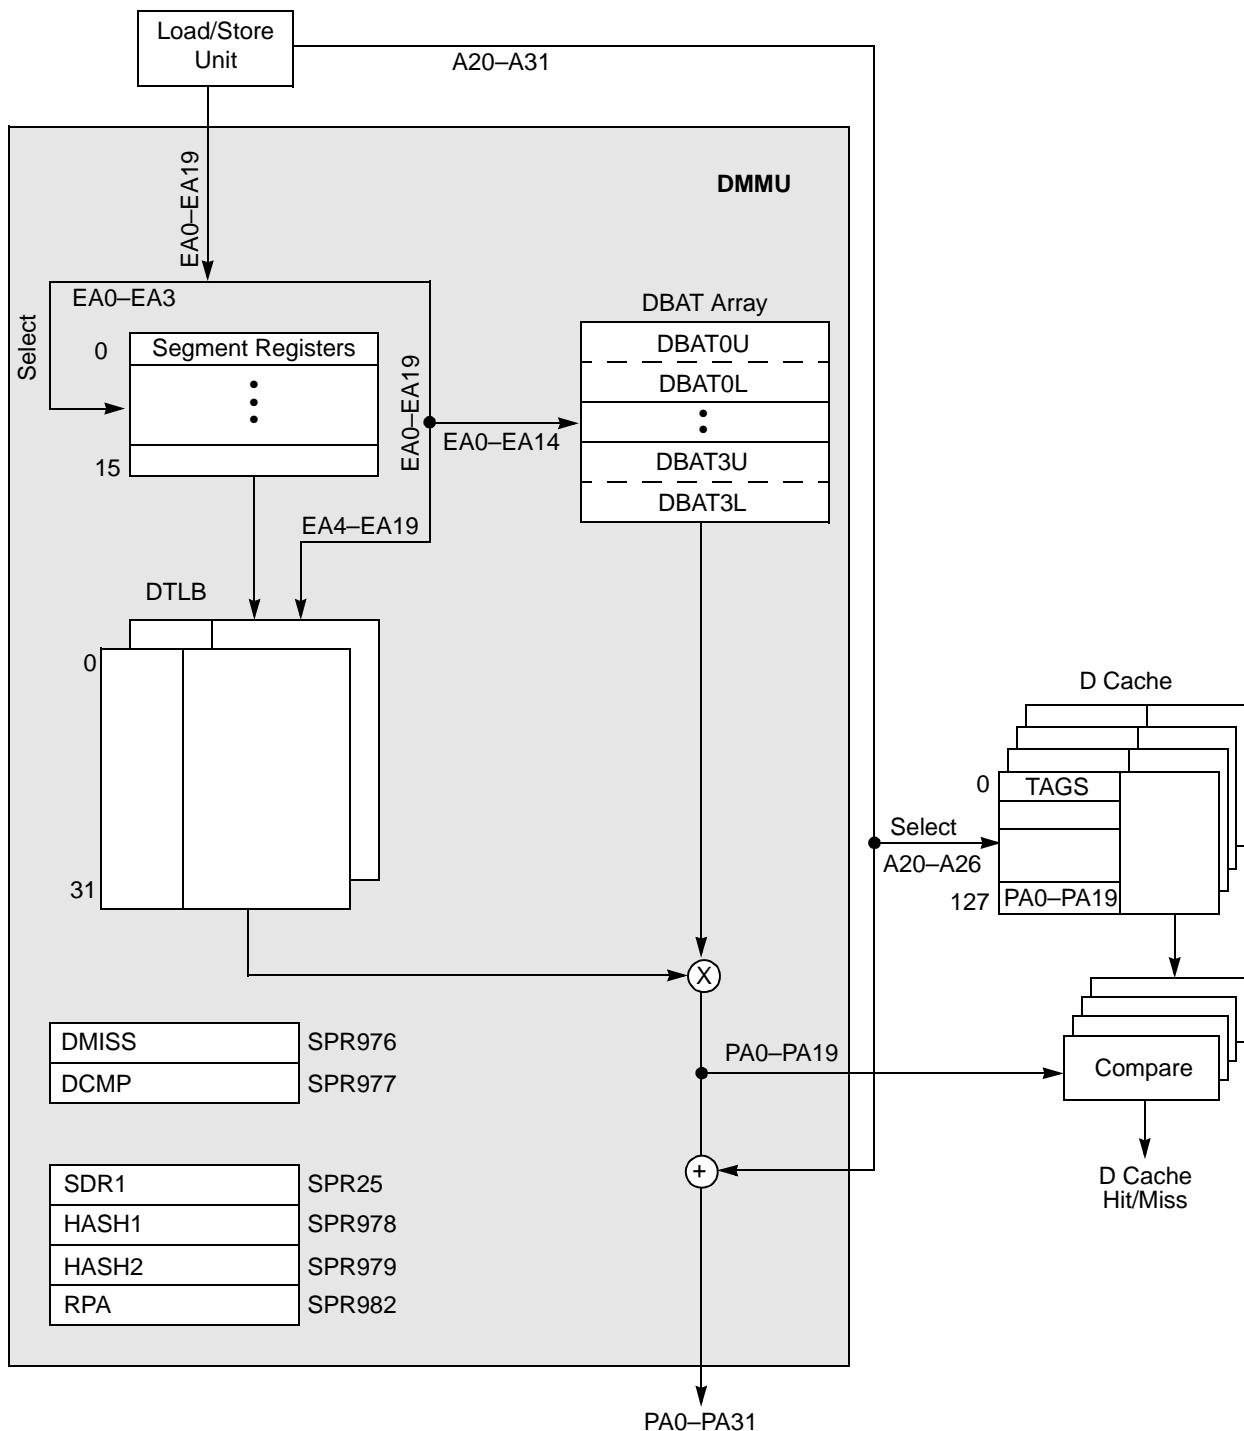


Figure 5-3. DMMU Block Diagram

5.1.3 Address Translation Mechanisms

Processors that implement the PowerPC architecture support the following four types of address translation:

- Page address translation—translates the page frame address for a 4-Kbyte page size.
- Block address translation—translates the block number for blocks that range in size from 128 Kbytes to 256 Mbytes.
- Direct-store interface address translation—used to generate direct-store interface accesses on the external bus; not implemented in the MPC603e.
- Real addressing mode translation—when address translation is disabled, the physical address is identical to the effective address.

Figure 5-4 shows the three implemented address translation mechanisms provided by the MPC603e MMUs. The segment descriptors shown in the figure control the page address translation mechanism. When an access uses page address translation, the appropriate segment descriptor is required. In 32-bit implementations, one of the 16 on-chip segment registers (which contain segment descriptors) is selected by the 4 highest-order effective address bits.

A control bit in the corresponding segment descriptor then determines if the access is to memory (memory-mapped) or to the direct-store interface space (selected when the direct-store translation control bit (T bit) in the corresponding segment descriptor is set). Note that the direct-store interface is present only for compatibility with existing I/O devices that use this interface. When an access is determined to be to the direct-store interface space, the MPC603e takes a DSI exception as described in Section 4.5.3, “DSI Exception (0x00300)” if it is a data access, and takes an ISI exception as described in Section 4.5.4, “ISI Exception (0x00400)” if it is an instruction access.

For memory accesses translated by a segment descriptor, the interim virtual address is generated using the information in the segment descriptor. Page address translation corresponds to the conversion of this virtual address into the 32-bit physical address used by the memory subsystem. In most cases, the physical address for the page resides in an on-chip TLB and is available for quick access. However, if the page address translation misses in an on-chip TLB, the MMU causes a search of the page tables in memory (using the virtual address information and a hashing function) to locate the required physical address. When this occurs, the MPC603e vectors to exception handlers that search the page tables with software.

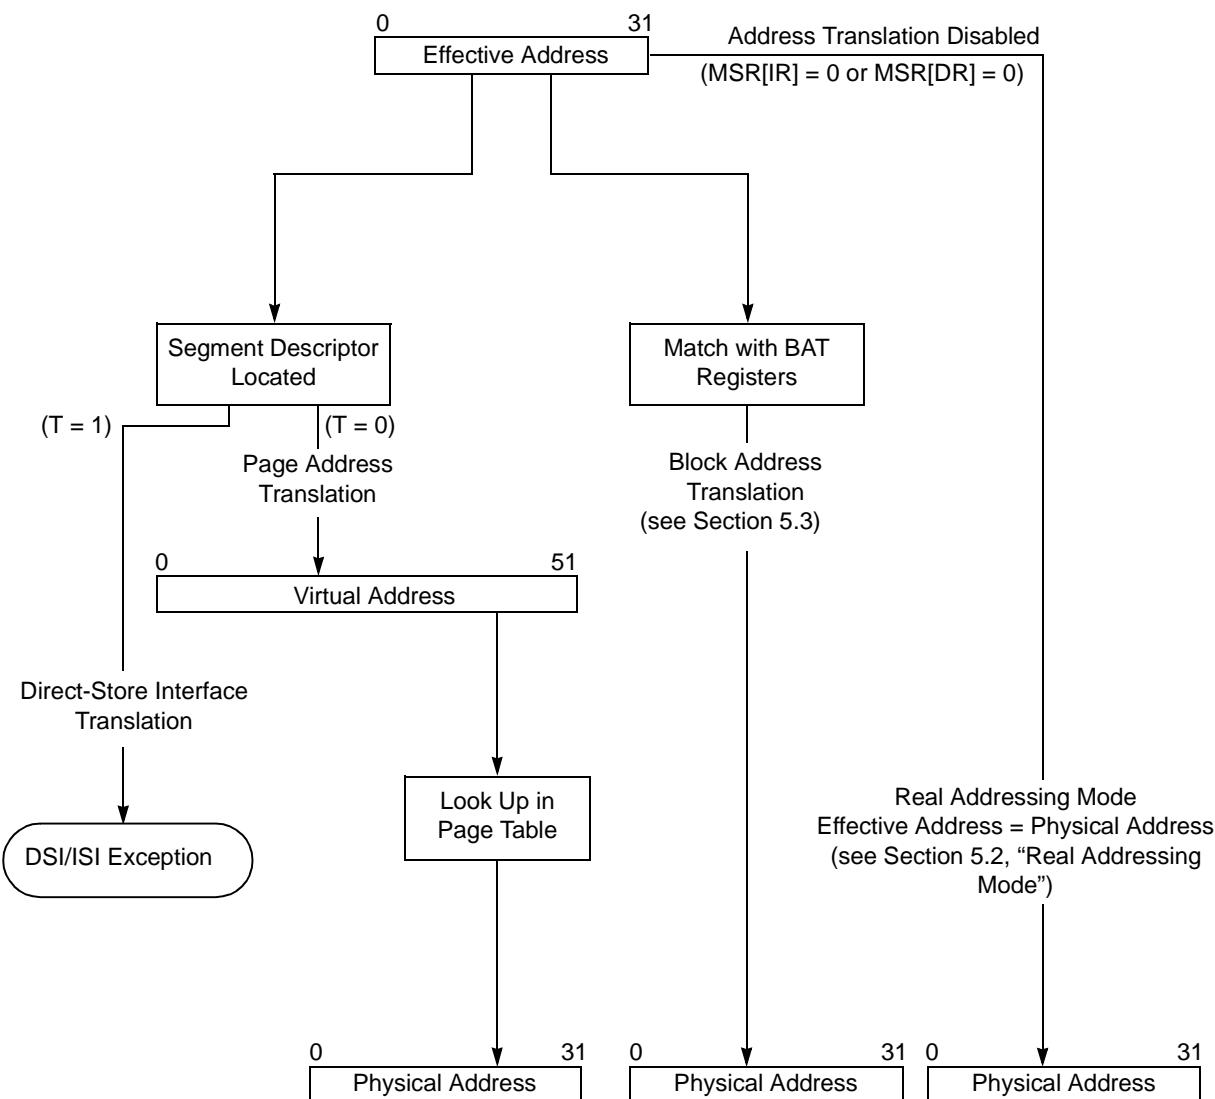


Figure 5-4. Address Translation Types

Block address translation occurs in parallel with page address translation and is similar to page address translation; however, fewer higher-order effective address bits are translated into physical address bits (more lower-order address bits (at least 17) are untranslated to form the offset into a block). Also, instead of segment descriptors and a TLB, block address translations use the on-chip BAT registers as a BAT array. If an effective address matches the corresponding field of a BAT register, the information in the BAT register is used to generate the physical address; in this case, the results of the page translation (occurring in parallel) are ignored (even if the segment corresponds to the direct-store interface space).

Real addressing mode translation occurs when address translation is disabled; in this case the physical address generated is identical to the effective address. Instruction and data address translation is enabled with the MSR[IR] and MSR[DR] bits, respectively. Thus, when the processor generates an access, and the corresponding address translation enable

bit in MSR (MSR[IR] for instruction accesses and MSR[DR] for data accesses) is cleared, the resulting physical address is identical to the effective address and all other translation mechanisms are ignored.

5.1.4 Memory Protection Facilities

In addition to the translation of effective addresses to physical addresses, the MMUs provide access protection of supervisor areas from user access and can designate areas of memory as read-only, as well as, no-execute or guarded. Table 5-2 shows the eight protection options supported by the MMUs for pages.

Table 5-2. Access Protection Options for Pages

Option	User Read		User Write	Supervisor Read		Supervisor Write
	I-Fetch	Data		I-Fetch	Data	
Supervisor-only	—	—	—	✓	✓	✓
Supervisor-only-no-execute	—	—	—	—	✓	✓
Supervisor-write-only	✓	✓	—	✓	✓	✓
Supervisor-write-only-no-execute	—	✓	—	—	✓	✓
Both user/supervisor	✓	✓	✓	✓	✓	✓
Both user/supervisor-no-execute	—	✓	✓	—	✓	✓
Both read-only	✓	✓	—	✓	✓	—
Both read-only-no-execute	—	✓	—	—	✓	—

Note:

- ✓ access permitted.
- protection violation.

The operating system programs whether instructions can be fetched from an area of memory by appropriately using the no-execute option provided in the segment descriptor. Each of the remaining options is enforced, based on a combination of information in the segment descriptor and the page table entry. Thus, the supervisor-only option allows only read and write operations generated while the processor is operating in supervisor mode (corresponding to MSR[PR] = 0) to access the page. User accesses that map into a supervisor-only page cause an exception to be taken.

Finally, there is a facility in the VEA and OEA that allows pages or blocks to be designated as guarded, preventing out-of order accesses that may cause undesired side effects. For example, areas of the memory map that are used to control I/O devices can be marked as guarded so that accesses (for example, instruction prefetches) do not occur unless they are explicitly required by the program.

For more information on memory protection, see “Memory Protection Facilities,” in Chapter 7, “Memory Management,” in the *Programming Environments Manual*.

5.1.5 Page History Information

The MMUs of processors that support the PowerPC architecture also define referenced (R) and changed (C) bits in the page address translation mechanism that can be used as history information relevant to the page. This information can then be used by the operating system to determine the areas of memory to write back to disk when new pages must be allocated in main memory. While these bits are initially programmed by the operating system into the page table, the architecture specifies that the R and C bits may be maintained either by the processor hardware (automatically) or by some software-assist mechanism that updates these bits when required. The software table search routines used by the MPC603e set the R bit when a PTE is accessed; the MPC603e causes an exception (to vector to the software table search routines) when the C bit in the corresponding TLB entry requires updating.

5.1.6 General Flow of MMU Address Translation

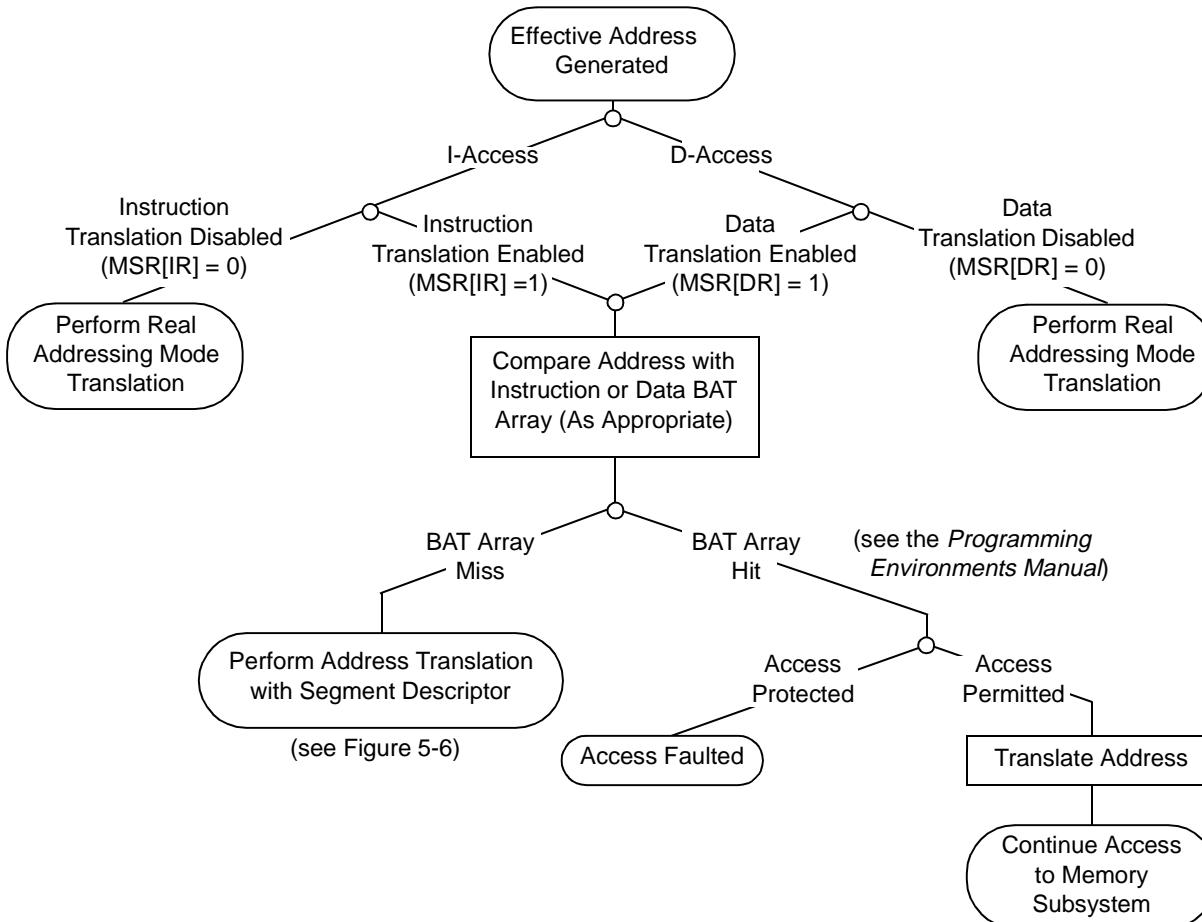
The following sections describe the general flow used by processors that implement the PowerPC architecture to translate effective addresses to virtual and then physical addresses.

5.1.6.1 Real Addressing Mode and Block Address Translation Selection

When an instruction or data access is generated and the corresponding instruction or data translation is disabled (MSR[IR] = 0 or MSR[DR] = 0), real addressing mode translation is used (physical address equals effective address) and the access continues to the memory subsystem as described in Section 5.2, “Real Addressing Mode.”

Figure 5-5 shows the flow used by the MMUs in determining whether to select real addressing mode, block address translation or to use the segment descriptor to select page address translation.

Note that if the BAT array search results in a hit, the access is qualified with the appropriate protection bits. If the access violates the protection mechanism, an exception (ISI or DS_I exception) is generated.



**Figure 5-5. General Flow of Address Translation
(Real Addressing Mode and Block)**

5.1.6.2 Page Address Translation Selection

If address translation is enabled (real addressing mode not selected) and the effective address information does not match with a BAT array entry, then the segment descriptor must be located. Once the segment descriptor is located, the T bit in the segment descriptor selects whether the translation is to a page or to a direct-store interface segment as shown in Figure 5-6. Note that the MPC603e does not implement the direct-store interface, and accesses to these segments cause a DSI exception. In addition, Figure 5-6 also shows the way the no-execute protection is enforced; if the N bit in the segment descriptor is set and the access is an instruction fetch, the access is faulted as described in Chapter 7, “Memory Management,” in the *Programming Environments Manual*. Note that the figure shows the flow for these cases as described by the PowerPC OEA and, therefore, the TLB references are shown as optional. Since the MPC603e implements TLBs, these branches are valid, and described in more detail throughout this chapter.

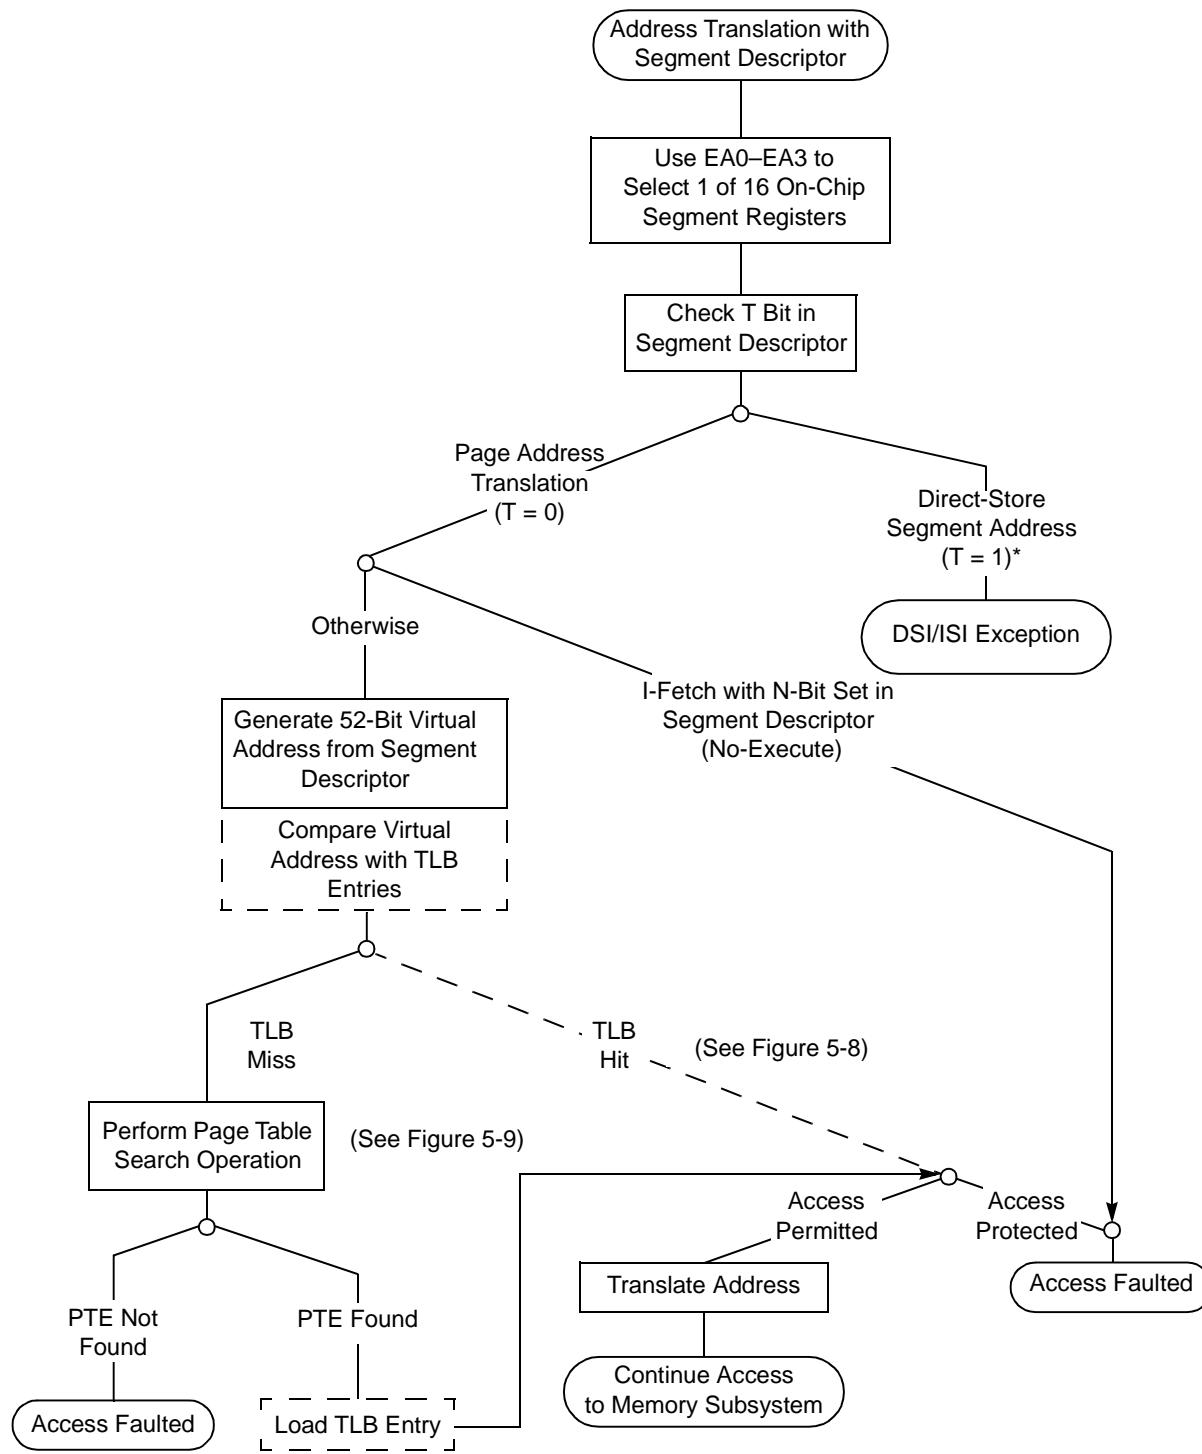


Figure 5-6. General Flow of Page and Direct-Store Interface Address Translation

If the T bit in the corresponding segment descriptor is zero, page address translation is selected. The information in the segment descriptor is then used to generate the 52-bit

virtual address. The virtual address is then used to identify the page address translation information (stored as page table entries (PTEs) in a page table in memory). For increased performance, the MPC603e has two TLBs to store recently-used PTEs on-chip.

If an access hits in the appropriate TLB, the page translation occurs and the physical address bits are forwarded to the memory subsystem. If the required PTE is not resident, the MMU requires a search of the page table. In this case, the MPC603e traps to one of three exception handlers for the system software to perform the page table search. If the PTE is successfully matched, a new TLB entry is created and the page translation is once again attempted. This time, the TLB is guaranteed to hit. Once the PTE is located, the access is qualified with the appropriate protection bits. If the access is a protection violation (not allowed), an exception (instruction access or data access) is generated.

If the PTE is not found by the table search operation, a page fault condition exists, and the TLB miss exception handlers synthesize either an ISI or DS1 exception to handle the page fault.

5.1.7 MMU Exceptions Summary

In order to complete any memory access, the effective address must be translated to a physical address. In the MPC603e, an MMU exception condition occurs if this translation fails for one of the following reasons:

- Page fault—There is no valid entry in the page table for the page specified by the effective address (and segment descriptor) and there is no valid BAT translation.
- An address translation is found but the access is not allowed by the memory protection mechanism.

Additionally, because the MPC603e relies on software to perform table search operations, the processor also takes an exception when:

- There is a miss in the corresponding (instruction or data) TLB.
- The page table requires an update to the changed (C) bit.

The state saved by the processor for each of these exceptions contains information that identifies the address of the failing instruction. Refer to Chapter 4, “Exceptions,” for a more detailed description of exception processing.

Because a page fault condition (PTE not found in the page tables in memory) is detected by the software that performs the table search operation (and not the MPC603e hardware), it does not cause MPC603e exception in the strictest sense in that exception processing as described in Chapter 4, “Exceptions,” does not occur. However, in order to maintain architectural compatibility with software written for other devices that implement the PowerPC architecture, the software that detects this condition should synthesize an exception by setting the appropriate bits in the DSISR or SRR1 and branching to the ISI or DS1 exception handler. Refer to Section 5.5.2, “Implementation-Specific Table Search

Operation,” for more information and examples of this exception software. The remainder of this chapter assumes that the table search software emulates this exception and refers to this condition as an exception.

The translation exception conditions defined by the OEA for 32-bit implementations cause either the ISI or the DS1 exception to be taken as shown in Table 5-3.

Table 5-3. Translation Exception Conditions

Condition	Description	Exception
Page fault (no PTE found)	No matching PTE found in page tables (and no matching BAT array entry)	I access: ISI exception ¹ SRR1[1] = 1
		D access: DS1 exception ¹ DSISR[1] = 1
Block protection violation	Conditions described for block in “Block Memory Protection” in Chapter 7, “Memory Management,” in the <i>Programming Environments Manual</i> .“	I access: ISI exception SRR1[4] = 1
		D access: DS1 exception DSISR[4] = 1
Page protection violation	Conditions described for page in “Page Memory Protection” in Chapter 7, “Memory Management,” in the <i>Programming Environments Manual</i> .	I access: ISI exception ² SRR1[4] = 1
		D access: DS1 exception ² DSISR[4] = 1
No-execute protection violation	Attempt to fetch instruction when SR[N] = 1	ISI exception SRR1[3] = 1
Instruction fetch from direct-store segment	Attempt to fetch instruction when SR[T] = 1	ISI exception SRR1[3] = 1
Data access to direct-store segment (including floating-point accesses) Note: This is an MPC603e-specific condition	Attempt to perform load or store (including floating-point load or store) when SR[T] = 1	DS1 exception DSISR[5] = 1
Instruction fetch from guarded memory with MSR[IR] = 1	Attempt to fetch instruction when MSR[IR] = 1 and either matching xBAT[G] = 1, or no matching BAT entry and PTE[G] = 1.	ISI exception SRR1[3] = 1

¹ The MPC603e hardware does not vector to these exceptions automatically. It is assumed that the software that performs the table search operations vectors to these exceptions and sets the appropriate bits when a page fault condition occurs.

² The table search software can also vector to these exception conditions.

In addition to the translation exceptions, there are other MMU-related conditions (some of them defined as implementation-specific and, therefore, not required by the architecture) that can cause an exception to occur in the MPC603e. These exception conditions map to the processor exception as shown in Table 5-4. For example, the MPC603e also defines three exception conditions to support software table searching. The only exception conditions that occur when MSR[DR] = 0 are the conditions that cause the alignment exception for data accesses. For more detailed information about the conditions that cause the alignment exception (in particular for string/multiple instructions), see Section 4.5.6, “Alignment Exception (0x00600).”

Table 5-4. Other MMU Exception Conditions

Condition	Description	Exception
TLB miss for an instruction fetch	No matching entry found in ITLB	Instruction TLB miss exception SRR1[13] = 1 MSR[14] = 1
TLB miss for a data load access	No matching entry found in DTLB for data load access	Data TLB miss on load exception SRR1[13] = 0 SRR1[15] = 1 MSR[14] = 1
TLB miss for a data store, or store and C = 0	No matching entry found in DTLB for data store access or matching DLTB entry has C = 0 and the access is a store	Data TLB miss on store exception, or store and C = 0 SRR1[13] = 0 SRR1[15] = 0 MSR[14] = 1
TLB miss for an instruction fetch	No matching entry found in ITLB	Instruction TLB miss exception SRR1[13] = 1 MSR[14] = 1
dcbz with W = 1 or I = 1	dcbz instruction to write-through or cache-inhibited segment or block	Alignment exception (not required by architecture for this condition)
dcbz when the data cache is locked	The dcbz instruction takes an alignment exception if the data cache is locked (HID0 bits 18 and 19) when it is executed	Alignment exception
Iwarx , stwcx. , eciwx , or ecowx instruction to direct-store segment	Reservation instruction or external control instruction when SR[T] = 1	DSI exception DSISR[5] = 1
Floating-point load or store to direct-store segment	FP memory access when SR[T] = 1	See data access to direct-store segment in Table 5-3
Load or store that results in a direct-store error	Does not occur in MPC603e	Does not apply
eciwx or ecowx attempted when external control facility disabled	eciwx or ecowx attempted with EAR[E] = 0	DSI exception DSISR[11] = 1
Imw , stmw , lswi , lswx , stswi , or stswx instruction attempted in little-endian mode	Imw , stmw , lswi , lswx , stswi , or stswx instruction attempted while MSR[LE] = 1.	Alignment exception
Operand misalignment	Translation enabled and operand is misaligned as described in Chapter 4, "Exceptions."	Alignment exception (some of these cases are implementation-specific)

Note that some exception conditions depend upon whether the memory area is set up as write-through (W = 1) or cache-inhibited (I = 1). These bits are described fully in "Memory/Cache Access Attributes," in Chapter 5, "Cache Model and Memory Coherency," of the *Programming Environments Manual*. Refer to Chapter 4, "Exceptions," and to Chapter 6, "Exceptions," in the *Programming Environments Manual* for a complete description of the SRR1 and DSISR bit settings for these exceptions.

5.1.8 MMU Instructions and Register Summary

The MMU instructions and registers provide the operating system with the ability to set up the block address translation areas and the page tables in memory.

Note that because the implementation of TLBs is optional, the instructions that refer to these structures are also optional. However, because these structures serve as caches of the page table, the architecture specifies a software protocol for maintaining coherency between these caches and the tables in memory whenever changes are made to the tables in memory. When the tables in memory are changed, the operating system purges these caches of the corresponding entries, allowing the translation caching mechanism to refetch from the tables when the corresponding entries are required.

Note that the MPC603e implements all TLB-related instructions except **tlbia**, which is treated as an illegal instruction. The MPC603e also uses some implementation-specific instructions to load two on-chip TLBs.

Because the MMU specification for these processors is so flexible, it is recommended that the software that uses these instructions and registers be encapsulated into subroutines to minimize the impact of migrating across the family of implementations.

Table 5-5 summarizes MPC603e instructions that specifically control the MMU. For more detailed information about the instructions, refer to Chapter 2, “Programming Model,” in this book and Chapter 8, “Instruction Set,” in the *Programming Environments Manual*.

Table 5-5. Instruction Summary—MMU Control

Instruction	Description
mtsr SR,rS	Move to Segment Register SR[SR#]← rS
mtsrin rS,rB	Move to Segment Register Indirect SR[rB[0–3]]← rS
mfsr rD,SR	Move from Segment Register rD← SR[SR#]
mfsrin rD,rB	Move from Segment Register Indirect rD← SR[rB[0–3]]
tlbie rB¹	TLB Invalidate Entry For effective address specified by rB, TLB[V]←0 The tlbie instruction invalidates both TLB entries indexed by the EA, and operates on both the instruction and data TLBs simultaneously invalidating four TLB entries. The index corresponds to bits 15–19 of the EA. Software must ensure that instruction fetches or memory references to the virtual pages specified by the tlbie instruction have been completed prior to executing the tlbie instruction.
tlbsync¹	TLB Synchronize Synchronizes the execution of all other tlbie instructions in the system. In the MPC603e, when the TLBISYNC signal is negated, instruction execution may continue or resume after the completion of a tlbsync instruction. When the TLBISYNC signal is asserted, instruction execution stops after the completion of a tlbsync instruction. For a complete description of the TLBISYNC signal, refer to Section 8.8.2, “TLBISYNC Input.”

Table 5-5. Instruction Summary—MMU Control (continued)

Instruction	Description
tbl1 (MPC603e-specific)	Load Instruction TLB Entry Loads the contents of the ICMP and RPA registers into the ITLB.
tbld (MPC603e-specific)	Load Data TLB Entry Loads the contents of the DCMP and RPA registers into the DTLB.

¹ These instructions are defined by the PowerPC architecture, but are optional.

Table 5-6 summarizes the registers that the operating system uses to program the MPC603e MMUs. These registers are accessible to supervisor-level software only. These registers are described in Chapter 2, “Register Set,” in the *Programming Environments Manual*. For MPC603e-specific registers, see Chapter 2, “Programming Model,” of this book.

Table 5-6. MMU Registers

Register	Description
Segment registers (SR0–SR15)	The sixteen 32-bit segment registers are present only in 32-bit implementations of the PowerPC architecture. The fields in the segment register are interpreted differently depending on the value of bit 0. The segment registers are accessed by the mtsrr , mtsrin , mfssr , and mfssrin instructions.
BAT registers (IBAT0U–IBAT3U, IBAT0L–IBAT3L, DBAT0U–DBAT3U, and DBAT0L–DBAT3L)	There are 16 BAT registers, organized as 4 pairs of instruction BAT registers (IBAT0U–IBAT3U paired with IBAT0L–IBAT3L) and 4 pairs of data BAT registers (DBAT0U–DBAT3U paired with DBAT0L–DBAT3L). The BAT registers are defined as 32-bit registers in 32-bit implementations. These are special-purpose registers that are accessed by the mtspr and mfsspr instructions.
SDR1	The SDR1 register specifies the variable used in accessing the page tables in memory. SDR1 is defined as a 32-bit register for 32-bit implementations. This is a special-purpose register that is accessed by the mtspr and mfsspr instructions.
Instruction TLB miss address and data TLB miss address registers (IMISS and DMISS)	When a TLB miss exception occurs, the IMISS or DMISS register contains the 32-bit effective address of the instruction or data access, respectively, that caused the miss. Note that the MPC603e always loads a big-endian address into the DMISS register. These registers are MPC603e-specific.
Primary and secondary hash address registers (HASH1 and HASH2)	The HASH1 and HASH2 registers contain the primary and secondary PTEG addresses that correspond to the address causing a TLB miss. These PTEG addresses are automatically derived by the MPC603e by performing the primary and secondary hashing function on the contents of IMISS or DMISS, for an ITLB or DTLB miss exception, respectively. These registers are MPC603e-specific.
Instruction and data PTE compare registers (ICMP and DCMP)	The ICMP and DCMP registers contain the word to be compared with the first word of a PTE in the table search software routine to determine if a PTE contains the address translation for the instruction or data access. The contents of ICMP and DCMP are automatically derived by the MPC603e when a TLB miss exception occurs. These registers are MPC603e-specific.
Required physical address register (RPA)	The system software loads a TLB entry by loading the second word of the matching PTE entry into the RPA register and then executing the tbl1 or tbld instruction (for loading the ITLB or DTLB, respectively). This register is MPC603e-specific.

Note that the MPC603e contains other features that do not specifically control the MPC603e MMU but are implemented to increase performance and flexibility. These are:

- Complete set of shadow segment registers for the instruction MMU. These registers are invisible to the programming model, as described in Section 5.4.3, “TLB Description.”
- Temporary GPR0–GPR3. These registers are available as **r0–r3** when MSR[TGPR] is set. The MPC603e automatically sets MSR[TGPR] whenever one of the three TLB miss exceptions occurs, allowing these exception handlers to have four registers that are used as scratchpad space, without having to save or restore this part of the machine state that existed when the exception occurred. Note that MSR[TGPR] is restored to the value in SRR1 when the **rfi** instruction is executed. Refer to Section 5.5.2, “Implementation-Specific Table Search Operation,” for code examples that take advantage of these registers.

In addition, the MPC603e also automatically saves the values of CR[CR0] of the executing context to SRR1[0–3] whenever one of the three TLB miss exceptions occurs. Thus, the exception handler can set CR[CR0] bits and branch accordingly in the exception handler routine, without having to save the existing CR[CR0] bits. However, the exception handler must restore these bits to CR[CR0] before executing the **rfi** instruction. There are also four other bits saved in SRR1 whenever a TLB miss exception occurs that give information about whether the access was an instruction or data access; and if it was a data access, whether it was for a load or a store instruction. Also, these bits give some information related to the protection attributes for the access, and which set in the TLB will be replaced when the next TLB entry is loaded. Refer to Section 5.5.2.1, “Resources for Table Search Operations,” for more information on these bits and their use.

5.2 Real Addressing Mode

If address translation is disabled (MSR[IR] = 0 or MSR[DR] = 0) for a particular access, the effective address is treated as the physical address and is passed directly to the memory subsystem as described in Chapter 7, “Memory Management,” in the *Programming Environments Manual*.

Note that the default WIMG bits (0b0011) cause data accesses to be considered cacheable ($I = 0$) and, thus, load and store accesses are weakly ordered. This is the case, even if the data cache is disabled in the HID0 register (as it is out of hard reset). If I/O devices require load and store accesses to occur in strict program order (strongly ordered), translation must be enabled so that the corresponding I bit can be set. Also, for instruction accesses, the default memory access mode bits (WIMG) are 0b0001. That is, instruction accesses are considered cacheable ($I = 0$), and the memory is guarded. Again, instruction cache accesses are considered cacheable even if the instruction cache is disabled in the HID0 register (as it is out of hard reset). The W and M bits have no effect on the instruction cache.

For information on the synchronization requirements for changes to MSR[IR] and MSR[DR], refer to “Synchronization Requirements for Special Registers and for Lookaside Buffers” in Chapter 2, “Register Set,” in the *Programming Environments Manual*.

5.3 Block Address Translation

The block address translation (BAT) mechanism in the OEA provides a way to map ranges of effective addresses larger than a single page into contiguous areas of physical memory. Such areas can be used for data that is not subject to normal virtual memory handling (paging), such as a memory-mapped display buffer or an extremely large array of numerical data.

The software model for block address translation in the MPC603e is described in Chapter 7, “Memory Management,” in the *Programming Environments Manual* for 32-bit implementations.

Implementation Note—The MPC603e BAT registers are not initialized by the hardware after the power-up or reset sequence. Consequently, all valid bits in both instruction and data BAT areas must be explicitly cleared before setting any BAT area for the first time and before enabling translation. Also, note that software must avoid overlapping blocks while updating a BAT area or areas. Even if translation is disabled, multiple BAT area hits (with the valid bits set) can corrupt the remaining portion (any bits except the valid bits) of the BAT registers.

Thus, multiple BAT hits (with valid bits set) are considered a programming error whether translation is enabled or disabled, and can lead to unpredictable results if translation is enabled, (or if translation is disabled, when translation is eventually enabled). For the case of unused BATs (if translation is to be enabled), it is sufficient precaution to simply clear the valid bits of the unused BAT entries.

5.4 Memory Segment Model

The MPC603e adheres to the memory segment model as defined in Chapter 7, “Memory Management,” in the *Programming Environments Manual* for 32-bit implementations. Memory in the PowerPC OEA is divided into 256-Mbyte segments. This segmented memory model provides a way to map 4-Kbyte pages of effective addresses to 4-Kbyte pages in physical memory (page address translation), while providing the programming flexibility afforded by a large virtual address space (52 bits).

The segment/page address translation mechanism may be superseded by the BAT mechanism described in Section 5.3, “Block Address Translation.” If not, the translation proceeds in the following two steps:

1. From effective address to the virtual address (which never exists as a specific entity, but can be considered to be the concatenation of the virtual page number and the byte offset within a page).
2. From virtual address to physical address.

This section highlights those areas of the memory segment model defined by the OEA that are specific to the MPC603e.

5.4.1 Page History Recording

Referenced (R) and changed (C) bits reside in each PTE to keep history information about the page. They are maintained by a combination of the MPC603e hardware and the table search software. The operating system uses this information to determine which areas of memory to write back to disk when new pages must be allocated in main memory. Referenced and changed recording is performed only for accesses made with page address translation and not for translations made with the BAT mechanism or for accesses that correspond to direct-store interface ($T = 1$) segments. Furthermore, R and C bits are maintained only for accesses made while address translation is enabled ($MSR[IR] = 1$ or $MSR[DR] = 1$).

In the MPC603e, the referenced and changed bits are updated as follows:

- For TLB hits, the C bit is updated according to Table 5-7.
- For TLB misses, when a table search operation is in progress to locate a PTE, the R and C bits are updated (set, if required) to reflect the status of the page based on this access.

Table 5-7. Table Search Operations to Update History Bits—TLB Hit Case

R and C Bits in TLB Entry	Processor Action	
00	Combination does not occur	
01	Combination does not occur	
10	Read: Write:	No special action Table search operation required to update C. Causes a data TLB miss on store exception.
11	No special action for read or write	

Table 5-7 shows that the status of the C bit in the TLB entry (in the case of a TLB hit) is what causes the processor to update the C bit in the PTE (the R bit is assumed to be set in the page tables if there is a TLB hit). Therefore, when software clears the R and C bits in

the page tables in memory, it must invalidate the TLB entries associated with the pages whose referenced and changed bits were cleared.

The MPC603e causes the R bit to be set for the execution of the **dcbt** or **dcbtst** instruction to that page (by causing a TLB miss exception to load the TLB entry in the case of a TLB miss). However, neither of these instructions causes the C bit to be set.

As defined by the PowerPC architecture, the referenced and changed bits are updated as if address translation were disabled (real addressing mode translation). Additionally, these updates should be performed with single-beat read and byte write transactions on the bus.

5.4.1.1 Referenced Bit

The referenced (R) bit of a page is located in the PTE in the page table. Every time a page is referenced (with a read or write access) and the R bit is zero, the R bit is then set in the page table. The OEA specifies that the referenced bit may be set immediately, or the setting may be delayed until the memory access is determined to be successful. Because the reference to a page is what causes a PTE to be loaded into the TLB, the referenced bit in all MPC603e TLB entries is effectively always set. The processor never automatically clears the referenced bit.

The referenced bit is only a hint to the operating system about the activity of a page. At times, the referenced bit may be set although the access was not logically required by the program, or even if the access was prevented by memory protection. Examples of this in these systems include the following:

- Fetching of instructions not subsequently executed
- Accesses generated by an **lswx** or **stswx** instruction with a zero length
- Accesses generated by a **stwex.** instruction when no store is performed because a reservation does not exist
- Accesses that cause exceptions and are not completed

5.4.1.2 Changed Bit

The changed bit of a page is located both in the PTE in the page table and in the copy of the PTE loaded into the TLB (if a TLB is implemented, as in the MPC603e). Whenever a data store instruction is executed successfully, if the TLB search (for page address translation) results in a hit, the changed bit in the matching TLB entry is checked. If it is already set, the processor does not change the C bit. If the TLB changed bit is zero, it is set and a table search operation is performed to also set the C bit in the corresponding PTE in the page table. The MPC603e causes a data TLB miss on store exception for this case so that the software can perform the table search operation for setting the C bit. Refer to Section 5.5.2, “Implementation-Specific Table Search Operation,” for an example code sequence that handles these conditions.

The changed bit (in both the TLB and PTE in the page tables) is set only when a store operation is allowed by the page memory protection mechanism and all conditional branches occurring earlier in the program have been resolved (such that the store is guaranteed to be in the execution path). Furthermore, the following conditions may cause the C bit to be set:

- The execution of an **stwcx.** instruction is allowed by the memory protection mechanism but a store operation is not performed because no reservation exists.
- The execution of an **stswx** instruction is allowed by the memory protection mechanism but a store operation is not performed because the specified length is zero.
- The store operation is not performed because an exception occurs before the store is performed.

Again, note that although the execution of the **dcbt** and **dcbtst** instructions may cause the R bit to be set, they never cause the C bit to be set.

5.4.1.3 Scenarios for Referenced and Changed Bit Recording

This section provides a summary of the model (defined by the OEA) that is used by processors for maintaining the referenced and changed bits. In some scenarios, the bits are guaranteed to be set by the processor, in some scenarios, the architecture allows that the bits may be set (not absolutely required), and in some scenarios, the bits are guaranteed to not be set.

In implementations that do not maintain the R and C bits in hardware (such as the MPC603e), software assistance is required. For these processors, the information in this section still applies, except that the software performing the updates is constrained to the rules described (that is, must set bits shown as guaranteed to be set and must not set bits shown as guaranteed to not be set).

Table 5-8 defines a prioritized list of the R and C bit settings for all scenarios. The entries in the table are prioritized from top to bottom, such that a matching scenario occurring closer to the top of the table takes precedence over a matching scenario closer to the bottom of the table. For example, if an **stwcx.** instruction causes a protection violation and there is no reservation, the C bit is not altered, as shown for the protection violation case. Note that in the table, load operations include those generated by load instructions, by the **eciwx** instruction, and by the cache management instructions that are treated as a load with respect to address translation. Similarly, store operations include those operations generated by store instructions, by the **ecowx** instruction, and by the cache management instructions that are treated as a store with respect to address translation. In the columns for the MPC603e, the combination of the MPC603e itself and the software used to search the page tables (described in Section 5.5.2, “Implementation-Specific Table Search Operation”) is assumed.

Table 5-8. Model for Guaranteed R and C Bit Settings

Priority	Scenario	R Bit Set		C Bit Set	
		OEA	MPC603e	OEA	MPC603e
1	No-execute protection violation	No	No	No	No
2	Page protection violation	Maybe	Yes	No	No
3	Out-of-order instruction fetch or load operation	Maybe	No	No	No
4	Out-of-order store operation for instructions that will cause no other kind of precise exception (in the absence of system-caused, imprecise, or floating-point assist exceptions)	Maybe ¹	No	No	No
5	All other out-of-order store operations	Maybe ¹	No	Maybe ¹	No
6	Zero-length load (lswx)	Maybe	Yes	No	No
7	Zero-length store (stswx)	Maybe ¹	Yes	Maybe ¹	Yes
8	Store conditional (stwcx.) that does not store	Maybe ¹	Yes	Maybe ¹	Yes
9	In-order instruction fetch	Yes ²	Yes	No	No
10	Load instruction or eciwx	Yes	Yes	No	No
11	Store instruction, ecowx or dcbz instruction	Yes	Yes	Yes	Yes
12	dcbt , dcbtst , dcbst , or dcbf instruction	Maybe	Yes	No	No
13	icbi instruction	Maybe ¹	No	No ¹	No
14	dcbi instruction	Maybe ¹	Yes	Maybe ¹	Yes

¹ If C is set, R is guaranteed to also be set.

² This includes the case when the instruction was fetched out-of-order and R was not set (does not apply for the MPC603e).

For more information, see “Page History Recording” in Chapter 7, “Memory Management,” of the *Programming Environments Manual*.

5.4.2 Page Memory Protection

The MPC603e implements page memory protection as it is defined in Chapter 7, “Memory Management,” in the *Programming Environments Manual*.

5.4.3 TLB Description

This section describes the hardware resources provided in the MPC603e to facilitate the page address translation process. Note that the hardware implementation of the MMU is not specified by the architecture, and while this description applies to the MPC603e, it does not necessarily apply to other processors that implement the PowerPC architecture.

5.4.3.1 TLB Organization

Because the MPC603e has two MMUs (IMMU and DMMU) that operate in parallel, some of the MMU resources are shared, and some are actually duplicated (shadowed) in each MMU to maximize performance. Figure 5-7 shows the relationships between these resources within both the IMMU and DMMU and how the various portions of the effective address are used in the address translation process.

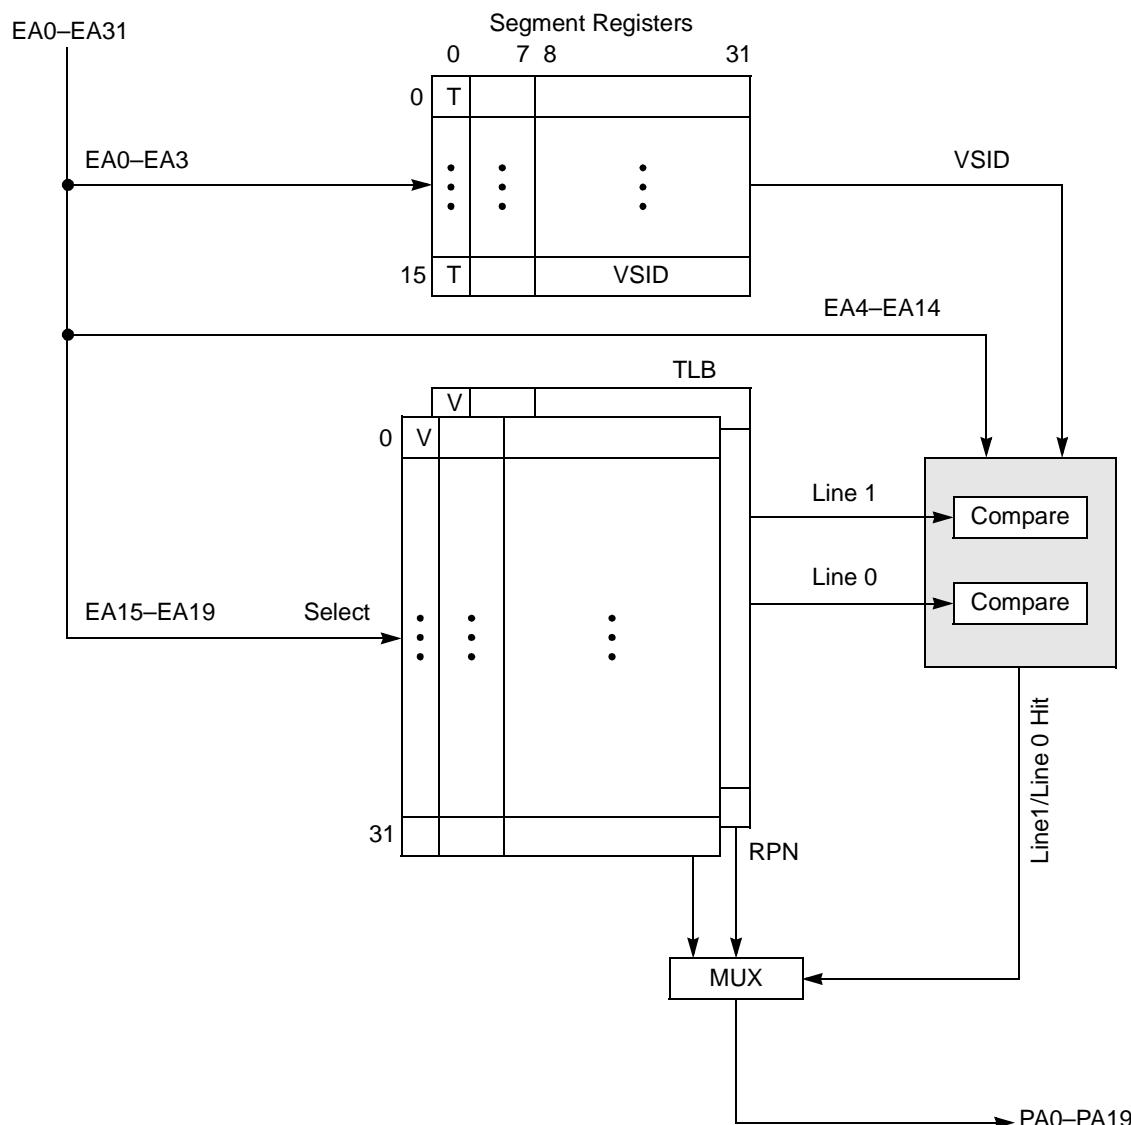


Figure 5-7. Segment Register and TLB Organization

While both MMUs can be accessed simultaneously (both sets of segment registers and TLBs can be accessed in the same clock), when there is an exception condition, only one exception is reported at a time. ITLB miss exceptions are reported when there are no more instructions to be dispatched or retired (the pipeline is empty). Refer to Chapter 6, “Instruction Timing,” for more detailed information about the internal pipelines and the reporting of exceptions.

As TLB entries are on-chip copies of PTEs in the page tables in memory, they are similar in structure. TLB entries consist of two words; the high-order word contains the VSID and API fields of the high-order word of the PTE and the low-order word contains the RPN, the C bit, the WIMG bits and the PP bits (as in the low-order word of the PTE). In order to uniquely identify a TLB entry as the required PTE, the TLB entry also contains five more bits of the page index, EA[10–14] (in addition to the API bits of the PTE).

When an instruction or data access occurs, the effective address is routed to the appropriate MMU. EA[0–3] select 1 of the 16 segment registers and the remaining effective address bits and the virtual address from the segment register is passed to the TLB. EA[15–19] then select two entries in the TLB; the valid bit is checked and EA[10–14], the VSID, and API fields (EA[4–9]) for the access are then compared with the corresponding values in the TLB entries. If one of the entries hits, the PP bits are checked for a protection violation, and the C bit is checked. If these bits do not cause an exception, the RPN value is passed to the memory subsystem and the WIMG bits are then used as attributes for the access.

Although address translation is disabled on a reset condition, the valid bits of the BAT array and TLB entries are not automatically cleared. Thus, TLB entries must be explicitly cleared by the system software (with the **tlbie** instruction) before address translation is enabled. Also, note that the segment registers do not have a valid bit, and so they should also be initialized before translation is enabled.

5.4.3.2 TLB Entry Invalidation

For processors, such as the MPC603e, that implement TLB structures to maintain on-chip copies of the PTEs that are resident in physical memory, the optional **tlbie** instruction provides a way to invalidate the TLB entries. Note that the execution of the **tlbie** instruction in the MPC603e invalidates four entries—both the ITLB entries indexed by EA[15–19] and both the indexed entries of the DTB.

The architecture allows **tlbie** to optionally enable a TLB invalidate signaling mechanism in hardware so that other processors also invalidate their resident copies of the matching PTE. The MPC603e does not signal the TLB invalidation to other processors and does not perform any action when a TLB invalidation is performed by another processor.

The **tlbsync** instruction causes instruction execution to stop if the TLBISYNC signal is also asserted. If TLBISYNC is negated, instruction execution may continue or resume after the completion of a **tlbsync** instruction. Section 8.8.2, “TLBISYNC Input,” describes the TLB synchronization mechanism in further detail.

The **tlbia** instruction is not implemented on the MPC603e and when its opcode is encountered, an illegal instruction program exception is generated. To invalidate all entries of both TLBs, 32 **tlbie** instructions must be executed, incrementing the value in EA[15–19] by 1 each time. See Chapter 8, “Instruction Set,” in the *Programming Environments Manual* for detailed information about the **tlbie** instruction.

5.4.4 Page Address Translation Summary

Figure 5-8 provides the detailed flow for the page address translation mechanism. The figure includes the checking of the N bit in the segment descriptor and then expands on the TLB Hit branch of Figure 5-6. The detailed flow for the TLB Miss branch is described in Section 5.5.1, “Page Table Search Operation—Conceptual Flow.” Note that as in the case of block address translation, if the **dcbz** instruction is attempted to be executed either in write-through mode or as cache-inhibited ($W = 1$ or $I = 1$), the alignment exception is generated. The checking of memory protection violation conditions for page address translation is described in Chapter 7, “Memory Management,” in the *Programming Environments Manual* for 32-bit implementations.

5.5 Page Table Search Operation

As stated earlier, the operating system must synthesize the table search algorithm for setting up the tables. The MPC603e TLB miss exception handlers also use this algorithm (with the assistance of some hardware-generated values) to load TLB entries when TLB misses occur, as described in Section 5.5.2, “Implementation-Specific Table Search Operation.”

5.5.1 Page Table Search Operation—Conceptual Flow

The table search process for a processor that implements the PowerPC architecture varies slightly for 64- and 32-bit implementations. The main differences are the address ranges and PTE formats specified. An outline of the page table search process performed by a 32-bit implementation is as follows:

1. The 32-bit physical address of the primary PTEG is generated as described in Chapter 7, “Memory Management,” in the *Programming Environments Manual* for 32-bit implementations.
2. The first PTE (PTE0) in the primary PTEG is read from memory. PTE reads should occur with an implied WIM memory/cache mode control bit setting of 0b001. Therefore, they are considered cacheable and burst in from memory and placed in the cache.
3. The PTE in the selected PTEG is tested for a match with the virtual page number (VPN) of the access. The VPN is the VSID concatenated with the page index field of the virtual address. For a match to occur, the following must be true:
 - PTE[H] = 0
 - PTE[V] = 1
 - PTE[VSID] = VA[0–23]
 - PTE[API] = VA[24–29]

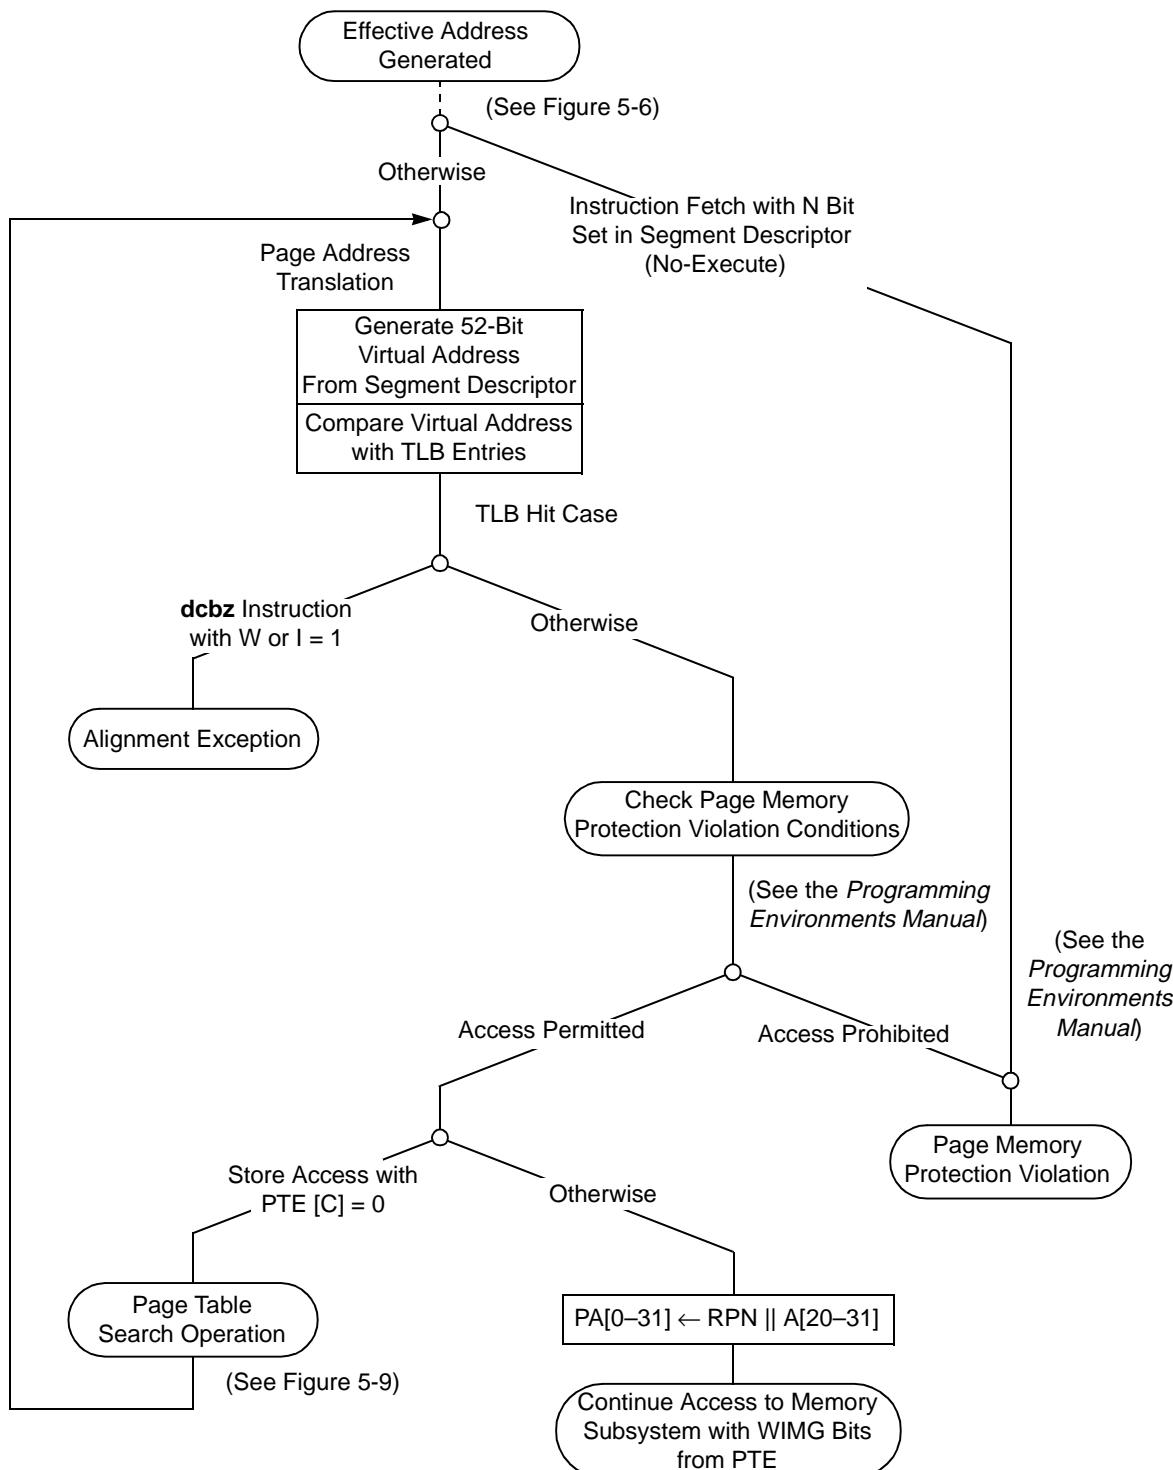


Figure 5-8. Page Address Translation Flow for 32-Bit Implementations—TLB Hit

4. If a match is not found, step 3 is repeated for each of the other seven PTEs in the primary PTEG. If a match is found, the table search process continues as described in step 8. If a match is not found within the eight PTEs of the primary PTEG, the address of the secondary PTEG is generated.
5. The first PTE (PTE0) in the secondary PTEG is read from memory. Again, because PTE reads typically have a WIM bit combination of 0b001, an entire cache line is burst into the on-chip cache.
6. The PTE in the selected secondary PTEG is tested for a match with the virtual page number (VPN) of the access. For a match to occur, the following must be true:
- PTE[H] = 1
 - PTE[V] = 1
 - PTE[VSID] = VA[0–23]
 - PTE[API] = VA[24–29]
7. If a match is not found, step 6 is repeated for each of the other seven PTEs in the secondary PTEG.
8. If a match is found, the PTE is written into the on-chip TLB and the R bit is updated in the PTE in memory (if necessary). If there is no memory protection violation, the C bit is also updated in memory and the table search is complete.
9. If no match is found in the eight PTEs of the secondary PTEG, the search fails and a page fault exception condition occurs (either an ISI exception or a DS1 exception). Note that the software routines that implement this algorithm must synthesize this condition by appropriately setting the SRR1 or DSISR and branching to the ISI or DS1 handler routine.

Reads from memory for table search operations should be performed as global (but not exclusive), cacheable operations, and can be loaded into the on-chip cache.

Figure 5-9 and Figure 5-10 provide conceptual flow diagrams of primary and secondary page table search operations as described in the OEA for 32-bit processors. Recall that the architecture allows implementations to perform the page table search operations automatically (in hardware) or with software assistance, as is the case with the MPC603e. Also, the elements in the figure that apply to TLBs are shown as optional because TLBs are not required by the architecture.

Figure 5-9 shows the case of a **dcbz** instruction that is executed with W = 1 or I = 1, and that the R bit may be updated in memory (if required) before the operation is performed or the alignment exception occurs. The R bit may also be updated by a memory protection violation.

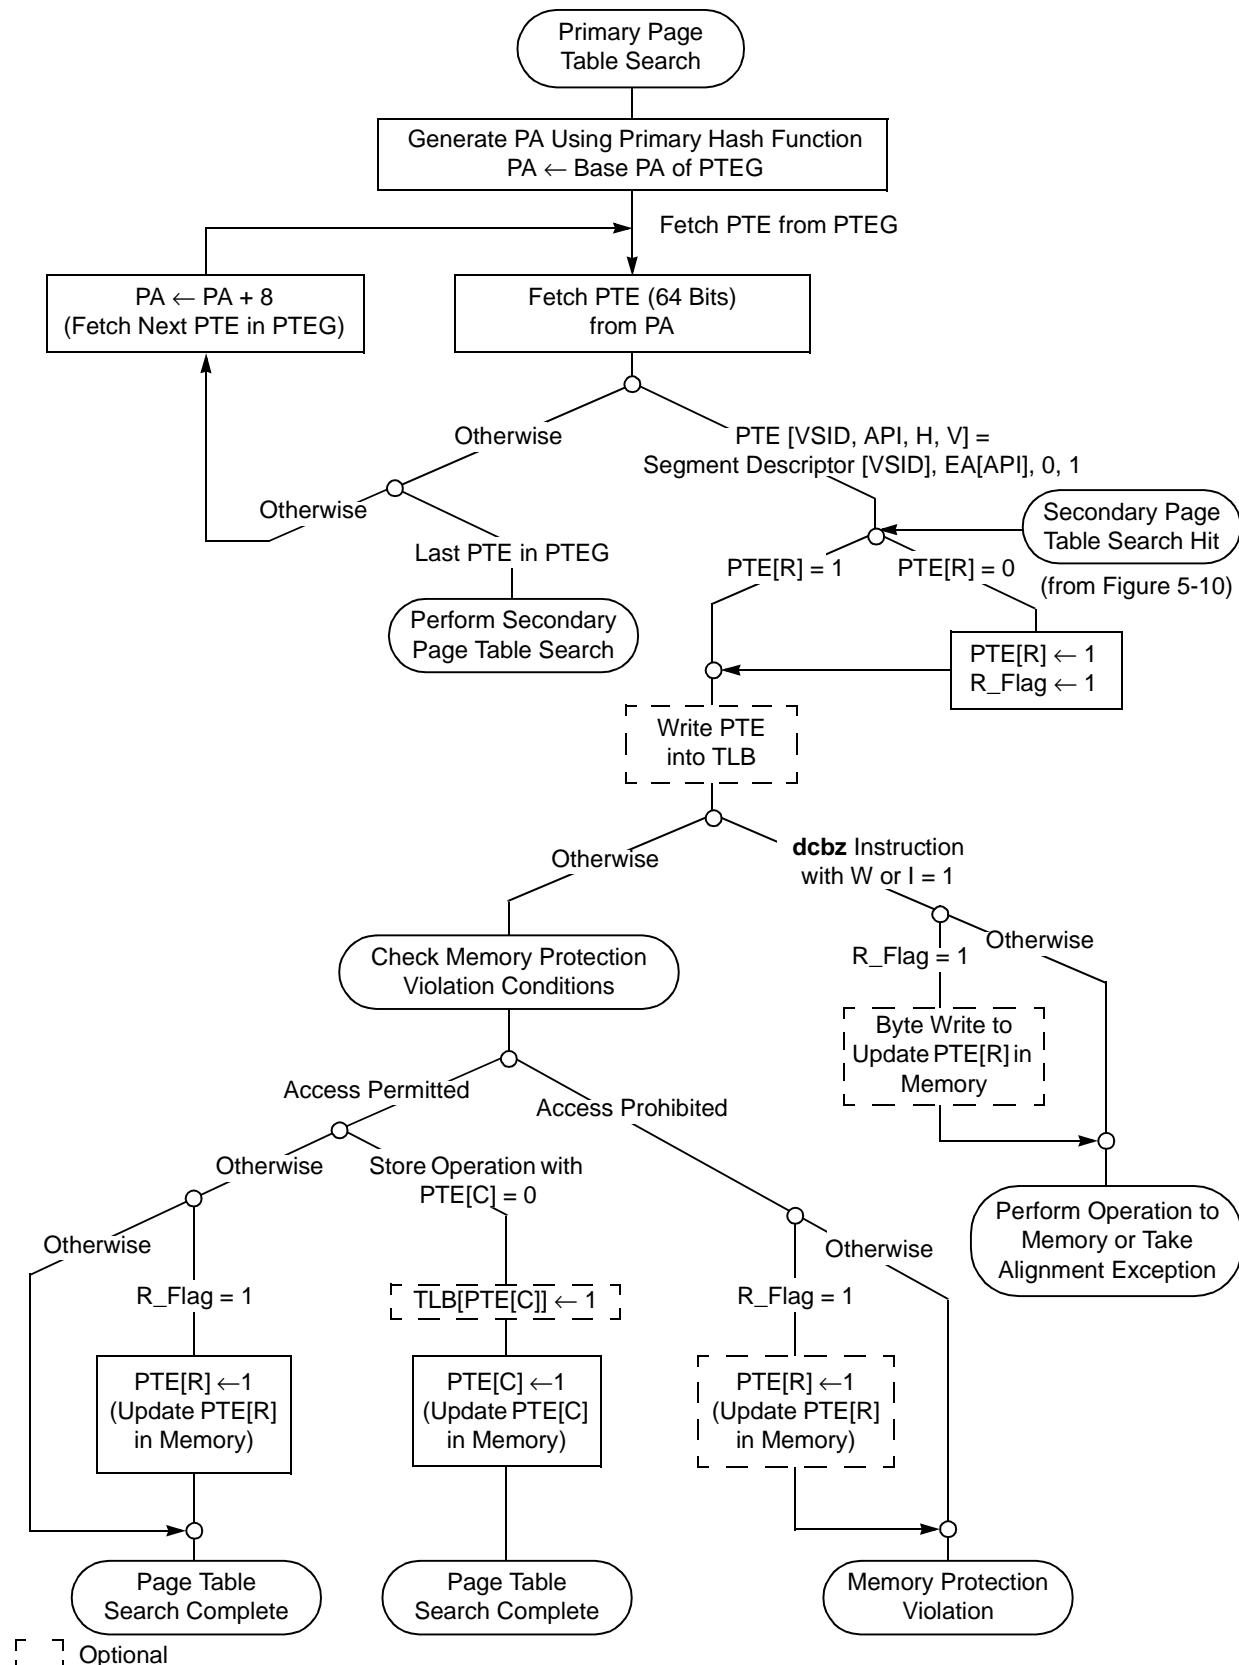


Figure 5-9. Primary Page Table Search—Conceptual Flow

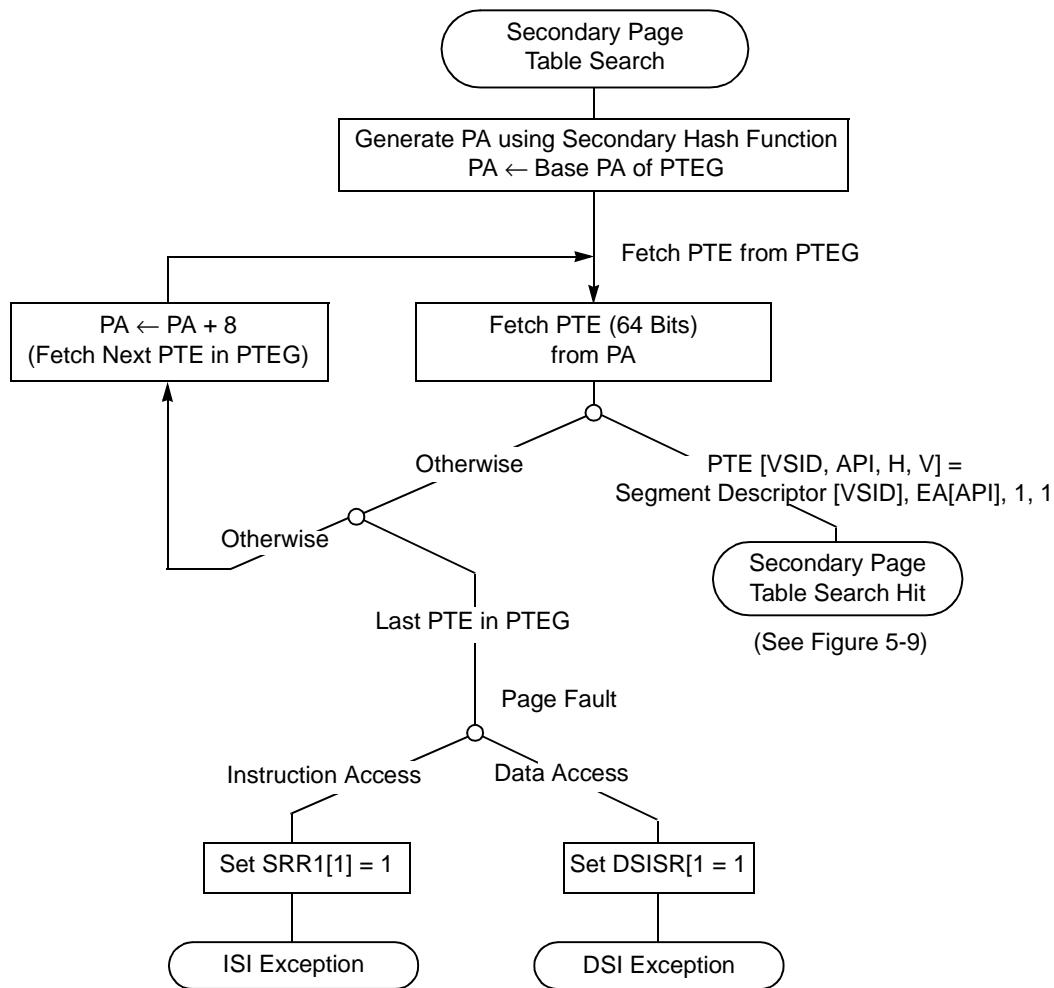


Figure 5-10. Secondary Page Table Search Flow—Conceptual Flow

5.5.2 Implementation-Specific Table Search Operation

The MPC603e has a set of implementation-specific registers, exceptions, and instructions that facilitate very efficient software searching of the page tables in memory. This section describes those resources and provides three example code sequences that can be used in a MPC603e system for an efficient search of the translation tables in software. These three code sequences can be used as handlers for the three exceptions requiring access to the PTEs in the page tables in memory—instruction TLB miss, data TLB miss on load, and data TLB miss on store exceptions.

5.5.2.1 Resources for Table Search Operations

In addition to setting up the translation page tables in memory, the system software must assist the processor in loading PTEs into the on-chip TLBs. When a required TLB entry is not found in the appropriate TLB, the processor vectors to one of the three TLB miss exception handlers so that the software can perform a table search operation and load the

TLB. When this occurs, the processor automatically saves information about the access and the executing context. Table 5-9 provides a summary of the implementation-specific exceptions, registers, and instructions that can be used by the TLB miss exception handler software in MPC603e systems. Refer to Chapter 4, “Exceptions,” for more information about exception processing.

Table 5-9. Implementation-Specific Resources for Table Search Operations

Resource	Name	Description
Exceptions	Instruction TLB miss exception (vector offset 0x1000)	No matching entry found in ITLB.
	Data TLB miss on load exception (vector offset 0x1100)	No matching entry found in DTLB for a load data access.
	Data TLB miss on store exception—also caused when changed bit must be updated (vector offset 0x1200)	No matching entry found in DTLB for a store data access or matching DLTB entry has C = 0 and access is a store.
Registers	IMISS and DMISS	When a TLB miss exception occurs, the IMISS or DMISS register contains the 32-bit effective address of the instruction or data access that caused the miss exception.
	ICMP and DCMP	The ICMP and DCMP registers contain the word to be compared with the first word of a PTE in the table search software routine to determine if a PTE contains the address translation for the instruction or data access. The contents of ICMP and DCMP are automatically derived by the MPC603e when a TLB miss exception occurs.
	HASH1 and HASH2	The HASH1 and HASH2 registers contain the primary and secondary PTEG addresses that correspond to the address causing a TLB miss. These PTEG addresses are automatically derived by the MPC603e by performing the primary and secondary hashing function on the contents of IMISS or DMISS, for an ITLB or DTLB miss exception, respectively.
	RPA	The system software loads a TLB entry by loading the second word of the matching PTE entry into the RPA register and then executing the tlbli or tblld instruction (for loading the ITLB or DTLB, respectively).
Instructions	tlbli rB	Loads the contents of the ICMP and RPA registers into the ITLB entry selected by <ea> and SRR1[WAY].
	tblld rB	Loads the contents of the DCMP and RPA registers into the DTLB entry selected by <ea> and SRR1[WAY].

In addition, the MPC603e contains the following features that do not specifically control the MPC603e MMU, but that are implemented to increase performance and flexibility in the software table search routines whenever one of the three TLB miss exceptions occurs:

- Temporary GPR0–GPR3. These registers are available as **r0–r3** when MSR[TGPR] is set. The MPC603e automatically sets MSR[TGPR] for these cases, allowing these exception handlers to have four registers that are used as scratchpad space, without having to save or restore this part of the machine state that existed when the exception occurred. Note that MSR[TGPR] is cleared when the **rfi** instruction is

executed because the old MSR value (with MSR[TGPR] = 0) saved in SRR1 is restored. Refer to Section 5.5.2.2, “Software Table Search Operation,” for code examples that take advantage of these registers.

- Also, the MPC603e automatically saves the values of CR[CR0] of the executing context to SRR1[0–3]. Thus, the exception handler can set CR[CR0] bits and branch accordingly in the exception handler routine, without having to save the existing CR[CR0] bits. However, the exception handler must restore these bits to CR[CR0] before executing the **rfi** instruction.
- SRR1[D/I] identifies an instruction or data miss, and SRR1[L/S] identifies a load or store miss. SRR1[WAY] identifies the associativity class of the TLB entry selected for replacement by the LRU algorithm. The software can change this value, effectively overriding the replacement algorithm. The SRR1[KEY] bit is used by the table search software to determine if there is a protection violation associated with the access (useful on data write misses for determining if the C bit should be updated in the table). Table 5-10 summarizes the SRR1 bits updated whenever one of the three TLB miss exceptions occurs.

Table 5-10. Implementation-Specific SRR1 Bits

Bits	Name	Function
0–3	CRF0	Condition register field 0 bits
12	KEY	Key for TLB miss (either K _s or K _p from segment register, depending on whether the access is a user or supervisor access).
13	D/I	Set if instruction TLB miss
14	WAY	Next TLB set to be replaced (set per LRU)
15	S/L	Set if data TLB miss was for a load instruction

The key bit saved in SRR1 is derived as follows.

Select KEY from segment register:

If MSR[PR] = 0, KEY = K_s

If MSR[PR] = 1, KEY = K_p

The rest of this section describes the format of the implementation-specific SPRs used by the TLB miss exception handlers. These registers can be accessed by supervisor-level instructions only. Because DMISS, IMISS, DCMP, ICMP, HASH1, HASH2, and RPA are used to access the translation tables for software table search operations, they should only be accessed when address translation is disabled (MSR[IR] = 0 and MSR[DR] = 0). Note that MSR[IR] and MSR[DR] are cleared whenever an exception occurs.

5.5.2.1.1 Data and Instruction TLB Miss Address Registers (DMISS and IMISS)

The DMISS and IMISS registers have the same format as shown in Figure 5-11. They are loaded automatically upon a data or instruction TLB miss. The DMISS and IMISS contain

the effective page address of the access which caused the TLB miss exception. The contents are used by the processor when calculating the values of HASH1 and HASH2, and by the **tblld** and **tblli** instructions when loading a new TLB entry. Note that the MPC603e always loads a big-endian address into the DMISS register. These registers are read-only to the software.

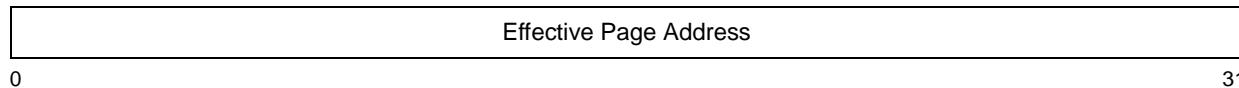


Figure 5-11. DMISS and IMISS Registers

5.5.2.1.2 Data and Instruction TLB Compare Registers (DCMP and ICMP)

The DCMP and ICMP registers are shown in Figure 5-12. These registers contain the first word in the required PTE. The contents are constructed automatically from the contents of the segment registers and the effective address (DMISS or IMISS) when a TLB miss exception occurs. Each PTE read from the tables in memory during the table search process should be compared with this value to determine whether or not the PTE is a match. Upon execution of a **tblld** or **tblli** instruction, the contents of the DCMP or ICMP register is loaded into the first word of the selected TLB entry.

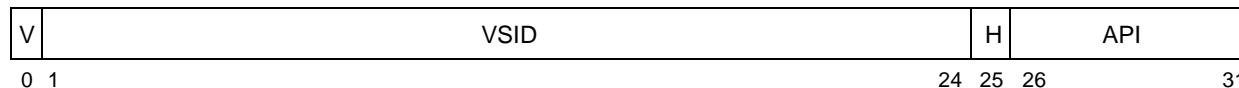


Figure 5-12. DCMP and ICMP Registers

Table 5-11 describes the bit settings for the DCMP and ICMP registers.

Table 5-11. DCMP and ICMP Bit Settings

Bits	Name	Description
0	V	Valid bit. Set by the processor on a TLB miss exception.
1–24	VSID	Virtual segment ID. Copied from VSID field of corresponding segment register.
25	H	Hash function identifier. Cleared by the processor on a TLB miss exception.
26–31	API	Abbreviated page index. Copied from API of effective address.

5.5.2.1.3 Primary and Secondary Hash Address Registers (HASH1 and HASH2)

HASH1 and HASH2 contain the physical addresses of the primary and secondary PTEGs for the access that caused the TLB miss exception. Only bits 7–25 differ between them. For convenience, the processor automatically constructs the full physical address by routing bits 0–6 of SDR1 into HASH1 and HASH2 and clearing the lower six bits. These registers are read-only and are constructed from the contents of the DMISS or IMISS register. The format for HASH1 and HASH2 is shown in Figure 5-13.

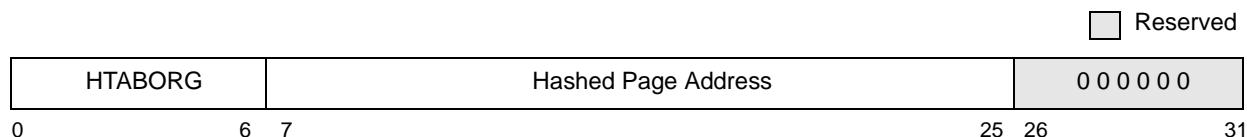
**Figure 5-13. HASH1 and HASH2 Registers**

Table 5-12 describes the bit settings of the HASH1 and HASH2 registers.

Table 5-12. HASH1 and HASH2 Bit Settings

Bits	Name	Description
0–6	HTABORG[0–6]	Copy of the upper 7 bits of the HTABORG field from SDR1
7–25	Hashed page address	Address bits 7–25 of the PTEG to be searched
26–31	—	Reserved

5.5.2.1.4 Required Physical Address Register (RPA)

The RPA is shown in Figure 5-14. During a page table search operation, the software must load the RPA with the second word of the correct PTE. When the **tblld** or **tblli** instruction is executed, data from the IMISS and ICMP (or DMISS and DCMP) and the RPA registers is merged and loaded into the selected TLB entry. The TLB entry is selected by the effective address of the access (loaded by the table search software from the DMISS or IMISS register) and SRR1[WAY].

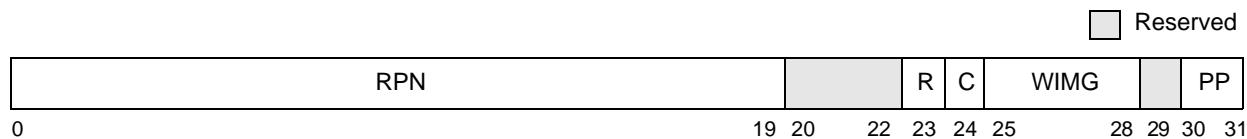
**Figure 5-14. Required Physical Address (RPA) Register**

Table 5-13 describes the bit settings of the RPA register.

Table 5-13. RPA Bit Settings

Bits	Name	Description
0–19	RPN	Physical page number from PTE
20–22	—	Reserved
23	R	Referenced bit from PTE
24	C	Changed bit from PTE
25–28	WIMG	Memory/cache access attribute bits
29	—	Reserved
30–31	PP	Page protection bits from PTE

5.5.2.2 Software Table Search Operation

When a TLB miss occurs, the instruction or data MMU loads IMISS or DMISS, with the effective address of the access. The processor completes all instructions ahead of the instruction that caused the exception, status information is saved in SRR1, and one of the three TLB miss exceptions is taken. In addition, the processor loads ICMP or DCMP with the value to be compared with the first word of PTEs in the tables in memory.

The software should then access the first PTE at the address pointed to by HASH1. The first word of the PTE should be loaded and compared to the contents of DCMP or ICMP. If there is a match, the required PTE has been found and the second word of the PTE is loaded from memory into RPA. Then the **tbl1** or **tbld** instruction is executed, which loads the contents of ICMP or DCMP and RPA into the selected TLB entry. The TLB entry is selected by the effective address of the access and SRR1[WAY].

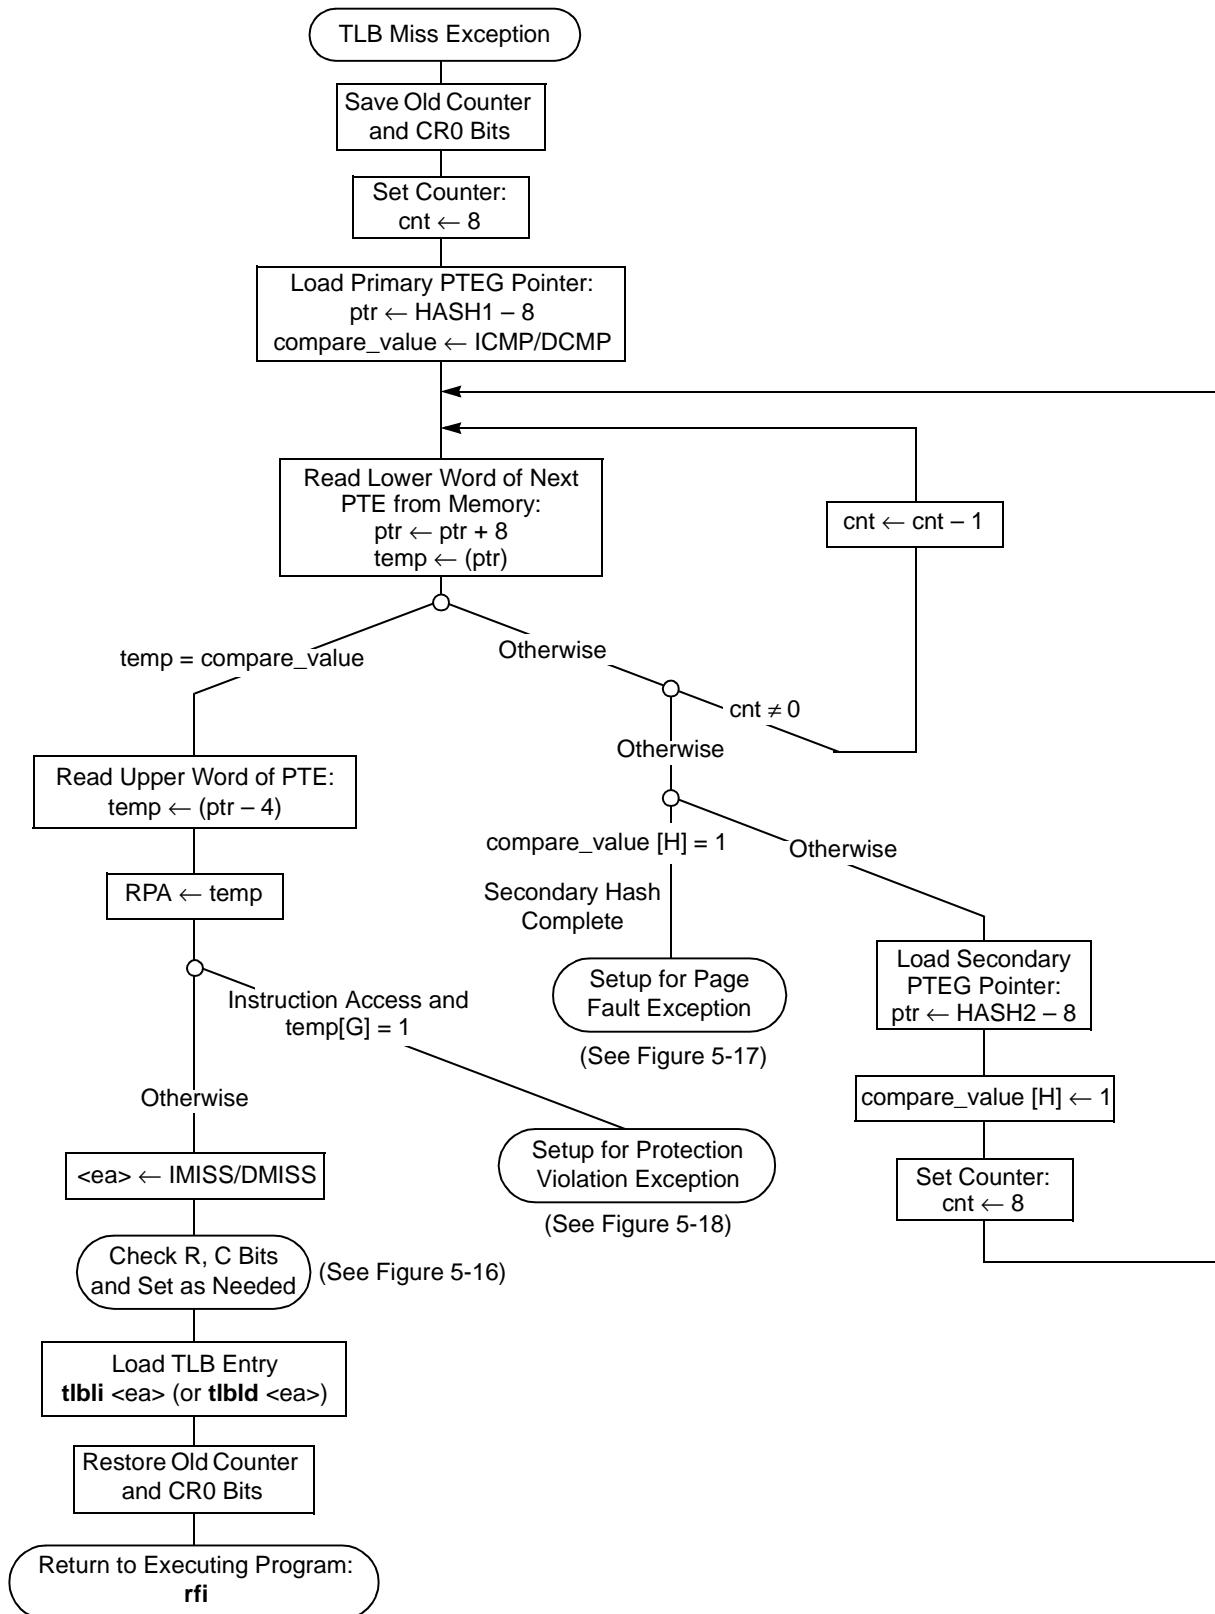
If the comparison does not match, the PTEG address is incremented to point to the next PTE in the table and the above sequence is repeated. If none of the eight PTEs in the primary PTEG matches, the sequence is then repeated using the secondary PTEG (at the address contained in HASH2).

If the PTE is also not found in the eight entries of the secondary page table, a page fault condition exists and a page fault exception must be synthesized. Thus, the appropriate bits must be set in SRR1 (or DSISR) and the TLB miss handler must branch to either the ISI or DS1 exception handler, which handles the page fault condition.

This section provides a flow diagram outlining some example software that can be used to handle the three TLB miss exceptions and sample assembly language that implements that flow.

5.5.2.2.1 Flow for Example Exception Handlers

Figure 5-15 shows the flow for the example TLB miss exception handlers. The flow shown is common for the three exception handlers, except that the IMISS and ICMP registers are used for the instruction TLB miss exception while the DMISS and DCMP registers are used for the two data TLB miss exceptions. Also, for the cases of store instructions that cause either a TLB miss or require a table search operation to update the C bit, the flow shows that the C bit is set in both the TLB entry and PTE in memory. Finally, in the case of a page fault (no PTE found in the table search operation), the setup for the ISI or DS1 exception is slightly different for these two cases.


Figure 5-15. Flow for Example Software Table Search Operation

The flow for checking the R and C bits and setting them appropriately is shown in Figure 5-16.

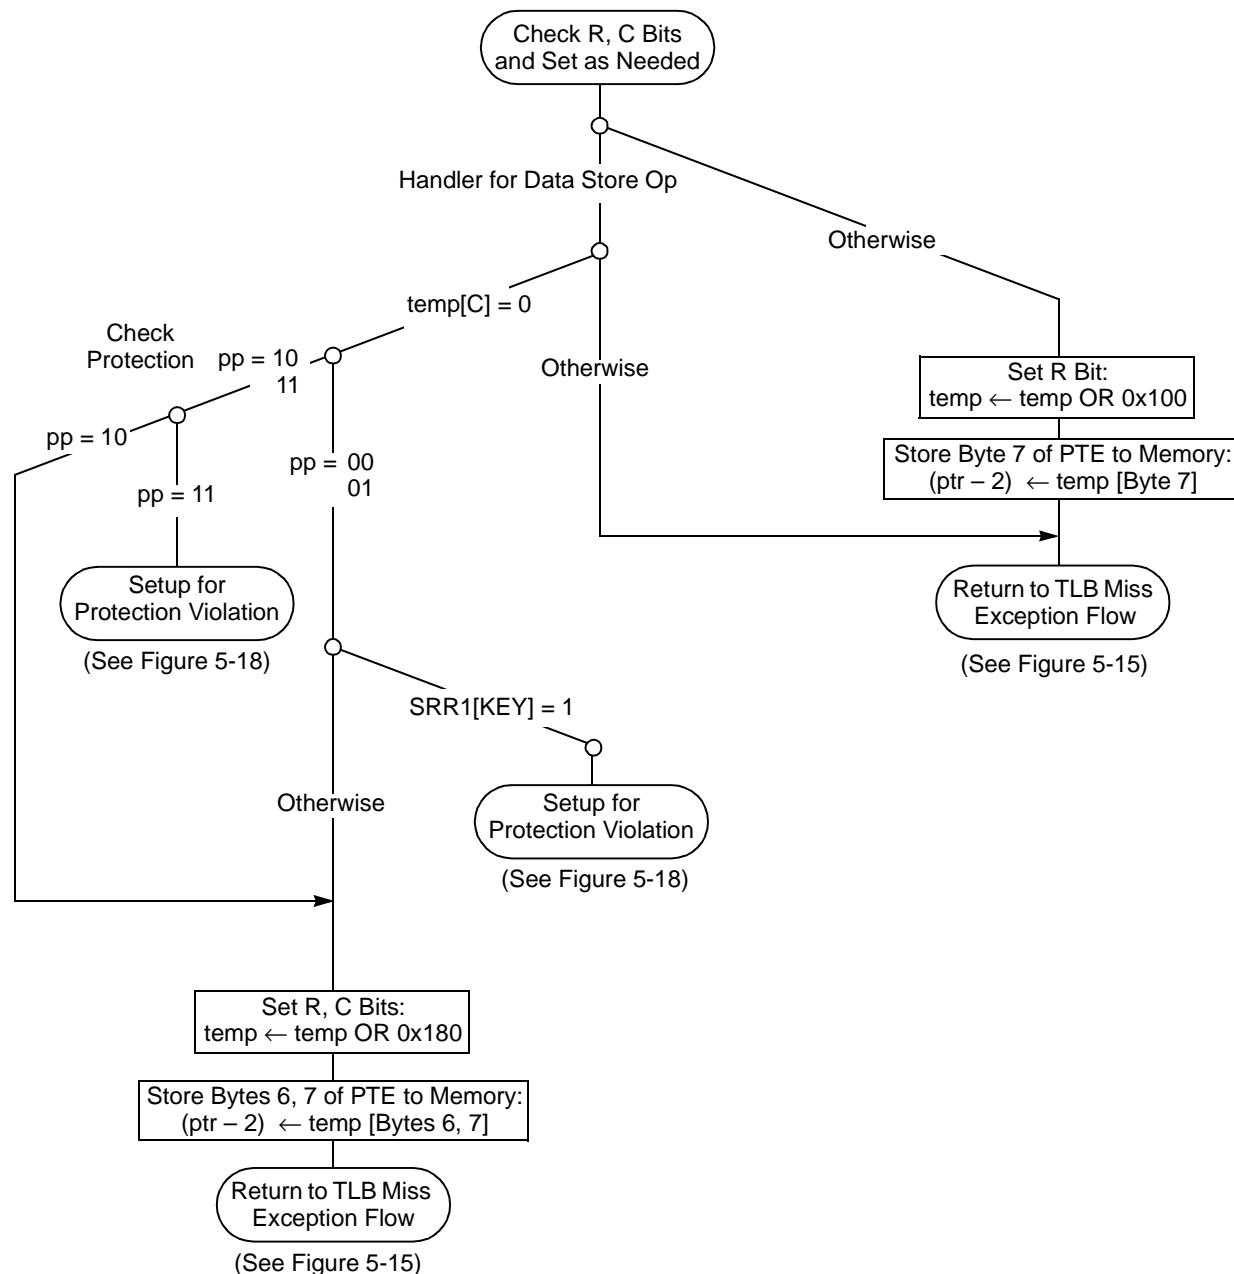


Figure 5-16. Check and Set R and C Bit Flow

Figure 5-17 shows the flow for synthesizing a page fault exception when no PTE is found.

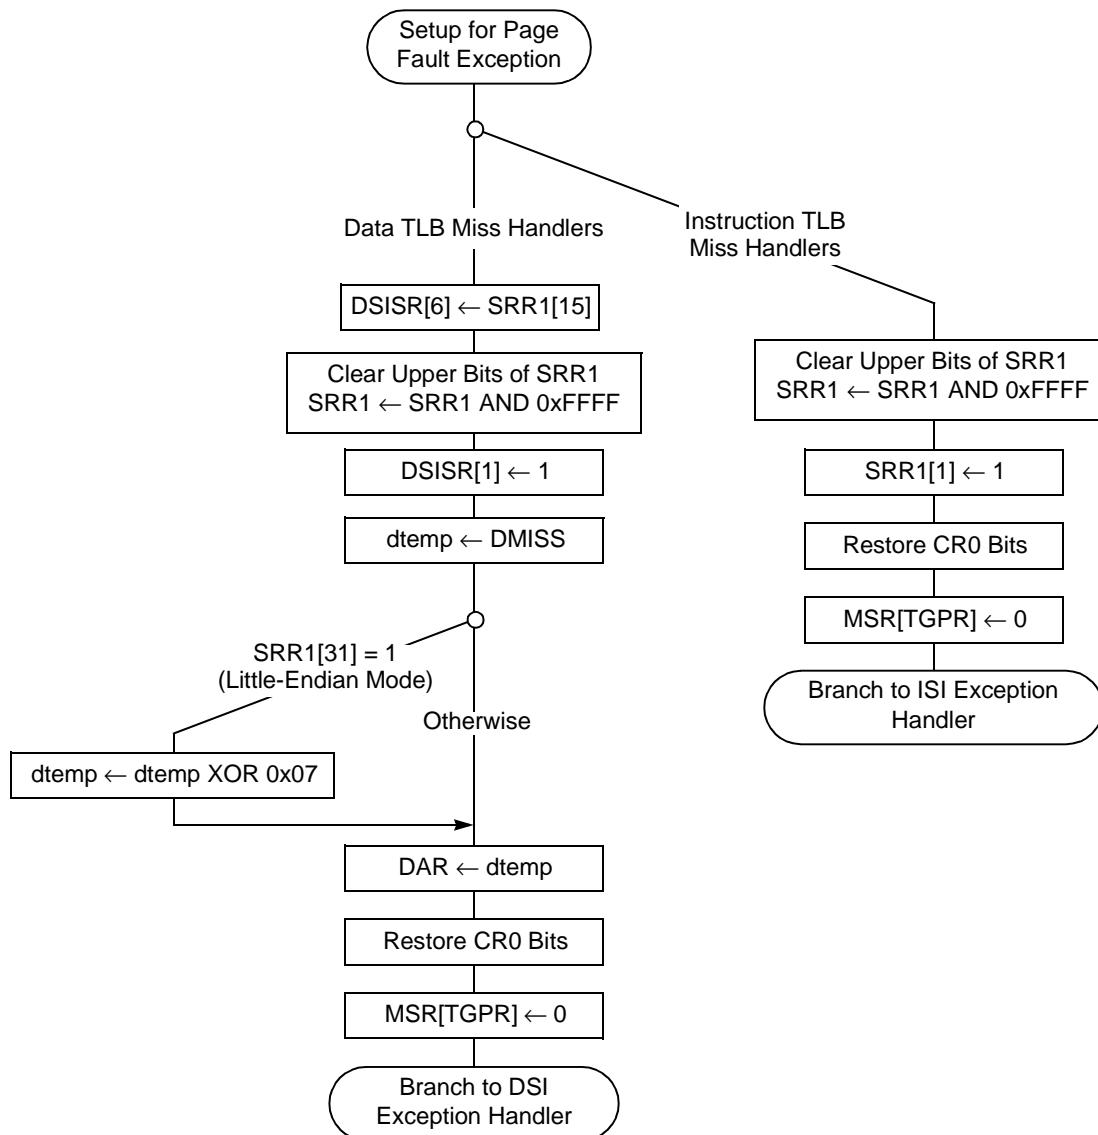


Figure 5-17. Page Fault Setup Flow

Figure 5-18 shows the flow for managing the cases of a TLB miss on an instruction access to guarded memory, and a TLB miss when C = 0 and a protection violation exists. The setup

for these protection violation exceptions is very similar to that of page fault conditions (as shown in Figure 5-17) except that different bits in SRR1 (and DSISR) are set.

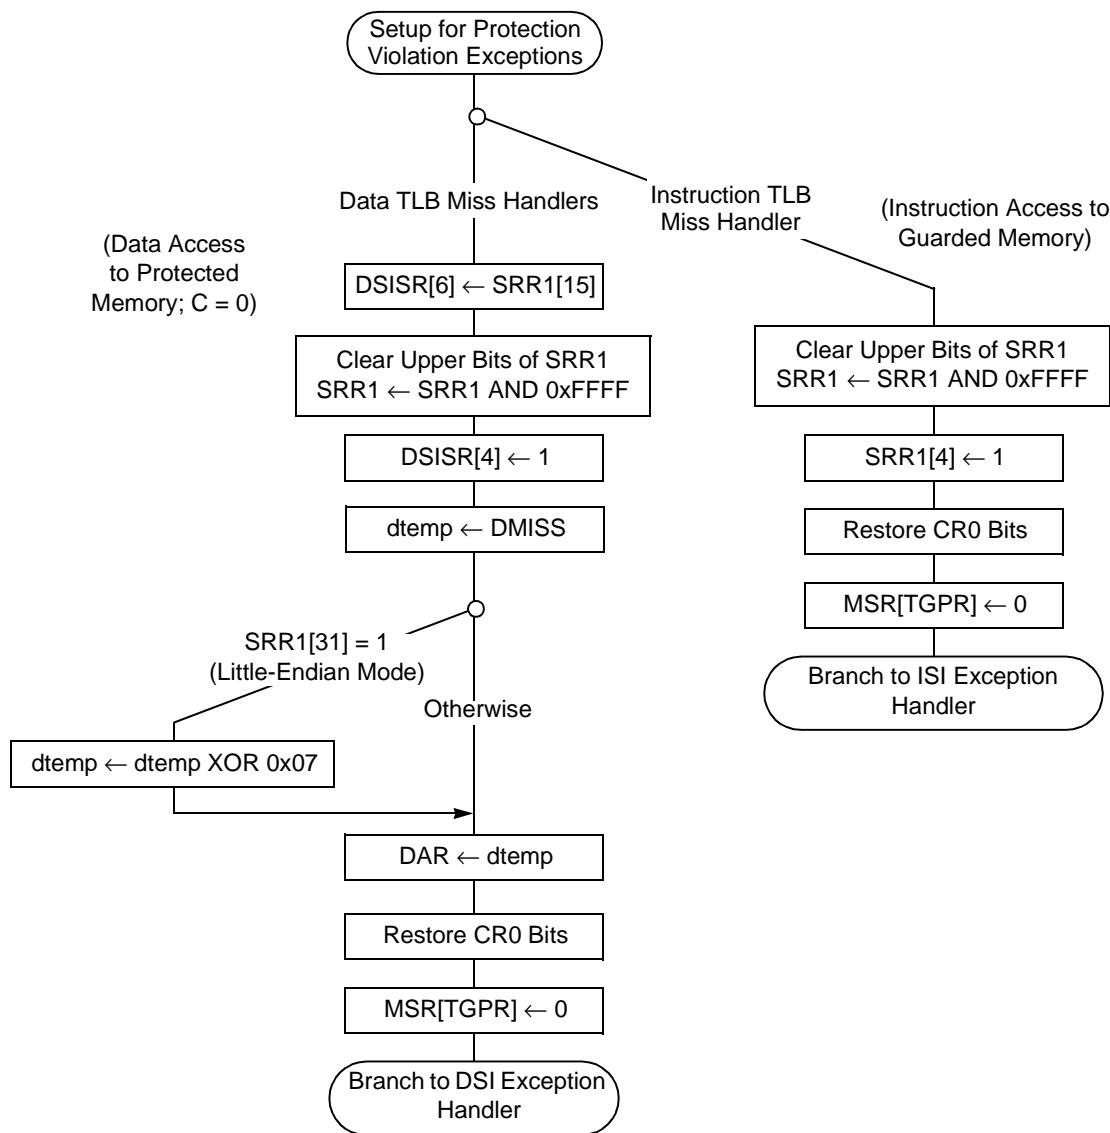


Figure 5-18. Setup for Protection Violation Exceptions

5.5.2.2.2 Code for Example Exception Handlers

This section provides assembly language examples that implement the flow diagrams described above. Note that although these routines fit into a few cache lines, they are supplied only as functional examples; they could be further optimized for faster performance.

```

# TLB software load for MPC603e
#
# New Instructions:
#     tlbld           - write the dtlb with the pte in rpa reg

```

```

#      tblli           - write the itlb with the pte in rpa reg
# New SPRS
#      dmiss          - address of dstream miss
#      imiss          - address of istream miss
#      hash1          - address primary hash PTEG address
#      hash2          - returns secondary hash PTEG address
#      iCmp           - returns the primary istream compare value
#      dCmp           - returns the primary dstream compare value
#      rpa            - the second word of pte used by tlbzx
#
# gpr r0..r3 are shadowed
#
# there are three flows.
#      tlbDataMiss   - tlb miss on data load
#      tlbCeq0        - tlb miss on data store or store with tlb change bit
#                      == 0
#      tlbInstrMiss  - tlb miss on instruction fetch
#+
# place labels for rel branches
#-
#.machine PPC_MPC603e
.set   r0, 0
.set   r1, 1
.set   r2, 2
.set   r3, 3
.set   dMiss, 1010
.set   dCmp, 1011
.set   hash1, 1012
.set   hash2, 1013
.set   iMiss, 1014
.set   iCmp, 1015
.set   rpa, 1010
.set   c0, 0
.set   dar, 19
.set   dsisr, 18
.set   srr0, 26
.set   srr1, 27
.
.csect tlbmiss[PR]
vec0:
.globl vec0
.org   vec0+0x300

vec300:
.org   vec0+0x400
vec400:
#+
# Instruction TB miss flow
# Entry:
#      Vec = 1000
#      srr0      -> address of instruction that missed
#      srr1      -> 0:3=cr0 4=lru way bit 16:31 = saved MSR
#      msr<tgpr> -> 1

```

```

#      iMiss      -> ea that missed
#      iCmp       -> the compare value for the va that missed

#      hash1      -> pointer to first hash pteg
#      hash2      -> pointer to second hash pteg
#
# Register usage:
#      r0 is saved counter
#      r1 is junk
#      r2 is pointer to pteg
#      r3 is current compare value

.org    vec0+0x1000

tlbInstrMiss:
    mfspr   r2, hash1      # get first pointer
    addi    r1, 0, 8        # load 8 for counter
    mfctr   r0              # save counter
    mfspr   r3, iCmp       # get first compare value
    addi    r2, r2, -8      # pre dec the pointer
im0:   mtctr   r1              # load counter
im1:   lwzu    r1, 8(r2)     # get next pte
    cmp     c0, r1, r3      # see if found pte
    bdnzf   0, im1         # dec count br if cmp ne and if count not zero

    bne     instrSecHash   # if not found set up second hash or exit
    l      r1, +4(r2)       # load tlb entry lower-word
    andi.  r3, r1, 8        # check G bit
    bne     doISIp          # if guarded, take an ISI
    mtctr   r0              # restore counter
    mfspr   r0, iMiss       # get the miss address for the tlbli
    mfspr   r3, srr1        # get the saved cr0 bits
    mtcrf   0x80, r3        # restore CR0
    mtspr   rpa, r1          # set the pte
    ori     r1, r1, 0x100    # set reference bit
    srw     r1, r1, 8        # get byte 7 of pte
    tlbli   r0              # load the itlb
    stb     r1, +6(r2)       # update page table
    rfi     r1               # return to executing program

#+
# Register usage:
#      r0 is saved counter
#      r1 is junk
#      r2 is pointer to pteg
#      r3 is current compare value
#-

instrSecHash:
    andi.  r1, r3, 0x0040  # see if we have done second hash
    bne    doISI            # if so, go to ISI exception
    mfspr   r2, hash2       # get the second pointer
    ori     r3, r3, 0x0040  # change the compare value
    addi    r1, 0, 8        # load 8 for counter
    addi    r2, r2, -8      # pre dec for update on load
    b      im0              # try second hash

```

```

#+
# entry Not Found: synthesize an ISI exception
# guarded memory protection violation: synthesize an ISI exception
# Entry:
#      r0 is saved counter
#      r1 is junk
#      r2 is pointer to pteg
#      r3 is current compare value
#
doISIP:
    mfspr    r3, srr1      # get srr1
    andi.   r2,r3,0xfffff  # clean upper srr1
    addis   r2, r2, 0x0800  # or in srr<4> = 1 to flag prot violation
    b      isil:
doISI:
    mfspr    r3, srr1      # get srr1
    andi.   r2, r3, 0xfffff # clean srr1
    addis   r2, r2, 0x4000  # or in srr1<1> = 1 to flag pte not found
isil   mtctr    r0          # restore counter
    mtspr    srr1, r2       # set srr1
    mfmsr    r0          # get msr
    xoris   r0, r0, 0x8000 # flip the msr<tgpr> bit
    mtcrf   0x80, r3       # restore CR0
    mtmsr    r0          # flip back to the native gprs
    b      vec400         # go to instr. access exception
#
#+
# Data TLB miss flow
# Entry:
#      Vec = 1100
#      srr0      -> address of instruction that caused data tlb miss
#      srr1      -> 0:3=cr0 4=lru way bit 5=1 if store 16:31 = saved MSR
#      msr<tgpr> -> 1
#      dMiss     -> ea that missed
#      dCmp      -> the compare value for the va that missed
#      hash1     -> pointer to first hash pteg
#      hash2     -> pointer to second hash pteg
#
# Register usage:
#      r0 is saved counter
#      r1 is junk
#      r2 is pointer to pteg
#      r3 is current compare value
#-
.csect  tlbmiss[PR]
.org    vec0+0x1100

tlbDataMiss:
    mfspr    r2, hash1      # get first pointer
    addi     r1, 0, 8        # load 8 for counter
    mfctr    r0          # save counter
    mfspr    r3, dCmp       # get first compare value
    addi     r2, r2, -8      # pre dec the pointer

```

Chapter 5. Memory Management

```

dm0:    mtctr   r1          # load counter
dm1:    lwzu     r1, 8(r2)   # get next pte
        cmp      c0, r1, r3   # see if found pte
        bdnzf   0, dm1       # dec count br if cmp ne and if count not zero
        bne     dataSecHash  # if not found set up second hash or exit
        l      r1, +4(r2)    # load tlb entry lower-word
        mtctr   r0          # restore counter
        mfspr   r0, dMiss    # get the miss address for the tlbld
        mfspr   r3, srr1     # get the saved cr0 bits
        mtcrf   0x80, r3     # restore CR0
        mtspr   rpa, r1      # set the pte
        ori     r1, r1, 0x100  # set reference bit
        srw     r1, r1, 8     # get byte 7 of pte
        tlbld   r0          # load the dtlb
        stb     r1, +6(r2)    # update page table
        rfi      # return to executing program

#+
# Register usage:
#     r0 is saved counter
#     r1 is junk
#     r2 is pointer to pteg
#     r3 is current compare value
#-

dataSecHash:
        andi   r1, r3, 0x0040  # see if we have done second hash
        bne    doDSI           # if so, go to DSI exception
        mfspr  r2, hash2       # get the second pointer
        ori    r3, r3, 0x0040  # change the compare value
        addi   r1, 0, 8         # load 8 for counter
        addi   r2, r2, -8       # pre dec for update on load
        b     dm0              # try second hash
#
#+
# C=0 in dtlb and dtlb miss on store flow
# Entry:
#     Vec = 1200
#     srr0    -> address of store that caused the exception
#     srr1    -> 0:3=cr0 4=lru way bit 5=1 16:31 = saved MSR
#     msr<tgpr> -> 1
#     dMiss   -> ea that missed
#     dCmp    -> the compare value for the va that missed
#     hash1   -> pointer to first hash pteg
#     hash2   -> pointer to second hash pteg
#
# Register usage:
#     r0 is saved counter
#     r1 is junk
#     r2 is pointer to pteg
#     r3 is current compare value
#-

```

```

.csect  tlbmiss[PR]
.org   vec0+0x1200

tlbCeq0:
    mfspr  r2, hash1      # get first pointer
    addi   r1, 0, 8        # load 8 for counter
    mfctr  r0              # save counter
    mfspr  r3, dCmp       # get first compare value
    addi   r2, r2, -8      # pre dec the pointer
    ceq0:  mtctr  r1        # load counter
    ceq1:  lwzu   r1, 8(r2)  # get next pte
    cmp    c0, r1, r3      # see if found pte
    bdnzf 0, ceq1         # dec count br if cmp ne and if count not zero
    bne   cEq0SecHash     # if not found set up second hash or exit
    l     r1, +4(r2)       # load tlb entry lower-word
    andi. r3,r1,0x80       # check the C-bit
    beq   cEq0ChkProt     # if (C==0) go check protection modes
    ceq2:  mtctr  r0        # restore counter
    mfspr  r0, dMiss       # get the miss address for the tlbld
    mfspr  r3, srr1       # get the saved cr0 bits
    mtcrf 0x80, r3        # restore CR0
    mtspr  rpa, r1         # set the pte
    tlbld  r0              # load the dtlb
    rfi

#+
# Register usage:
#     r0 is saved counter
#     r1 is junk
#     r2 is pointer to pteg
#     r3 is current compare value
#-

cEq0SecHash:
    andi. r1, r3, 0x0040  # see if we have done second hash
    bne   doDSI            # if so, go to DSI exception
    mfspr  r2, hash2       # get the second pointer
    ori    r3, r3, 0x0040  # change the compare value
    addi   r1, 0, 8        # load 8 for counter
    addi   r2, r2, -8      # pre dec for update on load
    b     ceq0              # try second hash

#+
# entry found and PTE(c-bit==0):
# (check protection before setting PTE(c-bit))
# Register usage:
#     r0 is saved counter
#     r1 is PTE entry
#     r2 is pointer to pteg
#     r3 is trashed
#-

cEq0ChkProt:
    rlwinm. r3,r1,30,0,1    # test PP
    bge-   chk0             # if (PP==00 or PP==01) goto chk0:

```

```

        andi.    r3,r1,1          # test PP[0]
        beq+    chk2              # return if PP[0]==0
        b      doDSIp           # else DSIP

chk0:   mfspr    r3,srr1      # get old msr
        andis.   r3,r3,0x0008  # test the KEY bit (SRR1-bit 12)
        beq     chk2              # if (KEY==0) goto chk2:
        b      doDSIp           # else DSIP

chk2:   ori      r1, r1, 0x180 # set reference and change bit
        sth     r1, 6(r2)       # update page table
        b      ceq2              # and back we go
#
#+
# entry Not Found: synthesize a DSI exception
# Entry:
#     r0 is saved counter
#     r1 is junk
#     r2 is pointer to pteg
#     r3 is current compare value
#
doDSI:
        mfspr    r3, srr1      # get srr1
        rlwinm   r1, r3, 9,6,6  # get srr1<flag> to bit 6 for load/store, zero
                                rest
        addis    r1, r1, 0x4000  # or in dsisr<1> = 1 to flag pte not found
        b      dsil:

doDSIp:
        mfspr    r3, srr1      # get srr1
        rlwinm   r1, r3, 9,6,6  # get srr1<flag> to bit 6 for load/store, zero
                                rest
        addis    r1, r1, 0x0800  # or in dsisr<4> = 1 to flag prot violation

dsil:   mtctr    r0          # restore counter
        andi.    r2, r3, 0xffff  # clear upper bits of srr1
        mtspr    srr1, r2        # set srr1
        mtspr    dsisr, r1        # load the dsisr
        mfspr    r1, dMiss      # get miss address
        rlwinm. r2,r2,0,31,31  # test LE bit
        beq     dsil2:           # if little endian then:
        xor     r1,r1,0x07        # de-mung the data address

dsi2:   mtspr    dar, r1      # put in dar
        mfmsr    r0          # get msr
        xoris   r0, r0, 0x2      # flip the msr<tgpr> bit
        mtcrf   0x80, r3        # restore CR0
        mtmsr   r0          # flip back to the native gprs
        b      vec300            # branch to DSI exception

```

5.5.3 Page Table Updates

TLBs are defined as noncoherent caches of the PTEs. TLB entries must be flushed explicitly with the TLB invalidate entry instruction (**tlbie**) whenever the corresponding PTE is modified. Because the MPC603e is intended primarily for uniprocessor environments, it does not provide coherency checking for TLBs between multiple processors. If the MPC603e is used in a multiprocessor environment where TLB coherency is required, synchronization must be implemented in software.

Processors may write referenced and changed bits with unsynchronized, atomic byte store operations. Note that each V, R, and C bits reside in a distinct byte of a PTE. Therefore, extreme care must be taken to use byte writes when updating only one of these bits.

Explicitly altering certain MSR bits (using the **mtmsr** instruction), PTEs, or certain system registers, may have the side effect of changing the effective or physical addresses from which the current instruction stream is being fetched. This kind of side effect is defined as an implicit branch. Implicit branches are not supported and an attempt to perform one causes boundedly-undefined results. Therefore, PTEs must not be changed in a manner that causes an implicit branch. Chapter 2, “Register Set,” in the *Programming Environments Manual*, lists the possible implicit branch conditions that can occur when system registers and MSR bits are changed.

5.5.4 Segment Register Updates

Synchronization requirements for using the move to segment register instructions (**mts** and **mtsrin**) are described in “Synchronization Requirements for Special Registers and for Lookaside Buffers” in Chapter 2, “Register Set,” in the *Programming Environments Manual*.



Chapter 6

Instruction Timing

This chapter describes how the MPC603e microprocessor fetches, dispatches, and executes instructions and how it reports the results of instruction execution. It gives detailed descriptions of how the MPC603e execution units work, and how those units interact with other parts of the processor, such as the instruction fetching mechanism, register files, and caches. It gives examples of instruction sequences, showing potential bottlenecks and how to minimize their effects. Finally, it includes tables that identify the unit that executes each instruction implemented on the MPC603e, the latency for each instruction, and other information that is useful for the assembly language programmer.

6.1 Terminology and Conventions

This section provides an alphabetical glossary of terms used in this chapter. These definitions are provided as a review of commonly used terms and as a way to point out specific ways these terms are used in this chapter.

- Branch prediction—The process of guessing whether a branch will be taken. Such predictions can be correct or incorrect; the term predicted as it is used here does not imply that the prediction is correct (successful). The PowerPC architecture defines a means for static branch prediction as part of the instruction encoding.
- Branch resolution—The determination of whether a branch is taken or not taken. A branch is said to be resolved when the processor can determine which instruction path to take. If the branch is resolved as predicted, the instructions following the predicted branch that may have been speculatively executed can complete (see completion). If the branch is not resolved as predicted, instructions on the mispredicted path, and any results of speculative execution, are purged from the pipeline and fetching continues from the nonpredicted path.
- Completion—Completion occurs when an instruction has finished executing, written back any results, and is removed from the completion queue (CQ). When an instruction completes, it is guaranteed that this instruction and all previous instructions can cause no exceptions.
- Fall-through (branch fall-through)—A not-taken branch. On the MPC603e, fall-through branch instructions are removed from the instruction stream at dispatch. That is, these instructions are allowed to fall through the instruction queue through

the dispatch mechanism, without either being passed to an execution unit and or given a position in the CQ.

- Fetch—The process of bringing instructions from memory (such as a cache or system memory) into the instruction queue.
- Finish—Finishing occurs in the last cycle of execution. In this cycle, the CQ entry is updated to indicate that the instruction has finished executing.
- Folding (branch folding)—The replacement of a branch instruction with target instructions and any instructions along the not-taken path, when a branch is either taken or predicted as taken.
- Latency—The number of clock cycles necessary to execute an instruction and make ready the results of that execution for a subsequent instruction.
- Pipeline—In the context of instruction timing, the term pipeline refers to the interconnection of the stages. The events necessary to process an instruction are broken into several cycle-length tasks to allow work to be performed on several instructions simultaneously—analogous to an assembly line. As an instruction is processed, it passes from one stage to the next. When it does, the stage becomes available for the next instruction.

Although an individual instruction may take many cycles to complete (the number of cycles is called instruction latency), pipelining makes it possible to overlap the processing so that the throughput (number of instructions completed per cycle) is greater than if pipelining were not implemented.

- Program order—The order of instructions in an executing program. More specifically, this term is used to refer to the original order in which program instructions are fetched into the instruction queue from the cache.
- Rename register—Temporary buffers used by instructions that have finished execution but have not completed.
- Reservation station—A buffer between the dispatch and execute stages that allows instructions to be dispatched even though the results of instructions on which the dispatched instruction may depend are not available.
- Retirement—Removal of the completed instruction from the CQ.
- Stage—The term stage is used in two different senses, depending on whether the pipeline is being discussed as a physical entity or a sequence of events. In the latter case, a stage is an element in the pipeline during which certain actions are performed, such as decoding the instruction, performing an arithmetic operation, or writing back the results. A stage is typically described as taking a processor clock cycle to perform its operation; however, some events (such as dispatch and write-back) happen instantaneously, and may be thought to occur at the end of the stage.

An instruction can spend multiple cycles in one stage. An integer multiply, for example, takes multiple cycles in the execute stage. When this occurs, subsequent instructions may stall.

In some cases, an instruction may also occupy more than one stage simultaneously, especially in the sense that a stage can be seen as a physical resource—for example, when instructions are dispatched they are assigned a place in the CQ at the same time they are passed to the execute stage. They can be said to occupy both the complete and execute stages in the same clock cycle.

- Stall—An occurrence when an instruction cannot proceed to the next stage.
- Superscalar—A superscalar processor is one that can dispatch multiple instructions concurrently from a conventional linear instruction stream. In a superscalar implementation, multiple instructions can be in the same stage at the same time.
- Throughput—A measure of the number of instructions that are processed per cycle. For example, a series of double-precision floating-point multiply instructions has a throughput of one instruction per clock cycle.
- Write-back—Write-back (in the context of instruction handling) occurs when a result is written from the rename registers into the architectural registers (typically the GPRs and FPRs). Results are written back at completion time or are moved into the write-back buffer. Results in the write-back buffer cannot be flushed. If an exception occurs, these buffers must write back before the exception is taken.

6.2 Instruction Timing Overview

The MPC603e design minimizes average instruction execution latency, the number of clock cycles it takes to fetch, decode, dispatch, and execute instructions and make the results available for a subsequent instruction. Some instructions, such as loads and stores, access memory and require additional clock cycles between the execute phase and the write-back phase. These latencies vary depending on whether the access is to cacheable or noncacheable memory, whether it hits in the L1 cache, whether the cache access generates a write-back to memory, whether the access causes a snoop hit from another device that generates additional activity, and other conditions that affect memory accesses.

The MPC603e implements many features to improve throughput, such as pipelining, superscalar instruction dispatch, branch folding, removal of fall-through branches, two-level speculative branch handling, and multiple execution units that operate independently and in parallel.

As an instruction of load/store and floating-point units passes from stage to stage in a pipelined system, the following instruction can follow through the stages as the former instruction vacates them, allowing several instructions to be processed simultaneously. While it may take several cycles for an instruction to pass through all the stages, when the pipeline has been filled, one instruction can complete its work on every clock cycle.

Figure 6-1 represents a generic pipelined execution unit.

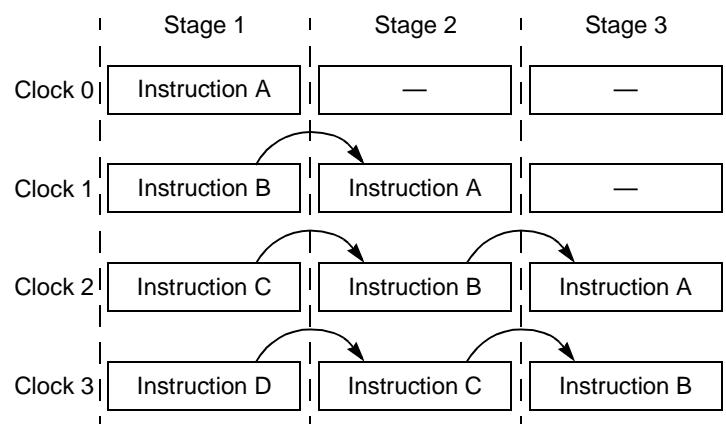


Figure 6-1. Pipelined Execution Unit

The entire path that instructions take through the fetch, decode/dispatch, execute, complete, and write-back stages is considered the MPC603e master pipeline, and two of the MPC603e execution units (FPU and LSU) are also multiple-stage pipelines.

The MPC603e contains the following execution units that operate independently and in parallel:

- Branch processing unit (BPU)
- 32-bit integer unit (IU)—executes all integer instructions
- 64-bit floating-point unit (FPU)
- Load/store unit (LSU)
- System register unit (SRU)

The MPC603e can retire two instructions on every clock cycle. In general, the MPC603e processes instructions in four stages—fetch, decode/dispatch, execute, and complete as shown in Figure 6-2. Note that the example of a pipelined execution unit in Figure 6-1 is similar to the three-stage FPU pipeline in Figure 6-2.

The instruction pipeline stages are described as follows:

- The instruction fetch stage includes the clock cycles necessary to request instructions from the memory system and the time the memory system takes to respond to the request. Instruction fetch timing depends on many variables, such as whether the instruction is in the branch target instruction cache, or in the on-chip instruction cache. Instruction fetch timing increases when it is necessary to fetch instructions from system memory. The variables that affect fetch timing include the processor-to-bus clock ratio, the amount of bus traffic, and whether any cache coherency operations are required.

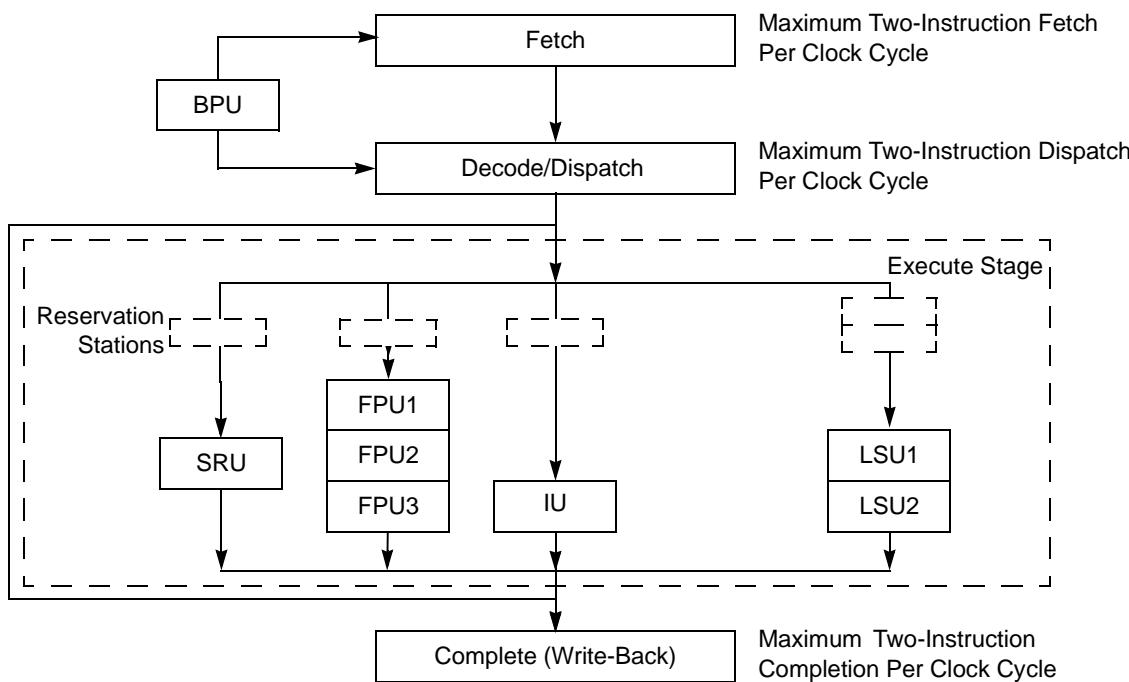


Figure 6-2. Superscalar/Pipeline Diagram

Because there are so many variables, unless otherwise specified, the instruction timing examples below assume optimal performance and that the instructions are available in the instruction queue in the same clock cycle that they are requested. The fetch stage ends when the instruction is dispatched.

- The decode/dispatch stage consists of the time it takes to fully decode the instruction and dispatch it from the instruction queue to the appropriate execution unit. Instruction dispatch requires the following:
 - Instructions can be dispatched only from the two lowest instruction queue entries, IQ0 and IQ1.
 - A maximum of two instructions can be dispatched per clock cycle.
 - Only one instruction can be dispatched to each execution unit per clock cycle.
 - There must be a vacancy in the specified execution unit.
 - A rename register must be available for each destination operand specified by the instruction.
 - For an instruction to dispatch, the appropriate execution unit must be available and there must be an open position in the CQ. If no entry is available, the instruction remains in the IQ.
- The execute stage consists of the time between dispatch to the execution unit (or reservation station) and the point at which the instruction vacates the execution unit. Most integer instructions have a one-cycle latency; results of these instructions can be used in the clock cycle after an instruction enters the execution unit. However,

integer multiply and divide instructions take multiple clock cycles to complete. The IU can process all integer instructions.

The LSU and FPU are pipelined, as shown in Figure 6-2.

- The complete (complete/write-back) pipeline stage maintains the correct architectural machine state and commits it to the architectural registers at the proper time. If the completion logic detects an instruction containing an exception status, all following instructions are canceled, their execution results in rename registers are discarded, and the correct instruction stream is fetched.

The complete stage ends when the instruction is retired. Two instructions can be retired per cycle. Instructions are retired only from the two lowest CQ entries, CQ0 and CQ1.

The notation conventions used in the instruction timing examples are as follows:

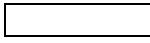
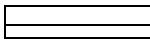
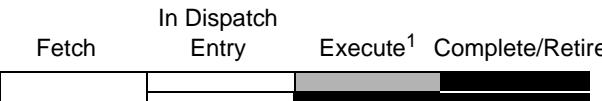
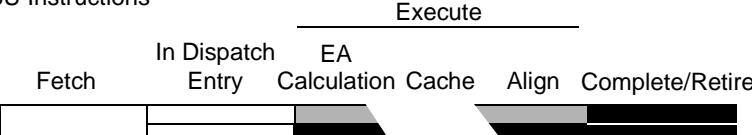
-  Fetch—The fetch stage includes the time between when an instruction is requested and when it is brought into the instruction queue. This latency can vary greatly, depending on whether the instruction is in the on-chip cache or system memory (in which case latency can be affected by bus speed and traffic on the system bus, and address translation dispatches). Therefore, in the examples in this chapter, the fetch stage is usually idealized; that is, an instruction is usually shown to be in the fetch stage when it is a valid instruction in the instruction queue. The instruction queue has six entries, IQ0–IQ5.
-  In dispatch entry (IQ0/IQ1)—Instructions can be dispatched from IQ0 and IQ1. Because dispatch is instantaneous, it is perhaps more useful to describe it as an event that marks the point in time between the last cycle in the fetch stage and the first cycle in the execute stage.
-  Execute—The operations specified by an instruction are being performed by the appropriate execution unit. The black stripe is a reminder that the instruction occupies an entry in the CQ, described in Figure 6-3.
-  Complete—The instruction is in the CQ. In the final stage, the results of the executed instruction are written back and the instruction is retired. The CQ has five entries, CQ0–CQ4.
-  In retirement entry—Completed instructions can be retired from CQ0 and CQ1. Like dispatch, retirement is an event that in this case occurs at the end of the final cycle of the complete stage.

Figure 6-3 shows the stages of MPC603e execution units.

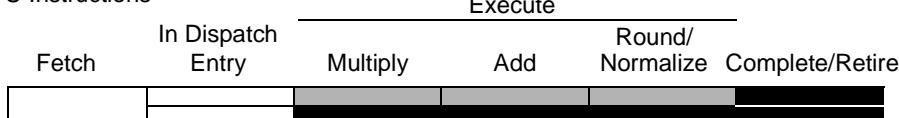
IU/SRU Instructions



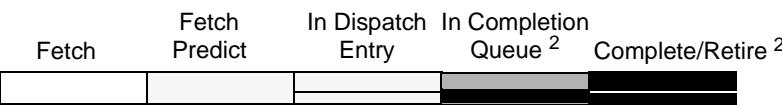
LSU Instructions



FPU Instructions



BPU Instructions



¹ Several integer instructions, such as multiply and divide instructions, require multiple cycles in the execute stage.

² Only those branch instructions that update the LR or CTR take an entry in the completion queue.

Figure 6-3. MPC603e Microprocessor Pipeline Stages

6.3 Timing Considerations

The MPC603e is a superscalar processor; as many as three instructions can be dispatched to the execution units (one branch instruction to the branch processing unit, and two instructions dispatched from the dispatch queue to the other execution units) during each clock cycle. Only one instruction can be dispatched to each execution unit.

Although instructions appear to the programmer to execute in program order, the MPC603e improves performance by executing multiple instructions at a time, using hardware to manage dependencies. When an instruction is dispatched, the register file provides the source data to the execution unit. The register files and rename register have sufficient bandwidth to allow dispatch of two instructions per clock under most conditions.

The MPC603e BPU decodes and executes branches immediately after they are fetched. When a conditional branch cannot be resolved due to a CR data dependency, the branch direction is predicted and execution continues from the predicted path. If the prediction is incorrect, the following steps are taken:

1. The instruction queue is purged and fetching continues from the correct path.
2. Any instructions ahead of the predicted branch in the CQ are allowed to complete.

3. Instructions after the mispredicted branch are purged.
4. Dispatching resumes from the correct path.

After an execution unit executes an instruction, it places resulting data into the appropriate GPR or FPR rename register. The results are then stored into the correct GPR or FPR during the write-back stage. If a subsequent instruction needs the result as a source operand, it is made available simultaneously to the appropriate execution unit, which allows a data-dependent instruction to be decoded and dispatched without waiting to read the data from the register file. Branch instructions that update either the LR or CTR write back their results in a similar fashion.

The following section describes this process in greater detail.

6.3.1 General Instruction Flow

As many as two instructions can be fetched into the instruction queue (IQ) in a single clock cycle. Instructions enter the IQ and are dispatched to the various execution units from the dispatch queue. The IQ is a six-entry queue, which together with the CQ is the backbone of the master pipeline for the microprocessor. The MPC603e tries to keep the IQ full at all times.

The number of instructions requested in a clock cycle is determined by the number of vacant spaces in the IQ during the previous clock cycle. This is shown in the examples in this chapter. Although the IQ can accept as many as two new instructions in a single clock cycle and even if there are more than two spaces available on the current clock cycle, if only one IQ entry was vacant on the previous cycle, only one instruction is fetched. Typically, instructions are fetched from the on-chip instruction cache. If the instruction request hits in the on-chip instruction cache, it can usually present the first two instructions of the new instruction stream in the next clock cycle, giving enough time for the next pair of instructions to be fetched from the cache with no idle cycles. Instructions not in the instruction cache are fetched from system memory.

Branch instructions that do not update the LR or CTR are removed from the instruction stream either by branch folding or removal of fall-through branch instructions, as described in Section 6.4.1.1, “Branch Folding.” Branch instructions that update the LR or CTR are treated as if they require dispatch (even though they are not dispatched to an execution unit in the process). They are assigned a position in the CQ to ensure that the CTR and LR are updated sequentially.

All other instructions are dispatched from IQ0 and IQ1. The dispatch rate depends on the availability of resources such as the execution units, rename registers, and CQ entries, and on the serializing behavior of some instructions. Instructions are dispatched in program order; an instruction in IQ1 can be dispatched at the same time as one in IQ0, but it cannot be dispatched ahead of one in IQ0.

Figure 6-4 shows the paths taken by instructions.

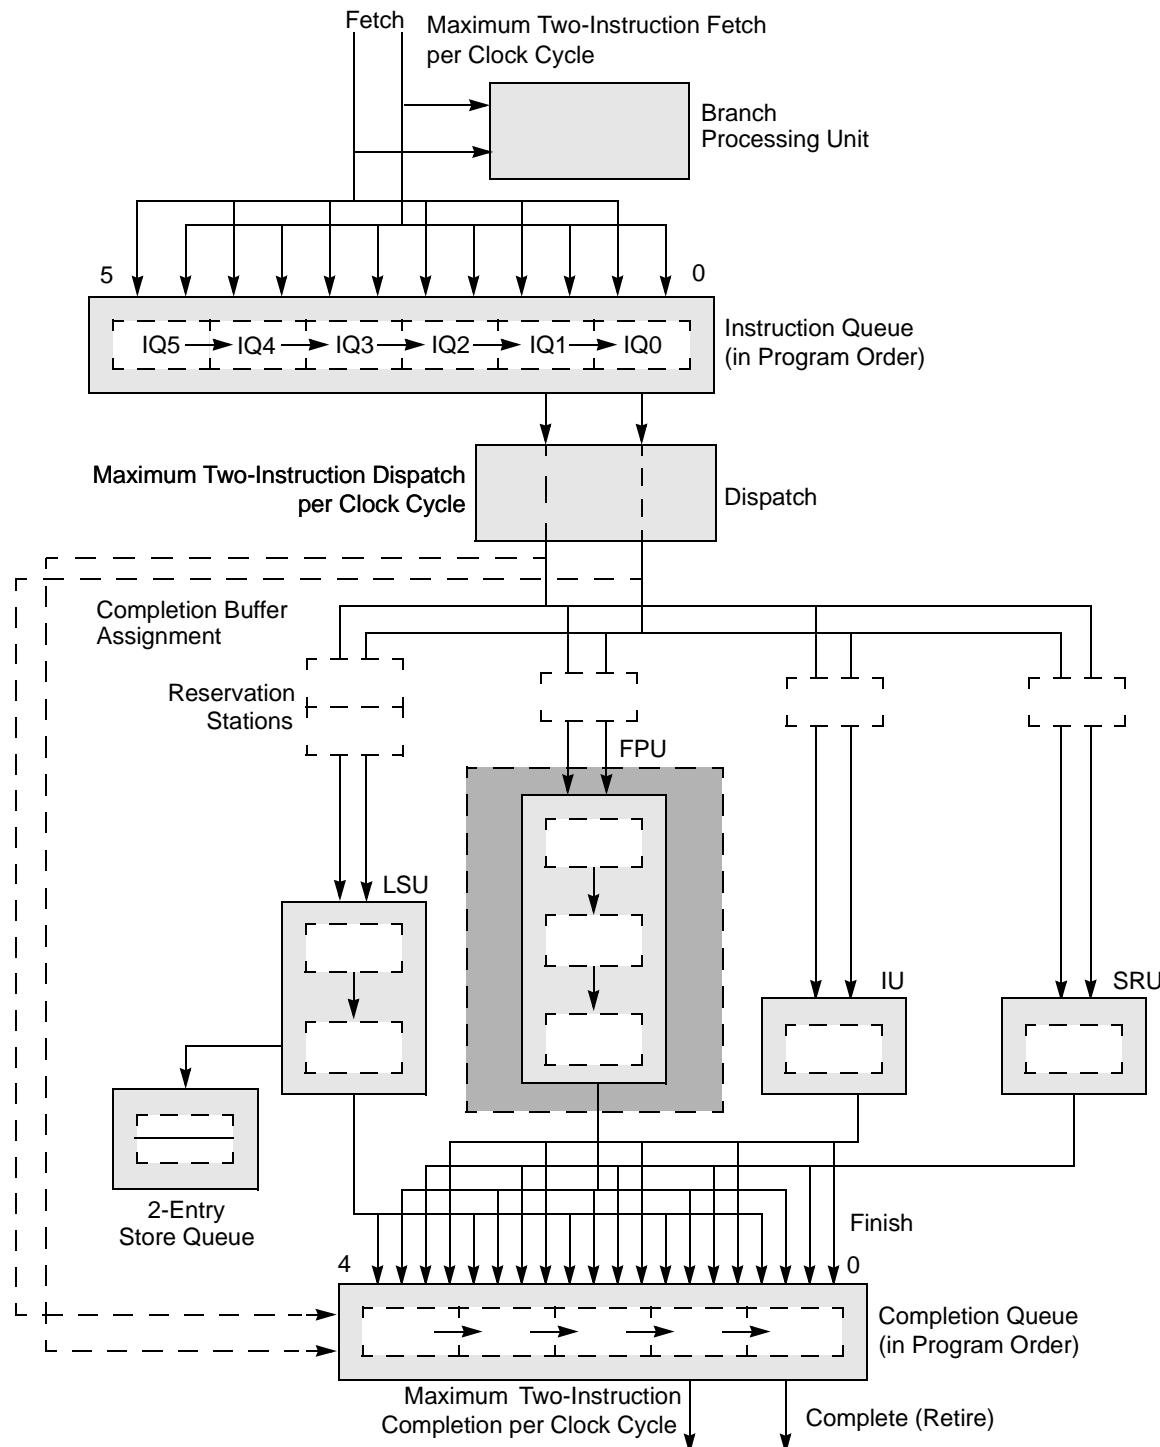


Figure 6-4. Instruction Flow Diagram

Instruction state and all information required for completion is kept in the five-entry, FIFO completion queue. A completion queue entry is allocated for each instruction when it is dispatched to an execute unit; if no entry is available, the dispatch unit stalls. A maximum

of two instructions per cycle may be completed and retired from the completion queue, and the flow of instructions can stall when a longer-latency instruction reaches the last position in the completion queue. Store instructions and instructions executed by the FPU and SRU (with the exception of integer add and compare instructions) can only be retired from the last position in the completion queue. Subsequent instructions cannot be completed and retired until that longer-latency instruction completes and retires. Examples of this are shown in Section 6.3.2.2, “Cache Hit,” and Section 6.3.2.3, “Cache Miss.”

The rate of instruction completion is also affected by the ability to write instruction results from the rename registers to the architected registers. The MPC603e can perform two write-back operations from the rename registers to the GPRs each clock cycle, but can perform only one write-back per cycle to the CR, FPR, LR, and CTR.

6.3.2 Instruction Fetch Timing

Instruction fetch latency depends on the fetch hits of the on-chip instruction cache. If no hit occurs, a memory transaction is required, in which case fetch latency is affected by bus traffic, bus clock speed, and memory translation. These conditions are discussed in the following sections.

6.3.2.1 Cache Arbitration

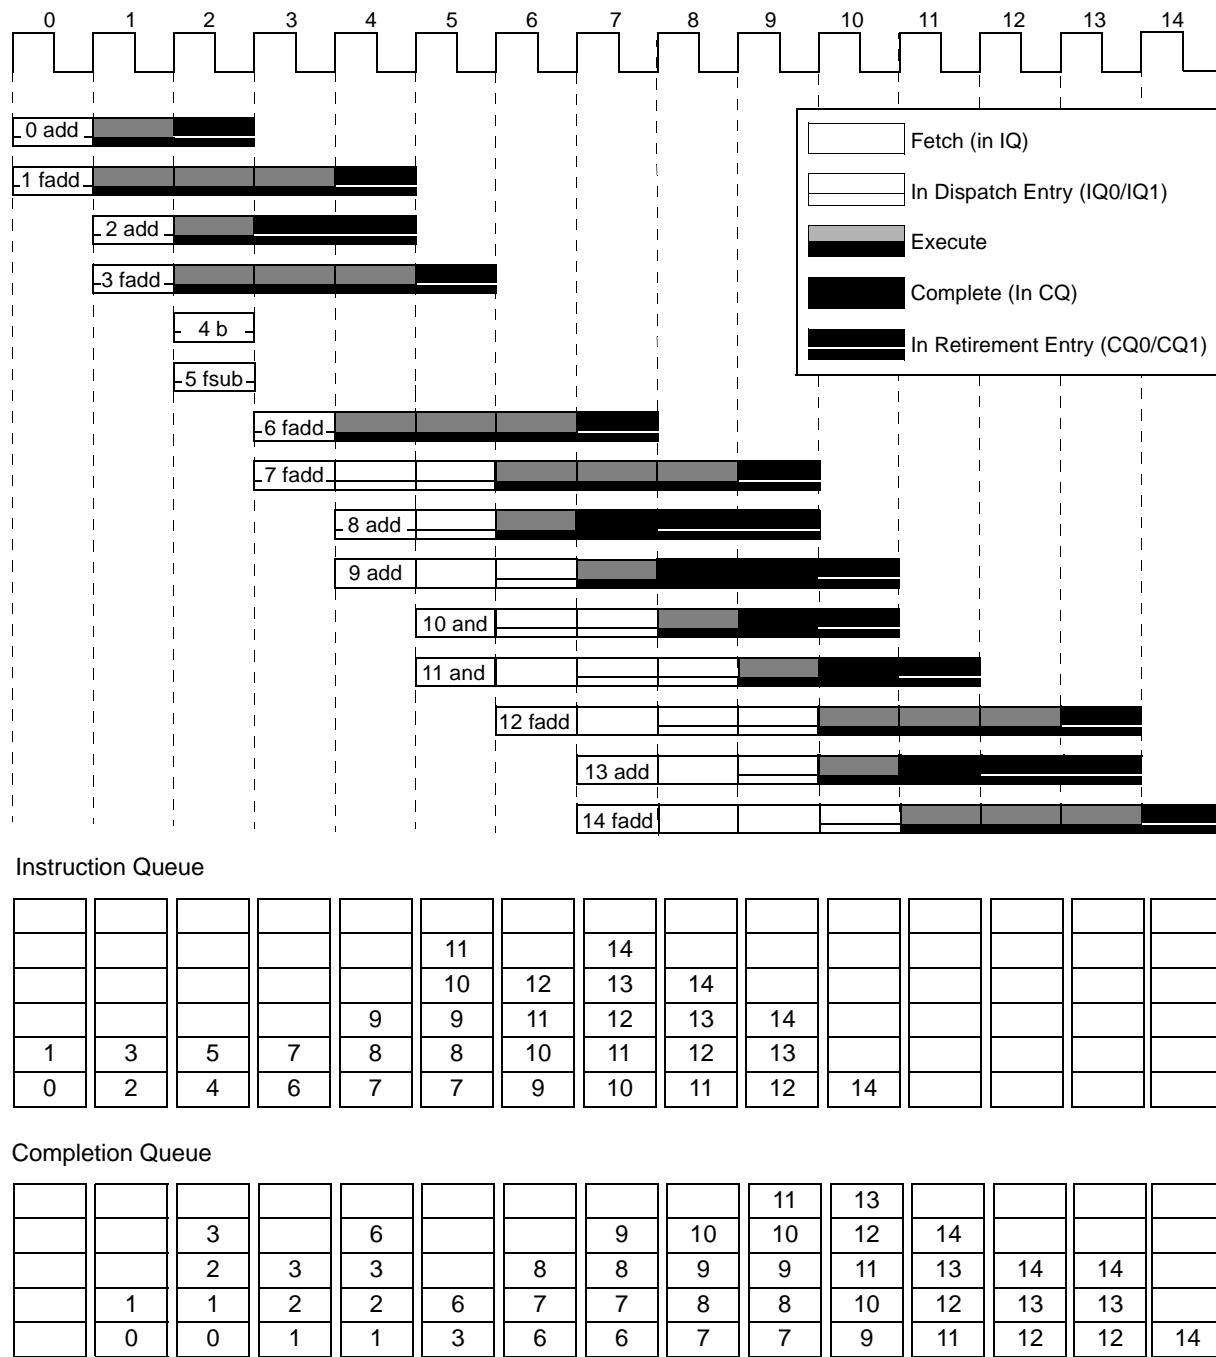
When the fetcher requests instructions from the cache, two things may happen. If the instruction cache is idle and the requested instructions are present, they are provided on the next clock cycle. However, if the instruction cache is busy due to a cache-line-reload operation, instructions cannot be fetched until that operation completes.

6.3.2.2 Cache Hit

An instruction fetch that hits the instruction cache takes only one clock cycle after the request for as many as two instructions to enter the IQ. Note that the cache is not blocked to internal accesses until a cache reload completes (hits under misses). The critical double word is written simultaneously to the cache and forwarded to the requesting unit, minimizing stalls due to load delays.

Figure 6-5 shows a simple example of instruction fetching that hits in the on-chip cache. This example uses a series of integer **add**, **fadd**, and double-precision floating-point add instructions to show how the number of instructions to be fetched is determined, how program order is maintained by the IQ and CQ, how instructions are dispatched and retired in pairs (maximum), and how the FPU pipeline functions. The following instruction sequence is examined:

```
0    add
1    fadd
2    add
3    fadd
```


Figure 6-5. Instruction Timing—Cache Hit

```

4  br 6
5  fsub
6  fadd
7  fadd
8  add
9  add
10 and
11 and
12 fadd
13 add
14 fadd
15 .
16 .
17 .

```

The instruction timing for this example is described cycle-by-cycle as follows:

0. In cycle 0, instructions 0 and 1 are fetched from the instruction cache and are placed in the two entries in the instruction queue (IQ0 and IQ1), where they can be dispatched on the next clock cycle.
1. In cycle 1, instructions 0 and 1 are dispatched to the IU and FPU, respectively. Notice that for instructions to be dispatched, they must be assigned positions in the CQ. In this case, because the CQ is empty, instructions 0 and 1 take the two lowest CQ entries (CQ0 and CQ1). Instructions 2 and 3 are fetched from the instruction cache.
2. At least two IQ positions were available in the IQ in cycle 1, so in cycle 2, instructions 4 and 5 are fetched. Instruction 4 is a branch unconditional instruction that resolves immediately as taken. Because the branch is taken and does not update CTR or LR, it can be folded from the IQ. Instruction 0 completes, writes back its results, and vacates the CQ by the end of the clock cycle. Instruction 1 enters the second FPU execute stage, instruction 2 enters the single-stage IU, and instruction 3 is dispatched into the first FPU stage.
3. In cycle 3, target instructions 6 and 7 are fetched, replacing the folded **b** instruction 4 and instruction 5. Instruction 1 enters the last FPU execute stage, instruction 2 has executed but must remain in the CQ until instruction 1 completes. Note that it can make its results available to subsequent instructions, but cannot be removed from the CQ. Instruction 3 passes into the last FPU execute stage. Note that all three FPU stages are full. To allow for the potential need for denormalization, the dispatch logic prevents instruction 7 (**fadd**) from being dispatched in the next clock cycle.
4. In cycle 4, target instructions 8 and 9 are fetched. Instruction 1 completes in cycle 4, allowing instruction 2, which had finished executing in the previous clock cycle, to be removed from the CQ. Instruction 6 replaces instruction 3 in the first stage of the FPU. Also, as will be shown in cycle 5, a single-cycle stall occurs when the FPU pipeline is full.

5. In cycle 5, instruction 3 completes, instruction 6 continues through the FPU pipeline, and although the first stage of the FPU pipeline is free, instruction 7 cannot be dispatched because of the potential need for one of the previous floating-point instructions to require denormalization. Because instruction 7 cannot be dispatched neither can instruction 8. This dispatch stall causes the instruction queue to become full when instructions 10 and 11 are fetched.
6. In cycle 6, instruction 12 is fetched. Instruction 7 is dispatched to the first FPU stage, so instruction 8 can also be dispatched to the IU. Instructions 9 and 10 move to IQ0 and IQ1, but because instructions 9, 10, and 11 are integer instructions, only one instruction is dispatched in each of the next 2 clock cycles. Note that moving instruction 12 (**fadd**) up further in the program flow would improve dispatch throughput.
7. In cycle 7, instruction 6 completes, instruction 7 is in the second FPU execute stage, and although instruction 8 has executed, it must wait for instruction 7 to complete. Instruction 9 dispatches to the IU. Instructions 10 and 11 move down in the IQ. Fetching resumes with instructions 13 and 14.
8. In cycle 8, instruction 7 is in the third FPU execute stage. Instructions 8 and 9 have executed and they remain in the CQ until instruction 7 completes. Instruction 10 is dispatched to the IU.
9. In cycle 9, instruction 7 completes, allowing instruction 8 to complete. Because the CQ is full, instructions 12 and 13 cannot be dispatched.
10. In cycle 10, instructions 9 and 10 complete. Instruction 11 has executed but cannot exit the CQ from CQ2. Instructions 12 and 13 are dispatched to the FPU and IU, respectively. Instruction 14 drops into IQ0.
11. In cycle 11, instruction 11 completes and instruction 12 is in the second FPU execute stage. Instruction 13 has executed but must remain in the CQ until instruction 12 completes. Instruction 14 enters the first FPU execute stage.

6.3.2.3 Cache Miss

Figure 6-6 shows an instruction fetch that misses the on-chip cache and shows how that fetch affects the instruction dispatch. Note that a processor/bus clock ratio of 1:2 is used. The same instruction sequence is used as in Section 6.3.2.2, “Cache Hit.”

A cache miss extends the latency of the fetch stage, so in this example, the fetch stage represents not only the time the instruction spends in the IQ but also the time required for the instruction to be loaded from system memory, beginning in clock cycle 3.

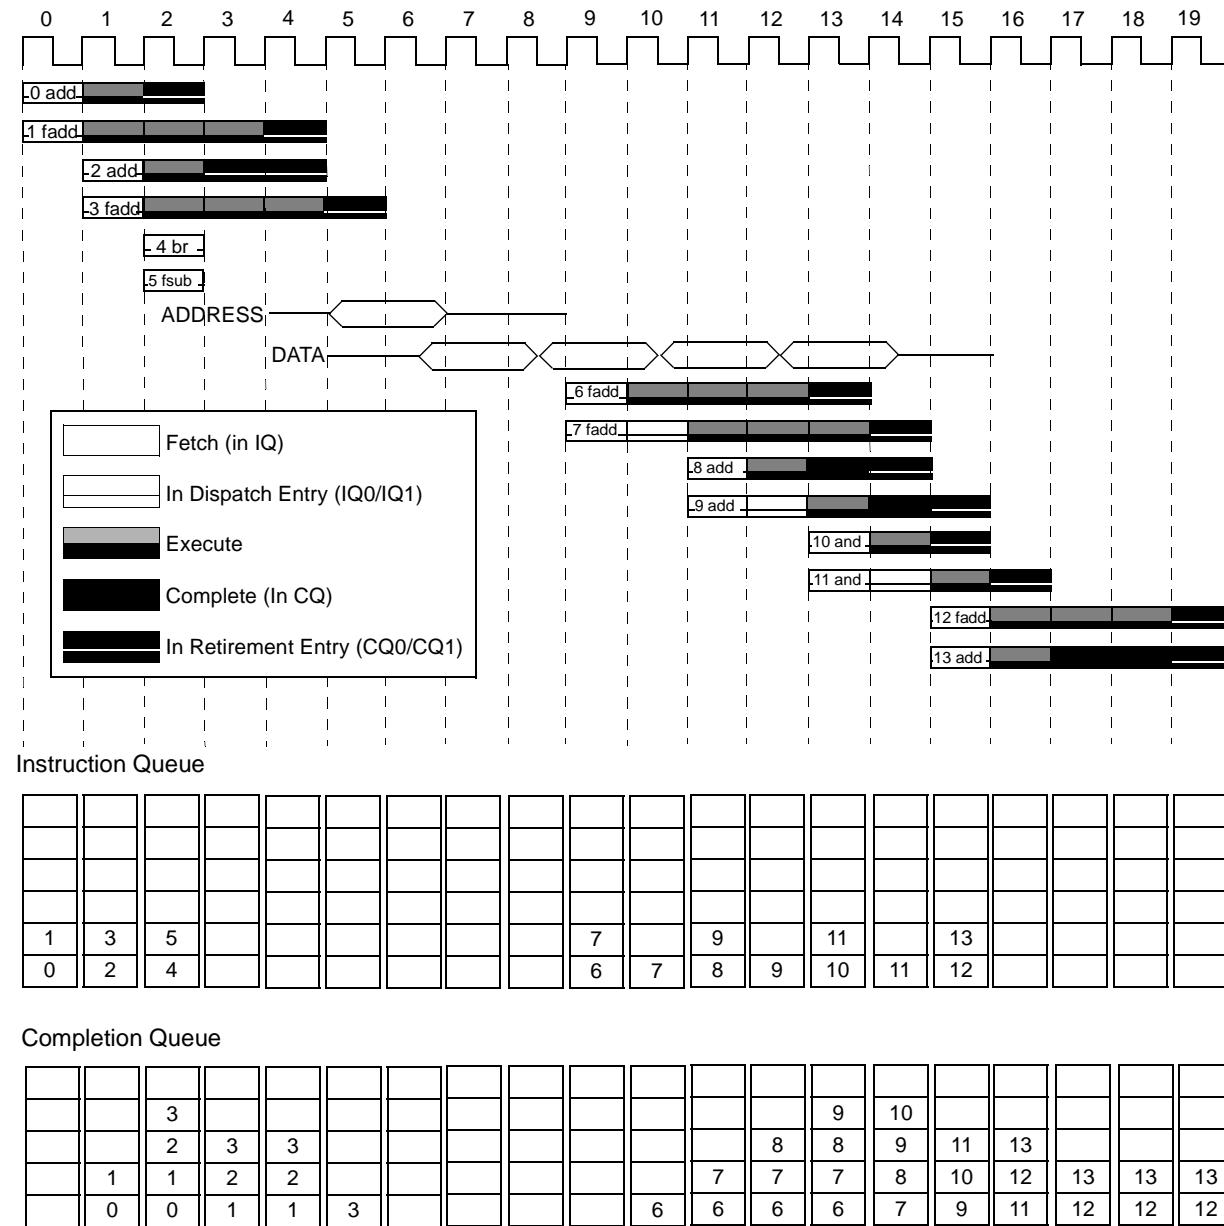


Figure 6-6. Instruction Timing—Cache Miss

During clock cycle 2, the target instruction for the **b** instruction is not in the instruction cache; therefore, a memory access must occur. During clock cycle 5, the address of the block of instructions is sent to the system bus. During clock cycle 9, two instructions (64 bits) are returned from memory on the first beat and are forwarded both to the cache and instruction fetcher.

6.3.3 Instruction Dispatch and Completion Considerations

Several factors affect the ability of the MPC603e to dispatch instructions at a peak rate of two per cycle—the availability of the execution unit, destination rename registers, and completion queue, as well as the handling of completion-serialized instructions. Several of these limiting factors are illustrated in the previous instruction timing examples.

To reduce dispatch unit stalls due to instruction data dependencies, the MPC603e provides a single-entry reservation station for the FPU, SRU, and each IU, and a two-entry reservation station for the LSU. If a data dependency keeps an instruction from starting execution, that instruction is dispatched to the reservation station associated with its execution unit (and the rename registers are assigned), thereby freeing the positions in the instruction queue so instructions can be dispatched to other execution units. Execution begins during the same clock cycle that the rename buffer is updated with the data the instruction is dependent on.

If both instructions in IQ0 and IQ1 require the same execution unit, the instruction in IQ1 cannot be dispatched until the first instruction proceeds through the pipeline and provides the subsequent instruction with a vacancy in the requested execution unit.

The completion unit maintains program order after instructions are dispatched, guaranteeing in-order completion and a precise exception model. Completing an instruction implies committing execution results to the architected destination registers. In-order completion ensures the correct architectural state when the MPC603e must recover from a mispredicted branch or an exception.

The MPC603e can execute instructions out-of-order, but in-order completion by the completion unit ensures a precise exception mechanism. Program-related exceptions are signaled when the instruction causing the exception reaches the last position in the completion queue. Prior instructions are allowed to complete before the exception is taken.

6.3.3.1 Rename Register Operation

To avoid contention for a given register file location, the MPC603e provides rename registers for holding instruction results before the completion commits them to the architected register. There are five GPR rename registers, four FPR rename registers, and one each for the CR, LR, and CTR.

When an instruction dispatches to its execution unit, any required rename registers are allocated for the results of that instruction. If an instruction is dispatched to the reservation station associated with an execution unit due to a data dependency, the dispatcher also provides a tag to the execution unit identifying the rename register that forwards the required data at completion. When the source data reaches the rename register, execution can begin.

Instruction results are transferred from rename registers to architected registers when an instruction is retired from the CQ after any associated exceptions are handled and any predicted branch conditions preceding it in the CQ are resolved. If a branch prediction is incorrect, the instructions following the branch are flushed from the CQ and any results of those instructions are flushed from the rename registers.

6.3.3.2 Instruction Serialization

Although the MPC603e can dispatch and complete two instructions per cycle, serializing instructions can be used to limit dispatch and completion to one instruction per cycle. Serialization falls into three categories—completion, dispatch, and refetch serialization, which are described as follows:

- Completion serialized instructions are held in the execution unit until all prior instructions in the completion unit have been retired. Completion serialization is used for instructions that access or modify a resource for which no rename register exists. Results from these instructions are not available or forwarded for subsequent instructions until the serializing instruction is retired. Instructions that are completion serialized are as follows:
 - Instructions (with the exception of integer add and compare instructions) executed by the system register unit (SRU)
 - Floating-point instructions that access or modify the FPSCR or CR (**mtfsb1**, **mcrfs**, **mtfsfi**, **mffs**, and **mtfsf**)
 - Instructions that manage caches and TLBs
 - Instructions that directly access the GPRs (load and store multiple word and load and store string instructions)
 - Instructions defined by the architecture to have synchronizing behavior
- Dispatch serialized inhibit the dispatching of subsequent instructions until the serializing instruction is retired. Dispatch serialization is used for instructions that access renamed resources used by the dispatcher, and for instructions requiring refetch serialization, including the following:
 - The load multiple instructions, **lmw**, **lswi**, and **lswx**
 - The **mtspr(XER)** and **mcrxr** instructions
 - The synchronizing instructions, **sync**, **isync**, **mtmsr**, **rfi**, and **sc**
- Refetch serialized instructions inhibit dispatching of subsequent instructions and force the refetching of subsequent instructions after the serializing instructions are retired. The context synchronizing instruction, **isync**, is refetch serializing.

6.3.3.3 Execution Unit Considerations

As previously noted, the MPC603e can dispatch and retire two instructions per clock cycle. The peak dispatch rate is affected by the availability of execution units on each clock cycle.

For an instruction to be dispatched, the required execution unit must be available. The dispatcher monitors the availability of all execution units and suspends instruction dispatch if the required execution unit is unavailable. An execution unit may not be available if it can accept and execute only one instruction per cycle or if an execution unit's pipeline becomes full, which may occur if instruction execution takes more clock cycles than the number of pipeline stages in the unit and additional instructions are dispatched to that unit to fill the remaining pipeline stages.

6.4 Execution Unit Timings

The following sections describe instruction timing considerations for each execution unit.

6.4.1 Branch Processing Unit Execution Timing

Flow control operations (conditional branches, unconditional branches, and traps) are typically expensive to execute in most machines because they disrupt normal flow in the instruction stream. When a change in program flow occurs, the IQ must be reloaded with the target instruction stream. During this time the execution units will be idle. However, previously dispatched instructions will continue to execute while the new instruction stream makes its way into the IQ.

Performance features such as branch folding and static branch prediction help minimize penalties associated with flow control operations. The timing for branch instruction execution is determined by many factors including the following:

- Whether the branch requires prediction
- Whether the branch is predicted as taken or not taken
- Whether the branch is taken
- Whether the target instruction stream is in the on-chip cache
- Whether the prediction is correct

6.4.1.1 Branch Folding

When a branch instruction is encountered by the fetcher, the BPU immediately tries to pull that instruction out of the instruction stream and resolve it. When the BPU removes the branch instruction from the stream, the subsequent instruction is shifted down to take the place of the removed branch instruction. This technique is called branch folding. Often, it eliminates the penalties of flow control instructions because instruction execution proceeds as though the branch were never there.

If the folded branch instruction changes program flow (the branch is said to be taken), the BPU immediately requests the instructions at the new target from the on-chip cache. In most cases, the new instructions arrive in the IQ before any bubbles are introduced into the execution units. If the folded branch does not change program flow (the branch is not taken), the branch instruction is already removed and execution continues as if there were never a branch in the original sequence.

When a conditional branch cannot be resolved due to a CR data dependency, the branch is executed by means of static branch prediction and instruction fetching proceeds down the predicted path. If the prediction is incorrect when the branch is resolved, the IQ and all subsequently executed instructions are purged, instructions executed before the predicted branch are allowed to complete, and instruction fetching resumes down the correct path.

There are several situations where instruction sequences create dependencies that prevent a branch instruction from being resolved immediately, thereby causing execution of the subsequent instruction stream based on the predicted outcome of the branch instruction. The instruction sequences, and the resulting action of the branch instruction is described as follows:

- An **mtspr(LR)** followed by a **bclr**—Fetching is stopped and the branch waits for the **mtspr** to execute.
- An **mtspr(CTR)** followed by a **bcctr**—Fetching is stopped and the branch waits for the **mtspr** to execute.
- An **mtspr(CTR)** followed by a **bc(CTR)**—Fetching is stopped and the branch waits for the **mtspr** to execute.
- A **bc(CTR)** followed by another **bc(CTR)**—Fetching is stopped and the second branch waits for the first to be completed.
- A **bc(CTR)** followed by a **bcctr**—Fetching is stopped and the **bcctr** waits for the first to be completed.
- A branch($LK = 1$) followed by a branch($LK = 1$)—Fetching is stopped and the second branch waits for the first to be completed. (Note: A **bl** instruction does not have to wait for a branch($LK = 1$) to complete.)
- A **bc(based-on-CR)** waiting for resolution due to a CR-dependency followed by a **bc(based-on-CR)**—Fetching is stopped and the second branch waits for the first CR-dependency to be resolved. (Note: Branch conditions can be a function of the CTR and CR; if the CTR condition is sufficient to resolve the branch, then a CR-dependency is ignored.)

6.4.1.2 Static Branch Prediction

Static branch prediction allows software (for example, compilers) to give a hint to the machine hardware about the direction the branch is likely to take. When a branch instruction encounters a data dependency, the BPU waits for the required condition code to

become available. Rather than stalling instruction dispatch until the source operand is ready, the MPC603e predicts the likely path and instructions are fetched and executed along that path. When the branch operand becomes available, the branch is evaluated. If the prediction is correct, program flow continues along that path uninterrupted; otherwise, the processor backs up and program flow resumes along the correct path.

If the target address of the branch (link or count register) is modified by an instruction that appears before the branch instruction, the BPU waits until the target address is available.

The MPC603e executes through one level of prediction. The microprocessor may not predict a branch if a prior branch instruction is still unresolved.

The number of instructions that can be executed after branch prediction is limited by the fact that instructions in the predicted stream cannot update the register files or memory until the branch is resolved. That is, instructions may be dispatched and executed, but cannot reach the write-back stage in the completion unit, instead, it stalls in the completion queue. When CQ is full, no more instructions can be dispatched.

In the case of a misprediction, the MPC603e is able to redirect the machine state rather effortlessly because the programming model has not been updated. When a branch is found to be mispredicted, all instructions that were dispatched subsequent to the predicted branch instruction are simply flushed from the completion queue, and their results flushed from the rename registers. No architected register state needs to be restored because no architected register state was modified by the instructions following the unresolved predicted branch.

6.4.1.2.1 Predicted Branch Timing Examples

Figure 6-7 shows how both taken and non-taken branches are handled and how the MPC603e handles both correct and incorrect predictions. The example shows the timing for the following instruction sequence:

0	add
1	add
2	bc
3	mulhw
4	bc T0
5	fadd
6	and
T0	add
T1	add
T2	add
T3	add
T4	and
T5	or

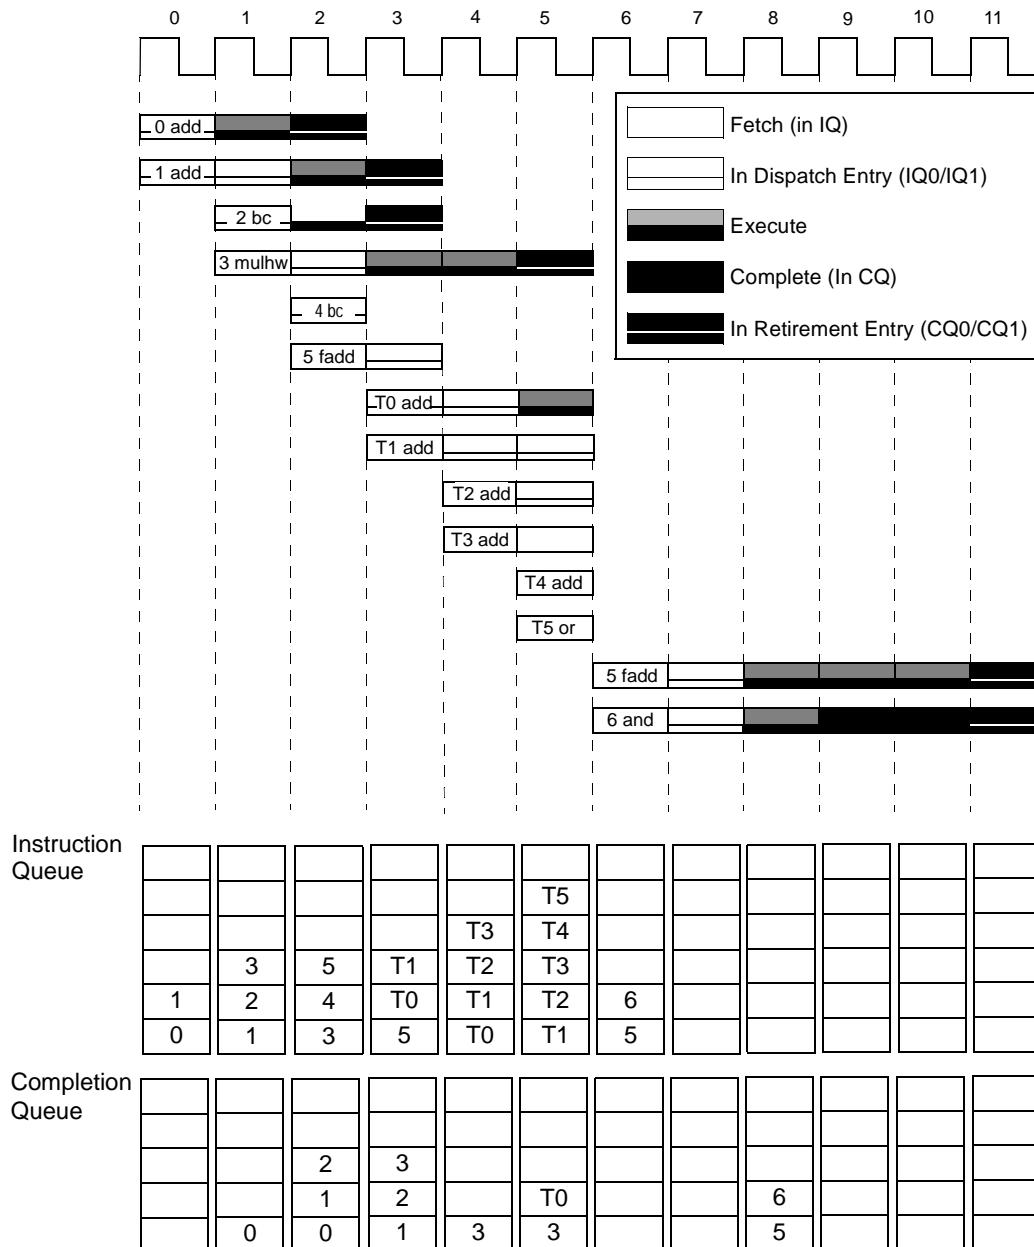


Figure 6-7. Branch Instruction Timing

0. During clock cycle 0, instructions 0 and 1 are dispatched in the beginning of clock cycle 1.
1. In clock cycle 1, instructions 2 and 3 are fetched in the IQ. Instruction 2 is a branch instruction that updates the CTR and instruction 3 is a **mulhw** instruction on which instruction 4 depends. Instruction 0 enters the IU. Instruction 1 has a single-cycle stall.

2. In clock cycle 2, instructions 4 (a second **bc** instruction) and 5 are fetched. The second **bc** instruction is predicted as taken. It can be folded, but it cannot be resolved until instruction 3 writes back. Instruction 0 completes at the end of this cycle. Instruction 1 is dispatched to the IU. Instruction 2 takes entry in the CQ.
3. In clock cycle 3, target instructions T0 and T1 are fetched. Instructions 1 and 2 complete, instruction 4 has been folded, and instruction 5 has been flushed from the IQ. Instruction 3 is assigned to CQ2.
4. In clock cycle 4, target instructions T2 and T3 are fetched. IU instructions T0 and T1 have multiple stalls as one execution possible in a clock cycle. Instruction 3 is assigned to CQ0.
5. In clock cycle 5, instruction 3, on which the second branch instruction depended, writes back and the branch prediction is proven incorrect. Even though T0 is in CQ0, where it could be written back, it is not because the prediction was incorrect. All target instructions are flushed from their positions in the pipeline at the end of this clock cycle, as there are many results in the rename registers.

After one clock cycle required to refetch the original instruction stream, instruction 5, the same instruction that was fetched in clock cycle 2, is brought back into the IQ from the instruction cache, along with one other.

6.4.2 Integer Unit Execution Timing

The integer unit executes all integer and bit-field computational instructions. Many of these instructions execute in a single clock cycle. The integer unit has one execute stage so when a multiple-cycle integer instruction is executed, no other integer instructions can also begin to execute. See Table 6-4 for integer instruction execution timing.

6.4.3 Floating-Point Unit Execution Timing

The FPU on the MPC603e executes all floating-point computational instructions. The LSU performs integer floating-point loads and stores. Execution of most floating-point instructions is pipelined within the FPU, allowing up to three instructions to be executing in the FPU concurrently. While most floating-point instructions execute with 3- or 4-cycle latency, and 1- or 2-cycle throughput, 3 instructions (**fdivs**, **fdiv**, and **fres**) execute with latencies of 18 to 33 cycles. The **fdivs**, **fdiv**, **fres**, **mtfsb0**, **mtfsb1**, **mtfsfi**, **mffs**, and **mtfsf** instructions block the floating-point unit pipeline until they complete execution, and thereby inhibit the dispatch of additional floating-point instructions. With the exception of the **mcrfs** instruction, all floating-point instructions will immediately forward their CR results to the BPU for fast branch resolution without waiting for the instruction to be retired by the completion unit, and the CR updated. See Table 6-5 for floating-point instruction execution timing.

6.4.4 Load/Store Unit Execution Timing

The LSU executes all floating-point and integer loads and stores. It also executes other instructions that address memory. The execution of most load and store instructions is pipelined. The LSU has two pipeline stages; the first is for effective address calculation and MMU translation, and the second is for accessing the physically addressed memory. Load and store instructions have a 2-cycle latency and 1-cycle throughput.

If operands are misaligned, additional latency may be required either for an alignment exception to be taken or for additional bus accesses. Load instructions that miss in the cache prevent subsequent cache accesses during the cache line refill. See Table 6-6 for load and store instruction execution timing.

6.4.5 System Register Unit Execution Timing

Most SRU instructions access or modify nonrenamed registers, or directly access renamed registers. They generally execute in a serial manner. Results from these instructions are not available or forwarded for use by subsequent instructions until the instruction completes and is retired. The SRU can also execute the integer instructions **addi**, **addis**, **add**, **addo**, **cmpi**, **cmp**, **cmpli**, and **cmpl** without serialization and in parallel with another integer instruction. Refer to Section 6.3.3.2, “Instruction Serialization,” for additional information on serializing instructions and Table 6-2, Table 6-3, and Table 6-4 for SRU instruction execution timing.

6.5 Memory Performance Considerations

Due to the MPC603e instruction throughput of three instructions per clock cycle, lack of data bandwidth can become a performance bottleneck. For the MPC603e to approach its potential performance levels, it must be able to read and write data quickly and efficiently. If there are many processors in a system environment, one processor may experience long memory latencies while another bus master (for example, a direct-memory access controller) is using the external bus.

To alleviate this possible contention, the MPC603e provides three memory update modes—copy-back, write-through, and cache-inhibit. Each page of memory is specified to be in one of these modes. If a page is in copy-back mode, data being stored to that page is written only to the on-chip cache. If a page is in write-through mode, writes to that page update the on-chip cache on hits and always update main memory. If a page is cache-inhibited, data in that page will never be stored in the on-chip cache. All three of these modes of operation have advantages and disadvantages. A decision as to which mode to use depends on the system environment as well as the application.

The following sections describe how performance is impacted by each memory update mode. For details about the operation of the on-chip cache and the memory update modes, see Chapter 3, “Instruction and Data Cache Operation.”

6.5.1 Copy-Back Mode

When data is stored in a location marked as copy back, store operations for cacheable data do not necessarily cause an external bus cycle to update memory. Instead, memory updates only occur on modified line replacements, cache flushes, or when another processor attempts to access a specific address for which there is a corresponding modified cache entry. For this reason, copy-back mode may be preferred when external bus bandwidth is a potential bottleneck—for example, in a multiprocessor environment. Copy-back mode is also well suited for data that is closely coupled to a processor, such as local variables.

If more than one device uses data stored in a page marked as copy back, snooping must be enabled to allow copy-back operations and cache invalidations of modified data. The MPC603e implements snooping hardware to prevent other devices from accessing invalid data. When bus snooping is enabled, the processor monitors the transactions of the other devices. For example, if another device accesses a memory location and its memory-coherent (M) bit is set and the MPC603e on-chip cache has a modified value for that address, the processor preempts the bus transaction and updates memory with the cache data. If the cache contents associated with the snooped address are unmodified, the MPC603e invalidates the cache block. The other device can then attempt an access to the updated address. See Chapter 3, “Instruction and Data Cache Operation.”

Copy-back mode provides complete cache/memory coherency as well as maximizing available external bus bandwidth.

6.5.2 Write-Through Mode

Store operations to memory in write-through mode always update memory as well as the on-chip cache (on cache hits). Write-through mode is used when the data in the cache must always agree with external memory (for example, video memory), when shared (global) data may be used frequently, or when allocation of a cache line on a cache miss is undesirable. Automatic copy back of cached data is not performed if that data is from a memory page marked as write-through mode because valid cache data always agrees with memory.

Stores to memory that are in write-through mode may cause a decrease in performance. Each time a store is performed to memory in write-through mode, the bus is busy for the extra clock cycles required to update memory; therefore, load operations that miss the on-chip cache must wait while the external store operation completes.

6.5.3 Cache-Inhibited Accesses

Data for a page marked cache-inhibited cannot be stored in the on-chip cache.

Areas of the memory map can be cache-inhibited by the operating system. If a cache-inhibited access hits in the on-chip cache, the corresponding cache line is

invalidated. If the line is marked modified, it is copied back to memory before being invalidated.

In summary, the copy-back mode allows both load and store operations to use the on-chip cache. The write-through mode allows load operations to use the on-chip cache, but store operations cause a memory access and a cache update if the data is already in the cache. Lastly, the cache-inhibited mode causes memory access for both loads and stores.

6.6 Instruction Scheduling Guidelines

The performance of the MPC603e can be improved by avoiding resource conflicts and promoting parallel utilization of execution units through efficient instruction scheduling. Instruction scheduling on the MPC603e can be improved by observing the following guidelines:

- Implement good static branch prediction (setting of y bit in BO field).
- When branch prediction is uncertain, or an even probability, predict fall through.
- To reduce mispredictions, separate the instruction that sets CR bits from the branch instruction that evaluates them; separation by more than nine instructions ensures that the CR bits will be immediately available for evaluation.
- When branching conditionally to a location specified by count registers (CTRs) or link registers (LRs), or when branching conditionally based on the value in the count register, separate the **mtspr** instruction that initializes the CTR or LR from the branch instruction performing the evaluation. Separation of the branch and **mtspr** instruction by more than nine instructions ensures the register values will be immediately available for use by the branch instruction.
- Schedule instructions such that they can dual dispatch.
- Schedule instructions to minimize stalls when an execution unit is busy.
- Avoid using serializing instructions.
- Schedule instructions to avoid dispatch stalls due to renamed resource limitations.
 - Only five instructions can be in execute-complete stage at any one time.
 - Only five GPR destinations can be in execute-complete-deallocate stage at any one time. Note that load with update address instructions use two destination registers.
 - Only four FPR destinations can be in execute-complete-deallocate stage at any one time.

6.6.1 Branch, Dispatch, and Completion Unit Resource Requirements

This section describes the specific resources required to avoid stalls during branch resolution, instruction dispatching, and instruction completion.

6.6.1.1 Branch Resolution Resource Requirements

The following is a list of branch instructions and the resources required to avoid stalling the fetch unit in the course of branch resolution:

- The **bclr** instruction requires LR availability.
- The **bcctr** instruction requires CTR availability.
- Branch and link instructions require shadow LR availability.
- The branch conditional on counter decrement and CR condition requires CTR availability or the CR condition must be false, and the MPC603e cannot be executing instructions following an unresolved predicted branch when the branch is encountered by the BPU.
- The branch conditional on CR condition cannot be executed following an unresolved predicted branch instruction.

6.6.1.2 Dispatch Unit Resource Requirements

The following is a list of resources required to avoid stalls in the dispatch unit. Note that the two dispatch buffers, IQ0 and IQ1, are at the bottom of the dispatch queue:

- Requirements for dispatching from IQ0 are as follows:
 - Needed execution unit available
 - Needed GPR rename registers available
 - Needed FPR rename registers available
 - Completion queue is not full
 - Instruction is dispatch serialized and completion buffer is empty
 - A dispatch serialized instruction is not currently being executed
- Requirements for dispatching from IQ1 are as follows:
 - Instruction in IQ0 must dispatch
 - Instruction dispatched by IQ0 is not dispatch serialized
 - Needed execution unit is available (after dispatch from IQ0)
 - Needed GPR rename registers are available (after dispatch from IQ0)
 - Needed FPR rename register is available (after dispatch from IQ0)

- Completion queue is not full (after dispatch from IQ0)
- Instruction dispatched from IQ1 is not dispatch serialized

6.6.1.3 Completion Unit Resource Requirements

The following is a list of resources required to avoid stalls in the completion unit; note that the two completion buffers are described as CQ0 and CQ1, where CQ0 is the entry at the end of the queue:

- Requirements for completing an instruction from CQ0 are as follows:
 - Instruction in CQ0 must be finished
 - Instruction in CQ0 must not follow an unresolved predicted branch
 - Instruction in CQ0 must not cause an exception
- Requirements for completing an instruction from CQ1 are as follows:
 - Instruction in CQ0 must complete in same cycle
 - Instruction in CQ1 must be finished
 - Instruction in CQ1 must not follow an unresolved predicted branch
 - Instruction in CQ1 must not cause an exception
 - Instruction in CQ1 must be an integer or load instruction
 - Number of CR updates from both CQ0 and CQ1 must not exceed one
 - Number of GPR updates from both CQ0 and CQ1 must not exceed two
 - Number of FPR updates from both CQ0 and CQ1 must not exceed one

6.7 Instruction Latency Summary

Table 6-1 through Table 6-6 list the latencies associated with each instruction executed by the MPC603e. Note that the instruction latency tables contain no 64-bit architected instructions. These instructions will trap to an illegal instruction exception handler when encountered. Recall that the term latency is defined as the total time it takes to execute an instruction and make ready the results of that instruction.

Table 6-1 provides the latencies for the branch instructions.

Table 6-1. Branch Instructions

Mnemonic	Primary	Extended	Unit	Latency (in Cycles) ¹
bc[l][a]	16	—	BPU	1
b[l][a]	18	—	BPU	1
bclr[l]	19	016	BPU	1
bcctr[l]	19	528	BPU	1

¹ These operations may be folded for an effective cycle time of 0.

Table 6-2 provides the latencies for the system register instructions.

Table 6-2. System Register Instructions

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
sc	17	--1	SRU	3
rfi	19	050	SRU	3
isync	19	150	SRU	1&
mfmsr	31	083	SRU	1
mtmsr	31	146	SRU	2
mtsrr	31	210	SRU	2
mtsrin	31	242	SRU	2
mfspr (not I/DBATs)	31	339	SRU	1
mfspr (DBATs)	31	339	SRU	3&
mfspr (IBATs)	31	339	SRU	3&
mtspr (not IBATs)	31	467	SRU	2 (XER-&)
mtspr (IBATs)	31	467	SRU	2&
mfsr	31	595	SRU	3&
sync	31	598	SRU	1&
mfsrin	31	659	SRU	3&
eieio	31	854	SRU	1
mftb	31	371	SRU	1
mttb	31	467	SRU	1

Note: Cycle times marked with & require a variable number of cycles due to serialization.

Table 6-3 provides the latencies for the condition register logical instructions.

Table 6-3. Condition Register Logical Instructions

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
mcrf	19	000	SRU	1
crnor	19	033	SRU	1
crandc	19	129	SRU	1
crxor	19	193	SRU	1
crnand	19	225	SRU	1
crand	19	257	SRU	1
creqv	19	289	SRU	1
crorc	19	417	SRU	1
cror	19	449	SRU	1
mfcr	31	019	SRU	1

Table 6-3. Condition Register Logical Instructions (continued)

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
mtcrf	31	144	SRU	1
mcrxr	31	512	SRU	1&

Note: Cycle times marked with & require a variable number of cycles due to serialization.

Table 6-4 provides the latencies for the integer instructions.

Table 6-4. Integer Instructions

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
twi	03	—	Integer	2
mulli	07	—	Integer	2,3
subfic	08	—	Integer	1
cmpli	10	—	Integer & SRU	1^
cmpi	11	—	Integer & SRU	1^
addic	12	—	Integer	1
addic.	13	—	Integer	1
addi	14	—	Integer & SRU	1
addis	15	—	Integer & SRU	1
rlwimi[.]	20	—	Integer	1
rlwinm[.]	21	—	Integer	1
rlwnm[.]	23	—	Integer	1
ori	24	—	Integer	1
oris	25	—	Integer	1
xori	26	—	Integer	1
xoris	27	—	Integer	1
andi.	28	—	Integer	1
andis.	29	—	Integer	1
cmp	31	000	Integer & SRU	1^
tw	31	004	Integer	2
subfc[o][.]	31	008	Integer	1
addc[o][.]	31	010	Integer	1
mulhwu[.]	31	011	Integer	2,3,4,5,6
slw[.]	31	024	Integer	1
cntlzw[.]	31	026	Integer	1
and[.]	31	028	Integer	1

Table 6-4. Integer Instructions (continued)

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
cmpl	31	032	Integer & SRU	1^
subf[.]	31	040	Integer	1
andc[.]	31	060	Integer	1
mulhw[.]	31	075	Integer	2,3,4,5
neg[o][.]	31	104	Integer	1
nor[.]	31	124	Integer	1
subfe[o][.]	31	136	Integer	1
adde[o][.]	31	138	Integer	1
subfze[o][.]	31	200	Integer	1
addze[o][.]	31	202	Integer	1
subfme[o][.]	31	232	Integer	1
addme[o][.]	31	234	Integer	1
mull[o][.]	31	235	Integer	2,3,4,5
add[o][.]	31	266	Integer & SRU ¹	1
eqv[.]	31	284	Integer	1
xor[.]	31	316	Integer	1
orc[.]	31	412	Integer	1
or[.]	31	444	Integer	1
divwu[o][.]	31	459	Integer	20
nand[.]	31	476	Integer	1
divw[o][.]	31	491	Integer	20
srw[.]	31	536	Integer	1
sraw[.]	31	792	Integer	1
srawi[.]	31	824	Integer	1
extsh[.]	31	922	Integer	1
extsb[.]	31	954	Integer	1

Note: ^ indicates that the cycle time immediately forwards their CR results to the BPU for fast branch resolution.

¹ The SRU can only execute the **add** and **add[o]** instructions.

Table 6-5 provides the latencies for the floating-point instructions.

Table 6-5. Floating-Point Instructions

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
fdivs[.]	59	018	FPU	18^
fsubs[.]	59	020	FPU	1-1-1^
fadds[.]	59	021	FPU	1-1-1^
fres[.]	59	024	FPU	18^
fmuls[.]	59	025	FPU	1-1-1^
fmsubs[.]	59	028	FPU	1-1-1^
fmaddss[.]	59	029	FPU	1-1-1^
fnmsubs[.]	59	030	FPU	1-1-1^
fnmadds[.]	59	031	FPU	1-1-1^
fcmpu	63	000	FPU	1-1-1^
frsp[.]	63	012	FPU	1-1-1^
fctiw[.]	63	014	FPU	1-1-1^
fctiwz[.]	63	015	FPU	1-1-1^
fdiv[.]	63	018	FPU	33^
fsub[.]	63	020	FPU	1-1-1^
fadd[.]	63	021	FPU	1-1-1^
fsel[.]	63	023	FPU	1-1-1^
fmul[.]	63	025	FPU	2-1-1^
frsqrtte[.]	63	026	FPU	1-1-1^
fmsub[.]	63	028	FPU	2-1-1^
fmadd[.]	63	029	FPU	2-1-1^
fnmsub[.]	63	030	FPU	2-1-1^
fnmadd[.]	63	031	FPU	2-1-1^
fcmwo	63	032	FPU	1-1-1^
mtfsb1[.]	63	038	FPU	1-1-1&^
fneg[.]	63	040	FPU	1-1-1^
mcrfs	63	064	FPU	1-1-1&
mtfsb0[.]	63	070	FPU	1-1-1&^
fmr[.]	63	072	FPU	1-1-1^
mtfsfil[.]	63	134	FPU	1-1-1&^
fnabs[.]	63	136	FPU	1-1-1^
fabs[.]	63	264	FPU	1-1-1^
mffs[.]	63	583	FPU	1-1-1&^

Table 6-5. Floating-Point Instructions (continued)

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
mtfsf[.]	63	711	FPU	1-1-1&^

Notes: Cycle times marked with & require a variable number of cycles due to completion serialization.

Cycle times marked with ^ immediately forward their CR results to the BPU for fast branch resolution.

Cycle times marked with a - specify the number of clock cycles in each pipeline stage. Instructions with a single entry in the cycles column are not pipelined.

Table 6-6 provides latencies for the load and store instructions.

Table 6-6. Load and Store Instructions

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
lwarx	31	020	LSU	2:1
lwzx	31	023	LSU	2:1
dcbst	31	054	LSU	2/5&
lwzux	31	055	LSU	2:1
dcbf	31	086	LSU	2/5&
lbzx	31	087	LSU	2:1
lbzux	31	119	LSU	2:1
stwcx.	31	150	LSU	8
stwx	31	151	LSU	2:1
stwux	31	183	LSU	2:1
stbx	31	215	LSU	2:1
dcbtst	31	246	LSU	2
stbux	31	247	LSU	2:1
dcbt	31	278	LSU	2
lhzx	31	279	LSU	2:1
tlbie	31	306	LSU	3&
eciwx	31	310	LSU	2:1
lhzux	31	311	LSU	2:1
lhax	31	343	LSU	2:1
lhaux	31	375	LSU	2:1
sthx	31	407	LSU	2:1
ecowx	31	438	LSU	2:1
sthux	31	439	LSU	2:1
dcbi	31	470	LSU	2&
lswx	31	533	LSU	2 + n&

Chapter 6. Instruction Timing

Table 6-6. Load and Store Instructions (continued)

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
lwbrx	31	534	LSU	2:1
lfsx	31	535	LSU	2:1
tlbsync	31	566	LSU	2&
lfsux	31	567	LSU	2:1
lswi	31	597	LSU	2 + n&
lfdx	31	599	LSU	2:1
lfdux	31	631	LSU	2:1
stswx	31	661	LSU	1 + n&
stwbrx	31	662	LSU	2:1
stfsx	31	663	LSU	2:1
stfsux	31	695	LSU	2:1
stswi	31	725	LSU	1 + n&
stfdx	31	727	LSU	2:1
stfdux	31	759	LSU	2:1
lhbrx	31	790	LSU	2:1
sthbrx	31	918	LSU	2:1
tlbld	31	978	LSU	2&
icbi	31	982	LSU	3&
stfiwx	31	983	LSU	2:1
tlbli	31	1010	LSU	3&
dcbz	31	1014	LSU	10&
lwz	32	—	LSU	2:1
lwzu	33	—	LSU	2:1
lbz	34	—	LSU	2:1
lbzu	35	—	LSU	2:1
stw	36	—	LSU	2:1
stwu	37	—	LSU	2:1
stb	38	—	LSU	2:1
stbu	39	—	LSU	2:1
lhz	40	—	LSU	2:1
lhzu	41	—	LSU	2:1
lha	42	—	LSU	2:1
lhau	43	—	LSU	2:1
sth	44	—	LSU	2:1
sthу	45	—	LSU	2:1

Table 6-6. Load and Store Instructions (continued)

Mnemonic	Primary	Extended	Unit	Latency (in Cycles)
lmw	46	—	LSU	$2 + n\&$
stmw	47	—	LSU	$1 + n\&$
lfs	48	—	LSU	2:1
lfsu	49	—	LSU	2:1
lfd	50	—	LSU	2:1
lfdū	51	—	LSU	2:1
stfs	52	—	LSU	2:1
stfsu	53	—	LSU	2:1
stfd	54	—	LSU	2:1
stfdu	55	—	LSU	2:1

Notes: Cycle times marked with & require a variable number of cycles due to serialization.

Cycle times marked with a / specify hit and miss times for cache management instructions that require conditional bus activity.

Cycle times marked with a : specify cycles of total latency and throughput.

Load and store multiple and string instruction cycles are shown as a fixed number of cycles plus a variable number of cycles where n is the number of words accessed by the instruction.



Chapter 7

Signal Descriptions

This chapter describes the MPC603e microprocessor's external signals. It contains a concise description of individual signals, showing behavior when the signal is asserted and negated and when the signal is an input and an output.

NOTE

A bar over a signal name indicates that the signal is active low—for example, $\overline{\text{ARTRY}}$ (address retry) and $\overline{\text{TS}}$ (transfer start). Active-low signals are referred to as asserted (active) when they are low and negated when they are high. Signals that are not active-low, such as AP[0:3] (address bus parity signals) and TT[0:4] (transfer type signals) are referred to as asserted when they are high and negated when they are low.

The MPC603e signals are grouped as follows:

- Address arbitration signals—The MPC603e uses these signals to arbitrate for address bus mastership.
- Address transfer start signals—These signals indicate that a bus master has begun a transaction on the address bus.
- Address transfer signals—These signals, consisting of the address bus, address parity, and address parity error signals, are used to transfer the address and to ensure the integrity of the transfer.
- Transfer attribute signals—These signals provide information about the type of transfer, such as the transfer size and whether the transaction is bursted, write-through, or cache-inhibited.
- Address transfer termination signals—These signals are used to acknowledge the end of the address phase of the transaction. They also indicate whether a condition exists that requires the address phase to be repeated.
- Data arbitration signals—The MPC603e uses these signals to arbitrate for data bus mastership.
- Data transfer signals—These signals, consisting of the data bus, data parity, and data parity error signals, are used to transfer the data and to ensure the integrity of the transfer.

- Data transfer termination signals—Data termination signals are required after each data beat in a data transfer. In a single-beat transaction, the data termination signals also indicate the end of the tenure. In burst accesses, the data termination signals apply to individual beats and indicate the end of the tenure only after the final data beat. They also indicate whether a condition exists that requires the data phase to be repeated.
- System status signals—These signals include the external interrupt signal, checkstop signals, and both soft- and hard-reset signals. These signals are used to interrupt and, under various conditions, to reset the processor.
- JTAG/COP interface signals—The JTAG (IEEE 1149.1) interface and common on-chip processor (COP) unit provides a serial interface to the system for performing monitoring and boundary tests.
- Processor status—These signals include the memory reservation, machine quiesce control, time base enable, and TLBISYNC signals.
- Clock signals—These signals provide for system clock input and frequency control.

7.1 Signal Configuration

Figure 7-1 shows how the MPC603e microprocessor's signals are grouped.

NOTE

A pinout showing actual pin numbers is included in the *MPC603e RISC Microprocessor Hardware Specifications*.

7.2 Signal Descriptions

This section describes individual MPC603e signals, grouped according to Figure 7-1. Note that the following sections are intended to provide a quick summary of signal functions. Chapter 8, “System Interface Operation,” describes many of these signals in greater detail, both with respect to how individual signals function and how groups of signals interact.

7.2.1 Address Bus Arbitration Signals

The address arbitration signals are a collection of input and output signals the MPC603e uses to request the address bus, recognize when the request is granted, and indicate to other devices when mastership is granted. For a detailed description of how these signals interact, see Section 8.3.1, “Address Bus Arbitration.”

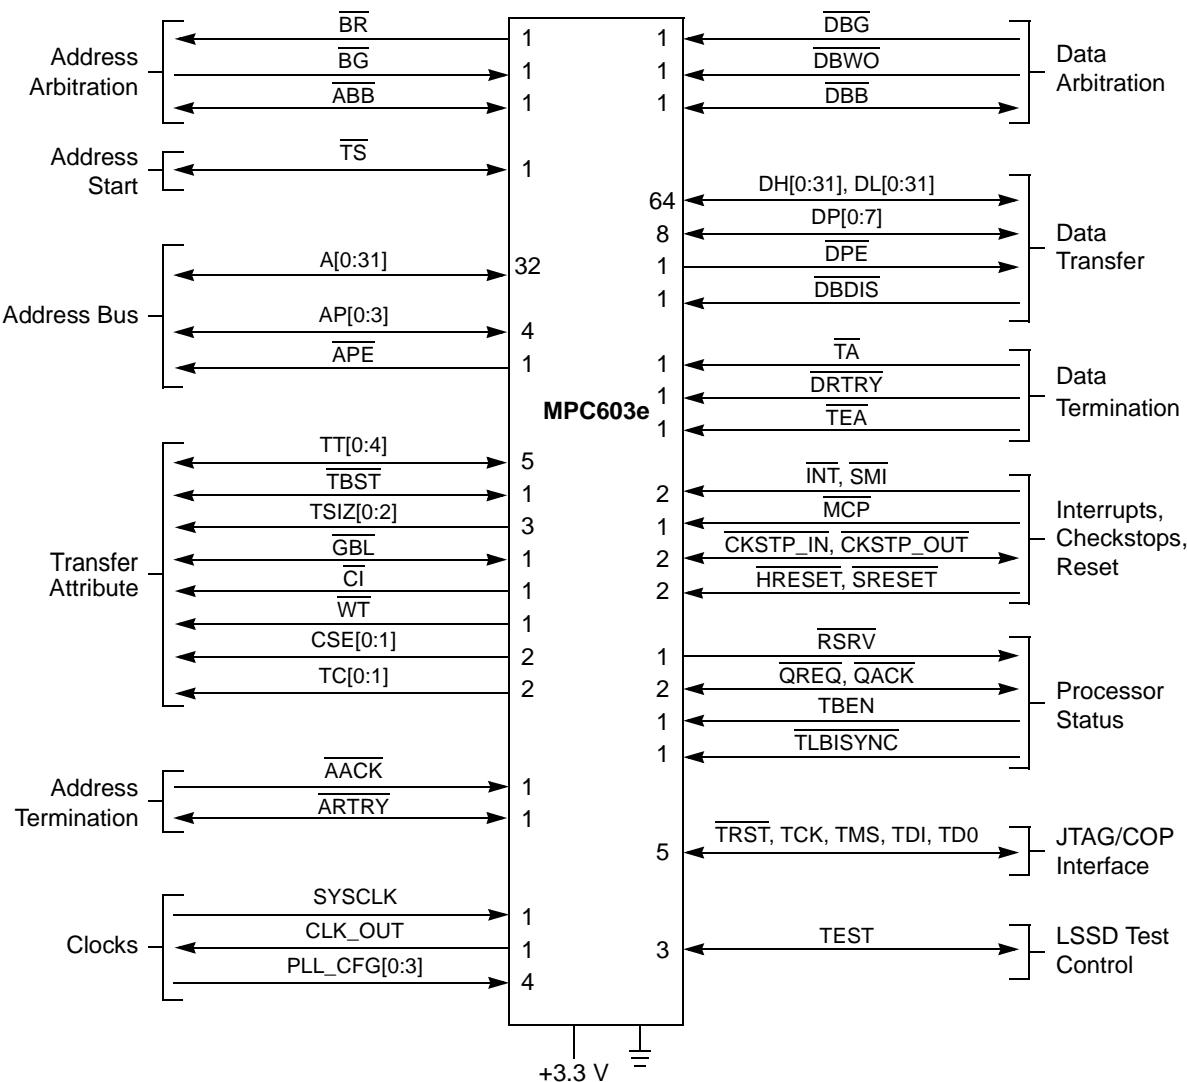


Figure 7-1. Signal Groups

7.2.1.1 Bus Request (**BR**)—Output

The **BR** signal is an output signal on the MPC603e. Following are the state meaning and timing comments for the **BR** output.

State Meaning

Asserted—Indicates that the MPC603e is requesting mastership of the address bus. Note that **BR** may be asserted for one or more cycles, and then negated due to an internal cancellation of the bus request such as a load hit in the touch load buffer. See Section 8.3.1, “Address Bus Arbitration.”

Negated—Indicates that the MPC603e is not requesting the address bus. The MPC603e may have no bus operation pending, it may be

parked, or the ARTRY input was asserted on the previous bus clock cycle.

Timing Comments

Assertion—Occurs when the MPC603e is not parked and a bus transaction is needed. This may occur even if the two possible pipeline accesses have occurred. BR is also asserted for one cycle during the execution of a **dcbz** instruction, and during the execution of a load instruction that hits in the touch load buffer.

Negation—Occurs for at least one bus clock cycle after an accepted, qualified bus grant (see BG and ABB), even if another transaction is pending. It is also negated for at least one bus clock cycle when the assertion of ARTRY is detected on the bus.

7.2.1.2 Bus Grant (BG)—Input

The BG signal is an input signal on the MPC603e. Following are the state meaning and timing comments for the BG input.

State Meaning

Asserted—Indicates that the MPC603e may, with the proper qualification, assume mastership of the address bus. A qualified bus grant occurs when BG is asserted and ABB and ARTRY (after AACK) are not asserted. ABB and ARTRY are driven by the MPC603e or other bus masters. If the MPC603e is parked, BR need not be asserted for the qualified bus grant. See Section 8.3.1, “Address Bus Arbitration.”

Negated—Indicates that the MPC603e is not the next potential address bus master.

Timing Comments

Assertion—May occur at any time to indicate the MPC603e is free to use the address bus. After the MPC603e assumes bus mastership, it does not check for a qualified bus grant again until the cycle during which the address bus tenure is completed (assuming it has another transaction to run). The MPC603e does not accept a BG in the cycles between the assertion of any TS and AACK.

Negation—May occur anytime to indicate the MPC603e cannot use the bus. The MPC603e may still assume bus mastership on the bus clock cycle of the negation of BG because during the previous cycle BG indicated to the MPC603e that it was free to take mastership (if qualified).

7.2.1.3 Address Bus Busy (ABB)

The ABB signal is both an input and output signal.

7.2.1.3.1 Address Bus Busy (\overline{ABB})—Output

Following are the state meaning and timing comments for the \overline{ABB} output.

State Meaning	Asserted—Indicates that the MPC603e is the address bus master. See Section 8.3.1, “Address Bus Arbitration.”
	Negated—Indicates that the MPC603e is not using the address bus. If \overline{ABB} is negated during the bus clock cycle following a qualified bus grant, the MPC603e did not accept mastership, even if \overline{BR} was asserted. This can occur if a potential transaction is aborted internally before the transaction is started.
Timing Comments	Assertion—Occurs on the bus clock cycle following a qualified \overline{BG} that is accepted by the processor (see Negated). Negation—Occurs for a minimum of one-half bus clock cycle following the assertion of \overline{AACK} . If \overline{ABB} is negated during the bus clock cycle following a qualified bus grant, the MPC603e did not accept mastership, even if \overline{BR} was asserted. High Impedance—Occurs after \overline{ABB} is negated.

7.2.1.3.2 Address Bus Busy (\overline{ABB})—Input

Following are the state meaning and timing comments for the \overline{ABB} input.

State Meaning	Asserted—Indicates that the address bus is in use. This condition effectively blocks the MPC603e from assuming address bus ownership, regardless of the \overline{BG} input; see Section 8.3.1, “Address Bus Arbitration.”
	Negated—Indicates that the address bus is not owned by another bus master and that it is available to the MPC603e when accompanied by a qualified bus grant.
Timing Comments	Assertion—May occur when the MPC603e must be prevented from using the address bus (and the processor is not currently asserting \overline{ABB}). Negation—May occur whenever the MPC603e can use the address bus.

7.2.2 Address Transfer Start Signals

Address transfer start signals are input and output signals that indicate that an address bus transfer has begun. The transfer start (\overline{TS}) signal identifies the operation as a memory transaction.

For detailed information about how \overline{TS} interacts with other signals, refer to Section 8.3.2, “Address Transfer.”

7.2.2.1 Transfer Start (\overline{TS})

The \overline{TS} signal is both an input and output signal on the MPC603e.

7.2.2.1.1 Transfer Start (\overline{TS})—Output

Following are the state meaning and timing comments for the \overline{TS} output.

State Meaning

Asserted—Indicates that the MPC603e has begun a memory bus transaction and that the address bus and transfer attribute signals are valid. When asserted with the appropriate TT[0:4] signals, it is also an implied data bus request for a memory transaction (unless it is an address-only operation).

Negated—Indicates that no bus transaction is occurring during normal operation.

Timing Comments

Assertion—C coinides with the assertion of \overline{ABB} .

Negation—Occurs one bus clock cycle after \overline{TS} is asserted.

High Impedance—C coinides with the negation of \overline{ABB} .

7.2.2.1.2 Transfer Start (\overline{TS})—Input

Following are the state meaning and timing comments for the \overline{TS} input.

State Meaning

Asserted—Indicates that another master has begun a bus transaction and that the address bus and transfer attribute signals are valid for snooping (see \overline{GBL}).

Negated—Indicates that no bus transaction is occurring.

Timing Comments

Assertion—May occur during the assertion of \overline{ABB} .

Negation—Must occur one bus clock cycle after \overline{TS} is asserted.

7.2.3 Address Transfer Signals

The address transfer signals are used to transmit the address and to generate and monitor parity for the address transfer. For a detailed description of how these signals interact, refer to Section 8.3.2, “Address Transfer.”

7.2.3.1 Address Bus (A[0:31])

The address bus (A[0:31]) consists of 32 signals that are both input and output signals.

7.2.3.1.1 Address Bus (A[0:31])—Output

Following are the state meaning and timing comments for the A[0:31] outputs.

State Meaning	Asserted/Negated—Represents the physical address (real address in the architecture specification) of the data to be transferred. On burst transfers, the address bus presents the double-word-aligned address containing the critical code or data that missed the cache on a read operation, or the first double word of the cache line on a write operation. Note that the address output during burst operations is not incremented. See Section 8.3.2, “Address Transfer.”
Timing Comments	Assertion/Negation—Occurs on the bus clock cycle after a qualified bus grant (coincides with assertion of \overline{ABB} and \overline{TS}). High Impedance—Occurs one bus clock cycle after \overline{AACK} is asserted.

7.2.3.1.2 Address Bus (A[0:31])—Input

Following are the state meaning and timing comments for the A[0:31] inputs.

State Meaning	Asserted/Negated—Represents the physical address of a snoop operation.
Timing Comments	Assertion/Negation—Must occur on the same bus clock cycle as the assertion of \overline{TS} ; is sampled by MPC603e only on this cycle.

7.2.3.2 Address Bus Parity (AP[0:3])

The AP[0:3] signals are both input and output signals reflecting 1 bit of odd-byte parity for each of the 4 bytes of address when a valid address is on the bus.

7.2.3.2.1 Address Bus Parity (AP[0:3])—Output

Following are the state meaning and timing comments for the AP[0:3] outputs.

State Meaning	Asserted/Negated—Represents odd parity for each of 4 bytes of the physical address for a transaction. Odd parity means that an odd number of bits, including the parity bit, are driven high. The signal assignments correspond to the following:
AP0	A[0:7]
AP1	A[8:15]
AP2	A[16:23]
AP3	A[24:31]

For more information, see Section 8.3.2.1, “Address Bus Parity.”

- Timing Comments** Assertion/Negation—The same as A[0:31].
High Impedance—The same as A[0:31].

7.2.3.2.2 Address Bus Parity (AP[0:3])—Input

Following are the state meaning and timing comments for the AP[0:3] inputs.

- State Meaning** Asserted/Negated—Represents odd parity for each of 4 bytes of the physical address for snooping operations. Detected even parity causes the processor to take a machine check exception or enter the checkstop state if address parity checking is enabled in the HID0 register; see Section 2.1.2.1, “Hardware Implementation Registers (HID0 and HID1).” (See also the $\overline{\text{APE}}$ signal description.)
- Timing Comments** Assertion/Negation—The same as A[0:31].

7.2.3.3 Address Parity Error ($\overline{\text{APE}}$)—Output

The $\overline{\text{APE}}$ signal is an output signal on the MPC603e. Note that $\overline{\text{APE}}$ is an open-drain type output, and requires an external pull-up resistor (for example, 10 k Ω to V_{DD}) to assure proper negation of $\overline{\text{APE}}$. Following are the state meaning and timing comments for the $\overline{\text{APE}}$ signal on the MPC603e. The $\overline{\text{APE}}$ signal will not be asserted if address parity checking is disabled (HID0[EBA] cleared to 0). For more information, see Section 8.3.2.1, “Address Bus Parity.”

- State Meaning** Asserted—Indicates incorrect address bus parity has been detected by the MPC603e on a snoop ($\overline{\text{GBL}}$ asserted).
Negated—Indicates that the MPC603e has not detected a parity error (even parity) on the address bus.
- Timing Comments** Assertion—Occurs on the second bus clock cycle after $\overline{\text{TS}}$ is asserted.
High Impedance—Occurs on the third bus clock cycle after $\overline{\text{TS}}$ is asserted.

7.2.4 Address Transfer Attribute Signals

The transfer attribute signals are a set of signals that further characterize the transfer—such as the size of the transfer, whether it is a read or write operation, and whether it is a burst or single-beat transfer. For a detailed description of how these signals interact, see Section 8.3.2, “Address Transfer.”

7.2.4.1 Transfer Type (TT[0:4])

The TT[0:4] signals consist of five input/output signals on the MPC603e. For a complete description of TT[0:4] signals and for transfer type encodings, see Table 7-1.

7.2.4.1.1 Transfer Type (TT[0:4])—Output

Following are the state meaning and timing comments for the TT[0:4] outputs.

State Meaning Asserted/Negated—Indicates the type of transfer in progress.

Timing Comments Assertion/Negation/High Impedance—The same as A[0:31].

7.2.4.1.2 Transfer Type (TT[0:4])—Input

Following are the state meaning and timing comments for the TT[0:3] inputs.

State Meaning Asserted/Negated—Indicates the type of transfer in progress (see Table 7-2).

Timing Comments Assertion/Negation—The same as A[0:31].

Table 7-1 describes the transfer encodings for a MPC603e bus master.

Table 7-1. Transfer Encoding for the Bus Master

MPC603e Bus Master Transaction	Transaction Source	TT0	TT1	TT2	TT3	TT4	60x Bus Specification Command	Transaction
N/A	N/A	0	0	0	0	0	Clean block	Address only
N/A	N/A	0	0	1	0	0	Flush block	Address only
N/A	N/A	0	1	0	0	0	sync	Address only
Address only	dcbz	0	1	1	0	0	Kill block	Address only
N/A	N/A	1	0	0	0	0	eieio	Address only
Single-beat write (nonGBL)	ecowx	1	0	1	0	0	External control word write	Single-beat write
N/A	N/A	1	1	0	0	0	TLB invalidate	Address only
Single-beat read (nonGBL)	eciwx	1	1	1	0	0	External control word read	Single-beat read
N/A	N/A	0	0	0	0	1	lwarx Reservation set	Address only
N/A	N/A	0	0	1	0	1	Reserved	—
N/A	N/A	0	1	0	0	1	tlbsync	Address only
N/A	N/A	0	1	1	0	1	icbi	Address only
N/A	N/A	1	X	X	0	1	Reserved	—

Table 7-1. Transfer Encoding for the Bus Master (continued)

MPC603e Bus Master Transaction	Transaction Source	TT0	TT1	TT2	TT3	TT4	60x Bus Specification Command	Transaction
Single-beat write	Caching-inhibited or write-through store	0	0	0	1	0	Write-with-flush	Single-beat write or burst
Burst (nonGBL)	Cast out, or snoop copy back	0	0	1	1	0	Write-with-kill	Single-beat write or burst
Single-beat read	Caching-inhibited load or instruction fetch	0	1	0	1	0	Read	Single-beat read or burst
Burst	Load miss, store miss, or instruction fetch	0	1	1	1	0	Read-with-intent-to-modify	Burst
Single-beat write	stwcx.	1	0	0	1	0	Write-with-flush-atomic	Single-beat write
N/A	N/A	1	0	1	1	0	Reserved	N/A
Single-beat read	lwarx (caching-inhibited load)	1	1	0	1	0	Read-atomic	Single-beat read or burst
Burst	lwarx (load miss)	1	1	1	1	0	Read-with-intent-to-modify-atomic	Burst
N/A	N/A	0	0	0	1	1	Reserved	—
N/A	N/A	0	0	1	1	1	Reserved	—
N/A	N/A	0	1	0	1	1	Read-with-no-intent-to-cache	Single-beat read or burst
N/A	N/A	0	1	1	1	1	Reserved	—
N/A	N/A	1	X	X	1	1	Reserved	—

Table 7-2 describes the 60x bus specification transfer encodings and the MPC603e bus snoop response on an address hit.

Table 7-2. Snoop Hit Response

60x Bus Specification Command	Transaction	TT0	TT1	TT2	TT3	TT4	MPC603e Bus Snooper; Action on Hit
Clean block	Address only	0	0	0	0	0	N/A
Flush block	Address only	0	0	1	0	0	N/A
sync	Address only	0	1	0	0	0	N/A
Kill block	Address only	0	1	1	0	0	Kill, cancel reservation
eieio	Address only	1	0	0	0	0	N/A

Table 7-2. Snoop Hit Response (continued)

60x Bus Specification Command	Transaction	TT0	TT1	TT2	TT3	TT4	MPC603e Bus Snooper; Action on Hit
External control word write	Single-beat write	1	0	1	0	0	N/A
TLB Invalidate	Address only	1	1	0	0	0	N/A
External control word read	Single-beat read	1	1	1	0	0	N/A
Iwarx reservation set	Address only	0	0	0	0	1	N/A
Reserved	—	0	0	1	0	1	N/A
tlbsync	Address only	0	1	0	0	1	N/A
icbi	Address only	0	1	1	0	1	N/A
Reserved	—	1	X	X	0	1	N/A
Write-with-flush	Single-beat write or burst	0	0	0	1	0	Flush, cancel reservation
Write-with-kill	Single-beat write or burst	0	0	1	1	0	Kill, cancel reservation
Read	Single-beat read or burst	0	1	0	1	0	Clean or flush
Read-with-intent-to-modify	Burst	0	1	1	1	0	Flush
Write-with-flush-atomic	Single-beat write	1	0	0	1	0	Flush, cancel reservation
Reserved	N/A	1	0	1	1	0	N/A
Read-atomic	Single-beat read or burst	1	1	0	1	0	Clean or flush
Read-with-intent-to-modify-atomic	Burst	1	1	1	1	0	Flush
Reserved	—	0	0	0	1	1	N/A
Reserved	—	0	0	1	1	1	N/A
Read-with-no-intent-to-cache	Single-beat read or burst	0	1	0	1	1	Clean
Reserved	—	0	1	1	1	1	N/A
Reserved	—	1	X	X	1	1	N/A

When HID0[ABE] is set, the PID7t-603e performs address-only bus transactions with the encodings shown in Table 7-3.

Table 7-3. Implementation-Specific Transfer Encoding

TT0	TT1	TT2	TT3	TT4	PID7t-603e Transaction	Transaction	Transaction Source
0	0	0	0	0	Clean block	Address only	dcbst
0	0	1	0	0	Flush block	Address only	dcbf
0	1	1	0	0	Kill block	Address only	dcbz, dcbi

7.2.4.2 Transfer Size (TSIZ[0:2])—Output

The TSIZ[0:2] signals consist of three output signals on the MPC603e. Following are the state meaning and timing comments for the TSIZ[0:2] outputs.

State Meaning

Asserted/Negated—For memory accesses, these signals along with $\overline{\text{TBST}}$, indicate the data transfer size for the current bus operation, as shown in Table 7-4. Table 8-5 shows how the transfer size signals are used with the address signals for aligned transfers. Table 8-6 shows how the transfer size signals are used with the address signals for misaligned transfers.

For external control instructions (**eciwx** and **ecowx**), TSIZ[0:2] are used to output bits 29–31 of the external access register (EAR), which are used to form the resource ID ($\overline{\text{TBST}}||\text{TSIZ}[0:2]$).

Timing Comments

Assertion/Negation—The same as A[0:31].

High Impedance—The same as A[0:31].

Table 7-4. Data Transfer Size

TBST	TSIZ[0:2]	Transfer Size
Asserted	010	Burst (32 bytes)
Negated	000	8 bytes
Negated	001	1 byte
Negated	010	2 bytes
Negated	011	3 bytes
Negated	100	4 bytes
Negated	101	5 bytes
Negated	110	6 bytes
Negated	111	7 bytes

7.2.4.3 Transfer Burst ($\overline{\text{TBST}}$)

The $\overline{\text{TBST}}$ signal is an input/output signal.

7.2.4.3.1 Transfer Burst ($\overline{\text{TBST}}$)—Output

Following are the state meaning and timing comments for the $\overline{\text{TBST}}$ output.

State Meaning

Asserted—Indicates that a burst transfer is in progress.

Negated—Indicates that a burst transfer is not in progress.

For external control instructions (**eciwx** and **ecowx**), $\overline{\text{TBST}}$ is used to output EAR[28], which is used to form the resource ID ($\overline{\text{TBST}}||\text{TSIZ}[0:2]$).

- Timing Comments** Assertion/Negation—The same as A[0:31].
High Impedance—The same as A[0:31].

7.2.4.3.2 Transfer Burst ($\overline{\text{TBST}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{TBST}}$ input.

- State Meaning** Asserted/Negated—Used when snooping for single-beat reads (read with no intent to cache).
- Timing Comments** Assertion/Negation—The same as A[0:31].

7.2.4.4 Transfer Code (TC[0:1])—Output

The TC[0:1] consists of two output signals on the MPC603e. Following are the state meaning and timing comments for the TC[0:1] outputs.

- State Meaning** Asserted/Negated—Represents a special encoding for the transfer in progress (see Table 7-5).
- Timing Comments** Assertion/Negation—The same as A[0:31].
High Impedance—The same as A[0:31].

Table 7-5. Encodings for TC[0:1] Signals

TC(0:1)	Read	Write
0 0	Data transaction	Any write
0 1	Touch load	—
1 0	Instruction fetch	—
1 1	Reserved	—

7.2.4.5 Cache Inhibit ($\overline{\text{CI}}$)—Output

Following are the state meaning and timing comments for the $\overline{\text{CI}}$ output.

- State Meaning** Asserted—Indicates that a single-beat transfer is not cached, reflecting the setting of the I bit for the block or page that contains the address of the current transaction.
Negated—Indicates that a burst transfer allocates a line in the MPC603e data cache.
- Timing Comments** Assertion/Negation—The same as A[0:31].
High Impedance—The same as A[0:31].

7.2.4.6 Write-Through (\overline{WT})—Output

Following are the state meaning and timing comments for the \overline{WT} output.

State Meaning Asserted—Indicates that a single-beat transaction is write-through, reflecting the value of the W bit for the block or page that contains the address of the current transaction.

Negated—Indicates that a transaction is not write-through.

Timing Comments Assertion/Negation—The same as A[0:31].

High Impedance—The same as A[0:31].

7.2.4.7 Global (\overline{GBL})

The \overline{GBL} signal is an input/output signal on the MPC603e.

7.2.4.7.1 Global (\overline{GBL})—Output

Following are the state meaning and timing comments for the \overline{GBL} output.

State Meaning Asserted—Indicates that a transaction is global, reflecting the setting of the M bit for the block or page that contains the address of the current transaction (except in the case of copy-back operations and instruction fetches, which are nonglobal).

Negated—Indicates that a transaction is not global.

Timing Comments Assertion/Negation—The same as A[0:31].

High Impedance—The same as A[0:31].

7.2.4.7.2 Global (\overline{GBL})—Input

Following are the state meaning and timing comments for the \overline{GBL} input.

State Meaning Asserted—Indicates that a transaction must be snooped by the MPC603e.

Negated—Indicates that a transaction is not snooped by the MPC603e.

Timing Comments Assertion/Negation—The same as A[0:31].

7.2.4.8 Cache Set Entry (CSE[0:1])—Output

Following are the state meaning and timing comments for the CSE[0:1] outputs.

State Meaning Asserted/Negated—Represents the cache replacement set element for the current transaction reloading into or writing out of the cache. Can be used with the address bus and the transfer attribute signals to externally track the state of each cache line in the MPC603e cache.

Note that CSE[0:1] are not meaningful during data cache touch load operations.

Timing Comments Assertion/Negation—The same as A[0:31].

High Impedance—The same as A[0:31].

7.2.5 Address Transfer Termination Signals

The address transfer termination signals are used to indicate either that the address phase of the transaction has completed successfully or must be repeated, and when it should be terminated. For detailed information about how these signals interact, see Section 8.3.3, “Address Transfer Termination.”

7.2.5.1 Address Acknowledge (AACK)—Input

Following are the state meaning and timing comments for the AACK input.

State Meaning Asserted—Indicates that the address phase of a transaction is complete. The address bus will go to a high-impedance state on the next bus clock cycle. The MPC603e samples ARTRY on the bus clock cycle following the assertion of AACK.

Negated—(During ABB) indicates that the address bus and transfer attributes must remain driven.

Timing Comments Assertion—May occur as early as the bus clock cycle after TS is asserted (unless MPC603e is configured for 1:1 or 1.5:1 clock modes, when AACK can be asserted no sooner than the second cycle following the assertion of TS—one address wait state); assertion can be delayed to allow adequate address access time for slow devices. For example, if an implementation supports slow snooping devices, an external arbiter can postpone the assertion of AACK.

Negation—Must occur one bus clock cycle after the assertion of AACK.

7.2.5.2 Address Retry (ARTRY)

The ARTRY signal is both an input and output signal on the MPC603e.

7.2.5.2.1 Address Retry (ARTRY)—Output

Following are the state meaning and timing comments for the ARTRY output.

State Meaning Asserted—Indicates that the MPC603e detects a condition in which a snooped address tenure must be retried. If the MPC603e needs to update memory as a result of the snoop that caused the retry, the

MPC603e asserts \overline{BR} the second cycle after \overline{AACK} if \overline{ARTRY} is asserted.

High Impedance—Indicates that the MPC603e does not need the snooped address tenure to be retried.

Timing Comments Assertion—Asserted the third bus cycle following the assertion of \overline{TS} if a retry is required.

Negation—Occurs the second bus cycle after the assertion of \overline{AACK} . Since this signal may be simultaneously driven by multiple devices, it negates in a unique fashion. First the buffer goes to high impedance for a minimum of one-half processor cycle (dependent on the clock mode), then it is driven negated for one bus cycle before returning to high impedance.

This special method of negation may be disabled by setting precharge disable in HID0.

7.2.5.2.2 Address Retry (\overline{ARTRY})—Input

Following are the state meaning and timing comments for the \overline{ARTRY} input.

State Meaning Asserted—If the MPC603e is the address bus master, \overline{ARTRY} indicates that the MPC603e must retry the preceding address tenure and immediately negate \overline{BR} (if asserted). If the associated data tenure has already started, the MPC603e will also abort the data tenure immediately, even if the burst data has been received. If the MPC603e is not the address bus master, this input indicates that the MPC603e should immediately negate \overline{BR} for one bus clock cycle following the assertion of \overline{ARTRY} by the snooping bus master to allow an opportunity for a copy-back operation to main memory. Note that the subsequent address presented on the address bus may not be the same one associated with the assertion of \overline{ARTRY} .

Negated/High Impedance—Indicates that the MPC603e does not need to retry the last address tenure.

Timing Comments Assertion—May occur as early as the second cycle following the assertion of \overline{TS} , and must occur by the bus clock cycle immediately following the assertion of \overline{AACK} if an address retry is required.

Negation—Must occur during the second cycle after the assertion of \overline{AACK} .

7.2.6 Data Bus Arbitration Signals

Like the address bus arbitration signals, data bus arbitration signals maintain an orderly process for determining data bus mastership. Note that there is no data bus arbitration signal

equivalent to the address bus arbitration signal \overline{BR} (bus request), because, except for address-only transactions, \overline{TS} implies data bus requests. For a detailed description on how these signals interact, see Section 8.4.1, “Data Bus Arbitration.”

One special signal, \overline{DBWO} , allows the MPC603e to be configured dynamically to write data out of order with respect to read data. For detailed information about using \overline{DBWO} , see Section 8.10, “Using \overline{DBWO} .”

7.2.6.1 Data Bus Grant (\overline{DBG})—Input

Following are the state meaning and timing comments for the \overline{DBG} input.

State Meaning	Asserted—Indicates that the MPC603e may, with the proper qualification, assume mastership of the data bus. The MPC603e derives a qualified data bus grant when \overline{DBG} is asserted and \overline{DBB} , \overline{DRTRY} , and \overline{ARTRY} are negated; that is, the data bus is not busy (\overline{DBB} is negated), there is no outstanding attempt to retry the current data tenure (\overline{DRTRY} is negated), and there is no outstanding attempt to perform an \overline{ARTRY} of the associated address tenure. Negated—Indicates that the MPC603e must hold off its data tenures.
Timing Comments	Assertion—May occur any time to indicate the MPC603e is free to take data bus mastership. It is not sampled until \overline{TS} is asserted. Negation—May occur at any time to indicate the MPC603e cannot assume data bus mastership.

7.2.6.2 Data Bus Write Only (\overline{DBWO})—Input

Following are the state meaning and timing comments for the \overline{DBWO} input.

State Meaning	Asserted—Indicates that the MPC603e may run the data bus tenure for an outstanding write address even if a read address is pipelined before the write address. Refer to Section 8.10, “Using \overline{DBWO} ,” for detailed instructions for using \overline{DBWO} . Negated—Indicates that the MPC603e must run the data bus tenures in the same order as the address tenures.
Timing Comments	Assertion—Must occur no later than a qualified \overline{DBG} for an outstanding write tenure. \overline{DBWO} is only recognized by the MPC603e on the clock of a qualified \overline{DBG} . If no write requests are pending, the MPC603e ignores \overline{DBWO} and assumes data bus ownership for the next pending read request. Negation—May occur any time after a qualified \overline{DBG} and before the next assertion of \overline{DBG} .

7.2.6.3 Data Bus Busy (\overline{DBB})

The \overline{DBB} signal is both an input and output signal on the MPC603e.

7.2.6.3.1 Data Bus Busy (\overline{DBB})—Output

Following are the state meaning and timing comments for the \overline{DBB} output.

State Meaning Asserted—Indicates that the MPC603e is the data bus master. The MPC603e always assumes data bus mastership if it needs the data bus and is given a qualified data bus grant (see \overline{DBG}).

Negated—Indicates that the MPC603e is not using the data bus.

Timing Comments Assertion—Occurs during the bus clock cycle following a qualified \overline{DBG} .

Negation—Occurs for a minimum of one-half bus clock cycle (dependent on clock mode) following the assertion of the final \overline{TA} .

High Impedance—Occurs after \overline{DBB} is negated.

7.2.6.3.2 Data Bus Busy (\overline{DBB})—Input

Following are the state meaning and timing comments for the \overline{DBB} input.

State Meaning Asserted—Indicates that another device is the bus master.

Negated—Indicates that the data bus is free (with proper qualification, see \overline{DBG}) for use by the MPC603e.

Timing Comments Assertion—Must occur when the MPC603e must be prevented from using the data bus.

Negation—May occur whenever the data bus is available.

7.2.7 Data Transfer Signals

Like the address transfer signals, the data transfer signals are used to transmit data and to generate and monitor parity for the data transfer. For a detailed description of how the data transfer signals interact, see Section 8.4.3, “Data Transfer.”

7.2.7.1 Data Bus (DH[0:31], DL[0:31])

The data bus (DH[0:31] and DL[0:31]) consists of 64 signals that are both input and output on the MPC603e. Following are the state meaning and timing comments for DH and DL.

State Meaning The data bus has two halves—data bus high (DH) and data bus low (DL). See Table 7-6 for the data bus lane assignments.

Timing Comments The data bus is driven once for noncached transactions and four times for cache transactions (bursts).

Table 7-6. Data Bus Lane Assignments

Data Bus Signals	Byte Lane
DH[0:7]	0
DH[8:15]	1
DH[16:23]	2
DH[24:31]	3
DL[0:7]	4
DL[8:15]	5
DL[16:23]	6
DL[24:31]	7

7.2.7.1.1 Data Bus (DH[0:31], DL[0:31])—Output

Following are the state meaning and timing comments for the DH and DL outputs.

State Meaning Asserted/Negated—Represents the state of data during a data write. Byte lanes not selected for data transfer will not supply valid data.

Timing Comments Assertion/Negation—Initial beat coincides with \overline{DBB} and, for bursts, transitions on the bus clock cycle following each assertion of \overline{TA} .

High Impedance—Occurs on the bus clock cycle after the final assertion of \overline{TA} .

7.2.7.1.2 Data Bus (DH[0:31], DL[0:31])—Input

Following are the state meaning and timing comments for the DH and DL inputs.

State Meaning Asserted/Negated—Represents the state of data during a data read transaction.

Timing Comments Assertion/Negation—Data must be valid on the same bus clock cycle that \overline{TA} is asserted.

7.2.7.2 Data Bus Parity (DP[0:7])

The eight DP[0:7] signals on the MPC603e are both output and input signals.

7.2.7.2.1 Data Bus Parity (DP[0:7])—Output

Following are the state meaning and timing comments for the DP outputs.

State Meaning Asserted/Negated—Represents odd parity for each of 8 bytes of data write transactions. Odd parity means that an odd number of bits,

including the parity bit, are driven high. The signal assignments are listed in Table 7-7.

- Timing Comments**
- Assertion/Negation—The same as DL[0:31].
 - High Impedance—The same as DL[0:31].

Table 7-7. DP[0:7] Signal Assignments

Signal Name	Signal Assignments
DP0	DH[0:7]
DP1	DH[8:15]
DP2	DH[16:23]
DP3	DH[24:31]
DP4	DL[0:7]
DP5	DL[8:15]
DP6	DL[16:23]
DP7	DL[24:31]

7.2.7.2.2 Data Bus Parity (DP[0:7])—Input

Following are the state meaning and timing comments for the DP inputs.

- State Meaning**
- Asserted/Negated—Represents odd parity for each byte of read data. Parity is checked on all data byte lanes, regardless of the size of the transfer. Detected even parity causes a checkstop if data parity errors are enabled in the HID0 register. (See \overline{DPE} .)

- Timing Comments**
- Assertion/Negation—The same as DL[0:31].

7.2.7.3 Data Parity Error (\overline{DPE})—Output

The \overline{DPE} signal is an output signal (output-only) on the MPC603e. Note that \overline{DPE} is an open-drain type output, and requires an external pull-up resistor (for example, 10 k Ω to V_{DD}) to assure proper negation of \overline{DPE} . Following are the state meaning and timing comments for the \overline{DPE} output.

- State Meaning**
- Asserted—Indicates incorrect data bus parity.
 - Negated—Indicates correct data bus parity.
- Timing Comments**
- Assertion—Occurs on the second bus clock cycle after \overline{TA} is asserted to the MPC603e, unless \overline{TA} is canceled by an assertion of \overline{DRTRY} .
 - High Impedance—Occurs on the third bus clock cycle after \overline{TA} is asserted to the MPC603e.

7.2.7.4 Data Bus Disable (DBDIS)—Input

The DBDIS signal is an input signal (input-only) on the MPC603e. Following are the state meaning and timing comments for the DBDIS input.

State Meaning	Asserted—Indicates (for a write transaction) that the MPC603e must release the data bus and the data bus parity to high impedance during the following cycle. The data tenure remains active, <u>DBB</u> remains driven, and the transfer termination signals are still monitored by the MPC603e. Negated—Indicates the data bus should remain normally driven. <u>DBDIS</u> is ignored during read transactions.
Timing Comments	Assertion/Negation—May be asserted on any clock cycle when the MPC603e is driving, or will be driving the data bus; may remain asserted multiple cycles.

7.2.8 Data Transfer Termination Signals

Data termination signals are required after each data beat in a data transfer. Note that in a single-beat transaction, the data termination signals also indicate the end of the tenure. While in burst accesses, the data termination signals apply to individual beats and indicate the end of the tenure only after the final data beat.

For a detailed description of how these signals interact, see Section 8.4.4, “Data Transfer Termination.”

7.2.8.1 Transfer Acknowledge (TA)—Input

Following are the state meaning and timing comments for the TA input.

State Meaning	Asserted—Indicates that a single-beat data transfer completed successfully or that a data beat in a burst transfer completed successfully (unless <u>DRTRY</u> is asserted on the next bus clock cycle). Note that <u>TA</u> must be asserted for each data beat in a burst transaction, and must be asserted during assertion of <u>DRTRY</u> . For more information, see Section 8.4.4, “Data Transfer Termination.” Negated—(During <u>DBB</u>) indicates that, until <u>TA</u> is asserted, the MPC603e must continue to drive the data for the current write or must wait to sample the data for reads.
Timing Comments	Assertion—Must not occur before <u>AACK</u> for the current transaction (if the address retry mechanism is to be used to prevent invalid data from being used by the processor); otherwise, assertion may occur at any time during the assertion of <u>DBB</u> . The system can withhold

assertion of $\overline{\text{TA}}$ to indicate that the MPC603e should insert wait states to extend the duration of the data beat.

Negation—Must occur after the bus clock cycle of the final (or only) data beat of the transfer. For a burst transfer, the system can assert $\overline{\text{TA}}$ for one bus clock cycle and then negate it to advance the burst transfer to the next beat and insert wait states during the next beat. (Note: When the MPC603e is configured for 1:1 clock mode and is performing a burst read into the data cache, the MPC603e requires one wait state between the assertion of $\overline{\text{TS}}$ and the first assertion of $\overline{\text{TA}}$ for that transaction. If no-DRTRY mode is also selected, the MPC603e requires two wait states for 1:1 clock mode, or one wait state for 1.5:1 clock mode.)

7.2.8.2 Data Retry ($\overline{\text{DRTRY}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{DRTRY}}$ input.

State Meaning	Asserted—Indicates that the MPC603e must invalidate the data from the previous read operation. Negated—Indicates that data presented with $\overline{\text{TA}}$ on the previous read operation is valid. Note that $\overline{\text{DRTRY}}$ is ignored for write transactions.
Timing Comments	Assertion—Must occur during the bus clock cycle immediately after $\overline{\text{TA}}$ is asserted if a retry is required. $\overline{\text{DRTRY}}$ may be held asserted for multiple bus clock cycles. When $\overline{\text{DRTRY}}$ is negated, data must have been valid on the previous clock with $\overline{\text{TA}}$ asserted. Negation—Must occur during the bus clock cycle after a valid data beat. This may occur several cycles after $\overline{\text{DBB}}$ is negated, effectively extending the data bus tenure. Start-Up— $\overline{\text{DRTRY}}$ is sampled at the negation of $\overline{\text{HRESET}}$; if $\overline{\text{DRTRY}}$ is asserted, no-DRTRY mode is selected. If $\overline{\text{DRTRY}}$ is negated at start-up, $\overline{\text{DRTRY}}$ is enabled.

7.2.8.3 Transfer Error Acknowledge ($\overline{\text{TEA}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{TEA}}$ input.

State Meaning	Asserted—Indicates that a bus error occurred. Causes a machine check exception (and possibly causes the processor to enter checkstop state if machine check enable bit is cleared (MSR[ME] = 0)). For more information, see Section 4.5.2.2, “Checkstop State (MSR[ME] = 0).” Assertion terminates the current transaction; that is, assertion of $\overline{\text{TA}}$ and $\overline{\text{DRTRY}}$ are ignored. The
----------------------	---

assertion of $\overline{\text{TEA}}$ causes the negation/high impedance of $\overline{\text{DBB}}$ in the next clock cycle. However, data entering the GPR or the cache are not invalidated.

Negated—Indicates that no bus error was detected.

Timing Comments Assertion—May be asserted while $\overline{\text{DBB}}$ is asserted, and the cycle after $\overline{\text{TA}}$ during a read operation. $\overline{\text{TEA}}$ should be asserted for one cycle only.

Negation— $\overline{\text{TEA}}$ must be negated no later than the negation of $\overline{\text{DBB}}$.

7.2.9 System Status Signals

Most system status signals are input signals that indicate when exceptions are received, when checkstop conditions have occurred, and when the MPC603e must be reset. The MPC603e generates the output signal $\overline{\text{CKSTP_OUT}}$ when it detects a checkstop condition. For a detailed description of these signals, see Section 8.7, “Interrupt, Checkstop, and Reset Signals.”

7.2.9.1 Interrupt ($\overline{\text{INT}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{INT}}$ input.

State Meaning Asserted—The MPC603e initiates an interrupt if MSR[EE] is set; otherwise, the MPC603e ignores the interrupt. To guarantee that the MPC603e takes the external interrupt, $\overline{\text{INT}}$ must be held asserted until the MPC603e takes the interrupt.

Negated—Indicates that normal operation should proceed. See Section 8.7.1, “External Interrupts.”

Timing Comments Assertion—May occur at any time and may be asserted asynchronously to the input clocks. The $\overline{\text{INT}}$ input is level-sensitive. Negation—Should not occur until interrupt is taken.

7.2.9.2 System Management Interrupt ($\overline{\text{SMI}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{SMI}}$ input.

State Meaning Asserted—The MPC603e initiates a system management interrupt exception if MSR[EE] is set; otherwise, the MPC603e ignores the exception condition. The system must hold $\overline{\text{SMI}}$ asserted until the exception is taken.

Negated—Indicates that normal operation should proceed. See Section 8.7.1, “External Interrupts.”

Timing Comments	Assertion—May occur at any time and may be asserted asynchronously to the input clocks. The $\overline{\text{SMI}}$ input is level-sensitive. Negation—Should not occur until exception is taken.
------------------------	--

7.2.9.3 Machine Check Interrupt ($\overline{\text{MCP}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{MCP}}$ input.

State Meaning	Asserted—The MPC603e initiates a machine check interrupt exception if MSR[ME] and HID0[EMCP] are set; if MSR[ME] is cleared and HID0[EMCP] is set, the MPC603e terminates operation by internally gating off all clocks, and releasing all outputs (except $\overline{\text{CKSTP_OUT}}$) to the high-impedance state. If HID0[EMCP] is cleared, the MPC603e ignores the interrupt condition. $\overline{\text{MCP}}$ must be held asserted for two bus clock cycles. Negated—Indicates that normal operation should proceed. See Section 8.7.1, “External Interrupts.”
Timing Comments	Assertion—May occur at any time and may be asserted asynchronously to the input clocks. $\overline{\text{MCP}}$ is negative edge-sensitive. Negation—May be negated two bus cycles after assertion.

7.2.9.4 Checkstop Input ($\overline{\text{CKSTP_IN}}$)—Input

Following are the state meaning and timing comments for the $\overline{\text{CKSTP_IN}}$ input.

State Meaning	Asserted—Indicates that the MPC603e must terminate operation by internally gating off all clocks, and release all outputs (except $\overline{\text{CKSTP_OUT}}$) to the high-impedance state. Once $\overline{\text{CKSTP_IN}}$ is asserted, it must remain asserted until the system has been reset. Negated—Indicates that normal operation should proceed. See Section 8.7.2, “Checkstops.”
Timing Comments	Assertion—May occur at any time and may be asserted asynchronously to the input clocks. Negation—May occur anytime after $\overline{\text{CKSTP_OUT}}$ is asserted.

7.2.9.5 Checkstop Output ($\overline{\text{CKSTP_OUT}}$)—Output

The $\overline{\text{CKSTP_OUT}}$ signal is output only on the MPC603e. Note that $\overline{\text{CKSTP_OUT}}$ is an open-drain type output, and requires an external pull-up resistor (for example, 10 k Ω to V_{DD}) to assure proper negation of $\overline{\text{CKSTP_OUT}}$. Following are the state meaning and timing comments for the $\overline{\text{CKSTP_OUT}}$ output.

State Meaning	Asserted—Indicates that the MPC603e has detected a checkstop condition and has ceased operation. Negated—Indicates that the MPC603e is operating normally. See Section 8.7.2, “Checkstops.”
Timing Comments	Assertion—May occur at any time and may be asserted asynchronously to the input clocks. Negation—Is negated upon assertion of <u>HRESET</u> .

7.2.9.6 Reset Signals

There are two reset signals on the MPC603e—hard reset (HRESET) and soft reset (SRESET). Descriptions of the reset signals are as follows.

7.2.9.6.1 Hard Reset (HRESET)—Input

The HRESET input must be used at power-on to properly reset the processor. Following are the state meaning and timing comments for the HRESET input.

State Meaning	Asserted—Initiates a complete hard reset operation when this input transitions from asserted to negated. Causes a reset exception as described in Section 4.5.1.1, “Hard Reset and Power-On Reset.” Output drivers are released to high impedance within five clocks after the assertion of <u>HRESET</u> . Negated—Indicates that normal operation should proceed. See Section 8.7.3, “Reset Inputs.”
----------------------	---

Timing Comments	Assertion—May occur at any time and may be asserted asynchronously to the MPC603e input clock; must be held asserted for a minimum of 255 clock cycles after the PLL lock time has been met. Refer to the appropriate hardware specifications for further timing comments. Negation—May occur any time after the minimum reset pulse width has been met.
------------------------	---

This input has additional functionality in certain test modes.

7.2.9.6.2 Soft Reset (SRESET)—Input

The SRESET signal is input only. Following are the state meaning and timing comments for the SRESET input.

State Meaning	Asserted—Initiates processing for a reset exception as described in Section 4.5.1.2, “Soft Reset.”
----------------------	--

	Negated—Indicates that normal operation should proceed. See Section 8.7.3, “Reset Inputs.”
Timing Comments	Assertion—May occur at any time and may be asserted asynchronously to the MPC603e input clock. <u>SRESET</u> is negative edge-sensitive. Negation—May be negated two bus cycles after assertion.
	This input has additional functionality in certain test modes.

7.2.9.7 Processor Status Signals

Processor status signals indicate the state of the processor. This includes the memory reservation, machine quiesce control, time base enable, and TLBISYNC signals.

7.2.9.7.1 Quiescent Request (QREQ)

Following are the state meaning and timing comments for the QREQ signal.

State Meaning	Asserted—Indicates that the MPC603e is requesting all bus activity normally required to be snooped to terminate or to pause so the MPC603e may enter a quiescent (low-power) state. Once the MPC603e enters a quiescent state, it no longer snoops bus activity. Negated—Indicates that the MPC603e is not making a request to enter the quiescent state.
Timing Comments	Assertion/Negation—May occur on any cycle. <u>QREQ</u> remains asserted for the duration of the quiescent state.

7.2.9.7.2 Quiescent Acknowledge (QACK)

Following are the state meaning and timing comments for the QACK input.

State Meaning	Asserted—Indicates that all bus activity that requires snooping has terminated or paused, and that the MPC603e may enter the quiescent (or low-power) state. Negated—Indicates that the MPC603e may not enter a quiescent state and must continue snooping the bus.
Timing Comments	Assertion/Negation—May occur on any cycle following the assertion of <u>QREQ</u> , and must be held asserted for a minimum of one bus clock cycle. Start-Up— <u>QACK</u> is sampled at the negation of <u>HRESET</u> to select reduced-pinout mode; if <u>QACK</u> is asserted at start-up, reduced-pinout mode is disabled.

7.2.9.7.3 Reservation (\overline{RSRV})—Output

Following are the state meaning and timing comments for the \overline{RSRV} output.

State Meaning	Asserted/Negated—Represents the state of the reservation coherency bit in the reservation address register that is used by the lwarx and stwcx. instructions. See Section 8.8.1, “Support for the lwarx/stwcx. Instruction Pair.”
Timing Comments	Assertion/Negation—Occurs synchronously with respect to bus clock cycles. The execution of an lwarx instruction sets the internal reservation condition.

7.2.9.7.4 Time Base Enable (TBEN)—Input

Following are the state meanings and timing comments for the TBEN input.

State Meaning	Asserted—Indicates that the time base should continue clocking. This input is essentially a count enable control for the time base counter. Negated—Indicates the time base should stop clocking.
Timing Comments	Assertion/Negation—May occur on any cycle.

7.2.9.7.5 TLBI Sync ($\overline{TLBISYNC}$)

Following are the state meaning and timing comments for the $\overline{TLBISYNC}$ input.

State Meaning	Asserted—Indicates that instruction execution should stop after execution of a tlbsync instruction. Negated—Indicates that the instruction execution may continue or resume after the completion of a tlbsync instruction.
Timing Comments	Assertion/Negation—May occur on any cycle. Start-Up— $\overline{TLBISYNC}$ is sampled at the negation of \overline{HRESET} to select the 32-bit data bus mode; if $\overline{TLBISYNC}$ is negated at start-up, the 32-bit mode is disabled and the default 64-bit mode is selected.

7.2.10 COP/Scan Interface

The MPC603e has extensive on-chip test capability including the following:

- Built-in instruction and data cache self-test (BIST)
- Debug control/observation (COP)
- Boundary scan (IEEE 1149.1 compliant interface)
- LSSD test control

The BIST hardware is not exercised as part of the power-on reset (POR) sequence. The COP and boundary scan logic are not used under typical operating conditions.

Detailed discussion of the MPC603e test functions is beyond the scope of this document; however, sufficient information has been provided to allow the system designer to disable the test functions that would impede normal operation.

The COP/scan interface is shown in Figure 7-2. For more information, see Section 8.9, “IEEE 1149.1-Compliant Interface.”

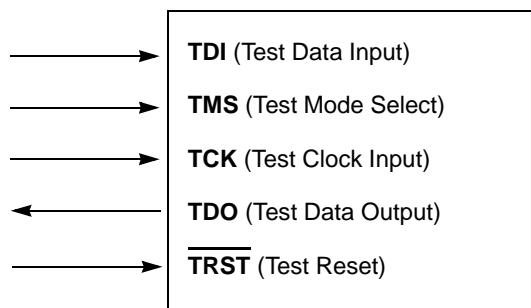


Figure 7-2. IEEE 1149.1-Compliant Boundary Scan Interface

7.2.11 Pipeline Tracking Support

The MPC603e provides for nonintrusive instruction pipeline tracking. Setting HID0[EICE] causes the address parity and data parity signals to be redefined as outputs providing pipeline tracking information. These signals toggle at the CPU clock rate and have special loading and timing requirements when in this mode.

Table 7-8 shows the outputs when HID0[EICE] is set.

Table 7-8. Pipeline Tracking Outputs

Bits	Function	Encoding
DP[0:1]	Fetch	00 None 01 Two 10 One 11 Branch
DP[2:3]	Retire	00 None 01 Two 10 One 11 Exception
DP[4:5]	Fold	00 None 01 First 10 Second 11 Both

Table 7-8. Pipeline Tracking Outputs (continued)

Bits	Function	Encoding
DP[6:7]	Prediction	00 Nonspec 01 Spec_2nd 10 Spec_both 11 Flush_spec
AP[0:3]	FEA	FEA[20:23]

Given the object code, these signals provide sufficient information to track instruction execution (except for register indirect branches). Register indirect branches may be tracked either by examining and matching potential target streams (nonintrusive but not always resolvable), or by forcing register indirect branch targets to be fetched externally by setting HID0[FBIOD].

Setting HID0[EICE] also directs the processor clock to the CLK_OUT signal, providing a synchronizing clock to the pipeline tracking outputs.

7.2.12 Clock Signals

The clock signal inputs of the MPC603e determine the system clock frequency and provide a flexible clocking scheme that allows the processor to operate at an integer multiple of the system clock frequency.

Refer to the appropriate hardware specifications for exact timing relationships of the clock signals.

7.2.12.1 System Clock (SYSCLK)—Input

The MPC603e requires a single system clock (SYSCLK) input. This input sets the frequency of operation for the bus interface. Internally, the MPC603e uses a phase-locked loop (PLL) circuit to generate a master clock for all of the CPU circuitry (including the bus interface circuitry) which is phase-locked to the SYSCLK input. The master clock may be set to an integer or half-integer multiple (1:1, 1.5:1, 2:1, 2.5:1, 3:1, 3.5:1, or 4:1) of the SYSCLK frequency allowing the CPU core to operate at an equal or greater frequency than the bus interface.

State Meaning	Asserted/Negated—The SYSCLK input is the primary clock input for the MPC603e, and represents the bus clock frequency for MPC603e bus operation. Internally, the MPC603e may be operating at an integer or half-integer multiple of the bus clock frequency.
Timing Comments	Duty cycle—Refer to the appropriate hardware specifications for timing comments.

Note: SYSCLK is used as the frequency reference for the internal PLL clock generator, and must not be suspended or varied during normal operation to ensure proper PLL operation.

7.2.12.2 Test Clock (CLK_OUT)—Output

The MPC603e provides a CLK_OUT signal for test purposes that allows the monitoring of the processor and bus clock frequencies. The frequency of the CLK_OUT signal is determined by the configuration of HID0[SBCLK, ECLK], as shown in Table 7-9. Note that the PID7t-603e CLK_OUT signal is driven at the processor frequency during the assertion of $\overline{\text{HRESET}}$; when $\overline{\text{HRESET}}$ is negated, CLK_OUT enters the default high-impedance state.

Table 7-9. CLK_OUT Signal Configuration

HID0[SBCLK]	HID0[ECLK]	CLK_OUT Output State
0	0	High impedance
0	1	Processor clock frequency
1	0	Half-bus clock frequency
1	1	Bus clock frequency

The CLK_OUT signal is an output-only signal on the MPC603e. Following are the state meaning and timing comments for the CLK_OUT output.

- State Meaning** Asserted/Negated—Provides PLL clock output for PLL testing and monitoring. The CLK_OUT signal clocks at either the processor clock frequency, bus clock frequency, or half-bus clock frequency if enabled by the appropriate bits in the HID0 register; the default state of CLK_OUT is high impedance. CLK_OUT is provided for testing purposes only.
- Timing Comments** Assertion/Negation—Refer to the appropriate hardware specifications for timing comments.

7.2.12.3 PLL Configuration (PLL_CFG[0:3])—Input

The PLL is configured by PLL_CFG[0:3]. For a given SYSCLK (bus) frequency, the PLL configuration signals set the internal CPU frequency of operation. Table 7-10 shows the PLL configuration.

Following are the state meaning and timing comments for the PLL_CFG[0:3] input.

- State Meaning** Asserted/Negated—Configures the operation of the PLL and the internal processor clock frequency. Settings are based on the desired bus and internal frequency of operation.

Timing Comments Assertion/Negation—Must remain stable during operation; should only be changed during the assertion of $\overline{\text{HRESET}}$ or during sleep mode. These bits may be read through HID1[PC0–PC3].

Table 7-10. PLL Configuration^{1, 2}

PLL_CFG[0:3]	Bus, CPU, and PLL Frequencies							
	CPU/ SYSCLK Ratio	Bus 16.6 MHz	Bus 20 MHz	Bus 25 MHz	Bus 33.3 MHz	Bus 40 MHz	Bus 50 MHz	Bus 66.6 MHz
0000	1:1	—	—	—	—	—	—	66.6 (133)
0001	1:1	—	—	—	33.3 (133)	40 (160)	50 (200)	—
0010	1:1	16.6 (133)	20 (160)	25 (200)	—	—	—	—
1100	1.5:1	—	—	—	—	—	75 (150)	100 (200)
0100	2:1	—	—	—	66.6 (133)	80 (160)	100 (200)	—
0101	2:1	33.3 (133)	40 (160)	50 (200)	—	—	—	—
0110	2.5:1	—	—	—	83.3 (166)	100 (200)	—	—
1000	3:1	—	—	75 (150)	100 (200)	—	—	—
1110	3.5:1	—	70 (140)	87.5 (175)	—	—	—	—
1010	4:1	66.6 (133)	80 (160)	100 (200)	—	—	—	—
0011	PLL bypass							
1111	Clock off							

¹ Some PLL configurations may select bus, CPU, or PLL frequencies that are not useful, not supported, or not tested for by the MPC603e. For complete, up-to-date information, refer to the appropriate hardware specifications. PLL frequencies, shown in parentheses, should not fall below 133 MHz, and should not exceed 200 MHz.

² In PLL-bypass mode, the SYSCLK input signal clocks the internal processor directly, and the bus is set for 1:1 mode operation. In clock-off mode, no clocking occurs inside the MPC603e regardless of the SYSCLK input.

7.2.13 Power and Ground Signals

The MPC603e provides the following connections for power and ground:

- V_{DD} and OV_{DD} —The V_{DD} and OV_{DD} signals provide the connection for the supply voltage. On the MPC603e, there is no electrical distinction between the V_{DD} and OV_{DD} signals.
- AV_{DD} —The AV_{DD} power signal provides power to the clock generation PLL. See the appropriate hardware specifications for information on how to use this signal.
- GND and OGND—The GND and OGND signals provide the connection for grounding the MPC603e. On the MPC603e, there is no electrical distinction between the GND and OGND signals.

Chapter 8

System Interface Operation

This chapter describes the MPC603e microprocessor's bus interface and its operation. It shows how the MPC603e signals, defined in Chapter 7, "Signal Descriptions," interact to perform address and data transfers.

8.1 Overview

The system interface prioritizes requests for bus operations from the instruction and data caches and performs bus operations per the MPC603e bus protocol. It includes address register queues, prioritization logic, and the bus control unit. The system interface latches snoop addresses for snooping in the data cache and address register queues, snoops for direct-store reply operations and reservations controlled by the Load Word and Reserve Indexed (**Iwarx**) and Store Word Conditional Indexed (**stwcx**) instructions, and maintains the touch load address for the cache. The interface allows one level of pipelining; that is, with certain restrictions discussed later, there can be two outstanding transactions at any given time. Accesses are prioritized with load operations preceding store operations.

Instructions are automatically fetched from the memory system into the instruction unit where they are dispatched to the execution units or forwarded to the branch processing unit at a peak rate of three instructions per clock (see Section 6.3, "Timing Considerations"). Conversely, load and store instructions explicitly specify the movement of operands to and from the integer and floating-point register files and the memory system.

When the MPC603e encounters an instruction or data access, it calculates the logical address (effective address in the architecture specification) and uses the low-order address bits to check for a hit in the on-chip, 16-Kbyte instruction and data caches. During cache lookup, the instruction and data memory management units (MMUs) use the higher-order address bits to calculate the virtual address, allowing them to calculate the physical address (real address in the architecture specification). The physical address bits are then compared with the corresponding cache tag bits to determine if a cache hit occurred. If the access misses in the corresponding cache, the physical address is used to access system memory.

In addition to loads, stores, and instruction fetches, the MPC603e performs software table search operations following TLB misses, cache cast-out operations when least recently used (LRU) cache lines are written to memory after a cache miss, and cache-line snoop

push-out operations when a modified cache line experiences a snoop hit from another bus master.

Figure 8-1 shows the address path from the execution units and instruction fetcher, through the translation logic to the caches and system interface logic.

The MPC603e uses separate address and data buses and a variety of control and status signals for performing reads and writes. The address bus is 32 bits wide and the data bus can be configured to be 32 or 64 bits wide. The interface is synchronous—all MPC603e inputs are sampled at and all outputs are driven from the rising edge of the bus clock. The bus can run at the full processor-clock frequency or at an integer division of the processor-clock speed. Although the MPC603e operates at 3.3 V, all the I/O signals are 5.0-V TTL-compatible.

8.1.1 Operation of the Instruction and Data Caches

The MPC603e provides independent instruction and data caches. Each cache is a physically-addressed, 16-Kbyte cache with four-way set-associativity. Both caches consist of 128 sets of four 8-word cache lines.

Because the on-chip data cache is a write-back primary cache, the predominant type of transaction is burst-read memory operations, followed by burst-write memory operations and single-beat (noncacheable or write-through) memory read and write operations. Additionally, there can be address-only operations, variants of the burst and single-beat operations such as, global memory operations that are snooped and atomic memory operations, and address retry activity (such as, when a snooped read access hits a modified line in the cache).

Because the MPC603e data cache tags are single ported, simultaneous load or store and snoop accesses cause resource contention. Snoop accesses have the highest priority and are given first access to the tags, unless the snoop access coincides with a tag write. In this case, the snoop is retried and must re-arbitrate for cache access. Loads or stores deferred due to snoop accesses are performed on the clock cycle following the snoop.

The MPC603e supports a three-state coherency protocol that implements the modified, exclusive, and invalid (MEI) cache states. The protocol is a subset of the MESI (modified/exclusive/ shared/invalid) four-state protocol and operates coherently in systems that contain four-state caches. With the exception of the **dcbz** instruction, the MPC603e does not broadcast cache control instructions. The cache control instructions are intended for the management of the local cache but not for other caches in the system.

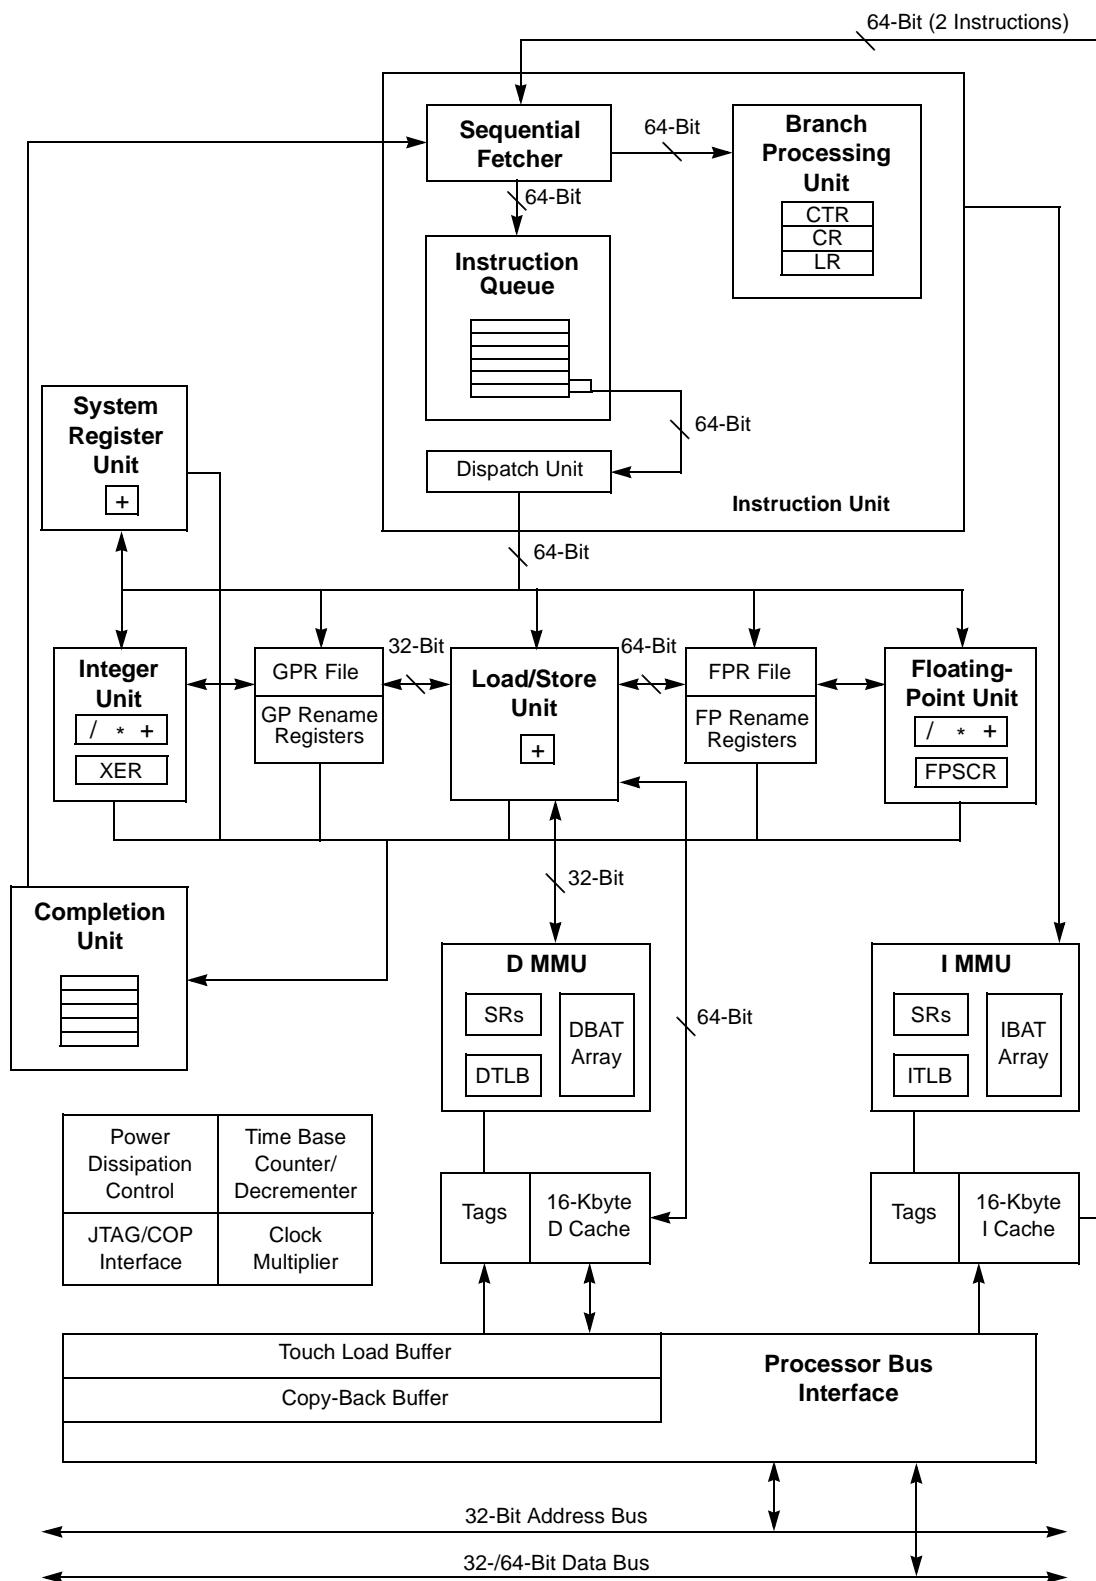


Figure 8-1. MPC603e Microprocessor Block Diagram

Cache lines in the MPC603e are loaded in four beats of 64 bits each (or eight beats of 32 bits each when operating in a 32-bit bus mode). The burst load is performed as a critical-double-word-first operation. The cache that is being loaded is blocked to internal accesses until the load completes (that is, no hits under misses). The critical double word is simultaneously written to the cache and forwarded to the requesting unit, minimizing stalls due to load delays.

Cache lines are selected for replacement based on an LRU algorithm. Each time a cache line is accessed, it is tagged as the most recently used line of the set. When a miss occurs, if all lines in the set are marked as valid, the LRU line is replaced with the new data. When data to be replaced is in the modified state, the modified data is written into a write-back buffer while the missed data is being read from memory. When the load completes, the MPC603e then pushes the replaced line from the write-back buffer to main memory in a burst write operation.

8.1.2 Operation of the System Interface

Memory accesses can occur in single-beat (1 to 8 bytes) and four-beat (32 bytes) burst data transfers when the MPC603e is configured with a 64-bit data bus. When the MPC603e is in the optional 32-bit data bus mode, memory accesses can occur in single-beat (1 to 4 bytes), two-beat (8 bytes), and eight-beat (32 bytes) bursts. The address and data buses are independent for memory accesses to support pipelining and split transactions. The MPC603e can pipeline as many as two transactions and has limited support for out-of-order split-bus transactions.

Access to the system interface is granted through an external arbitration mechanism that allows devices to compete for bus mastership. This arbitration mechanism is flexible, allowing the MPC603e to be integrated into systems that implement various fairness and bus-parking procedures to avoid arbitration overhead.

Typically, memory accesses are weakly ordered—sequences of operations, including load/store string and multiple instructions, do not necessarily complete in the order they begin—maximizing the bus efficiency without sacrificing data coherency. The MPC603e allows load operations to precede store operations (except when a dependency exists). In addition, the MPC603e can be configured to reorder high-priority store operations ahead of lower-priority store operations. Because the processor can dynamically optimize run-time ordering of load/store traffic, overall performance is improved.

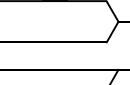
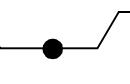
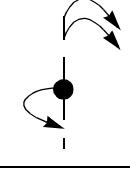
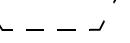
Note that the Synchronize (**sync**) instruction can be used to enforce strong ordering.

The following sections describe how the MPC603e interface operates, providing detailed timing diagrams that illustrate how the signals interact. A collection of more general timing diagrams are included as examples of typical bus operations.

Table 8-1 is a legend of the conventions used in the timing diagrams.

This is a synchronous interface—all MPC603e input signals are sampled and output signals are driven on the rising edge of the bus clock cycle (see the *MPC603e RISC Microprocessor Hardware Specifications* for exact timing information).

Table 8-1. Timing Diagram Legend

Feature	Example	Description
Gray	AP0	MPC603e input while the MPC603e is the bus master.
Bold	BR	MPC603e output while the MPC603e is the bus master.
Plain	Data	MPC603e input or output while the MPC603e is the bus master.
+	ADDR+	MPC603e output (grouped: here, address plus attributes)
lowercase	<u>qual_bg</u>	Internal MPC603e signal inaccessible to the user, but used in diagrams to clarify operations.
Curled arrow		Dependency.
Zig-zag		Indication that some clocks may have been skipped.
Unshaded		A valid output or input signal or bus that can be in any of the possible states indicated.
Shaded		MPC603e nonsampled input or indeterminately driven output among the possible states indicated.
Dot		Signal with sample point
Dot on dotted vertical line		A sampled condition (dot on high or low state) with multiple dependencies
Dotted signal		Timing for a signal had it been asserted.

8.1.2.1 Optional 32-Bit Data Bus Mode

The MPC603e supports an optional 32-bit data bus mode, which differs from the 64-bit data bus mode only in the byte lanes involved in the transfer and the number of data beats performed. A data tenure in the 32-bit data bus mode takes one, two, or eight beats

depending on the transfer size and the cache mode for the address. For additional information, see Section 8.6.1, “32-Bit Data Bus Mode.”

8.1.3 Direct-Store Accesses

The MPC603e does not support the extended transfer protocol for accesses to the direct-store storage space. If SR[T] is set, the memory access is a direct-store access. An attempt to access to a direct-store segment results in a DS1 exception.

8.2 Memory Access Protocol

Figure 8-2 shows that the address and data tenures are distinct from one another and that both consist of three phases—arbitration, transfer, and termination. Address and data tenures are independent (indicated in Figure 8-2 by the fact that the data tenure begins before the address tenure ends), which allows split-bus transactions to be implemented at the system level in multiprocessor systems. Figure 8-2 shows a data transfer that consists of a single-beat transfer of as many as 64 bits. Four-beat burst transfers of 32-byte cache lines require data transfer termination signals for each beat of data.

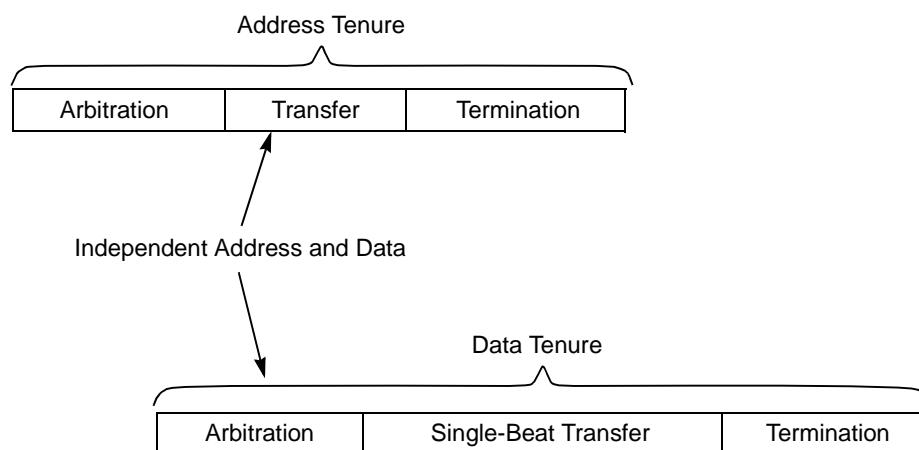


Figure 8-2. Overlapping Tenures on the Bus for a Single-Beat Transfer

The basic functions of the address and data tenures are as follows:

- Address tenure
 - Arbitration: During arbitration, address bus arbitration signals are used to gain address bus mastership.
 - Transfer: After the MPC603e is the address bus master, it transfers the address on the address bus. The address signals and the transfer attribute signals control the address transfer. The address parity and address parity error signals ensure the integrity of the address transfer.

- Termination: After the address transfer, the system signals that the address tenure is complete or that it must be repeated.
- Data tenure
 - Arbitration: To begin the data tenure, the MPC603e arbitrates for data bus mastership.
 - Transfer: After the MPC603e is the data bus master, it samples the data bus for read operations or drives the data bus for write operations. The data parity and data parity error signals ensure the integrity of the transfer.
 - Termination: Data termination signals are required after each beat. Note that in a single-beat transaction, the data termination signals also indicate the end of the tenure, while in burst accesses, the data termination signals apply to individual beats and indicate the end of the tenure only after the final data beat.

The MPC603e generates an address-only bus transfer during the execution of the **dcbz** instruction, and uses only the address bus with no data transfer involved. Additionally, the MPC603e retry capability provides an efficient snooping protocol for systems with multiple memory systems (including caches) that must remain coherent.

8.2.1 Arbitration Signals

Arbitration for both address and data bus mastership is performed by a central, external arbiter and, minimally, by the arbitration signals shown in Section 7.2.1, “Address Bus Arbitration Signals.” Most arbiter implementations require additional signals to coordinate bus master/slave/snooping activities. Note that address bus busy (\overline{ABB}) and data bus busy (\overline{DBB}) are bidirectional signals. These signals are inputs unless the MPC603e has mastership of one or both of the respective buses; they must be connected high through pull-up resistors so that they remain negated when no devices have control of the buses.

The following list describes the address arbitration signals:

- \overline{BR} (bus request)—Assertion indicates that the MPC603e is requesting mastership of the address bus.
- \overline{BG} (bus grant)—Assertion indicates that the MPC603e may, with the proper qualification, assume mastership of the address bus. A qualified bus grant occurs when \overline{BG} is asserted and \overline{ABB} and \overline{ARTRY} are negated.
If the MPC603e is parked, \overline{BR} need not be asserted for the qualified bus grant.
- \overline{ABB} (address bus busy)—Assertion by the MPC603e indicates that the MPC603e is the address bus master.

The following list describes the data arbitration signals:

- \overline{DBG} (data bus grant)—Indicates that the MPC603e may, with the proper qualification, assume mastership of the data bus. A qualified data bus grant occurs when \overline{DBG} is asserted while \overline{DBB} , \overline{DRTRY} , and \overline{ARTRY} are negated.

\overline{DBB} is driven by the current bus master, \overline{DRTRY} is driven only from the bus, and \overline{ARTRY} is driven from the bus, but only for the address tenure associated with the current data tenure (that is, not from another address tenure).

- \overline{DBWO} (data bus write only)—Assertion indicates that the MPC603e may perform the data bus tenure for an outstanding write address even if a read address is pipelined before the write address. If \overline{DBWO} is asserted, the MPC603e assumes data bus mastership for a pending data bus write operation; the MPC603e takes the data bus for a pending read operation if this input is asserted along with \overline{DBG} and no write is pending. Care must be taken with \overline{DBWO} to ensure the desired write is queued (for example, a cache-line snoop push-out operation).
- \overline{DBB} (data bus busy)—Assertion by the MPC603e indicates that the MPC603e is the data bus master. The MPC603e always assumes data bus mastership if it needs the bus and is given a qualified data bus grant (see \overline{DBG}).

For more detailed information on the arbitration signals, refer to Section 7.2.1, “Address Bus Arbitration Signals,” and Section 7.2.6, “Data Bus Arbitration Signals.”

8.2.2 Address Pipelining and Split-Bus Transactions

The MPC603e protocol provides independent address and data bus capability to support pipelined and split-bus transaction system organizations. Address pipelining allows the address tenure of a new bus transaction to begin before the data tenure of the current transaction has finished. Split-bus transaction capability allows other bus activity to occur (either from the same master or from different masters) between the address and data tenures of a transaction.

While this capability does not inherently reduce memory latency, support for address pipelining and split-bus transactions can greatly improve effective bus/memory throughput. For this reason, these techniques are most effective in shared-memory multiprocessor implementations where bus bandwidth is an important measurement of system performance.

External arbitration is required in systems in which multiple devices must compete for the system bus. The design of the external arbiter affects pipelining by regulating address bus grant (\overline{BG}), data bus grant (\overline{DBG}), and address acknowledge (\overline{AACK}) signals. For example, a one-level pipeline is enabled by asserting \overline{AACK} to the current address bus master and granting mastership of the address bus to the next requesting master before the current data bus tenure has completed. Two address tenures can occur before the current data bus tenure completes.

The MPC603e can pipeline its own transactions to a depth of one level (intraprocessor pipelining); however, the MPC603e bus protocol does not constrain the maximum number of levels of pipelining that can occur on the bus between multiple masters (interprocessor

pipelining). The external arbiter must control the pipeline depth and synchronization between masters and slaves.

In a pipelined implementation, data bus tenures are kept in strict order with respect to address tenures. However, external hardware can further decouple the address and data buses, allowing the data tenures to occur out of order with respect to the address tenures. This requires some form of system tag to associate the out-of-order data transaction with the proper originating address transaction (not defined for the MPC603e interface). Individual bus requests and data bus grants from each processor can be used by the system to implement tags to support interprocessor, out-of-order transactions.

The MPC603e supports a limited intraprocessor out-of-order, split-transaction capability via the data bus write only (DBWO) signal. For more information concerning the use of DBWO, see Section 8.10, “Using DBWO.”

8.3 Address Bus Tenure

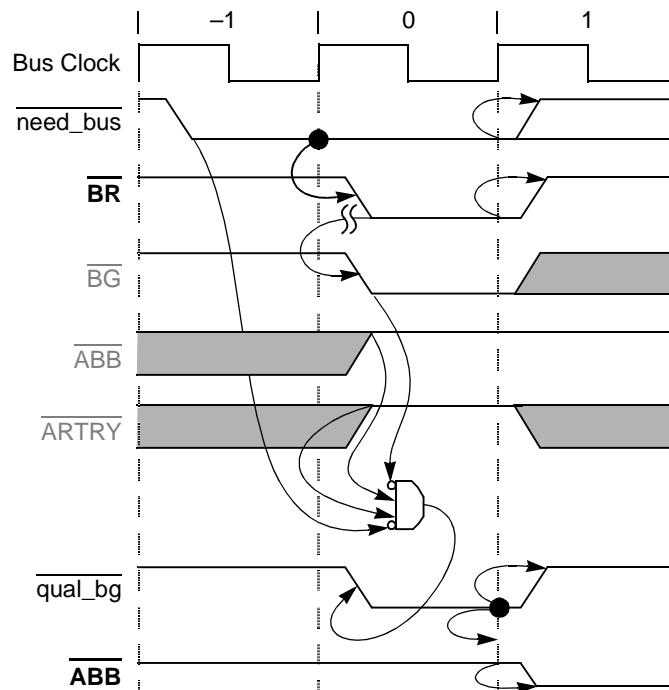
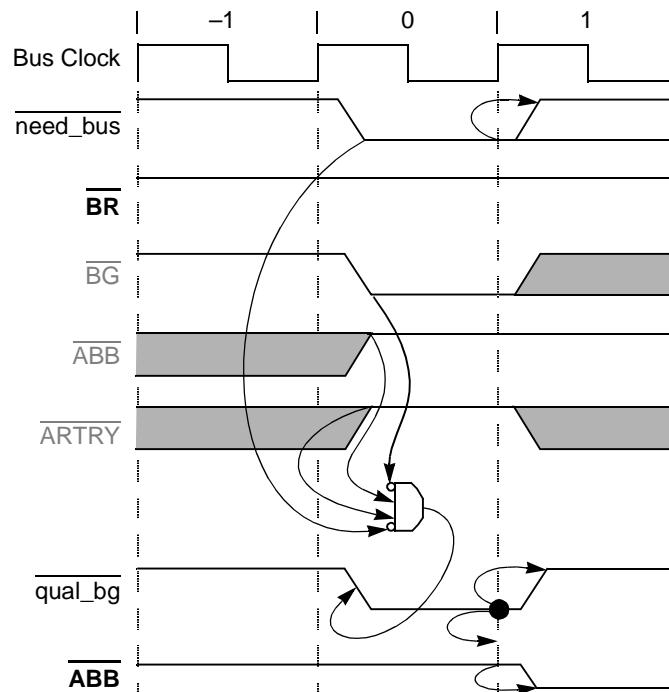
This section describes the address bus arbitration, transfer, and termination phases.

8.3.1 Address Bus Arbitration

When the MPC603e needs access to the external bus and is not parked (BG is negated), it asserts bus request (BR) until it is granted mastership of the bus and the bus is available (see Figure 8-3). The external arbiter must grant master-elect status to the potential master by asserting the bus grant (BG) signal. The MPC603e requesting the bus determines that the bus is available when the ABB input is negated (the address bus is not busy), BG is asserted, and the address retry (ARTRY) input is negated. This is referred to as a qualified bus grant. The potential master assumes address bus mastership by asserting ABB when it receives a qualified bus grant.

External arbiters must allow only one device at a time to be the address bus master. In implementations where no other device can be a master, BG can be grounded (always asserted) to continually grant mastership of the address bus to the MPC603e.

If the MPC603e asserts BR before the external arbiter asserts BG, the MPC603e is considered to be unparked, as shown in Figure 8-3. Figure 8-4 shows the parked case, where a qualified bus grant exists on the clock edge following a need_bus condition. Notice that the bus clock cycle required for arbitration is eliminated if the MPC603e is parked, reducing overall memory latency for a transaction. The MPC603e always negates ABB for at least one bus clock cycle after AACK is asserted, even if it is parked and has another transaction pending.

**Figure 8-3. Address Bus Arbitration****Figure 8-4. Address Bus Arbitration Showing Bus Parking**

Typically, bus parking is provided to the device that was the most recent bus master. However, system designers may choose other schemes, such as providing unrequested bus grants in situations where it is easy to correctly predict the next device requesting bus mastership.

When the MPC603e receives a qualified bus grant, it assumes address bus mastership by asserting \overline{ABB} and negating \overline{BR} . Meanwhile, the MPC603e drives the address for the requested access onto the address bus and asserts \overline{TS} to indicate the start of a new transaction.

When designing external bus arbitration logic, note that the MPC603e may assert \overline{BR} without using the bus after it receives the qualified bus grant. For example, in a system using bus snooping, if the MPC603e asserts \overline{BR} to perform a replacement copy-back operation, another device can invalidate that line before the MPC603e is granted mastership of the bus. Once the MPC603e is granted the bus, it no longer needs to perform the copy-back operation; therefore, the MPC603e does not assert \overline{ABB} and does not use the bus for the copy-back operation. Note that the MPC603e asserts \overline{BR} for at least 1 clock cycle in these instances.

8.3.2 Address Transfer

During the address transfer, the physical address and all attributes of the transaction are transferred from the bus master to slave devices. Snooping logic may monitor the transfer to enforce cache coherency; see discussion about snooping in Section 8.3.3, “Address Transfer Termination.”

The signals used in the address transfer include the following signal groups:

- Address transfer start signal: Transfer start (\overline{TS})
- Address transfer signals: Address bus ($A[0:31]$), address parity ($AP[0:3]$), and address parity error (\overline{APE}).
- Address transfer attribute signals: Transfer type ($TT[0:4]$), transfer code ($TC[0:1]$), transfer size ($TSIZ[0:2]$), transfer burst (\overline{TBST}), cache inhibit (\overline{CI}), write-through (\overline{WT}), global (\overline{GBL}), and cache set element ($CSE[0:1]$).

Figure 8-5 shows that the timing for all of these signals, except \overline{TS} and \overline{APE} , is identical. All of the address transfer and address transfer attribute signals are combined into the ADDR+ grouping in Figure 8-5. The \overline{TS} signal indicates that the MPC603e has begun an address transfer and that the address and transfer attributes are valid (within the context of a synchronous bus). The MPC603e always asserts \overline{TS} coincident with \overline{ABB} . As an input, \overline{TS} need not coincide with the assertion of \overline{ABB} on the bus (that is, \overline{TS} can be asserted with, or on, a subsequent clock cycle after \overline{ABB} is asserted; the MPC603e tracks this transaction correctly).

In Figure 8-5, the address transfer occurs during bus clock cycles 1 and 2 (arbitration occurs in bus clock cycle 0 and the address transfer is terminated in bus clock 3). In this diagram,

the address bus termination input, $\overline{\text{AACK}}$, is asserted to the MPC603e on the bus clock following assertion of $\overline{\text{TS}}$ (as shown by the dependency line). This is the minimum duration of the address transfer for the MPC603e; the duration can be extended by delaying the assertion of $\overline{\text{AACK}}$ for one or more bus clocks.

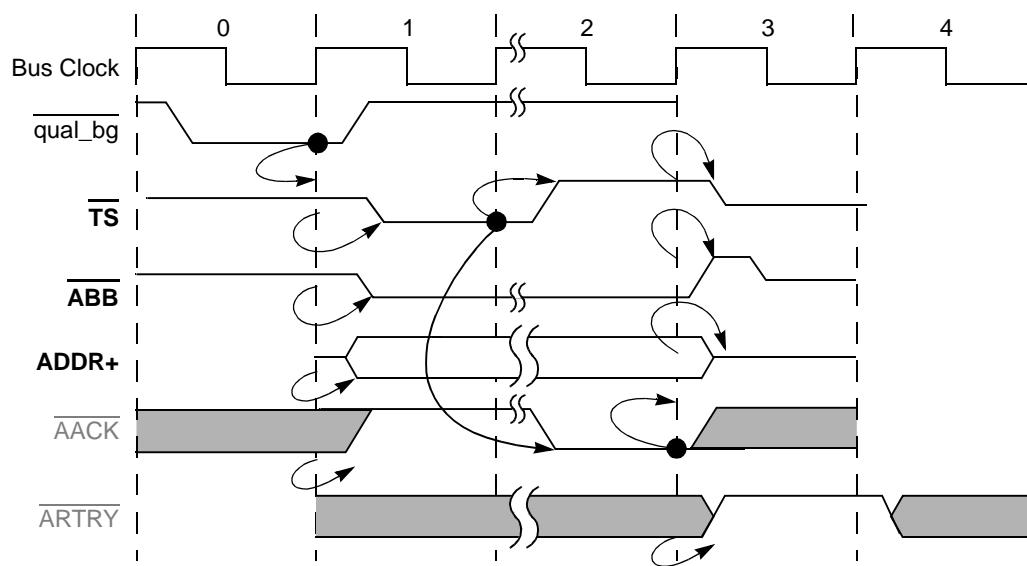


Figure 8-5. Address Bus Transfer

8.3.2.1 Address Bus Parity

The MPC603e always generates 1 bit of correct odd-byte parity for each of the 4 bytes of address when a valid address is on the bus. The calculated values are placed on the AP[0:3] outputs when the MPC603e is the address bus master. If the MPC603e is not the master, and $\overline{\text{TS}}$ and $\overline{\text{GBL}}$ are asserted together (qualified condition for snooping memory operations), the calculated values are compared with the AP[0:3] inputs. If there is an error and address parity checking is enabled (HID0[EBA] is set), the $\overline{\text{APE}}$ output is asserted. An address bus parity error causes a checkstop condition if MSR[ME] is cleared. For more information about checkstop conditions, see Chapter 4, “Exceptions.”

8.3.2.2 Address Transfer Attribute Signals

The transfer attribute signals include several encoded signals such as the transfer type (TT[0:4]) signals, transfer burst (TBST) signal, transfer size (TSIZ[0:2]) signals, and transfer code (TC[0:1]) signals. Section 7.2.4, “Address Transfer Attribute Signals,” describes the encodings for the address transfer attribute signals.

8.3.2.2.1 Transfer Type (TT[0:4]) Signals

Snooping logic should fully decode the transfer type signals if the $\overline{\text{GBL}}$ signal is asserted. Slave devices can sometimes use the individual transfer type signals without fully decoding

the group. For a complete description of the encoding for transfer type signals TT[0:4], refer to Table 7-1 and Table 7-2.

8.3.2.2.2 Transfer Size (TSIZ[0:2]) Signals

The transfer size signals (TSIZ[0:2]) indicate the size of the requested data transfer as shown in Table 8-2. The TSIZ[0:2] signals may be used along with $\overline{\text{TBST}}$ and A[29:31] to determine which portion of the data bus contains valid data for a write transaction or which portion of the bus should contain valid data for a read transaction. Note that for a burst transaction (as indicated by the assertion of $\overline{\text{TBST}}$), TSIZ[0:2] are always set to 0b010. Therefore, if $\overline{\text{TBST}}$ is asserted, the memory system should transfer a total of 8 words (32 bytes), regardless of the TSIZ[0:2] encoding.

Table 8-2. Transfer Size Signal Encodings

$\overline{\text{TBST}}$	TSIZ0	TSIZ1	TSIZ2	Transfer Size
Asserted	0	1	0	8-word burst
Negated	0	0	0	8 bytes
Negated	0	0	1	1 byte
Negated	0	1	0	2 bytes
Negated	0	1	1	3 bytes
Negated	1	0	0	4 bytes
Negated	1	0	1	5 bytes (N/A)
Negated	1	1	0	6 bytes (N/A)
Negated	1	1	1	7 bytes (N/A)

The basic coherency size of the bus is defined to be 32 bytes (corresponding to one cache line). Data transfers that cross an aligned, 32-byte boundary either must present a new address onto the bus at that boundary (for coherency consideration) or must operate as noncoherent data with respect to the MPC603e. The MPC603e never generates a bus transaction with a transfer size of 5, 6, or 7 bytes.

8.3.2.3 Burst Ordering During Data Transfers

During burst data transfer operations, 32 bytes of data (one cache line) are transferred to or from the cache in order. Burst write transfers are always performed zero double word first, but because burst reads are performed critical-double-word-first, a burst-read transfer may not start with the first double word of the cache line, and the cache line fill may wrap around the end of the cache line. This section describes the burst ordering for the MPC603e when operating in either the 64- or 32-bit bus mode.

Table 8-3 describes the burst ordering when the MPC603e is configured with a 64-bit data bus.

Table 8-3. Burst Ordering—64-Bit Bus

Data Transfer	For Starting Address:			
	A[27:28] = 00	A[27:28] = 01	A[27:28] = 10	A[27:28] = 11
First data beat	DW0	DW1	DW2	DW3
Second data beat	DW1	DW2	DW3	DW0
Third data beat	DW2	DW3	DW0	DW1
Fourth data beat	DW3	DW0	DW1	DW2

Note: A[29:31] are always 0b000 for burst transfers by the MPC603e.

Table 8-4 describes the burst ordering when the MPC603e is configured with a 32-bit bus.

Table 8-4. Burst Ordering—32-Bit Bus

Data Transfer	For Starting Address:			
	A[27:28] = 00	A[27:28] = 01	A[27:28] = 10	A[27:28] = 11
First data beat	DW0-U	DW1-U	DW2-U	DW3-U
Second data beat	DW0-L	DW1-L	DW2-L	DW3-L
Third data beat	DW1-U	DW2-U	DW3-U	DW0-U
Fourth data beat	DW1-L	DW2-L	DW3-L	DW0-L
Fifth data beat	DW2-U	DW3-U	DW0-U	DW1-U
Sixth data beat	DW2-L	DW3-L	DW0-L	DW1-L
Seventh data beat	DW3-U	DW0-U	DW1-U	DW2-U
Eighth data beat	DW3-L	DW0-L	DW1-L	DW2-L

Note: A[29:31] are always 0b000 for burst transfers by the MPC603e. U and L represent the upper and lower word of the double word, respectively.

8.3.2.4 Effect of Alignment in Data Transfers (64-Bit Bus)

Table 8-5 lists the aligned transfers that can occur on the MPC603e bus when configured with a 64-bit width. These are transfers in which the data is aligned to an address that is an integer multiple of the size of the data. For example, Table 8-5 shows that 1-byte data is always aligned; however, for a 4-byte word to be aligned, it must be oriented on an address that is a multiple of 4.

Table 8-5. Aligned Data Transfers (64-Bit Bus)

Transfer Size	TSIZ0	TSIZ1	TSIZ2	A[29:31]	Data Bus Byte Lanes ¹							
					0	1	2	3	4	5	6	7
Byte	0	0	1	000	A	—	—	—	—	—	—	—
	0	0	1	001	—	A	—	—	—	—	—	—
	0	0	1	010	—	—	A	—	—	—	—	—
	0	0	1	011	—	—	—	A	—	—	—	—
	0	0	1	100	—	—	—	—	A	—	—	—
	0	0	1	101	—	—	—	—	—	A	—	—
	0	0	1	110	—	—	—	—	—	—	A	—
	0	0	1	111	—	—	—	—	—	—	—	A
Half word	0	1	0	000	A	A	—	—	—	—	—	—
	0	1	0	010	—	—	A	A	—	—	—	—
	0	1	0	100	—	—	—	—	A	A	—	—
	0	1	0	110	—	—	—	—	—	—	A	A
Word	1	0	0	000	A	A	A	A	—	—	—	—
	1	0	0	100	—	—	—	—	A	A	A	A
Double word	0	0	0	000	A	A	A	A	A	A	A	A

¹ A: These entries indicate the byte portions of the requested operand that are read or written during that bus transaction.

—: These entries are not required and are ignored during read transactions and are driven with undefined data during all write transactions.

The MPC603e supports misaligned memory operations, although their use may substantially degrade performance. Misaligned memory transfers address memory that is not aligned to the size of the data being transferred (such as, a word read of an odd byte address). Although most of these operations hit in the primary cache (or generate burst memory operations if they miss), the MPC603e interface supports misaligned transfers within a word (32-bit aligned) boundary, as shown in Table 8-6. Note that the 4-byte transfer in Table 8-6 is only one example of misalignment. As long as the attempted transfer does not cross a word boundary, the MPC603e can transfer the data on the misaligned address (for example, a half-word read from an odd byte-aligned address). An attempt to address data that crosses a word boundary requires two bus transfers to access the data. Note that an attempt to load or store a floating-point operand that is not word-aligned will result in a floating-point alignment exception. For more information, refer to Section 4.5.6, “Alignment Exception (0x00600).”

Table 8-6. Misaligned Data Transfers (4-Byte Examples)

Transfer Size (4 Bytes)	TSIZ[0:2]	A[29:31]	Data Bus Byte Lanes							
			0	1	2	3	4	5	6	7
Aligned	1 0 0	0 0 0	A	A	A	A	—	—	—	—
Misaligned	First access	0 1 1	0 0 1	—	A	A	A	—	—	—
		0 0 1	1 0 0	—	—	—	—	A	—	—
Misaligned	First access	0 1 0	0 1 0	—	—	A	A	—	—	—
		0 1 0	1 0 0	—	—	—	—	A	A	—
Misaligned	First access	0 0 1	0 1 1	—	—	—	A	—	—	—
		0 1 1	1 0 0	—	—	—	—	A	A	A
Aligned	1 0 0	1 0 0	—	—	—	—	A	A	A	A
Misaligned	First access	0 1 1	1 0 1	—	—	—	—	—	A	A
		0 0 1	0 0 0	A	—	—	—	—	—	—
Misaligned	First access	0 1 0	1 1 0	—	—	—	—	—	—	A
		0 1 0	0 0 0	A	A	—	—	—	—	—
Misaligned	First access	0 0 1	1 1 1	—	—	—	—	—	—	—
		0 1 1	0 0 0	A	A	A	—	—	—	—

Notes:

A: Byte lane used.

—: Byte lane not used.

Due to the performance degradations associated with misaligned memory operations, they are best avoided. In addition to the double-word straddle boundary condition, the address translation logic can generate substantial exception overhead when the load/store multiple and load/store string instructions access misaligned data. It is strongly recommended that software attempt to align code and data where possible.

8.3.2.5 Effect of Alignment in Data Transfers (32-Bit Bus)

The aligned data transfer cases for 32-bit data bus mode are shown in Table 8-7. All of the transfers require a single data beat (if caching-inhibited or write-through) except for double-word cases which require two data beats. The double-word case is only generated by the MPC603e for load or store double operations to/from the floating-point GPRs. All caching-inhibited instruction fetches are performed as word operations.

Table 8-7. Aligned Data Transfers (32-Bit Bus Mode)

Transfer Size	TSIZ0	TSIZ1	TSIZ2	A[29:31]	Data Bus Byte Lanes							
					0	1	2	3	4	5	6	7
Byte	0	0	1	000	A	—	—	—	x	x	x	x
	0	0	1	001	—	A	x	—	x	x	x	x
	0	0	1	010	—	—	A	—	x	x	x	x
	0	0	1	011	—	—	—	A	x	x	x	x
	0	0	1	100	A	—	—	—	x	x	x	x
	0	0	1	101	—	A	—	—	x	x	x	x
	0	0	1	110	—	—	A	—	x	x	x	x
	0	0	1	111	—	—	—	A	x	x	x	x
Half word	0	1	0	000	A	A	—	—	x	x	x	x
	0	1	0	010	—	—	A	A	x	x	x	x
	0	1	0	100	A	A	—	—	x	x	x	x
	0	1	0	110	—	—	A	A	x	x	x	x
Word	1	0	0	000	A	A	A	A	x	x	x	x
	1	0	0	100	A	A	A	A	x	x	x	x
Double word	0	0	0	000	A	A	A	A	x	x	x	x
Second beat	0	0	0	000	A	A	A	A	x	x	x	x

Notes:

A: Byte lane used.

—: Byte lane not used.

x: Byte lane not used in 32-bit bus mode.

Misaligned data transfers when the MPC603e is configured with a 32-bit data bus operate in the same way as when configured with a 64-bit data bus, with the exception that only the DH[0:31] data bus is used. See Table 8-8 for an example of a 4-byte misaligned transfer starting at each possible byte address within a double word.

Table 8-8. Misaligned 32-Bit Data Bus Transfer (4-Byte Examples)

Transfer Size (4 Bytes)	TSIZ[0:2]	A[29:31]	Data Bus Byte Lanes							
			0	1	2	3	4	5	6	7
Aligned	1 0 0	0 0 0	A	A	A	A	X	X	X	X
Misaligned	First access	0 1 1	0 0 1		A	A	A	X	X	X
		0 0 1	1 0 0	A	—	—	—	X	X	X
Misaligned	First access	0 1 0	0 1 0	—	—	A	A	X	X	X
		0 1 0	1 0 0	A	A	—	X	X	X	X
Misaligned	First access	0 0 1	0 1 1	—	—	—	A	X	X	X
		0 1 1	1 0 0	A	A	A	—	X	X	X
Aligned	1 0 0	1 0 0	A	A	A	A	X	X	X	X
Misaligned	First access	0 1 1	1 0 1	—	A	A	A	X	X	X
		0 0 1	0 0 0	A	—	—	—	X	X	X
Misaligned	First access	0 1 0	1 1 0	—	—	A	A	X	X	X
		0 1 0	0 0 0	A	A	—	—	X	X	X
Misaligned	First access	0 0 1	1 1 1	—	—	—	A	X	X	X
		0 1 1	0 0 0	A	A	A	—	X	X	X

Notes:

A: Byte lane used.

—: Byte lane not used.

x: Byte lane not used in 32-bit bus mode.

8.3.2.5.1 Alignment of External Control Instructions

The size of the data transfer associated with the **eciwx** and **ecowx** instructions is always 4 bytes. However, if either is misaligned and crosses any word boundary, the MPC603e generates two bus operations, each smaller than 4 bytes. For the first bus operation, bits A[29:31] equal bits 29–31 of the effective address of the instruction, which is 0b101, 0b110, or 0b111. The size associated with the first bus operation will be 3, 2, or 1 bytes, respectively. For the second bus operation, bits A[29:31] equal 0b000 and the size associated with the operation is 1, 2, or 3 bytes, respectively. For both operations, TBST and TSIZ[0:2] are redefined to specify the resource ID (RID). The resource ID is copied from bits 28–31 of the EAR. For **eciwx/ecowx** operations, TBST is high if EAR[28] is set. The size of the second bus operation cannot be deduced from the operation itself; the system must determine how many bytes were transferred on the first bus operation to determine the size of the second operation.

Furthermore, the two bus operations associated with such a misaligned external control instruction are not atomic. That is, the MPC603e may initiate other types of memory operations between the two transfers. Also, the two bus operations associated with a

misaligned **ecowx** may be interrupted by an **eciwx** bus operation, and vice versa. The MPC603e does guarantee that the two operations associated with a misaligned **ecowx** will not be interrupted by another **ecowx** operation, and likewise for **eciwx**.

Because a misaligned external control address is considered a programming error, the system may choose to assert **TEA** or otherwise cause an exception when a misaligned external control bus operation occurs.

8.3.2.6 Transfer Code (TC[0:1]) Signals

The TC0 and TC1 signals provide supplemental information about the corresponding address. Note that the TC x signals can be used with the TT[0:4] and **TBST** signals to further define the current transaction.

Table 8-9 shows the encodings of the TC0 and TC1 signals.

Table 8-9. Transfer Code Encoding

TC[0:1]	Read	Write
0 0	Data transaction	Any write
0 1	Touch load	N/A
1 0	Instruction fetch	N/A
1 1	(Reserved)	N/A

8.3.3 Address Transfer Termination

The address tenure of a bus operation is terminated when completed with the assertion of **AACK**, or retried with the assertion of **ARTRY**. The MPC603e does not terminate the address transfer until the **AACK** (address acknowledge) input is asserted; therefore, the system can extend the address transfer phase by delaying the assertion of **AACK** to the MPC603e. **AACK** can be asserted as early as the bus clock cycle following **TS** (see Figure 8-6), which allows a minimum address tenure of two bus cycles. However, when the MPC603e clock is configured for 1:1 or 1.5:1 processor-to-bus clock mode, the **ARTRY** snoop response cannot be determined in the minimum allowed address tenure period. Thus, in a system with two or more MPC603e processors using 1:1 or 1.5:1 clock mode, **AACK** must not be asserted until the third clock of the address tenure (one address wait state) to allow the snooping MPC603e processors an opportunity to assert **ARTRY** on the bus. For other clock configurations (2:1, 2.5:1, 3:1, 3.5:1, and 4:1), the **ARTRY** snoop response can be determined in the minimum address tenure period, and **AACK** may be asserted as early as the second bus clock of the address tenure. As shown in Figure 8-6, these signals are asserted for one bus clock cycle, three-stated for half of the next bus clock cycle, driven high until the following bus cycle, and finally three-stated. Note that **AACK** must be asserted for only one bus clock cycle.

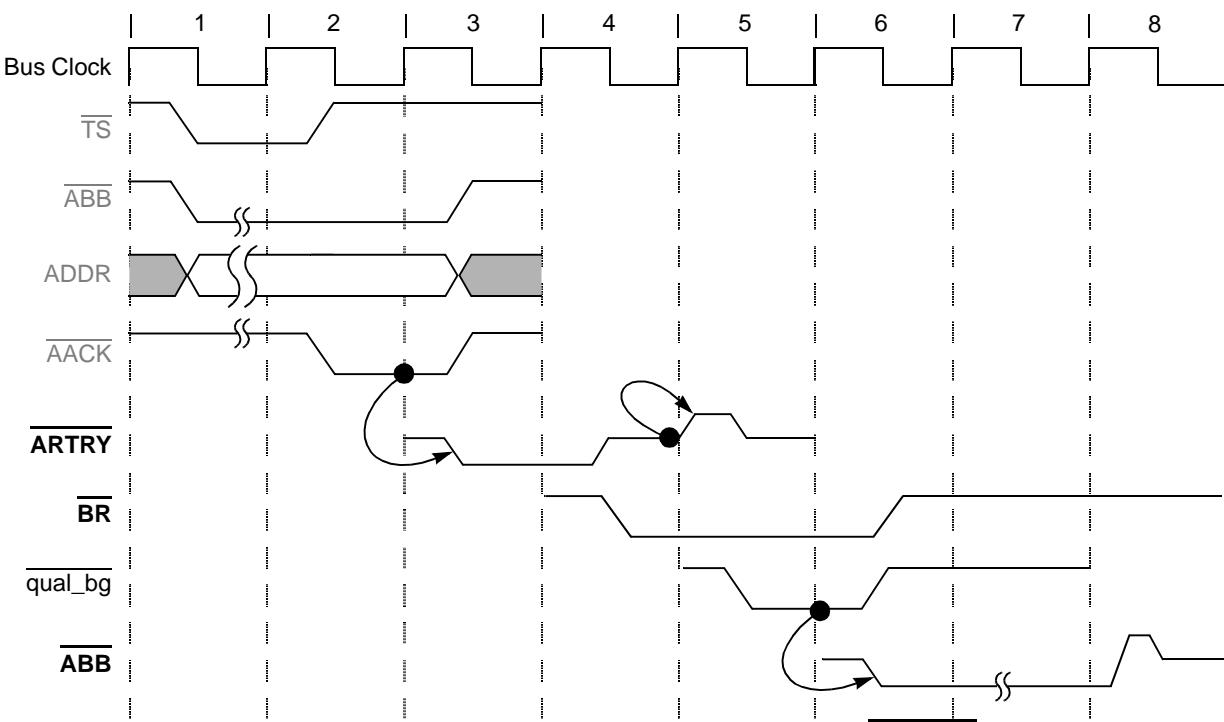


Figure 8-6. Snooped Address Cycle with ARTRY

The address transfer can be terminated with the requirement to retry if ARTRY is asserted anytime during the address tenure and through the cycle following AACK. The assertion causes the entire transaction (address and data tenure) to be rerun. As a snooping device, the MPC603e asserts ARTRY for a snooped transaction that hits modified data in the data cache that must be written back to memory, or if the snooped transaction could not be serviced. As a bus master, the MPC603e responds to an assertion of ARTRY by aborting the bus transaction and re-requesting the bus. Note that after recognizing an assertion of ARTRY and aborting the transaction in progress, the MPC603e is not guaranteed to run the same transaction the next time it is granted the bus due to internal reordering of load and store operations.

If an address retry is required, the ARTRY response is asserted by a bus snooping device as early as the second cycle after the assertion of TS (or until the third cycle following TS if 1:1 or 1.5:1 processor to bus clock ratio is selected). Once asserted, ARTRY must remain asserted through the cycle after the assertion of AACK. The assertion of ARTRY during the cycle after the assertion of AACK is referred to as a qualified ARTRY. An earlier assertion of ARTRY during the address tenure is referred to as an early ARTRY.

As a bus master, the MPC603e recognizes either an early or qualified ARTRY and prevents the data tenure associated with the retried address tenure. If the data tenure has already begun, the MPC603e aborts and terminates the data tenure immediately even if the burst data has been received. If the assertion of ARTRY is received up to or on the bus cycle following the first (or only) assertion of TA for the data tenure, the MPC603e ignores the first data beat, and if it is a load operation, does not forward data internally to the cache and

execution units. If $\overline{\text{ARTRY}}$ is asserted after the first (or only) assertion of $\overline{\text{TA}}$, improper operation of the bus interface may result.

During the clock of a qualified $\overline{\text{ARTRY}}$, the MPC603e also determines if it should negate $\overline{\text{BR}}$ and ignore $\overline{\text{BG}}$ on the following cycle. On the following cycle, only the snooping master that asserted $\overline{\text{ARTRY}}$ and needs to perform a snoop copy-back operation is allowed to assert $\overline{\text{BR}}$. This guarantees the snooping master an opportunity to request and be granted the bus before the just-retried master can restart its transaction. Note that a nonclocked bus arbiter may detect the assertion of address bus request by the bus master that asserted $\overline{\text{ARTRY}}$, and return a qualified bus grant one cycle earlier than shown in Figure 8-6.

8.4 Data Bus Tenure

This section describes the data bus arbitration, transfer, and termination phases defined by the MPC603e memory access protocol. The phases of the data tenure are identical to those of the address tenure, underscoring the symmetry in the control of the two buses.

8.4.1 Data Bus Arbitration

Data bus arbitration uses the data arbitration signal group— $\overline{\text{DBG}}$, $\overline{\text{DBWO}}$, and $\overline{\text{DBB}}$. Additionally, the combination of $\overline{\text{TS}}$ and $\text{TT}[0:4]$ provides information about the data bus request to external logic.

The $\overline{\text{TS}}$ signal is an implied data bus request from the MPC603e; the arbiter must qualify $\overline{\text{TS}}$ with the transfer type (TT) encodings to determine if the current address transfer is an address-only operation, which does not require a data bus transfer (see Figure 8-6). If the data bus is needed, the arbiter grants data bus mastership by asserting the $\overline{\text{DBG}}$ input to the MPC603e. As with the address bus arbitration phase, the MPC603e must qualify the $\overline{\text{DBG}}$ input with a number of input signals before assuming bus mastership, as shown in Figure 8-7.

A qualified data bus grant can be expressed as the following:

Qualified data bus grant = $\overline{\text{DBG}}$ asserted while $\overline{\text{DBB}}$, $\overline{\text{DRTRY}}$, and $\overline{\text{ARTRY}}$ (associated with the data bus operation) are negated.

When a data tenure overlaps with its associated address tenure, a qualified $\overline{\text{ARTRY}}$ assertion coincident with a data bus grant signal does not result in data bus mastership ($\overline{\text{DBB}}$ is not asserted). Otherwise, the MPC603e always asserts $\overline{\text{DBB}}$ on the bus clock cycle after recognition of a qualified data bus grant. Since the MPC603e can pipeline transactions, there may be an outstanding data bus transaction when a new address transaction is retried. In this case, the MPC603e becomes the data bus master to complete the previous transaction.

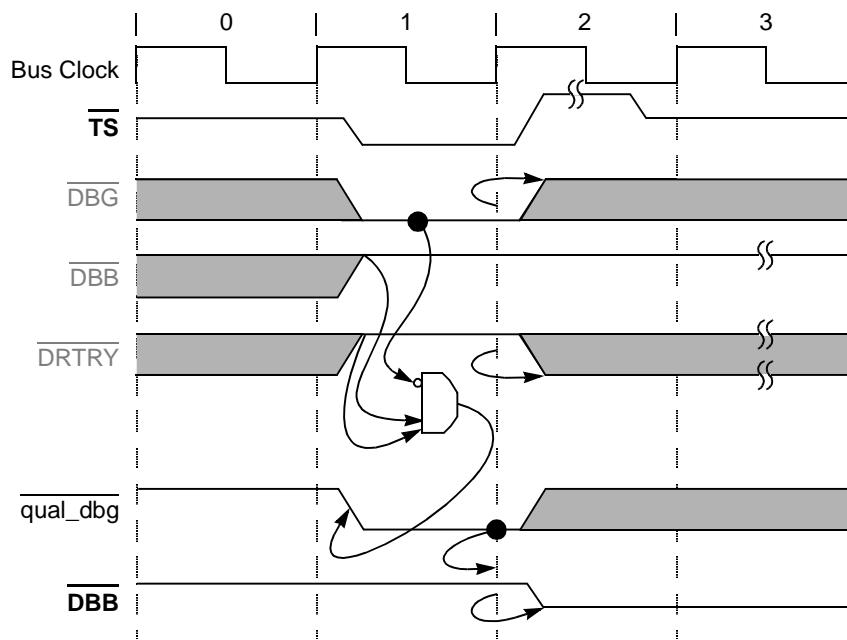


Figure 8-7. Data Bus Arbitration

8.4.1.1 Using the DBB Signal

The **DBB** signal should be connected between masters if data tenure scheduling is left to the masters. Optionally, the memory system can control data tenure scheduling directly with **DBG**. However, it is possible to ignore the **DBB** signal in the system if the **DBB** input is not used as the final data bus allocation control between data bus masters, and if the memory system can track the start and end of the data tenure. If **DBB** is not used to signal the end of a data tenure, **DBG** is only asserted to the next bus master the cycle before the cycle that the next bus master may actually begin its data tenure, rather than asserting it earlier (usually during another master's data tenure) and allowing the negation of **DBB** to be the final gating signal for a qualified data bus grant. Even if **DBB** is ignored in the system, the MPC603e always recognizes its own assertion of **DBB** and requires one cycle after data tenure completion to negate its own **DBB** before recognizing a qualified data bus grant for another data tenure. If **DBB** is ignored in the system, it must still be connected to a pull-up resistor on the MPC603e to ensure proper operation.

8.4.2 Data Bus Write Only

As a result of address pipelining, the MPC603e may have up to two data tenures queued to perform when it receives a qualified **DBG**. Generally, the data tenures should be performed in strict order (the same order) as their address tenures were performed. The MPC603e, however, also supports a limited out-of-order capability with the data bus write only (**DBWO**) input. When recognized on the clock of a qualified **DBG**, **DBWO** may direct the MPC603e to perform the next pending data write tenure even if a pending read tenure

would have normally been performed first. For more information on the operation of $\overline{\text{DBWO}}$, refer to Section 8.10, “Using $\overline{\text{DBWO}}$.”

If the MPC603e has any data tenures to perform, it always accepts data bus mastership to perform a data tenure when it recognizes a qualified $\overline{\text{DBG}}$. If $\overline{\text{DBWO}}$ is asserted with a qualified $\overline{\text{DBG}}$ and no write tenure is queued to run, the MPC603e still takes mastership of the data bus to perform the next pending read data tenure.

Generally, $\overline{\text{DBWO}}$ should only be used to allow a copy-back operation (burst write) to occur before a pending read operation. If $\overline{\text{DBWO}}$ is used for single-beat write operations, it may negate the effect of the **eieio** instruction by allowing a write operation to precede a program-scheduled read operation.

8.4.3 Data Transfer

The data transfer signals include DH[0:31], DL[0:31], DP[0:7], and $\overline{\text{DPE}}$. For memory accesses, the DH and DL signals form a 64-bit data path for read and write operations.

The MPC603e transfers data in either single- or four-beat burst transfers when configured with a 64-bit data bus; when configured with a 32-bit data bus, the MPC603e performs one-, two-, and eight-beat data transfers. Single-beat operations can transfer from 1 to 8 bytes at a time and can be misaligned; see Section 8.3.2.4, “Effect of Alignment in Data Transfers (64-Bit Bus).” Burst operations always transfer eight words and are aligned on eight-word address boundaries. Burst transfers can achieve significantly higher bus throughput than single-beat operations.

The type of transaction initiated by the MPC603e depends on whether the code or data is cacheable and, for store operations, whether the cache is considered in write-back or write-through mode, which software controls on either a page or block basis. Burst transfers support cacheable operations only; that is, memory structures must be marked as cacheable (and write-back for data store operations) in the respective page or block descriptor to take advantage of burst transfers.

The MPC603e output $\overline{\text{TBST}}$ indicates to the system whether the current transaction is a single- or four-beat transfer (except during **eciwx/ecowx** transactions, when it signals the state of EAR[28]). A burst transfer has an assumed address order. For load or store operations that miss in the cache (and are marked as cacheable and, for stores, write-back by the MMU), the MPC603e uses the double-word-aligned address associated with the critical code or data that initiated the transaction. This minimizes latency by allowing the critical code or data to be forwarded to the processor before the rest of the cache line is filled. For all other burst operations, however, the cache line is transferred beginning with the eight-word-aligned data.

The MPC603e does not directly support dynamic interfacing to subsystems with less than a 64-bit data path. It does, however, provide a static 32-bit data bus mode; for more information, see Section 8.1.2.1, “Optional 32-Bit Data Bus Mode.”

8.4.4 Data Transfer Termination

Four signals are used to terminate data bus transactions— $\overline{\text{TA}}$, $\overline{\text{DRTRY}}$ (data retry), $\overline{\text{TEA}}$ (transfer error acknowledge), and $\overline{\text{ARTRY}}$. The $\overline{\text{TA}}$ signal indicates normal termination of data transactions. It must always be asserted on the bus cycle coincident with the data that it is qualifying. It may be withheld by the slave for any number of clocks until valid data is ready to be supplied or accepted. $\overline{\text{DRTRY}}$ indicates invalid read data in the previous bus clock cycle. $\overline{\text{DRTRY}}$ extends the current data beat and does not terminate it. If it is asserted after the last (or only) data beat, the MPC603e negates $\overline{\text{DBB}}$ but still considers the data beat active and waits for another assertion of $\overline{\text{TA}}$. $\overline{\text{DRTRY}}$ is ignored on write operations. $\overline{\text{TEA}}$ indicates a nonrecoverable bus error event. Upon receiving a final (or only) termination condition, the MPC603e always negates $\overline{\text{DBB}}$ for one cycle.

If $\overline{\text{DRTRY}}$ is asserted by the memory system to extend the last (or only) data beat past the negation of $\overline{\text{DBB}}$, the memory system should three-state the data bus on the clock after the final assertion of $\overline{\text{TA}}$, even though it will negate $\overline{\text{DRTRY}}$ on that clock. This is to prevent a potential momentary data bus conflict if a write access begins on the following cycle.

The $\overline{\text{TEA}}$ signal is used to signal a nonrecoverable error during the data transaction. It may be asserted on any cycle during $\overline{\text{DBB}}$, or on the cycle after a qualified $\overline{\text{TA}}$ during a read operation, except when no- $\overline{\text{DRTRY}}$ mode is selected (where no- $\overline{\text{DRTRY}}$ mode cancels checking the cycle after $\overline{\text{TA}}$). The assertion of $\overline{\text{TEA}}$ terminates the data tenure immediately, even if in the middle of a burst; however, it does not prevent incorrect data that has just been acknowledged with $\overline{\text{TA}}$ from being written into the MPC603e cache or GPRs. The assertion of $\overline{\text{TEA}}$ initiates either a machine check exception or a checkstop condition based on the setting of the MSR.

An assertion of $\overline{\text{ARTRY}}$ causes the data tenure to be terminated immediately if the $\overline{\text{ARTRY}}$ is for the address tenure associated with the data tenure in operation. If $\overline{\text{ARTRY}}$ is connected for the MPC603e, the earliest allowable assertion of $\overline{\text{TA}}$ to the MPC603e is directly dependent on the earliest possible assertion of $\overline{\text{ARTRY}}$ to the MPC603e; see Section 8.3.3, “Address Transfer Termination.”

If the MPC603e clock is configured for 1:1 or 1.5:1 (processor clock to bus clock ratio) mode and the MPC603e is performing a burst read into its data cache, at least one wait state must be provided between the assertion of $\overline{\text{TS}}$ and the first assertion of $\overline{\text{TA}}$ for that transaction. If no- $\overline{\text{DRTRY}}$ mode is also selected, at least two wait states must be provided. The wait states are required due to possible resource contention in the data cache caused by a block replacement (or cast out) required in connection with the new linefill. These wait states may be provided by withholding the assertion of $\overline{\text{TA}}$ to the MPC603e for that data tenure, or by withholding $\overline{\text{DBG}}$ to the MPC603e, thereby delaying the start of the data tenure. This restriction applies only to burst reads into the data cache when configured in 1:1 or 1.5:1 clock modes. (It does not apply to instruction fetches, write operations, noncacheable read operations, or non-1:1 or non-1.5:1 clock modes.)

8.4.4.1 Normal Single-Beat Termination

Normal termination of a single-beat data read operation occurs when $\overline{\text{TA}}$ is asserted by a responding slave. The $\overline{\text{TEA}}$ and $\overline{\text{DRTRY}}$ signals must remain negated during the transfer (see Figure 8-8).

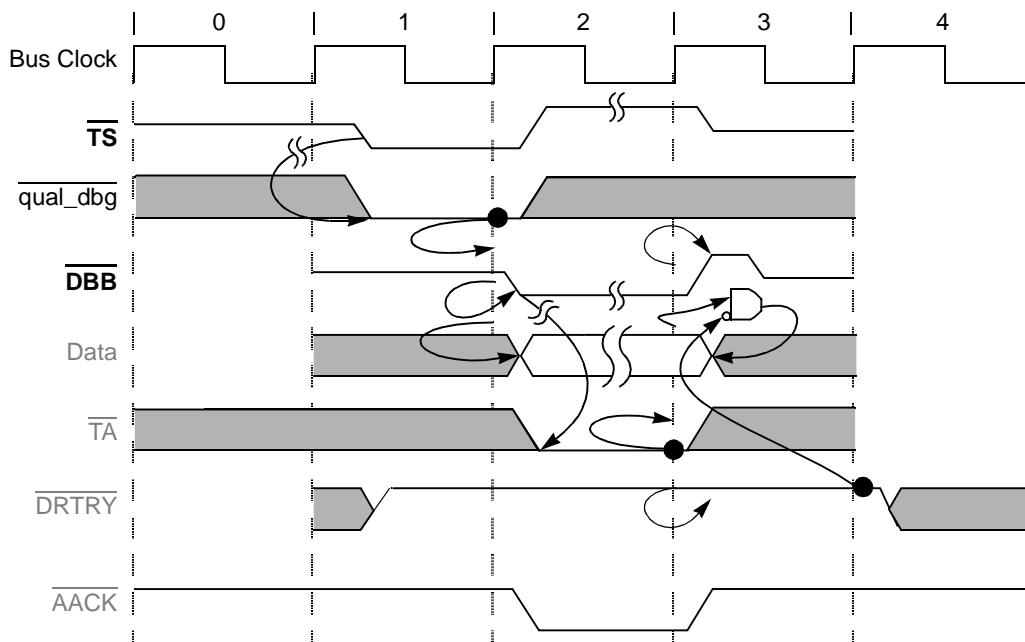


Figure 8-8. Normal Single-Beat Read Termination

The $\overline{\text{DRTRY}}$ signal is not sampled during data writes, as shown in Figure 8-9.

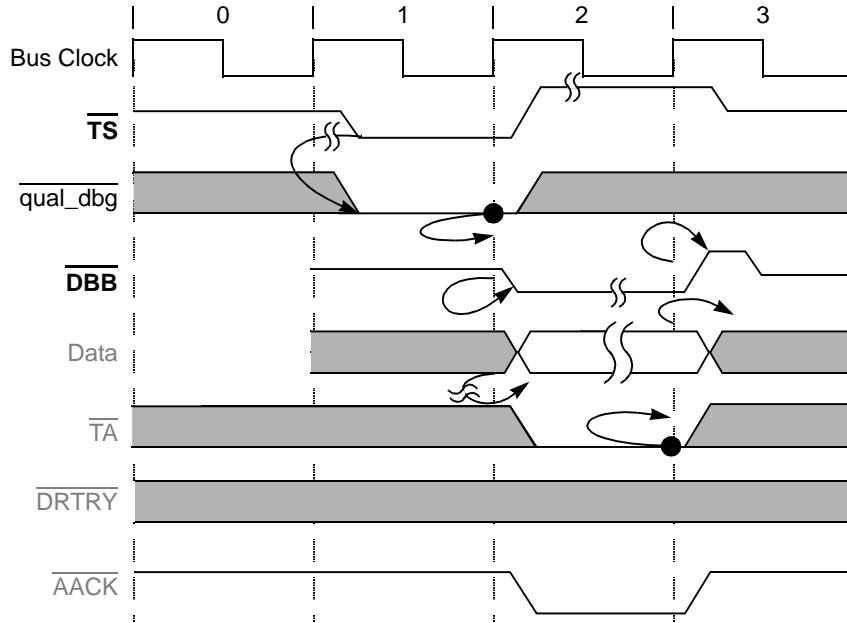


Figure 8-9. Normal Single-Beat Write Termination

8.4.4.2 Normal Burst Termination

Normal termination of a burst transfer occurs when $\overline{\text{TA}}$ is asserted for four bus clock cycles, as shown in Figure 8-10. The bus clock cycles in which $\overline{\text{TA}}$ is asserted need not be consecutive, thus allowing pacing of the data transfer beats. For read bursts to terminate successfully, $\overline{\text{TEA}}$ and $\overline{\text{DRTRY}}$ must remain negated during the transfer. For write bursts, $\overline{\text{TEA}}$ must remain negated for a successful transfer. $\overline{\text{DRTRY}}$ is ignored during data writes.

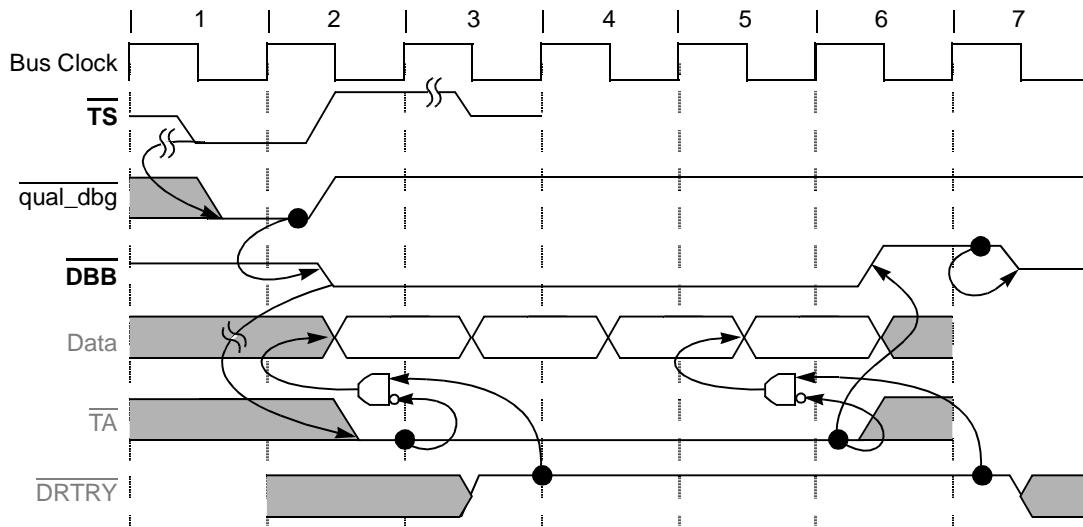


Figure 8-10. Normal Burst Transaction

For read bursts, $\overline{\text{DRTRY}}$ may be asserted one bus clock cycle after $\overline{\text{TA}}$ is asserted to signal that the data presented with $\overline{\text{TA}}$ is invalid and that the processor must wait for the negation of $\overline{\text{DRTRY}}$ before forwarding data to the processor (see Figure 8-11). Thus, a data beat can be terminated by a predicted branch with $\overline{\text{TA}}$ and then one bus clock cycle later confirmed with the negation of $\overline{\text{DRTRY}}$. The $\overline{\text{DRTRY}}$ signal is valid only for read transactions. $\overline{\text{TA}}$ must be asserted on the bus clock cycle before the first bus clock cycle of the assertion of $\overline{\text{DRTRY}}$; otherwise the results are undefined.

The $\overline{\text{DRTRY}}$ signal extends data bus mastership such that other processors cannot use the data bus until $\overline{\text{DRTRY}}$ is negated. Therefore, in the example shown in Figure 8-11, $\overline{\text{DBB}}$ cannot be asserted until bus clock cycle 5. This is true for both read and write operations even though $\overline{\text{DRTRY}}$ does not extend bus mastership for write operations.

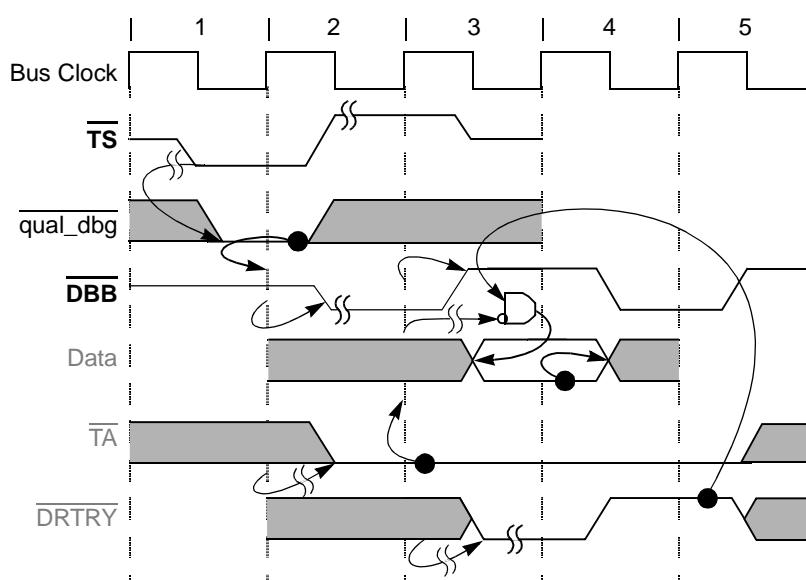


Figure 8-11. Termination with $\overline{\text{DRTRY}}$

Figure 8-12 shows the effect of using $\overline{\text{DRTRY}}$ during a burst read. It also shows the effect of using $\overline{\text{TA}}$ to pace the data transfer rate. Notice that in bus clock cycle 3 of Figure 8-12, $\overline{\text{TA}}$ is negated for the second data beat. The MPC603e data pipeline does not proceed until bus clock cycle 4 when $\overline{\text{TA}}$ is reasserted.

Note that $\overline{\text{DRTRY}}$ is useful for systems that implement predicted forwarding of data such as those with direct-mapped, second-level caches where hit/miss is determined on the following bus clock cycle, or for parity- or ECC-checked memory systems.

Note that $\overline{\text{DRTRY}}$ may not be implemented on other processors.

8.4.4.3 Data Transfer Termination Due to a Bus Error

The $\overline{\text{TEA}}$ signal indicates that a bus error occurred. It may be asserted while $\overline{\text{DBB}}$ (and/or $\overline{\text{DRTRY}}$ for read operations) is asserted. Asserting $\overline{\text{TEA}}$ to the MPC603e terminates the transaction; that is, further assertions of $\overline{\text{TA}}$ and $\overline{\text{DRTRY}}$ are ignored and $\overline{\text{DBB}}$ is negated; see Figure 8-12.

Assertion of the $\overline{\text{TEA}}$ signal causes a machine check exception (and possibly a checkstop condition within the MPC603e). For more information, see Section 4.5.2, “Machine Check Exception (0x00200).” Note also that the MPC603e does not implement a synchronous error capability for memory accesses. This means that the exception instruction pointer does not point to the memory operation that caused the assertion of $\overline{\text{TEA}}$, but to the instruction about to be executed (perhaps several instructions later). However, assertion of $\overline{\text{TEA}}$ does not invalidate data entering the GPR or the cache. Additionally, the corresponding address of the access that caused $\overline{\text{TEA}}$ to be asserted is not latched by the MPC603e. To recover, the exception handler must determine and remedy the cause of the

TEA, or the MPC603e must be reset; therefore, this function should only be used to flag fatal system conditions to the processor (such as parity or uncorrectable ECC errors).

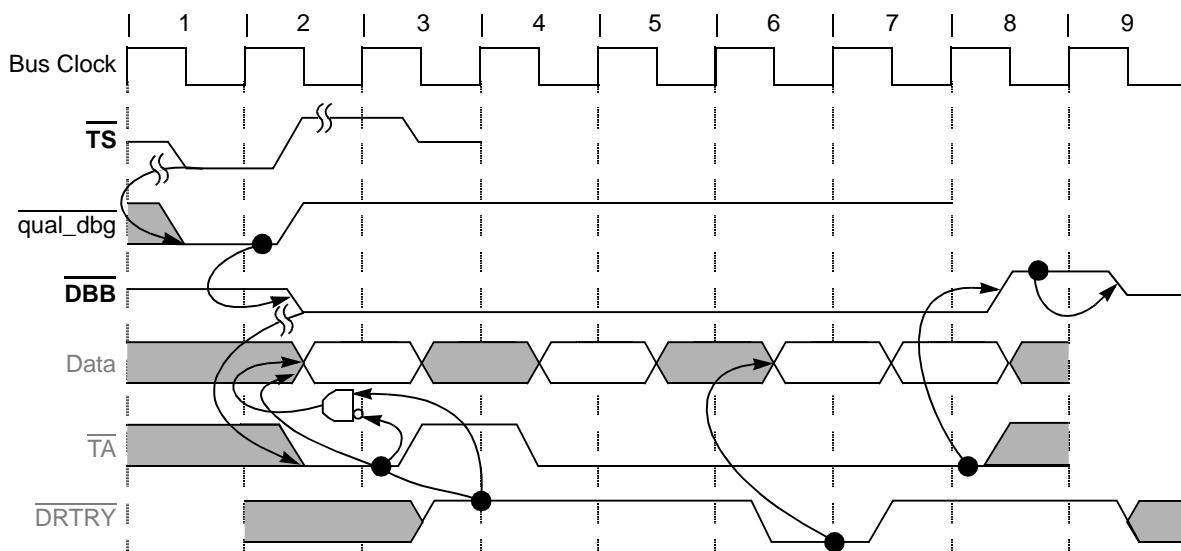


Figure 8-12. Read Burst with TA Wait States and DRTRY

After the MPC603e has committed to run a transaction, that transaction must eventually complete. Address retry causes the transaction to be restarted; **TA** wait states and **DRTRY** assertion for reads delay termination of individual data beats. Eventually, however, the system must either terminate the transaction or assert the **TEA** signal (and vector the MPC603e into a machine check exception.) For this reason, care must be taken to check for the end of physical memory and the location of certain system facilities to avoid memory accesses that result in the generation of machine check exceptions.

Note that **TEA** generates a machine check exception depending on MSR[ME]. Clearing the machine check exception enable control bits leads to a true checkstop condition (instruction execution halted and processor clock stopped).

8.4.5 Memory Coherency—MEI Protocol

The MPC603e provides dedicated hardware to provide memory coherency by snooping bus transactions. The address retry capability enforces the three-state, MEI cache-coherency protocol (see Figure 8-13).

The global (**GBL**) output signal indicates whether the current transaction must be snooped by other snooping devices on the bus. Address bus masters assert **GBL** to indicate that the current transaction is a global access (that is, an access to memory shared by more than one device). If **GBL** is not asserted for the transaction, that transaction is not snooped. When other devices detect the **GBL** input asserted, they must respond by snooping the broadcast address.

Normally, \overline{GBL} reflects the M-bit value specified for the memory reference in the corresponding translation descriptor. Note that care must be taken to minimize the number of pages marked as global, because the retry protocol discussed in the previous section is used to enforce coherency and can require significant bus bandwidth.

When the MPC603e is not the address bus master, \overline{GBL} is an input. The MPC603e snoops a transaction if \overline{TS} and \overline{GBL} are asserted together in the same bus clock cycle (this is a qualified snooping condition). No snoop update to the MPC603e cache occurs if the snooped transaction is not marked global. This includes invalidation cycles.

When the MPC603e detects a qualified snoop condition, the address associated with the \overline{TS} is compared against the data cache tags. Snooping completes if no hit is detected. However, if the address hits in the cache, the MPC603e reacts according to the MEI protocol shown in Figure 8-13, assuming the WIM bits are set to write-back, caching-allowed, and coherency-enforced modes (WIM = 001).

The MPC603e on-chip data cache is implemented as a four-way set-associative cache. To facilitate external monitoring of the internal cache tags, the cache set entry (CSE[0:1]) signals indicate which cache set is being replaced on read operations. Note that these signals are valid only for MPC603e burst operations; for all other bus operations, CSE[0:1] should be ignored.

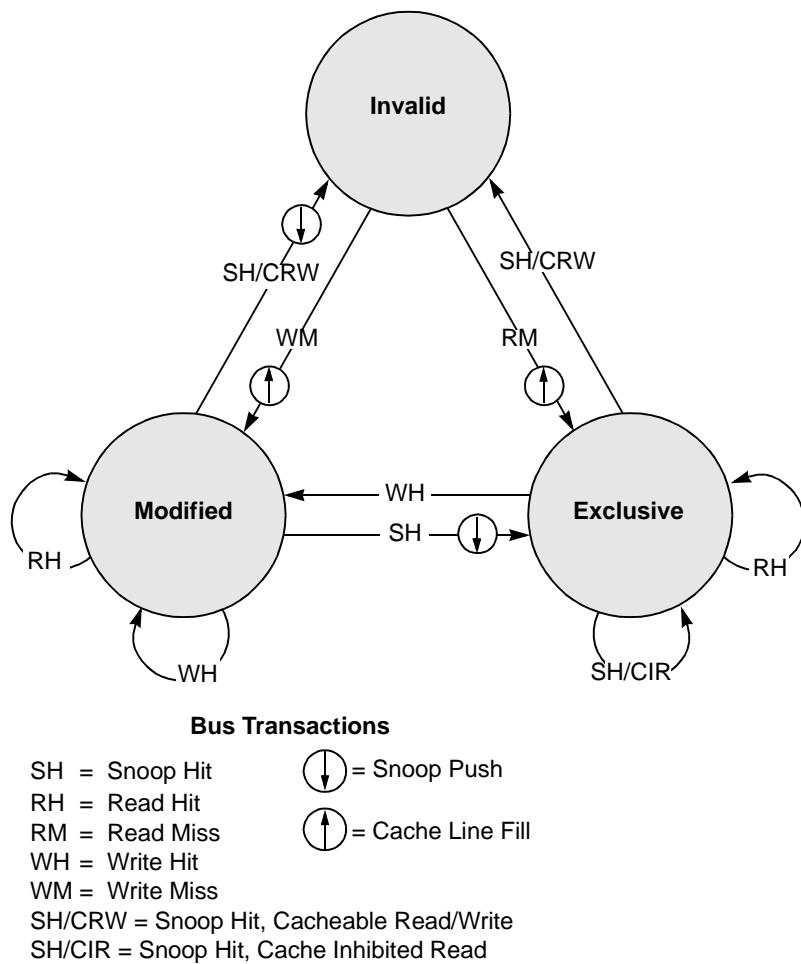


Figure 8-13. MEI Cache Coherency Protocol—State Diagram (WIM = 001)

Table 8-10 shows the CSE encodings.

Table 8-10. CSE[0:1] Signals

CSE[0:1]	Cache Set Element
00	Set 0
01	Set 1
10	Set 2
11	Set 3

8.5 Timing Examples

This section shows timing diagrams for various scenarios. Figure 8-14 illustrates the fastest single-beat reads possible for the MPC603e. This figure shows both minimal latency and maximum single-beat throughput. By delaying the data bus tenure, the latency increases, but, because of split-transaction pipelining, the overall throughput is not affected unless the data bus latency causes the third address tenure to be delayed.

Note that all bidirectional signals are three-stated between bus tenures.

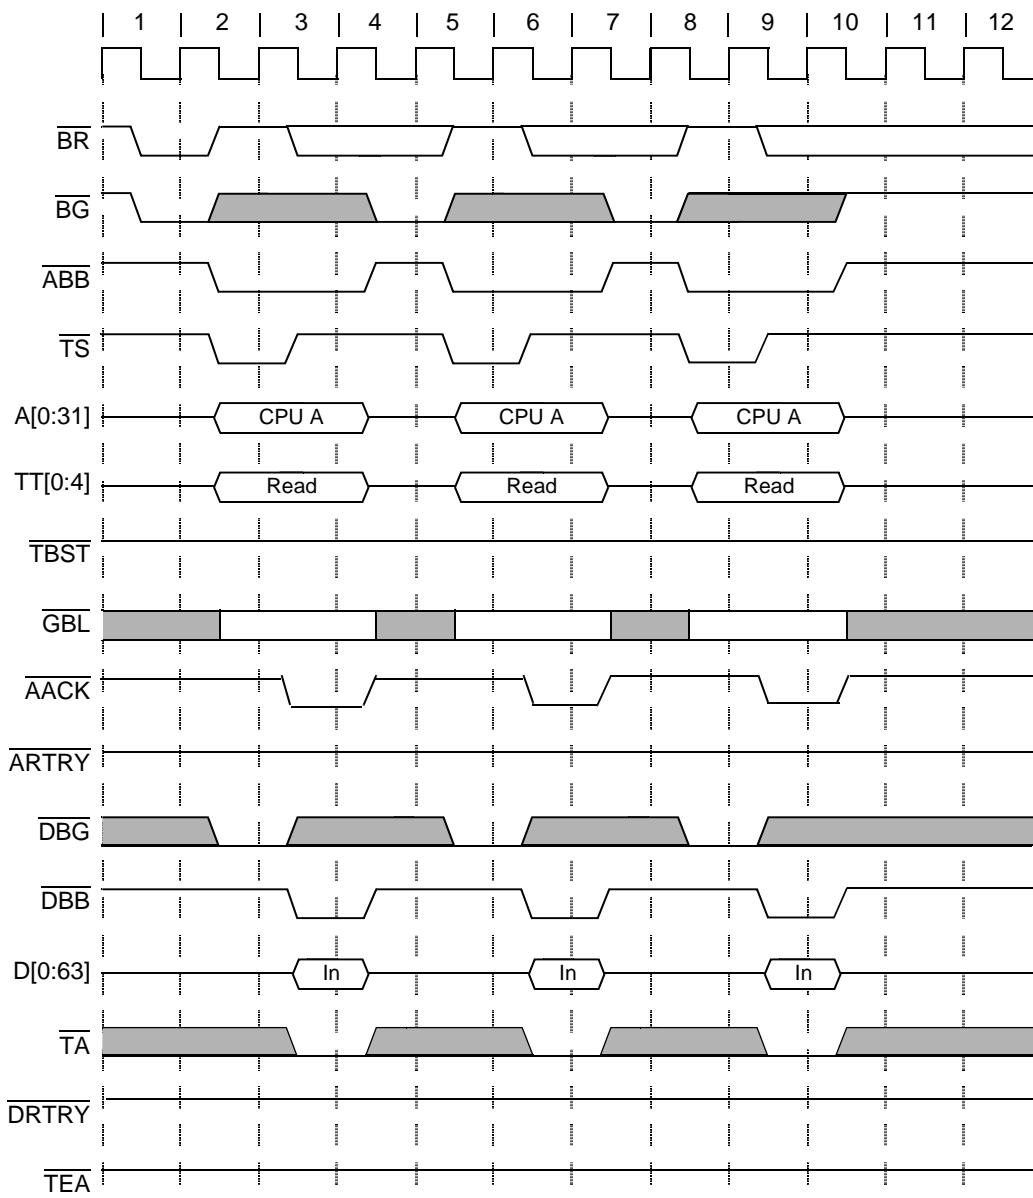


Figure 8-14. Fastest Single-Beat Reads

Figure 8-15 illustrates the fastest single-beat writes supported by the MPC603e. All bidirectional signals are three-stated between bus tenures.

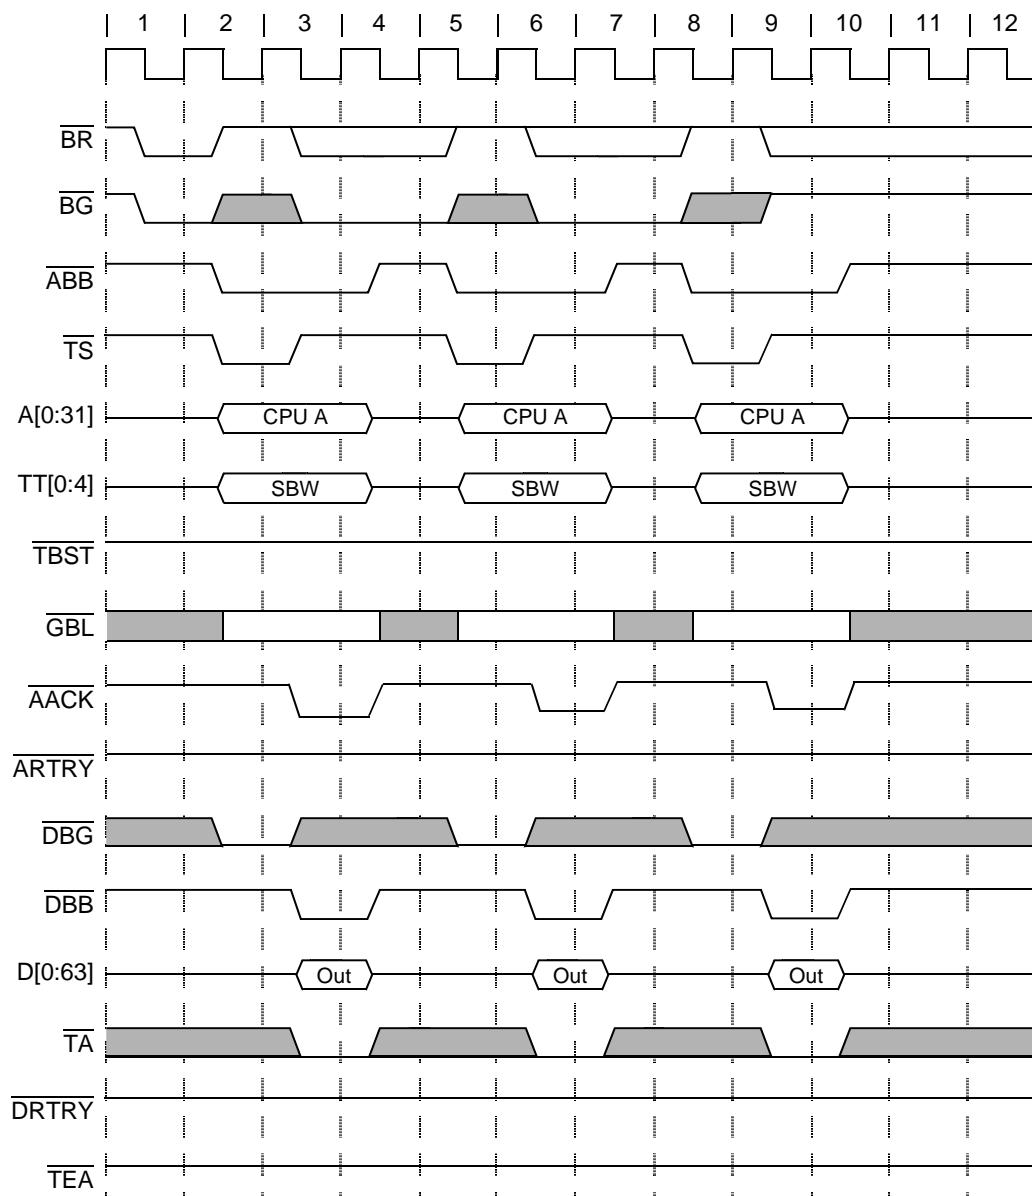


Figure 8-15. Fastest Single-Beat Writes

Figure 8-16 shows three ways to delay single-beat reads showing data-delay controls:

- The \overline{TA} signal can remain negated to insert wait states in clock cycles 3 and 4.
- For the second access, \overline{DBG} could have been asserted in clock cycle 6.
- In the third access, \overline{DRTRY} is asserted in clock cycle 11 to flush the previous data.

Note that all bidirectional signals are three-stated between bus tenures. The pipelining shown in Figure 8-16 can occur if the second access is not another load (for example, an instruction fetch).

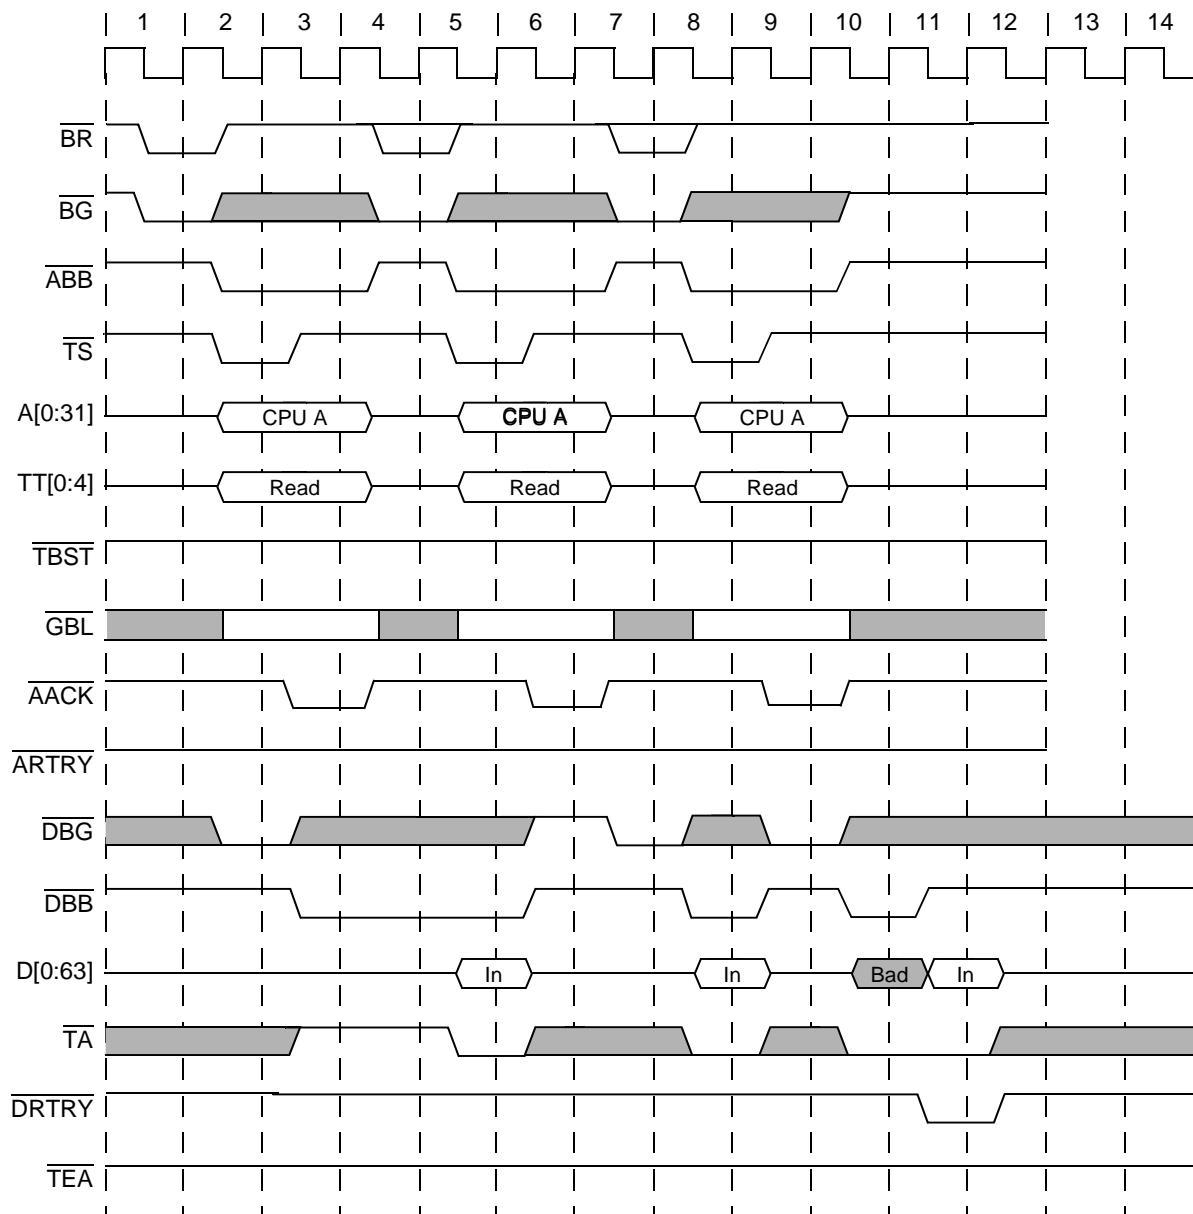


Figure 8-16. Single-Beat Reads Showing Data-Delay Controls

Figure 8-17 shows data-delay controls in a single-beat write operation. Note that all bidirectional signals are three-stated between bus tenures. Data transfers are delayed in the following ways:

- The **TA** signal is held negated to insert wait states in clocks 3 and 4.
- In clock 6, **DBG** is held negated, delaying the start of the data tenure.

The last access is not delayed (**DRTRY** is valid only for read operations).

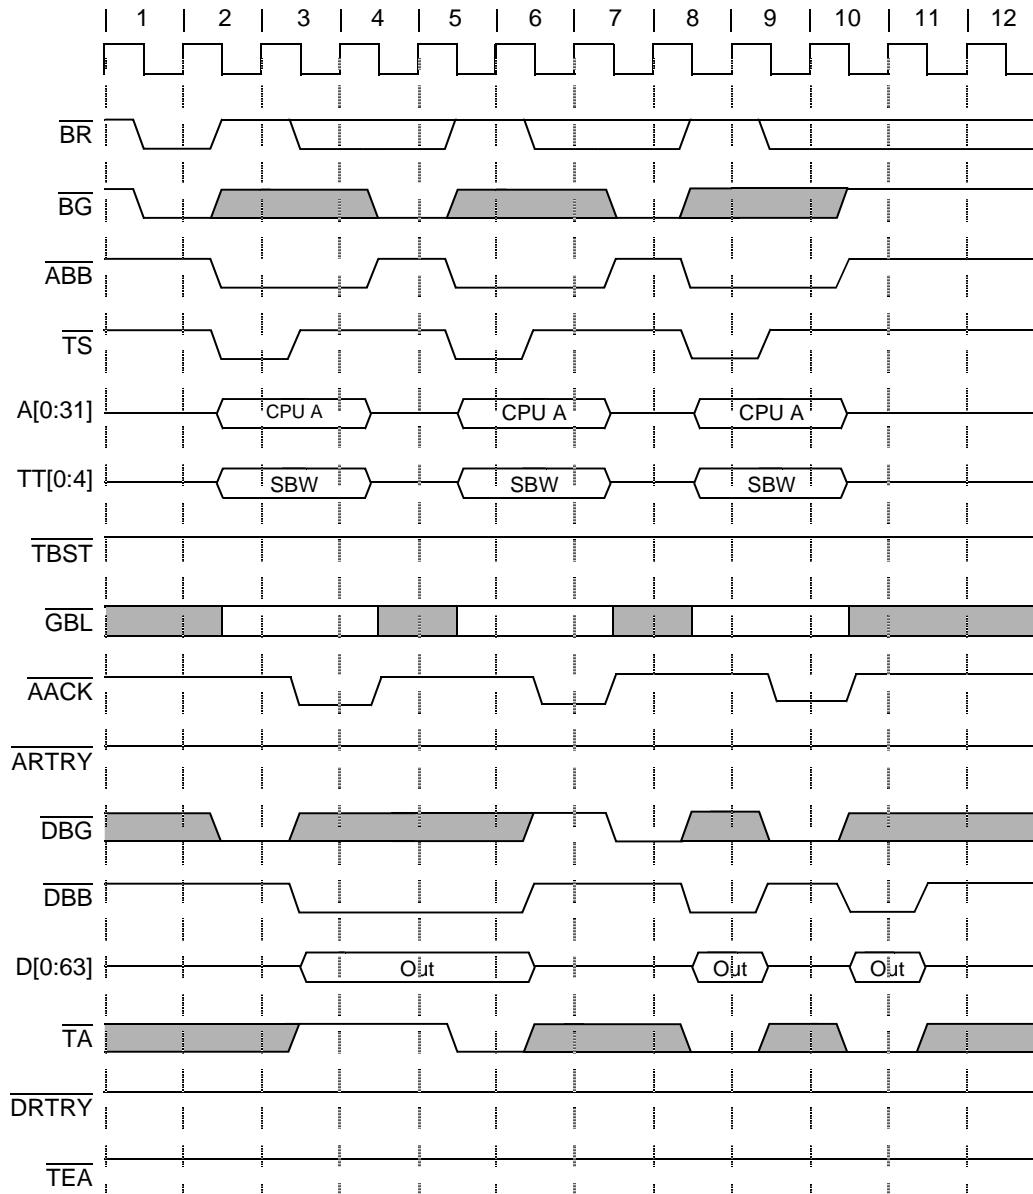


Figure 8-17. Single-Beat Writes Showing Data-Delay Controls

Figure 8-18 shows the use of data-delay controls with burst transfers. Note that all bidirectional signals are three-stated between bus tenures. Note the following:

- The first data beat of bursted read data (clock 0) is the critical quad word.
- The write burst shows the use of \overline{TA} signal negation to delay the third data beat.
- The final read burst shows the use of \overline{DRTRY} on the third data beat.
- The address for the third transfer is delayed until the first transfer completes.

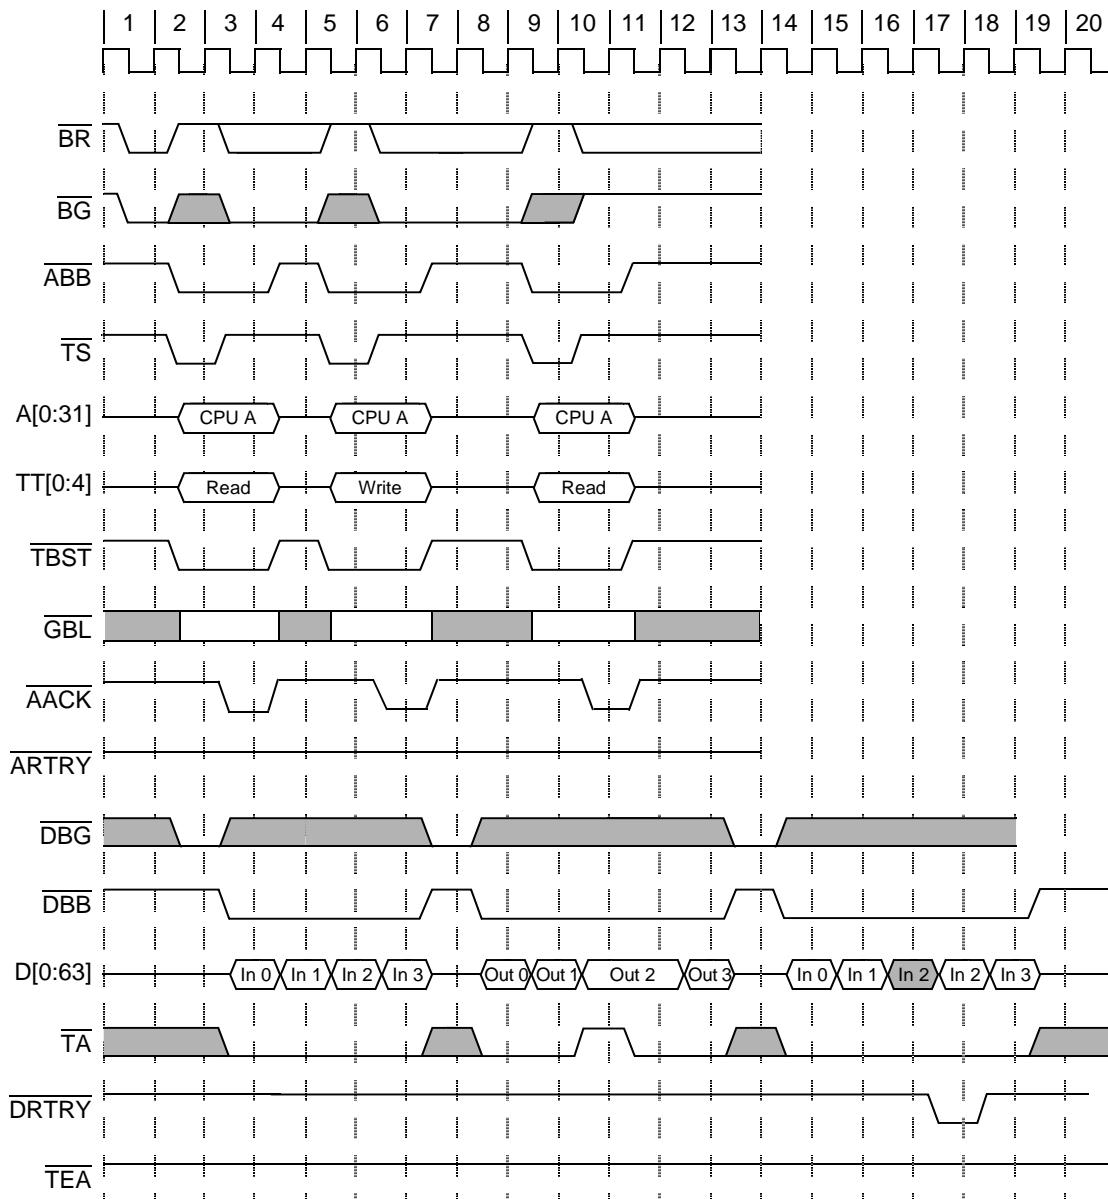


Figure 8-18. Burst Transfers with Data-Delay Controls

Figure 8-19 shows the use of the $\overline{\text{TEA}}$ signal. Note that all bidirectional signals are three-stated between bus tenures. Note the following:

- The first data beat of the read burst (in clock 0) is the critical quad word.
- The $\overline{\text{TEA}}$ signal truncates the burst write transfer on the third data beat.
- The MPC603e eventually causes an exception to be taken on the $\overline{\text{TEA}}$ event.

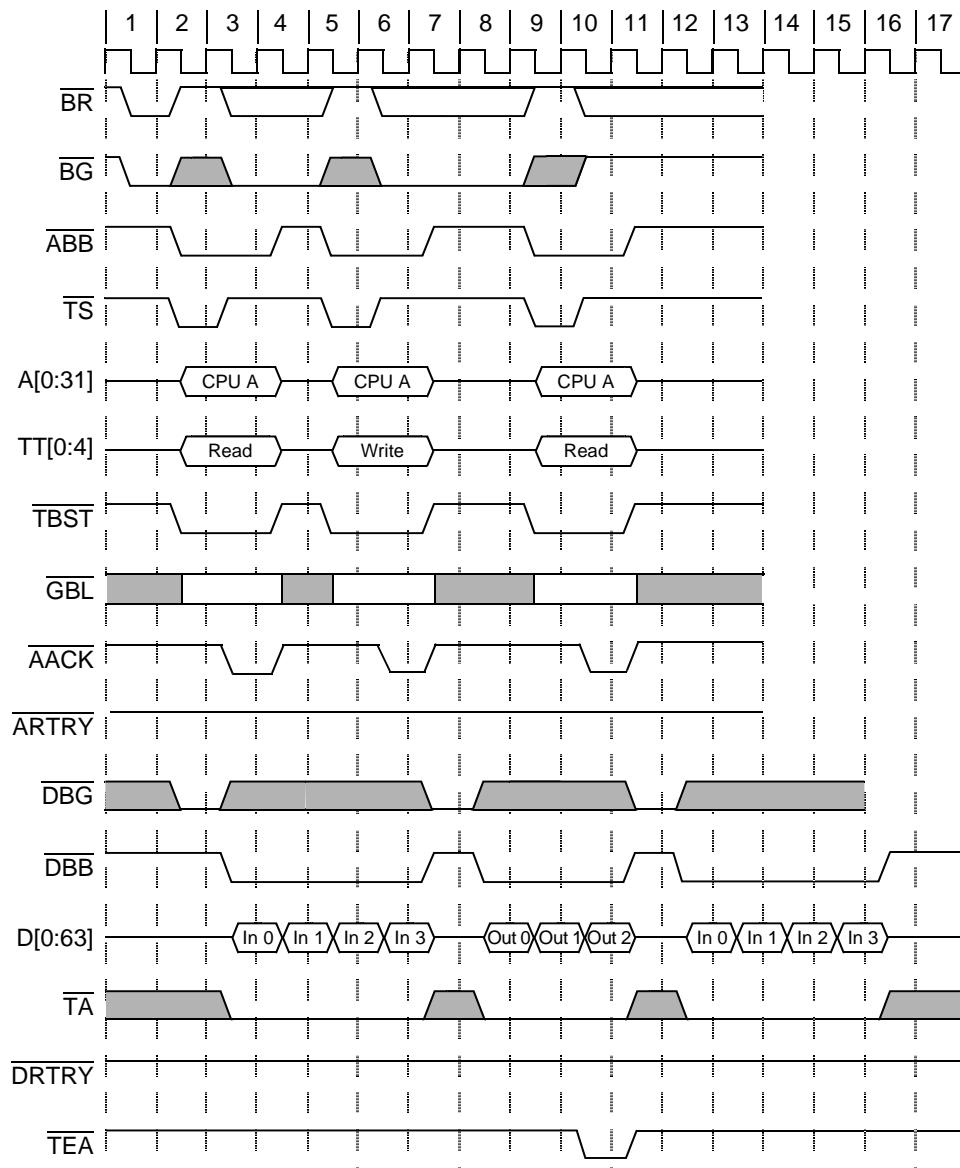


Figure 8-19. Use of Transfer Error Acknowledge (TEA)

8.6 Optional Bus Configurations

The MPC603e supports three optional bus configurations that are selected by the assertion or negation of $\overline{\text{DRTRY}}$, $\overline{\text{TLBISYNC}}$, and $\overline{\text{QACK}}$ during the negation of $\overline{\text{HRESET}}$. The operation and selection of the optional bus configurations are described in the following sections.

8.6.1 32-Bit Data Bus Mode

The MPC603e supports an optional 32-bit data bus mode, which differs from the 64-bit data bus mode only in the byte lanes involved in the transfer and the number of data beats performed. When in 32-bit data bus mode, only byte lanes 0 through 3 are used corresponding to DH0:DH31 and DP0:DP3. Byte lanes 4 through 7 corresponding to DL0:DL31 and DP4:DP7 are never used in this mode. The unused data bus signals are not sampled by the MPC603e during read operations, and they are driven low during write operations.

A data tenure in the 32-bit data bus mode takes one, two, or eight beats depending on the transfer size and the cache mode for the address. Data transactions of one or two data beats are performed for caching-inhibited load/store or write-through store operations. These transactions do not assert the $\overline{\text{TBST}}$ signal even though a two-beat burst may be performed (having the same $\overline{\text{TBST}}$ and TSIZ[0:2] encodings as the 64-bit data bus mode). Single-beat data transactions are performed for bus operations of 4 bytes or less, and double-beat data transactions are performed for 8-byte operations only. The MPC603e only generates an 8-byte operation for a double-word-aligned load or store double operation to or from the floating-point GPRs. All cache-inhibited instruction fetches are performed as word (single-beat) operations.

Data transactions of eight data beats are performed for burst operations that load into or store from the MPC603e internal caches. These transactions transfer 32 bytes in the same way as in 64-bit data bus mode, asserting the $\overline{\text{TBST}}$ signal, and signaling a transfer size of 2 (TSIZ[0:2] = 0b010).

The same bus protocols apply for arbitration, transfer, and termination of the address and data tenures in the 32-bit data bus mode as they apply to the 64-bit data bus mode. Late $\overline{\text{ARTRY}}$ cancellation of the data tenure applies on the bus clock after the first data beat is acknowledged (after the first $\overline{\text{TA}}$) for word or smaller transactions, or on the bus clock after the second data beat is acknowledged (after the second $\overline{\text{TA}}$) for double-word or burst operations (or coincident with respective $\overline{\text{TA}}$ if no- $\overline{\text{DRTRY}}$ mode is selected).

An example of an eight-beat data transfer while the MPC603e is in 32-bit data bus mode is shown in Figure 8-20.

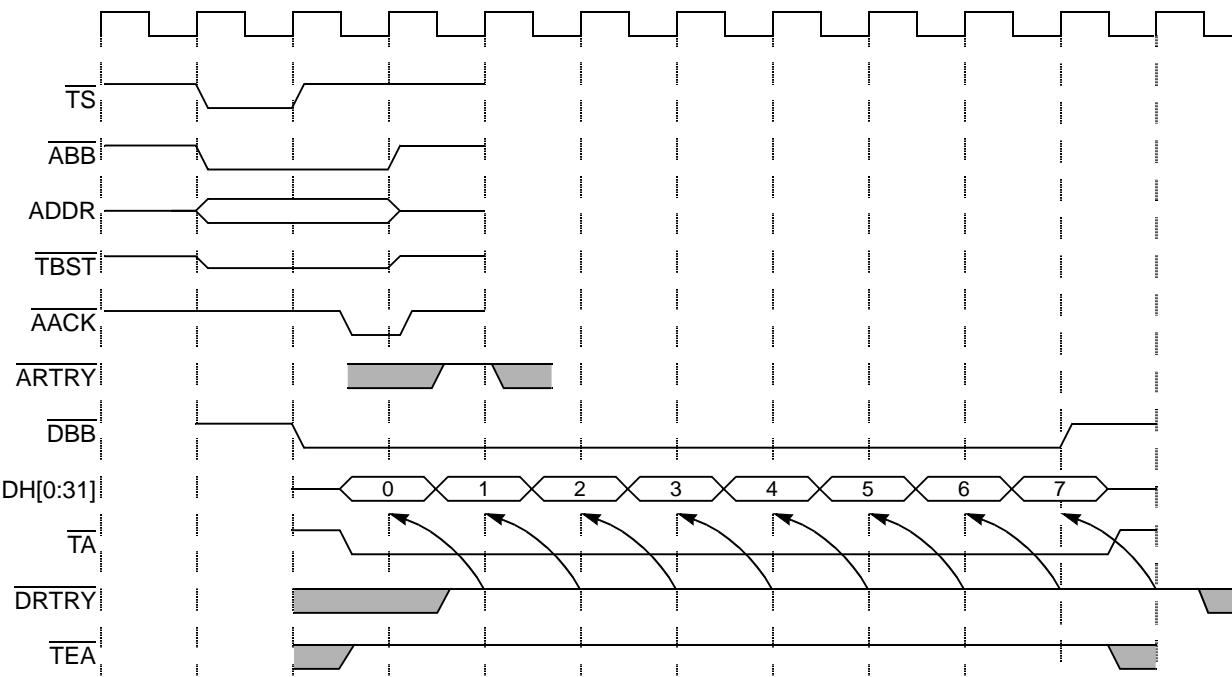


Figure 8-20. 32-Bit Data Bus Transfer (Eight-Beat Burst)

An example of a two-beat data transfer (with **DRTRY** asserted during each data tenure) is shown in Figure 8-21.

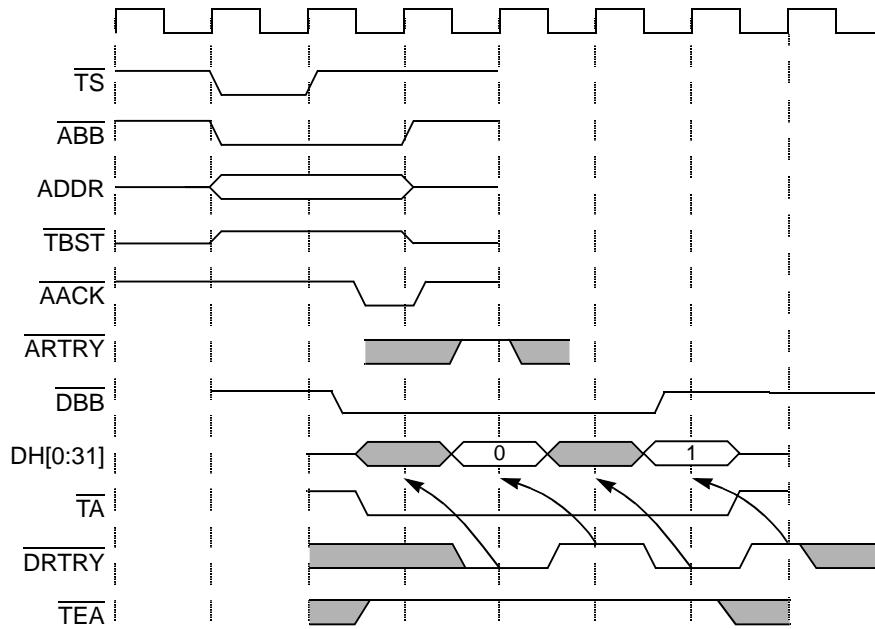


Figure 8-21. 32-Bit Data Bus Transfer (Two-Beat Burst with DRTRY)

The MPC603e selects a 64- or 32-bit data bus mode at startup by sampling the state of the **TLBISYNC** signal at the negation of **HRESET**. If the **TLBISYNC** signal is negated at the

negation of $\overline{\text{HRESET}}$, 64-bit data mode is entered by the MPC603e. If $\overline{\text{TLBISYNC}}$ is asserted at the negation of $\overline{\text{HRESET}}$, 32-bit data mode is entered.

8.6.2 No-DRTRY Mode

The MPC603e supports an optional mode to disable the use of the data retry function provided through $\overline{\text{DRTRY}}$. The no- $\overline{\text{DRTRY}}$ mode allows the forwarding of data during load operations to the processor core one bus cycle sooner than in the normal bus protocol.

The bus protocol specifies that, during load operations, the memory system can normally cancel data that was read by the master on the bus cycle after $\overline{\text{TA}}$ was asserted. This late cancellation protocol requires the MPC603e to hold any loaded data at the bus interface for one additional bus clock to verify that the data is valid before forwarding it to the processor core. For systems that do not implement the $\overline{\text{DRTRY}}$ function, the MPC603e provides an optional no- $\overline{\text{DRTRY}}$ mode that eliminates this one-cycle stall during all load operations, and allows for the forwarding of data to the internal CPU immediately when $\overline{\text{TA}}$ is recognized.

When the MPC603e is in no- $\overline{\text{DRTRY}}$ mode, data can no longer be canceled the cycle after it is acknowledged by an assertion of $\overline{\text{TA}}$. Data is immediately forwarded to the processor core, and any attempt at late cancellation by the system may cause improper operation by the MPC603e.

When the MPC603e is following normal bus protocol, data may be canceled the bus cycle after $\overline{\text{TA}}$ by either of two means—late cancellation by $\overline{\text{DRTRY}}$, or late cancellation by $\overline{\text{ARTRY}}$. When no- $\overline{\text{DRTRY}}$ mode is selected, both cancellation cases must be disallowed in the system design for the bus protocol.

When no- $\overline{\text{DRTRY}}$ mode is selected, the system must ensure that $\overline{\text{DRTRY}}$ is not asserted to the MPC603e which may cause improper operation of the bus interface. The system must also ensure that an assertion of $\overline{\text{ARTRY}}$ by a snooping device must occur before or coincident with the first assertion of $\overline{\text{TA}}$ to the MPC603e, but not on the cycle after the first assertion of $\overline{\text{TA}}$.

Other than the inability to cancel data that was read by the master on the bus cycle after $\overline{\text{TA}}$ was asserted, the bus protocol for the MPC603e is identical to that for the basic transfer bus protocols described in this chapter, as well as for 32-bit data bus mode.

The MPC603e selects the desired $\overline{\text{DRTRY}}$ mode at startup by sampling the state of the $\overline{\text{DRTRY}}$ signal itself at the negation of $\overline{\text{HRESET}}$. If $\overline{\text{DRTRY}}$ is negated at the negation of $\overline{\text{HRESET}}$, normal operation is selected. If $\overline{\text{DRTRY}}$ is asserted at the negation of $\overline{\text{HRESET}}$, no- $\overline{\text{DRTRY}}$ mode is selected.

8.6.3 Reduced-Pinout Mode

The MPC603e provides an optional reduced-pinout mode which idles the switching of numerous signals for reduced power consumption. The DL[0:31], DP[0:7], AP[0:3], \overline{APE} , \overline{DPE} , and \overline{RSRV} signals are disabled when the reduced-pinout mode is selected. Note that the 32-bit data bus mode is implicitly selected when the reduced-pinout mode is enabled.

In reduced-pinout mode, the bidirectional and output signals disabled are always driven low during the periods when they would normally have been driven by the MPC603e. The open-drain outputs (\overline{APE} and \overline{DPE}) are always three-stated. The bidirectional inputs are always turned-off at the input receivers of the MPC603e and are not sampled.

The MPC603e selects either full-pinout or reduced-pinout mode at startup by sampling the state of the \overline{QACK} signal at the negation of \overline{HRESET} . If \overline{QACK} is asserted at the negation of \overline{HRESET} , full-pinout mode is selected by the MPC603e. If \overline{QACK} is negated at the negation of \overline{HRESET} , reduced-pinout mode is selected.

8.7 Interrupt, Checkstop, and Reset Signals

This section describes external interrupts, checkstop operations, and hard and soft reset inputs.

8.7.1 External Interrupts

Asserting the external interrupt input signals (\overline{INT} , \overline{SMI} , and \overline{MCP}) of the MPC603e eventually forces the processor to take the external interrupt vector, or the system management interrupt vector if the MSR[EE] is set, or the machine check interrupt if MSR[ME] and HID0[EMCP] are set.

8.7.2 Checkstops

Asserting the MPC603e has two checkstop input signals— $\overline{CKSTP_IN}$ (non-maskable) and \overline{MCP} (enabled when MSR[ME] is cleared and HID0[EMCP] is set), and a checkstop output ($\overline{CKSTP_OUT}$). If $\overline{CKSTP_IN}$ or \overline{MCP} is asserted, the MPC603e halts operations by gating off all internal clocks. The MPC603e asserts $\overline{CKSTP_OUT}$ if $\overline{CKSTP_IN}$ is asserted.

If $\overline{CHECKSTOP}$ is asserted by the MPC603e, it has entered the checkstop state and processing has halted internally. The $\overline{CHECKSTOP}$ signal can be asserted for various reasons including receiving a \overline{TEA} signal and detection of external parity errors. For more information about checkstop state, see Section 4.5.2.2, “Checkstop State (MSR[ME] = 0).”

8.7.3 Reset Inputs

The MPC603e has two reset inputs, described as follows:

- **HRESET** (hard reset)—**HRESET** is used for power-on reset sequences, or for situations in which the MPC603e must go through the entire cold-start sequence of internal hardware initializations.
- **SRESET** (soft reset)—The soft reset input provides warm reset capability. This input can be used to avoid forcing the MPC603e to complete the cold start sequence.

When either reset input is negated, the processor attempts to fetch code from the system reset exception vector. The vector is located at offset 0x00100 from the exception prefix (all zeros or ones, depending on the setting of the exception prefix bit in the machine state register (MSR[IP])). The IP bit is set for **HRESET**.

8.7.4 System Quiesce Control Signals

The system quiesce control signals (**QREQ** and **QACK**) allow the processor to enter a low power state, and bring bus activity to a quiescent state in an orderly fashion.

The system quiesce state is entered by configuring the processor to assert the **QREQ** output. This signal allows the system to terminate or pause any bus activities that are normally snooped. When the system is ready to enter the system quiesce state, it asserts **QACK**. At this time, the MPC603e may enter a quiescent (low-power) state during which it stops snooping bus activity.

8.8 Processor State Signals

This section describes the MPC603e support for atomic update and memory through the use of the **Iwarx/stwcx.** instruction pair and includes a description of the MPC603e **TLBISYNC** input.

8.8.1 Support for the Iwarx/stwcx. Instruction Pair

The Load Word and Reserve Indexed (**Iwarx**) and the Store Word Conditional Indexed (**stwcx.**) instructions provide a means for atomic memory updating. Memory can be updated atomically by setting a reservation on the load and checking that the reservation is still valid before the store is performed. In the MPC603e, the reservations are made on behalf of aligned, 32-byte sections of the memory address space.

The reservation (**RSRV**) output signal is driven synchronously with the bus clock and reflects the status of the reservation coherency bit in the reservation address buffer (see Section 3.9, “Instruction and Data Cache Operation” for more information). See Section 7.2.9.7.3, “Reservation (**RSRV**)—Output,” for information about timing.

8.8.2 TLBISYNC Input

The **TLBISYNC** input allows for the hardware synchronization of changes to MMU tables when the MPC603e and another DMA master share the same MMU translation tables in system memory. It is asserted by a DMA master when it is using shared addresses that could be changed in the MMU tables by the MPC603e during the DMA master's tenure.

Asserting the **TLBISYNC** input to the MPC603e prevents it from completing any instructions past a **tlbsync** instruction. Generally, during the execution of an **eciwx** or **ecowx** instruction by the MPC603e, the selected DMA device should assert the MPC603e **TLBISYNC** signal and keep it asserted during its DMA tenure if it is using a shared translation address. Subsequent instructions should include a **sync** and **tlbsync** instruction before any MMU table changes are performed. This prevents the MPC603e from making table changes disruptive to the other master during the DMA period.

8.9 IEEE 1149.1-Compliant Interface

The MPC603e boundary-scan interface is a fully-compliant implementation of the IEEE 1149.1 standard. This section describes the MPC603e IEEE 1149.1 (JTAG) interface.

8.9.1 IEEE 1149.1 Interface Description

The MPC603e has five dedicated JTAG signals (described in Table 8-11). The TDI and TDO scan ports are used to scan instructions, as well as data, into the various scan registers for JTAG operations. The scan operation is controlled by the test access port (TAP) controller, which in turn is controlled by the TMS input sequence. The scan data is latched in at the rising edge of TCK.

Table 8-11. IEEE Interface Pin Descriptions

Signal Name	Input/Output	Weak Pullup Provided	IEEE 1149.1 Function
TDI	Input	Yes	Serial scan input signal
TDO	Output	No	Serial scan output signal
TMS	Input	Yes	TAP controller mode signal
TCK	Input	Yes	Scan clock
TRST	Input	Yes	TAP controller reset

TRST is a JTAG optional signal used to reset the TAP controller asynchronously. The **TRST** signal assures that the JTAG logic does not interfere with the normal operation of the chip, and can be asserted coincident with the **HRESET**.

The PID7t-603e implements the JTAG/COP in the same manner as does the PID6-603e implementation with the exception of the introduction of the 33-bit Run_N counter register

in which the most-significant 32 bits form a 32-bit counter. The function of the least-significant bit remains unchanged. The Run_N counter is used by the COP to control the number of processor cycles that the processor runs before halting.

8.10 Using DBWO

The MPC603e supports split-transaction pipelined transactions. It supports a limited out-of-order capability for its own pipelined transactions through the $\overline{\text{DBWO}}$ signal. When recognized on the clock of a qualified $\overline{\text{DBG}}$, the assertion of $\overline{\text{DBWO}}$ directs the MPC603e to perform the next pending data write tenure (if any), even if a pending read tenure would have normally been performed because of address pipelining. The $\overline{\text{DBWO}}$ signal does not reorder write tenures with respect to other write tenures from the same MPC603e. It only allows that a write tenure be performed ahead of a pending read tenure from the same MPC603e.

In general, an address tenure on the bus is followed strictly in order by its associated data tenure. Transactions pipelined by the MPC603e complete strictly in order. However, the MPC603e can run bus transactions out of order only when the external system allows the MPC603e to perform a cache-line-snoop-push-out operation (or other write transaction, if pending in the MPC603e write queues) between the address and data tenures of a read operation through the use of $\overline{\text{DBWO}}$. This effectively envelopes the write operation within the read operation. Figure 8-22 shows how $\overline{\text{DBWO}}$ is used to perform an enveloped write transaction.

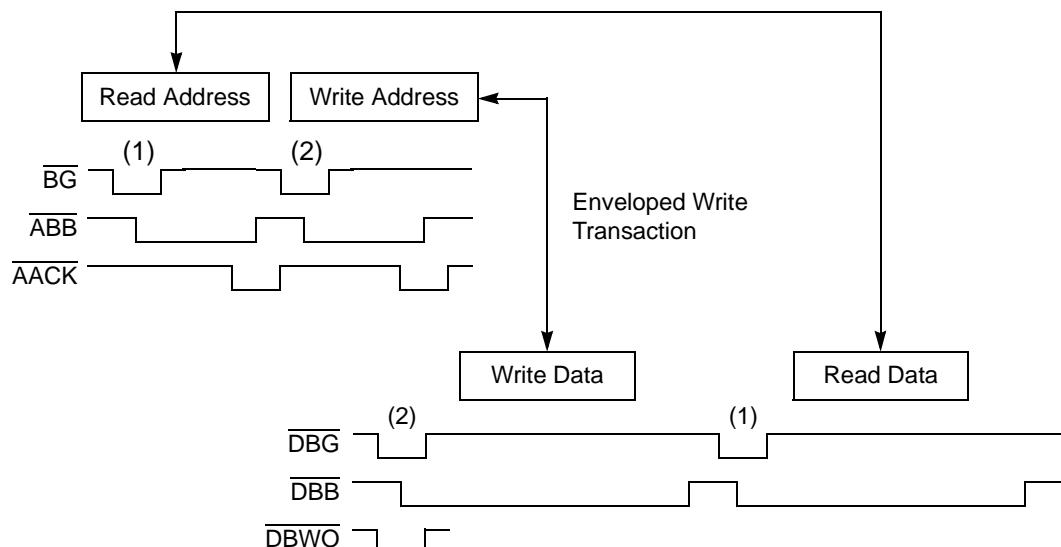


Figure 8-22. DBWO Transaction

Note that although the MPC603e can pipeline any write transaction behind the read transaction, special care should be used when using the enveloped write feature. It is envisioned that most system implementations will not need this capability; for these

applications, $\overline{\text{DBWO}}$ should remain negated. In systems where this capability is needed, $\overline{\text{DBWO}}$ should be asserted under the following scenario:

1. The MPC603e initiates a read transaction (either single-beat or burst) by completing the read address tenure with no address retry.
2. Then, the MPC603e initiates a write transaction by completing the write address tenure, with no address retry.
3. At this point, if $\overline{\text{DBWO}}$ is asserted with a qualified data bus grant to the MPC603e, the MPC603e asserts $\overline{\text{DBB}}$ and drives the write data onto the data bus, out of order with respect to the address pipeline. The write transaction concludes with the MPC603e negating $\overline{\text{DBB}}$.
4. The next qualified data bus grant signals the MPC603e to complete the outstanding read transaction by latching the data on the bus. This assertion of $\overline{\text{DBG}}$ should not be accompanied by an asserted $\overline{\text{DBWO}}$.

Any number of bus transactions by other bus masters can be attempted between any of these steps.

Note the following regarding $\overline{\text{DBWO}}$:

- $\overline{\text{DBWO}}$ can be asserted if no data bus read is pending, but it has no effect on write ordering.
- The ordering and presence of data bus writes is determined by the writes in the write queues at the time $\overline{\text{BG}}$ is asserted for the write address (not $\overline{\text{DBG}}$). If a particular write is desired (for example, a cache-line-snoop-push-out operation), then $\overline{\text{BG}}$ must be asserted after that particular write is in the queue and it must be the highest priority write in the queue at that time. A cache-line-snoop-push-out operation may be the highest priority write, but more than one may be queued.
- Because more than one write may be in the write queue when $\overline{\text{DBG}}$ is asserted for the write address, more than one data bus write may be enveloped by a pending data bus read.

The arbiter must monitor bus operations and coordinate the various masters and slaves with respect to the use of the data bus when $\overline{\text{DBWO}}$ is used. Individual $\overline{\text{DBG}}$ signals associated with each bus device should allow the arbiter to synchronize both pipelined and split-transaction bus organizations. Individual $\overline{\text{DBG}}$ and $\overline{\text{DBWO}}$ signals provide a primitive form of source-level tagging for the granting of the data bus.

Note that use of $\overline{\text{DBWO}}$ allows some operation-level tagging with respect to the MPC603e and the use of the data bus.

Chapter 9

Power Management

The MPC603e microprocessor is the first microprocessor specifically designed for low-power operation. It provides both automatic and program-controllable power reduction modes for progressive reduction of power consumption. This chapter describes the hardware support provided by the MPC603e for power management.

9.1 Overview

The MPC603e has explicit power management features that are described in this chapter. Note that the design of the MPC603e is fully static, allowing the internal processor state to be preserved when no internal clock is present. The PID7t-603e implementation offers the following enhancements to the original MPC603e family:

- Lower-power design
- 2.5-volt core and 3.3-volt I/O

The device drivers must be modified for power management as operating systems service I/O requests by system calls to the device drivers. When a device driver is called to reduce the power of a device, it needs to be able to check the power state of the device, save the device configuration parameters, and put the device into a power-saving mode. Furthermore, every time the device driver is called, it needs to check the power status of the device and restore the device to the full-on state, if the device is in a power-saving mode.

9.2 Dynamic Power Management

Dynamic power management (DPM) automatically powers up and down the individual execution units of the MPC603e, based on the contents of the instruction stream. For example, if no floating-point instructions are being executed, the floating-point unit is automatically powered down. Power is not actually removed from the execution unit; instead, each execution unit has an independent clock input, which is automatically controlled on a clock-by-clock basis. Because CMOS circuits consume negligible power when they are not switching, stopping the clock to an execution unit effectively eliminates its power consumption. The operation of DPM is completely transparent to software or any external hardware. Dynamic power management is enabled by setting HID0[11] on power-up following a hard reset sequence (HRESET).

9.3 Programmable Power Modes

Hardware can enable a power management state through external asynchronous interrupts. The hardware interrupt causes the transfer of program flow to interrupt handler code. The appropriate mode is then set by the software. The MPC603e provides a separate interrupt and interrupt vector for power management—the system management interrupt (\overline{SMI}). The MPC603e also contains a decrementer timer that allows it to enter the nap or doze mode for a predetermined period and then return to full power operation through the decrementer interrupt exception.

The MPC603e provides four power modes selectable by setting the appropriate control bits in the MSR and HID0. The four power modes are described briefly as follows:

- Full-power—This is the default power state of the MPC603e. The MPC603e is fully powered and the internal functional units are operating at the full processor clock speed. If the dynamic power management mode is enabled, functional units that are idle will automatically enter a low-power state without affecting performance, software execution, or external hardware.
- Doze—All the functional units of the MPC603e are disabled, except for the time base/decrementer registers and the bus snooping logic. When the processor is in doze mode, an external asynchronous interrupt, system management interrupt, decrementer exception, hard or soft reset, or machine check brings the MPC603e into the full-power state. The MPC603e in doze mode maintains the phase-locked loop (PLL) in a fully powered state and is locked to the system external clock input (SYSCLK) so a transition to the full-power state takes only a few processor clock cycles.
- Nap—The nap mode further reduces power consumption by disabling bus snooping, leaving only the time base register and the PLL in a powered state. The MPC603e returns to the full-power state upon receipt of an external asynchronous interrupt, system management interrupt, decrementer exception, hard or soft reset, or machine check input (\overline{MCP}) signal. A return to full-power state from a nap state takes only a few processor clock cycles.
- Sleep—Sleep mode reduces power consumption to a minimum by disabling all internal functional units, after which external system logic may disable the PLL and SYSCLK. Returning the MPC603e to the full-power state requires the enabling of the PLL and SYSCLK, followed by the assertion of an external asynchronous interrupt, system management interrupt, hard or soft reset, or \overline{MCP} signal after the time required to relock the PLL.

Note that the MPC603e cannot switch from one power management mode to another without first returning to full-on mode. The nap and sleep modes disable bus snooping; therefore, a hardware handshake is provided to ensure coherency before the MPC603e enters these power management modes. Table 9-1 summarizes the four power states for the core.

Table 9-1. MPC603e Programmable Power Modes

PM Mode	Functioning Units	Activation Method	Full-Power Wake-Up Method
Full power	All units active	—	—
Full power (with DPM)	Requested logic by demand	By instruction dispatch	—
Doze	<ul style="list-style-type: none"> • Bus snooping • Data cache as needed • Decrementer timer 	Controlled by SW	External asynchronous exceptions Decrementer interrupt Reset
Nap	Decrementer timer	Controlled by hardware and software	External asynchronous exceptions Decrementer interrupt Reset
Sleep	None	Controlled by hardware and software	External asynchronous exceptions Reset

9.3.1 Power Management Modes

The following sections describe the characteristics of the MPC603e power management modes, the requirements for entering and exiting the various modes, and the system capabilities provided by the MPC603e while the power management modes are active.

9.3.1.1 Full-Power Mode with DPM Disabled

Full-power mode with DPM disabled is selected when the DPM enable bit in HID0[11] is cleared. The following characteristics apply:

- Default state following power-up and $\overline{\text{HRESET}}$
- All functional units are operating at full processor speed at all times

9.3.1.2 Full-Power Mode with DPM Enabled

Full-power mode with DPM enabled ($\text{HID0}[11] = 1$) provides on-chip power management without affecting the functionality or performance of the MPC603e as follows:

- Required functional units are operating at full processor speed
- Functional units are clocked only when needed
- No software or hardware intervention required after mode is set
- Software/hardware and performance transparent

9.3.1.3 Doze Mode

Doze mode disables most functional units but maintains cache coherency by enabling the bus interface unit and snooping. A snoop hit causes the MPC603e to enable the data cache, copy the data back to memory, disable the cache, and fully return to the doze mode.

Doze mode is characterized by the following features:

- Most functional units disabled
- Bus snooping and time base/decrementer still enabled
- PLL running and locked to SYSCLK

To enter the doze mode, the following conditions must occur:

- Set doze bit (HID0[8] = 1), MSR[POW] is set
- MPC603e enters doze mode after several processor clocks

To return to full-power mode, the following conditions must occur:

- Assert $\overline{\text{INT}}$, $\overline{\text{MCP}}$, $\overline{\text{SMI}}$, or decrementer interrupts
- Hard reset or soft reset
- Transition to full-power state occurs only after a few processor cycles

9.3.1.4 Nap Mode

The nap mode disables the MPC603e except for the processor PLL and time base/decrementer. The time base can be used to restore the core to a full-on state after a specified period.

Because bus snooping is disabled for nap and sleep mode, a hardware handshake using the quiesce request ($\overline{\text{QREQ}}$) and quiesce acknowledge ($\overline{\text{QACK}}$) signals are required to maintain data coherency. The MPC603e will assert the $\overline{\text{QREQ}}$ signal to indicate that it is ready to disable bus snooping.

When the system logic has ensured that snooping is no longer necessary, it allows the processor to enter the nap (or sleep) mode and causes the assertion of the MPC603e $\overline{\text{QACK}}$ output signal for the duration of the nap mode period.

Nap mode is characterized by the following features:

- Time base/decrementer still enabled
- Most functional units disabled (including bus snooping)
- PLL running and locked to SYSCLK

To enter the nap mode, the following conditions must occur:

- Set nap bit (HID0[9] = 1), MSR[POW] is set
- MPC603e asserts $\overline{\text{QREQ}}$
- System asserts $\overline{\text{QACK}}$
- The processor enters nap mode after several processor clocks

To return to full-power mode, the following conditions must occur:

- Assert $\overline{\text{INT}}$, $\overline{\text{MCP}}$, $\overline{\text{SMI}}$, or decrementer interrupts

- Transition to full-power takes only a few processor cycles
- Hard reset or soft reset

9.3.1.5 Sleep Mode

Sleep mode consumes the least amount of power of the four modes, since all functional units are disabled. To conserve the maximum amount of power, the PLL and internal SYSCLK signals can be disabled. Due to the fully static design of the MPC603e, the internal processor state is preserved when no internal clock is present. Because the time base and decrementer are disabled while the MPC603e is in sleep mode, the time base contents must be updated from an external time base following sleep mode, if accurate time-of-day maintenance is required.

Before entering sleep mode, the MPC603e asserts $\overline{\text{QREQ}}$ to indicate that it is ready to disable bus snooping. When the system has ensured that snooping is no longer necessary, the system logic allows the MPC603e to enter sleep mode by asserting $\overline{\text{QACK}}$ for the duration of the sleep mode period.

Sleep mode is characterized by the following features:

- All functional units disabled (including bus snooping and time base)
- All nonessential input receivers disabled
- Internal clock regenerators disabled
- PLL and SYSCLK can be disabled

To enter sleep mode, the following conditions must occur:

- Set sleep bit ($\text{HID0}[10] = 1$), $\text{MSR}[\text{POW}]$ is set
- MPC603e asserts $\overline{\text{QREQ}}$
- System logic asserts $\overline{\text{QACK}}$
- MPC603e enters sleep mode after several processor clocks

To return to full-power mode, the following conditions must occur:

- Assert $\overline{\text{INT}}$, $\overline{\text{MCP}}$, or $\overline{\text{SMI}}$
- Hard reset or soft reset

To return to full-power mode after PLL and SYSCLK are disabled in sleep mode, the following conditions must occur:

- Enable SYSCLK
- Reconfigure PLL into desired processor clock mode
- System logic waits for PLL startup and relock time (100 μsec)
- System logic asserts one of the sleep recovery signals (for example, $\overline{\text{INT}}$ or $\overline{\text{SMI}}$)

9.3.2 Power Management Software Considerations

Because the MPC603e is a dual issue processor with out-of-order execution capability, care must be taken in how the power management modes are entered. Furthermore, nap and sleep modes require all outstanding bus operations to be completed before the power management mode is entered. Section 9.4, “Example Code Sequence for Entering Processor Sleep Mode,” provides an example software sequence for putting the MPC603e into sleep mode.

Normally, during system configuration time, one of the power management modes would be selected by setting the appropriate HID0 mode bit. Later on, the power management mode is invoked by setting MSR[POW]. To ensure a clean transition into and out of the power management mode, set MSR[EE] and execute the following code sequence:

```

sync
mtmsr[POW = 1]
isync
loop: b loop

```

9.4 Example Code Sequence for Entering Processor Sleep Mode

The following is a sample code sequence for entering MPC603e sleep mode 1.

```

*****
# set up MPC603e HID0 power management bits
*****
*****processor HID and external interrupt initialization*****
#
# set up HID registers for the various processors of this family
# hid setup taken from minix's mpxPowerPC.s

        mfspr    r31, pvr          # pvr reg
        srawi    r31, r31, 16
resetTest603:
        cmpi    0, 0, r31, 3
        bne     cr0, endHIDSetup

        addi    r0, r0, 0
        oris    r0, r0, 0x8000      # enable machine check pin EMCP
        oris    r0, r0, 0x0010      # enable dynamic power mgmt DPM
        oris    r0, r0, 0x0020      # enable SLEEP power mode
        ori     r0, r0, 0x8000      # enable the Icache ICE
        ori     r0, r0, 0x4000      # enable the Dcache DCE
        ori     r0, r0, 0x0800      # invalidate Icache ICFI
        ori     r0, r0, 0x0400      # invalidate Dcache DCFI
        mtspr    hid0, r0
        isync
*****
# then when the processor is in a loop, force an SMI interrupt

```

```
*****
.orig 0x00001400          # System Management Interrupt

# force big-endian mode
    stw      r0,0x05f8,r0      # need nop every second inst.
    stw      r0,0x05fc,r0
    mfmsr   r0
    ori     r0,r0,r0
    ori     r0,r0,0x0001      # force big-endian LE bit
    ori     r0,r0,r0
    xori   r0,r0,0x0001      # force big-endian LE bit
    ori     r0,r0,r0
    mtmsr   r0
    ori     r0,r0,r0
    isync
    ori     r0,r0,r0

# save off additional registers to be corrupted
    stw      r20,0x05f4,r0
    mfspr   r21, srr0          # put srr0 in r21
    stw      r21,0x05f0,r0      # put r21 in 0x05f0
    mfspr   r22, srr1          # put srr1 in r22
    stw      r22,0x05ec,r0      # put r22 in 0x05ec
    stw      r23,0x05e8,r0
    mfcr    r23
    stw      r23,0x05e4,r0
    xor     r0,r0,r0

*****
# set msr pow bit to go into sleep mode

    sync
    mfmsr r5                  # get MSR
    addis r3, r0, 0x0004        # turn on POW bit
    ori   r3, r3, 0x0000        # turn on ME bit 19
    or    r5, r3, r5
    mtmsr r5
    isync

    addis r20, r0, 0x0000
    ori   r20, r20, 0x0002
stay_here:
    addic. r20, r20, -1        # subtract 1 from r20 and set cc
    bgt cr0, stay_here        # loop if positive

#
# restore corrupted registers
    lwz      r23,0x05e4,r0
    mtcrf   0xff,r23
    lwz      r23,0x05e8,r0
    lwz      r22,0x05ec,r0
    mtspr   srr1, r22
    lwz      r21,0x05f0,r0
```

Chapter 9. Power Management



```
mtspr      srr0, r21
lwz        r20,0x05f4,r0
lwz        r0,0x05fc,r0

sync
rfi

*****  
# to get out of sleep mode, do a Soft Reset  
*****  
  
.orig 0x00000100          # Reset handler in low memory  
  
# force big-endian mode
    stw      r0,0x05f8,r0      # need nop every second inst.
    stw      r0,0x05fc,r0
    mfmsr   r0
    ori     r0,r0,r0
    ori     r0,r0,0x0001      # force big-endian LE bit
    ori     r0,r0,r0
    xori   r0,r0,0x0001      # force big-endian LE bit
    ori     r0,r0,r0
    mtmsr   r0
    ori     r0,r0,r0
    isync
    ori     r0,r0,r0

# save off additional registers to be corrupted
    stw      r20,0x05f4,r0
    stw      r21,0x05f0,r0
    stw      r22,0x05ec,r0
    stw      r23,0x05e8,r0
    mfcr   r23
    stw      r23,0x05e4,r0
    xor     r0,r0,r0

# restore corrupted registers
    lwz      r23,0x05e4,r0
    mtcrf  0xff,r23
    lwz      r23,0x05e8,r0
    lwz      r22,0x05ec,r0
    lwz      r21,0x05f0,r0
    lwz      r20,0x05f4,r0
    lwz      r0,0x05fc,r0

sync
rfi
*****
```

Appendix A

PowerPC Instruction Set Listings

This appendix lists the MPC603e microprocessor's instruction set as well as the additional PowerPC instructions not implemented in the MPC603e. Instructions are sorted by mnemonic, opcode, function, and form. Also included is a quick reference table that contains general information, such as the architecture level, privilege level, and form, and indicates if the instruction is 64-bit and optional.

Note that split fields representing the concatenation of sequences from left to right, are shown in lowercase. For more information refer to Chapter 8, "Instruction Set," in the *Programming Environments Manual*.

The following key applies to the tables in this appendix.

Key: Reserved Bits Instruction Not Implemented in the MPC603e

A.1 Instructions Sorted by Mnemonic

Table A-1 lists the instructions implemented in the PowerPC architecture in alphabetical order by mnemonic.

Table A-1. Complete Instruction List Sorted by Mnemonic

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
addx	31		D		A		B		OE								266									Rc		
addcx	31		D		A		B		OE								10									Rc		
adde	31		D		A		B		OE								138									Rc		
addi	14		D		A												SIMM											
addic	12		D		A												SIMM											
addic.	13		D		A												SIMM											
addis	15		D		A												SIMM											
addmex	31		D		A		0 0 0 0 0		OE							234									Rc			
addzx	31		D		A		0 0 0 0 0		OE							202									Rc			
andx	31		S		A		B									28										Rc		
andcx	31		S		A		B									60										Rc		
andi.	28		S		A											UIMM												

Table A-1. Complete Instruction List Sorted by Mnemonic (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
andis.	29	S			A																							
bx	18														LI										AA	LK		
bctx	16	BO					BI																		AA	LK		
bcctrx	19	BO					BI				0	0	0	0	0											LK		
bclrx	19	BO					BI				0	0	0	0	0											LK		
cmp	31	crfD	0	L			A				B													0		0		
cmpli	11	crfD	0	L			A																		SIMM			
cmpl	31	crfD	0	L			A				B													32		0		
cmpli	10	crfD	0	L			A																		UIMM			
cntlzdx ¹	31	S			A		0	0	0	0														58		Rc		
cntlzwx	31	S			A		0	0	0	0														26		Rc		
crand	19	crbD					crbA				crbB													257		0		
crandc	19	crbD					crbA				crbB													129		0		
creqv	19	crbD					crbA				crbB													289		0		
crnand	19	crbD					crbA				crbB													225		0		
crnor	19	crbD					crbA				crbB													33		0		
cror	19	crbD					crbA				crbB													449		0		
crorc	19	crbD					crbA				crbB													417		0		
crxor	19	crbD					crbA				crbB													193		0		
dcbf	31	0	0	0	0		A				B													86		0		
dcbi ²	31	0	0	0	0		A				B													470		0		
dcbst	31	0	0	0	0		A				B													54		0		
dcbt	31	0	0	0	0		A				B													278		0		
dcbtst	31	0	0	0	0		A				B													246		0		
dcbz	31	0	0	0	0		A				B													1014		0		
divdx ¹	31	D			A		B				OE													489		Rc		
divdux ¹	31	D			A		B				OE													457		Rc		
divwx	31	D			A		B				OE													491		Rc		
divwux	31	D			A		B				OE													459		Rc		
eciwx	31	D			A		B																	310		0		
ecowx	31	S			A		B																	438		0		
eieio	31	0	0	0	0		0	0	0	0	0	0											854		0			
eqvx	31	S			A		B																	284		Rc		
extsbx	31	S			A		0	0	0	0													954		Rc			
extshx	31	S			A		0	0	0	0													922		Rc			
extswx ¹	31	S			A		0	0	0	0													986		Rc			
fabsx	63	D			0	0	0	0	0	0	B												264		Rc			

Table A-1. Complete Instruction List Sorted by Mnemonic (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	
faddx	63		D			A		B			0 0 0 0 0			21			Rc												
faddsx	59		D			A		B			0 0 0 0 0			21			Rc												
fcfidx¹	63		D			0 0 0 0 0		B									846										Rc		
fcmpo	63		crfD	0 0		A		B										32									0		
fcmpu	63		crfD	0 0		A		B										0									0		
fctididx¹	63		D			0 0 0 0 0		B									814										Rc		
fctidzx¹	63		D			0 0 0 0 0		B									815										Rc		
fctiwx	63		D			0 0 0 0 0		B									14										Rc		
fctiwzx	63		D			0 0 0 0 0		B									15										Rc		
fdivx	63		D			A		B			0 0 0 0 0			18			Rc												
fdivsx	59		D			A		B			0 0 0 0 0			18			Rc												
fmaddx	63		D			A		B			C			29			Rc												
fmaddsx	59		D			A		B			C			29			Rc												
fmxrx	63		D			0 0 0 0 0		B						72			Rc												
fmsubx	63		D			A		B			C			28			Rc												
fmsubsx	59		D			A		B			C			28			Rc												
fmulx	63		D			A		0 0 0 0 0			C			25			Rc												
fmulsx	59		D			A		0 0 0 0 0			C			25			Rc												
fnabsx	63		D			0 0 0 0 0		B						136			Rc												
fnegx	63		D			0 0 0 0 0		B						40			Rc												
fnmaddrx	63		D			A		B			C			31			Rc												
fnmaddsx	59		D			A		B			C			31			Rc												
fnmsubrx	63		D			A		B			C			30			Rc												
fnmsubsx	59		D			A		B			C			30			Rc												
fresx³	59		D			0 0 0 0 0		B			0 0 0 0 0			24			Rc												
frspx	63		D			0 0 0 0 0		B						12			Rc												
frsqrtex³	63		D			0 0 0 0 0		B			0 0 0 0 0			26			Rc												
fselx³	63		D			A		B			C			23			Rc												
fsqrtx³	63		D			0 0 0 0 0		B			0 0 0 0 0			22			Rc												
fsqrtsx³	59		D			0 0 0 0 0		B			0 0 0 0 0			22			Rc												
fsubx	63		D			A		B			0 0 0 0 0			20			Rc												
fsubsx	59		D			A		B			0 0 0 0 0			20			Rc												
icbi	31		0 0 0 0 0			A		B						982			0												
isync	19		0 0 0 0 0			0 0 0 0 0		0 0 0 0 0						150			0												
lbz	34		D			A								d															
lbzu	35		D			A								d															
lbuux	31		D			A		B						119			0												

Appendix A. PowerPC Instruction Set Listings

Table A-1. Complete Instruction List Sorted by Mnemonic (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lbzx	31		D		A		B											87									0	
ld¹	58		D		A													ds									0	
ldarx¹	31		D		A		B											84									0	
ldu¹	58		D		A													ds									1	
ldux¹	31		D		A		B											53									0	
idx¹	31		D		A		B											21									0	
lfd	50		D		A													d										
lfdyu	51		D		A													d										
lfdux	31		D		A		B											631									0	
lfdx	31		D		A		B											599									0	
lfs	48		D		A													d										
lfsu	49		D		A													d										
lfsux	31		D		A		B											567									0	
lfsx	31		D		A		B											535									0	
lha	42		D		A													d										
lhau	43		D		A													d										
lhaux	31		D		A		B											375									0	
lhax	31		D		A		B											343									0	
lhbrx	31		D		A		B											790									0	
lhz	40		D		A													d										
lhzu	41		D		A													d										
lhzux	31		D		A		B											311									0	
lhzx	31		D		A		B											279									0	
lmw⁴	46		D		A													d										
lswi⁴	31		D		A		NB											597									0	
lswx⁴	31		D		A		B											533									0	
lwa¹	58		D		A												ds									2		
lwaxr	31		D		A		B											20									0	
lwaux¹	31		D		A		B											373									0	
lwax¹	31		D		A		B											341									0	
lwbrx	31		D		A		B											534									0	
lwz	32		D		A													d										
lwzu	33		D		A													d										
lwzux	31		D		A		B											55									0	
lwzx	31		D		A		B											23									0	
mcrf	19	crfD	0 0		crfS	0 0		0 0 0 0 0										0									0	
mcrfs	63	crfD	0 0		crfS	0 0		0 0 0 0 0										64									0	

Table A-1. Complete Instruction List Sorted by Mnemonic (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
mcrxr	31		crfD	00		00000		00000											512								0	
mfcr	31		D			00000		00000											19								0	
mffsx	63		D			00000		00000											583								Rc	
mfmsr ²	31		D			00000		00000											83								0	
mfsp ⁵	31		D						spr										339								0	
mfsr ²	31		D		0	SR		00000											595								0	
mfsrin ²	31		D			00000		B											659								0	
mftb	31		D					tbr											371								0	
mtcrf	31		S		0		CRM		0										144								0	
mtfsb0x	63		crbD			00000		00000											70								Rc	
mtfsb1x	63		crbD			00000		00000											38								Rc	
mtfsfx	63	0		FM				0	B										711								Rc	
mtfsfix	63		crfD	00		00000		IMM	0										134								Rc	
mtmsr ²	31		S			00000		00000											146								0	
mtspr ⁵	31		S					spr											467								0	
mts ²	31		S		0	SR		00000											210								0	
mtsarin ²	31		S			00000		B											242								0	
mulhdx ¹	31		D			A		B	0										73								Rc	
mulhdux ¹	31		D			A		B	0										9								Rc	
mulhw x	31		D			A		B	0										75								Rc	
mulhwux	31		D			A		B	0										11								Rc	
mulldx ¹	31		D			A		B	OE										233								Rc	
mulli	7		D			A													SIMM									
mullw x	31		D			A		B	OE										235								Rc	
nandx	31		S			A		B											476								Rc	
negx	31		D			A		00000	OE										104								Rc	
norx	31		S			A		B											124								Rc	
orx	31		S			A		B											444								Rc	
orcx	31		S			A		B											412								Rc	
ori	24		S			A													UIMM									
oris	25		S			A													UIMM									
rfi ²	19		00000		00000		00000		00000										50								0	
rldclx ¹	30		S			A		B										mb		8							Rc	
rldcrx ¹	30		S			A		B										me		9							Rc	
rldiclx ¹	30		S			A		sh										mb		2	sh	Rc						
rldicrx ¹	30		S			A		sh										mb		0	sh	Rc						
rldicrx ¹	30		S			A		sh										me		1	sh	Rc						

Appendix A. PowerPC Instruction Set Listings

Table A-1. Complete Instruction List Sorted by Mnemonic (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
rldimix ¹	30		S		A		sh			mb		3	sh	Rc														
rlwimix	20		S		A		SH			MB		ME															Rc	
rlwinmx	21		S		A		SH			MB		ME															Rc	
rlwnmx	23		S		A		B			MB		ME															Rc	
sc	17		00000		00000			0000000000000000																1	0			
sbia ^{1, 2, 3}	31		00000		00000		00000					498														0		
sbie ^{1, 2, 3}	31		00000		00000			B				434														0		
sidx ¹	31		S		A		B					27															Rc	
slwx	31		S		A		B					24															Rc	
sradx ¹	31		S		A		B					794															Rc	
sradix ¹	31		S		A		sh					413		sh	Rc													
srawx	31		S		A		B					792															Rc	
srawix	31		S		A		SH					824															Rc	
srwdx ¹	31		S		A		B					539															Rc	
srwx	31		S		A		B					536															Rc	
stb	38		S		A							d																
stbu	39		S		A							d																
stbux	31		S		A		B					247														0		
stbx	31		S		A		B					215															0	
std ¹	62		S		A							ds															0	
stdcx. ¹	31		S		A		B					214															1	
stdu ¹	62		S		A							ds															1	
stdux ¹	31		S		A		B					181															0	
stdx ¹	31		S		A		B					149															0	
stfd	54		S		A							d																
stfdyu	55		S		A							d																
stfdux	31		S		A		B					759															0	
stfdx	31		S		A		B					727															0	
stfiwx ³	31		S		A		B					983															0	
stfs	52		S		A							d																
stfsu	53		S		A							d																
stfsux	31		S		A		B					695															0	
stfsx	31		S		A		B					663															0	
sth	44		S		A							d																
sthbrx	31		S		A		B					918															0	
sthru	45		S		A							d																
sthux	31		S		A		B					439															0	

Table A-1. Complete Instruction List Sorted by Mnemonic (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
sthx	31		S		A		B											407									0	
stmw ⁴	47		S		A													d										
stswi ⁴	31		S		A			NB										725									0	
stswx ⁴	31		S		A			B										661									0	
stw	36		S		A													d										
stwbrx	31		S		A		B											662									0	
stwcx.	31		S		A		B											150									1	
stwu	37		S		A													d										
stwux	31		S		A		B											183									0	
stwx	31		S		A		B											151									0	
subfx	31		D		A		B		OE									40									Rc	
subfcx	31		D		A		B		OE									8									Rc	
subfex	31		D		A		B		OE									136									Rc	
subfic	08		D		A													SIMM										
subfmex	31		D		A			00000		OE								232									Rc	
subfzex	31		D		A			00000		OE								200									Rc	
sync	31	00000		00000		00000			00000									598									0	
td ¹	31	TO		A		B												68									0	
tdi ¹	02	TO		A														SIMM										
tlbia ^{2, 3}	31	00000		00000		00000			00000									370									0	
tlbie ^{2, 3}	31	00000		00000		00000			B									306									0	
tlbld ^{2, 6}	31	00000		00000		00000			B									978									0	
tlbli ^{2, 6}	31	00000		00000		00000			B									1010									0	
tlbsync ^{2, 3}	31	00000		00000		00000			00000									566									0	
tw	31	TO		A		B												4									0	
twi	03	TO		A														SIMM										
xorx	31	S		A		B												316									Rc	
xori	26	S		A														UIMM										
xoris	27	S		A														UIMM										

¹ 64-bit instruction² Supervisor-level instruction³ Optional in the PowerPC architecture⁴ Load and store string or multiple instruction⁵ Supervisor- and user-level instruction⁶ Implementation-specific instruction**Appendix A. PowerPC Instruction Set Listings**

A.2 Instructions Sorted by Opcode

Table A-2 lists the instructions defined in the PowerPC architecture in numeric order by opcode.

Table A-2. Complete Instruction List Sorted by Opcode

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
tdi ¹	000010	TO			A																						SIMM	
twi	000011	TO			A																						SIMM	
mulli	000111	D			A																						SIMM	
subfic	001000	D			A																						SIMM	
cmpli	001010	crfD	0	L		A																					UIMM	
cmpi	001011	crfD	0	L		A																					SIMM	
addic	001100	D			A																						SIMM	
addic.	001101	D			A																						SIMM	
addi	001110	D			A																						SIMM	
addis	001111	D			A																						SIMM	
bcx	010000	BO			BI																						AA LK	
sc	010001	00000			00000																						1 0	
bx	010010																										AA LK	
mcrf	010011	crfD	0 0		crfS	0 0																					0	
bclrx	010011	BO			BI																						LK	
crnor	010011	crbD			crbA																						0	
rfi	010011	00000			00000																						0	
crandc	010011	crbD			crbA																						0	
isync	010011	00000			00000																						0	
crxor	010011	crbD			crbA																						0	
crnand	010011	crbD			crbA																						0	
crand	010011	crbD			crbA																						0	
creqv	010011	crbD			crbA																						0	
crorc	010011	crbD			crbA																						0	
cror	010011	crbD			crbA																						0	
bcctrx	010011	BO			BI																						LK	
rlwimx	010100	S			A																						Rc	
rlwinmx	010101	S			A																						Rc	
rlwnmx	010111	S			A																						Rc	
ori	011000	S			A																						UIMM	
oris	011001	S			A																						UIMM	
xori	011010	S			A																						UIMM	
xoris	011011	S			A																						UIMM	

Table A-2. Complete Instruction List Sorted by Opcode (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
andi.	0 1 1 1 0 0	S			A																							
andis.	0 1 1 1 0 1	S			A																							
rldiclx¹	0 1 1 1 1 0	S			A														mb		0 0 0	sh	Rc					
rldicrx¹	0 1 1 1 1 0	S			A														me		0 0 1	sh	Rc					
rldicx¹	0 1 1 1 1 0	S			A														mb		0 1 0	sh	Rc					
rldimix¹	0 1 1 1 1 0	S			A														mb		0 1 1	sh	Rc					
rldclx¹	0 1 1 1 1 0	S			A														mb		0 1 0 0 0		Rc					
rldcrx¹	0 1 1 1 1 0	S			A														me		0 1 0 0 1		Rc					
cmp	0 1 1 1 1 1	crfD	0	L		A													0 0 0 0 0 0 0 0 0 0 0 0		0							
tw	0 1 1 1 1 1	TO			A														0 0 0 0 0 0 0 1 0 0		0							
subfcx	0 1 1 1 1 1	D			A														OE	0 0 0 0 0 0 1 0 0 0		Rc						
mulhdux¹	0 1 1 1 1 1	D			A														0	0 0 0 0 0 0 1 0 0 1		Rc						
addcx	0 1 1 1 1 1	D			A														OE	0 0 0 0 0 0 1 0 1 0		Rc						
mulhwux	0 1 1 1 1 1	D			A														0	0 0 0 0 0 0 1 0 1 1		Rc						
mfcr	0 1 1 1 1 1	D			0 0 0 0 0														0 0 0 0 0 1 0 0 1 1		0							
lwarx	0 1 1 1 1 1	D			A														B	0 0 0 0 0 1 0 1 0 0		0						
Idx¹	0 1 1 1 1 1	D			A														B	0 0 0 0 0 1 0 1 0 1		0						
lwzx	0 1 1 1 1 1	D			A														B	0 0 0 0 0 1 0 1 1 1		0						
slwx	0 1 1 1 1 1	S			A														B	0 0 0 0 0 1 1 0 0 0		Rc						
cntlzwx	0 1 1 1 1 1	S			A														0 0 0 0 0	0 0 0 0 0 1 1 0 1 0		Rc						
sidx¹	0 1 1 1 1 1	S			A														B	0 0 0 0 0 1 1 0 1 1		Rc						
andx	0 1 1 1 1 1	S			A														B	0 0 0 0 0 1 1 1 0 0		Rc						
cmpl	0 1 1 1 1 1	crfD	0	L		A													B	0 0 0 0 1 0 0 0 0 0		0						
subfx	0 1 1 1 1 1	D			A														B	0 0 0 0 1 0 1 0 0 0		Rc						
Idux¹	0 1 1 1 1 1	D			A														B	0 0 0 0 1 1 0 1 0 1		0						
dcbst	0 1 1 1 1 1	0 0 0 0 0			A														B	0 0 0 0 1 1 0 1 1 0		0						
l wzux	0 1 1 1 1 1	D			A														B	0 0 0 0 1 1 0 1 1 1		0						
cntlzdx¹	0 1 1 1 1 1	S			A														0 0 0 0 0	0 0 0 0 1 1 1 0 1 0		Rc						
andcx	0 1 1 1 1 1	S			A														B	0 0 0 0 1 1 1 1 0 0		Rc						
td¹	0 1 1 1 1 1	TO			A														B	0 0 0 1 0 0 0 1 0 0		0						
mulhdx¹	0 1 1 1 1 1	D			A														B	0 0 0 1 0 0 1 0 0 1		Rc						
mulhwx	0 1 1 1 1 1	D			A														B	0 0 0 1 0 0 1 0 1 1		Rc						
mfmsr	0 1 1 1 1 1	D			0 0 0 0 0													0 0 0 0 0	0 0 0 1 0 1 0 0 1 1		0							
ldarx¹	0 1 1 1 1 1	D			A														B	0 0 0 1 0 1 0 1 0 0		0						
dcbf	0 1 1 1 1 1	0 0 0 0 0			A														B	0 0 0 1 0 1 0 1 1 0		0						
lbzx	0 1 1 1 1 1	D			A														B	0 0 0 1 0 1 0 1 1 1		0						
negx	0 1 1 1 1 1	D			A													0 0 0 0 0	OE	0 0 0 1 1 0 1 0 0 0		Rc						

Appendix A. PowerPC Instruction Set Listings

**For More Information On This Product,
Go to: www.freescale.com**

Table A-2. Complete Instruction List Sorted by Opcode (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lbzux	0 1 1 1 1 1		D		A		B						0 0 0 1 1 1 0 1 1 1		0													
norx	0 1 1 1 1 1		S		A		B						0 0 0 1 1 1 1 1 0 0		Rc													
subfex	0 1 1 1 1 1		D		A		B						OE	0 0 1 0 0 0 1 0 0 0		Rc												
addex	0 1 1 1 1 1		D		A		B						OE	0 0 1 0 0 0 1 0 1 0		Rc												
mtcrf	0 1 1 1 1 1		S	0		CRM		0					0 0 1 0 0 1 0 0 0 0		0													
mtmsr	0 1 1 1 1 1		S		0 0 0 0 0		0 0 0 0 0						0 0 1 0 0 1 0 0 1 0		0													
stdx¹	0 1 1 1 1 1		S		A		B						0 0 1 0 0 1 0 1 0 1		0													
stwcx.	0 1 1 1 1 1		S		A		B						0 0 1 0 0 1 0 1 1 0		1													
stwx	0 1 1 1 1 1		S		A		B						0 0 1 0 0 1 0 1 1 1		0													
stdux¹	0 1 1 1 1 1		S		A		B						0 0 1 0 1 1 0 1 0 1		0													
stwux	0 1 1 1 1 1		S		A		B						0 0 1 0 1 1 0 1 1 1		0													
subfzex	0 1 1 1 1 1		D		A		0 0 0 0 0						OE	0 0 1 1 0 0 1 0 0 0		Rc												
addzex	0 1 1 1 1 1		D		A		0 0 0 0 0						OE	0 0 1 1 0 0 1 0 1 0		Rc												
mtsrx	0 1 1 1 1 1		S	0	SR		0 0 0 0 0						0 0 1 1 0 1 0 0 1 0		0													
stdcx.¹	0 1 1 1 1 1		S		A		B						0 0 1 1 0 1 0 1 1 0		1													
stbx	0 1 1 1 1 1		S		A		B						0 0 1 1 0 1 0 1 1 1		0													
subfmex	0 1 1 1 1 1		D		A		0 0 0 0 0						OE	0 0 1 1 1 0 1 0 0 0		Rc												
mulld¹	0 1 1 1 1 1		D		A		B						OE	0 0 1 1 1 0 1 0 0 1		Rc												
addmex	0 1 1 1 1 1		D		A		0 0 0 0 0						OE	0 0 1 1 1 0 1 0 1 0		Rc												
mullwx	0 1 1 1 1 1		D		A		B						OE	0 0 1 1 1 0 1 0 1 1		Rc												
mtsrin	0 1 1 1 1 1		S		0 0 0 0 0		B						0 0 1 1 1 1 0 0 1 0		0													
dcbtst	0 1 1 1 1 1		0 0 0 0 0		A		B						0 0 1 1 1 1 0 1 1 0		0													
stbux	0 1 1 1 1 1		S		A		B						0 0 1 1 1 1 0 1 1 1		0													
addrx	0 1 1 1 1 1		D		A		B						OE	0 1 0 0 0 0 1 0 1 0		Rc												
dcbt	0 1 1 1 1 1		0 0 0 0 0		A		B						0 1 0 0 0 1 0 1 1 0		0													
lhzx	0 1 1 1 1 1		D		A		B						0 1 0 0 0 1 0 1 1 1		0													
eqvx	0 1 1 1 1 1		S		A		B						0 1 0 0 0 1 1 1 0 0		Rc													
tlbie^{2, 3}	0 1 1 1 1 1		0 0 0 0 0		0 0 0 0 0		B						0 1 0 0 1 1 0 0 1 0		0													
eciwx	0 1 1 1 1 1		D		A		B						0 1 0 0 1 1 0 1 1 0		0													
lhzux	0 1 1 1 1 1		D		A		B						0 1 0 0 1 1 0 1 1 1		0													
xorx	0 1 1 1 1 1		S		A		B						0 1 0 0 1 1 1 1 1 0 0		Rc													
mfspr⁴	0 1 1 1 1 1		D			spr							0 1 0 1 0 1 0 0 1 1		0													
lwax¹	0 1 1 1 1 1		D		A		B						0 1 0 1 0 1 0 1 0 1		0													
lhax	0 1 1 1 1 1		D		A		B						0 1 0 1 0 1 0 1 1 1		0													
tlbia^{2, 3}	0 1 1 1 1 1		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0						0 1 0 1 1 1 0 0 1 0		0													
mftb	0 1 1 1 1 1		D			tbr							0 1 0 1 1 1 0 0 1 1		0													
lwaux¹	0 1 1 1 1 1		D		A		B						0 1 0 1 1 1 0 1 0 1		0													

Table A-2. Complete Instruction List Sorted by Opcode (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lhaux	011111		D		A		B						0101110111														0	
sthx	011111		S		A		B						0110010111														0	
orcx	011111		S		A		B						0110011100														Rc	
sradix¹	011111		S		A		sh						1100111011		sh												Rc	
slbie^{1, 2, 3}	011111	00000		00000			B						0110110010														0	
ecowx	011111		S		A		B						0110110110														0	
sthux	011111		S		A		B						0110110111														0	
orx	011111		S		A		B						0110111100														Rc	
divdux¹	011111		D		A		B		OE				0111001001														Rc	
divwux	011111		D		A		B		OE				0111001011														Rc	
mtspr⁴	011111		S			spr							0111010011														0	
dcbi	011111	00000		A		B							0111010110														0	
nandx	011111		S		A		B						0111011100														Rc	
divdx¹	011111		D		A		B		OE				0111101001														Rc	
divwx	011111		D		A		B		OE				0111101011														Rc	
slbia^{1, 2, 3}	011111	00000		00000			00000						0111110010														0	
mcrxr	011111	crfD	00		00000		00000						1000000000														0	
lswx⁵	011111		D		A		B						1000010101														0	
lwbrx	011111		D		A		B						1000010110														0	
lfsx	011111		D		A		B						1000010111														0	
srwx	011111		S		A		B						1000011000														Rc	
srdx¹	011111		S		A		B						1000011011														Rc	
tibsync^{2, 3}	011111	00000		00000			00000						1000110110														0	
lfsux	011111		D		A		B						1000110111														0	
mfsr	011111		D	0	SR		00000						1001010011														0	
lswi⁵	011111		D		A		NB						1001010101														0	
sync	011111	00000		00000			00000						1001010110														0	
lfdx	011111		D		A		B						1001010111														0	
lfdtx	011111		D		A		B						1001110111														0	
mfsrin²	011111		D	00000			B						1010010011														0	
stswx⁵	011111		S		A		B						1010010101														0	
stwbrx	011111		S		A		B						1010010110														0	
stfsx	011111		S		A		B						1010010111														0	
stfsux	011111		S		A		B						1010110111														0	
stswi⁵	011111		S		A		NB						1011010101														0	
stfdx	011111		S		A		B						1011010111														0	
stfdtx	011111		S		A		B						1011110111														0	

Appendix A. PowerPC Instruction Set Listings

Table A-2. Complete Instruction List Sorted by Opcode (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lhbrx	0 1 1 1 1 1		D		A		B						1 1 0 0 0 1 0 1 1 0												0			
srawx	0 1 1 1 1 1		S		A		B						1 1 0 0 0 1 1 0 0 0												Rc			
sradx ¹	0 1 1 1 1 1		S		A		B						1 1 0 0 0 1 1 0 1 0												Rc			
srawix	0 1 1 1 1 1		S		A		SH						1 1 0 0 1 1 1 0 0 0												Rc			
eieio	0 1 1 1 1 1		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0						1 1 0 1 0 1 0 1 1 0											0				
sthbrx	0 1 1 1 1 1		S		A		B						1 1 1 0 0 1 0 1 1 0												0			
extshx	0 1 1 1 1 1		S		A		0 0 0 0 0						1 1 1 0 0 1 1 0 1 0												Rc			
extsbx	0 1 1 1 1 1		S		A		0 0 0 0 0						1 1 1 0 1 1 1 0 1 0												Rc			
tibld ^{2, 6}	0 1 1 1 1 1		0 0 0 0 0		0 0 0 0 0			B					1 1 1 1 0 1 0 0 1 0												0			
icbi	0 1 1 1 1 1		0 0 0 0 0		A		B						1 1 1 1 0 1 0 1 1 0												0			
stfiwx ³	0 1 1 1 1 1		S		A		B						1 1 1 1 0 1 0 1 1 1												0			
extsw ¹	0 1 1 1 1 1		S		A		0 0 0 0 0						1 1 1 1 0 1 1 0 1 0												Rc			
tiblid ^{2, 6}	0 1 1 1 1 1		0 0 0 0 0		0 0 0 0 0			B					1 1 1 1 1 1 0 0 1 0												0			
dcbz	0 1 1 1 1 1		0 0 0 0 0			A		B					1 1 1 1 1 1 0 1 1 0												0			
lwz	1 0 0 0 0 0		D		A								d															
lwzu	1 0 0 0 0 1		D		A								d															
lbz	1 0 0 0 1 0		D		A								d															
lbzu	1 0 0 0 1 1		D		A								d															
stw	1 0 0 1 0 0		S		A								d															
stwu	1 0 0 1 0 1		S		A								d															
stb	1 0 0 1 1 0		S		A								d															
stbu	1 0 0 1 1 1		S		A								d															
lhz	1 0 1 0 0 0		D		A								d															
lhzu	1 0 1 0 0 1		D		A								d															
lha	1 0 1 0 1 0		D		A								d															
lhau	1 0 1 0 1 1		D		A								d															
sth	1 0 1 1 0 0		S		A								d															
sthu	1 0 1 1 0 1		S		A								d															
lmw ⁵	1 0 1 1 1 0		D		A								d															
stmw ⁵	1 0 1 1 1 1		S		A								d															
lfs	1 1 0 0 0 0		D		A								d															
lfsu	1 1 0 0 0 1		D		A								d															
lfd	1 1 0 0 1 0		D		A								d															
lfdu	1 1 0 0 1 1		D		A								d															
stfs	1 1 0 1 0 0		S		A								d															
stfsu	1 1 0 1 0 1		S		A								d															
stfd	1 1 0 1 1 0		S		A								d															

Table A-2. Complete Instruction List Sorted by Opcode (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
stfd	1 1 0 1 1 1	S		A														d										
ld¹	1 1 1 0 1 0	D		A														ds			0 0							
ldu¹	1 1 1 0 1 0	D		A														ds			0 1							
lwa¹	1 1 1 0 1 0	D		A														ds			1 0							
fdivsx	1 1 1 0 1 1	D		A		B											0 0 0 0 0	1 0 0 1 0		Rc								
fsubsx	1 1 1 0 1 1	D		A		B											0 0 0 0 0	1 0 1 0 0		Rc								
faddsx	1 1 1 0 1 1	D		A		B											0 0 0 0 0	1 0 1 0 1		Rc								
fsqrtsx³	1 1 1 0 1 1	D		0 0 0 0 0		B											0 0 0 0 0	1 0 1 1 0		Rc								
fresx³	1 1 1 0 1 1	D		0 0 0 0 0		B											0 0 0 0 0	1 1 0 0 0		Rc								
fmulsx	1 1 1 0 1 1	D		A		0 0 0 0 0											C		1 1 0 0 1		Rc							
fmsubsx	1 1 1 0 1 1	D		A		B											C		1 1 1 0 0		Rc							
fmaddsx	1 1 1 0 1 1	D		A		B											C		1 1 1 0 1		Rc							
fnmsubsx	1 1 1 0 1 1	D		A		B											C		1 1 1 1 0		Rc							
fnmaddsx	1 1 1 0 1 1	D		A		B											C		1 1 1 1 1		Rc							
std¹	1 1 1 1 1 0	S		A													ds			0 0								
stdu¹	1 1 1 1 1 0	S		A													ds			0 1								
fcmpu	1 1 1 1 1 1	crfD	0 0		A		B										0 0 0 0 0 0 0 0 0 0 0		0									
frspx	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 0 0 0 1 1 0 0			Rc								
fctiwx	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 0 0 0 1 1 1 0											
fctiwzx	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 0 0 0 1 1 1 1			Rc								
fdivx	1 1 1 1 1 1	D		A		B											0 0 0 0 0	1 0 0 1 0		Rc								
fsubx	1 1 1 1 1 1	D		A		B											0 0 0 0 0	1 0 1 0 0		Rc								
faddx	1 1 1 1 1 1	D		A		B											0 0 0 0 0	1 0 1 0 1		Rc								
fsqrtx³	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 0 0	1 0 1 1 0		Rc								
fselx³	1 1 1 1 1 1	D		A		B											C		1 0 1 1 1		Rc							
fmulx	1 1 1 1 1 1	D		A		0 0 0 0 0											C		1 1 0 0 1		Rc							
frsqrte³	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 0 0	1 1 0 1 0		Rc								
fmsubx	1 1 1 1 1 1	D		A		B											C		1 1 1 0 0		Rc							
fmaddx	1 1 1 1 1 1	D		A		B											C		1 1 1 0 1		Rc							
fnmsubx	1 1 1 1 1 1	D		A		B											C		1 1 1 1 0		Rc							
fnmaddrx	1 1 1 1 1 1	D		A		B											C		1 1 1 1 1		Rc							
fcmwo	1 1 1 1 1 1	crfD	0 0		A		B										0 0 0 1 0 0 0 0 0		0									
mfsb1x	1 1 1 1 1 1	crbD			0 0 0 0 0		0 0 0 0 0										0 0 0 0 1 0 0 1 1 0			Rc								
fnegx	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 0 1 0 1 0 0 0			Rc								
mcrfs	1 1 1 1 1 1	crfD	0 0	crfS	0 0		0 0 0 0 0										0 0 0 1 0 0 0 0 0 0		0									
mfsb0x	1 1 1 1 1 1	crbD			0 0 0 0 0		0 0 0 0 0										0 0 0 1 0 0 0 1 1 0			Rc								
fmx	1 1 1 1 1 1	D		0 0 0 0 0		B											0 0 0 1 0 0 1 0 0 0			Rc								

Appendix A. PowerPC Instruction Set Listings

Table A-2. Complete Instruction List Sorted by Opcode (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
mtfsfix	1 1 1 1 1 1		crfD	0 0		0 0 0 0 0		IMM	0		0 0 1 0 0 0 0 1 1 0		Rc															
fnabsxv	1 1 1 1 1 1			D		0 0 0 0 0		B			0 0 1 0 0 0 1 0 0 0		Rc															
fabsx	1 1 1 1 1 1			D		0 0 0 0 0		B			0 1 0 0 0 0 1 0 0 0		Rc															
mffsx	1 1 1 1 1 1			D		0 0 0 0 0		0 0 0 0 0			1 0 0 1 0 0 0 1 1 1		Rc															
mtfsfx	1 1 1 1 1 1	0			FM			0		B		1 0 1 1 0 0 0 1 1 1		Rc														
fctidx¹	1 1 1 1 1 1			D		0 0 0 0 0		B			1 1 0 0 1 0 1 1 1 0		Rc															
fctidzx¹	1 1 1 1 1 1			D		0 0 0 0 0		B			1 1 0 0 1 0 1 1 1 1		Rc															
fcfidx¹	1 1 1 1 1 1			D		0 0 0 0 0		B			1 1 0 1 0 0 1 1 1 0		Rc															

¹ 64-bit instruction² Supervisor-level instruction³ Optional in the PowerPC architecture⁴ Supervisor- and user-level instruction⁵ Load and store string or multiple instruction⁶ MPC603e-implementation-specific instruction

A.3 Instructions Grouped by Functional Categories

Table A-3 through Table A-30 list the PowerPC instructions grouped by function.

Table A-3. Integer Arithmetic Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
addx	31		D		A		B		OE									266								Rc		
addcx	31		D		A		B		OE										10								Rc	
adde	31		D		A		B		OE										138								Rc	
addi	14		D		A														SIMM									
addic	12		D		A														SIMM									
addic.	13		D		A														SIMM									
addis	15		D		A														SIMM									
addmex	31		D		A		0 0 0 0 0		OE									234								Rc		
addzex	31		D		A		0 0 0 0 0		OE									202								Rc		
divdx¹	31		D		A		B		OE									489								Rc		
divdux¹	31		D		A		B		OE									457								Rc		
divwx	31		D		A		B		OE									491								Rc		
divwux	31		D		A		B		OE									459								Rc		
mulhdx¹	31		D		A		B		0									73								Rc		
mulhdux¹	31		D		A		B		0									9								Rc		
mulhwx	31		D		A		B		0									75								Rc		
mulhwux	31		D		A		B		0									11								Rc		
mulld¹	31		D		A		B		OE									233								Rc		
mulli	07		D		A													SIMM										
mullwx	31		D		A		B		OE									235								Rc		
negx	31		D		A		0 0 0 0 0		OE									104								Rc		
subfx	31		D		A		B		OE									40								Rc		
subfcx	31		D		A		B		OE									8								Rc		
subficx	08		D		A													SIMM										
subfex	31		D		A		B		OE									136								Rc		
subfmex	31		D		A		0 0 0 0 0		OE									232								Rc		
subfzex	31		D		A		0 0 0 0 0		OE									200								Rc		

¹ 64-bit instruction

Table A-4. Integer Compare Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
cmp	31		crfD	0	L		A		B				000000000000														0	
cmpi	11		crfD	0	L		A																				SIMM	
cmpl	31		crfD	0	L		A		B																	32	0	
cmpli	10		crfD	0	L		A																				UIMM	

Table A-5. Integer Logical Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
andx	31		S			A		B																		28	Rc	
andcx	31		S			A		B																		60	Rc	
andi.	28		S			A																				UIMM		
andis.	29		S			A																				UIMM		
cntlzdx¹	31		S			A		00000																	58	Rc		
cntlzwx	31		S			A		00000																	26	Rc		
eqvx	31		S			A		B																		284	Rc	
extsbx	31		S			A		00000																	954	Rc		
extshx	31		S			A		00000																	922	Rc		
extswx¹	31		S			A		00000																	986	Rc		
nandx	31		S			A		B																		476	Rc	
norx	31		S			A		B																		124	Rc	
orx	31		S			A		B																		444	Rc	
orcx	31		S			A		B																		412	Rc	
ori	24		S			A																				UIMM		
oris	25		S			A																				UIMM		
xorx	31		S			A		B																		316	Rc	
xori	26		S			A																				UIMM		
xoris	27		S			A																				UIMM		

¹ 64-bit instruction**Table A-6. Integer Rotate Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
rldclx¹	30		S			A		B						mb			8									Rc		
rldcrx¹	30		S			A		B						me			9									Rc		
rldicx¹	30		S			A		sh					mb			2	sh									Rc		
rldiclx¹	30		S			A		sh					mb			0	sh									Rc		
rldicrx¹	30		S			A		sh					me			1	sh									Rc		
rldimix¹	30		S			A		sh					mb			3	sh									Rc		

Table A-6. Integer Rotate Instructions (continued)

rlwimix	22	S	A	SH	MB	ME	Rc
rlwinmx	20	S	A	SH	MB	ME	Rc
rlwnmx	21	S	A	SH	MB	ME	Rc

¹ 64-bit instruction**Table A-7. Integer Shift Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
sldx¹	31		S		A		B											27									Rc	
slwx	31		S		A		B											24									Rc	
sradx¹	31		S		A		B											794									Rc	
sradix¹	31		S		A		sh											413		sh							Rc	
srawx	31		S		A		B											792									Rc	
srawix	31		S		A		SH											824									Rc	
srdx¹	31		S		A		B											539									Rc	
srwx	31		S		A		B											536									Rc	

¹ 64-bit instruction**Table A-8. Floating-Point Arithmetic Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
faddx	63		D		A		B										0 0 0 0 0			21						Rc		
faddsx	59		D		A		B										0 0 0 0 0			21						Rc		
fdivx	63		D		A		B										0 0 0 0 0			18						Rc		
fdivsx	59		D		A		B										0 0 0 0 0			18						Rc		
fmulx	63		D		A		0 0 0 0 0										C			25						Rc		
fmulsx	59		D		A		0 0 0 0 0										C			25						Rc		
fresx¹	59		D		0 0 0 0 0		B										0 0 0 0 0			24						Rc		
fsqrte¹	63		D		0 0 0 0 0		B										0 0 0 0 0			26						Rc		
fsubx	63		D		A		B										0 0 0 0 0			20						Rc		
fsubsx	59		D		A		B										0 0 0 0 0			20						Rc		
fselx¹	63		D		A		B										C			23						Rc		
fsqrtx¹	63		D		0 0 0 0 0		B										0 0 0 0 0			22						Rc		
fsqrtsx¹	59		D		0 0 0 0 0		B										0 0 0 0 0			22						Rc		

¹ Optional in the PowerPC architecture**Appendix A. PowerPC Instruction Set Listings**

Table A-9. Floating-Point Multiply-Add Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
fmaddx	63		D		A		B		C		29		Rc															
fmaddsx	59		D		A		B		C		29		Rc															
fmsubx	63		D		A		B		C		28		Rc															
fmsubsx	59		D		A		B		C		28		Rc															
fnmaddx	63		D		A		B		C		31		Rc															
fnmaddsx	59		D		A		B		C		31		Rc															
fnmsubx	63		D		A		B		C		30		Rc															
fnmsubsx	59		D		A		B		C		30		Rc															

Table A-10. Floating-Point Rounding and Conversion Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
fctidx¹	63		D		0	0	0	0	0	0	B			846		Rc												
fctidz¹	63		D		0	0	0	0	0	0	B			814		Rc												
fctidzx¹	63		D		0	0	0	0	0	0	B			815		Rc												
fctiw_x	63		D		0	0	0	0	0	0	B			14		Rc												
fctiwz_x	63		D		0	0	0	0	0	0	B			15		Rc												
frsp_x	63		D		0	0	0	0	0	0	B			12		Rc												

¹ 64-bit instruction**Table A-11. Floating-Point Compare Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
fcmpo	63		crfD		0	0		A			B				32		0											
fcmpu	63		crfD		0	0		A			B					0		0										

Table A-12. Floating-Point Status and Control Register Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
mcrfs	63		crfD		0	0	crfS	0	0	0	0	0	0	0	0	0	64		0									
mffsx	63		D		0	0	0	0	0	0	0	0	0	0	0	0	583		Rc									
mtfsb0_x	63		crbD		0	0	0	0	0	0	0	0	0	0	0	0	70		Rc									
mtfsb1_x	63		crbD		0	0	0	0	0	0	0	0	0	0	0	0	38		Rc									
mtfsfx	31	0		FM			0				B					711		Rc										
mtfsfix	63		crfD		0	0	0	0	0	0	IMM	0				134		Rc										

Table A-13. Integer Load Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lbz	34		D		A													d										
lbzu	35		D		A													d										
lbzux	31		D		A						B							119										0
lbzx	31		D		A					B								87										0
ld¹	58		D		A													ds										0
ldu¹	58		D		A													ds										1
ldux¹	31		D		A					B							53										0	
ldx¹	31		D		A				B								21										0	
lha	42		D		A													d										
lhau	43		D		A													d										
lhaux	31		D		A				B								375										0	
lhax	31		D		A			B									343										0	
lhz	40		D		A												d											
lhzu	41		D		A												d											
lhzux	31		D		A			B									311										0	
lhzx	31		D		A			B									279										0	
lwa¹	58		D		A												ds										2	
lwaux¹	31		D		A			B									373										0	
lwax¹	31		D		A			B									341										0	
lwz	32		D		A												d											
lwzu	33		D		A												d											
lwzux	31		D		A			B									55										0	
lwzx	31		D		A			B									23										0	

¹ 64-bit instruction**Table A-14. Integer Store Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
stb	38		S		A												d											
stbu	39		S		A												d											
stbux	31		S		A				B								247										0	
stbx	31		S		A			B									215										0	
std¹	62		S		A												ds										0	
stdu¹	62		S		A												ds										1	
stdux¹	31		S		A			B									181										0	
stdx¹	31		S		A			B									149										0	

Appendix A. PowerPC Instruction Set Listings

Table A-14. Integer Store Instructions (continued)

sth	44	S	A	d		
sth<u>u</u>	45	S	A	d		
sth<u>ux</u>	31	S	A	B	439	0
sth<u>x</u>	31	S	A	B	407	0
stw	36	S	A	d		
stw<u>u</u>	37	S	A	d		
stw<u>ux</u>	31	S	A	B	183	0
stw<u>x</u>	31	S	A	B	151	0

¹ 64-bit instruction

Table A-15. Integer Load and Store with Byte-Reverse Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lhbrx	31	D		A		B																					790	0
lwbrx	31	D		A		B																					534	0
sthbrx	31	S		A		B																					918	0
stwbrx	31	S		A		B																					662	0

Table A-16. Integer Load and Store Multiple Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lmw ¹	46	D		A																							d	
stmw ¹	47	S		A																							d	

¹ Load and store string or multiple instruction

Table A-17. Integer Load and Store String Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lswi ¹	31	D		A			NB																			597	0	
lswx ¹	31	D		A			B																			533	0	
stswi ¹	31	S		A			NB																			725	0	
stswx ¹	31	S		A			B																			661	0	

¹ Load and store string or multiple instruction

Table A-18. Memory Synchronization Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
eieio	31	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	854								0		
isync	19	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	150								0		
ldarx ¹	31	D			A			B										84								0		
lwarcx	31	D			A			B										20								0		
stdcx ¹	31	S			A			B										214								1		
stwcx.	31	S			A			B										150								1		
sync	31	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	598								0		

¹ 64-bit instruction**Table A-19. Floating-Point Load Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
lfd	50		D			A												d										
lfdyu	51		D			A												d										
lfdux	31		D			A			B									631								0		
lfdx	31		D			A			B									599								0		
lfs	48		D			A												d										
lfsu	49		D			A												d										
lfsux	31		D			A			B									567								0		
lfsx	31		D			A			B									535								0		

Table A-20. Floating-Point Store Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
stfd	54		S			A												d										
stfdyu	55		S			A												d										
stfdux	31		S			A			B									759								0		
stfdx	31		S			A			B									727								0		
stfiwx ¹	31		S			A			B									983								0		
stfs	52		S			A												d										
stfsu	53		S			A												d										
stfsux	31		S			A			B									695								0		
stfsx	31		S			A			B									663								0		

¹ Optional in the PowerPC architecture**Appendix A. PowerPC Instruction Set Listings**

Table A-21. Floating-Point Move Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
fabsx	63		D		0	0	0	0	0		B																Rc	
fmrx	63		D		0	0	0	0	0		B																Rc	
fnabsx	63		D		0	0	0	0	0		B																Rc	
fnegx	63		D		0	0	0	0	0		B																Rc	

Table A-22. Branch Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
bx	18																										AA	LK
btx	16		BO			BI																				AA	LK	
bcctrx	19		BO			BI										0	0	0	0								LK	
bclrx	19		BO			BI										0	0	0	0								LK	

Table A-23. Condition Register Logical Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
crand	19		crbD			crbA			crbB																	257	0	
crandc	19		crbD			crbA			crbB																	129	0	
creqv	19		crbD			crbA			crbB																	289	0	
crnand	19		crbD			crbA			crbB																	225	0	
crnor	19		crbD			crbA			crbB																	33	0	
cror	19		crbD			crbA			crbB																	449	0	
crorc	19		crbD			crbA			crbB																	417	0	
crxor	19		crbD			crbA			crbB																	193	0	
mcrf	19		crfD	0	0	crfS	0	0		0	0	0														000000000000	0	

Table A-24. System Linkage Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
rfi ¹	19		0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
sc	17		0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0

¹ Supervisor-level instruction

Table A-25. Trap Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
td ¹	31		TO		A		B						68														0	
tdi ¹	03		TO		A																							
tw	31		TO		A		B						4														0	
twi	03		TO		A																							

¹ 64-bit instruction**Table A-26. Processor Control Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
mcrxr	31		crfS	0	0	0	0	0	0	0	0	0	512													0		
mfcr	31		D		0	0	0	0	0	0	0	0	19													0		
mfmsr ¹	31		D		0	0	0	0	0	0	0	0	83													0		
mfsprr ²	31		D										339														0	
mftb	31		D										371														0	
mtcrf	31		S	0									144														0	
mtmsr ¹	31		S		0	0	0	0	0	0	0	0	146													0		
mtspr ²	31		D										467														0	

¹ Supervisor-level instruction² Supervisor- and user-level instruction**Table A-27. Cache Management Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
dcbf	31	0	0	0	0	0	0		A		B		86													0		
dcbi ¹	31	0	0	0	0	0	0		A		B		470													0		
dcbst	31	0	0	0	0	0	0		A		B		54													0		
dcbt	31	0	0	0	0	0	0		A		B		278													0		
dcbtst	31	0	0	0	0	0	0		A		B		246													0		
dcbz	31	0	0	0	0	0	0		A		B		1014													0		
icbi	31	0	0	0	0	0	0		A		B		982													0		

¹ Supervisor-level instruction**Appendix A. PowerPC Instruction Set Listings**

Table A-28. Segment Register Manipulation Instructions

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
mfsr ¹	31		D		0		SR		0 0 0 0 0													595		0				
mfsrin ¹	31		D				0 0 0 0 0						B										659		0			
mtsr ¹	31		S		0		SR		0 0 0 0 0													210		0				
mtsrin ¹	31		S				0 0 0 0 0						B										242		0			

¹ Supervisor-level instruction**Table A-29. Lookaside Buffer Management Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
slbia ^{1, 2, 3}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0																498		0			
slbie ^{1, 2, 3}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0						B										434		0			
tlbia ^{1, 3}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0																370		0			
tlbie ^{1, 3}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0					B											306		0			
tlbld ^{1, 4}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0					B											978		0			
tlbli ^{1, 4}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0					B											1010		0			
tlbsync ^{1, 3}	31		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0																566		0			

¹ Supervisor-level instruction² 64-bit instruction³ Optional in the PowerPC architecture⁴ MPC603e-implementation specific instruction**Table A-30. External Control Instructions**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
eciwx	31		D			A			B													310		0				
ecowx	31		S			A			B													438		0				

A.4 Instructions Sorted by Form

Table A-31 through Table A-45 list the PowerPC instructions grouped by form.

Table A-31. I-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD																LI									AA	LK	
Specific Instruction																												
bx		18															LI									AA	LK	

Table A-32. B-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD																BO		BI				BD			AA	LK	
Specific Instruction																												
bcx			16													BO		BI				BD			AA	LK		

Table A-33. SC-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD				0	0	0	0	0							0	0	0	0	0	0	0	0	0	0	0	1	0
Specific Instruction																												
sc				17											0	0	0	0	0	0	0	0	0	0	0	0	1	0

Table A-34. D-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD				D						A						d											
	OPCD				D						A							SIMM										
	OPCD				S						A						d											
	OPCD				S						A						UIMM											
	OPCD		crfD	0	L						A						SIMM											
	OPCD		crfD	0	L						A						UIMM											
	OPCD				TO						A						SIMM											

Table A-34. D-Form (continued)

Specific Instruction

addi	14	D		A	SIMM	
addic	12	D		A	SIMM	
addic.	13	D		A	SIMM	
addis	15	D		A	SIMM	
andi.	28	S		A	UIMM	
andis.	29	S		A	UIMM	
cmpi	11	crfD	0	L	A	SIMM
cmpli	10	crfD	0	L	A	UIMM
lbz	34	D		A	d	
lbzu	35	D		A	d	
lfd	50	D		A	d	
lfdु	51	D		A	d	
lfs	48	D		A	d	
lfsu	49	D		A	d	
lha	42	D		A	d	
lhau	43	D		A	d	
lhz	40	D		A	d	
lhzu	41	D		A	d	
lmw¹	46	D		A	d	
lwz	32	D		A	d	
lwzu	33	D		A	d	
mulli	7	D		A	SIMM	
ori	24	S		A	UIMM	
oris	25	S		A	UIMM	
stb	38	S		A	d	
stbu	39	S		A	d	
stfd	54	S		A	d	
stfdु	55	S		A	d	
stfs	52	S		A	d	
stfsu	53	S		A	d	
sth	44	S		A	d	
sthу	45	S		A	d	
stmw¹	47	S		A	d	
stw	36	S		A	d	
stwu	37	S		A	d	

Table A-34. D-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
subfic	08		D		A																							
tdi ²	02		TO		A																							
twi	03		TO		A																							
xori	26		S		A																							
xoris	27		S		A																							

¹ Load and store string or multiple instruction² 64-bit instruction**Table A-35. DS-Form**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		D		A																					XO		
	OPCD		S		A																					XO		

Specific Instructions

ld ¹	58	D		A																						0	
ldu ¹	58	D		A																						1	
lwa ¹	58	D		A																						2	
std ¹	62	S		A																						0	
stdu ¹	62	S		A																						1	

¹ 64-bit instruction**Table A-36. X-Form**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		D		A																				XO	0		
	OPCD		D		A																				XO	0		
	OPCD		D		0 0 0 0 0																			XO	0			
	OPCD		D		0 0 0 0 0																			XO	0			
	OPCD		D	0	SR																				XO	0		
	OPCD		S		A																				XO	Rc		
	OPCD		S		A																				XO	1		
	OPCD		S		A																				XO	0		
	OPCD		S		A																				NB	0		
	OPCD		S		A																				XO	Rc		
	OPCD		S		0 0 0 0 0																			XO	0			
	OPCD		S		0 0 0 0 0																			XO	0			
	OPCD		S	0	SR																				XO	0		

Appendix A. PowerPC Instruction Set Listings

Table A-36. X-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
------	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

	OPCD	S			A			SH			XO			Rc		
	OPCD	crfD	0	L	A			B			XO			0		
	OPCD	crfD	0 0		A			B			XO			0		
	OPCD	crfD	0 0		crfS	0 0		0 0 0 0 0			XO			0		
	OPCD	crfD	0 0			0 0 0 0 0			0 0 0 0 0			XO			0	
	OPCD	crfD	0 0			0 0 0 0 0			IMM		0	XO			Rc	
	OPCD	TO			A			B			XO			0		
	OPCD	D			0 0 0 0 0			B			XO			Rc		
	OPCD	D			0 0 0 0 0			0 0 0 0 0			XO			Rc		
	OPCD	crbD			0 0 0 0 0			0 0 0 0 0			XO			Rc		
	OPCD	0 0 0 0 0			A			B			XO			0		
	OPCD	0 0 0 0 0			0 0 0 0 0			B			XO			0		
	OPCD	0 0 0 0 0			0 0 0 0 0			0 0 0 0 0			XO			0		

Specific Instructions

andx	31	S			A			B			28			Rc
andcx	31	S			A			B			60			Rc
cmp	31	crfD	0	L	A			B			0			0
cmpl	31	crfD	0	L	A			B			32			0
cntlzdx ¹	31	S			A			0 0 0 0 0			58			Rc
cntlzwx	31	S			A			0 0 0 0 0			26			Rc
dcbf	31	0 0 0 0 0			A			B			86			0
dcbi ²	31	0 0 0 0 0			A			B			470			0
dcbst	31	0 0 0 0 0			A			B			54			0
dcbt	31	0 0 0 0 0			A			B			278			0
dcbtst	31	0 0 0 0 0			A			B			246			0
dcbz	31	0 0 0 0 0			A			B			1014			0
eciwx	31	D			A			B			310			0
ecowx	31	S			A			B			438			0
eieio	31	0 0 0 0 0			0 0 0 0 0			0 0 0 0 0			854			0
eqvx	31	S			A			B			284			Rc
extsbx	31	S			A			0 0 0 0 0			954			Rc
extshx	31	S			A			0 0 0 0 0			922			Rc
extswx ¹	31	S			A			0 0 0 0 0			986			Rc
fabsx	63	D			0 0 0 0 0			B			264			Rc
fcfidx ¹	63	D			0 0 0 0 0			B			846			Rc

Table A-36. X-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
fcmpo	63		crfD	0 0		A		B					32														0	
fcmpu	63		crfD	0 0		A		B					0														0	
fctidx¹	63		D		0 0 0 0 0		B						814														Rc	
fctidzx¹	63		D		0 0 0 0 0		B						815														Rc	
fctiwx	63		D		0 0 0 0 0		B						14														Rc	
fctiwzx	63		D		0 0 0 0 0		B						15														Rc	
fmrx	63		D		0 0 0 0 0		B						72														Rc	
fnabsx	63		D		0 0 0 0 0		B						136														Rc	
fnegx	63		D		0 0 0 0 0		B						40														Rc	
frspx	63		D		0 0 0 0 0		B						12														Rc	
icbi	31		0 0 0 0 0			A		B					982														0	
lbzux	31		D			A		B					119														0	
lbzx	31		D			A		B					87														0	
ldarx¹	31		D			A		B					84														0	
ldux¹	31		D			A		B					53														0	
Idx¹	31		D			A		B					21														0	
lfdux	31		D			A		B					631														0	
lfdx	31		D			A		B					599														0	
lfsux	31		D			A		B					567														0	
lfsx	31		D			A		B					535														0	
lhaux	31		D			A		B					375														0	
lhax	31		D			A		B					343														0	
lhbrx	31		D			A		B					790														0	
lhzux	31		D			A		B					311														0	
lhzx	31		D			A		B					279														0	
lswi³	31		D			A		NB					597														0	
lswx³	31		D			A		B					533														0	
lwaxr	31		D			A		B					20														0	
lwaux¹	31		D			A		B					373														0	
lwax¹	31		D			A		B					341														0	
lwbrx	31		D			A		B					534														0	
lwzux	31		D			A		B					55														0	
lwzx	31		D			A		B					23														0	
mcrfs	63		crfD	0 0		crfS	0 0	0 0 0 0 0					64													0		
mcrxr	31		crfD	0 0		0 0 0 0 0		0 0 0 0 0					512													0		
mfcr	31		D		0 0 0 0 0		0 0 0 0 0		0 0 0 0 0				19													0		

Appendix A. PowerPC Instruction Set Listings

Table A-36. X-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	
mffsx	63		D						0 0 0 0 0			0 0 0 0 0									583		Rc						
mfmsr ²	31		D						0 0 0 0 0			0 0 0 0 0									83		0						
mfsr ²	31		D				0		SR			0 0 0 0 0										595		0					
mfsrin ²	31		D						0 0 0 0 0			B										659		0					
mtfsb0x	63		crbD						0 0 0 0 0			0 0 0 0 0									70		Rc						
mtfsb1x	63		crfD						0 0 0 0 0			0 0 0 0 0									38		Rc						
mtfsfix	63		crbD	0 0					0 0 0 0 0			IMM	0									134		Rc					
mtmsr ²	31		S						0 0 0 0 0			0 0 0 0 0									146		0						
mtsri ²	31		S				0		SR			0 0 0 0 0										210		0					
mtsrin ²	31		S						0 0 0 0 0			B										242		0					
nandx	31		S						A			B										476		Rc					
norx	31		S						A			B										124		Rc					
orx	31		S						A			B										444		Rc					
orcx	31		S						A			B										412		Rc					
slbia ^{1, 2, 4}	31		0 0 0 0 0						0 0 0 0 0			0 0 0 0 0									498		0						
slbie ^{1, 2, 4}	31		0 0 0 0 0						0 0 0 0 0			B									434		0						
sldx ¹	31		S						A			B									27		Rc						
slwx	31		S						A			B									24		Rc						
sradx ¹	31		S						A			B									794		Rc						
srawx	31		S						A			B									792		Rc						
srawix	31		S						A			SH									824		Rc						
srdx ¹	31		S						A			B									539		Rc						
srwx	31		S						A			B									536		Rc						
stbx	31		S						A			B									247		0						
stbx	31		S						A			B									215		0						
stdcx. ¹	31		S						A			B									214		1						
stdux ¹	31		S						A			B									181		0						
stdx ¹	31		S						A			B									149		0						
stfdtx	31		S						A			B									759		0						
stfdx	31		S						A			B									727		0						
stfiwx ⁴	31		S						A			B									983		0						
stfsux	31		S						A			B									695		0						
stfsx	31		S						A			B									663		0						
sthbrx	31		S						A			B									918		0						
sthux	31		S						A			B									439		0						
sthx	31		S						A			B									407		0						

Table A-36. X-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
stswi ³	31		S			A							NB										725		0			
stswx ³	31		S			A							B										661		0			
stwbrx	31		S			A							B										662		0			
stwcx.	31		S			A							B										150		1			
stwux	31		S			A							B										183		0			
stwx	31		S			A							B										151		0			
sync	31		0	0	0	0	0						0	0	0	0	0						598		0			
td ¹	31		TO			A							B										68		0			
tlbia ^{2, 4}	31		0	0	0	0							0	0	0	0	0						370		0			
tlbie ^{2, 4}	31		0	0	0	0							0	0	0	0	0						306		0			
tlbld ^{2, 5}	31		0	0	0	0							B										978		0			
tlbli ^{2, 5}	31		0	0	0	0							B										1010		0			
tlbsync ^{2, 4}	31		0	0	0	0							0	0	0	0	0						566		0			
tw	31		TO			A							B										4		0			
xorx	31		S			A							B										316		Rc			

¹ 64-bit instruction² Supervisor- and user-level instruction³ Load and store string or multiple instruction⁴ Optional in the PowerPC architecture⁵ MPC603e-implementation specific instruction**Table A-37. XL-Form**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		BO			BI							0	0	0	0	0						XO		LK			
	OPCD		crbD			crbA							crbB											XO		0		
	OPCD		crfD	0	0	crfS	0	0					0	0	0	0	0						XO		0			
	OPCD		0	0	0	0							0	0	0	0	0						XO		0			

Specific Instructions

bcctrx	19		BO			BI							0	0	0	0	0						528		LK		
bclrx	19		BO			BI							0	0	0	0	0						16		LK		
crand	19		crbD			crbA							crbB											257		0	
crandc	19		crbD			crbA							crbB											129		0	
creqv	19		crbD			crbA							crbB											289		0	
crnand	19		crbD			crbA							crbB											225		0	
crnor	19		crbD			crbA							crbB											33		0	
cror	19		crbD			crbA							crbB											449		0	

Appendix A. PowerPC Instruction Set Listings

Table A-37. XL-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	
crorc	19																											417	0
crxor	19																											193	0
isync	19																											150	0
mcrf	19																											0	0
rfi ¹	19																											50	0

¹ Supervisor-level instruction

Table A-38. XFX-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		D																								XO	0
	OPCD		D			0																				XO	0	
	OPCD		S																							XO	0	
	OPCD		D																							XO	0	

Specific Instructions

mfsp ¹	31		D																								339	0
mftb	31		D																								371	0
mtcr	31		S			0																					144	0
mtspr ¹	31		D																								467	0

¹ Supervisor- and user-level instruction

Table A-39. XFL-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD	0																									XO	Rc

Specific Instructions

mtfsfx	63	0																									711	Rc
---------------	----	---	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	-----	----

Table A-40. XS-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		S																								XO	sh Rc

Specific Instructions

sradix ¹	31		S			A																					413	sh Rc
----------------------------	----	--	---	--	--	---	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	-----	-------

¹ 64-bit instruction

Table A-41. XO-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		D		A		B		OE			XO															Rc	
	OPCD		D		A		B		0			XO															Rc	
	OPCD		D		A			0 0 0 0 0		OE			XO														Rc	

Specific Instructions

addx	31	D	A	B	OE		266		Rc
addcx	31	D	A	B	OE		10		Rc
addex	31	D	A	B	OE		138		Rc
addmex	31	D	A	0 0 0 0 0	OE		234		Rc
addzex	31	D	A	0 0 0 0 0	OE		202		Rc
divdx¹	31	D	A	B	OE		489		Rc
divdux¹	31	D	A	B	OE		457		Rc
divwx	31	D	A	B	OE		491		Rc
divwux	31	D	A	B	OE		459		Rc
mulhdx¹	31	D	A	B	0		73		Rc
mulhdux¹	31	D	A	B	0		9		Rc
mulhwx	31	D	A	B	0		75		Rc
mulhwux	31	D	A	B	0		11		Rc
mulldx¹	31	D	A	B	OE		233		Rc
mullwx	31	D	A	B	OE		235		Rc
negx	31	D	A	0 0 0 0 0	OE		104		Rc
subfx	31	D	A	B	OE		40		Rc
subfcx	31	D	A	B	OE		8		Rc
subfex	31	D	A	B	OE		136		Rc
subfmex	31	D	A	0 0 0 0 0	OE		232		Rc
subfzex	31	D	A	0 0 0 0 0	OE		200		Rc

¹ 64-bit instruction**Table A-42. A-Form**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		D		A		B			0 0 0 0 0			XO													Rc		
	OPCD		D		A		B				C			XO												Rc		
	OPCD		D		A			0 0 0 0 0			C			XO												Rc		
	OPCD		D			0 0 0 0 0		B			0 0 0 0 0			XO													Rc	

Appendix A. PowerPC Instruction Set Listings

Table A-42. A-Form (continued)

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
------	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

Specific Instructions

faddx	63	D	A	B	0 0 0 0	21	Rc
faddsx	59	D	A	B	0 0 0 0	21	Rc
fdivx	63	D	A	B	0 0 0 0	18	Rc
fdivsx	59	D	A	B	0 0 0 0	18	Rc
fmaddx	63	D	A	B	C	29	Rc
fmaddsx	59	D	A	B	C	29	Rc
fmsubx	63	D	A	B	C	28	Rc
fmsubsx	59	D	A	B	C	28	Rc
fmulx	63	D	A	0 0 0 0	C	25	Rc
fmulsx	59	D	A	0 0 0 0	C	25	Rc
fnmaddx	63	D	A	B	C	31	Rc
fnmaddsx	59	D	A	B	C	31	Rc
fnmsubx	63	D	A	B	C	30	Rc
fnmsubsx	59	D	A	B	C	30	Rc
fresx ¹	59	D	0 0 0 0	B	0 0 0 0	24	Rc
frsqrtex ¹	63	D	0 0 0 0	B	0 0 0 0	26	Rc
fselx ¹	63	D	A	B	C	23	Rc
fsqrtx ¹	63	D	0 0 0 0	B	0 0 0 0	22	Rc
fsqrtsx ¹	59	D	0 0 0 0	B	0 0 0 0	22	Rc
fsubx	63	D	A	B	0 0 0 0	20	Rc
fsubsx	59	D	A	B	0 0 0 0	20	Rc

¹ Optional in the PowerPC architecture

Table A-43. M-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
------	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

OPCD	S	A	SH	MB	ME	Rc
OPCD	S	A	B	MB	ME	Rc

Specific Instructions

rlwimix	20	S	A	SH	MB	ME	Rc
rlwinmx	21	S	A	SH	MB	ME	Rc
rlwnmx	23	S	A	B	MB	ME	Rc

Table A-44. MD-Form

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		S		A		sh			mb		XO	sh	Rc														
	OPCD		S		A		sh			me		XO	sh	Rc														

Specific Instructions

ridicx ¹	30	S	A	sh	mb	2	sh	Rc
rldiclx ¹	30	S	A	sh	mb	0	sh	Rc
rldicrx ¹	30	S	A	sh	me	1	sh	Rc
rldimix ¹	30	S	A	sh	mb	3	sh	Rc

¹ 64-bit instruction**Table A-45. MDS-Form**

Name	0	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	OPCD		S		A		B			mb		XO		Rc														
	OPCD		S		A		B			me		XO		Rc														

Specific Instructions

rldclx ¹	30	S	A	B	mb	8		Rc
rldcrx ¹	30	S	A	B	me	9		Rc

¹ 64-bit instruction

A.5 Instruction Set Legend

Table A-46 provides general information on the PowerPC instruction set (such as the architectural level, privilege level, and form).

Table A-46. PowerPC Instruction Set Legend

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
addx	✓						XO
addcx	✓						XO
adde	✓						XO
addi	✓						D
addic	✓						D
addic.	✓						D
addis	✓						D
addmex	✓						XO
addzex	✓						XO
andx	✓						X
andcx	✓						X
andi.	✓						D
andis.	✓						D
bx	✓						I
bcx	✓						B
bcctrx	✓						XL
bclrx	✓						XL
cmp	✓						X
cmpi	✓						D
cmpl	✓						X
cmpli	✓						D
cntlzdx¹	✓				✓		X
cntlzwx	✓						X
crand	✓						XL
crandc	✓						XL
creqv	✓						XL
crnand	✓						XL
crnor	✓						XL
cror	✓						XL
crorc	✓						XL
crxor	✓						XL

Table A-46. PowerPC Instruction Set Legend (continued)

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
dcbf		✓					X
dcbi ²			✓	✓			X
dcbst		✓					X
dcbt		✓					X
dcbbst		✓					X
dcbz		✓					X
divdx ¹	✓				✓		XO
divdux ¹	✓				✓		XO
divwx	✓						XO
divwux	✓						XO
eciwx		✓				✓	X
ecowx		✓				✓	X
eieio		✓					X
eqvx	✓						X
extsbx	✓						X
extshx	✓						X
extswx ¹	✓				✓		X
fabsx	✓						X
faddx	✓						A
faddsx	✓						A
fcfidx ¹	✓				✓		X
fcmpo	✓						X
fcmpu	✓						X
fctididx ¹	✓				✓		X
fctidzx ¹	✓				✓		X
fctiwx	✓						X
fctiwzx	✓						X
fdivx	✓						A
fdivsx	✓						A
fmaddx	✓						A
fmaddsx	✓						A
fmxr	✓						X
fmsubx	✓						A
fmsubsx	✓						A
fmulx	✓						A

Appendix A. PowerPC Instruction Set Listings

Table A-46. PowerPC Instruction Set Legend (continued)

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
fmulsx	✓						A
fnabsx	✓						X
fnegx	✓						X
fmaddx	✓						A
fmaddsx	✓						A
fmsubx	✓						A
fmsubsx	✓						A
fresx³	✓					✓	A
frspx	✓						X
frsqrtex³	✓					✓	A
fselx³	✓					✓	A
fsqrtx³	✓					✓	A
fsqrtsx³	✓					✓	A
fsubx	✓						A
fsubsx	✓						A
icbi		✓					X
isync		✓					XL
lbz	✓						D
lbzu	✓						D
lbzux	✓						X
lbzx	✓						X
ld¹	✓				✓		DS
ldarx¹	✓				✓		X
ldu¹	✓				✓		DS
ldux¹	✓				✓		X
idx¹	✓				✓		X
lfd	✓						D
lfdu	✓						D
lfdux	✓						X
lfdx	✓						X
lfs	✓						D
lfsu	✓						D
lfsux	✓						X
lfsx	✓						X
lha	✓						D

Table A-46. PowerPC Instruction Set Legend (continued)

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
Ihau	✓						D
Ihaux	✓						X
Ihax	✓						X
Ihbx	✓						X
Ihz	✓						D
Ihzu	✓						D
Ihzux	✓						X
Ihzx	✓						X
Imw ⁴	✓						D
Iswi ⁴	✓						X
Iswx ⁴	✓						X
Iwa ¹	✓				✓		DS
Iwarx	✓						X
Iwaux ¹	✓				✓		X
Iwax ¹	✓				✓		X
Iwbrx	✓						X
Iwz	✓						D
Iwzu	✓						D
Iwzux	✓						X
Iwzx	✓						X
mcrf	✓						XL
mcrfs	✓						X
mcrxr	✓						X
mfcr	✓						X
mffsx	✓						X
mfmsr ²			✓	✓			X
mfsp ⁵	✓		✓	✓			XFX
mfsr ²			✓	✓			X
mfsrcin ²			✓	✓			X
mftb		✓					XFX
mtcrcf	✓						XFX
mtfsb0x	✓						X
mtfsb1x	✓						X
mtfsfx	✓						XFL
mtfsfix	✓						X

Appendix A. PowerPC Instruction Set Listings

Table A-46. PowerPC Instruction Set Legend (continued)

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
mtmsr ²			✓	✓			X
mtspr ⁵	✓		✓	✓			XFX
mtsr ²			✓	✓			X
mtsrin ²			✓	✓			X
mulhdx ¹	✓				✓		XO
mulhdux ¹	✓				✓		XO
mulhwx	✓						XO
mulhwux	✓						XO
mulldx ¹	✓				✓		XO
mulli	✓						D
mullwx	✓						XO
nandx	✓						X
negx	✓						XO
norx	✓						X
orx	✓						X
orcx	✓						X
ori	✓						D
oris	✓						D
rfi ²			✓	✓			XL
rldclx ¹	✓				✓		MDS
rldcrx ¹	✓				✓		MDS
rldicx ¹	✓				✓		MD
rldiclx ¹	✓				✓		MD
rldicrx ¹	✓				✓		MD
rldimix ¹	✓				✓		MD
rlwimix	✓						M
rlwinmx	✓						M
rlwnmx	✓						M
sc	✓		✓				SC
slbia ^{1, 2, 3}			✓	✓	✓	✓	X
slbie ^{1, 2, 3}			✓	✓	✓	✓	X
sldx ¹	✓				✓		X
slwx	✓						X
sradx ¹	✓				✓		X
sradix ¹	✓				✓		XS

Table A-46. PowerPC Instruction Set Legend (continued)

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
srawx	✓						X
srawix	✓						X
srdx ¹	✓				✓		X
srwx	✓						X
stb	✓						D
stbu	✓						D
stbux	✓						X
stbx	✓						X
std ¹	✓				✓		DS
stdcx. ¹	✓				✓		X
stdu ¹	✓				✓		DS
stdux ¹	✓				✓		X
stdx ¹	✓				✓		X
stfd	✓						D
stfdyu	✓						D
stfdux	✓						X
stfdx	✓						X
stfiwx ³	✓					✓	X
stfs	✓						D
stfsu	✓						D
stfsux	✓						X
stfsx	✓						X
sth	✓						D
sthbrx	✓						X
sthuy	✓						D
sthux	✓						X
sthx	✓						X
stmw ⁴	✓						D
stswi ⁴	✓						X
stswx ⁴	✓						X
stw	✓						D
stwbrx	✓						X
stwcx.	✓						X
stwu	✓						D
stwux	✓						X

Appendix A. PowerPC Instruction Set Listings

Table A-46. PowerPC Instruction Set Legend (continued)

	UISA	VEA	OEA	Supervisor Level	Optional	64-Bit	Form
stwx	✓						X
subfx	✓						XO
subfcx	✓						XO
subfex	✓						XO
subfic	✓						D
subfmex	✓						XO
subfzex	✓						XO
sync	✓						X
td¹	✓				✓		X
tdi¹	✓				✓		D
tlbia^{2, 3}			✓	✓		✓	X
tlbie^{2, 3}			✓	✓		✓	X
tblb^{2, 6}				✓			X
tblbi^{2, 6}				✓			X
tlbsync^{2, 3}			✓	✓			X
tw	✓						X
twi	✓						D
xorx	✓						X
xori	✓						D
xoris	✓						D

¹ 64-bit instruction² Supervisor-level instruction³ Optional in the PowerPC architecture⁴ Load and store string or multiple instruction⁵ Supervisor- and user-level instruction⁶ MPC603-implementation specific instruction

Appendix B Instructions Not Implemented

This appendix provides a list of the 32- and 64-bit instructions that are not implemented in the MPC603e microprocessor. It also provides a list of the 64-bit SPR encodings not implemented by the MPC603e. Note that any attempt to execute unimplemented instructions generates an illegal instruction exception.

Table B-1 lists the 32-bit instructions that are optional to the PowerPC architecture and not implemented by the MPC603e.

Table B-1. 32-Bit Instructions Not Implemented by the MPC603e

Mnemonic	Instruction
fsqrt	Floating Square Root (Double-Precision)
fsqrts	Floating Square Root Single
tlbia	TLB Invalidate All

Table B-2 provides a list of 64-bit instructions that are not implemented by 32-bit implementation such as the MPC603e microprocessor.

Table B-2. 64-Bit Instructions Not Implemented by the MPC603e

Mnemonic	Instruction
cntlzd	Count Leading Zeros Double Word
divd	Divide Double Word
divdu	Divide Double Word Unsigned
extsw	Extend Sign Word
fcfid	Floating Convert From Integer Double Word
fctid	Floating Convert to Integer Double Word
fctidz	Floating Convert to Integer Double Word with Round toward Zero
ld	Load Double Word
ldarx	Load Double Word and Reserve Indexed
ldu	Load Double Word with Update
ldux	Load Double Word with Update Indexed

Table B-2. 64-Bit Instructions Not Implemented by the MPC603e (continued)

Mnemonic	Instruction
Idx	Load Double Word Indexed
Iwa	Load Word Algebraic
Iwaux	Load Word Algebraic with Update Indexed
Iwax	Load Word Algebraic Indexed
mulld	Multiply Low Double Word
mulhd	Multiply High Double Word
mulhdu	Multiply High Double Word Unsigned
rldcl	Rotate Left Double Word then Clear Left
rldcr	Rotate Left Double Word then Clear Right
rldic	Rotate Left Double Word Immediate then Clear
rldicl	Rotate Left Double Word Immediate then Clear Left
rldicr	Rotate Left Double Word Immediate then Clear Right
rldimi	Rotate Left Double Word Immediate then Mask Insert
sibia	SLB Invalidate All
sibia	SLB Invalidate Entry
sld	Shift Left Double Word
srad	Shift Right Algebraic Double Word
sradi	Shift Right Algebraic Double Word Immediate
srd	Shift Right Double Word
std	Store Double Word
stdcx.	Store Double Word Conditional Indexed
stdu	Store Double Word with Update
stdux	Store Double Word Indexed with Update
stdx	Store Double Word Indexed
td	Trap Double Word
tdi	Trap Double Word Immediate

Table B-3 provides the 64-bit SPR encoding that is not implemented by the MPC603e microprocessor.

Table B-3. 64-Bit SPR Encoding Not Implemented by the MPC603e

SPR			Register Name	Access
Decimal	spr[5–9]	spr[0–4]		
280	01000	11000	ASR	Supervisor

Appendix C Revision History

This appendix provides a list of the major differences between the *MPC603e & EC603e RISC Microprocessor User's Manual*, Revision 0 and the *MPC603e RISC Microprocessor User's Manual*, Revision 3. Note that the list only covers the major changes to the user's manual.

C.1 Revision Changes From Revision 2 to Revision 3

Major changes to the *MPC603e RISC Microprocessor User's Manual* from Revision 2 to Revision 3 are as follows:

Section, Page	Change
4.5.2, 4-19	Replace the second paragraph in Section 4.5.2, "Machine Check Exception (0x00200):" Note that the 603e makes no attempt to force recoverability; however, it does guarantee the machine check exception is always taken immediately upon request, with a nonpredicted address saved in SRR0, regardless of the current machine state. Any pending stores in the completed store queue are canceled when the exception is taken. Software can use the machine check exception in a recoverable mode for checking bus configuration. For this case, a sync , load, sync instruction sequence is used. A subsequent machine check exception at the load address indicates a bus configuration problem and the processor is in a recoverable state.

With the following three paragraphs:

A machine check exception also occurs when an address or data parity error is detected on the bus and the address or data parity error is enabled in HID0. See Section 2.1.2.1, "Hardware Implementation Registers (HID0 and HID1)," for more information.

Note that the MPC603e makes no attempt to force recoverability on a machine check; however, it does guarantee that the machine check exception is always taken immediately upon request, with a

nonpredicted address saved in SRR0, regardless of the current machine state. Because pending stores in the completed store queue are not canceled when a machine check exception occurs, two consecutive stores that result in the assertion of $\overline{\text{TEA}}$ can cause the processor to checkstop. To prevent a checkstop in this case, a **sync** instruction must be placed between two stores that can result in assertion of $\overline{\text{TEA}}$.

Software can use the machine check exception in a recoverable mode to probe memory. For this case, a **sync**, load, **sync** instruction sequence is used. If the load access results in a system error (for example, the assertion of $\overline{\text{TEA}}$), the processor can handle this in a recoverable state. If the **sync** instruction is not used, a second access to the same address as the first load could cause the processor to enter the checkstop state.

5.5.2.2.2, 5-42

Replace the third line of the im1 segment in the tlbInstrMiss section with the following:

```
bdnzf 0, im1 # dec count br if cmp ne and if count not zero
```

5.5.2.2.2, 5-45

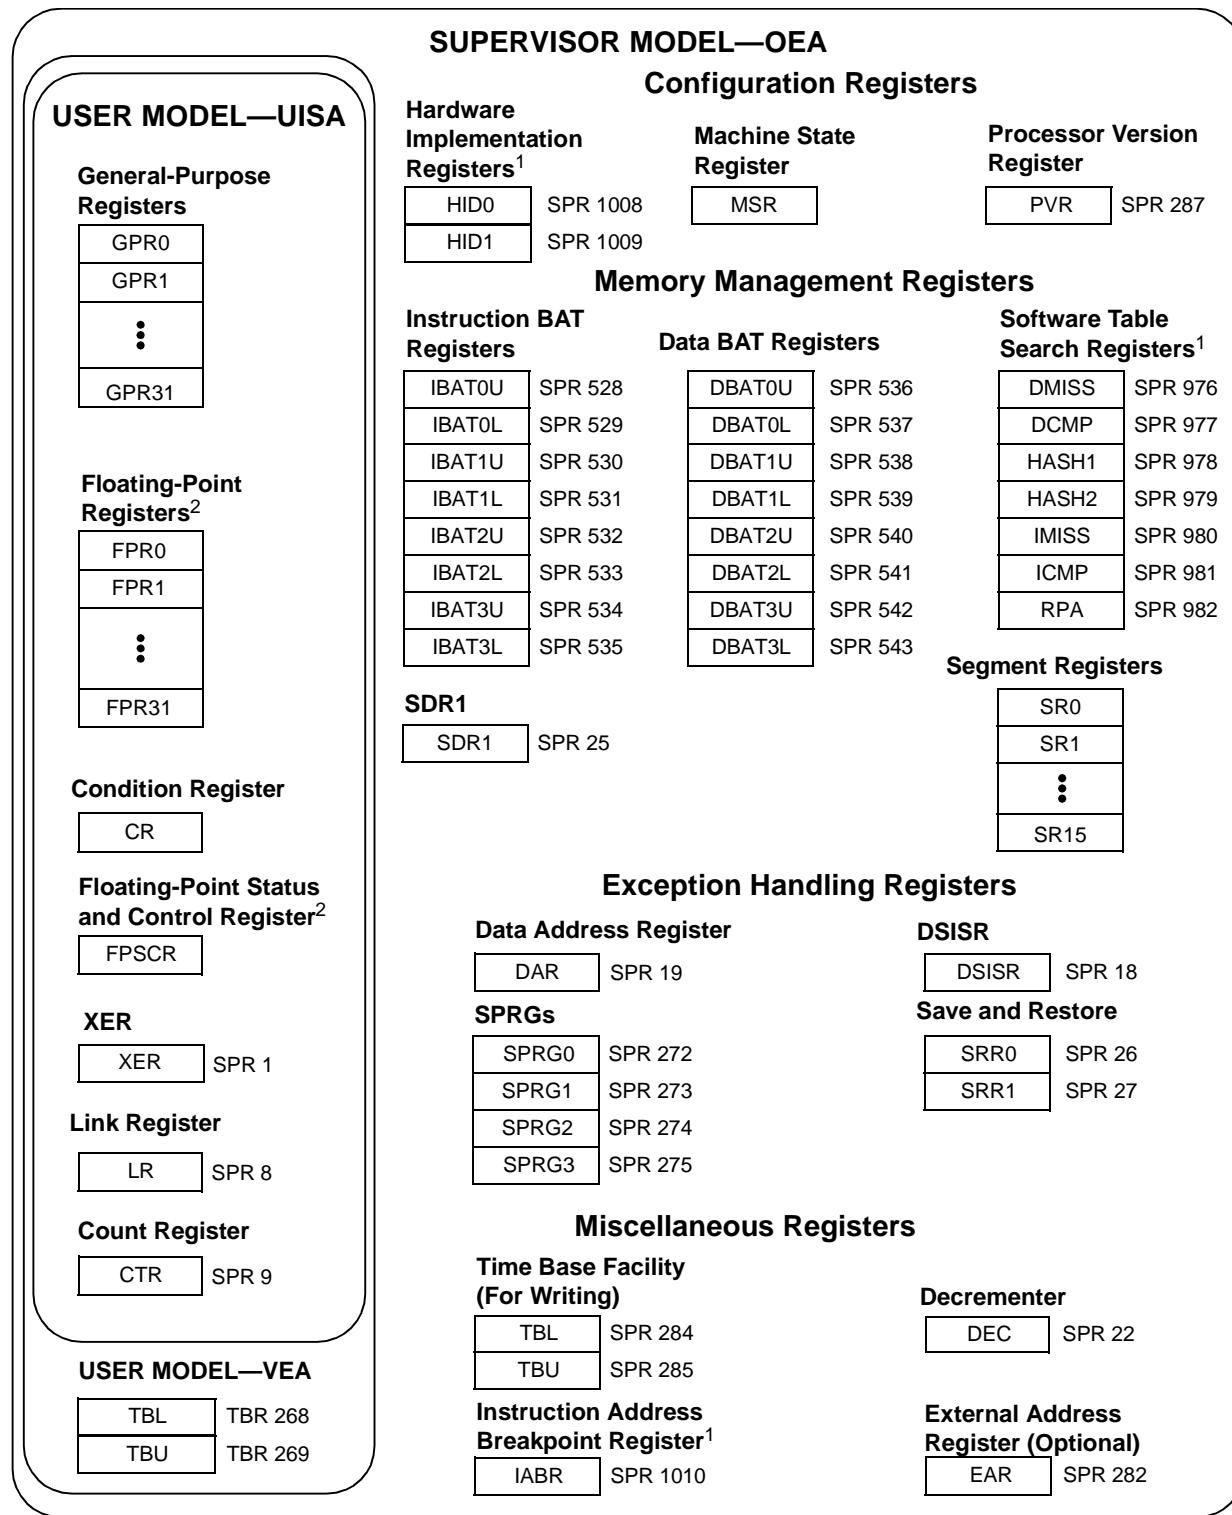
Replace the third line of the ceq1 segment in the tlbCeq0 section with the following:

```
bdnzf 0, ceq1 # dec count br if cmp ne and if count not zero
```

C.2 Revision Changes From Revision 1 to Revision 2

Major changes to the *MPC603e RISC Microprocessor User's Manual* from Revision 1 to Revision 2 are as follows:

Section, Page	Change
1.1.1, 1-6	The bus connecting the LSU to the GPRs should be 32 bits wide.
2.1.1, 2-3	Figure 2-1 in Revision 1 of the published manual should be replaced by the following figure:



Notes: ¹These registers are 603e-specific (PID6-603e and PID7v-603e) registers.
They may not be supported by other PowerPC processors.

²Not supported on the EC603e microprocessor.

Appendix C. Revision History

For More Information On This Product,
Go to: www.freescale.com

2.1.1, 2-4 Add the following text after the first paragraph of the sub-bullet for Floating-point registers (FPRs):

Before the **stfd** instruction is used to store the contents of an FPR to memory, the FPR must have been initialized after reset (explicitly loaded with any value) by using a floating point load instruction.

2.1.1, 2-5 Replace Table 2-1 with the following:

Table 2-1. MSR[POW] and MSR[TGPR] Bits

Bit	Name	Description
13	POW	<p>Power management enable (603e-specific)</p> <p>0 Disables programmable power modes (normal operation mode).</p> <p>1 Enables programmable power modes (nap, doze, or sleep mode).</p> <p>This bit controls the programmable power modes only; it has no effect on dynamic power management (DPM). MSR[POW] may be altered with an mtmsr instruction only. Also, when altering the POW bit, software may alter only this bit in the MSR and no others. The mtmsr instruction must be followed by a context-synchronizing instruction.</p> <p>See Chapter 9, “Power Management,” for more information on power management.</p>
14	TGPR	<p>Temporary GPR remapping (603e-specific)</p> <p>0 Normal operation</p> <p>1 TGPR mode. GPR0–GPR3 are remapped to TGPR0–TGPR3 for use by TLB miss routines. The contents of GPR0–GPR3 remain unchanged while MSR[TGPR] = 1. Attempts to use GPR4–GPR31 with MSR[TGPR] = 1 yield undefined results. When this bit is set, all instruction accesses to GPR0–GPR3 are mapped to TGPR0–TGPR3, respectively. The TGPR bit is set when either an instruction TLB miss, data TLB miss on load, or data TLB miss on store exception is taken. The TGPR bit is cleared by an rfi instruction.</p>

2.1.2.1, 2-8 Replace Table 2-2 with the HID0 bits descriptions as follows:

Table 2-2. HID0 Bit Functions

Bit	Name	Function
0	EMCP	<p>Enable <u>MCP</u>. The primary purpose of this bit is to mask out further machine check exceptions caused by assertion of <u>MCP</u>, similar to how MSR[EE] can mask external interrupts.</p> <p>0 Masks <u>MCP</u>. Asserting <u>MCP</u> does not generate a machine check exception or a checkstop.</p> <p>1 Asserting MCP causes checkstop if MSR[ME] = 0 or a machine check exception if ME = 1.</p>
1	—	Reserved
2	EBA	<p>Enable/disable 60x bus address parity checking</p> <p>0 Disables address parity checking.</p> <p>1 Allows a address parity error to cause a checkstop if MSR[ME] = 0 or a machine check exception if MSR[ME] = 1.</p> <p>EBA and EBD allow the processor to operate with memory subsystems that do not generate parity.</p>
3	EBD	<p>Enable 60x bus data parity checking</p> <p>0 Disables data parity checking.</p> <p>1 Allows a data parity error to cause a checkstop if MSR[ME] = 0 or a machine check exception if MSR[ME] = 1.</p> <p>EBA and EBD allow the processor to operate with memory subsystems that do not generate parity.</p>

Table 2-2. HID0 Bit Functions (continued)

Bit	Name	Function
4	BCLK	<u>CLK_OUT</u> output enable and clock type selection. Used in conjunction with HID0[ECLK] and the HRESET signal to configure CLK_OUT. See .
5	EICE	Enables in-circuit emulator outputs for pipeline tracking. See Section 7.2.11, “Pipeline Tracking Support,” for more information.
6	ECLK	<u>CLK_OUT</u> output enable and clock type selection. Used in conjunction with HID0[BCLK] and the HRESET signal to configure CLK_OUT. See .
7	PAR	Disable precharge of <u>ARTRY</u> . 0 Precharge of <u>ARTRY</u> enabled 1 Alters bus protocol slightly by preventing the processor from driving <u>ARTRY</u> to high (negated) state. If this is done, the system must restore the signals to the high state.
8	DOZE ¹	Doze mode enable. Operates in conjunction with MSR[POW]. 0 Doze mode disabled. 1 Doze mode enabled. Doze mode is invoked by setting MSR[POW] while this bit is set. In doze mode, the PLL, time base, and snooping remain active.
9	NAP ¹	Nap mode enable. Operates in conjunction with MSR[POW]. 0 Nap mode disabled. 1 Nap mode enabled. Doze mode is invoked by setting MSR[POW] while this bit is set. In nap mode, the PLL and the time base remain active.
10	SLEEP ¹	Sleep mode enable. Operates in conjunction with MSR[POW]. 0 Sleep mode disabled. 1 Sleep mode enabled. Sleep mode is invoked by setting MSR[POW] while this bit is set. <u>QREQ</u> is asserted to indicate that the processor is ready to enter sleep mode. If the system logic determines that the processor may enter <u>sleep</u> mode, the quiesce acknowledge signal, QACK, is asserted back to the processor. Once QACK assertion is detected, the processor enters sleep mode after several processor clocks. At this point, the system logic may turn off the PLL by first configuring PLL_CFG[0:3] to PLL bypass mode, then disabling SYCLK.
11	DPM ¹	Dynamic power management enable. 0 Dynamic power management is disabled. 1 Functional units enter a low-power mode automatically if the unit is idle. This does not affect operational performance and is transparent to software or any external hardware.
12–15	—	Reserved
16	ICE ²	Instruction cache enable 0 The instruction cache is neither accessed nor updated. All pages are accessed as if they were marked cache-inhibited (WIM = x1x). Potential cache accesses from the bus (snoop and cache operations) are ignored. In the disabled state for the L1 caches, the cache tag state bits are ignored and all accesses are propagated to the L2 cache or bus as single-beat transactions. For those transactions, however, CI reflects the original state determined by address translation regardless of cache disabled status. ICE is zero at power-up. 1 The instruction cache is enabled.
17	DCE ²	Data cache enable 0 The data cache is neither accessed nor updated. All pages are accessed as if they were marked cache-inhibited (WIM = x1x). Potential cache accesses from the bus (snoop and cache operations) are ignored. In the disabled state for the L1 caches, the cache tag state bits are ignored and all accesses are propagated to the L2 cache or bus as single-beat transactions. For those transactions, however, CI reflects the original state determined by address translation regardless of cache disabled status. DCE is zero at power-up. 1 The data cache is enabled.

Appendix C. Revision History

Table 2-2. HID0 Bit Functions (continued)

Bit	Name	Function
18	ILOCK ²	<p>Instruction cache lock</p> <p>0 Normal operation</p> <p>1 Instruction cache is locked. A locked cache supplies data normally on a hit, but are treated as a cache-inhibited transaction on a miss. On a miss, the transaction to the bus or the L2 cache is single-beat, however, CI still reflects the original state as determined by address translation independent of cache locked or disabled status.</p> <p>To prevent locking during a cache access, an isync instruction must precede the setting of ILOCK.</p>
19	DLOCK ²	<p>Data cache lock.</p> <p>0 Normal operation</p> <p>1 Data cache is locked. A locked cache supplies data normally on a hit but is treated as a cache-inhibited transaction on a miss. On a miss, the transaction to the bus or the L2 cache is single-beat, however, CI still reflects the original state as determined by address translation independent of cache locked or disabled status. A snoop hit to a locked L1 data cache performs as if the cache were not locked. A cache block invalidated by a snoop remains invalid until the cache is unlocked.</p> <p>To prevent locking during a cache access, a sync instruction must precede the setting of DLOCK.</p>
20	ICFI ²	<p>Instruction cache flash invalidate</p> <p>0 The instruction cache is not invalidated. The bit is cleared when the invalidation operation begins (usually the next cycle after the write operation to the register). The instruction cache must be enabled for the invalidation to occur.</p> <p>1 An invalidate operation is issued that marks the state of each instruction cache block as invalid without writing back modified cache blocks to memory. Cache access is blocked during this time. Bus accesses to the cache are signaled as a miss during invalidate-all operations. Setting ICFI clears all the valid bits of the blocks and the PLRU bits to point to way L0 of each set. For 603e processors, the proper use of the ICFI and DCFI bits is to set them and clear them with two consecutive mtspr operations.</p>
21	DCFI ²	<p>Data cache flash invalidate</p> <p>0 The data cache is not invalidated. The bit is cleared when the invalidation operation begins (usually the next cycle after the write operation to the register). The data cache must be enabled for the invalidation to occur.</p> <p>1 An invalidate operation is issued that marks the state of each data cache block as invalid without writing back modified cache blocks to memory. Cache access is blocked during this time. Bus accesses to the cache are signaled as a miss during invalidate-all operations. Setting DCFI clears all the valid bits of the blocks and the PLRU bits to point to way L0 of each set. For 603e processors, the proper use of the ICFI and DCFI bits is to set them and clear them with two consecutive mtspr operations.</p>
22–23	—	Reserved
24	IFEM	Instruction fetch enable M (PID7v-603e only). Enables the M bit on the bus. Used for instruction fetches.
25–26	—	Reserved
27	FBIOD	<p>Force branch indirect on bus.</p> <p>0 Register indirect branch targets are fetched normally</p> <p>1 Forces register indirect branch targets to be fetched externally</p>

Table 2-2. HID0 Bit Functions (continued)

Bit	Name	Function
28	ABE ²	<p>Address broadcast enable—controls whether certain address-only operations (such as cache operations) are broadcast on the 60x bus.</p> <p>0 Address-only operations affect only local caches and are not broadcast.</p> <p>1 Address-only operations are broadcast on the 60x bus. Affected instructions are dcbi, dcbf, and dcbst.</p> <p>Note that these cache control instruction broadcasts are not snooped by the PID7v-603e. Refer to Section 3.2.3, “Data Cache Control,” for more information.</p>
29–30	—	Reserved
31	NOOPTI	<p>No-op the data cache touch instructions.</p> <p>0 The dcbt and dcbtst instructions are enabled.</p> <p>1 The dcbt and dcbtst instructions are no-operated globally.</p>

¹ See Chapter 9, “Power Management,” for more information.

² See Chapter 3, “Instruction and Data Cache Operation,” for more information

2.1.2.1, 2-8 Also, add to this section the following: Table 2-3, the paragraph preceding it, and the sentence immediately following it.

Table 2-3 shows how HID0[BCLK], HID0[ECLK], and **HRESET** are used to configure CLK_OUT. See Section 7.2.12.2, “Test Clock (CLK_OUT)—Output,” for more information.

Table 2-3. HID0[BCLK] and HID0[ECLK] CLK_OUT Configuration

HRESET	HID0[ECLK]	HID0[BCLK]	CLK_OUT
Asserted	x	x	Bus
Negated	0	0	High impedance
Negated	0	1	Core clock frequency
Negated	1	0	Bus
Negated	1	1	Core clock frequency

HID0 can be accessed with **mtspr** and **mfsp** using SPR1008.

Then, after Figure 2-3, add the following:

HID1 can be accessed with **mfsp** using SPR1009.

2.1.2.7, 2-12 Remove this section. The Run_N counter is not an SPR.

2.3.5.3, 2-41 Add the following text to the end of this section:

Note that incoherency may occur if the following sequence of accesses hits the same cache block: a write-through, a **dcbz** instruction, a snoop. This occurs when the logical address for the **dcbz** and the write-through store are different but aliased to the same physical page.

Appendix C. Revision History

To avoid potential adverse effects, **dcbz** should not address write-through memory that can be accessed through multiple logical addresses. Explicit store instructions that write all zeroes should be used instead.

Note that broadcasting a sequence of **dcbz** instructions may cause snoop accesses to be retried indefinitely, which may cause the snoop originator to time out or may cause the snooped transaction to not complete. This can be avoided by disabling the broadcasting of **dcbz** by marking the memory space being addressed by the **dcbz** instruction as not global in the BAT or PTE.

2.3.5.4, 2-42

Replace the first four sentences of this section with the following:

The **eciwx** instruction provides an alternative way to map special devices. The MMU translation of the EA is not used to select the special device, as it is used in loads and stores. Rather, it is used as an address operand that is passed to the device over the address bus. Four other signals (the burst and size signals on the 60x bus) are used to select the device; these four signals output the 4-bit resource ID (RID) field in the EAR register. The **eciwx** instruction also loads a word from the data bus that is output by the special device.

2.3.6.3.1, 2-44

The section should be completely replaced with the following:

The supervisor-level cache management instruction in the PowerPC architecture, **dcbi**, should not be used on the 603e. The user-level **dcbf** instruction, described in Section 2.3.5.3, "Memory Control Instructions—VEA," and in Section 3.7, "Cache Control Instructions," should be used when the program needs to invalidate cache blocks. Note that the **dcbf** instruction causes modified blocks to be flushed to system memory if they are the target of a **dcbf** instruction whereas, by definition in the PowerPC architecture, the **dcbi** instruction only invalidates modified blocks.

3.1.3.2, 3-4

Change the last sentence of the paragraph to the following:

To prevent the cache from being enabled or disabled in the middle of a data access, an **isync** instruction should be issued before changing the value of ICE.

3.2.3.2, 3-6

Change the last sentence of the first paragraph to the following:

To prevent the cache from being enabled or disabled in the middle of a data access, a **sync** instruction should be issued before changing the value of DCE.

- 3.5.5.2, 3-14 Remove the third condition of the out-of-order load or store instruction to access guard memory:
“The target memory is cache-inhibited ($I = 1$), the load or store instruction is in the execution path, and it is guaranteed that no prior instructions can cause an exception.”
- 4.1, 4-3 The parenthetical clause in the second sentence of the first bullet should be “... (that is, all enabled floating-point exceptions are always precise on the 603e).”
- 4.1, 4-4 Figure 4-1 should be re-labeled as Table 4-1.
- 4.1, 4-5 At the end of the first sentence in the description of the floating-point unavailable exception, change “...when the floating-point available bit is disabled...” to “...when the floating-point available bit is cleared...”
Also, change “... Note that the EC603e microprocessor... execution of a floating-point...” to “Note that the EC603e microprocessor... execution of any floating-point...”
Also, the description for instruction translation miss should change from “... miss exception is caused when an effective address for an...” to “... miss exception is caused when the effective address for an...”
- 4.1, 4-6 The descriptions for Data load translation miss and Data store translation miss should change from “... miss exception is caused when an effective address for a...” to “... miss exception is caused when the effective address for a...”
- 4.2, 4-13 In Table 4-5 change the description of the FP bit from “...The processor prevents dispatch of floating-point instructions, including floating-point loads, stores, and moves, default...” to “...The processor prevents dispatch of floating-point instructions, including floating-point loads, stores, and moves; default...”
Bit 24 is “Reserved” rather than “Reserved. Full Function.”
- 4.2.2, 4-15 The first sentence of step 4 should be “The MSR is set as described in Table 4-7.”
- 4.5, 4-17 In Table 4-7, the setting for the IP bit in the System reset row should be a 1 instead of the “—” shown in the table.
- 4.5, 4-18 The first note for Table 4-7 should read as “The floating-point available bit is always cleared to 0 on the EC603e microprocessor.”

- 4.5.1.2, 4-20 The first sentence in the second paragraph should read as follows:
 Unlike a hard reset, no registers or latches are initialized; however, the instruction cache is disabled (HID0[ICE] = 0).
- The second sentence in the second paragraph should read as follows:
 After SRESET is recognized as asserted, the processor begins fetching instructions from the system reset routine at offset 0x0100.
- The third sentence should read as follows:
 When a soft reset occurs, registers are set as shown in Table 4-9 and HID0[ICE] is cleared.
- 4.5.5, 4-25 Replace the second sentence in the second paragraph with the following:
 The 603e allows the next instruction in program order to complete, including handling any exceptions that instruction may generate. However, the 603e blocks subsequent instructions from completing and allows any outstanding stores to occur to system memory.
- 4.5.5, 4-26 The last sentence of this section should be appended as follows:
 The interrupt handler must send a command to the device that asserted INT, acknowledging the interrupt and instructing the device to negate INT before the handler re-enables recognition of external interrupts.
- 4.5.6.1.1, 4-28 Remove the heading “4.5.6.1.1 Page Address Translation Access.” The text that follows this heading should just be part of section 4.5.6.1.
- 4.5.12, 4-34 In Table 4-16, the KEY bit was omitted and the polarity of bit 15 was reversed. Therefore, replace the SRR1 row with the following:

Table 4-16. Instruction and Data TLB Miss Exceptions— Register Settings (Partial)

SRR1	0–3 Loaded from condition register CR0 field
	4–11 Cleared
	12 KEY. Key for TLB miss (SR[Ks] or SR[Kp], depending on whether the access is a user or supervisor access).
	13 D/I. Data or instruction access 0 = Data TLB miss 1 = Instruction TLB miss
	14 WAY. Next TLB set to be replaced (set per LRU) 0 = Replace TLB associativity set 0 1 = Replace TLB associativity set 1
	15 S/L. Store or load data access 0 = Data TLB miss on load 1 = Data TLB miss on store (or C = 0)
	16–31 Loaded from MSR[16–31]

5.1.7, 5-16

In Table 5-4, replace the first three rows with the following table and delete the sixth row (**lwarx** or **stwcx.** with W = 1).

Table 5-4. Other MMU Exception Conditions

Condition	Description	Exception
TLB miss for an instruction fetch	No matching entry found in ITLB	Instruction TLB miss exception SRR1[13] = 1 MSR[14] = 1
TLB miss for a data load access	No matching entry found in DTLB for data load access	Data TLB miss on load exception SRR1[13] = 0 SRR1[15] = 1 MSR[14] = 1
TLB miss for a data store, or store and C = 0	No matching entry found in DTLB for data store access or matching DLTB entry has C = 0 and the access is a store	Data TLB miss on store exception, or store and C = 0 SRR1[13] = 0 SRR1[15] = 0 MSR[14] = 1

5.1.8, 5-18

In Table 5-5, add the following paragraph at the end of the **tlbie** description.

Software must ensure that instruction fetches or memory references to the virtual pages specified by the **tlbie** instruction have been completed prior to executing the **tlbie** instruction.

5.1.8, 5-18

In Table 5-5, add the following sentence at the end of the **tlbsync** description.

For a complete description of the **TLBISYNC** signal, refer to Section 8.8.2, “**TLBISYNC** Input.”

5.5.2.2.2, 5-42

Replace the third line of the im1 segment in the **tlbInstrMiss** section with the following:

```
bdnzf 0, im1 # dec count br if cmp ne and if count not zero
```

5.5.2.2.2, 5-45

Replace the third line of the ceq1 segment in the **tlbCeq0** section with the following:

```
bdnzf 0, ceq1 # dec count br if cmp ne and if count not zero
```

5.3, 5-20

Replace the Implementation Note for Block Address Translation with:

The 603e BAT registers are not initialized by the hardware after the power-up or reset sequence. Consequently, all valid bits in both instruction and data BAT areas must be explicitly cleared before setting any BAT area for the first time and before enabling translation. Also, note that software must avoid overlapping blocks while updating a BAT area or areas. Even if translation is disabled, multiple

Appendix C. Revision History

BAT area hits (with the valid bits set) can corrupt the remaining portion (any bits except the valid bits) of the BAT registers.

Thus, multiple BAT hits (with valid bits set) are considered a programming error whether translation is enabled or disabled and can lead to unpredictable results if translation is enabled (or if translation is disabled, when translation is eventually enabled). For the case of unused BATs (if translation is to be enabled) it is sufficient precaution to simply clear the valid bits of the unused BAT entries.”

5.4.3.1, 5-26

The second part of the second sentence in the first paragraph of that page should be deleted so that the sentence should read as follows:

ITLB miss exception conditions are reported when there are no more instructions to be dispatched or retired (the pipeline is empty).

5.4.3.1, 5-26

The very last sentence on this page should be edited as follows:

In order to uniquely identify a TLB entry as the required PTE, the TLB entry also contains...

5.4.3.1, 5-27

The phrase, “the valid entries are loaded and” should be deleted from the second sentence in the last paragraph of section 5.4.3.1.

5.4.4, 5-29

Figure 5-8 in the published manual incorrectly shows the loopback arrow on the left side pointing to the node above the word ‘otherwise.’ Replace Figure 5-8 with the following:

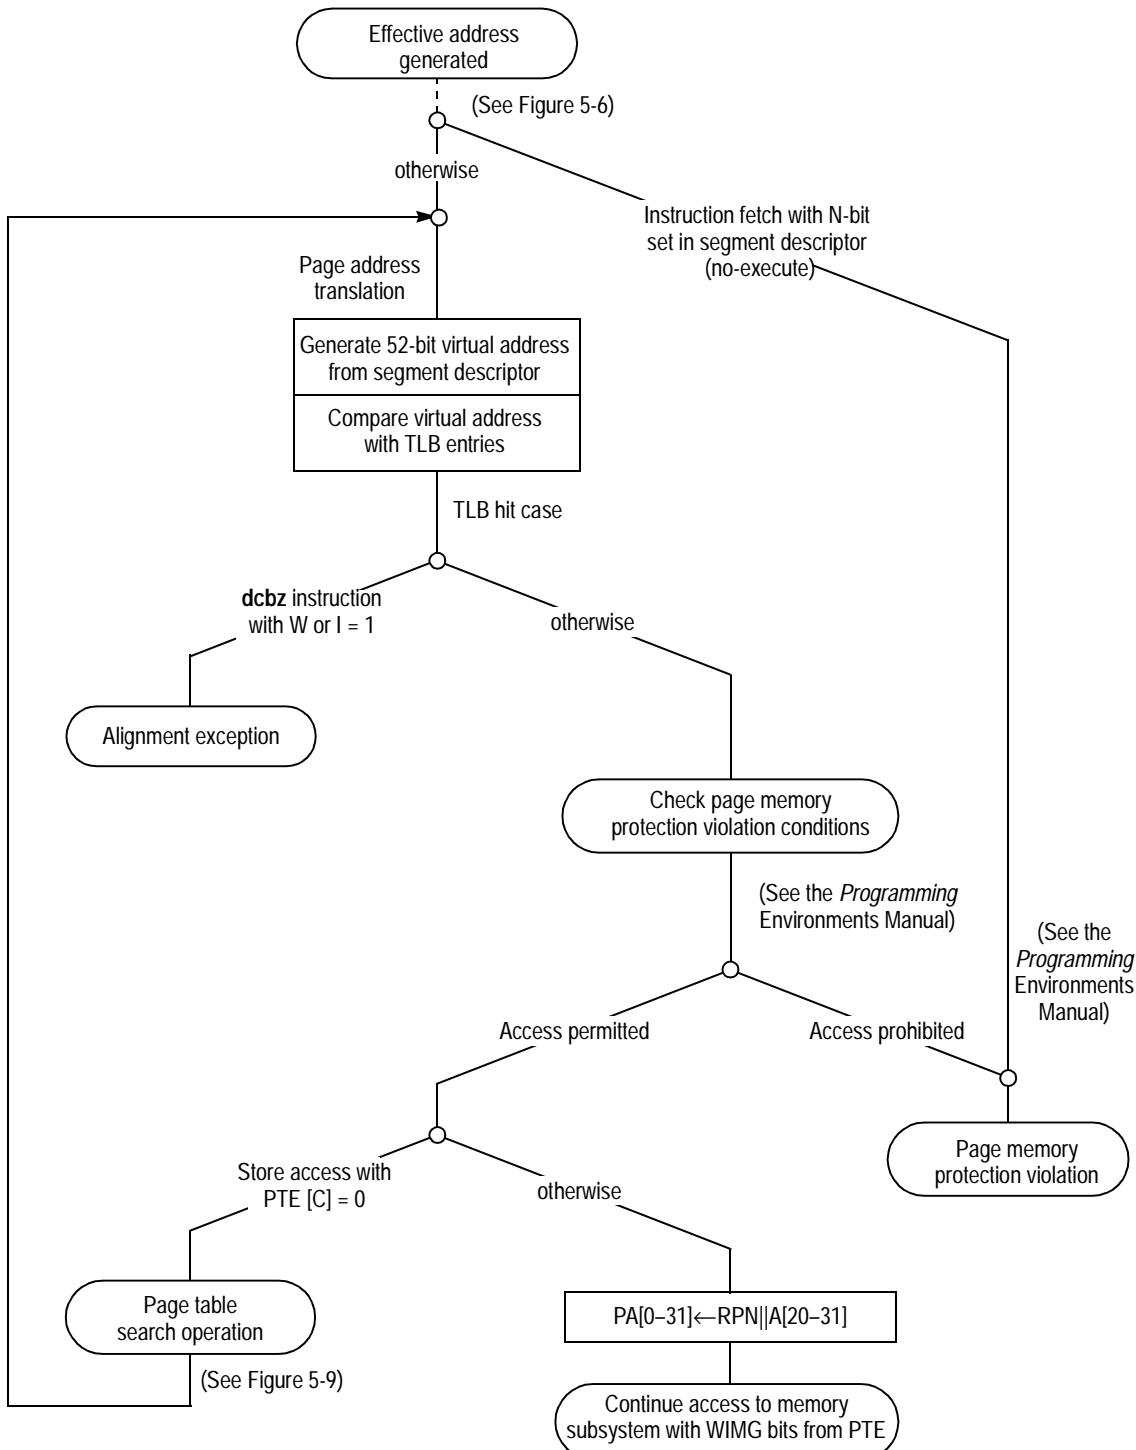


Figure 5-8. Page Address Translation Flow—TLB Hit

Appendix C. Revision History

5.5.2.2.2, 5-46 Replace the doISI segment with the following:

doISI:

```

mfspr r3, srr1      # get srr1
andi. r2, r3, 0xffff # clean srr1
addis r2, r2, 0x4000 # or in srr1<1> = 1 to flag pte not found

isil mtctr r0        # restore counter
mtspr srr1, r2       # set srr1
mfmsr r0             # get msr
xoris r0, r0, 0x8000 # flip the msr<tgpr> bit
mtcrf 0x80, r3       # restore CR0
mtmsr r0             # flip back to the native gprs
b     vec400          # go to instr. access exception
#

```

5.5.2.2.2, 5-47 Replace the third line of the dml segment with the following:

```
bdnzf 0, dml         # dec count br if cmp ne and if count not zero
```

5.5.2.2.2, 5-49 Replace the comment of the second line of the chk0 segment with the following:

```
andis. r3,r3,0x0008 # test the KEY bit (SRR1-bit 12)
```

5.5.2.2.2, 5-49 Replace the second line in the chk2 segment with the following

```
sth r1, 6(r2)        # update page table
```

5.5.2.2.2, 5-50 Replace the third line from the top of the page with the following

```
beq ds12:            # if little endian then:
```

8.2.1, 8-8 The second part of the DBG bullet should be edited to read as follows:
 “The DBB signal is driven by the current bus master. DRTRY is only driven from the bus; ARTRY is driven from the bus, but only for the address bus tenure...”

8.3.2.2.1, 8-13 The cross references at the end of this section should be to Table 7-1 and Table 7-2, instead of Table 8-1 and Table 8-2.

8.3.2.4, 8-17 The fifth entry in the TSIZ[0–2] column (second access for the second misaligned entry) should be 010 instead of 011.

8.3.3, 8-20 Replace sentences 3–5 with the following:

Although AACK can be asserted as early as the bus clock cycle following TS (see Figure 8-7), which allows a minimum address tenure of two bus cycles when the 603e clock is configured for 1:1 or 1.5:1 processor-to-bus clock mode, the ARTRY snoop response cannot be determined in the minimum allowed address tenure period. Thus in a system with two or more 603e processors using 1:1 or 1.5:1

clock mode, AACK must not be asserted until the third clock of the address tenure (one address wait state) to allow the snooping 603e processors an opportunity to assert ARTRY on the bus. For other clock configurations (2:1, 2.5:1, 3:1, 3.5:1, and 4:1), the ARTRY snoop response can be determined in the minimum address tenure period, and AACK may be asserted as early as the second bus clock of the address tenure.

8.4.4.1, 8-27

Insert a new section heading, 8.4.4.2, “Normal Burst Termination,” immediately after Figure 8-10.

8.7.4, 8-42

Replace the first sentence of the second paragraph with the following:

The system quiesce state is entered by configuring the processor to assert the QREQ output.

8.8.1, 8-42

The first sentence of the second paragraph should read as follows:

The reservation (RSRV) output signal is driven synchronously with the bus clock and reflects the status of the reservation coherency bit in the reservation address buffer (see Section 3.9, “Instruction and Data Cache Operation,” for more information).

9.2, 9-2

In the third sentence of the last paragraph on this page, the SMI should have an overbar (SMI).

9.2.1.4, 9-4

Change the fourth second-level bullet under the “Nap mode sequence” bullet to the following:

The processor enters nap mode after several processor clocks.

9.2.1.5, 9-5

The last sub-bullet should have overbars on INT and SMI.



Glossary of Terms and Abbreviations

The glossary contains an alphabetical list of terms, phrases, and abbreviations used in this book. Some of the terms and definitions included in the glossary are reprinted from IEEE Standard 754-1985, *IEEE Standard for Binary Floating-Point Arithmetic*, copyright ©1985 by the Institute of Electrical and Electronics Engineers, Inc. with the permission of the IEEE.

A

Architecture. A detailed specification of requirements for a processor or computer system. It does not specify details of how the processor or computer system must be implemented; instead it provides a template for a family of compatible *implementations*.

Asynchronous exception. *Exceptions* that are caused by events external to the processor's execution. In this document, the term *asynchronous exception* is used interchangeably with the word *interrupt*.

Atomic access. A bus access that attempts to be part of a read-write operation to the same address uninterrupted by any other access to that address (the term refers to the fact that the transactions are indivisible). The PowerPC architecture implements *atomic accesses* through the **Iwarx/stwcx**. instruction pair.

B

BAT (block address translation) mechanism. A software-controlled array that stores the available block address translations on-chip.

Beat. A single state on the MPC603e bus interface that may extend across multiple bus cycles. A MPC603e transaction can be composed of multiple address or data *beats*.

Biased exponent. An *exponent* whose range of values is shifted by a constant (bias). Typically a bias is provided to allow a range of positive values to express a range that includes both positive and negative values.

Big-endian. A byte-ordering method in memory where the address *n* of a word corresponds to the *most-significant byte*. In an addressed

memory word, the bytes are ordered (left to right) 0, 1, 2, 3, with 0 being the *most-significant byte*. See *Little-endian*.

Block. An area of memory that ranges from 128 Kbytes to 256 Mbytes whose size, translation, and protection attributes are controlled by the *BAT* mechanism.

Boundedly undefined. A characteristic of certain operation results that are not rigidly prescribed by the PowerPC architecture. Boundedly-undefined results for a given operation may vary among implementations and between execution attempts in the same implementation.

Although the architecture does not prescribe the exact behavior for when results are allowed to be *boundedly undefined*, the results of executing instructions in contexts where results are allowed to be *boundedly undefined* are constrained to ones that could have been achieved by executing an arbitrary sequence of defined instructions, in valid form, starting in the state the machine was in before attempting to execute the given instruction.

Branch folding. The replacement with target instructions of a branch instruction and any instructions along the not-taken path when a branch is either taken or predicted as taken.

Branch prediction. The process of guessing whether a branch will be taken. Such predictions can be correct or incorrect; the term ‘predicted’ as it is used here does not imply that the prediction is correct (successful). The PowerPC architecture defines a means for *static branch* prediction as part of the instruction encoding.

Branch resolution. The determination of whether a branch is taken or not taken. A branch is said to be resolved when the processor can determine which instruction path to take. If the branch is resolved as predicted, the instructions following the predicted branch that may have been speculatively executed can complete. If the branch is not resolved as predicted, instructions on the mispredicted path, and any results of speculative execution, are purged from the pipeline and fetching continues from the nonpredicted path.

Burst. A multiple-beat data transfer whose total size is typically equal to a cache block.

Bus clock. Clock that causes the bus state transitions.

Bus master. The owner of the address or data bus; the device that initiates or requests the transaction.

C

Cache. High-speed memory containing recently accessed data or instructions (subset of main memory).

Cache block. A small region of contiguous memory that is copied from memory into a *cache*. The size of a *cache block* may vary among processors; the maximum block size is one *page*. In PowerPC processors, *cache coherency* is maintained on a cache-block basis. Note that the term *cache block* is often used interchangeably with ‘cache line.’

Cache coherency. An attribute wherein an accurate and common view of memory is provided to all devices that share the same memory system. Caches are coherent if a processor performing a read from its cache is supplied with data corresponding to the most recent value written to memory or to another processor’s cache.

Cache flush. An operation that removes from a cache any data from a specified address range. This operation ensures that any modified data within the specified address range is written back to main memory. This operation is generated typically by a Data Cache Block Flush (**dcbf**) instruction.

Caching-inhibited. A memory update policy in which the cache is bypassed and the load or store is performed to or from main memory.

Cast out. A *cache block* that must be written to memory when a cache miss causes a *cache block* to be replaced.

Changed bit. One of two *page history bits* found in each *page table entry* (PTE). The processor sets the changed bit if any store is performed into the *page*. See also *Page access history bits* and *Referenced bit*.

Clean. An operation that causes a *cache block* to be written to memory, if modified, and then left in a valid, unmodified state in the cache.

Clear. To cause a bit or bit field to register a value of zero. See also *Set*.

Context synchronization. An operation that ensures that all instructions in execution complete past the point where they can produce an *exception*, that all instructions in execution complete in the context in which they began execution, and that all subsequent instructions are *fetched* and executed in the new context. Context

synchronization may result from executing specific instructions (such as **isync** or **rfi**) or when certain events occur (such as an *exception*).

Copy-back operation. A cache operation in which a cache line is copied back to memory to enforce cache coherency. Copy-back operations consist of snoop push-out operations and cache cast-out operations.

D

Denormalized number. A nonzero floating-point number whose exponent has a reserved value, usually the format's minimum, and whose explicit or implicit leading significand bit is zero.

Direct-mapped cache. A cache in which each main memory address can appear in only one location within the cache, operates more quickly when the memory request is a cache hit.

Direct-store segment access. An access to an I/O address space. The MPC603 defines separate memory-mapped and I/O address spaces, or segments, distinguished by the corresponding segment register T bit in the address translation logic of the MPC603. If the T bit is cleared, the memory reference is a normal memory-mapped access and can use the virtual memory management hardware of the MPC603. If the T bit is set, the memory reference is a direct-store access.

E

Effective address (EA). The 32-bit address specified for a load, store, or an instruction fetch. This address is then submitted to the MMU for translation to either a *physical memory* address or an I/O address.

Exception. A condition encountered by the processor that requires special, supervisor-level processing.

Exception handler. A software routine that executes when an exception is taken. Normally, the exception handler corrects the condition that caused the exception, or performs some other meaningful task (that may include aborting the program that caused the exception). The address for each exception handler is identified by an exception vector offset defined by the architecture and a prefix selected via the MSR.

Exclusive state. MEI state (E) in which only one caching device contains data that is also in system memory.

Execution synchronization. A mechanism by which all instructions in execution are architecturally complete before beginning execution (appearing to begin execution) of the next instruction. Similar to context synchronization but doesn't force the contents of the instruction buffers to be deleted and refetched.

Exponent. In the binary representation of a floating-point number, the exponent is the component that normally signifies the integer power to which the value two is raised in determining the value of the represented number. See also *Biased exponent*.

F

Feed-forwarding. A MPC603e feature that reduces the number of clock cycles that an execution unit must wait to use a register. When the source register of the current instruction is the same as the destination register of the previous instruction, the result of the previous instruction is routed to the current instruction at the same time that it is written to the register file. With feed-forwarding, the destination bus is gated to the waiting execution unit over the appropriate source bus, saving the cycles which would be used for the write and read.

Fetch. Retrieving instructions from either the cache or main memory and placing them into the instruction queue.

Floating-point register (FPR). Any of the 32 registers in the floating-point register file. These registers provide the source operands and destination results for floating-point instructions. Load instructions move data from memory to FPRs and store instructions move data from FPRs to memory. The FPRs are 64 bits wide and store floating-point values in double-precision format.

Floating-point unit. The functional unit in the MPC603e processor responsible for executing all floating-point instructions.

Flush. An operation that causes a cache block to be invalidated and the data, if modified, to be written to memory.

Fraction. In the binary representation of a floating-point number, the field of the *significand* that lies to the right of its implied binary point.

G

General-purpose register (GPR). Any of the 32 registers in the general-purpose register file. These registers provide the source operands and destination results for all integer data manipulation

instructions. Integer load instructions move data from memory to GPRs and store instructions move data from GPRs to memory.

Guarded. The guarded attribute pertains to out-of-order execution. When a page is designated as guarded, instructions and data cannot be accessed out-of-order.

H

Harvard architecture. An architectural model featuring separate caches and other memory management resources for instructions and data.

Hashing. An algorithm used in the *page table* search process.

I

IEEE 754. A standard written by the Institute of Electrical and Electronics Engineers that defines operations and representations of binary floating-point numbers.

Illegal instructions. A class of instructions that are not implemented for a particular PowerPC processor. These include instructions not defined by the PowerPC architecture. In addition, for 32-bit implementations, instructions that are defined only for 64-bit implementations are considered to be illegal instructions. For 64-bit implementations instructions that are defined only for 32-bit implementations are considered to be illegal instructions.

Implementation. A particular processor that conforms to the PowerPC architecture, but may differ from other architecture-compliant implementations for example in design, feature set, and implementation of *optional* features. The PowerPC architecture has many different implementations.

Imprecise exception. A type of *synchronous exception* that is allowed not to adhere to the precise exception model (see *Precise exception*). The PowerPC architecture allows only floating-point exceptions to be handled imprecisely.

Instruction queue. A holding place for instructions fetched from the current instruction stream.

Integer unit. The functional unit in the MPC603e responsible for executing all integer instructions.

In-order. An aspect of an operation that adheres to a sequential model. An operation is said to be performed in-order if, at the time that it is performed, it is known to be required by the sequential execution model. See *Out-of-order*.

Instruction latency. The total number of clock cycles necessary to execute an instruction and make ready the results of that instruction.

Interrupt. An external signal that causes the MPC603e to suspend current execution and take a predefined exception.

K

Key bits. A set of key bits referred to as K_s and K_p in each segment register and each BAT register. The key bits determine whether supervisor or user programs can access a *page* within that *segment* or *block*.

Kill. An operation that causes a *cache block* to be invalidated without writing any modified data to memory.

L

Latency. The number of clock cycles necessary to execute an instruction and make ready the results of that execution for a subsequent instruction.

L2 cache. See *Secondary cache*.

Least-significant bit (lsb). The bit of least value in an address, register, field, data element, or instruction encoding.

Least-significant byte (LSB). The byte of least value in an address, register, data element, or instruction encoding.

Little-endian. A byte-ordering method in memory where the address *n* of a word corresponds to the *least-significant byte*. In an addressed memory word, the bytes are ordered (left to right) 3, 2, 1, 0, with 3 being the *most-significant byte*. See *Big-endian*.

M

Mantissa. The decimal part of logarithm.

MEI (modified/exclusive/invalid). *Cache coherency* protocol used to manage caches on different devices that share a memory system. Note that the PowerPC architecture does not specify the implementation of a MEI protocol to ensure cache coherency.

Memory access ordering. The specific order in which the processor performs load and store memory accesses and the order in which those accesses complete.

Memory-mapped accesses. Accesses whose addresses use the page or block address translation mechanisms provided by the MMU and that occur externally with the bus protocol defined for memory.

Memory coherency. An aspect of caching in which it is ensured that an accurate view of memory is provided to all devices that share system memory.

Memory consistency. Refers to agreement of levels of memory with respect to a single processor and system memory (for example, on-chip cache, secondary cache, and system memory).

Memory management unit (MMU). The functional unit that is capable of translating an *effective* (logical) *address* to a physical address, providing protection mechanisms, and defining caching methods.

Modified state. MEI state (M) in which one, and only one, caching device has the valid data for that address. The data at this address in external memory is not valid.

Most-significant bit (msb). The highest-order bit in an address, registers, data element, or instruction encoding.

Most-significant byte (MSB). The highest-order byte in an address, registers, data element, or instruction encoding.

N

NaN. An abbreviation for not a number; a symbolic entity encoded in floating-point format. There are two types of NaNs—signaling NaNs and quiet NaNs.

No-op. No-operation. A single-cycle operation that does not affect registers or generate bus activity.

Normalization. A process by which a floating-point value is manipulated such that it can be represented in the format for the appropriate precision (single- or double-precision). For a floating-point value to be representable in the single- or double-precision format, the leading implied bit must be a 1.

O

OEA (operating environment architecture). The level of the architecture that describes PowerPC memory management model, supervisor-level registers, synchronization requirements, and the exception model. It also defines the time-base feature from a supervisor-level perspective. Implementations that conform to the PowerPC OEA also conform to the PowerPC UISA and VEA.

Optional. A feature, such as an instruction, a register, or an exception, that is defined by the PowerPC architecture but not required to be implemented.

Out-of-order. An aspect of an operation that allows it to be performed ahead of one that may have preceded it in the sequential model, for example, speculative operations. An operation is said to be performed out-of-order if, at the time that it is performed, it is not known to be required by the sequential execution model. See *In-order*.

Out-of-order execution. A technique that allows instructions to be issued and completed in an order that differs from their sequence in the instruction stream.

Overflow. A condition that occurs during arithmetic operations when the result cannot be stored accurately in the destination register(s). For example, if two 32-bit numbers are multiplied, the result may not be representable in 32 bits. Since the 32-bit registers of the MPC603e cannot represent this sum, an overflow condition occurs.

P

Page. A region in memory. The OEA defines a page as a 4-Kbyte area of memory, aligned on a 4-Kbyte boundary.

Page access history bits. The *changed* and *referenced* bits in the PTE keep track of the access history within the page. The referenced bit is set by the MMU whenever the page is accessed for a read or write operation. The changed bit is set when the page is stored into. See *Changed bit* and *Referenced bit*.

Page fault. A page fault is a condition that occurs when the processor attempts to access a memory location that does not reside within a *page* not currently resident in *physical memory*. On PowerPC processors, a page fault exception condition occurs when a matching, valid *page table entry* (PTE[V] = 1) cannot be located.

Page table. A table in memory is comprised of *page table entries*, or PTEs. It is further organized into eight PTEs per PTEG (page table entry group). The number of PTEGs in the page table depends on the size of the page table (as specified in the SDR1 register).

Page table entry (PTE). Data structures containing information used to translate *effective address* to physical address on a 4-Kbyte page basis. A PTE consists of 8 bytes of information in a 32-bit processor and 16 bytes of information in a 64-bit processor.

Park. The act of allowing a bus master to maintain bus mastership without having to arbitrate.

Physical memory. The actual memory that can be accessed through the system's memory bus.

Pipelining. A technique that breaks operations, such as instruction processing or bus transactions, into smaller distinct stages or tenures (respectively) so that a subsequent operation can begin before the previous one has completed.

Precise exceptions. A category of exception for which the pipeline can be stopped so instructions that preceded the faulting instruction can complete and subsequent instructions can be flushed and redispached after exception handling has completed. See *Imprecise exceptions*.

Primary opcode. The most-significant 6 bits (bits 0–5) of the instruction encoding that identifies the type of instruction.

Program order. The order of instructions in an executing program. More specifically, this term is used to refer to the original order in which program instructions are fetched into the instruction queue from the cache.

Protection boundary. A boundary between *protection domains*.

Protection domain. A protection domain is a segment, a virtual page, a BAT area, or a range of unmapped effective addresses. It is defined only when the appropriate relocate bit in the MSR (IR or DR) is 1.

Q

Quiesce. To come to rest. The processor is said to quiesce when an exception is taken or a **sync** instruction is executed. The instruction stream is stopped at the decode stage and executing instructions are allowed to complete to create a controlled context for instructions that may be affected by out-of-order, parallel execution. See *Context synchronization*.

Quiet NaN. A type of *Nan* that can propagate through most arithmetic operations without signaling exceptions. A quiet NaN is used to represent the results of certain invalid operations, such as invalid arithmetic operations on infinities or on NaNs, when invalid. See *Signaling NaN*.

R

rA. The **rA** instruction field is used to specify a GPR to be used as a source or destination.

rB. The **rB** instruction field is used to specify a GPR to be used as a source.

rD. The **rD** instruction field is used to specify a GPR to be used as a destination.

rS. The **rS** instruction field is used to specify a GPR to be used as a source.

Real address mode. An MMU mode when no address translation is performed and the *effective address* specified is the same as the physical address. The processor's MMU is operating in real address mode if its ability to perform address translation has been disabled through the MSR registers IR and/or DR bits.

Record bit. Bit 31 (or the **Rc** bit) in the instruction encoding. When it is set, updates the condition register (CR) to reflect the result of the operation.

Referenced bit. One of two *page history bits* found in each *page table entry*. The processor sets the *referenced bit* whenever the page is accessed for a read or write. See also *Page access history bits*.

Register indirect addressing. A form of addressing that specifies one GPR that contains the address for the load or store.

Register indirect with immediate index addressing. A form of addressing that specifies an immediate value to be added to the contents of a specified GPR to form the target address for the load or store.

Register indirect with index addressing. A form of addressing that specifies that the contents of two GPRs be added together to yield the target address for the load or store.

Rename register. Temporary buffers used by instructions that have finished execution but have not completed.

Reservation. The processor establishes a reservation on a *cache block* of memory space when it executes an **Iwarx** instruction to read a memory semaphore into a GPR.

Reservation station. A buffer between the dispatch and execute stages that allows instructions to be dispatched even though the results of instructions on which the dispatched instruction may depend are not available.

RISC (reduced instruction set computing). An *architecture* characterized by fixed-length instructions with nonoverlapping functionality and by a separate set of load and store instructions that perform memory accesses.

S

Scan interface. The MPC603e test interface.

Secondary cache. A cache memory that is typically larger and has a longer access time than the primary cache. A secondary cache may be shared by multiple devices. Also referred to as L2, or level-2, cache.

Set (v). To write a nonzero value to a bit or bit field; the opposite of *clear*. The term ‘set’ may also be used to generally describe the updating of a bit or bit field.

Set (n). A subdivision of a *cache*. Cacheable data can be stored in a given location in one of the sets, typically corresponding to its lower-order address bits. Because several memory locations can map to the same location, cached data is typically placed in the set whose *cache block* corresponding to that address was used least recently. See *Set-associative*.

Set-associative. Aspect of cache organization in which the cache space is divided into sections, called *sets*. The cache controller associates a particular main memory address with the contents of a particular set, or region, within the cache.

Shadowing. Shadowing allows a register to be updated by instructions that are executed out of order without destroying machine state information.

Signaling NaN. A type of *NaN* that generates an invalid operation program exception when it is specified as arithmetic operands. See *Quiet NaN*.

Significand. The component of a binary floating-point number that consists of an explicit or implicit leading bit to the left of its implied binary point and a fraction field to the right.

Simplified mnemonics. Assembler mnemonics that represent a more complex form of a common operation.

Slave. The device addressed by a master device. The slave is identified in the address tenure and is responsible for supplying or latching the requested data for the master during the data tenure.

Snooping. Monitoring addresses driven by a bus master to detect the need for coherency actions.

Snoop push. Response to a snooped transaction that hits a modified cache block. The cache block is written to memory and made available to the snooping device.

Split-transaction. A transaction with independent request and response tenures.

Split-transaction bus. A bus that allows address and data transactions from different processors to occur independently.

Stage. The term *stage* is used in two different senses, depending on whether the pipeline is being discussed as a physical entity or a sequence of events. In the latter case, a stage is an element in the pipeline during which certain actions are performed, such as decoding the instruction, performing an arithmetic operation, or writing back the results. Typically, the latency of a stage is one processor clock cycle. Some events, such as dispatch, write-back, and completion, happen instantaneously and may be thought to occur at the end of a stage. An instruction can spend multiple cycles in one stage. An integer multiply, for example, takes multiple cycles in the execute stage. When this occurs, subsequent instructions may stall. An instruction may also occupy more than one stage simultaneously, especially in the sense that a stage can be seen as a physical resource—for example, when instructions are dispatched they are assigned a place in the CQ at the same time they are passed to the execute stage. They can be said to occupy both the complete and execute stages in the same clock cycle.

Stall. An occurrence when an instruction cannot proceed to the next stage.

Static branch prediction. Mechanism by which software (for example, compilers) can hint to the machine hardware about the direction a branch is likely to take.

Superscalar machine. A machine that can issue multiple instructions concurrently from a conventional linear instruction stream.

Supervisor mode. The privileged operation state of a processor. In supervisor mode, software, typically the operating system, can access all control registers and can access the supervisor memory space, among other privileged operations.

Synchronization. A process to ensure that operations occur strictly *in order*. See *Context synchronization* and *Execution synchronization*.

Synchronous exception. An *exception* that is generated by the execution of a particular instruction or instruction sequence. There are two types of synchronous exceptions, *precise* and *imprecise*.

System memory. The physical memory available to a processor.

T

Tenure. The period of bus mastership. For the MPC603e, there can be separate address bus tenures and data bus tenures. A tenure consists of three phases: arbitration, transfer, and termination.

TLB (translation lookaside buffer). A cache that holds recently-used *page table entries*.

Throughput. The measure of the number of instructions that are processed per clock cycle.

Transaction. A complete exchange between two bus devices. A transaction is typically comprised of an address tenure and one or more data tenures, which may overlap or occur separately from the address tenure. A transaction may be minimally comprised of an address tenure only.

Transfer termination. Signal that refers to both signals that acknowledge the transfer of individual beats (of both single-beat transfer and individual beats of a burst transfer) and to signals that mark the end of the tenure.

U

UISA (user instruction set architecture). The level of the architecture to which user-level software should conform. The UISA defines the base user-level instruction set, user-level registers, data types, floating-point memory conventions and exception model as seen by user programs, and the memory and programming models.

Underflow. A condition that occurs during arithmetic operations when the result cannot be represented accurately in the destination register. For example, underflow can happen if two floating-point fractions are multiplied and the result requires a smaller *exponent* and/or *mantissa* than the single-precision format can provide. In other words, the result is too small to be represented accurately.

User mode. The operating state of a processor used typically by application software. In user mode, software can access only certain control registers and can access only user memory space. No privileged operations can be performed. Also referred to as problem state.

V

VEA (virtual environment architecture). The level of the *architecture* that describes the memory model for an environment in which multiple devices can access memory, defines aspects of the cache model, defines cache control instructions, and defines the time-base facility from a user-level perspective. *Implementations* that conform to the PowerPC VEA also adhere to the UISA, but may not necessarily adhere to the OEA.

Virtual address. An intermediate address used in the translation of an *effective address* to a physical address.

Virtual memory. The address space created using the memory management facilities of the processor. Program access to *virtual memory* is possible only when it coincides with *physical memory*.

W

Way. A location in the cache that holds a cache block, its tags and status bits.

Word. A 32-bit data element.

Write-back. A cache memory update policy in which processor write cycles are directly written only to the cache. External memory is updated only indirectly, for example, when a modified cache block is *cast out* to make room for newer data.

Write-through. A cache memory update policy in which all processor write cycles are written to both the cache and memory.



Freescale Semiconductor, Inc.

Freescale Semiconductor, Inc.

MPC603e RISC Microprocessor User's Manual

**For More Information On This Product,
Go to: www.freescale.com**

Index

A

AACK signal, 7-15
ABB signal, 7-4, 8-7
ABE (address broadcast enable) bit, 3-23
Address bus
 address tenure, 8-6
 address transfer
 An, 7-6
 AP_E, 7-8, 8-12
 AP_n, 7-7
 address transfer attribute
 CI, 7-13
 CSE_n, 7-14
 GBL, 7-14
 TBST, 7-12, 8-13
 TC_n, 7-13, 8-19
 TSIZ_n, 7-12, 8-13
 TT_n, 7-9, 8-12
 WT, 7-14
 address transfer start
 TS, 7-6, 8-11
 XATS (MPC603-specific), 1-6
 address transfer termination
 AACK, 7-15
 ARTRY, 3-21, 7-15
 terminating address transfer, 8-19
 arbitration signals, 7-2, 8-7
 bus arbitration
 ABB, 7-4, 8-7
 BG, 7-4, 8-7
 BR, 7-3, 8-7
 bus parking, 8-10
 Address calculation
 branch instructions, 2-36
 effective address, 2-20
 floating-point load and store, 2-34
 integer load and store, 2-30
 Address translation, *see* Memory management unit
 Addressing conventions
 addressing modes, 2-19
 alignment, 2-14
 Aligned data transfer, 2-14, 8-14, 8-18

Alignment

data transfers, 2-14, 8-14
 exception, 4-24, 5-15
 rules, 2-14
An signals, 7-6
AP_E signal, 7-8, 8-12
AP_n signals, 7-7
 Arbitration, system bus, 8-9, 8-21
ARTRY signal, 3-21, 7-15
 Atomic memory references
 stwcx., 2-39
 using **Iwarx/stwex.**, 3-20

B

BG signal, 7-4, 8-7
 Block address translation
 block address translation flow, 5-11
 selection of block address translation, 5-9
 Boundedly undefined, definition, 2-17
BR signal, 7-3, 8-7
 Branch folding, 6-17
 Branch instructions
 address calculation, 2-36
 branch instructions, 2-36, A-22
 condition register logical, 2-37, A-22
 system linkage, 2-43, A-22
 trap, 2-38, A-23
 Branch prediction, 6-1, 6-18
 Branch processing unit
 branch instruction timing, 6-20
 execution timing, 6-17
 latency, branch instructions, 6-26
 overview, 1-9
 Branch resolution
 definition, 6-1
 Burst data transfers
 32-bit data bus, 8-14
 64-bit data bus, 8-13
 transfers with data delays, timing, 8-35
 Burst transactions, 3-9
 Bus arbitration, *see* Data bus
 Bus configurations, 8-37, 8-39

Bus interface unit (BIU), 3-2
 Byte ordering
 default, 2-19
 Byte-reverse instructions, 2-32, A-20

C

Cache
 cache miss, 6-13
 characteristics, 3-1
 instructions, 2-41, 2-45, 3-22, A-23
 MEI state definition, 3-16
 organization, instruction/data, 3-3-3-8
 overview, 1-24
 Cache arbitration, 6-10
 Cache block push operation, 3-9
 Cache block, definition, 3-1
 Cache cast-out operation, 3-8
 Cache coherency
 actions on load operations, 3-19
 actions on store operations, 3-19
 copy-back operation, 3-12
 in single-processor systems, 3-19
 MEI protocol, 3-15
 out-of-order execution, 3-14
 overview, 3-3
 reaction to bus operations, 3-20
 WIMG bits, 3-10, 3-14, 8-28
 write-back mode, 3-12
 Cache hit, 6-10
 Cache management instructions, 2-41, 2-45, 3-22, A-23
 Cache operations
 basic data cache operations, 3-8
 data cache transactions, 3-9
 instruction cache fill operations, 3-4
 overview, 1-13, 3-1
 response to bus transactions, 3-20
 Cache unit
 memory performance, 6-22
 operation of the cache, 8-2
 overview, 3-1
 Cache-inhibited accesses (I bit)
 cache interactions, 3-10
 I-bit setting, 3-12
 timing considerations, 6-23
 Changed (C) bit maintenance
 recording, 5-11, 5-21-5-23
 Checkstop
 signal, 7-24, 8-40
 state, 4-21
 \overline{CI} signal, 7-13
 Classes of instructions, 2-17
 Clean block operation, 3-20

Clock signals
 CLK_OUT, 7-30
 PLL_CFGn, 7-30
 SYSCLK, 7-29
 Compare instructions, 2-28, A-16
 Completion
 definition, 6-1
 Completion considerations, 6-15
 Context synchronization, 2-21
 Conventions, xxxi, xxxv, 2-13
 COP/scan interface, 7-27
 Copy-back mode, 6-23
 CR logical instructions, 2-37
 CSEn signals, 7-14, 8-29

D

Data bus
 32-bit data bus mode, 8-37
 arbitration signals, 7-16, 8-7
 bus arbitration, 8-21
 data tenure, 8-7
 data transfer, 7-18, 8-23
 data transfer termination, 7-21, 8-24
 Data cache
 basic operations, 3-8
 broadcasting, 3-7
 bus transactions, 3-9
 cache control, 3-6
 configuration, 3-1
 DCFI, DCE, DLOCK bits, 3-6
 disabling, 3-7
 fill operations, 3-8
 locking, 3-7
 organization, 3-6
 touch load operations, 3-8
 touch load support, 3-8
 Data storage interrupt (DSI), *see* DSI exception
 Data TLB miss on load exception, 4-31
 Data TLB miss on store exception, 4-32
 Data transfers
 alignment, 2-14, 8-14
 burst ordering, 8-13
 eciwx and ecowx instructions, alignment, 8-18
 signals, 8-23
 DBB signal, 7-18, 8-8, 8-22
 $\overline{\text{DBDIS}}$ signal, 7-21
 $\overline{\text{DBG}}$ signal, 7-17, 8-7
 DBWO signal, 7-17, 8-8, 8-22, 8-43
 DCMP and ICMP registers, 5-34
 Decrementer interrupt, 4-28, 9-2
 Defined instruction class, 2-17
 DHn/DLn signals, 7-18

Direct address translation (translation disabled)
 data accesses, 3-11, 5-9, 5-11, 5-19
 instruction accesses, 3-11, 5-9, 5-11, 5-19
 Direct-store access on the MPC603e, 3-10
 Dispatch considerations, 6-15
 DMISS and IMISS registers, 5-33
DPE signal, 7-20
DPn signals, 7-19
DRTRY signal, 7-22, 8-24, 8-27
 DS1 exception, 4-21

E

Effective address calculation
 address translation, 5-3
 branches, 2-20, 2-36
 loads and stores, 2-20, 2-30, 2-34
 Error termination, 8-27
 Exceptions
 alignment exception, 4-24
 data TLB miss on load, 4-31
 data TLB miss on store, 4-32
 decrementer interrupt, 4-28
 DSI exception, 4-21
 enabling and disabling, 4-13
 exception classifications, 4-2
 exception processing, 4-9, 4-13
 external interrupt, 4-23
 FP unavailable exception, 4-28
 instruction address breakpoint, 4-32
 instruction related, 2-21
 instruction TLB miss, 4-30
 machine check exception, 4-19
 overview, 1-26
 program exception, 4-27
 register settings
 FPSCR, 4-27
 MSR, 4-16
 SRR0/SRR1, 4-10
 reset, 4-17
 returning from an exception handler, 4-14
 summary, 2-21
 system call, 4-29
 system management interrupt, 4-34
 trace exception, 4-29
 Execution synchronization, 2-21
 Execution units, 1-9
 External control instructions, 2-43, 8-18, A-24

F

Features list, 1-3
 Finish cycle, definition, 6-2
 Floating-point model
 FE0/FE1 bits, 4-13

FP arithmetic instructions, 2-26, A-17
 FP compare instructions, 2-28, A-18
 FP execution models, 2-13
 FP load instructions, 2-34, A-21
 FP move instructions, 2-29, A-22
 FP multiply-add instructions, 2-27, A-18
 FP rounding/conversion instructions, 2-28, A-18
 FP store instructions, 2-35, A-21
 FP unavailable exception, 4-28
 FPSCR instructions, 2-28, A-18

Floating-point unit

 execution timing, 6-21
 latency, FP instructions, 6-30
 overview, 1-10

Flow control instructions

 branch instruction address calculation, 2-36
 branch instructions, 2-37
 condition register logical, 2-37

Flush-block operation, 3-20

FPR0–FPR31, 2-2
 FPSCR instructions, 2-28, A-18

G

GBL signal, 7-14
GPR0–GPR31, 2-2
 Guarded memory bit (G bit)
 cache interactions, 3-10
 G-bit setting, 3-13

H

HASH1 and **HASH2** registers, 5-34
 Hashing functions
 primary PTEG, 5-30
 secondary PTEG, 5-31
HID0 (hardware implementation-dependent 0)
 registers
 nap bit, 9-4
HID0 register
 DCFI, DCE, DLOCK bits, 3-6
 doze bit, 9-4
 DPM enable bit, 9-3
 ICFI, ICE, ILOCK bits, 3-4
 PID7t-specific bits, 1-19, 3-23
HID1 register
 bit settings, 2-10
 PLL configuration, 2-10, 7-30
HRESET signal, 7-25

I

IEEE 1149.1-compliant interface, 8-42
 IFEM (instruction fetch enable) bit, 1-19
 Illegal instruction class, 2-18

- ILOCK control bit, 3-5
- Instruction address breakpoint exception, 4-32
- Instruction cache
 - cache control bits, 3-4
 - cache fill operations, 3-4
 - configuration, 3-2
 - organization, 3-4
- Instruction timing
 - examples
 - cache hit, 6-11, 6-14
 - execution unit, 6-17
 - instruction flow, 6-8
 - memory performance considerations, 6-22
 - overview, 1-32, 6-3
 - terminology, 6-1
- Instruction TLB miss exception, 4-30
- Instruction unit, 1-8
- Instructions
 - branch address calculation, 2-36
 - branch instructions, 2-36, A-22
 - cache management instructions, 2-41, 2-45, 3-22, A-23
 - classes, 2-17
 - condition register logical, 2-37, A-22
 - defined instructions, 2-17
 - external control, 2-43, A-24
 - floating-point
 - arithmetic, 2-26, A-17
 - compare, 2-28, A-18
 - FP load instructions, 2-34, A-21
 - FP move instructions, 2-29, A-22
 - FP status and control register, 2-28
 - FP store instructions, 2-35, A-21
 - FPSR instructions, 2-28, A-18
 - multiply-add, 2-27, A-18
 - rounding and conversion, 2-28, A-18
 - illegal instructions, 2-18
 - integer
 - arithmetic, 2-23, A-15
 - compare, 2-24, A-16
 - load, A-19
 - logical, 2-24, A-16
 - multiple, 2-32, A-20
 - rotate and shift, 2-25, A-16
 - store, 2-31, A-19
 - latency summary, 6-26
 - load and store
 - address generation, floating-point, 2-34
 - address generation, integer, 2-30
 - byte-reverse instructions, 2-32, A-20
 - integer load, 2-30
 - integer multiple instructions, 2-32, A-20
 - integer store, 2-31
 - string instructions, 2-33, A-20
 - memory control, 2-41, 2-45, 3-22, A-23
 - memory synchronization, 2-38, 2-41, A-21
 - MPC603e, instructions not implemented, B-1
 - MPC603e-specific instructions, 2-47
 - PowerPC instructions, list
 - form (format), A-25
 - function, A-15
 - legend, A-36
 - mnemonic, A-1
 - opcode, A-8
 - processor control, 2-38, 2-40, 2-44, A-23
 - reserved instructions, 2-19
 - segment register manipulation, 2-46, A-24
 - simplified mnemonics, 2-47
 - supervisor-level cache management, 2-45
 - support for **Iwarx/stwcx.**, 8-41
 - system linkage, 2-43, A-22
 - TLB management instructions, 2-46, A-24
 - trap instructions, 2-38, A-23
 - INT signal, 7-23, 8-40
 - Integer arithmetic instructions, 2-23, A-15
 - Integer compare instructions, 2-24, A-16
 - Integer load instructions, 2-30, A-19
 - Integer logical instructions, 2-24, A-16
 - Integer multiple instructions, 2-32, A-20
 - Integer rotate and shift instructions, 2-25, A-16
 - Integer store instructions, 2-31, A-19
 - Integer unit
 - execution timing, 6-21
 - latency, integer instructions, 6-28
 - overview, 1-9
 - Interrupt, external, 4-23
 - Interrupt, *see* Exceptions

K

Kill-block operation, 3-20

L

- Latency, 6-2, 6-26, 8-23
- Load operations
 - memory coherency actions, 3-19
- Load/store
 - address generation, 2-30, 2-34
 - byte-reverse instructions, 2-32, A-20
 - floating-point load instructions, 2-34, A-21
 - floating-point move instructions, 2-29, A-22
 - floating-point store instructions, 2-35, A-21
 - integer load instructions, 2-30, A-19
 - integer store instructions, 2-31, A-19
 - load/store multiple instructions, 2-32, A-20
 - memory synchronization instructions, 2-38, 2-41, A-21
 - string instructions, 2-33, A-20

- Load/store unit
 - execution timing, 6-22
 - latency, load and store instructions, 6-31
 - Logical addresses
 - translation into physical addresses, 5-1
 - lwarx/stwcx.**
 - atomic memory references, 3-20
 - support, 8-41
- M**
- Machine check exception
 - checkstop state, 4-21
 - register settings, 4-20
 - SRR1 bit settings, 4-10
 - machine check exception enabled, 4-20
 - MCP** signal, 7-24
 - MEI protocol
 - definition, MEI states, 3-16
 - enforcing memory coherency, 8-28
 - hardware considerations, 3-17
 - Memory accesses, 8-4
 - Memory coherency bit (M bit)
 - cache interactions, 3-10
 - I-bit setting, 3-12
 - M-bit setting, 3-12
 - timing considerations, 6-23
 - Memory control instructions
 - segment register manipulation, 2-46
 - TLB management, 2-46
 - user-level cache, 2-41, 2-45, 3-22
 - Memory management unit
 - address translation flow, 5-11
 - address translation mechanisms, 5-7, 5-11
 - block address translation, 5-9, 5-11, 5-20
 - block diagram, 5-5-5-7
 - direct address translation, 3-11, 5-9, 5-11, 5-19
 - exceptions, 5-14
 - features summary, 5-2
 - instructions and registers, 5-17
 - memory protection, 5-10
 - overview, 1-11, 1-30
 - page address translation, 5-8, 5-11, 5-27
 - page history status, 5-11, 5-21-5-24
 - page table search operation, 5-27
 - segment model, 5-20
 - software table search operation, 5-31, 5-36, 5-37
 - Memory synchronization
 - instructions, 2-38, 2-41, A-21
 - stwcx.**, 2-39
 - Memory/cache access modes
 - performance impact of copy-back mode, 6-23
 - see also* WIMG bits
 - Misaligned accesses, 2-14
 - Misaligned data transfer, 8-16
 - Move instructions, 2-29
 - MPC603e
 - features
 - hardware, 1-6
 - list of features, 1-3
 - PID7t-specific, 1-5
 - instructions, 2-47
 - overview, 1-1, 1-16
 - PID7t
 - features, 1-5
 - HID0 bits (PID7t-specific), 3-23
 - processor identification (PID) number definition, 1-1
 - MPC603-specific features, 1-6
 - MSR (machine state register)
 - bit settings, 4-11
 - DR/IR bit, 4-12
 - EE bit, 4-11
 - FE0/FE1 bits, 4-13
 - POW bit, 2-5, 4-11, C-4
 - RI bit, 4-14
 - settings due to exception, 4-16
 - TGPR bit, 2-5, 4-11, C-4
- N**
- No-**DRTRY** mode, 8-39
 - Nondenormalized mode, support, 2-26
- O**
- Operand conventions, 2-13
 - Operand placement and performance, 2-15
 - Operating environment architecture (OEA), xxvii, 1-15, 2-43
 - Optional instructions, A-36
- P**
- Page address translation
 - page address translation flow, 5-27
 - page size, 5-20
 - selection of page address translation, 5-8, 5-13
 - table search operation, 5-27
 - TLB organization, 5-25
 - Page history status
 - R and C bit recording, 5-11, 5-21-5-24
 - Page tables
 - resources for table search operations, 5-31
 - software table search operation, 5-31, 5-36
 - table search for PTE, 5-27
 - Performance considerations, memory, 6-22
 - Physical address generation
 - memory management unit, 5-1

PID7t-603e features, 1-5
 Pipeline
 instruction timing, definition, 6-2
 pipeline stages, 6-7
 superscalar/pipeline diagram, 6-5
 Pipelined execution unit, 6-4
 PLL configuration, 7-31
 Power management
 doze mode, 9-3
 doze, nap, sleep, DPM bits, 2-10
 full-power mode, 9-3
 nap mode, 9-4
 programmable power modes, 9-3
 sleep mode, 9-5
 software considerations, 9-6
 Power management modes, 1-14
 Power-on reset settings, 4-17
 PowerPC architecture
 instruction list, A-1, A-8, A-15
 levels of implementation, 1-15
 operating environment architecture (OEA), xxvii,
 1-15, 2-43
 user instruction set architecture (UISA), xxvii,
 1-15, 2-2
 virtual environment architecture (VEA), xxvii,
 1-15, 2-40
 Privilege levels
 supervisor-level cache instruction, 2-45
 Privileged state, *see* Supervisor mode
 Problem state, *see* User mode
 Process switching, 4-15
 Processor control instructions, 2-38, 2-40, 2-44, A-23
 Processor identification (PID) number definition, 1-1
 Program exception, 4-27
 Program order, definition, 6-2
 Programmable power states
 doze mode, 9-3
 full-power mode (DPM enabled/disabled), 9-3
 nap mode, 9-4
 sleep mode, 9-5
 Protection of memory areas
 no-execute protection, 5-12
 options available, 5-10
 protection violations, 5-14
 PTEGs (PTE groups), 5-27
 PTEs (page table entries), 5-27

Q

QACK signal, 7-26, 8-37, 8-40
QREQ signal, 7-26, 8-41
 Qualified bus grant, 8-7
 Qualified data bus grant, 8-21

R

Read operation, 3-20
 Read with intent to modify operation, 3-20
 Read-atomic operation, 3-20
 Real address (RA), *see* Physical address generation
 Real addressing mode, *see* Direct address translation
 Reduced-pinout mode, 8-40
 Referenced (R) bit maintenance
 recording, 5-11, 5-21-5-23, 5-29
 Registers
 configuration registers
 MSR, 2-4
 PVR, 2-4
 exception handling registers
 DAR, 2-5
 DSISR, 2-5
 SPRG0-SPRG3, 2-5
 SRR0, 2-5
 SRR1, 2-5
 implementation-specific registers
 DCMP/ICMP, 2-10
 DMISS/IMISS, 2-10
 HASH1/HASH2, 2-11
 HID0/HID1, 1-19, 2-6
 IABR, 2-12
 RPA, 2-12
 memory management registers
 BAT, 2-5
 SDR1, 2-5
 SR, 2-5
 supervisor-level
 BAT, 2-5
 DAR, 2-5
 DCMP and ICMP, 2-10, 5-34
 DEC, 2-6
 DMISS and IMISS, 2-10, 5-33
 DSISR, 2-5
 EAR, 2-6
 HASH1 and HASH2, 2-11, 5-34
 HID0 and HID1, 1-19, 2-6
 IABR, 2-12
 MSR, 2-4
 PVR, 2-4
 RPA, 2-12
 SDR1, 2-5
 SPRG0-SPRG3, 2-5
 SR, 2-5
 SRR0, 2-5
 SRR1, 2-5
 TB, 2-6
 user-level
 CR, 2-2
 CTR, 2-4

- FPR0–FPR31, 2-2
 - FPSCR, 2-2
 - GPR0–GPR31, 2-2
 - LR, 2-4
 - TB, 2-4
 - TGPR0–TGPR3, 5-32
 - XER, 2-4
 - Rename buffer, 6-2
 - Rename register operation, 6-15
 - Reservation station, 6-2
 - Reserved instruction class, 2-19
 - Reset
 - $\overline{\text{HRESET}}$ signal, 7-25
 - $\overline{\text{HRESET}}$ signal, 8-41
 - reset exception, 4-17
 - settings caused by hard reset, 4-17
 - $\overline{\text{SRESET}}$ signal, 7-25, 8-41
 - Retirement, definition, 6-2
 - Rotate and shift instructions, 2-25, A-16
 - RPA (required physical address), 5-35
 - $\overline{\text{RSRV}}$ signal, 7-27, 8-41, C-15
- S**
- Segment registers
 - SR manipulation instructions, 2-46, A-24
 - T bit, Glossary-4
 - Segmented memory model, *see* Memory management unit
 - Self-modifying code, 2-30
 - Serializing instructions, 6-16
 - Signals
 - $\overline{\text{AACK}}$, 7-15
 - $\overline{\text{ABB}}$, 7-4, 8-7
 - address arbitration, 7-2, 8-7
 - address transfer, 8-11
 - address transfer attribute, 8-12
 - An , 7-6
 - $\overline{\text{APE}}$, 7-8
 - APn , 7-7
 - $\overline{\text{ARTRY}}$, 7-15, 8-24
 - $\overline{\text{BG}}$, 7-4, 8-7
 - $\overline{\text{BR}}$, 7-3, 8-7
 - checkstop, 8-40
 - $\overline{\text{CI}}$, 7-13
 - $\overline{\text{CKSTP_IN}}$, 7-24
 - $\overline{\text{CKSTP_OUT}}$, 7-24
 - CLK_OUT , 7-30
 - configuration, 7-2
 - COP/scan interface, 7-27
 - CSEn , 7-14, 8-29
 - data arbitration, 8-7, 8-21
 - data transfer termination, 8-24
 - $\overline{\text{DBB}}$, 7-18, 8-8, 8-22
 - $\overline{\text{DBDIS}}$, 7-21
 - $\overline{\text{DBG}}$, 7-17, 8-7
 - $\overline{\text{DBWO}}$, 7-17, 8-8, 8-22, 8-43
 - $\overline{\text{DHn/DLn}}$, 7-18
 - $\overline{\text{DPE}}$, 7-20
 - $\overline{\text{DPn}}$, 7-19
 - $\overline{\text{DRTRY}}$, 7-22, 8-24, 8-27
 - $\overline{\text{GBL}}$, 7-14
 - $\overline{\text{HRESET}}$, 7-25
 - $\overline{\text{INT}}$, 7-23, 8-40
 - $\overline{\text{MCP}}$, 7-24
 - $\overline{\text{PLL_CFGn}}$, 7-30
 - $\overline{\text{QACK}}$, 7-26, 8-37, 8-40
 - $\overline{\text{QREQ}}$, 7-26, 8-41
 - Reset, 8-41
 - $\overline{\text{RSRV}}$, 7-27, 8-41, C-15
 - $\overline{\text{SMI}}$, 4-34, 7-23
 - $\overline{\text{SRESET}}$, 7-25, 8-41
 - $\overline{\text{TA}}$, 7-21
 - $\overline{\text{TBEN}}$, 7-27
 - $\overline{\text{TBST}}$, 7-12, 8-23
 - $\overline{\text{TCn}}$, 7-13, 8-19
 - $\overline{\text{TEA}}$, 7-22, 8-24, 8-27
 - $\overline{\text{TLBISYNC}}$, 7-27
 - $\overline{\text{TS}}$, 7-6
 - $\overline{\text{TSIZn}}$, 7-12, 8-13
 - $\overline{\text{TTn}}$, 7-9, 8-12
 - $\overline{\text{WT}}$, 7-14
 - $\overline{\text{XATS}}$ (MPC603-specific), 1-6
 - Single-beat reads with data delays, timing, 8-34
 - Single-beat transactions, 3-9
 - Single-beat transfer
 - reads with data delays, timing, 8-33
 - reads, timing, 8-31
 - termination, 8-25
 - writes, timing, 8-32
 - $\overline{\text{SMI}}$ signal, 4-34, 7-23
 - Snoop operation, 3-20, 6-23
 - Split-bus transaction, 8-8
 - SPR encodings not implemented in MPC603e, B-2
 - $\overline{\text{SRESET}}$ signal, 7-25
 - SRR0/SRR1 (status save/restore registers)
 - bit settings for machine check exception, 4-10
 - bit settings for table search operations, 4-10
 - Stall, definition, 6-3
 - Static branch prediction, 6-18
 - Store operations
 - memory coherency actions, 3-19
 - single-beat writes, 8-32
 - String instructions, 2-33, A-20
 - Superscalar, 6-3
 - Supervisor mode, *see* Privilege levels
 - Supervisor-level registers summary, 2-4
 - sync operation, 3-20

Synchronization
 context/execution synchronization, 2-20
 execution of **rfi**, 4-14
 memory synchronization instructions, 2-38, 2-41, A-21
SYCLK signal, 7-29
 System call exception, 4-29
 System interface
 overview, 1-33
 System linkage instructions, 2-43, A-22
 System management interrupt, 4-34, 9-2
 System quiesce control signals, 8-41
 System register unit
 execution timing, 6-22
 latency, CR logical instructions, 6-27
 latency, system register instructions, 6-27
 System status
CKSTP_IN, 7-24
CKSTP_OUT, 7-24
HRESET, 7-25
INT, 7-23
MCP, 7-24
QACK, 7-26
QREQ, 7-26
RSRV, 7-27
SMI, 7-23
SRESET, 7-25
TBEN, 7-27
TLBISYNC, 7-27

T

TA signal, 7-21
 Table search operations
 algorithm, 5-27
 software routines, 5-31
 software routines for the MPC603e, 5-36-5-46
 SRR1 bit settings, 4-10
 table search flow (primary and secondary), 5-29
TBEN signal, 7-27
TBST signal, 7-12, 8-13, 8-23
TC_n signals, 7-13, 8-19
TEA signal, 7-22, 8-27
 Termination, 8-19, 8-24
TGPR0–GPR3 registers, 5-32
 Throughput, 6-3
 Timing diagrams, interface
 address transfer signals, 8-11
 burst transfers with data delays, 8-35
 single-beat reads, 8-31
 single-beat reads with data delays, 8-33
 single-beat writes, 8-32
 single-beat writes with data delays, 8-34

use of **TEA**, 8-36
 using **DBWO**, 8-43
 Timing, instruction
 BPU execution timing, 6-17
 branch timing example, 6-20
 cache arbitration, 6-10
 cache hit, 6-10, 6-11, 6-14
 FPU execution timing, 6-21
 instruction dispatch, 6-15
 instruction flow, 6-8
 instruction scheduling guidelines, 6-24
 IU execution timing, 6-21
 latency summary, 6-26
 load/store unit execution timing, 6-22
 overview, 6-3
 SRU execution timing, 6-22
 stage, definition, 6-2
TLB
 description, 5-24
 invalidate, A-24
 invalidate (**tlbie** instruction), 5-26, 5-47
 TLB management instructions, 2-47, A-24
TLBISYNC signal, 7-27
 Trace exception, 4-29
 Transactions, data cache, 3-9
 Transfer, 8-11, 8-23
 Trap instructions, 2-38
TS signal, 7-6, 8-11
TSIZn signals, 7-12, 8-13
TTn signals, 7-9, 8-12

U

Use of **TEA**, timing, 8-36
 User mode, 4-1
 User instruction set architecture (UISA), xxvii, 1-15, 2-2
 User-level registers summary, 2-2
 Using **DBWO**, timing, 8-43

V

Virtual environment architecture (VEA), xxvii, 1-15, 2-40

W

WIMG bits, 3-10, 8-28
 Write-back, 6-3
 Write-back mode, 3-12
 Write-through mode (W bit)
 cache interactions, 3-10
 timing considerations, 6-23
 W-bit setting, 3-11
 Write-with-atomic operation, 3-20

Write-with-flush operation, 3-20

Write-with-kill operation, 3-20

WT signal, 7-14

X

$\overline{\text{XATS}}$ signal (MPC603-specific), 1-6



Freescale Semiconductor, Inc.

Freescale Semiconductor, Inc.

MPC603e RISC Microprocessor User's Manual

**For More Information On This Product,
Go to: www.freescale.com**

Overview	1
Programming Model	2
Instruction and Data Cache Operation	3
Exceptions	4
Memory Management	5
Instruction Timing	6
Signal Descriptions	7
System Interface Operation	8
Power Management	9
PowerPC Instruction Set Listings	A
Instructions Not Implemented	B
Revision History	C
Glossary of Terms and Abbreviations	GLO
Index	IND

- 1** Overview
- 2** Programming Model
- 3** Instruction and Data Cache Operation
- 4** Exceptions
- 5** Memory Management
- 6** Instruction Timing
- 7** Signal Descriptions
- 8** System Interface Operation
- 9** Power Management
- A** PowerPC Instruction Set Listings
- B** Instructions Not Implemented
- C** Revision History
- GLO** Glossary of Terms and Abbreviations
- IND** Index